

# **Dedup Est Machina**

+

# **Flip Feng Shui**

**Erik Bosman**



**WARNING**  
THIS PRESENTATION  
MAY CONTAIN POINTERS



# **Part I: Dedup est Machina**

# **Part I: Dedup est Machina**

## **Deduplication (software side-channel)**

# **Part I: Dedup est Machina**

**Deduplication  
(software side-channel)**

**+**

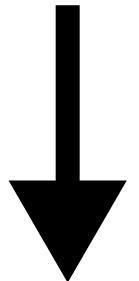
**Rowhammer  
(hardware bug)**

# **Part I: Dedup est Machina**

**Deduplication  
(software side-channel)**

**+**

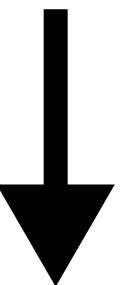
**Rowhammer  
(hardware bug)**



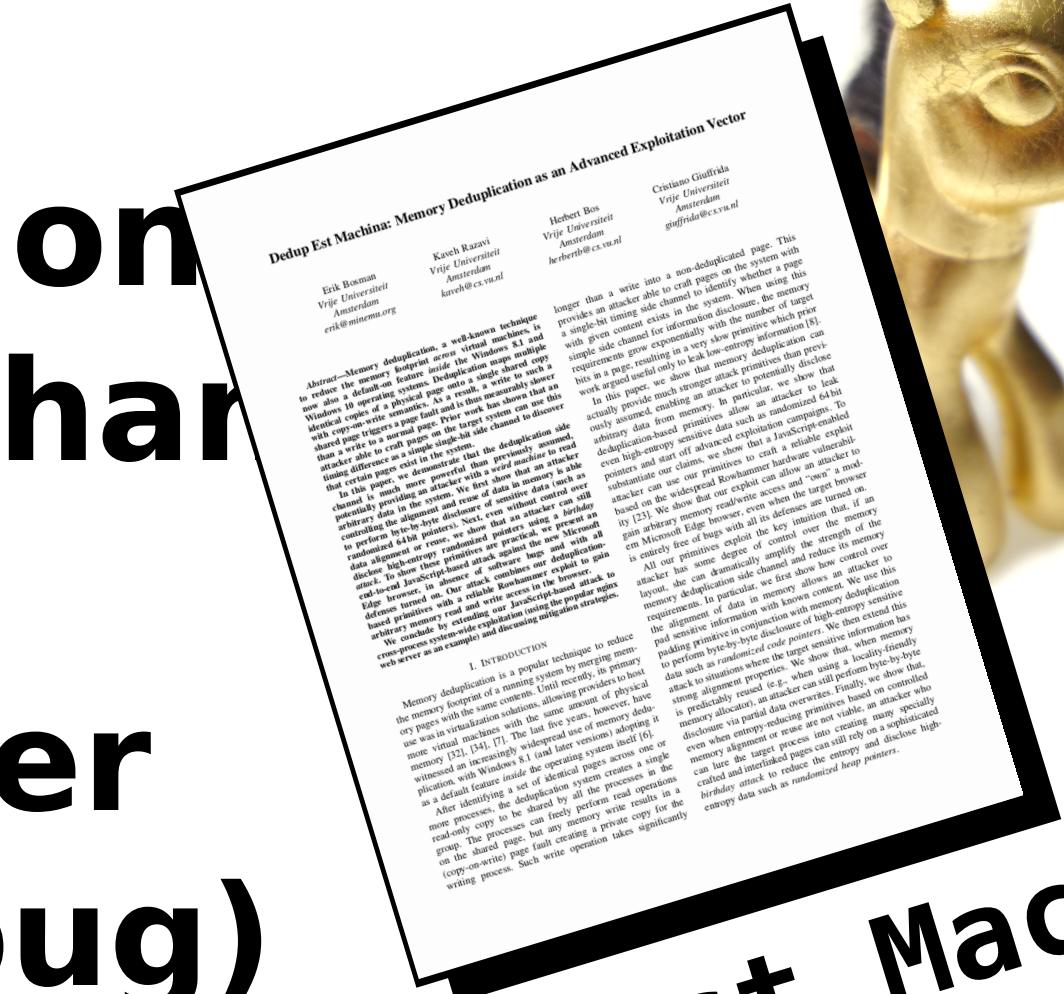
**Exploit MS Edge without software bugs  
(from JavaScript)**

# Part I: Dedup est Machina

Deduplication  
(software side-channel)  
+  
Rowhammer  
(hardware bug)



Exploit MS Edge without software bugs  
(from JavaScript)



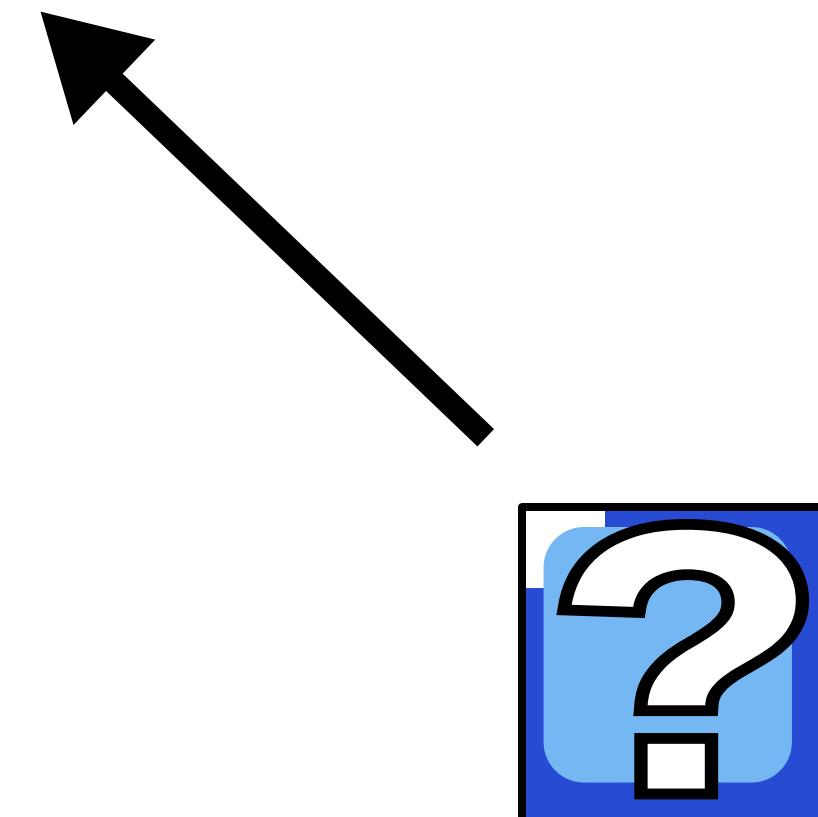
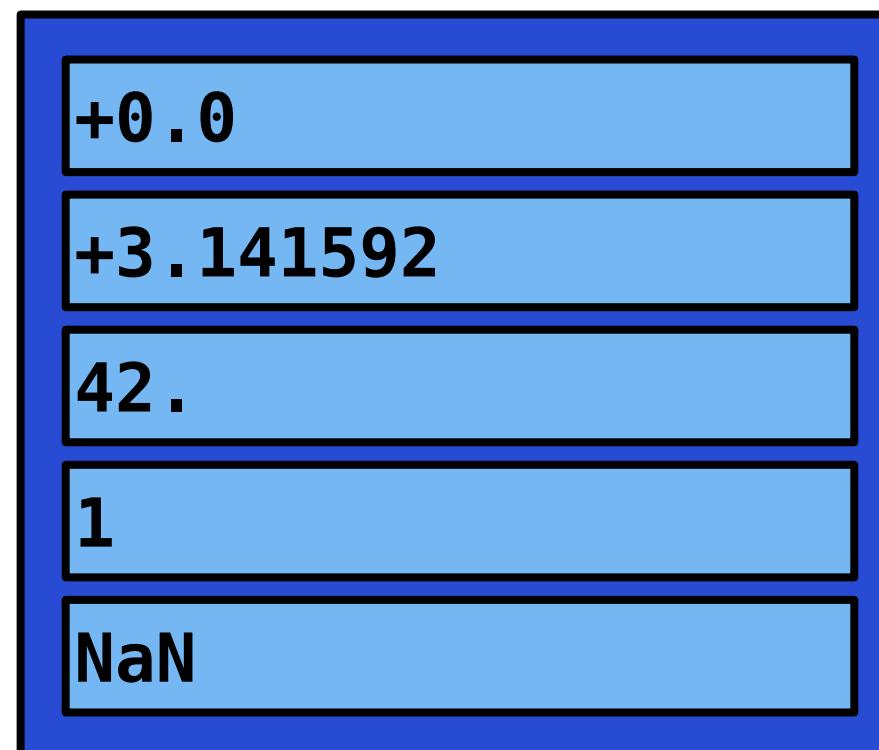
*Dedup est Machina*

# Outline:

## Deduplication

- leak heap & code addresses

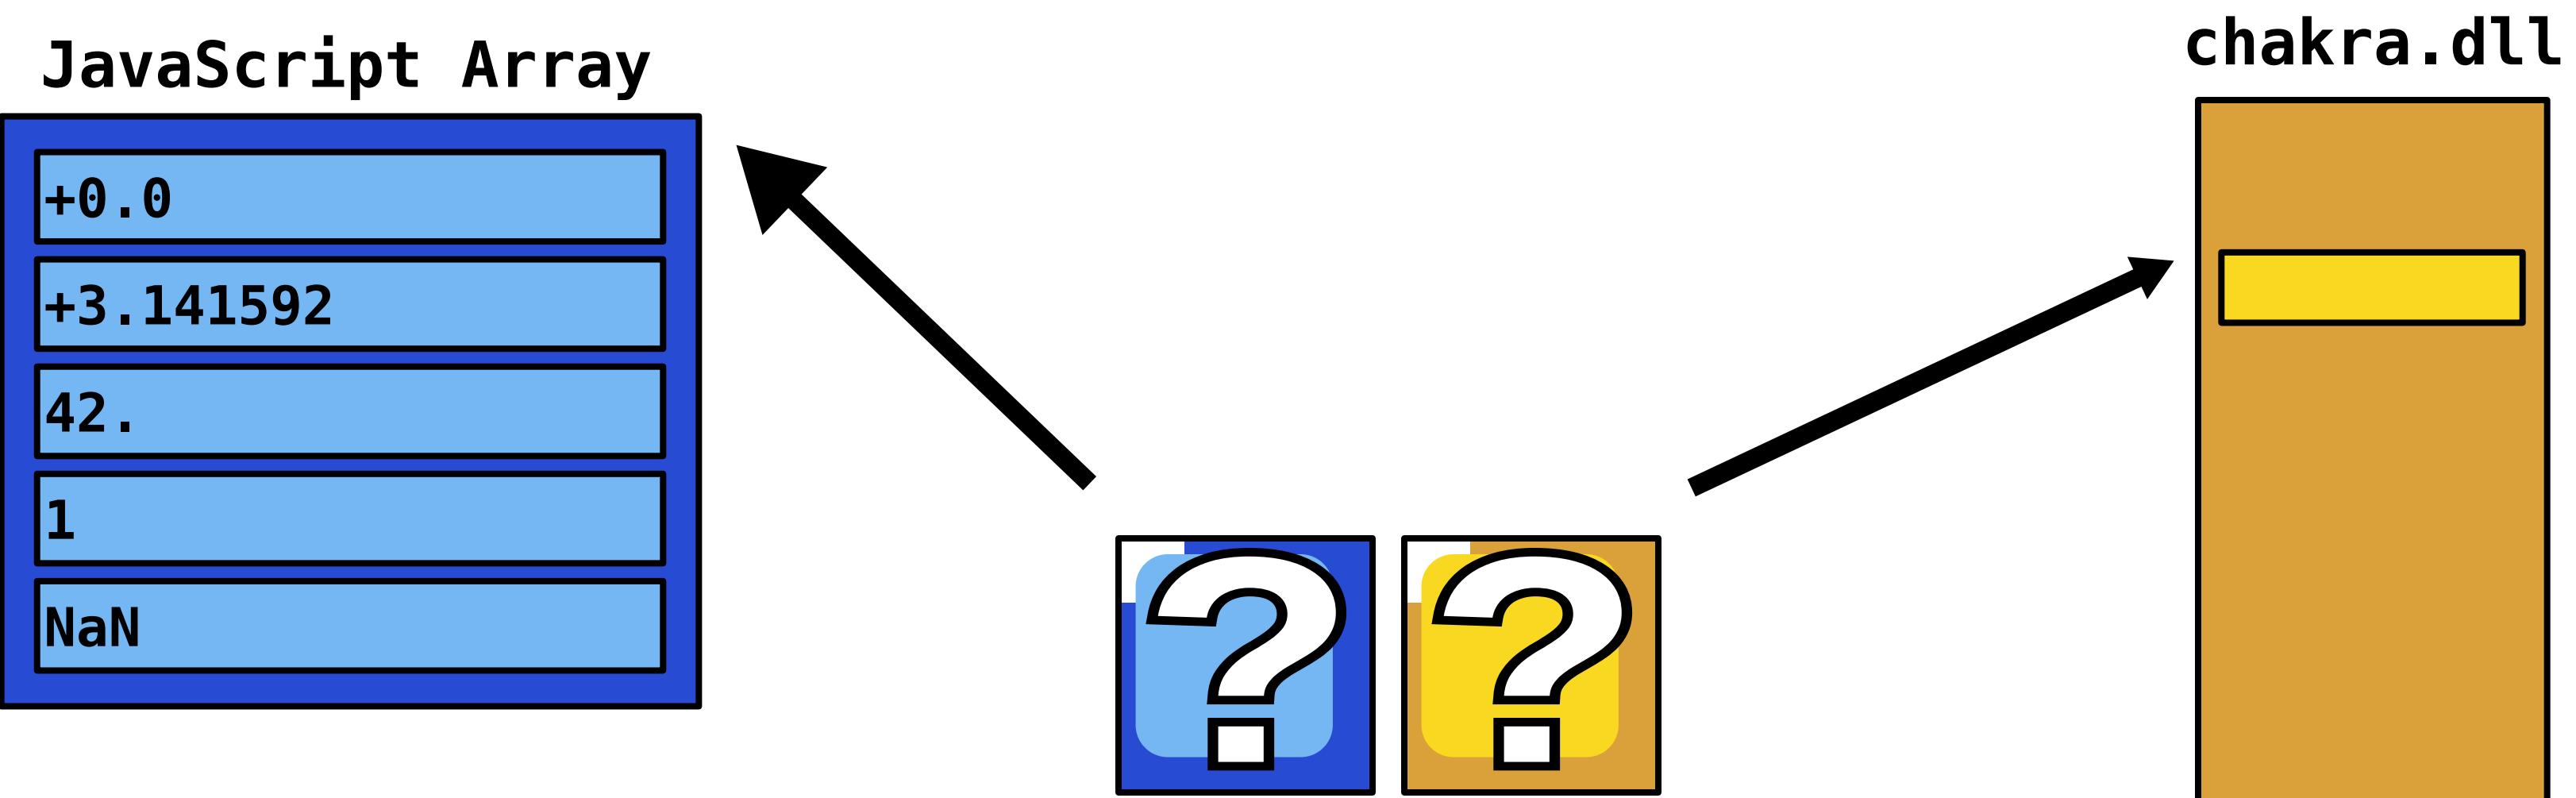
JavaScript Array



# Outline:

## Deduplication

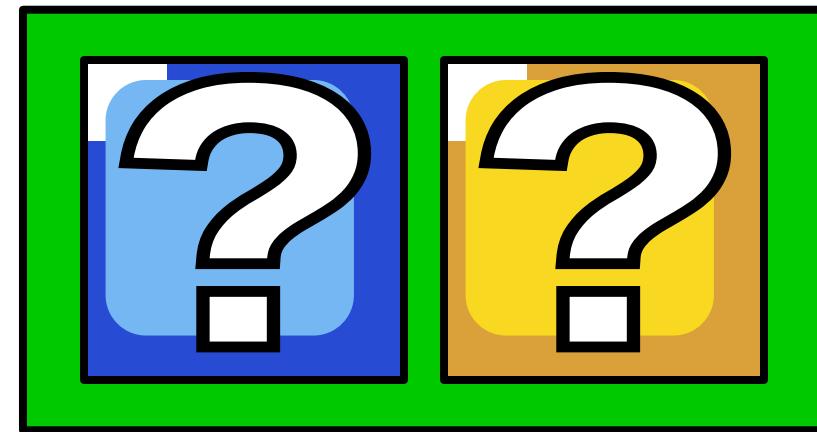
- leak heap & code addresses



# Outline:

## Deduplication

- leak heap & code addresses
- create a fake object



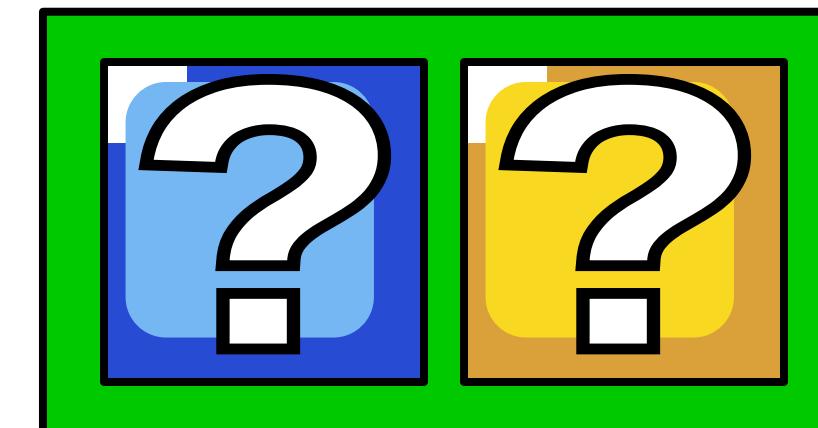
# Outline:

## Deduplication

- leak heap & code addresses
- create a fake object

## Rowhammer

- create reference to our fake object



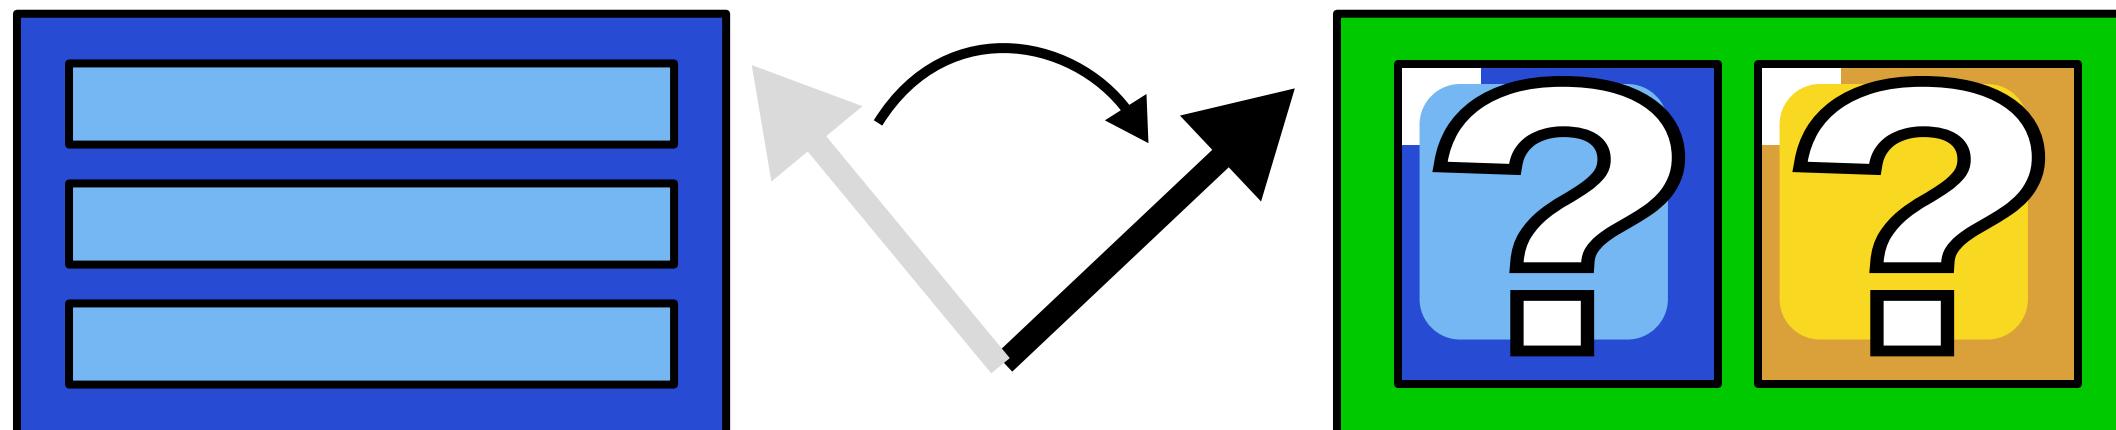
# Outline:

## Deduplication

- leak heap & code addresses
- create a fake object

## Rowhammer

- create reference to our fake object



# memory deduplication

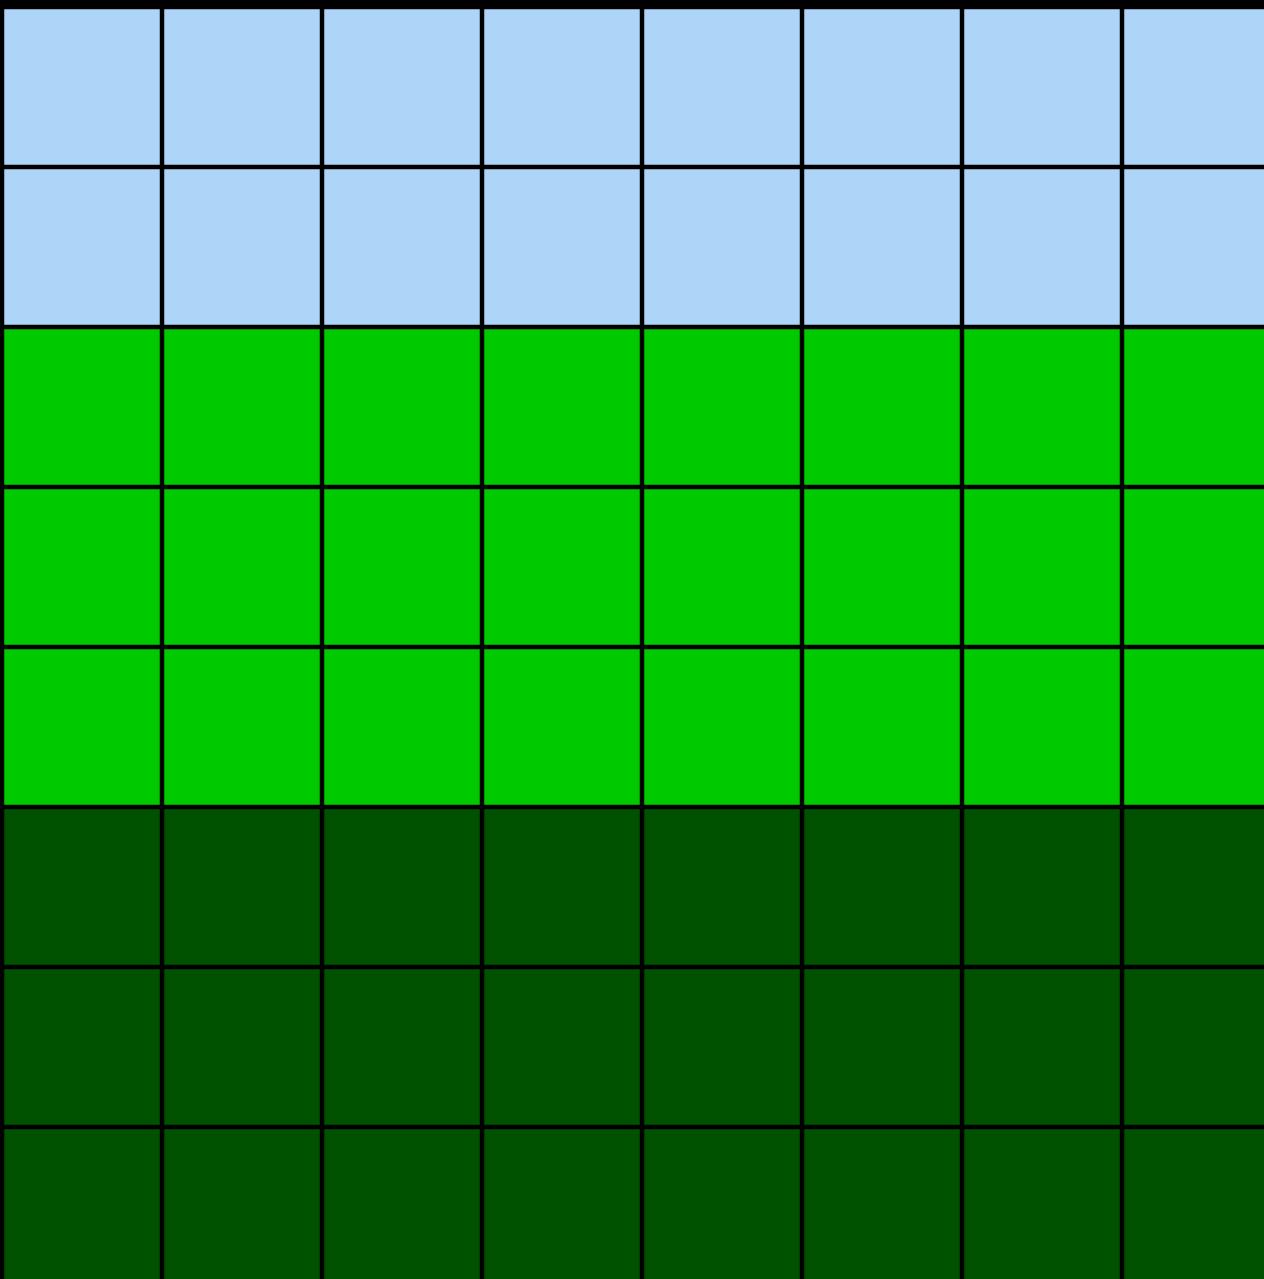
A method of reducing memory usage.

Used in virtualisation environments,

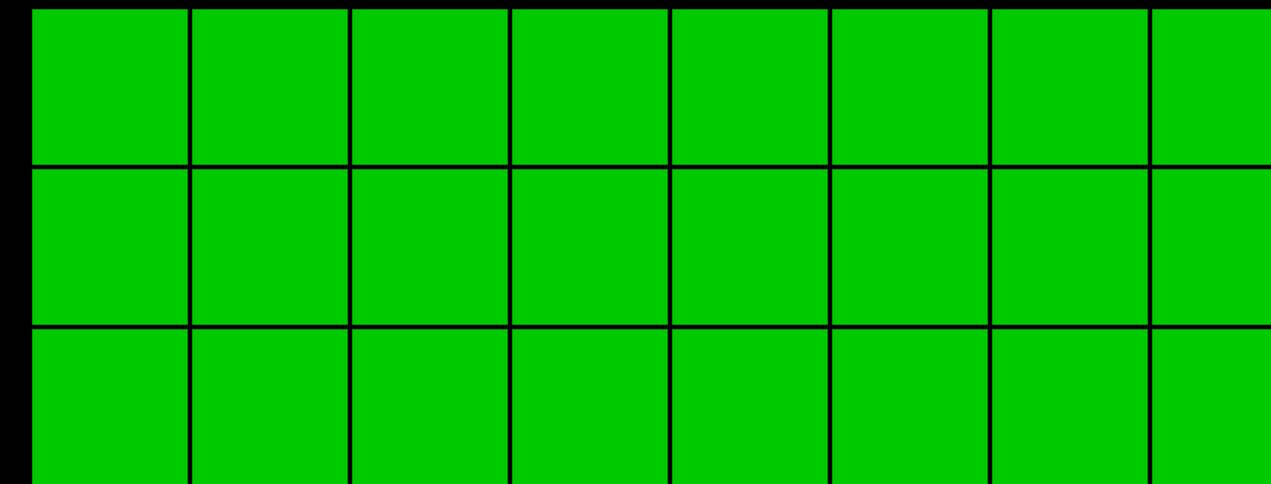
(was) also enabled by default on  
Windows 8.1 and 10.

# memory deduplication

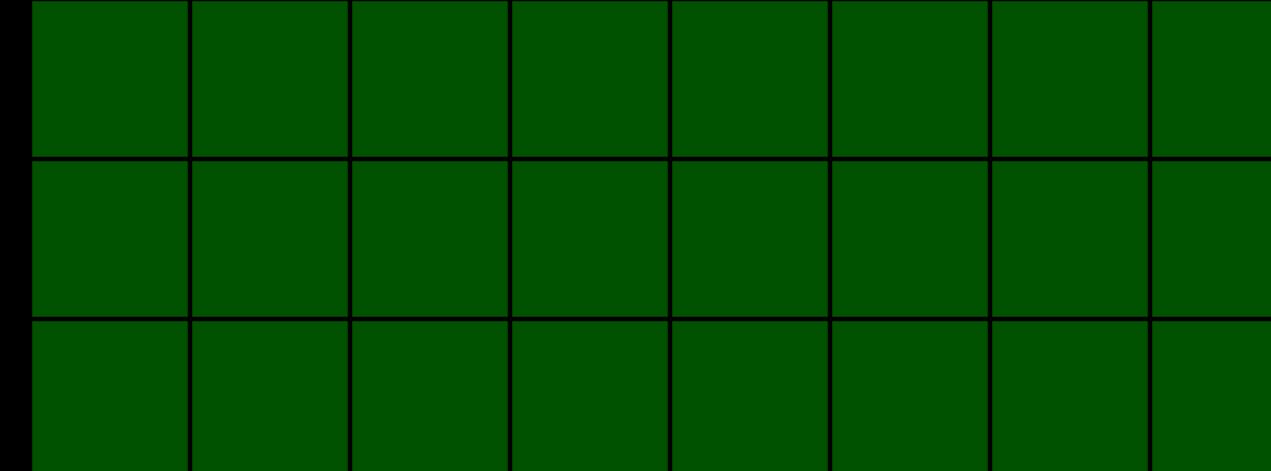
physical memory



process A

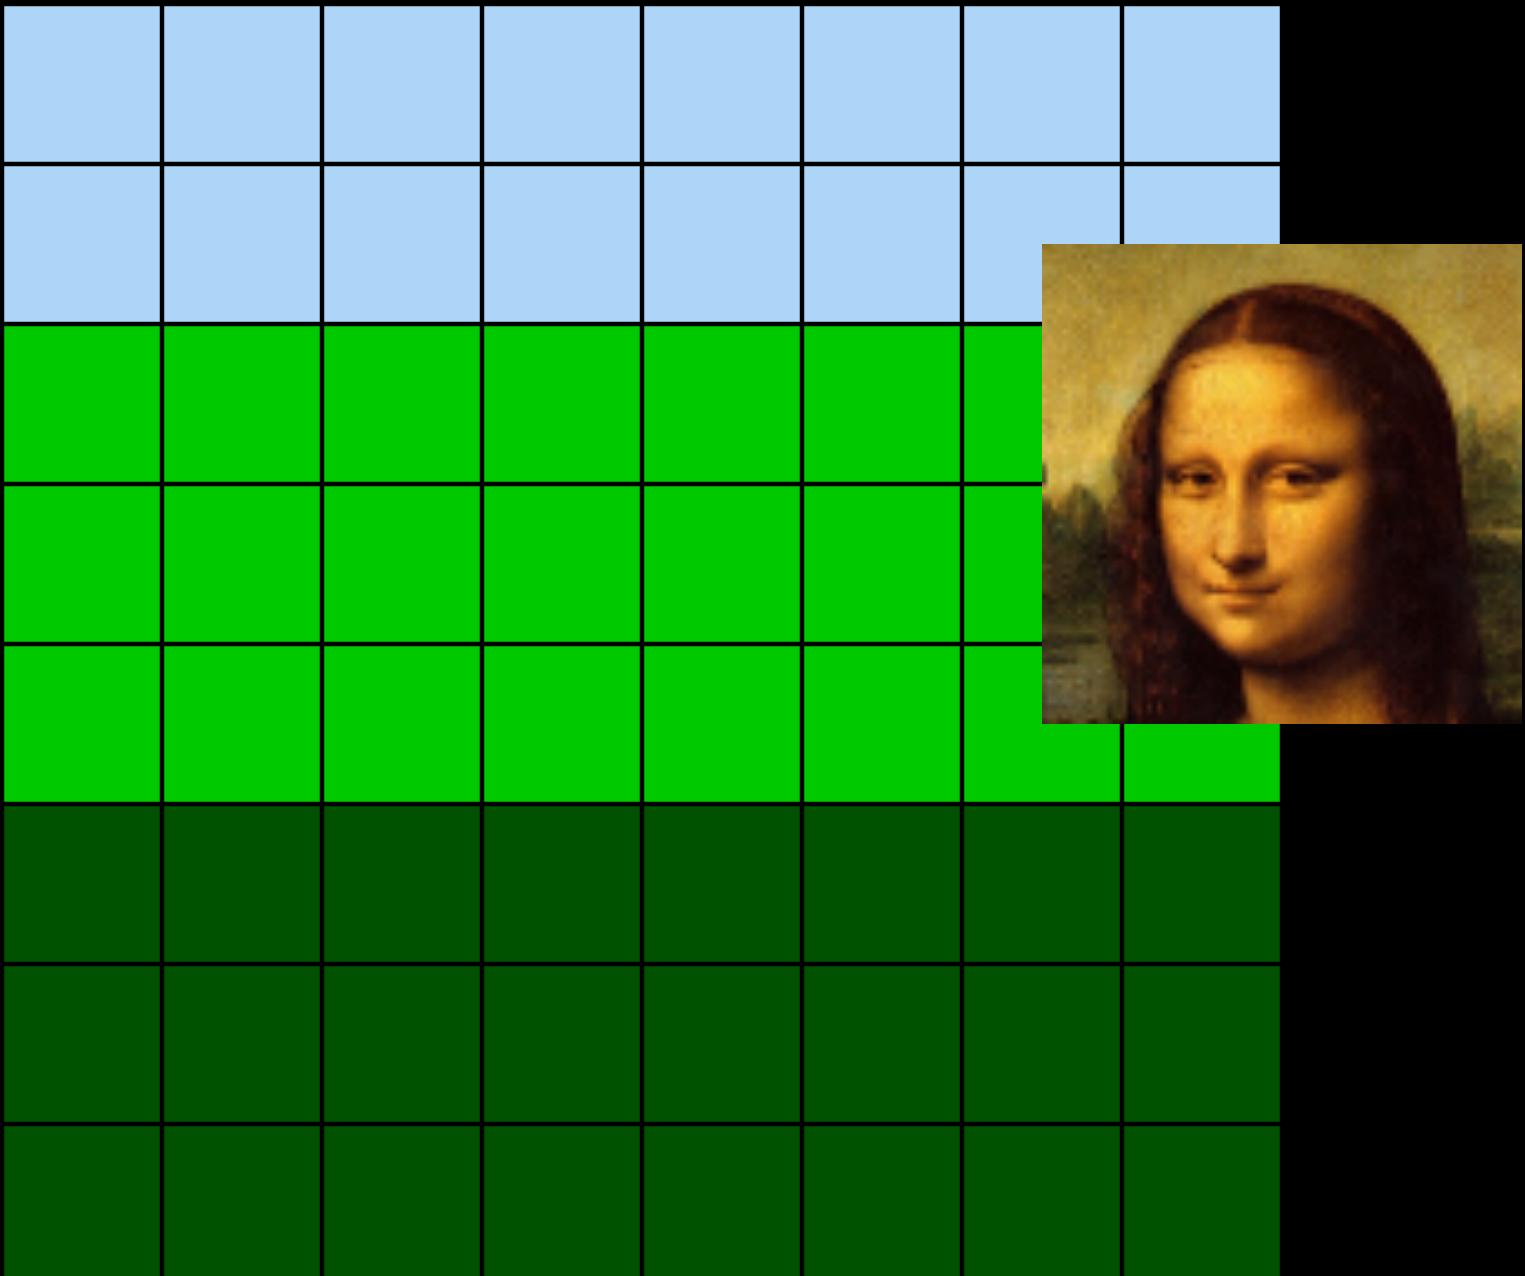


process B

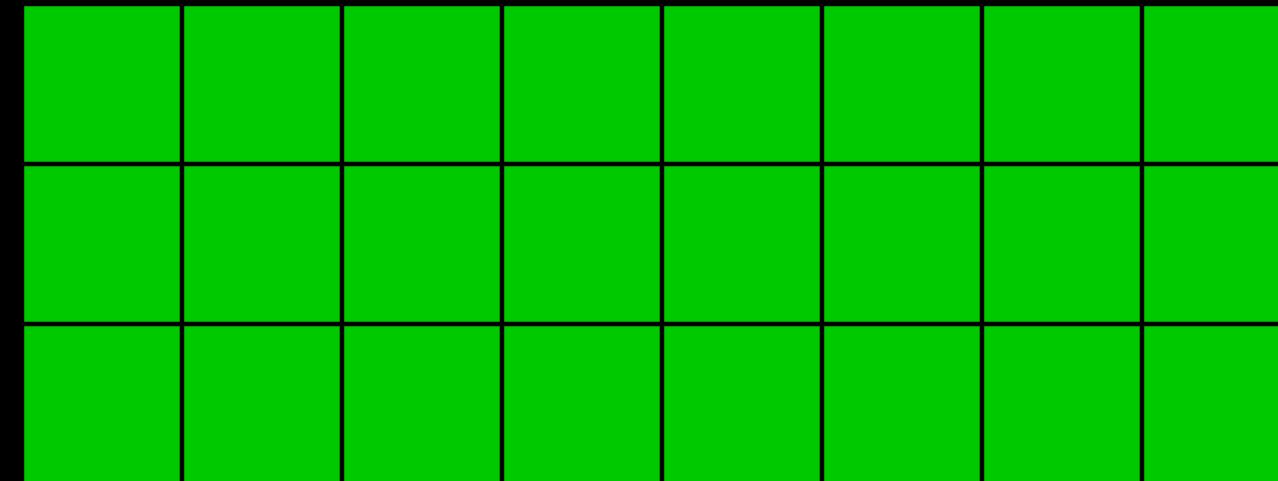


# memory deduplication

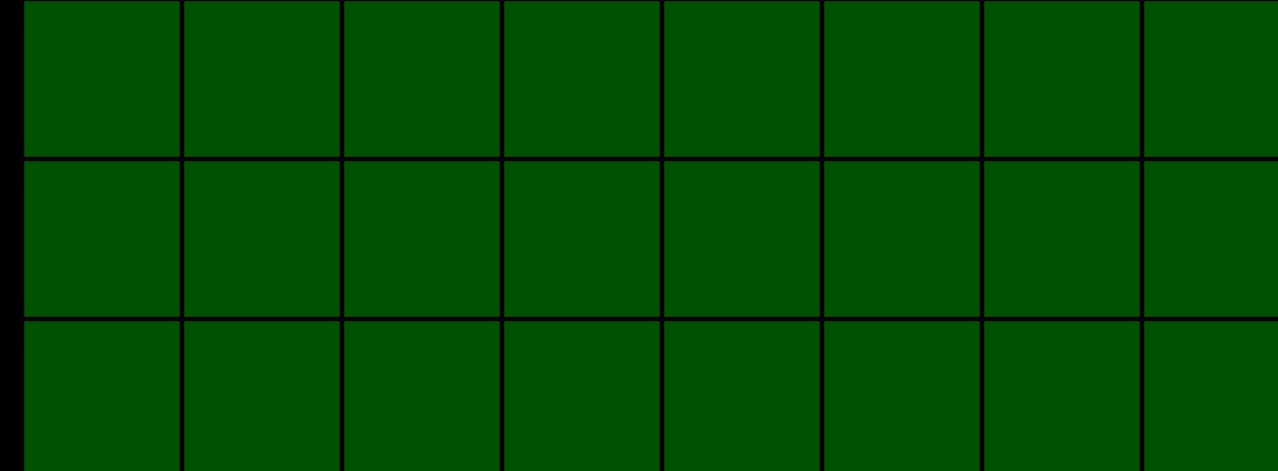
physical memory



process A

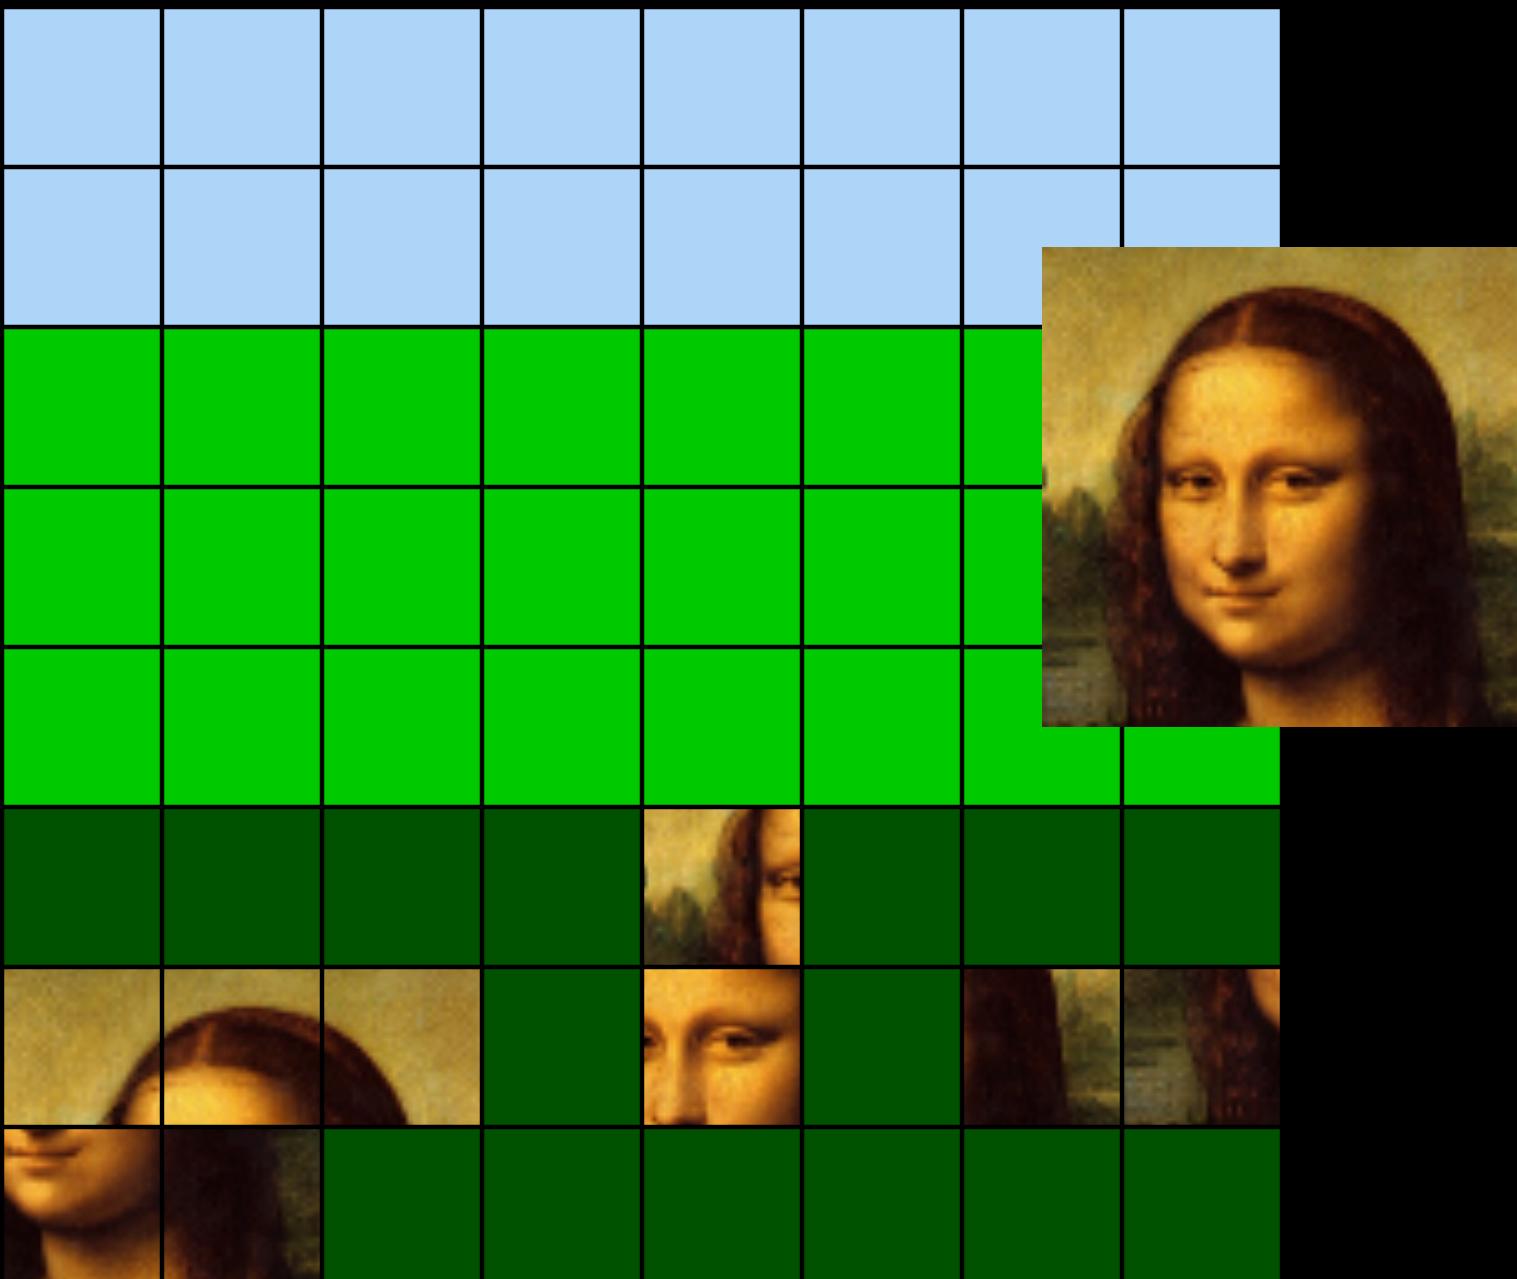


process B

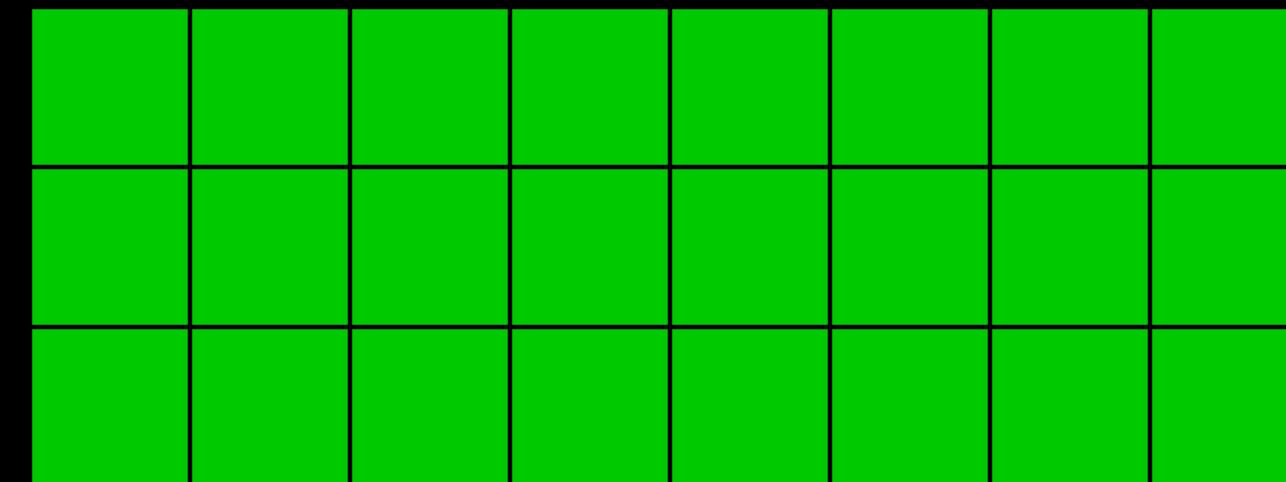


# memory deduplication

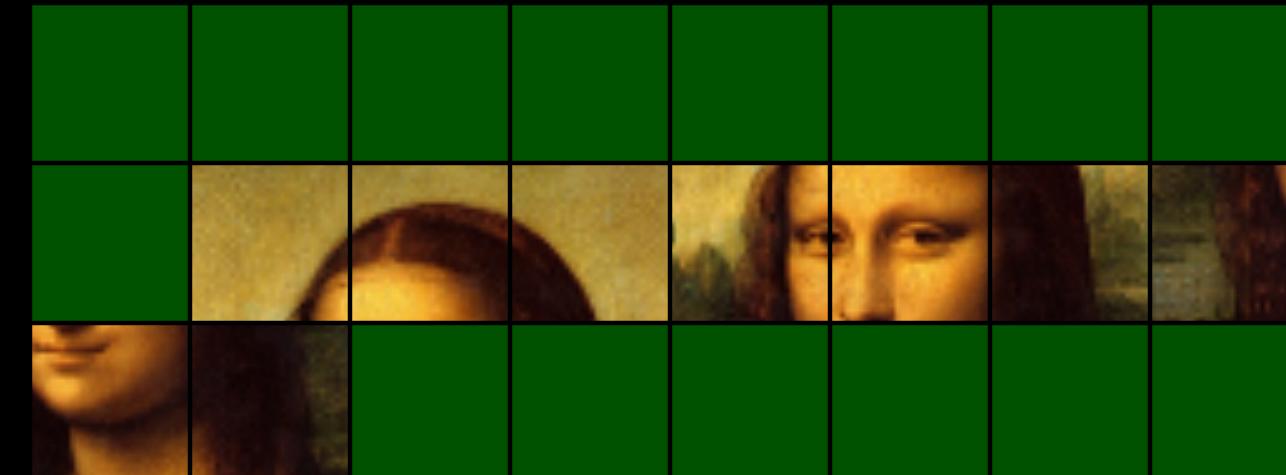
physical memory



process A

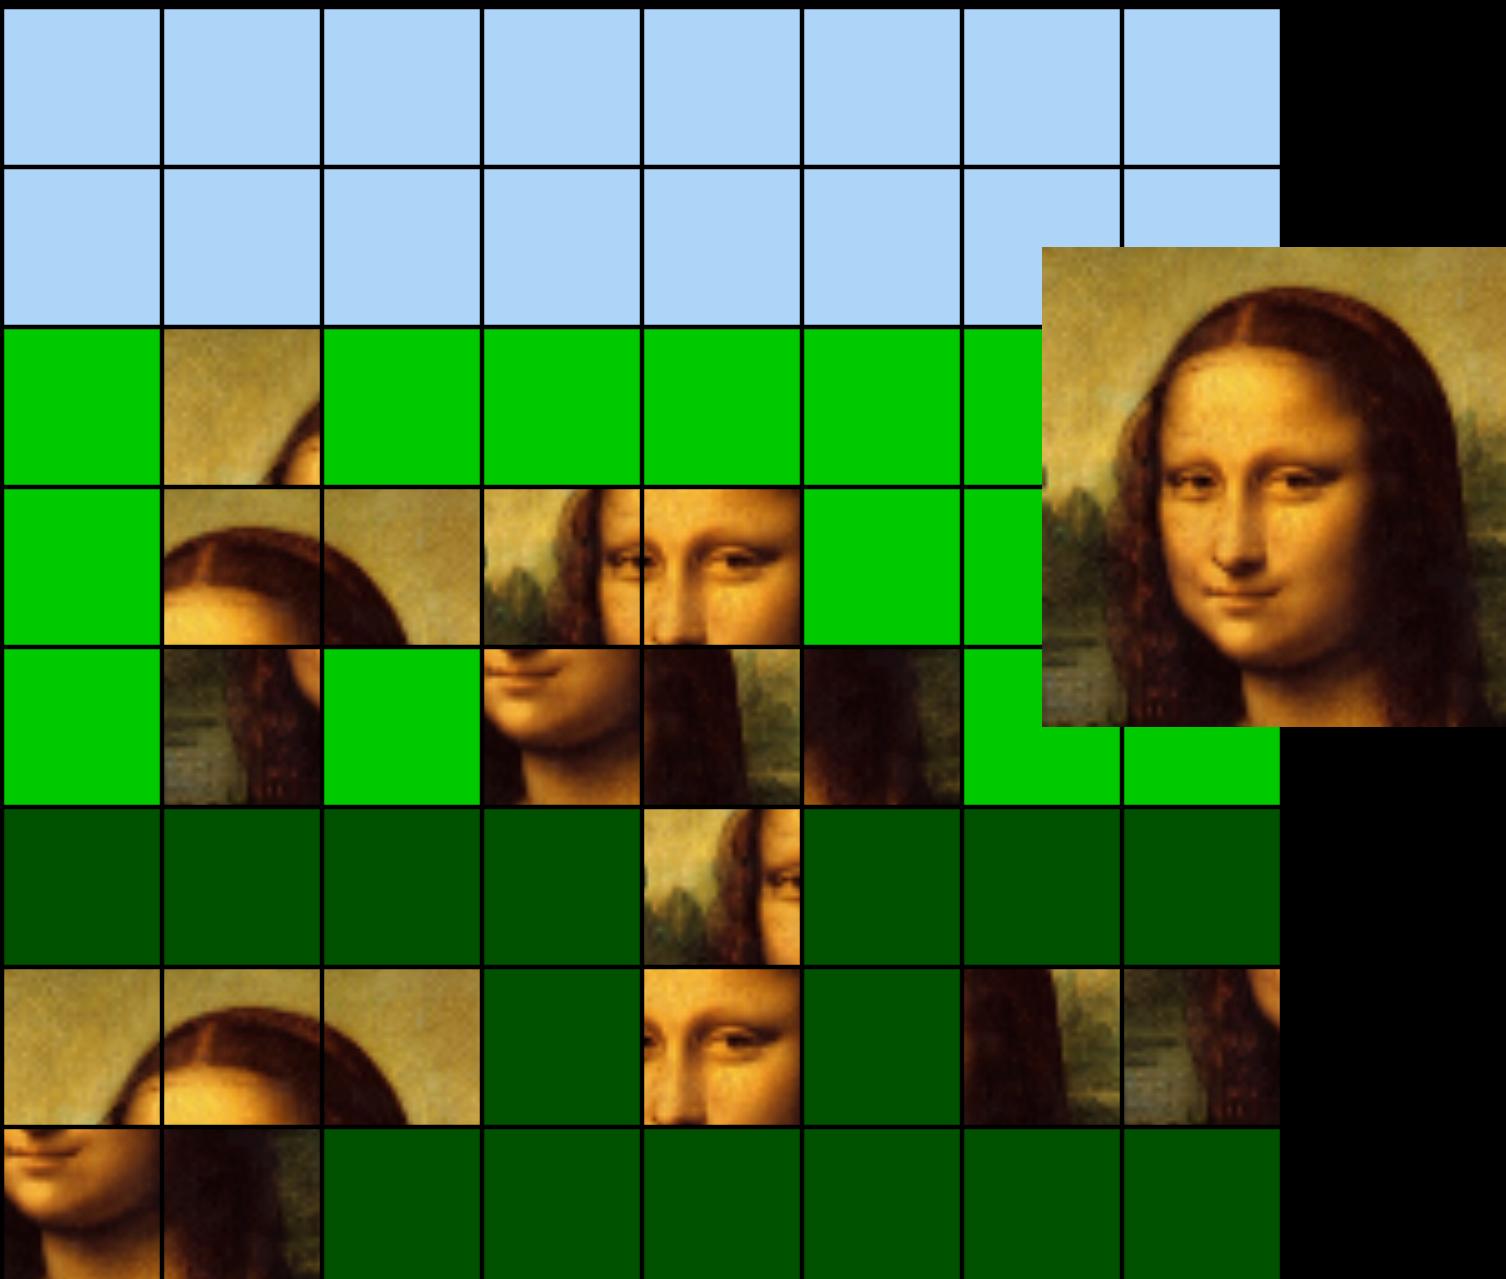


process B

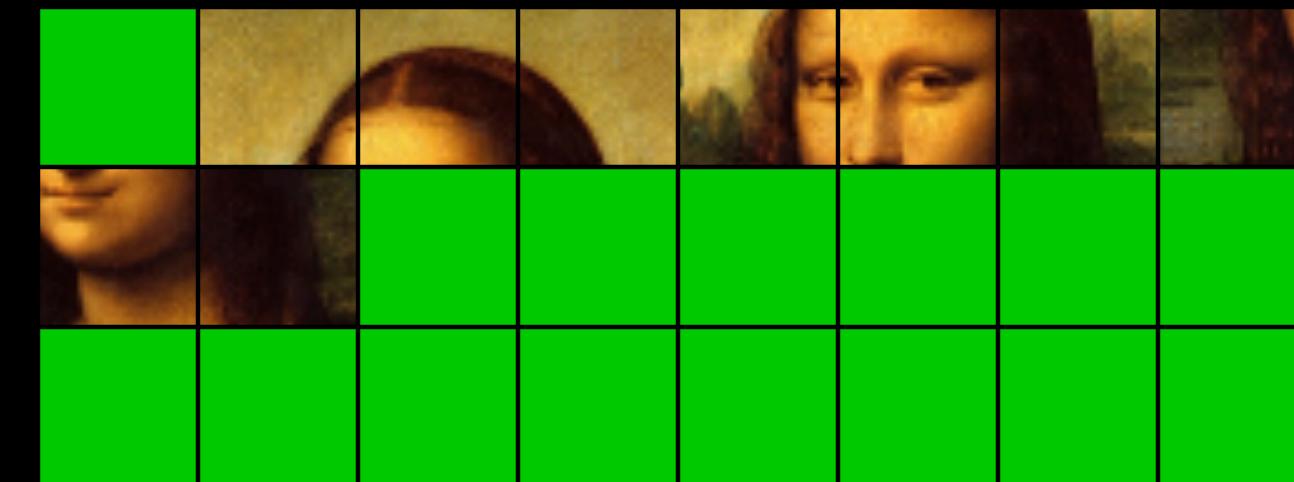


# memory deduplication

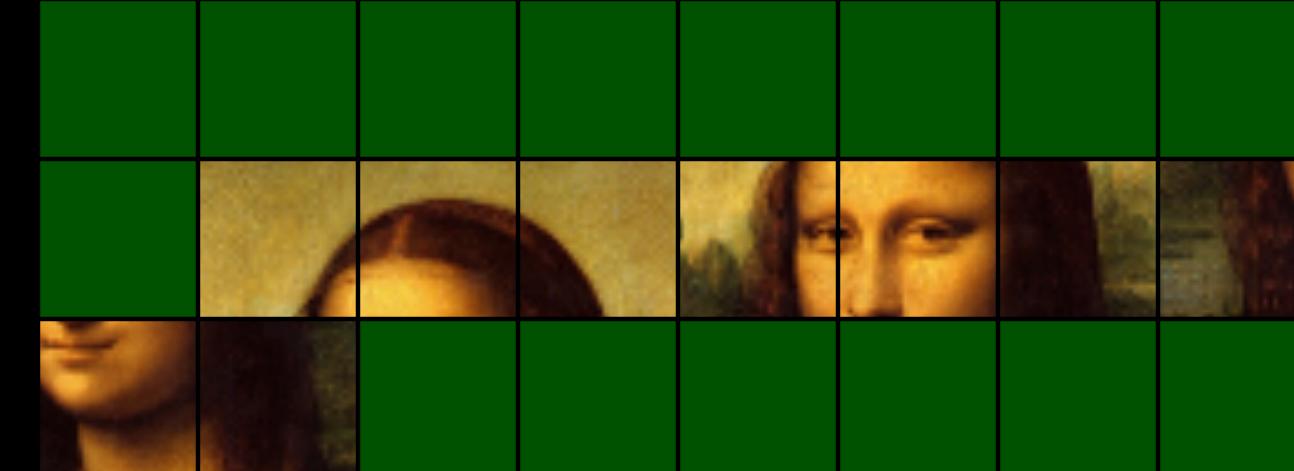
physical memory



process A

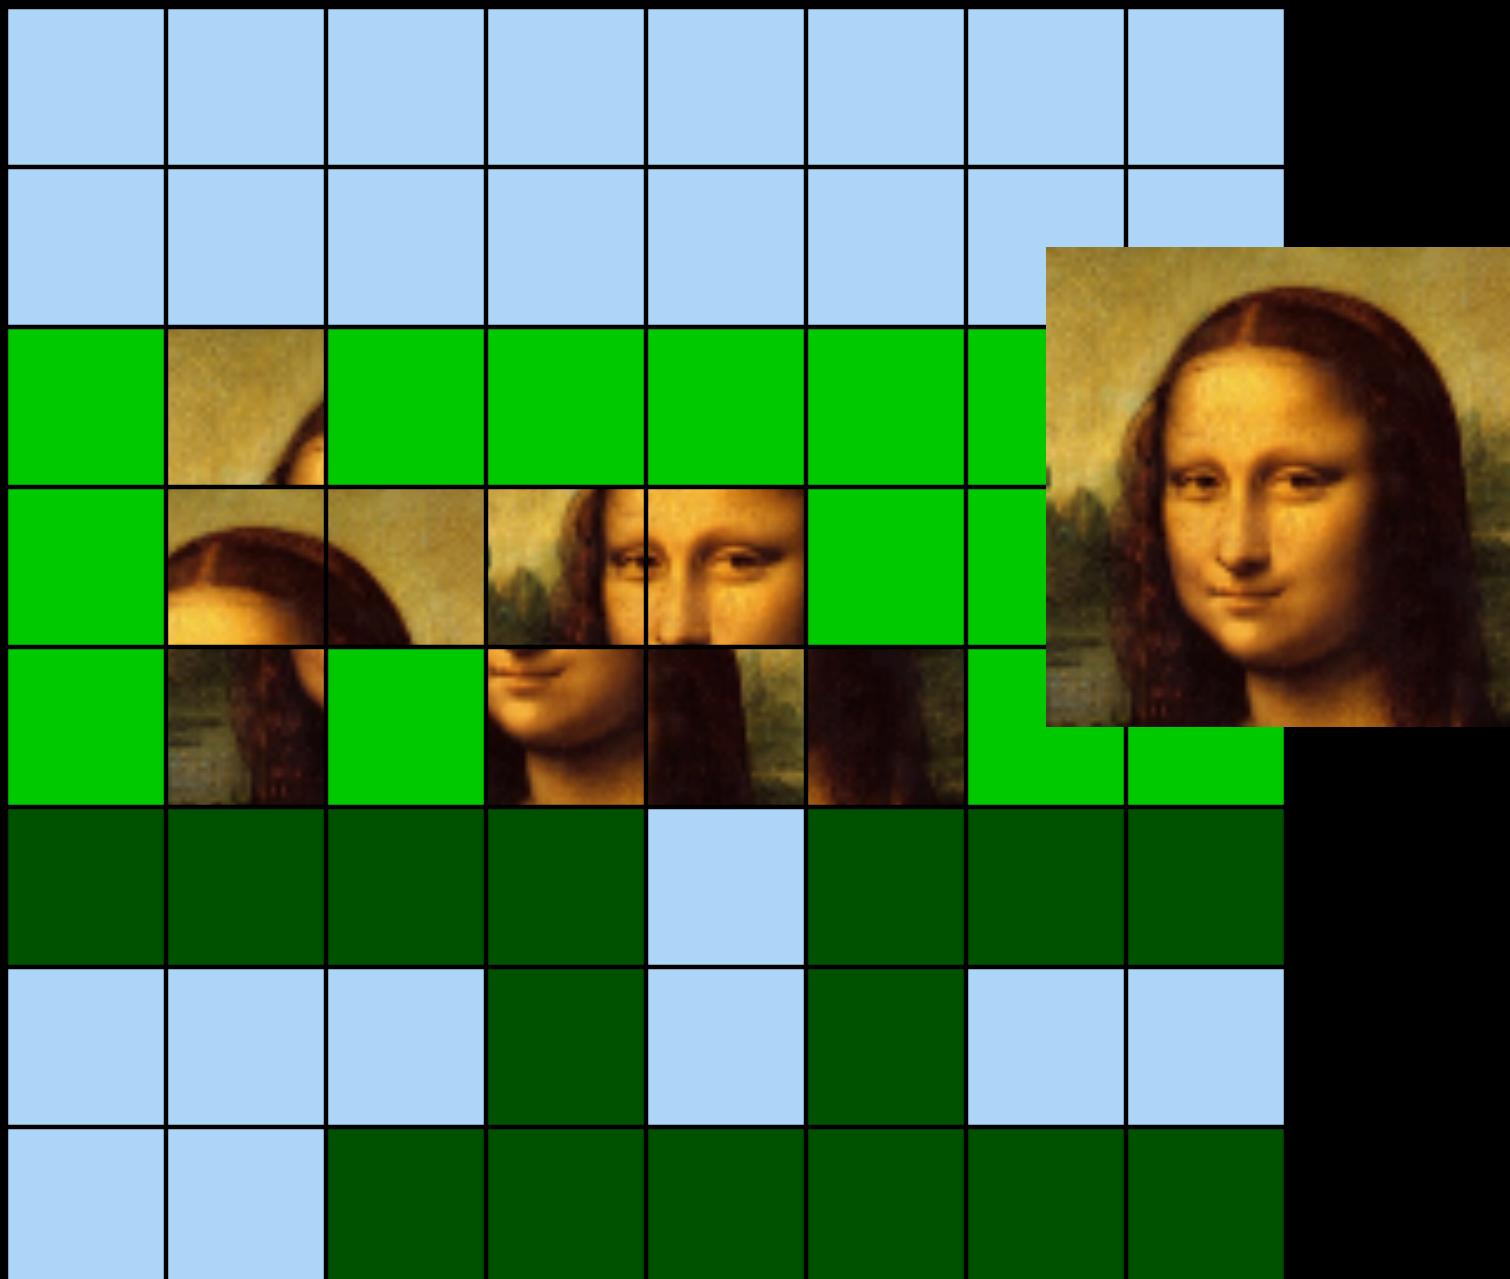


process B

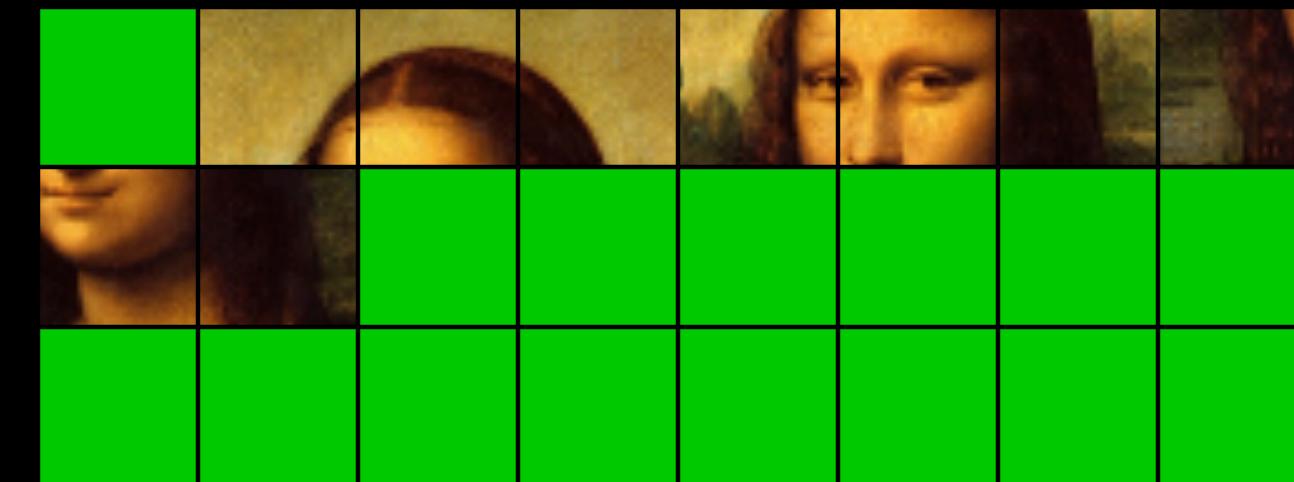


# memory deduplication

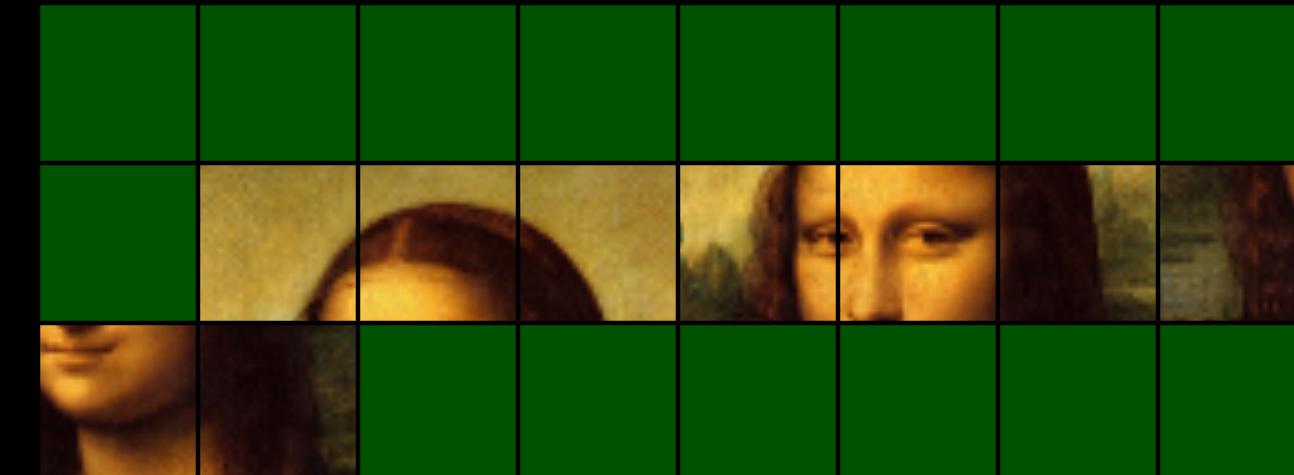
physical memory



process A

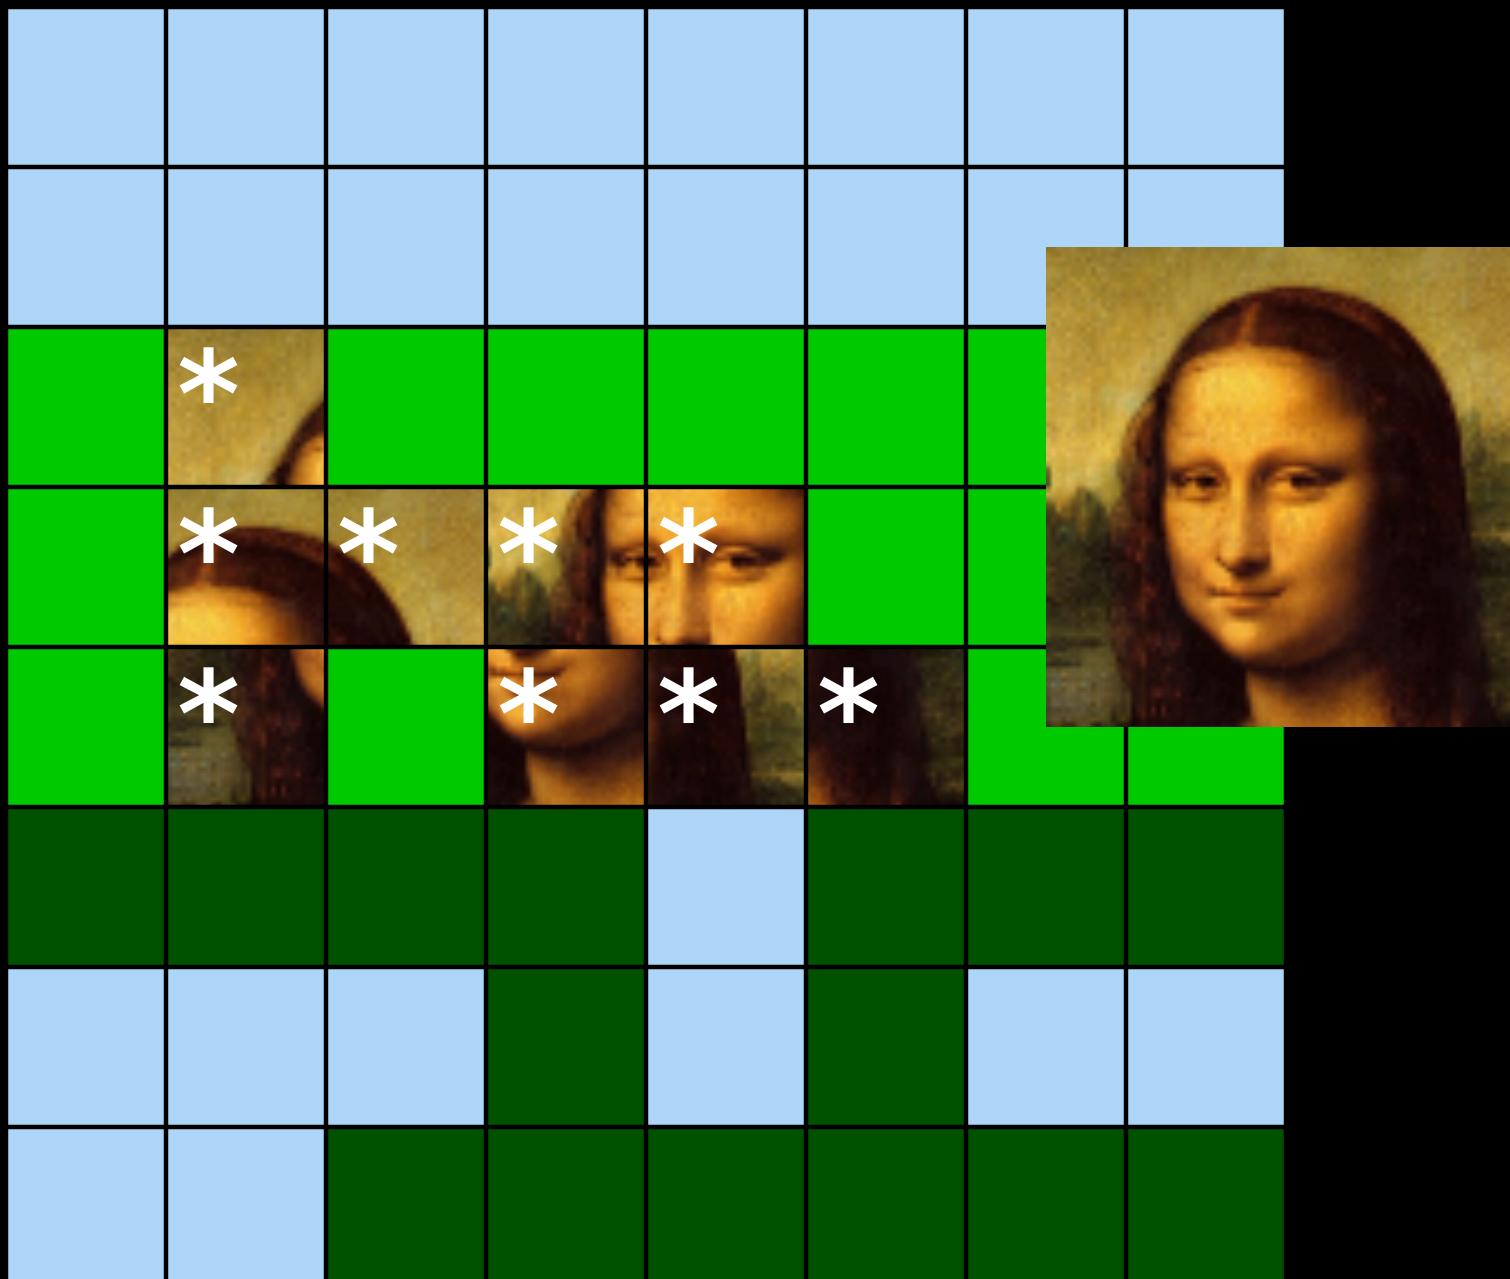


process B

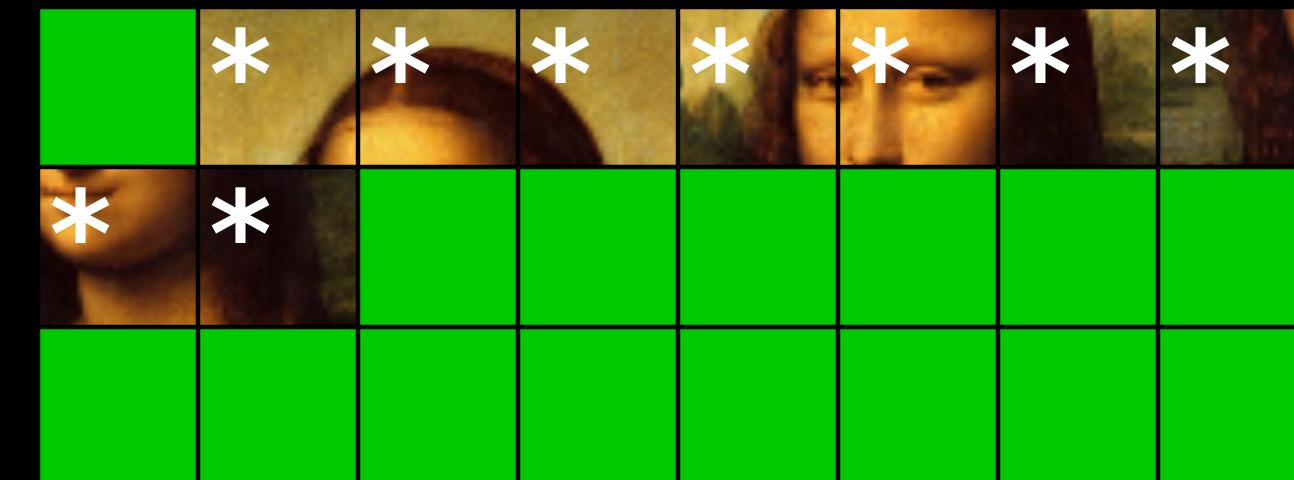


# memory deduplication

physical memory



process A



process B



# memory deduplication: The Problem

Deduplicated memory does not need  
to have the same origin.

(unlike `fork()`, file-backed memory)

An attacker can use deduplication  
as a side-channel

# deduplication side-channel attack

normal write



# deduplication side-channel attack

normal write

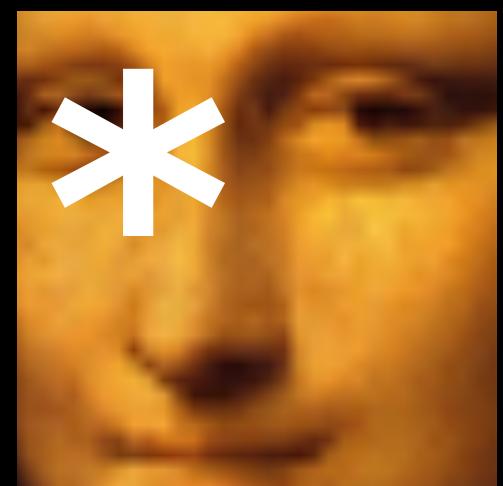


# deduplication side-channel attack

normal write



copy on write (due to deduplication)



# deduplication side-channel attack

normal write



copy on write (due to deduplication)

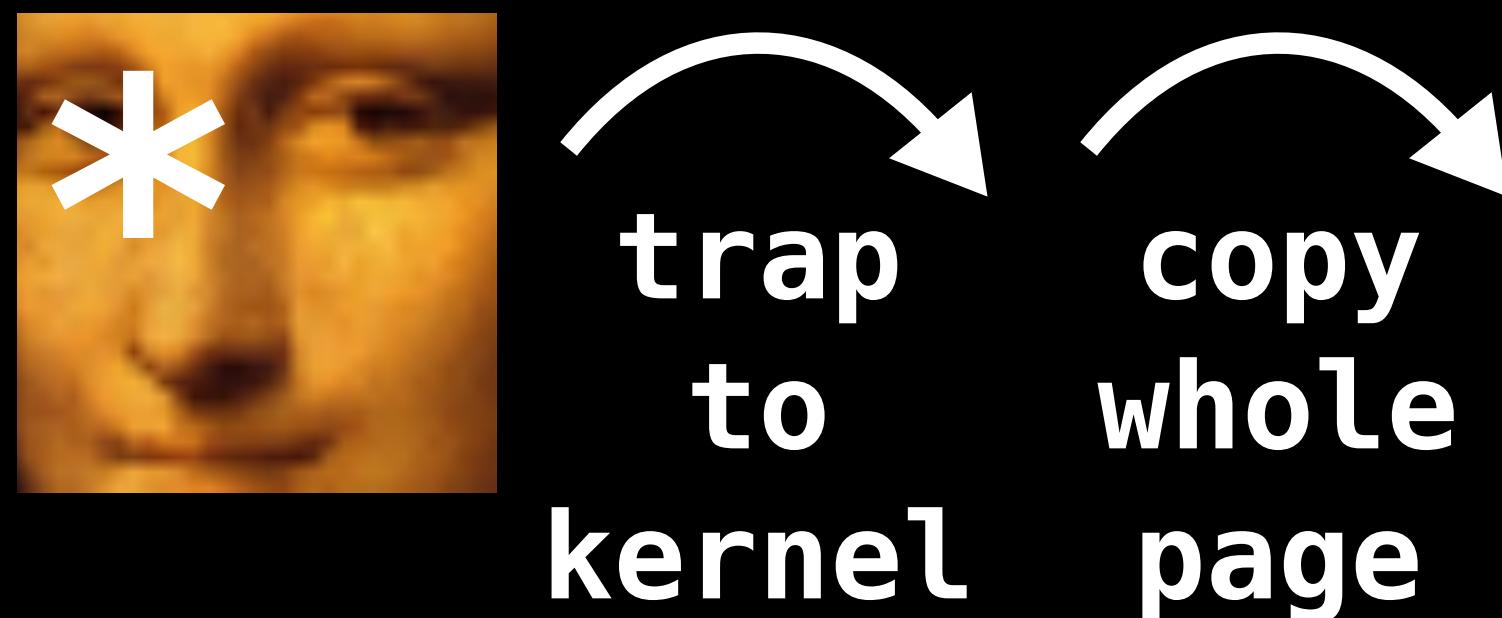


# deduplication side-channel attack

normal write



copy on write (due to deduplication)

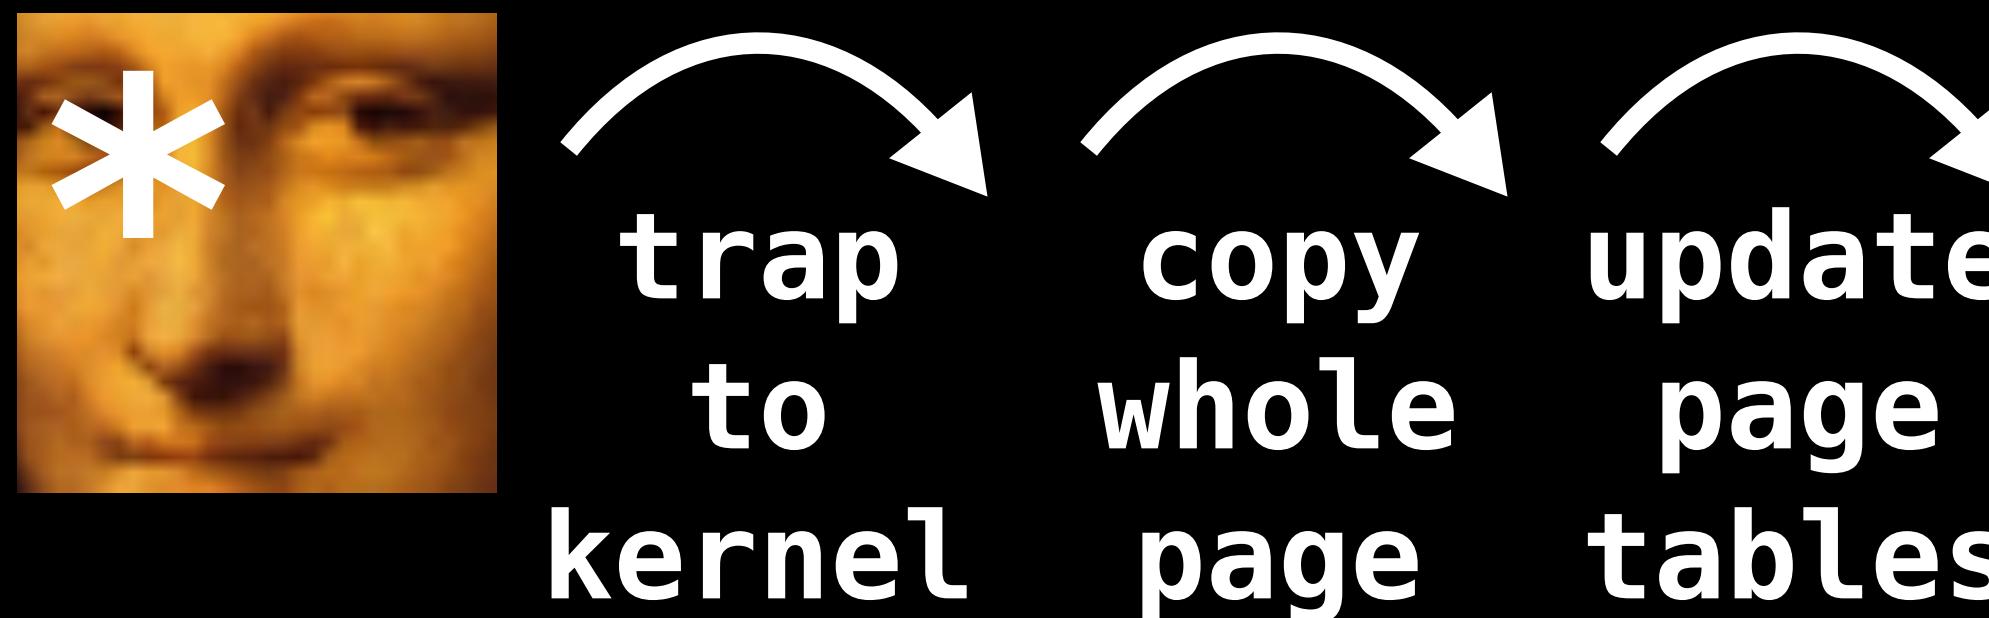


# deduplication side-channel attack

normal write



copy on write (due to deduplication)

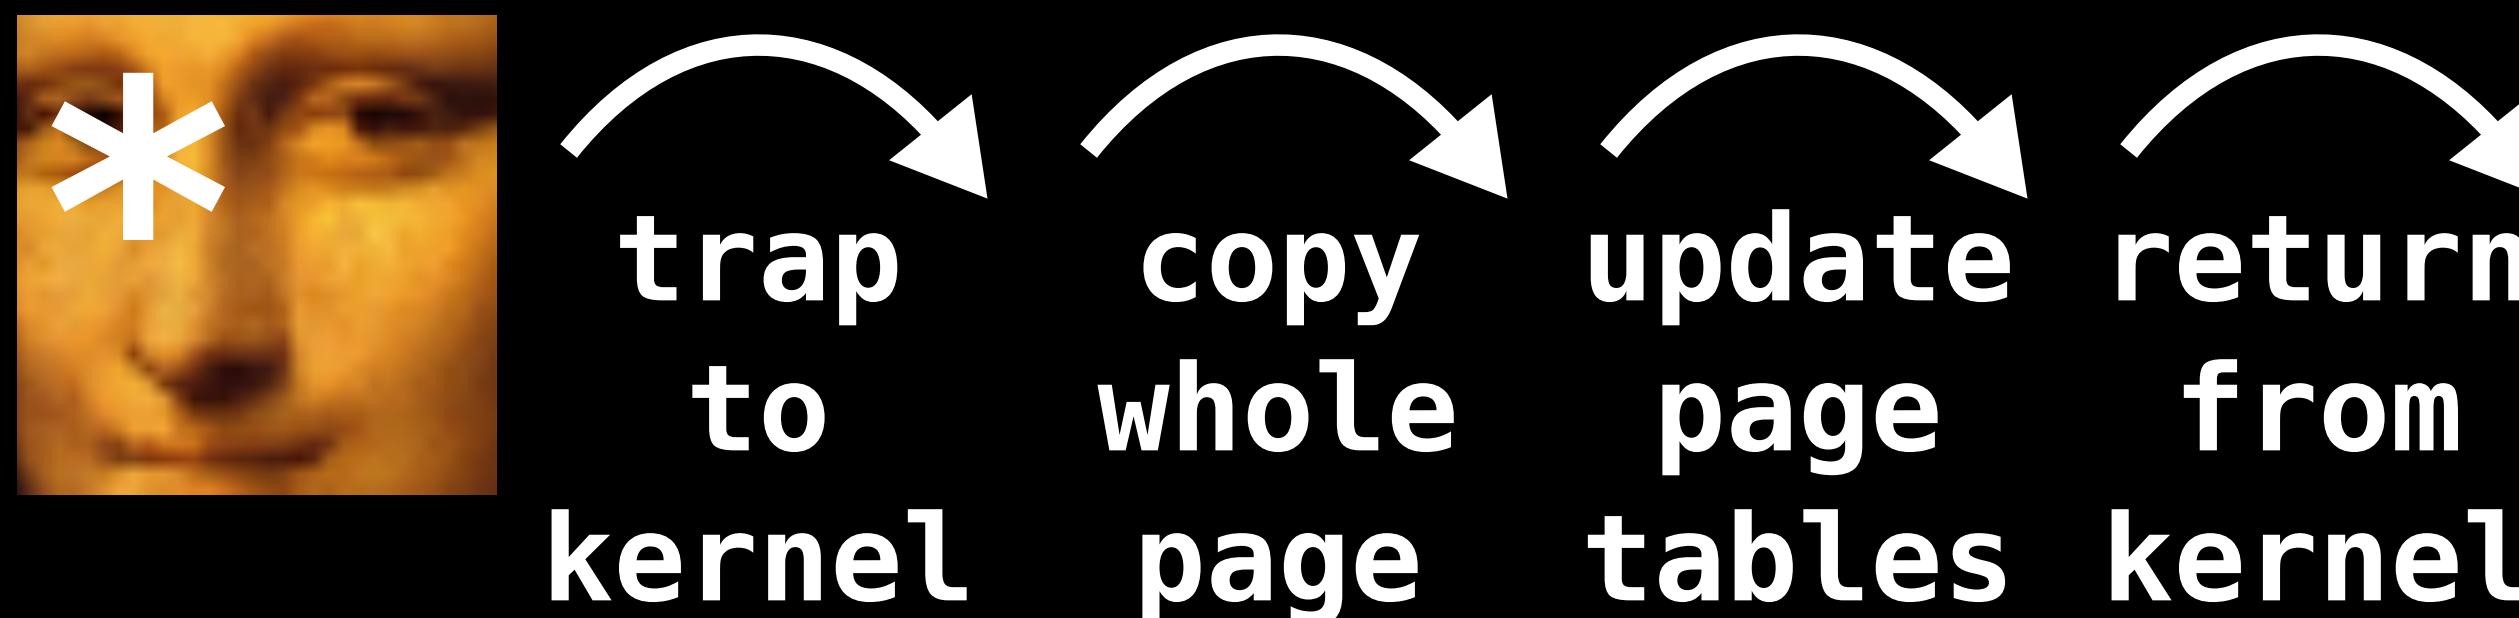


# deduplication side-channel attack

normal write



copy on write (due to deduplication)

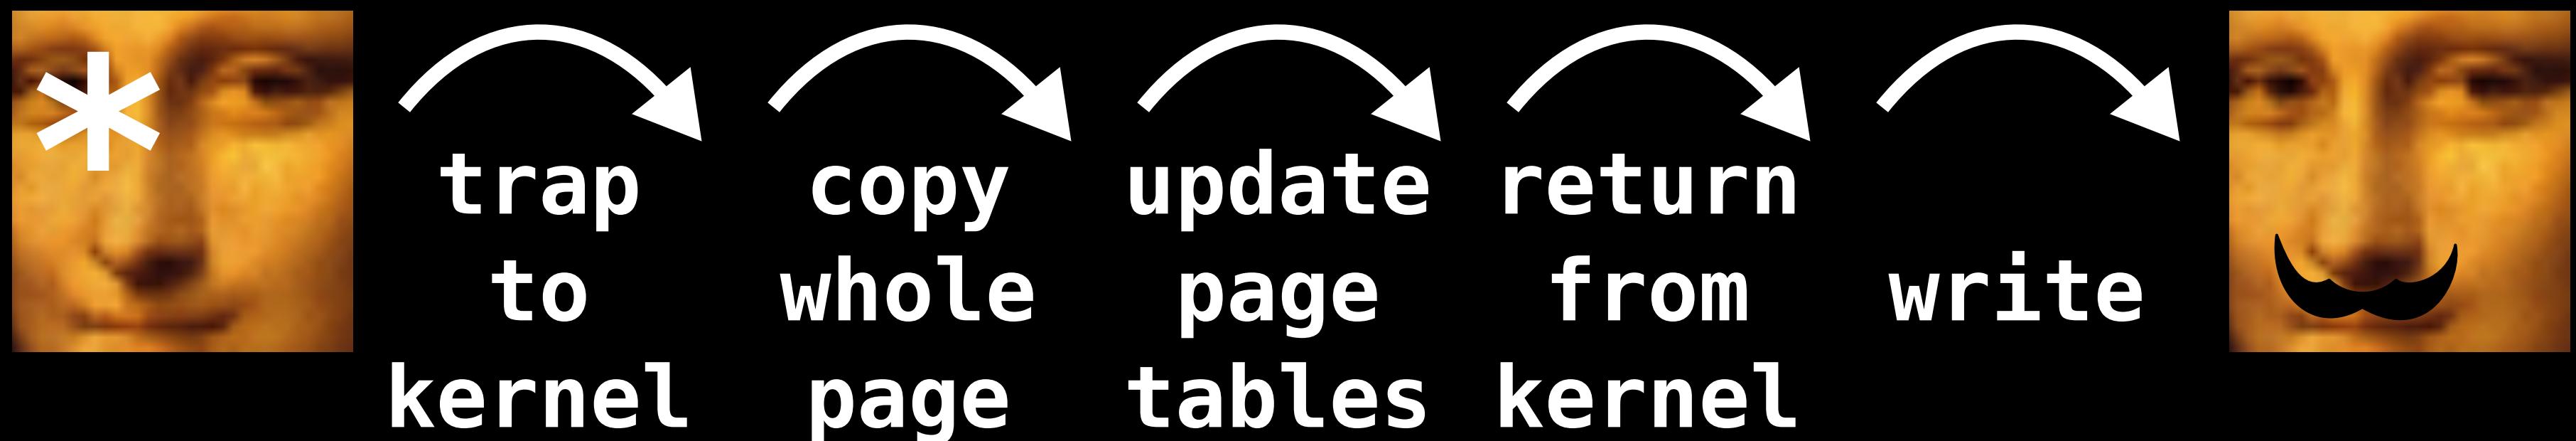


# deduplication side-channel attack

normal write



copy on write (due to deduplication)



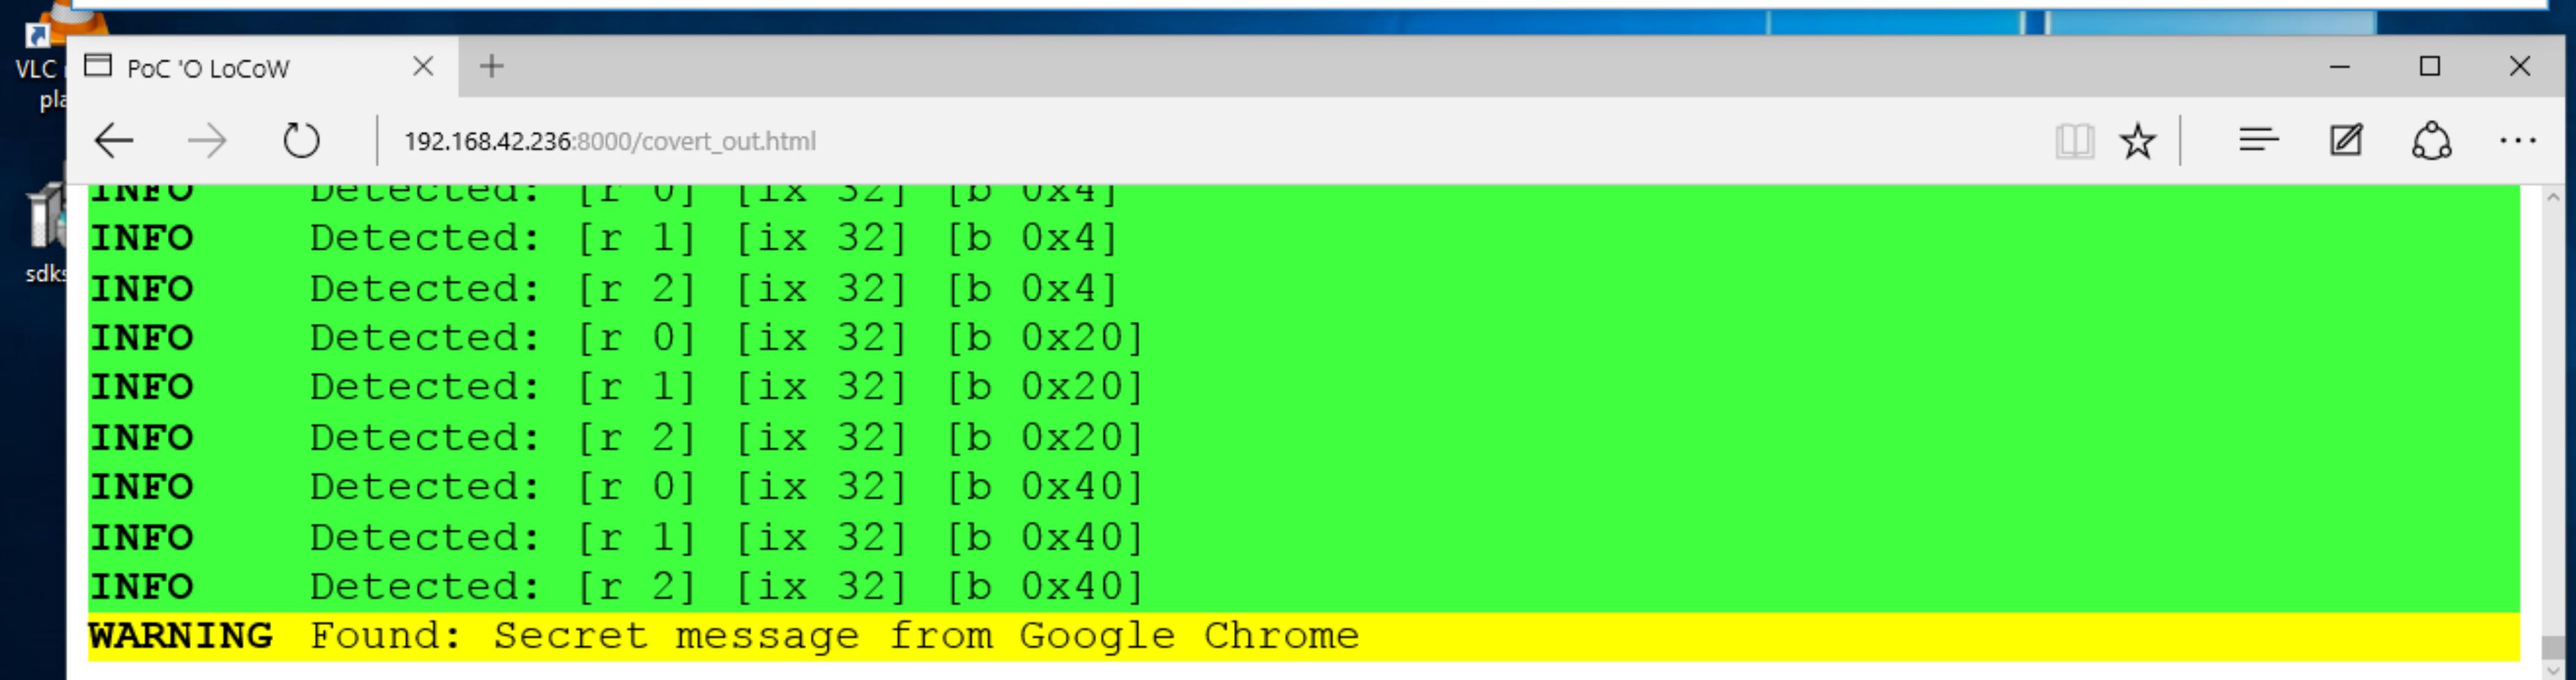
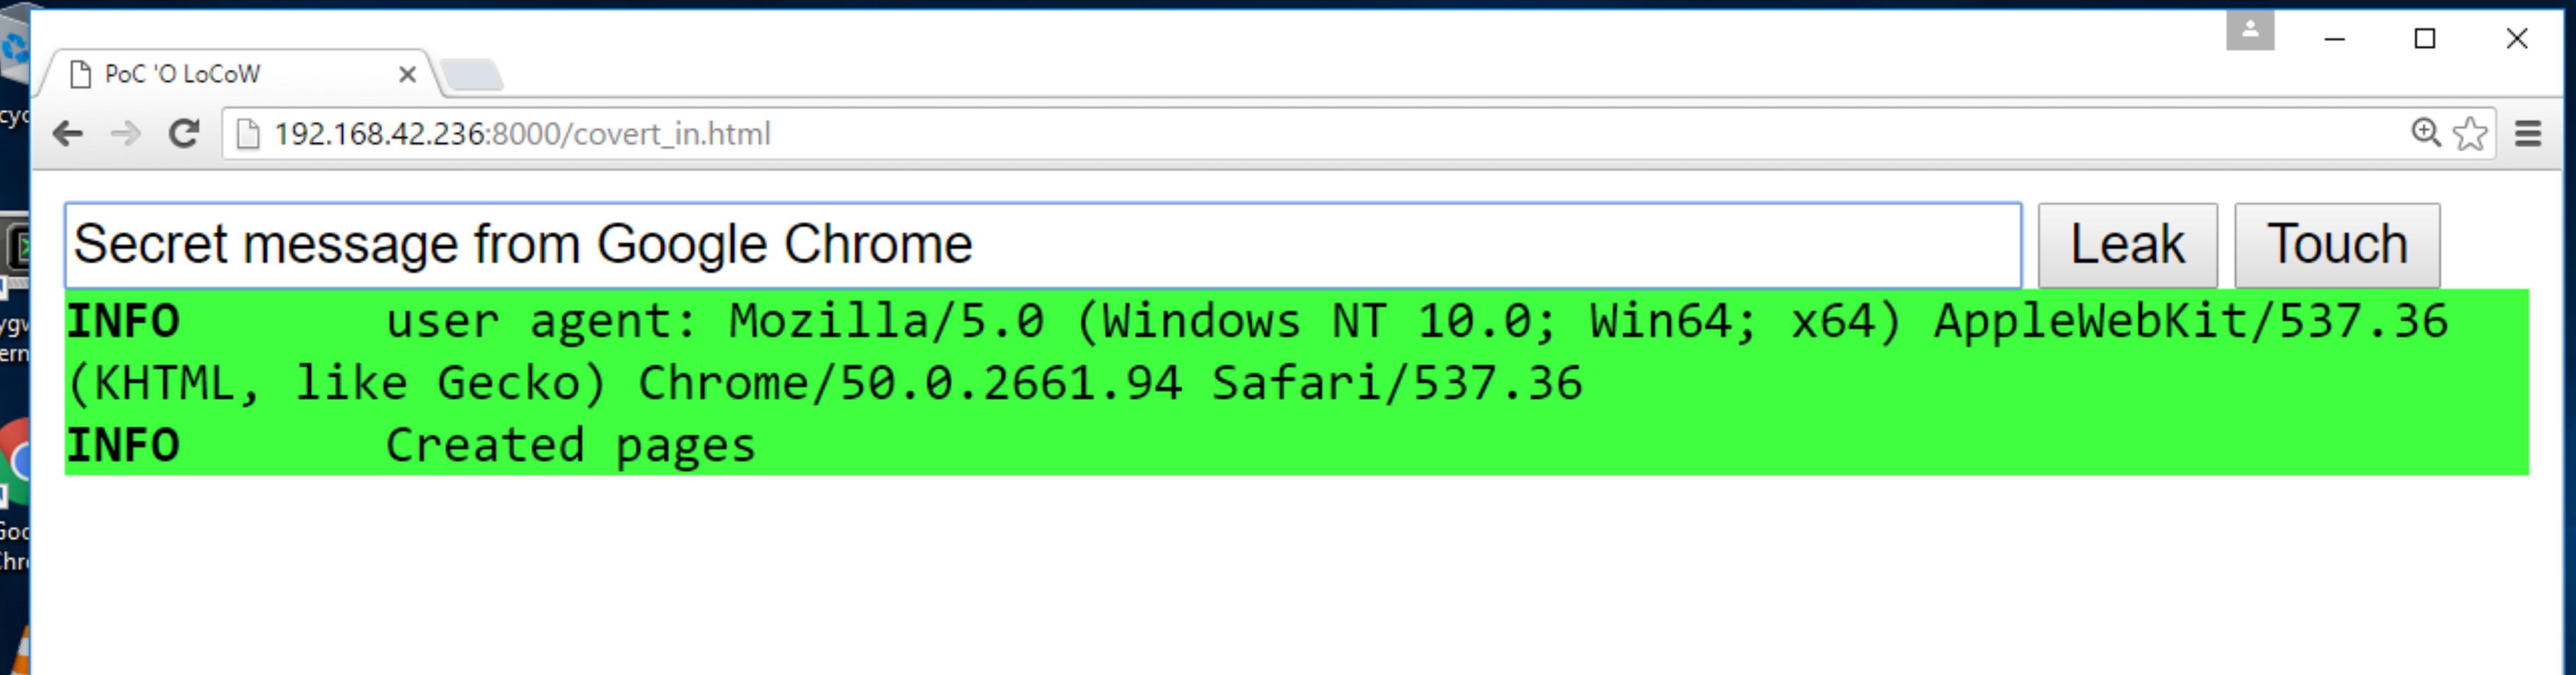
# Deduplication side-channel attack

A 1-bit side channel which is able to leak data across security boundaries

- > cross VM
- > cross-process
- > leak process data from javascript code

# Having fun with deduplication

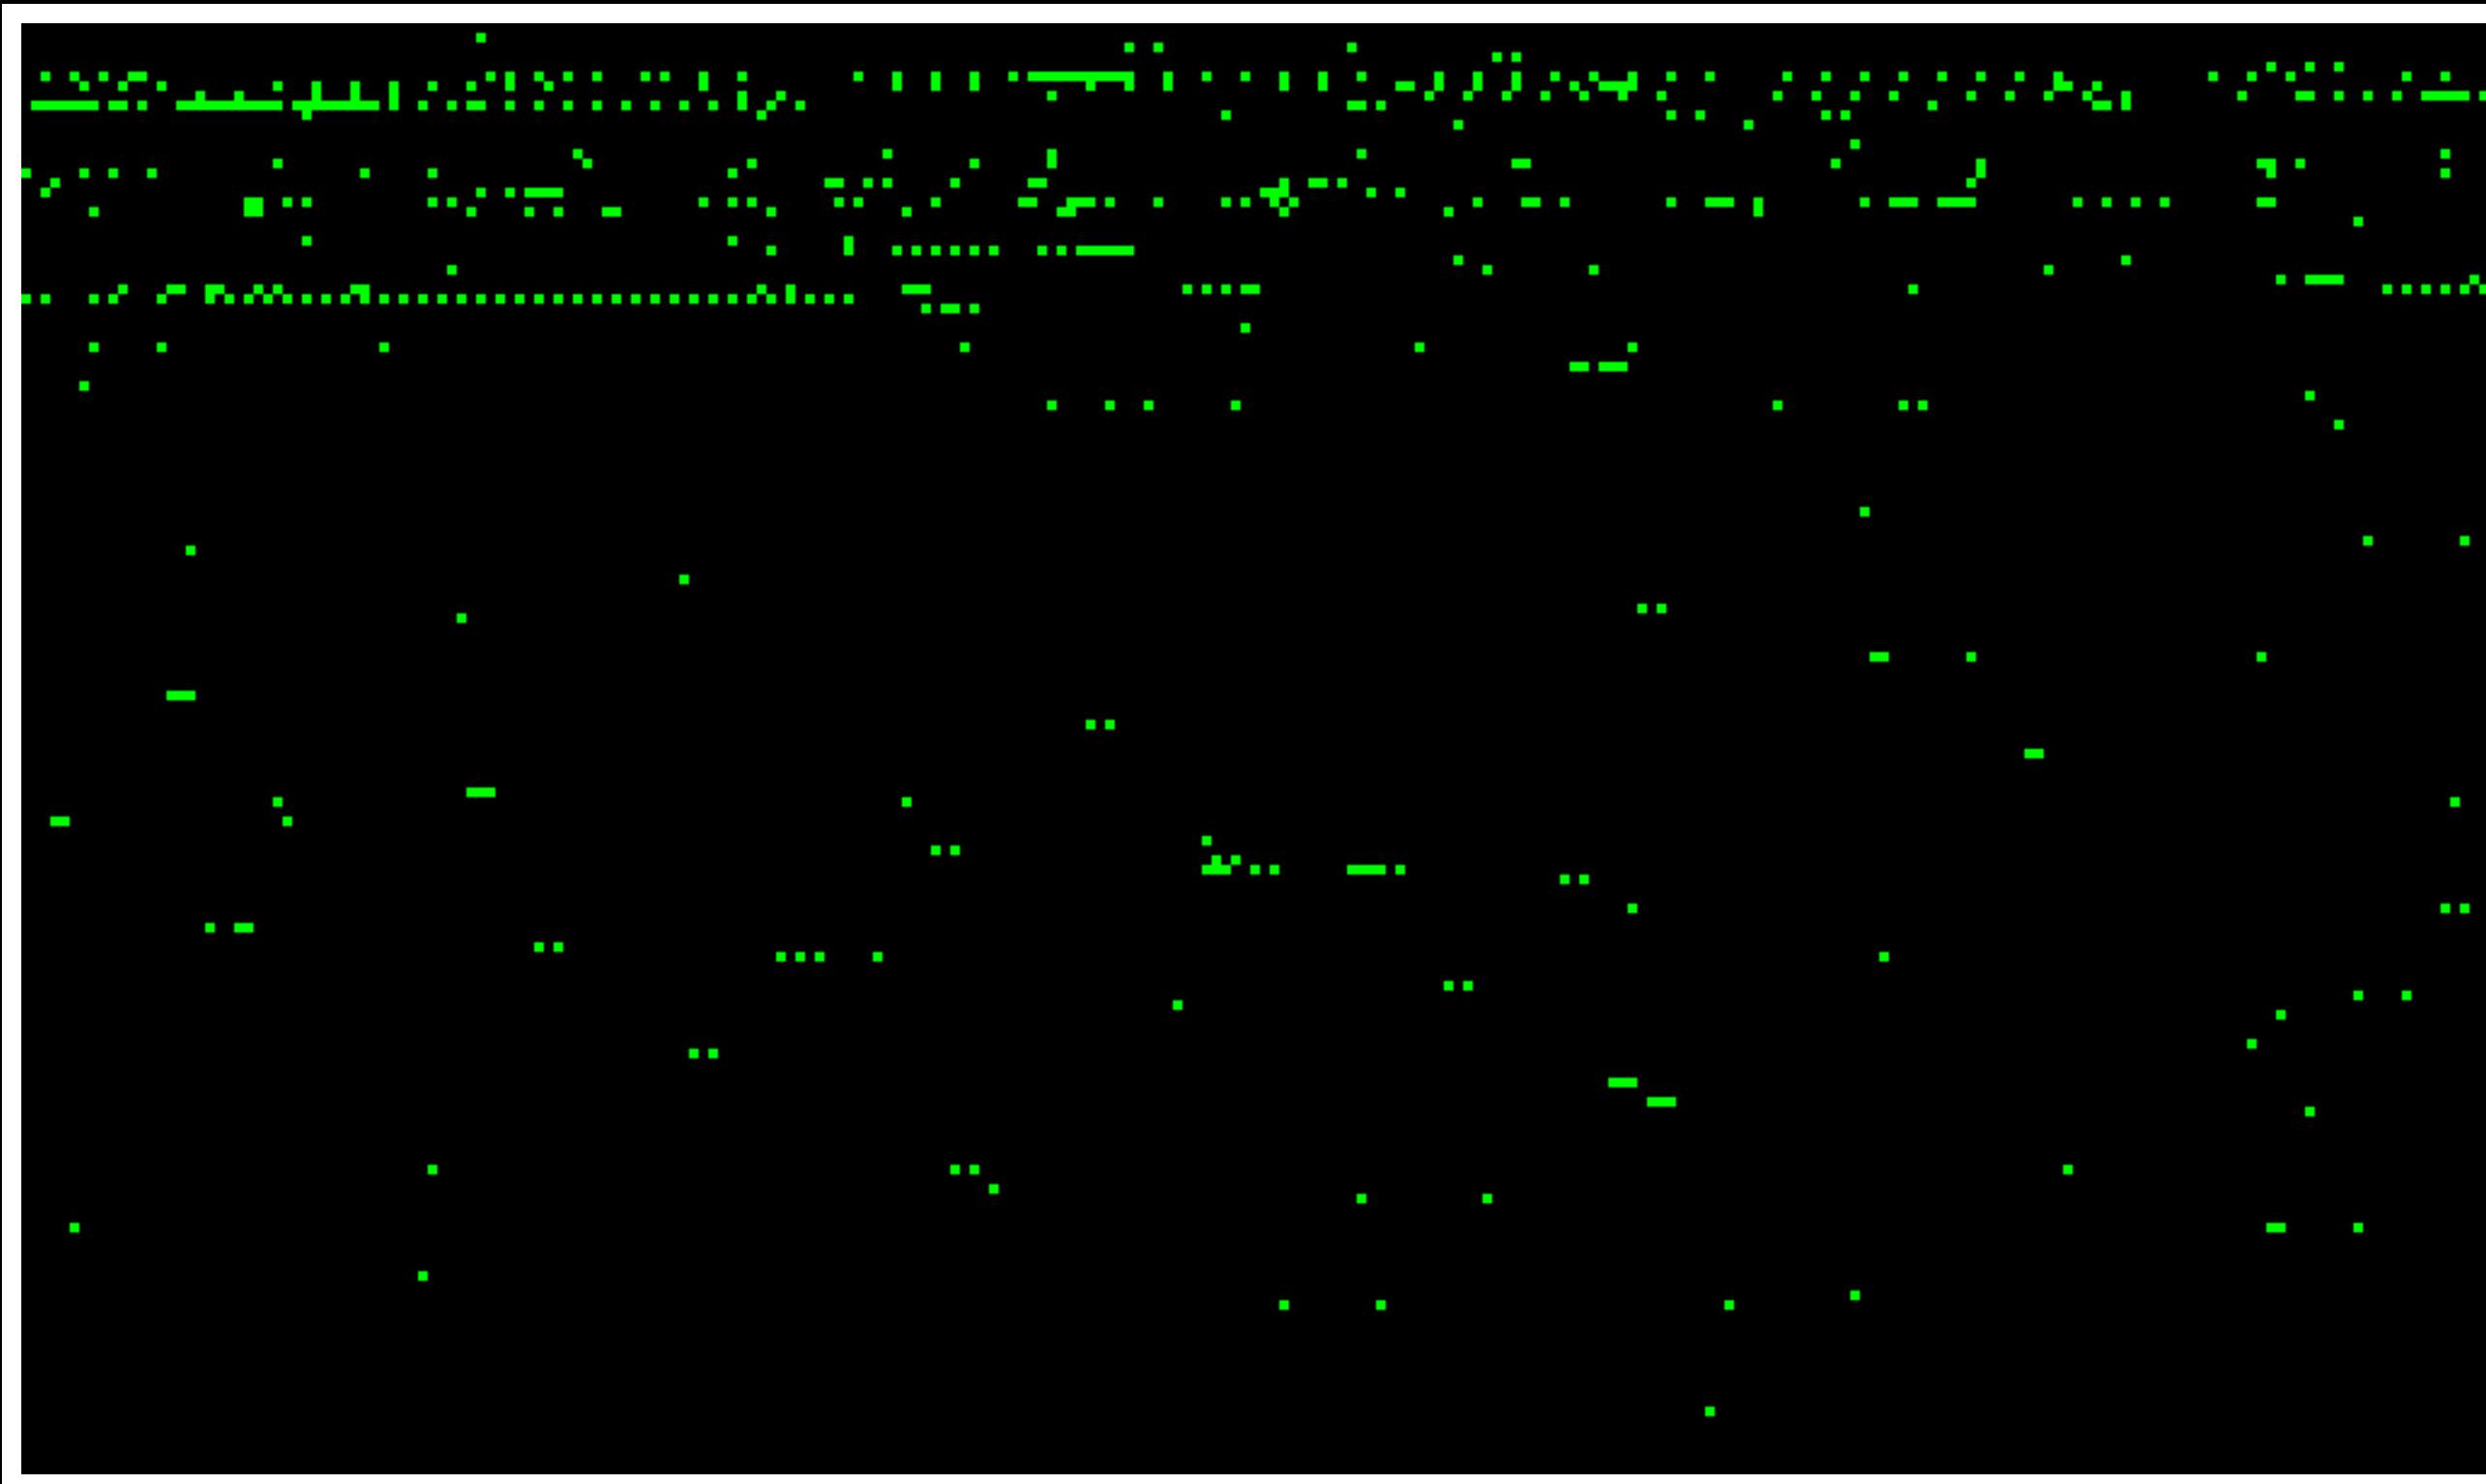
> covert channel



# Having fun with deduplication

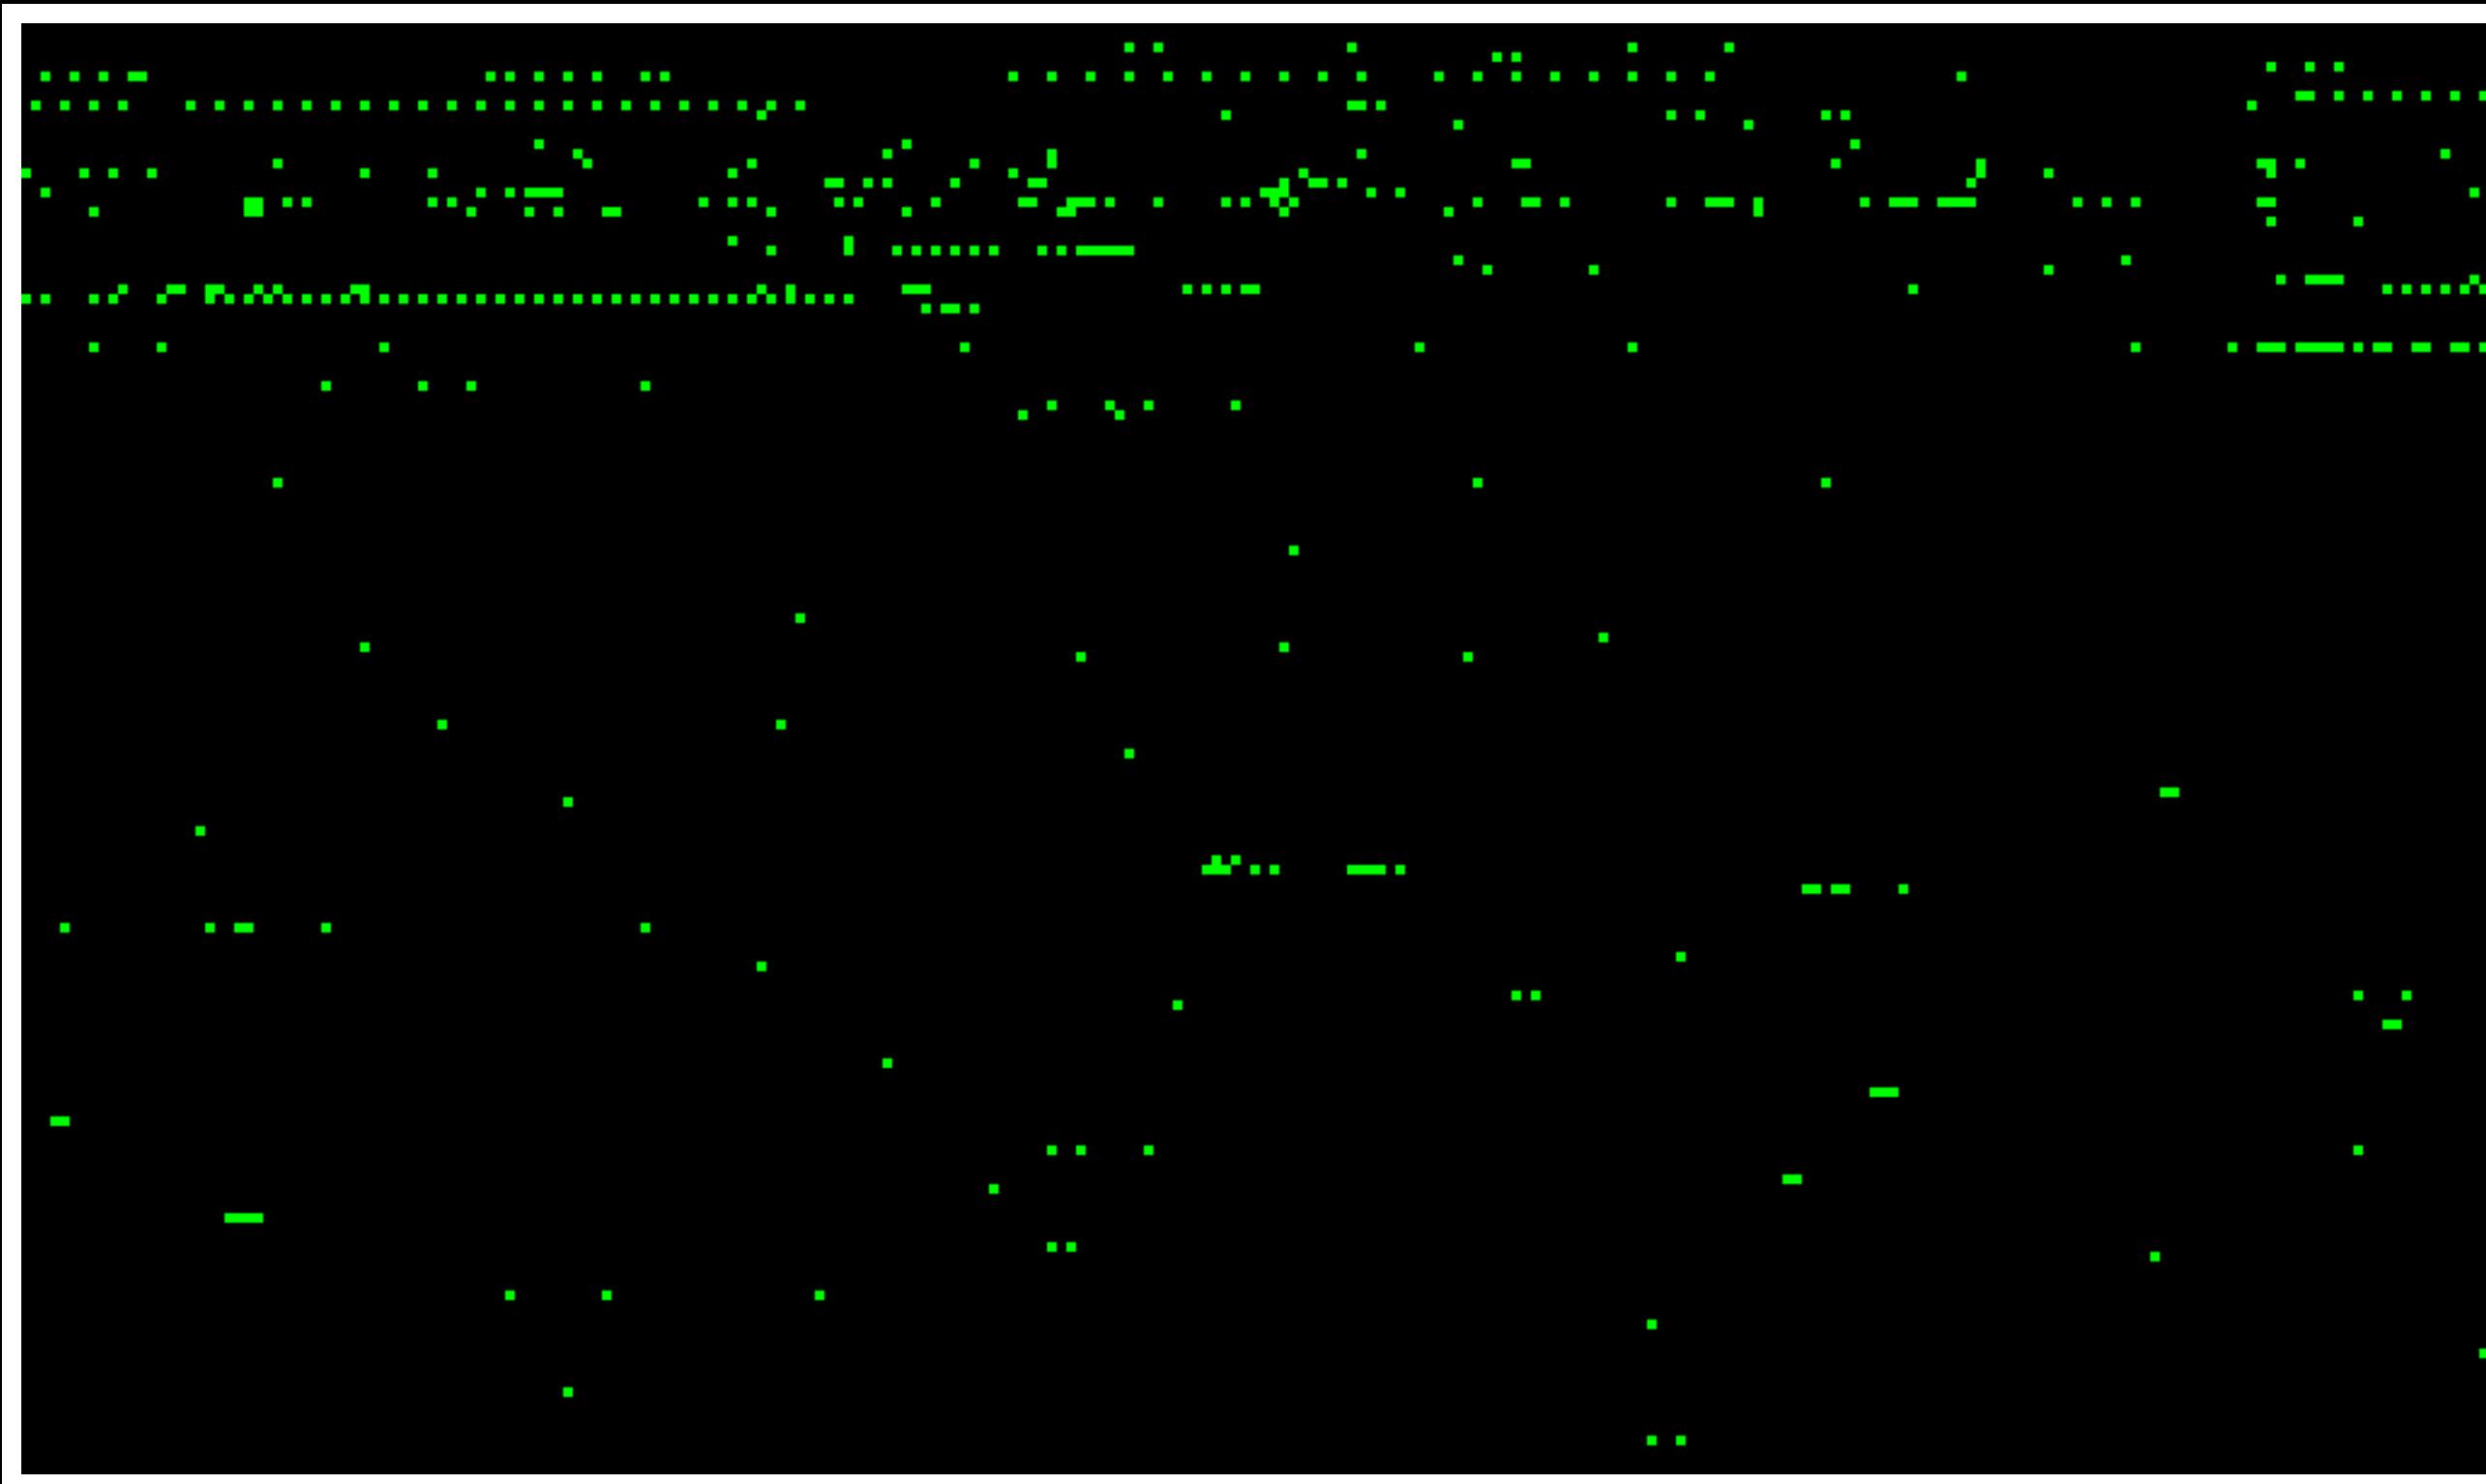
- > covert channel
- > detect running software

# Wordpad memory dump



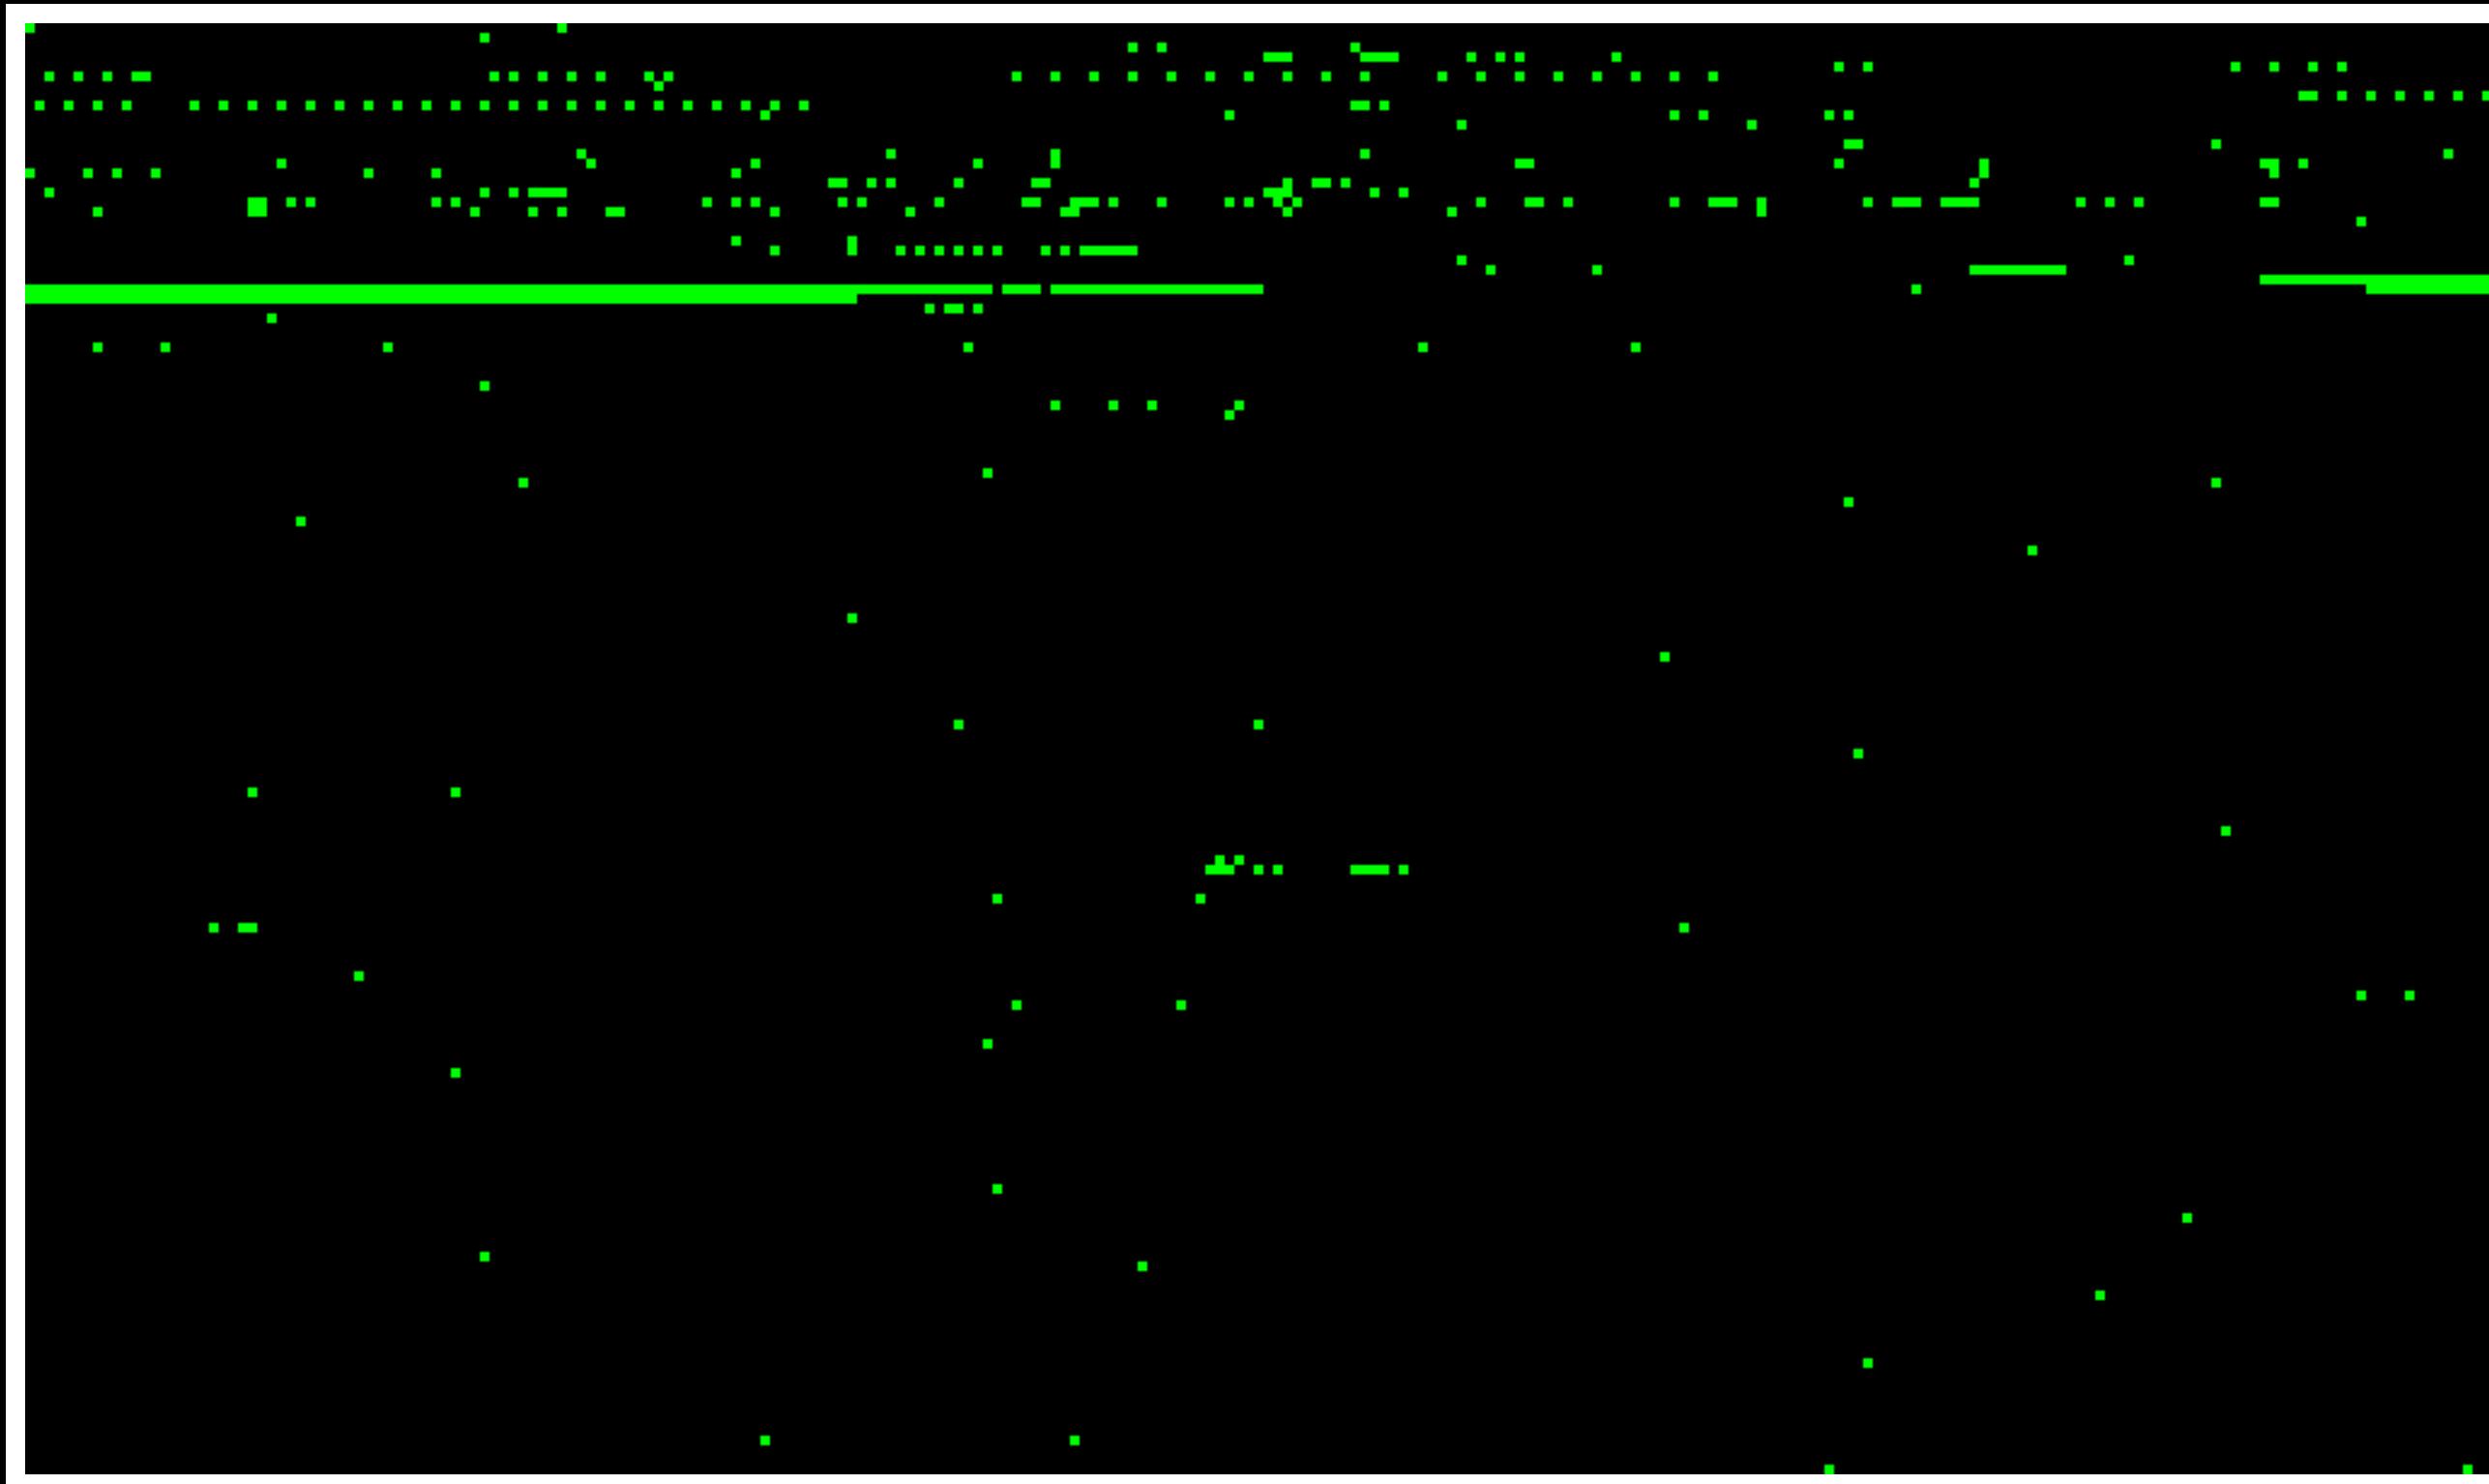
wordpad not running

# Wordpad memory dump



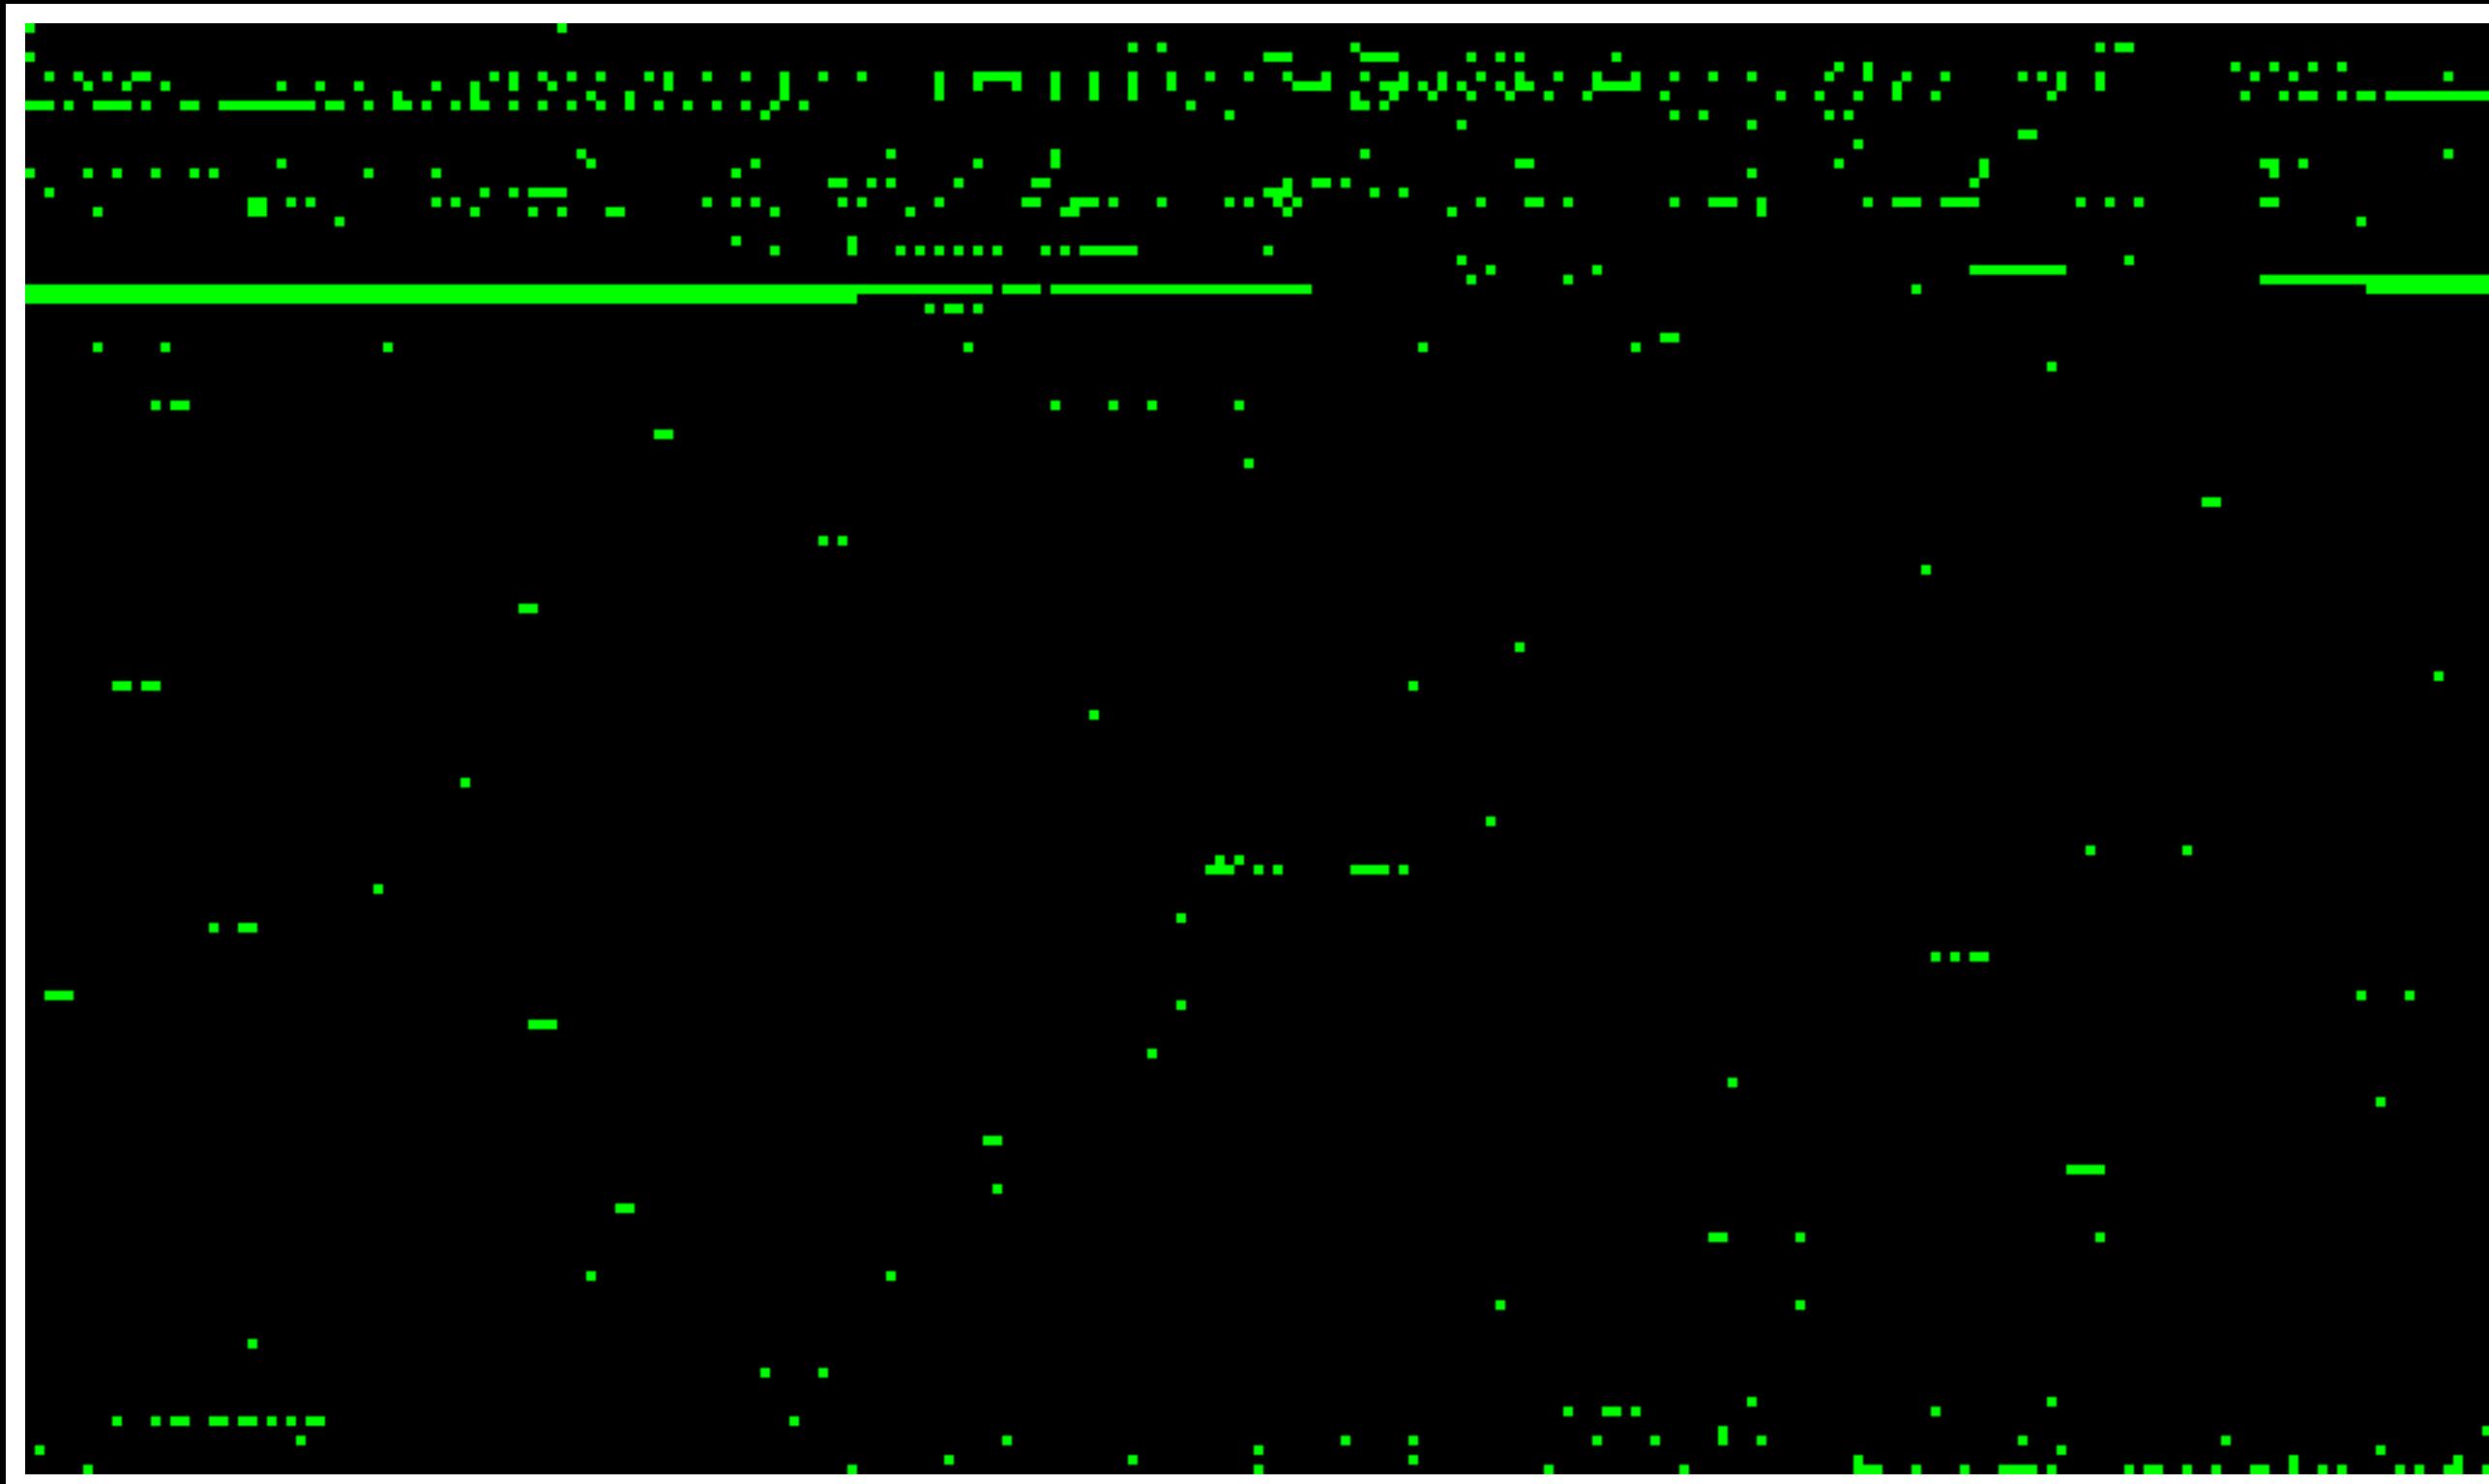
wordpad not running

# Wordpad memory dump



wordpad running

# Wordpad memory dump



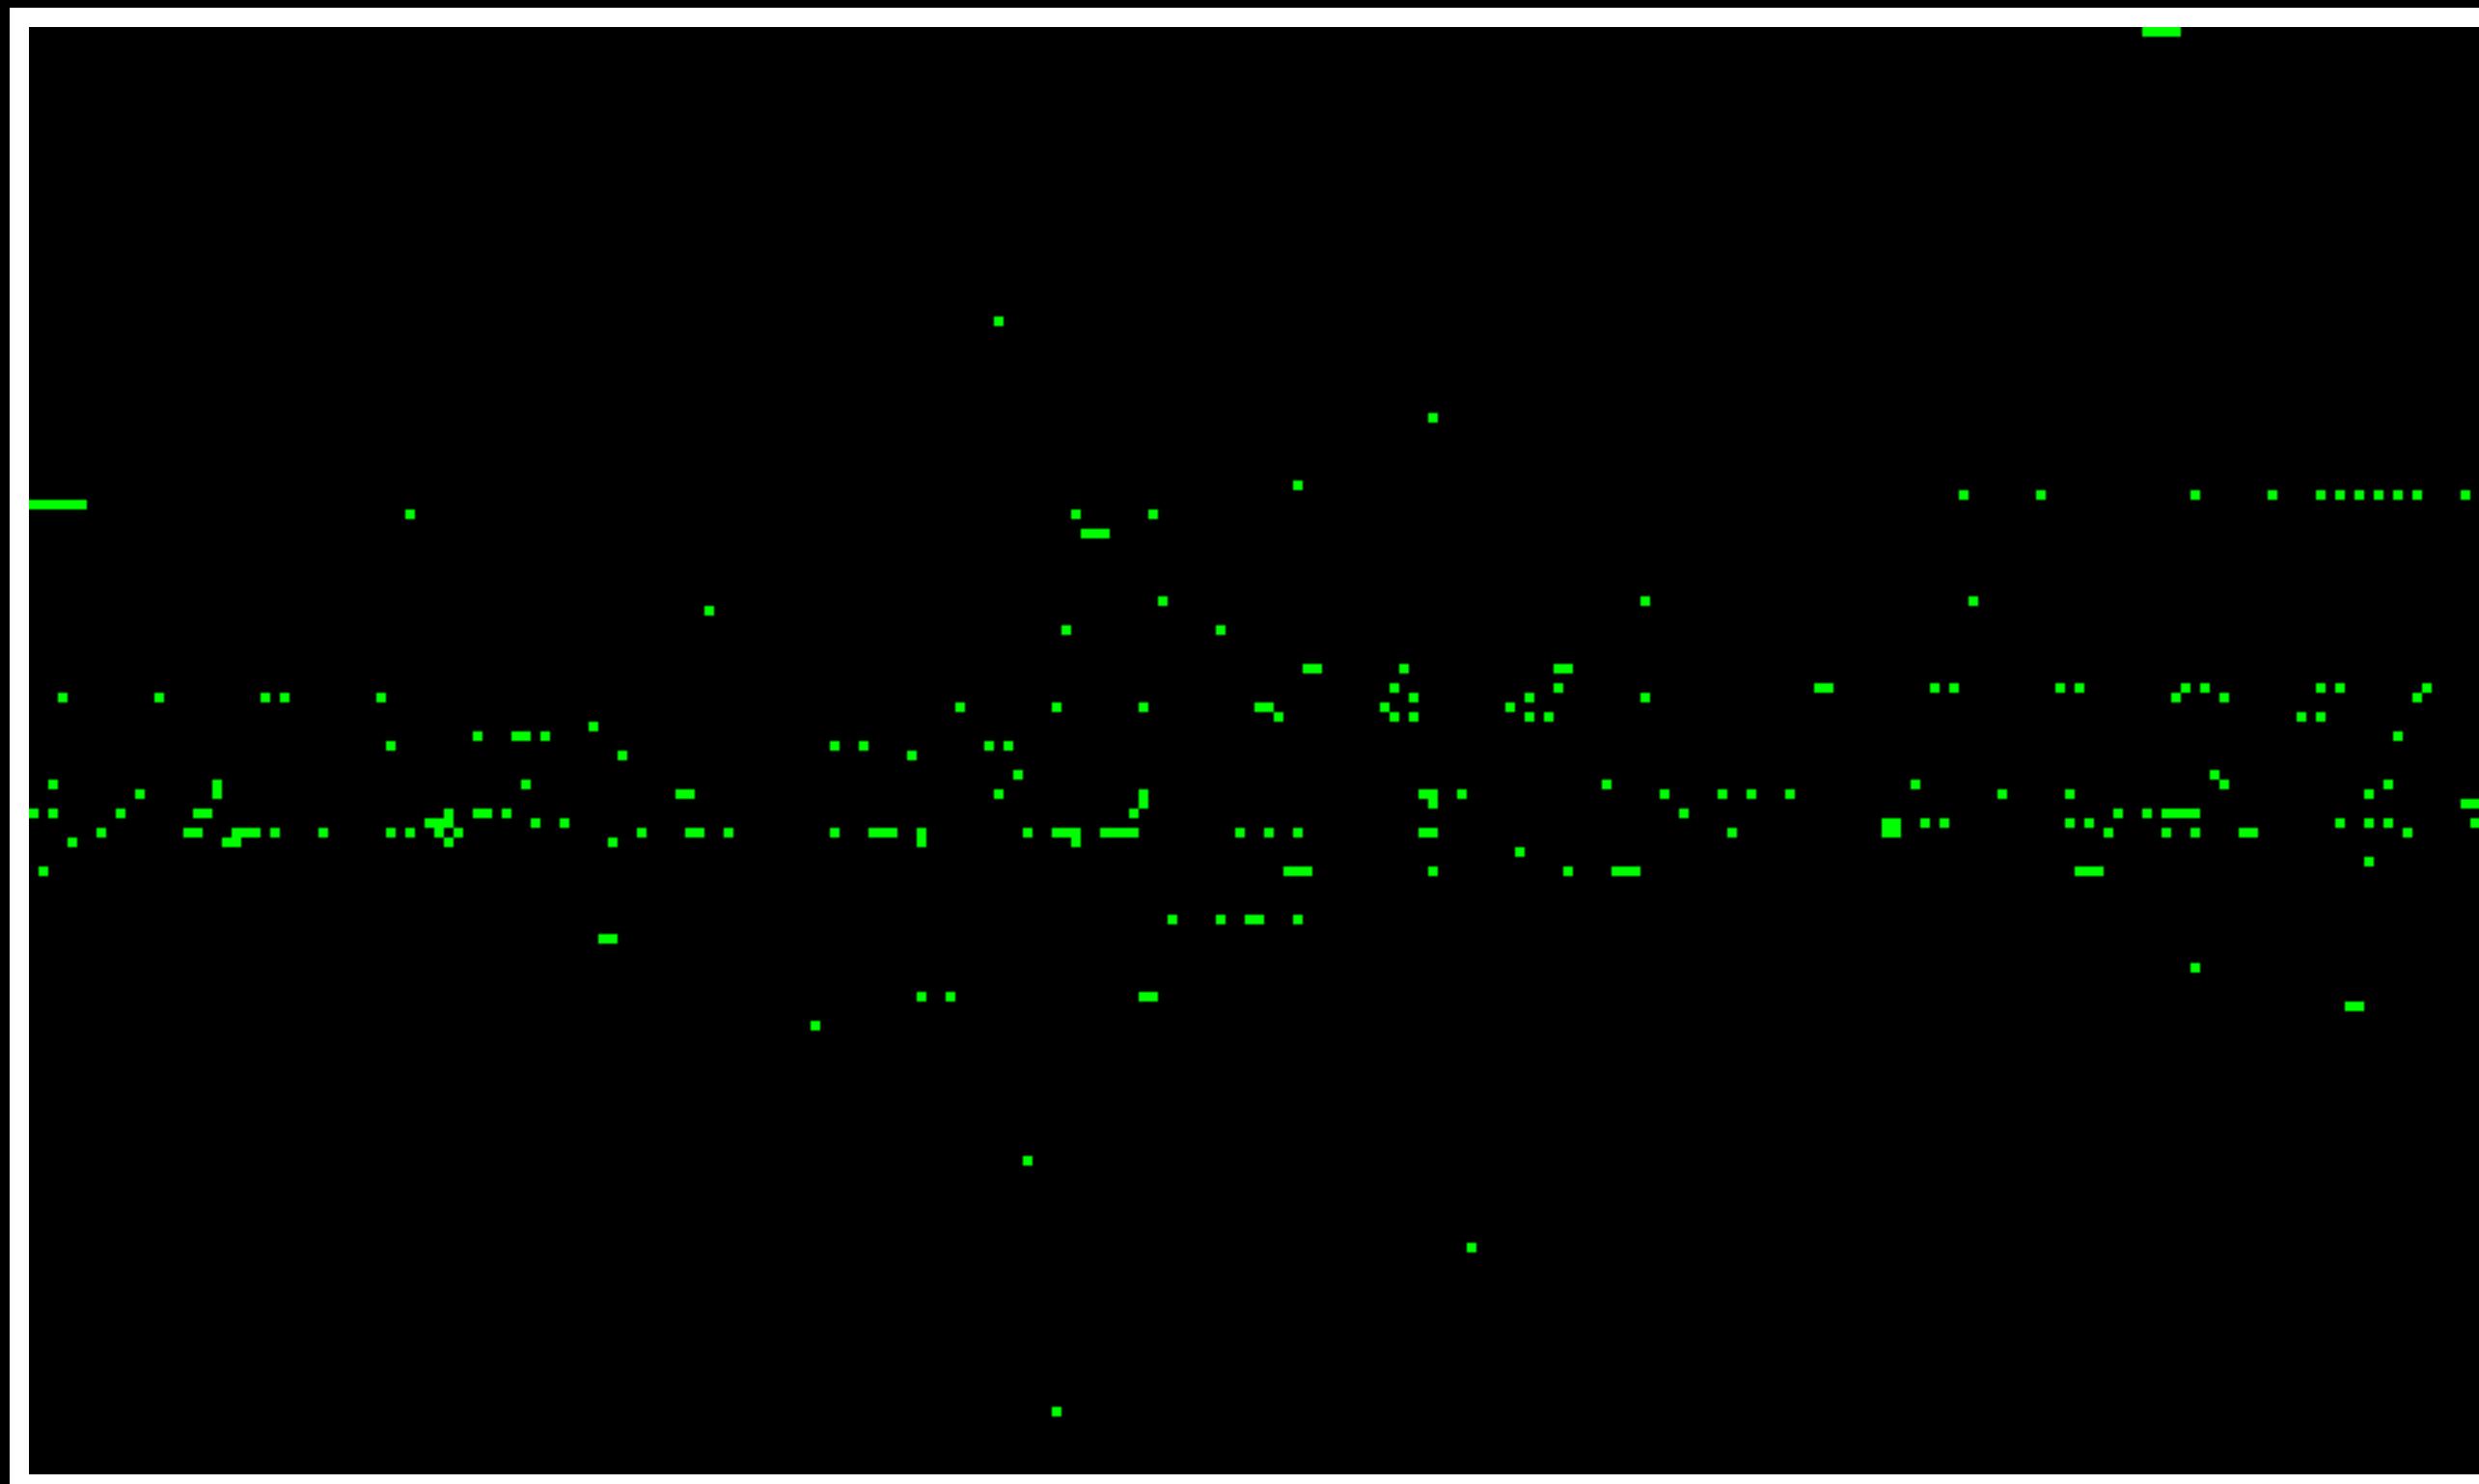
wordpad running

**Signal not as clear as expected,**

**Reason:**

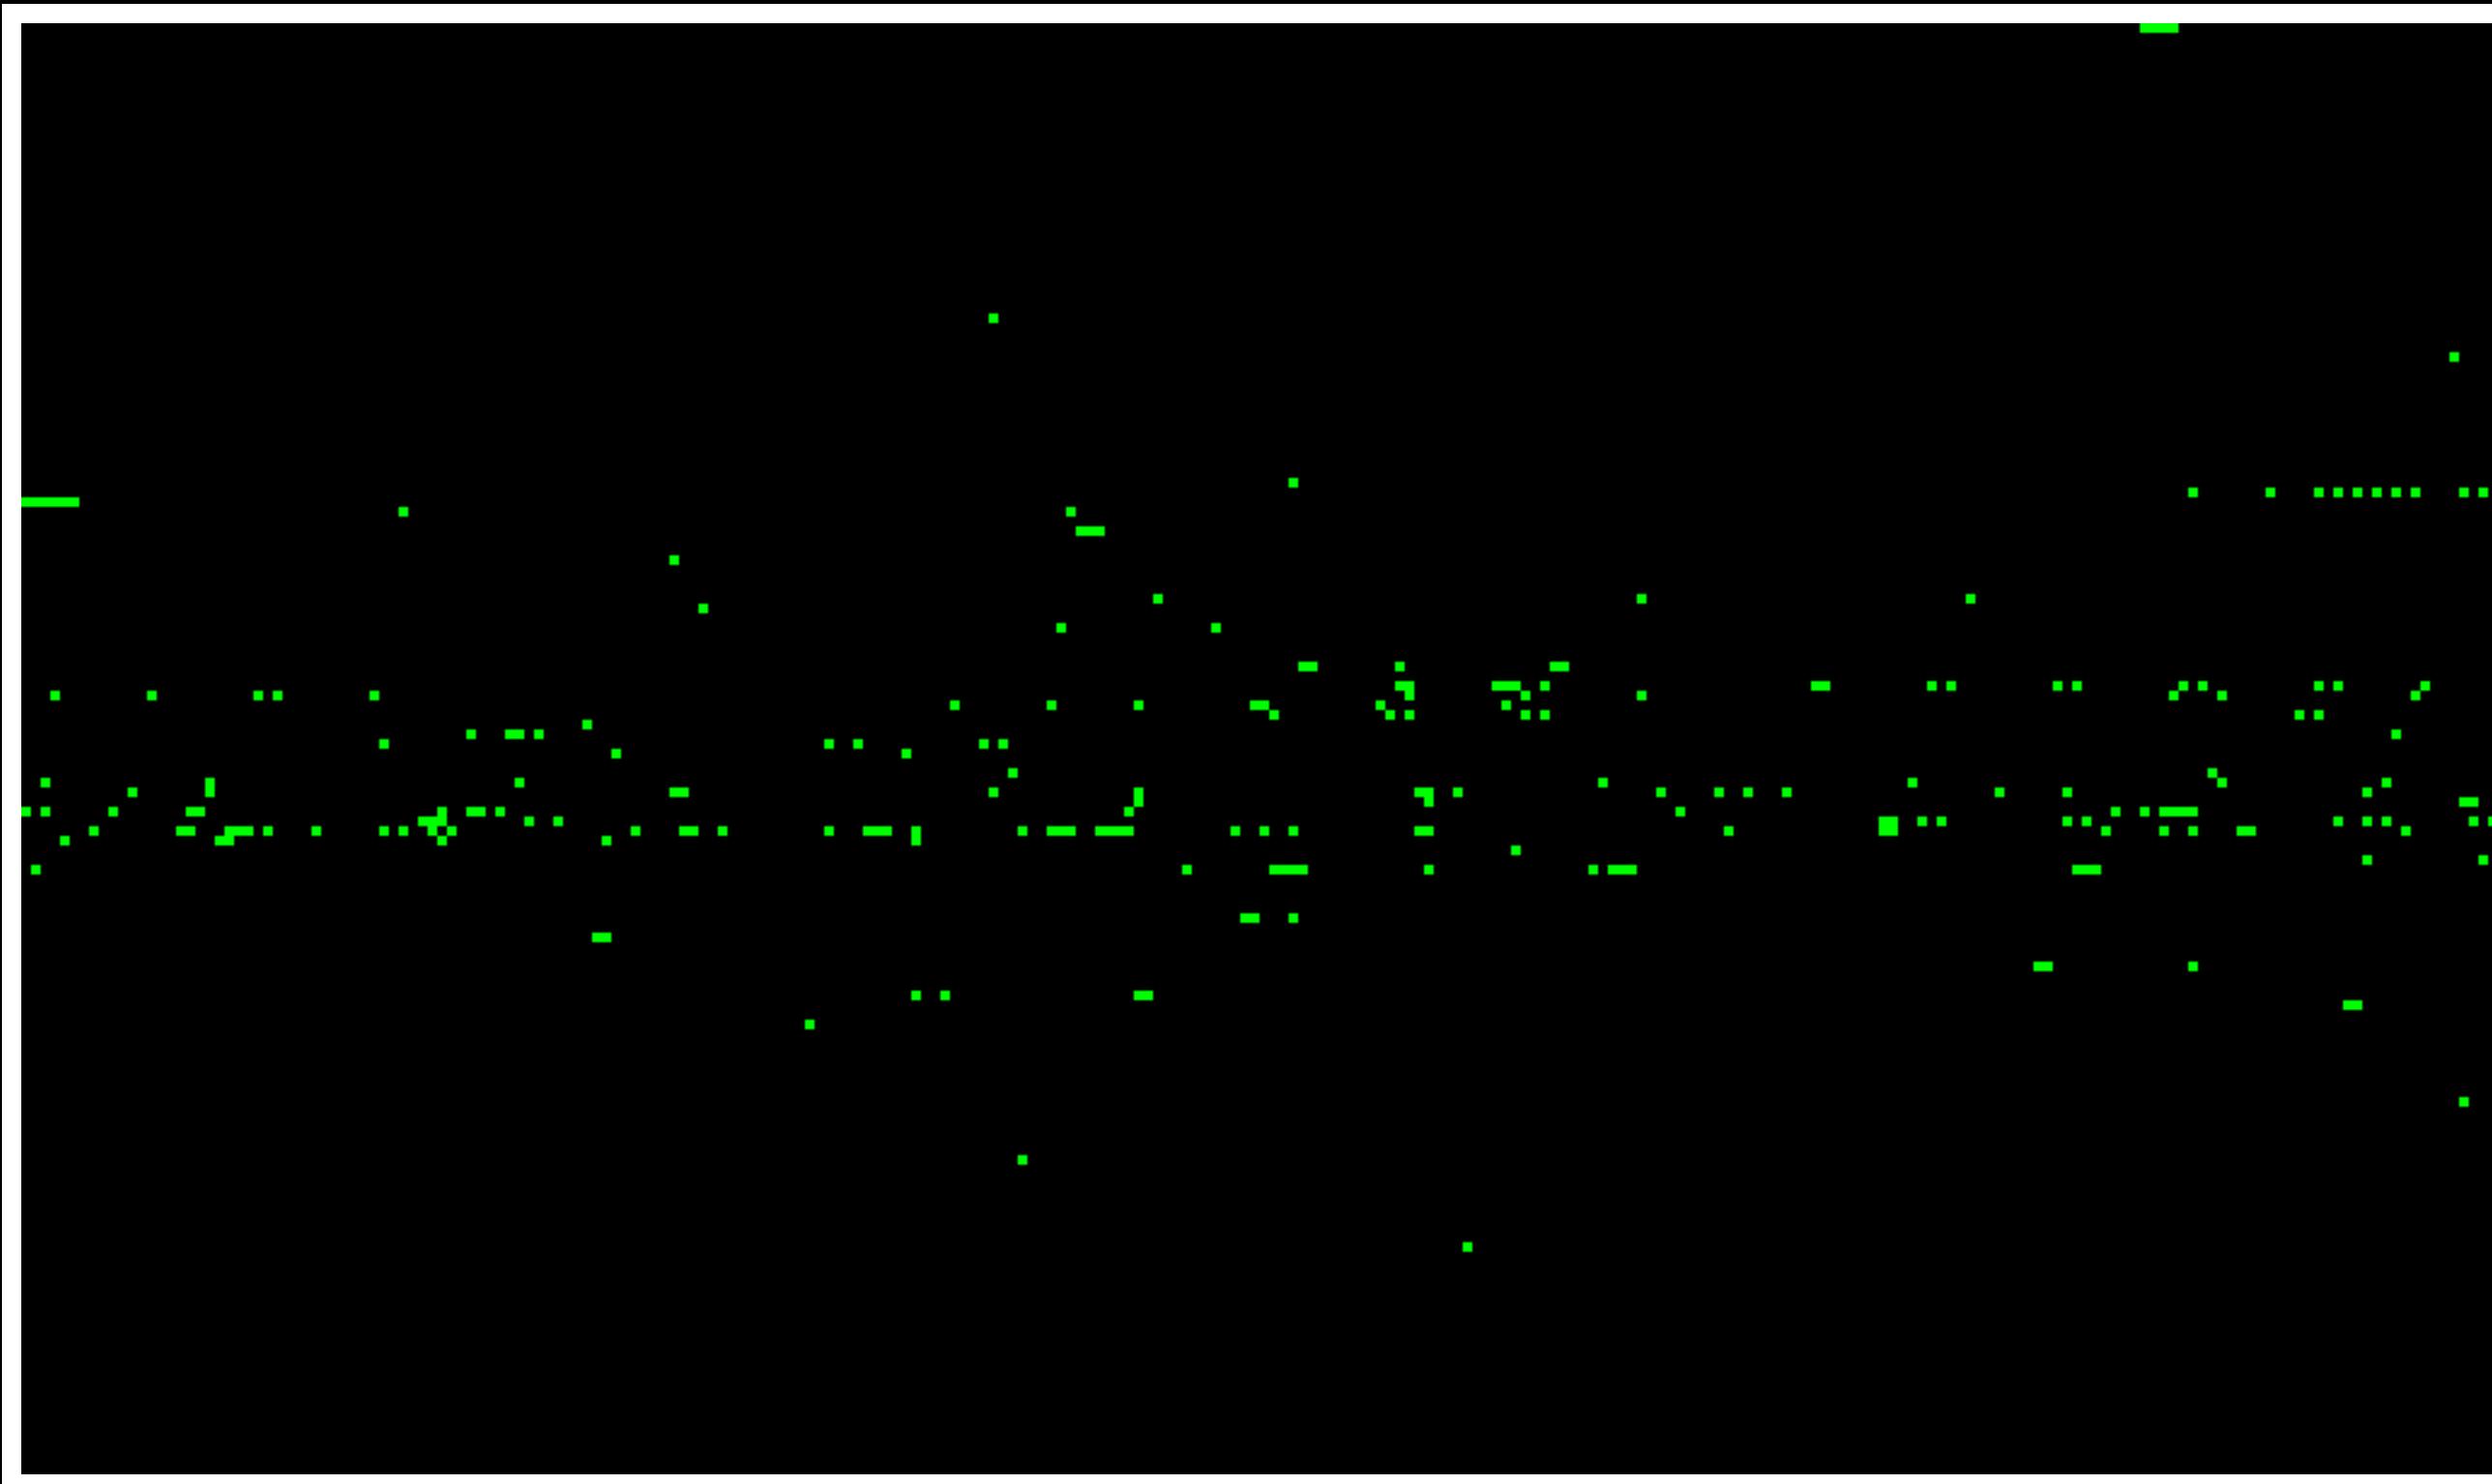
**file backed memory  
not deduplicated the same way  
on Windows .**

# Skype memory dump



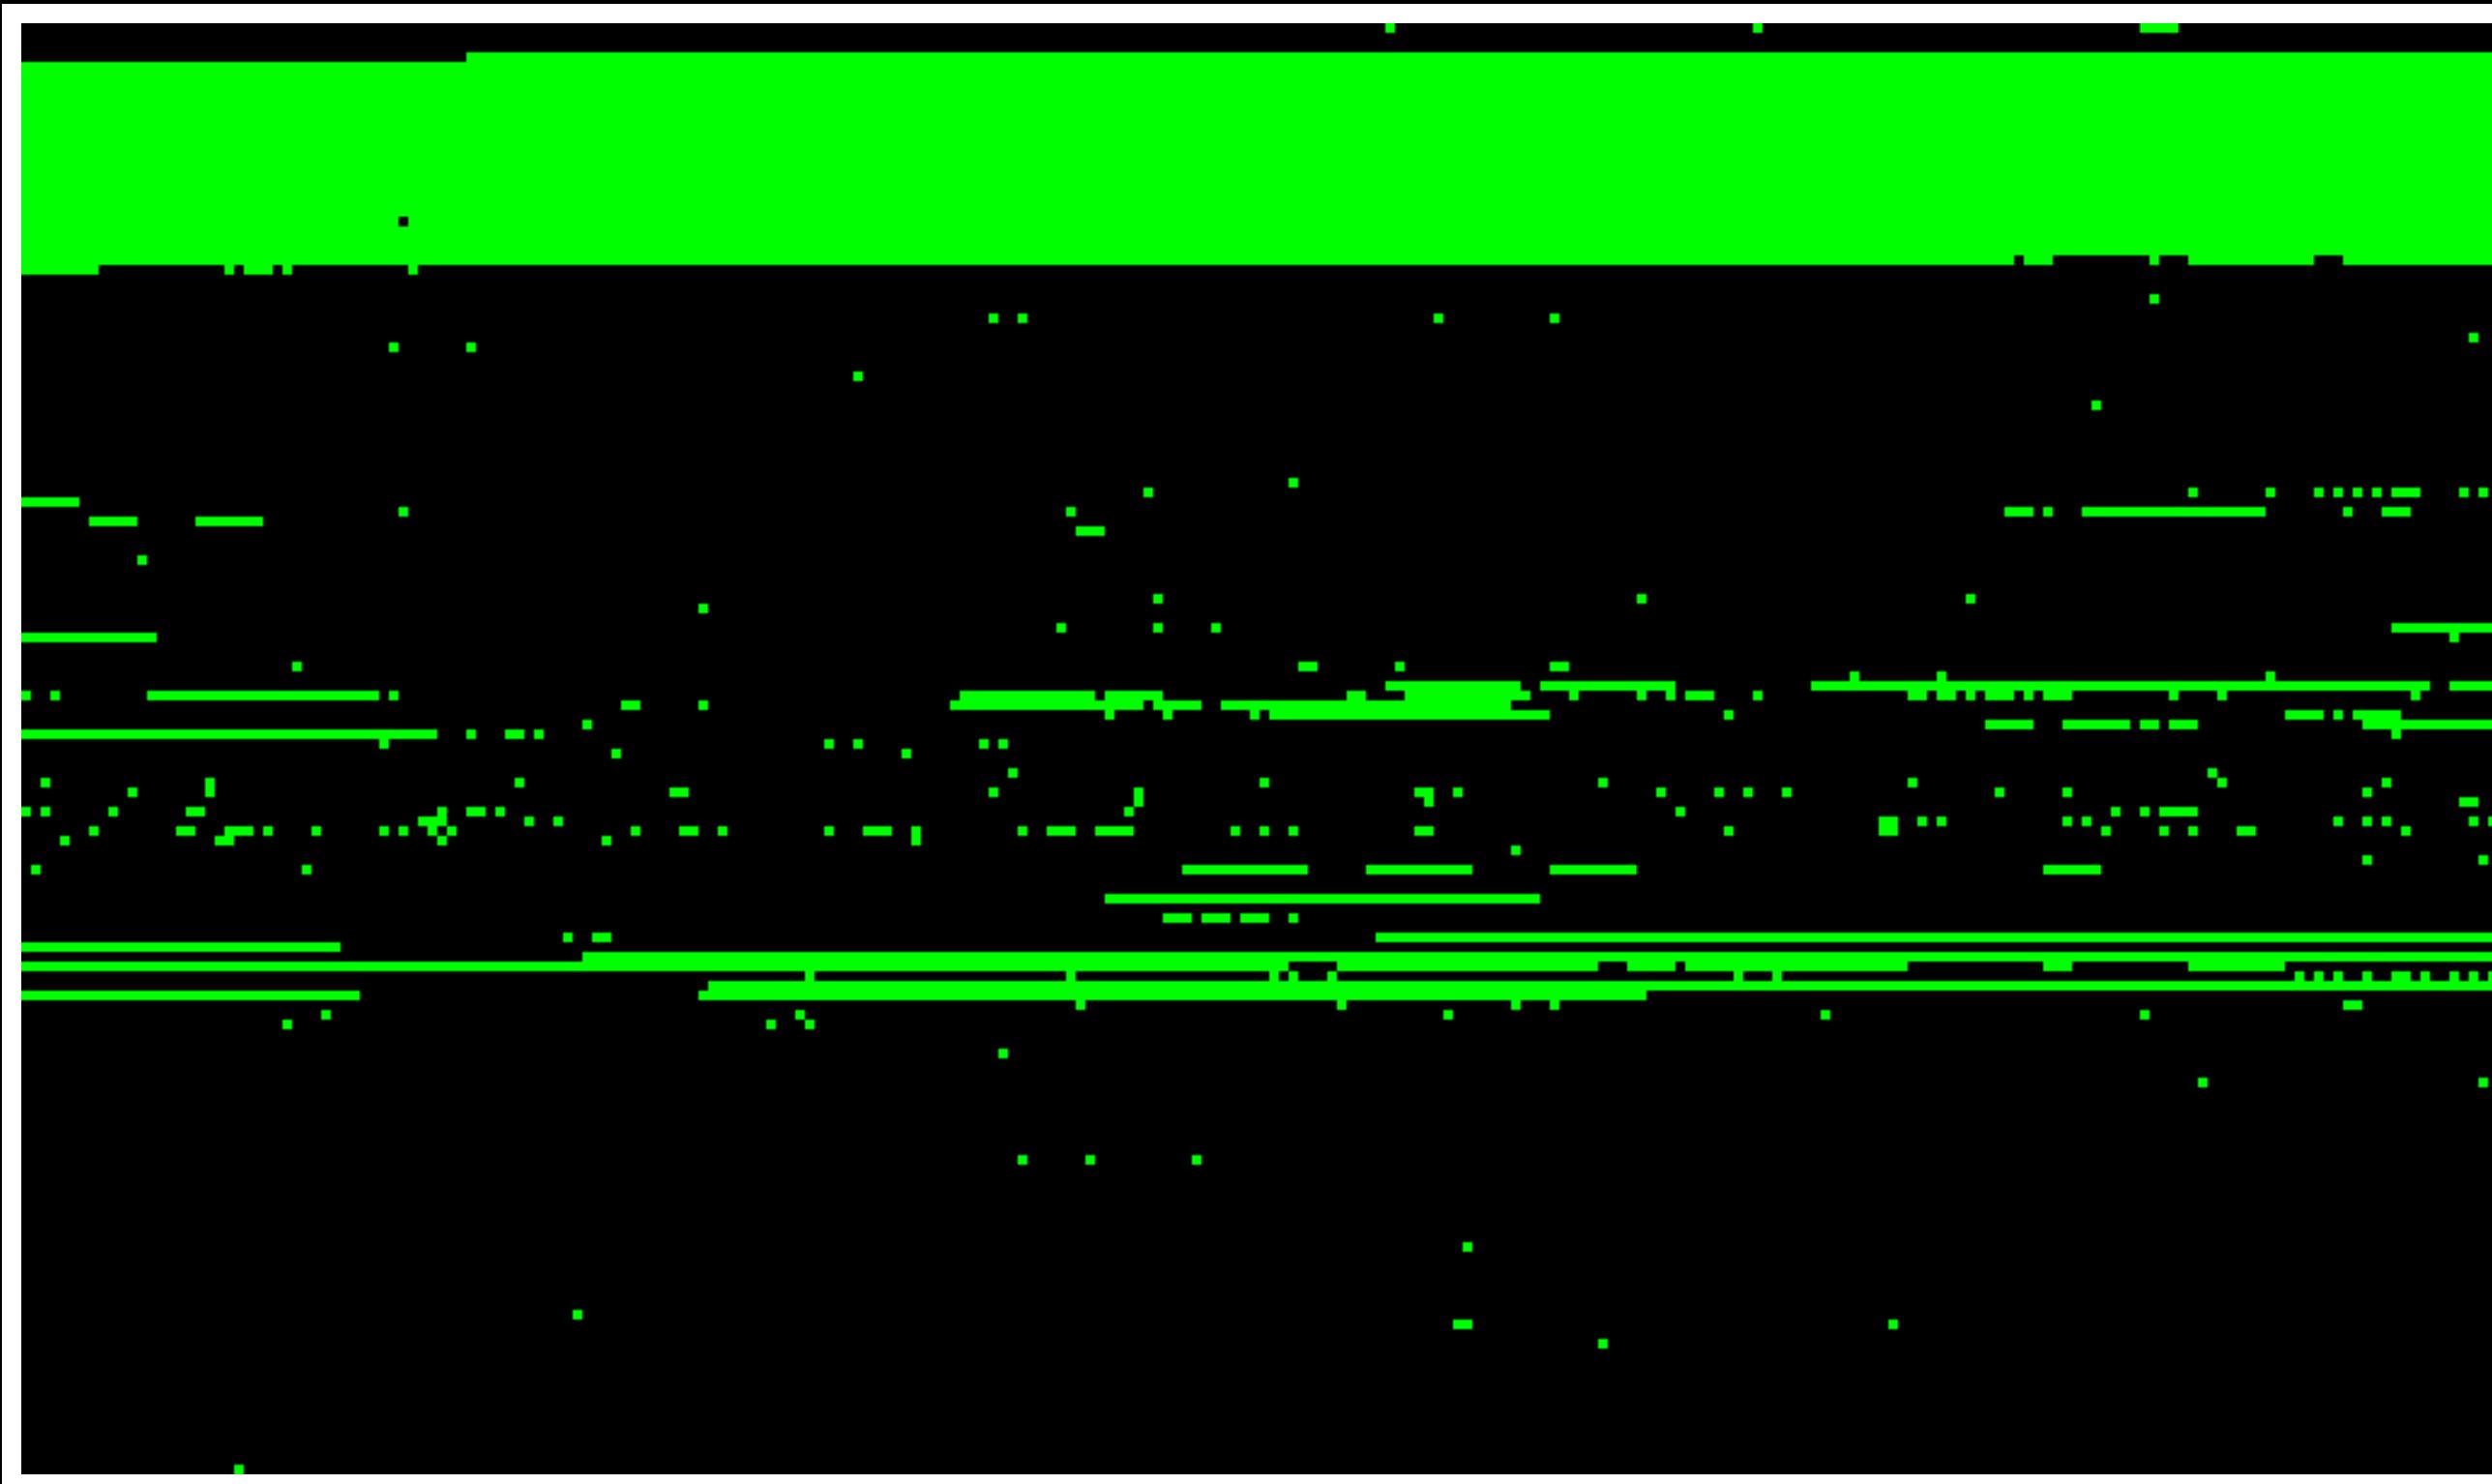
skype not running

# Skype memory dump



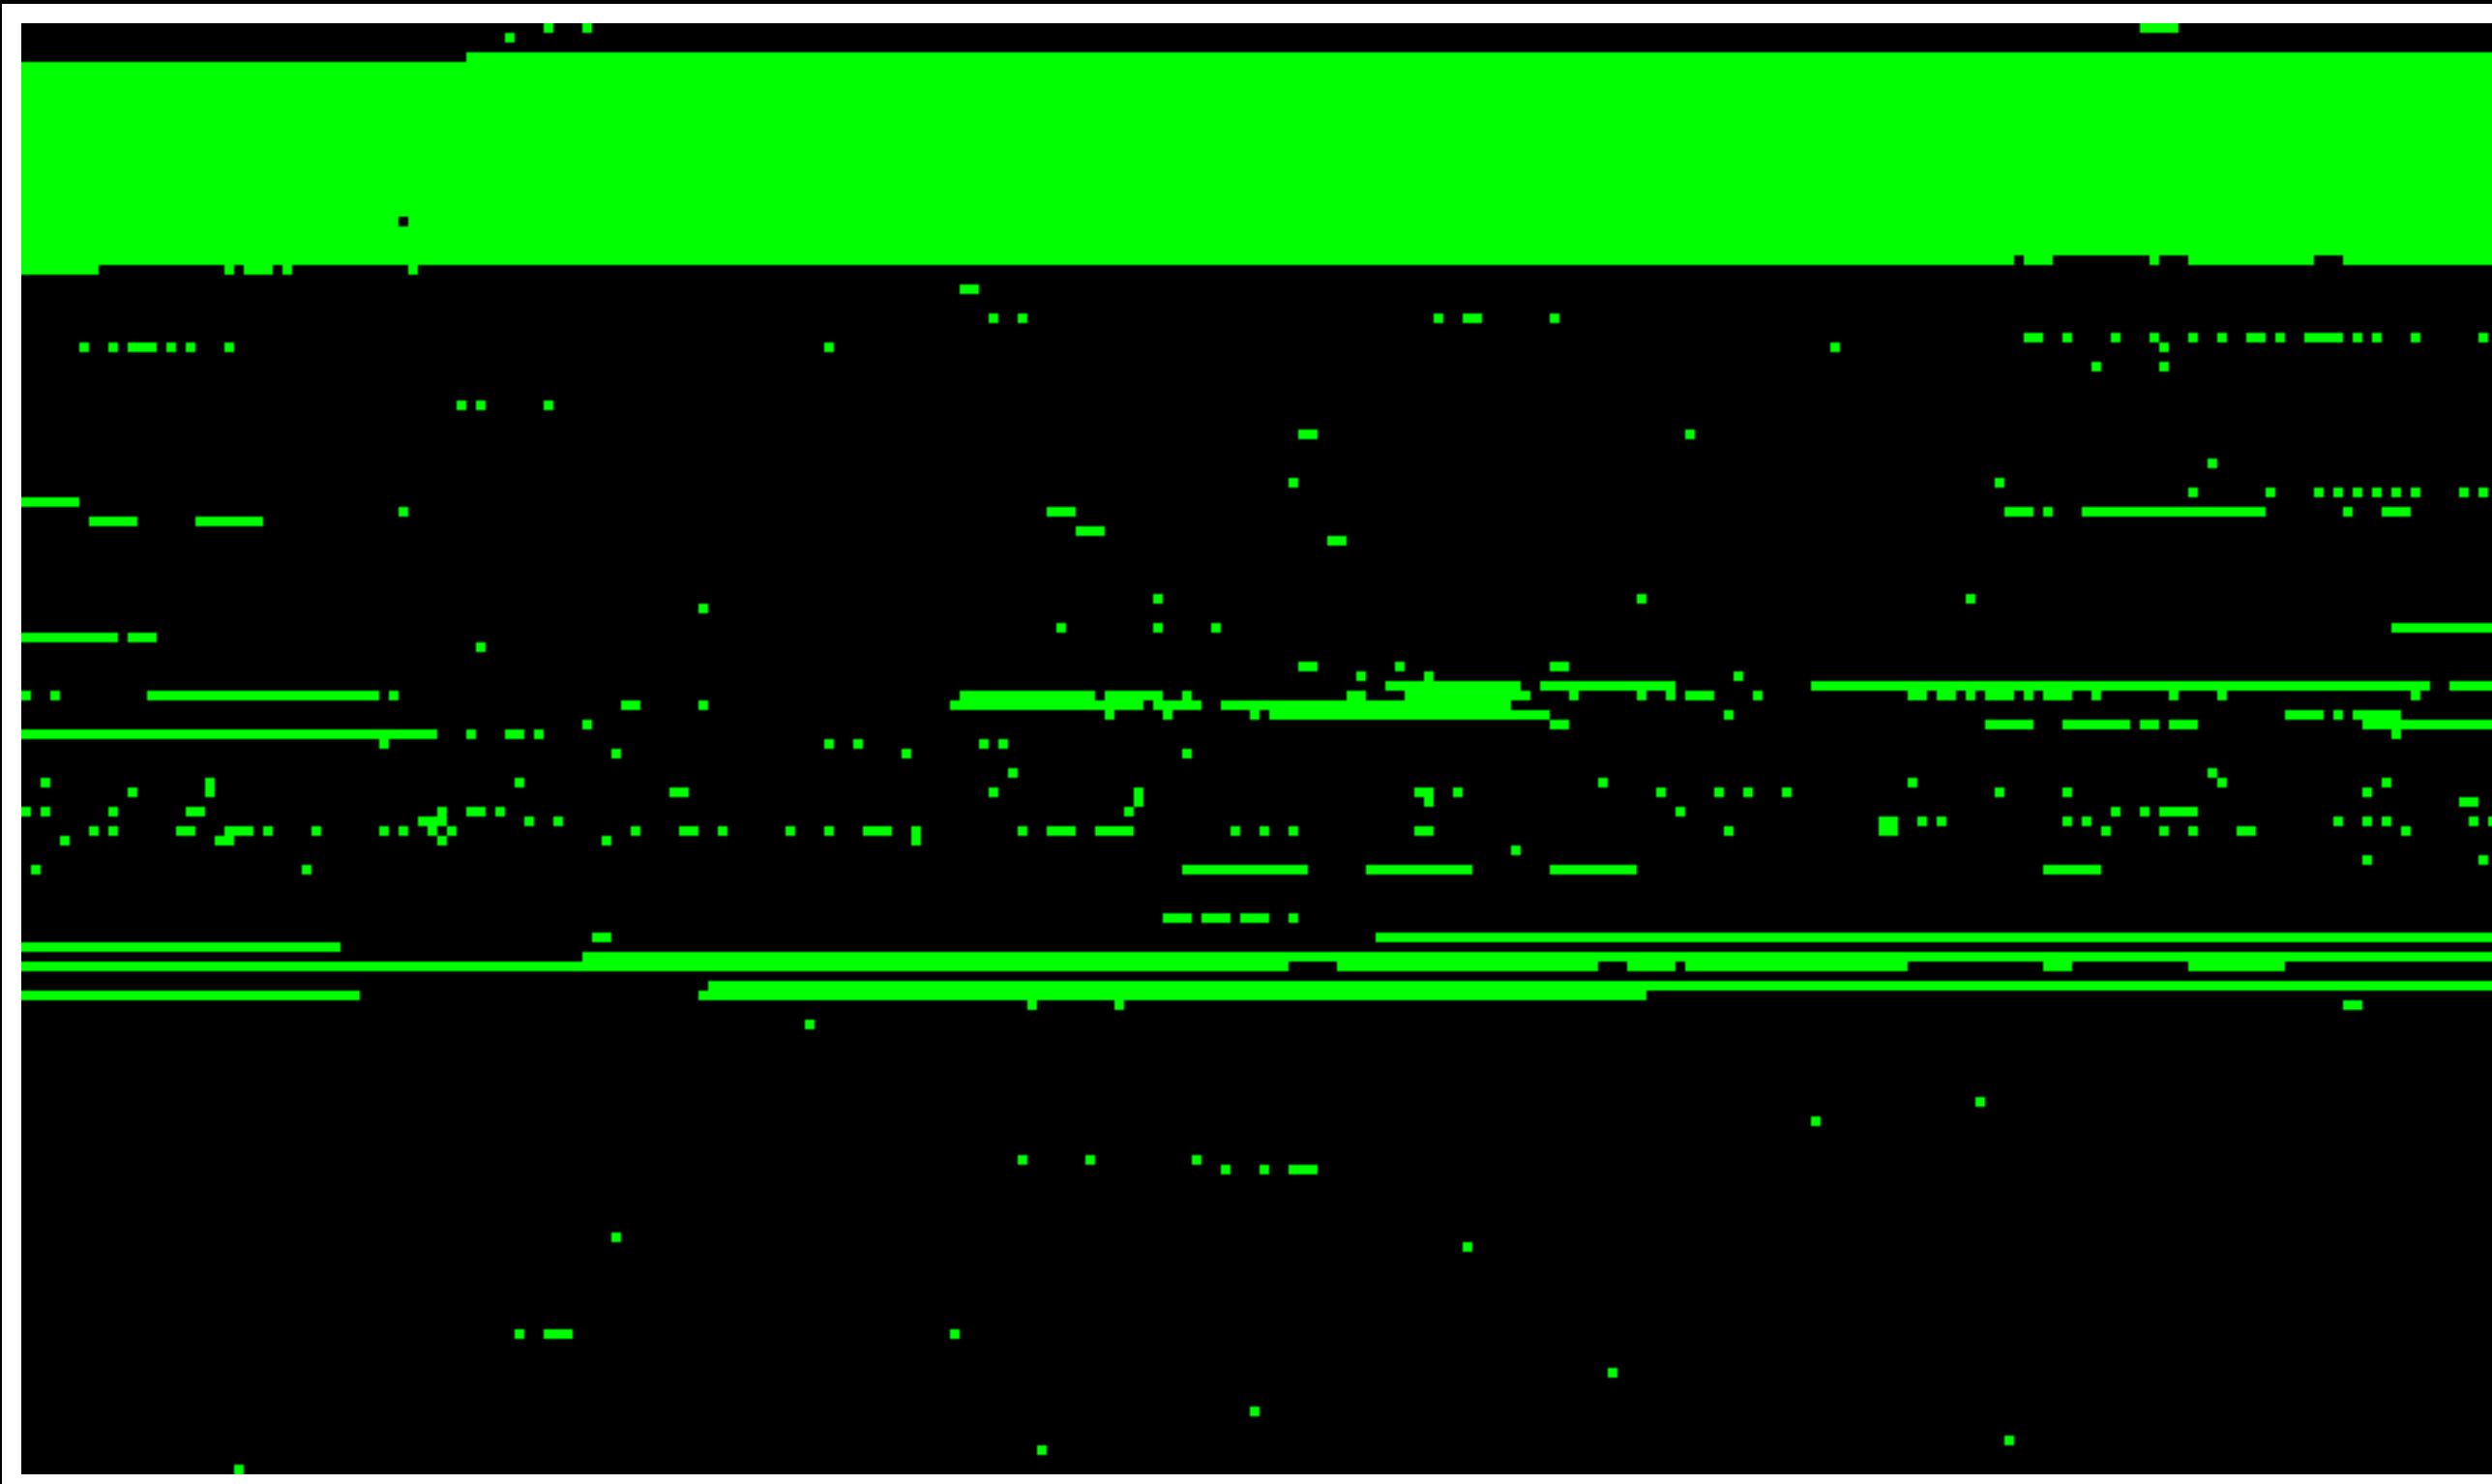
skype not running

# Skype memory dump



skype running

# Skype memory dump



skype running

For our Edge exploit,  
a single-bit, page-granularity  
info leak isn't enough

Can we generalize this to leaking arbitrary data, like an ASLR pointer or a password?

Leaking existing pages is slow and the gained information is limited.

What if we can manipulate the contents of the victim's memory to leak secrets hand-picked by the attacker.

# Challenge 1:

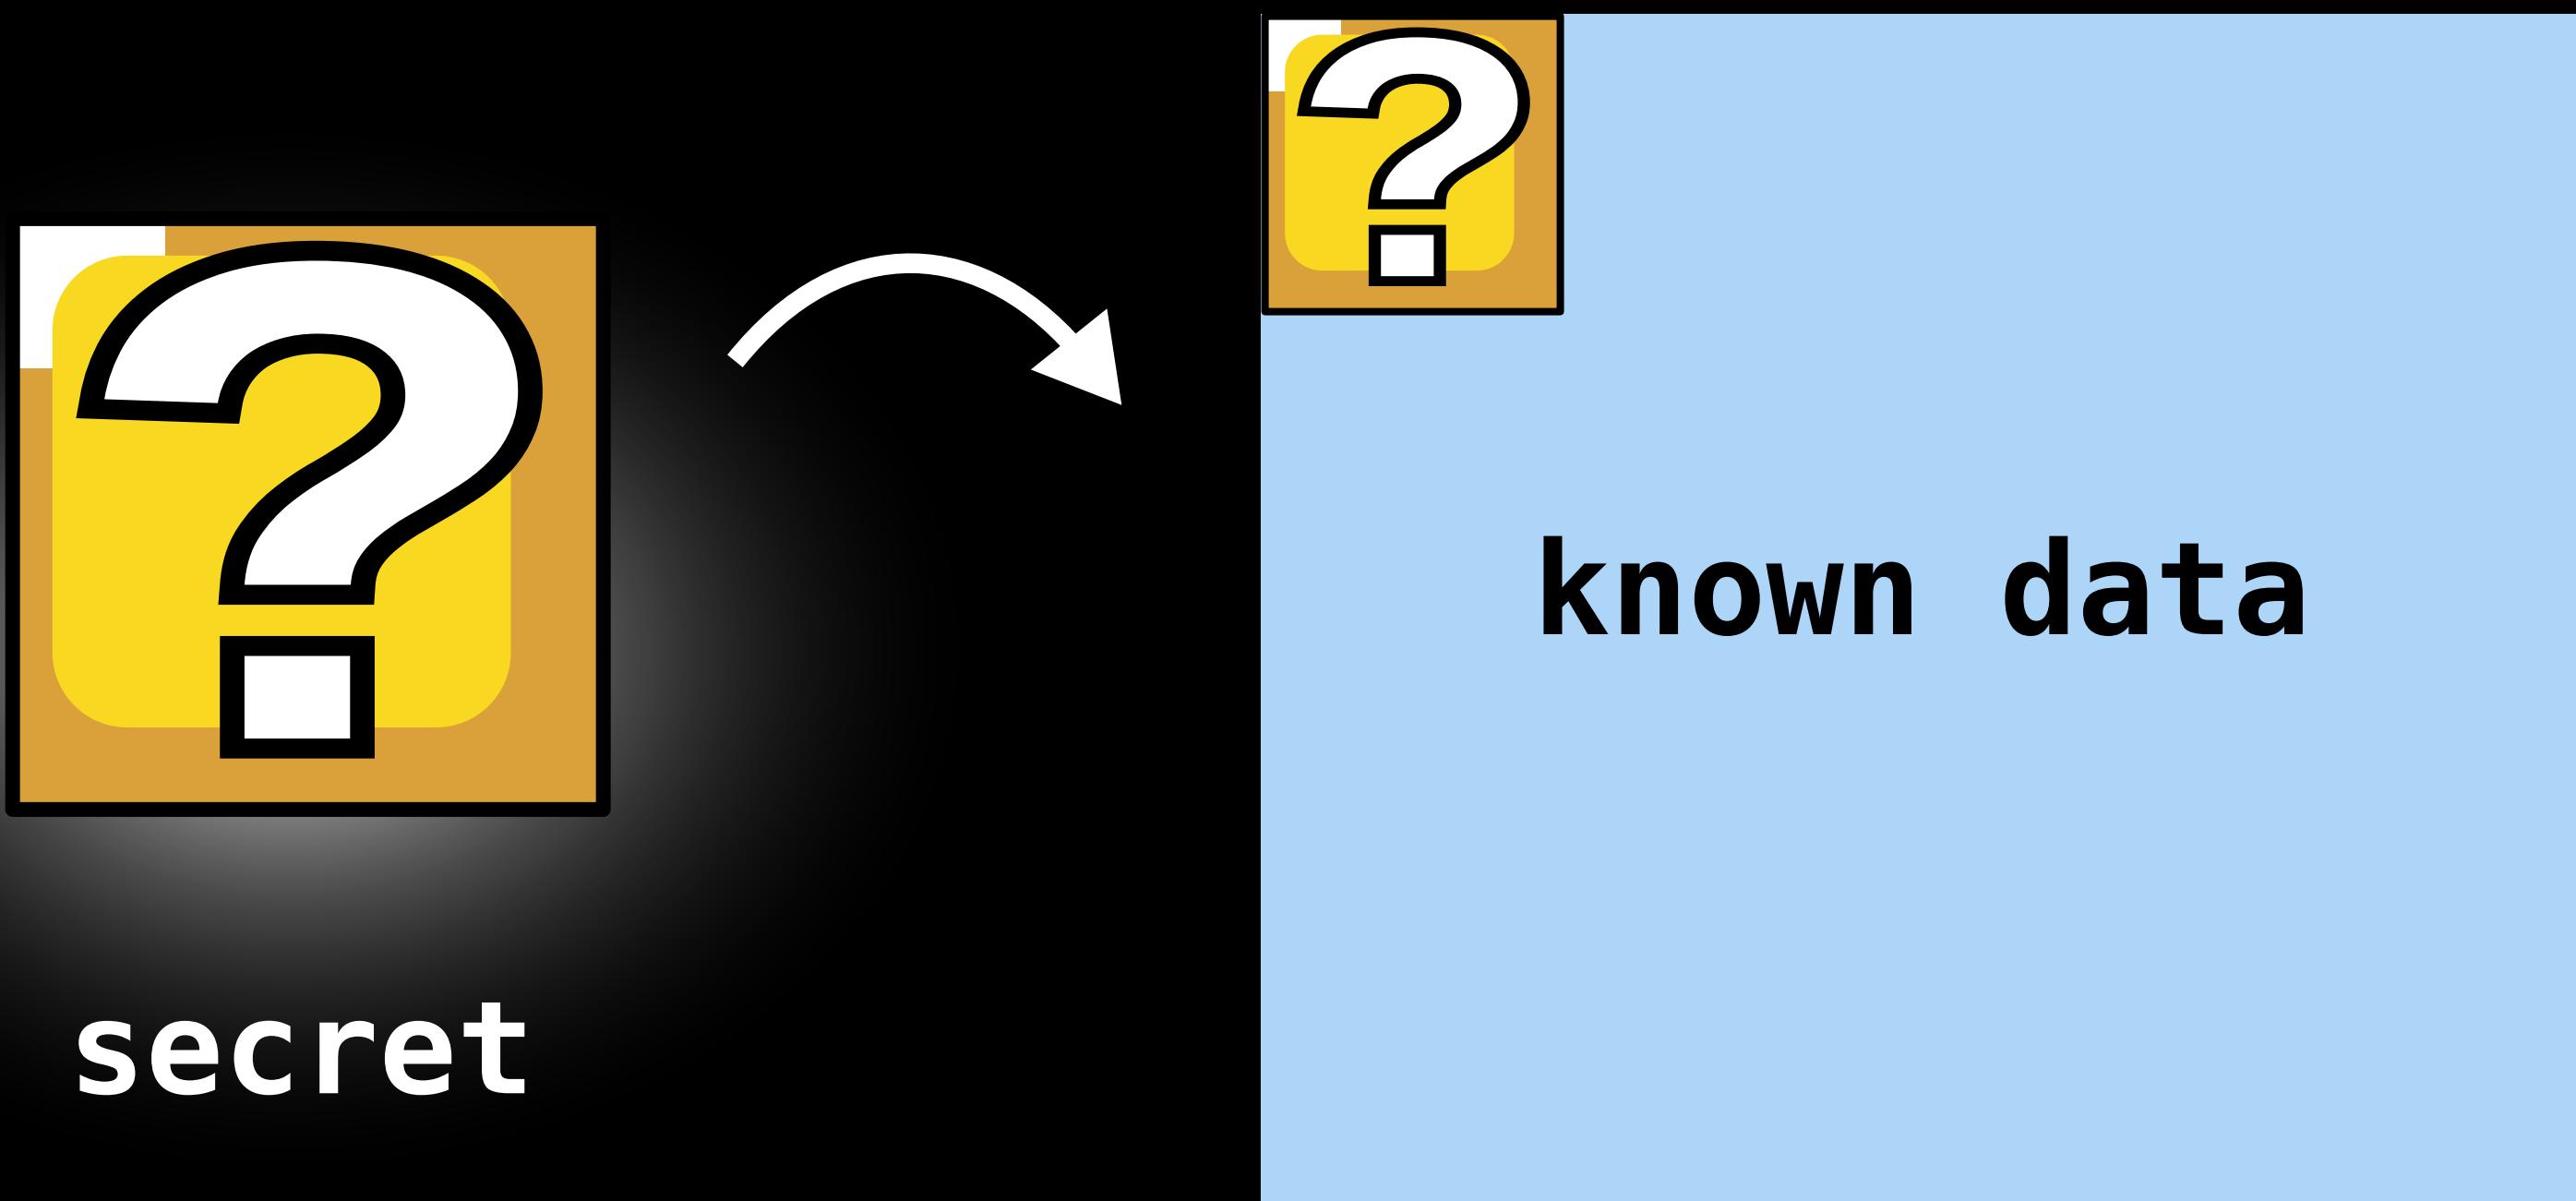
The secret we want to leak does not span an entire page.

# turning a secret into a page



secret

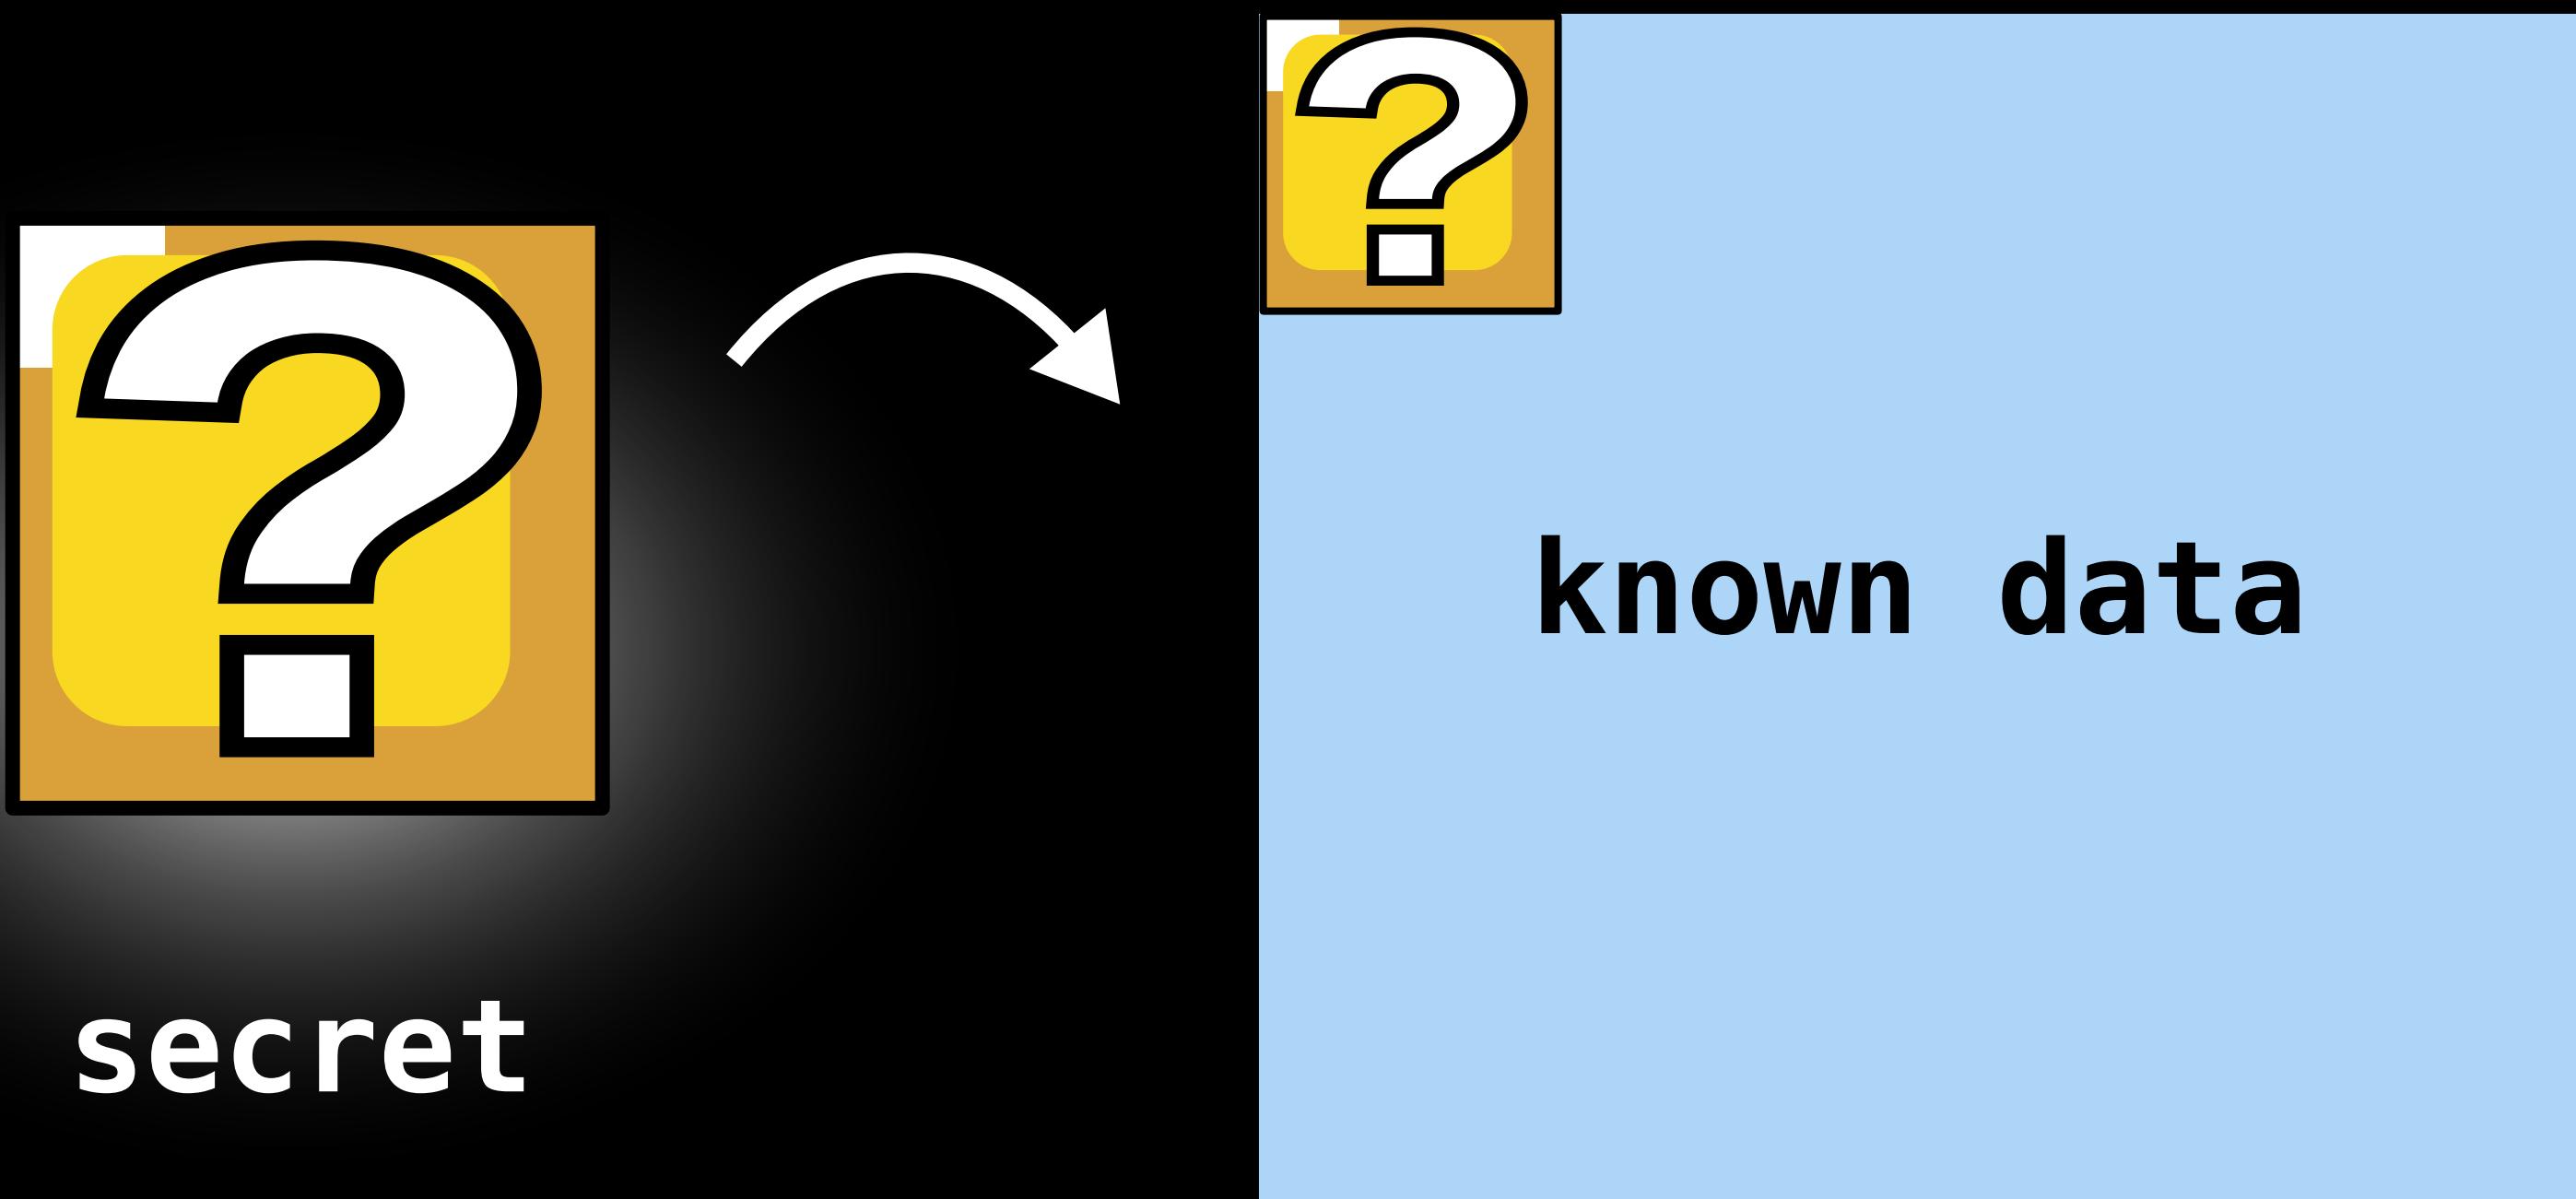
# turning a secret into a page



## Challenge 2:

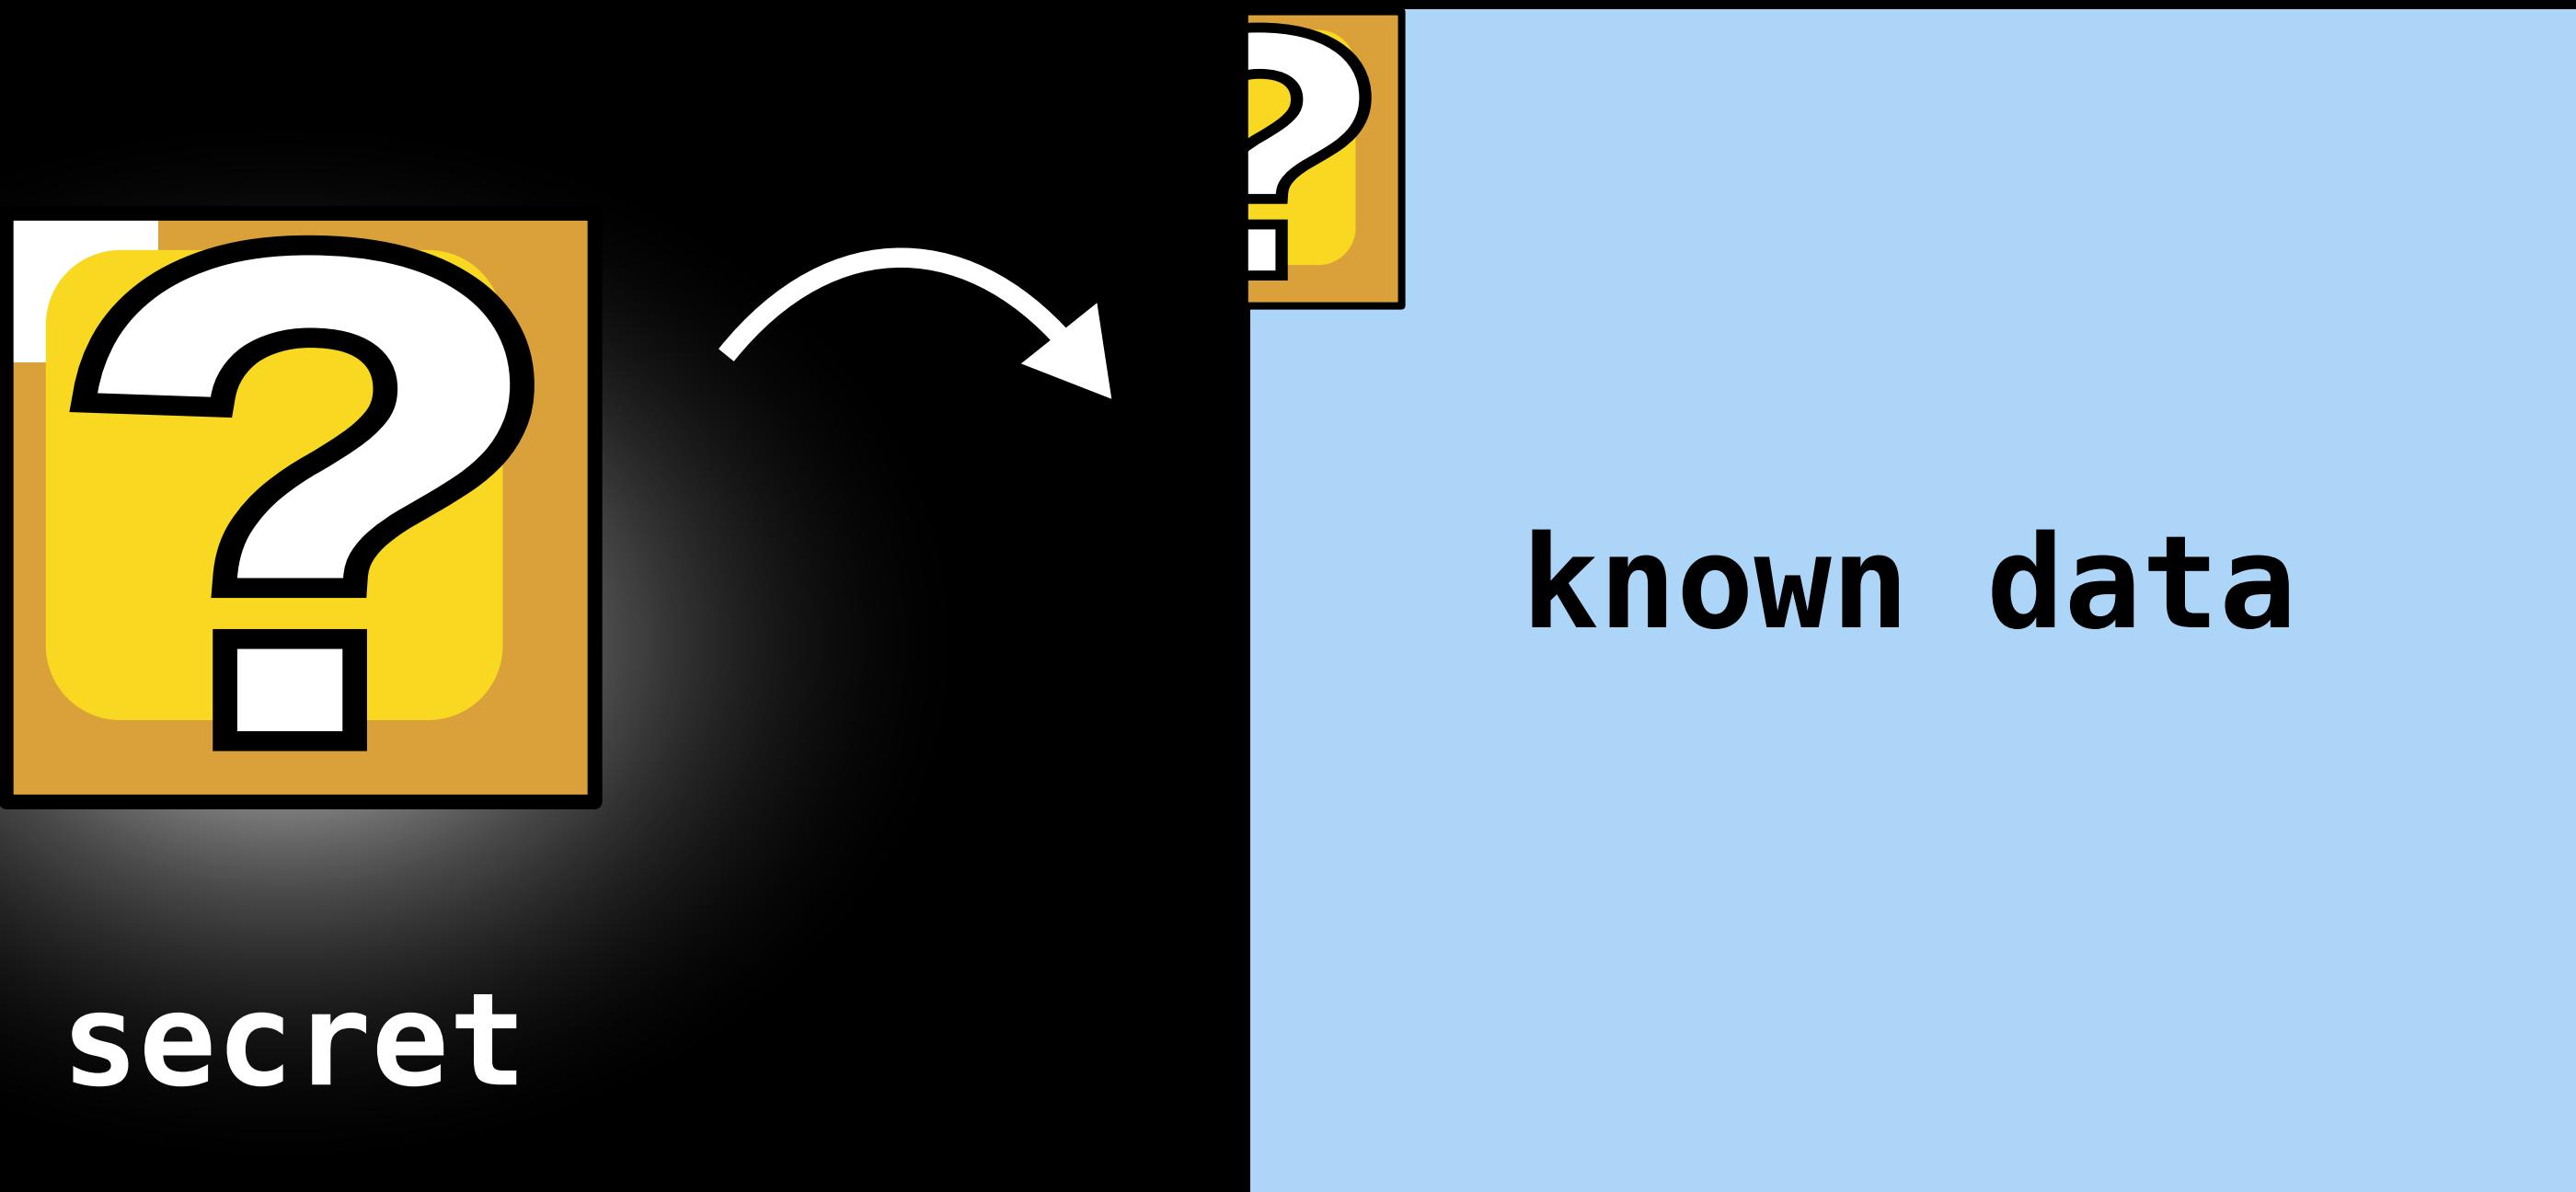
The secret we want to leak has too much entropy to leak all at once.

# primitive #1: alignment probing

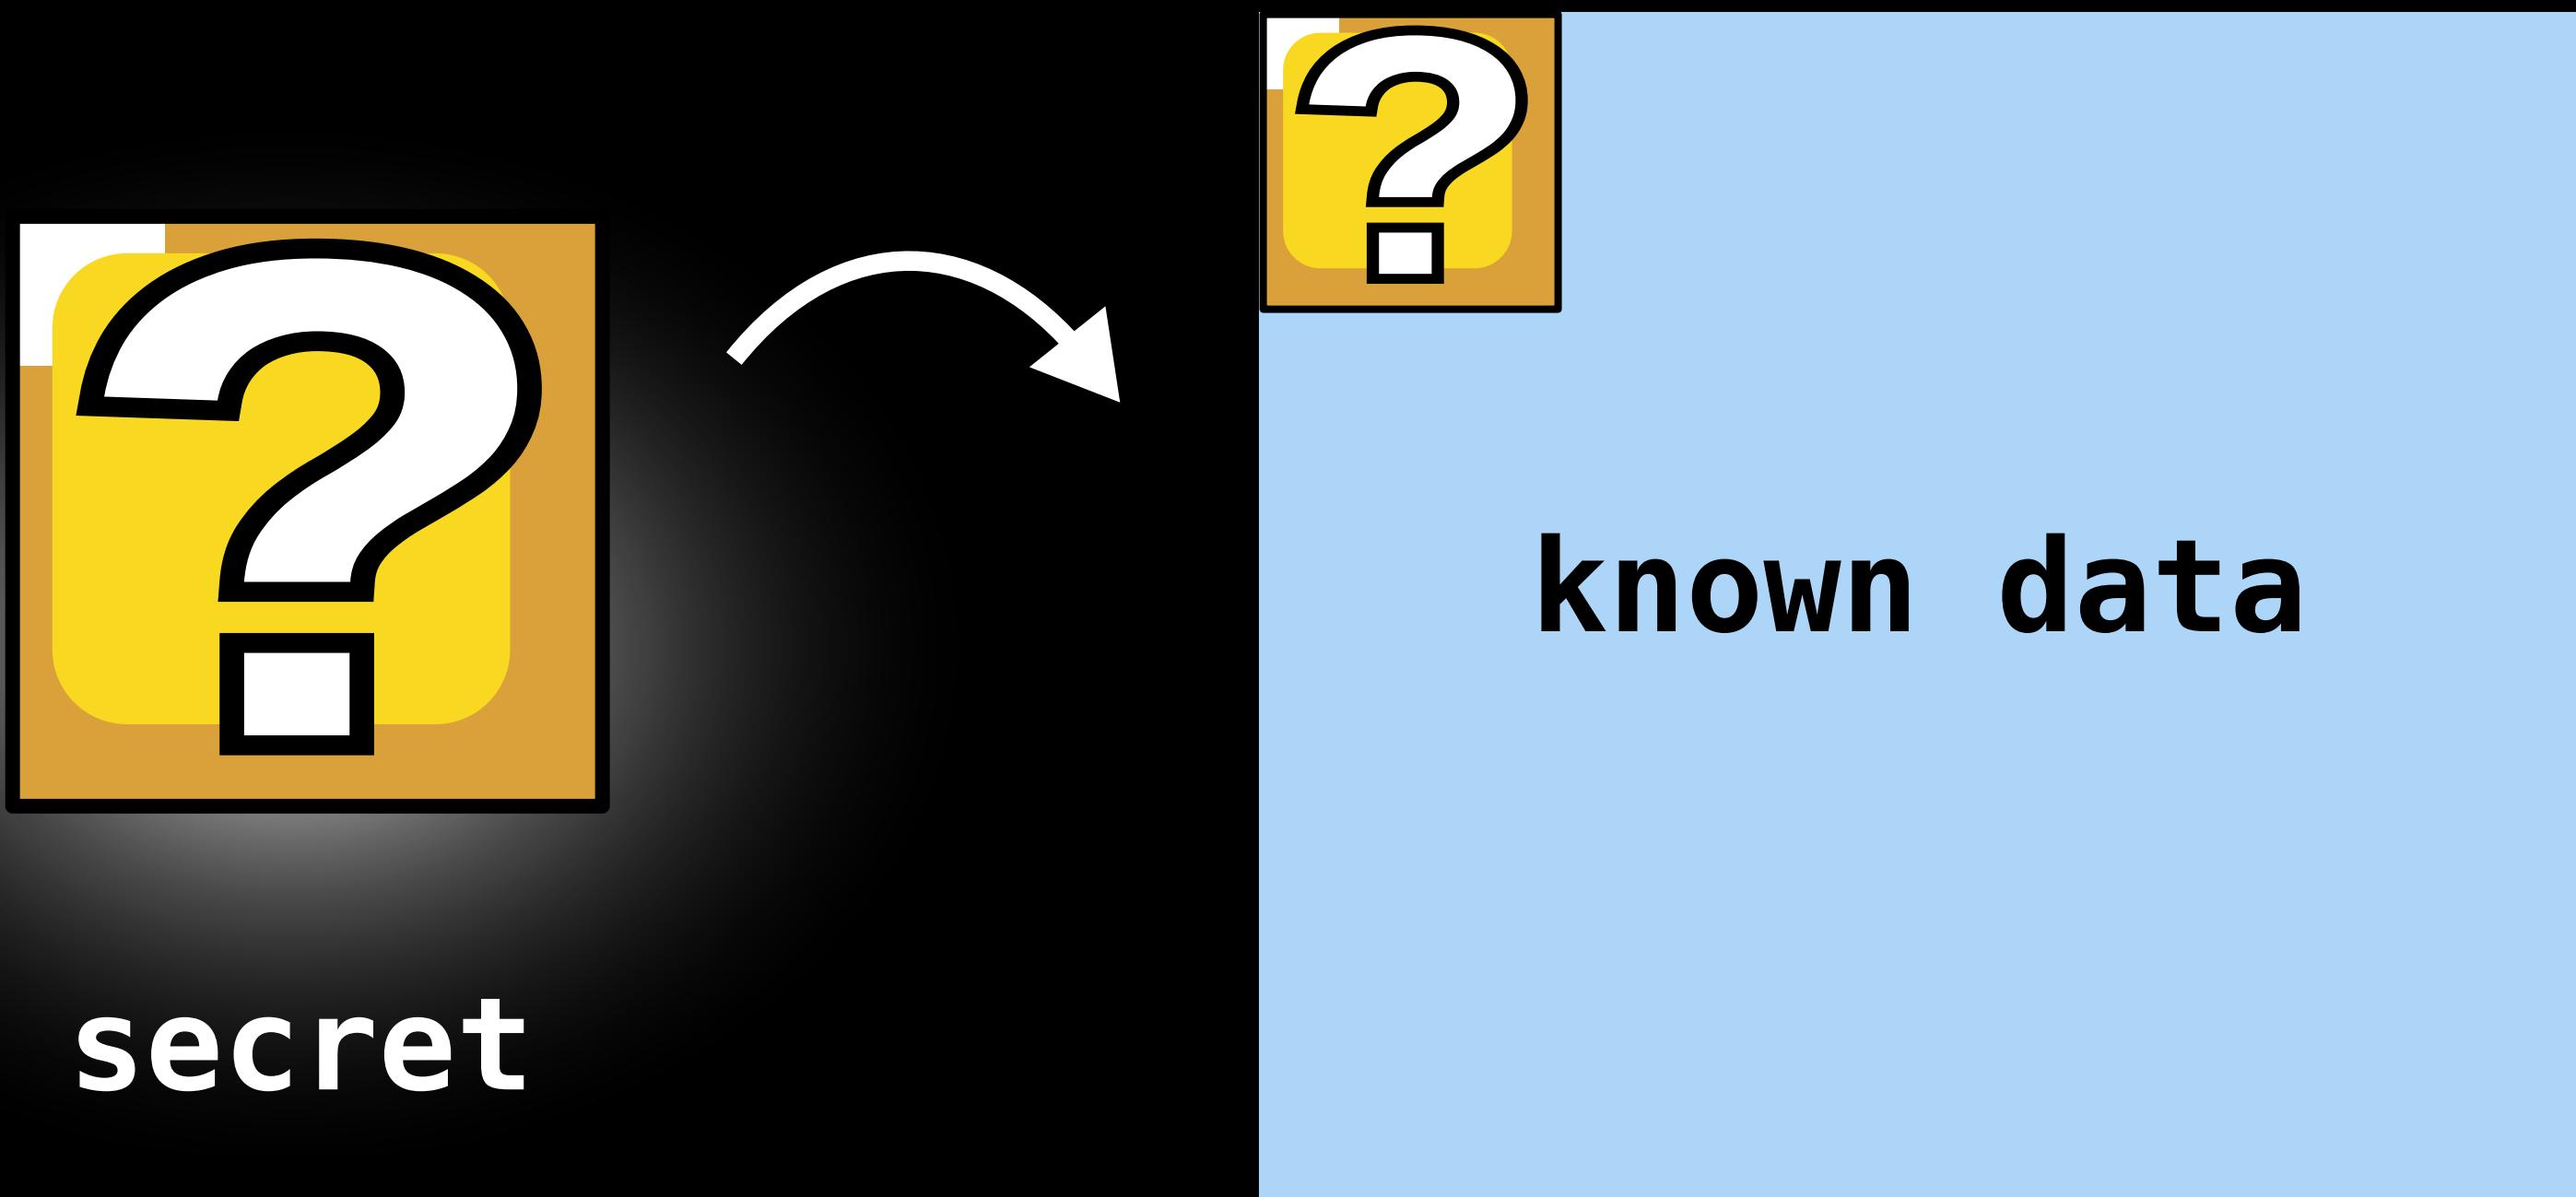


secret page

# primitive #1: alignment probing

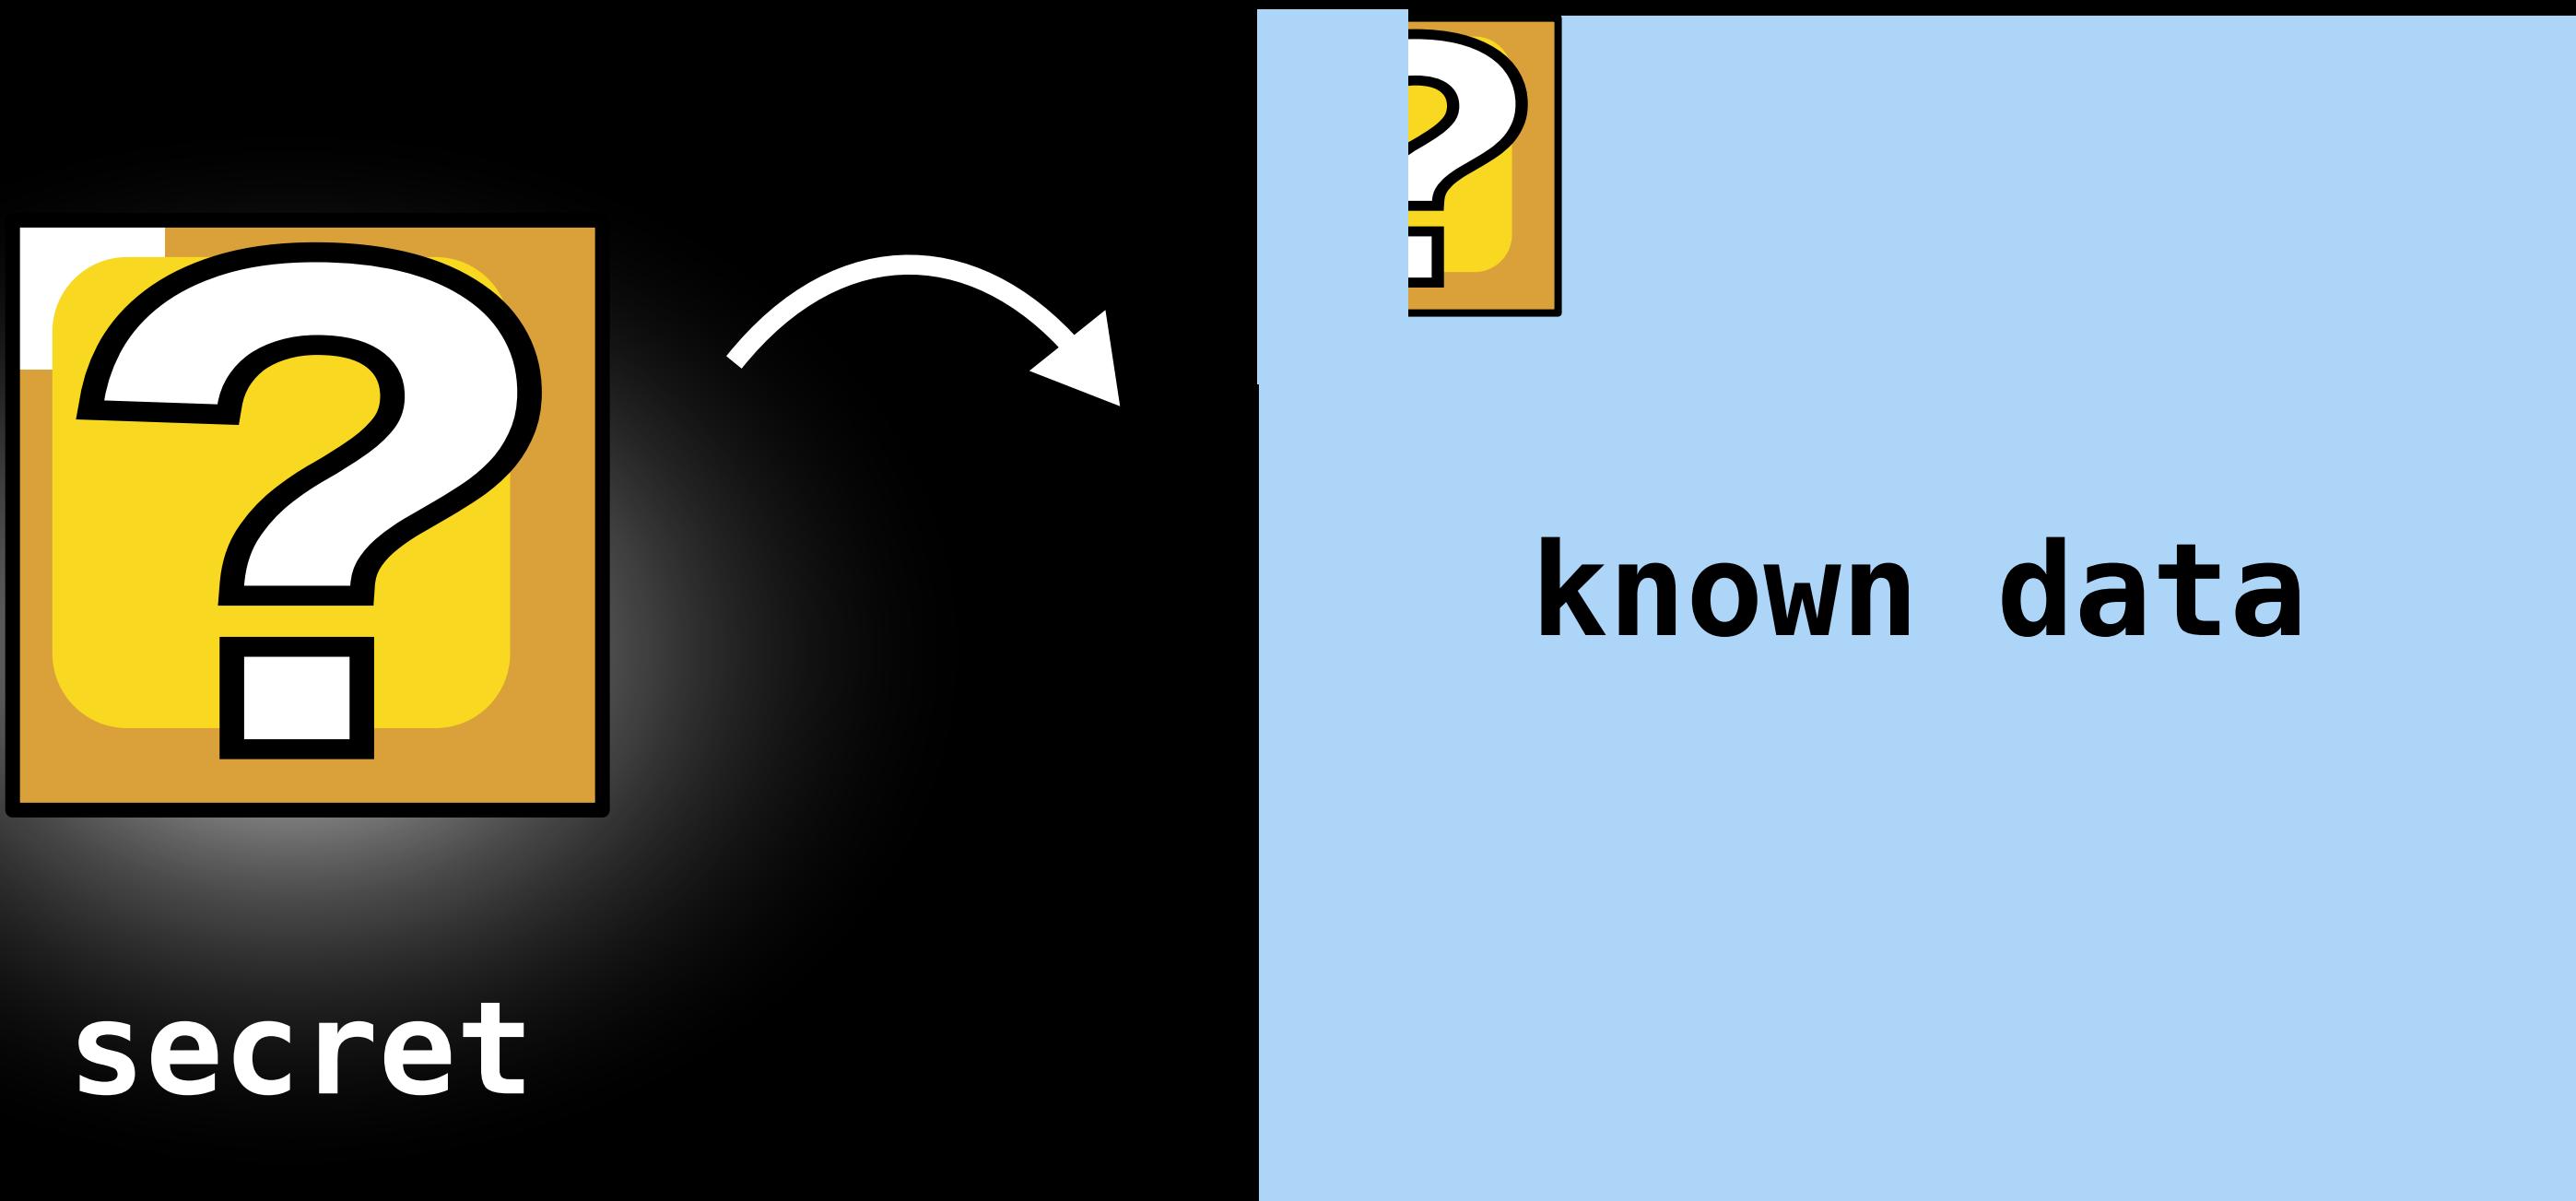


# primitive #2: partial reuse



secret page

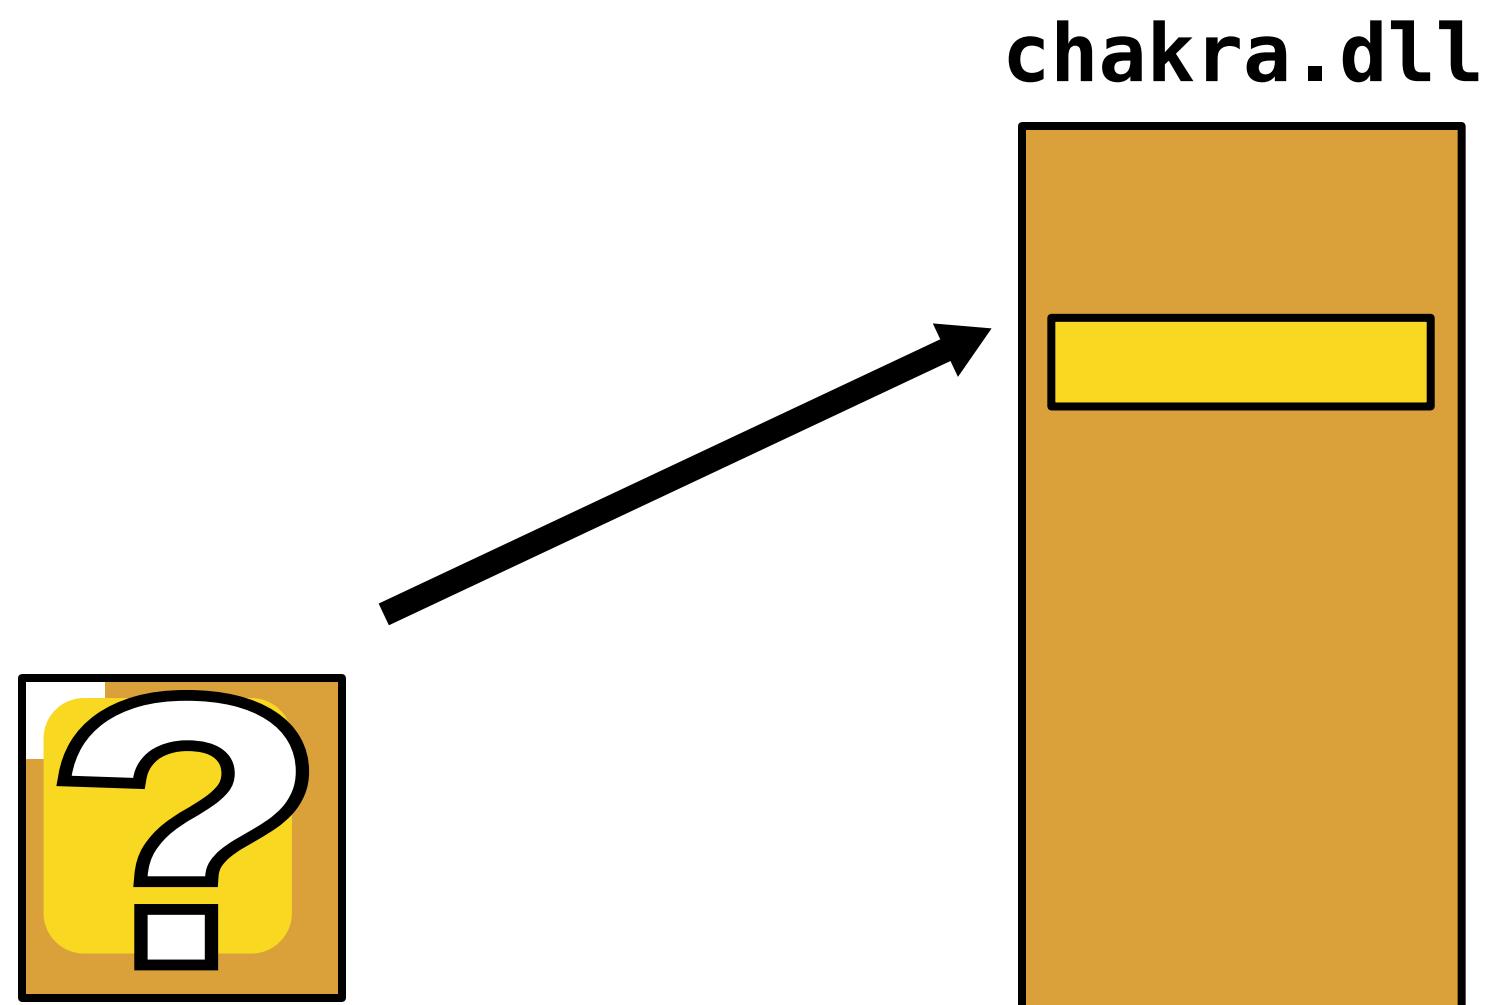
# primitive #2: partial reuse



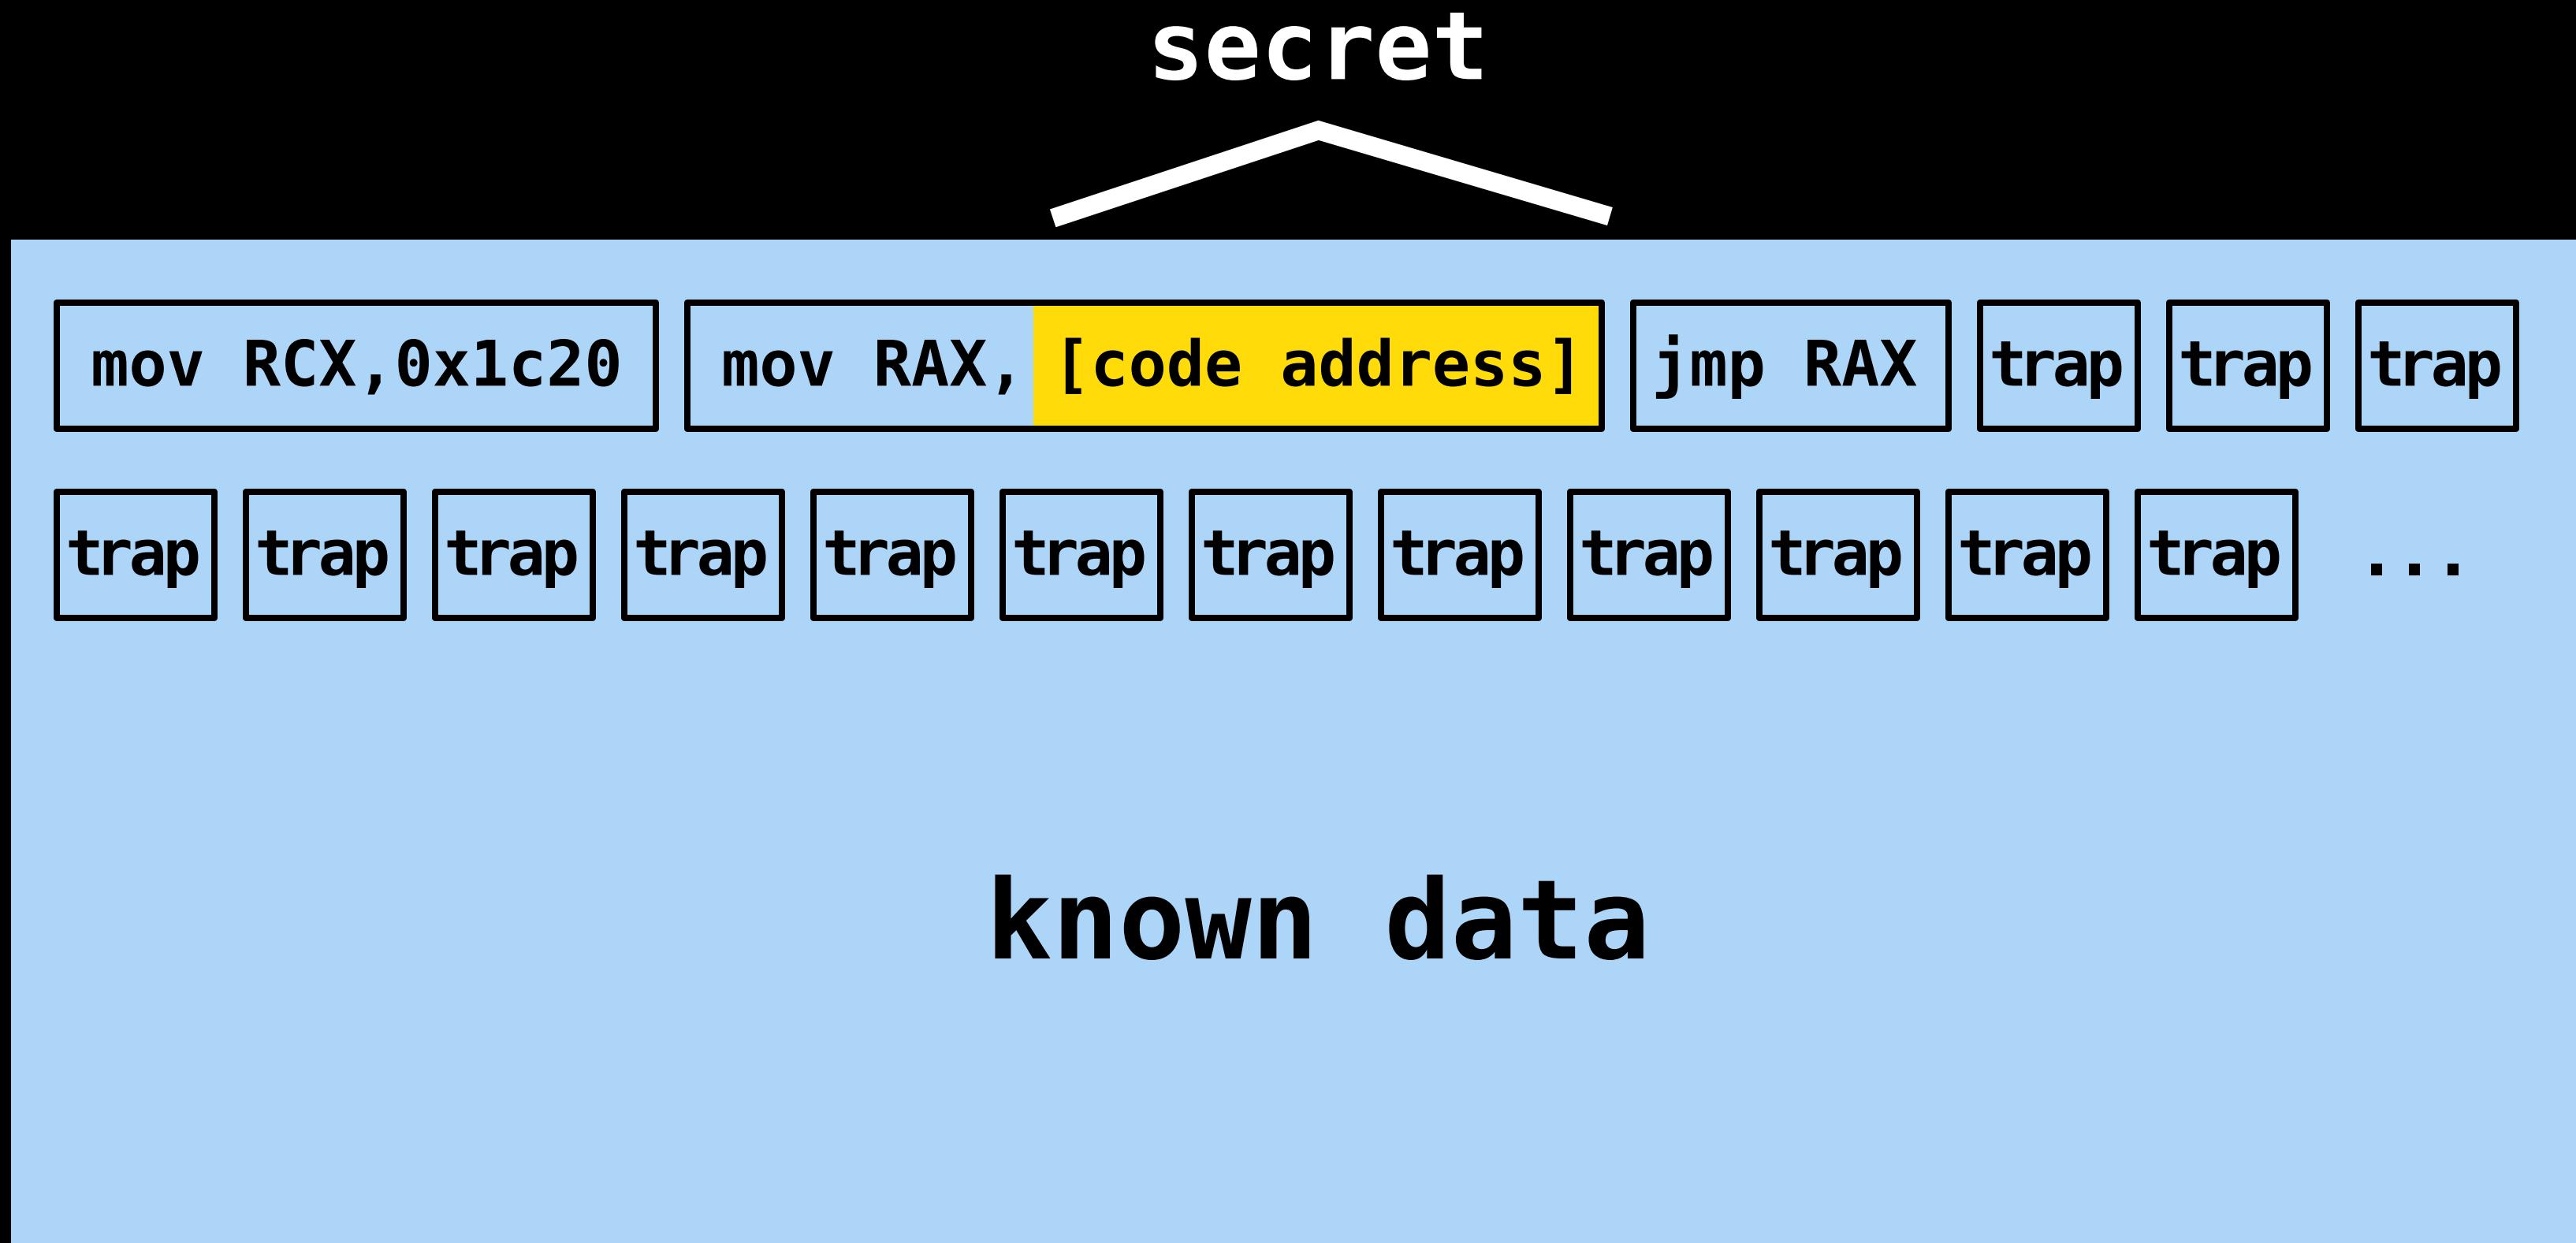
# Outline:

## Deduplication

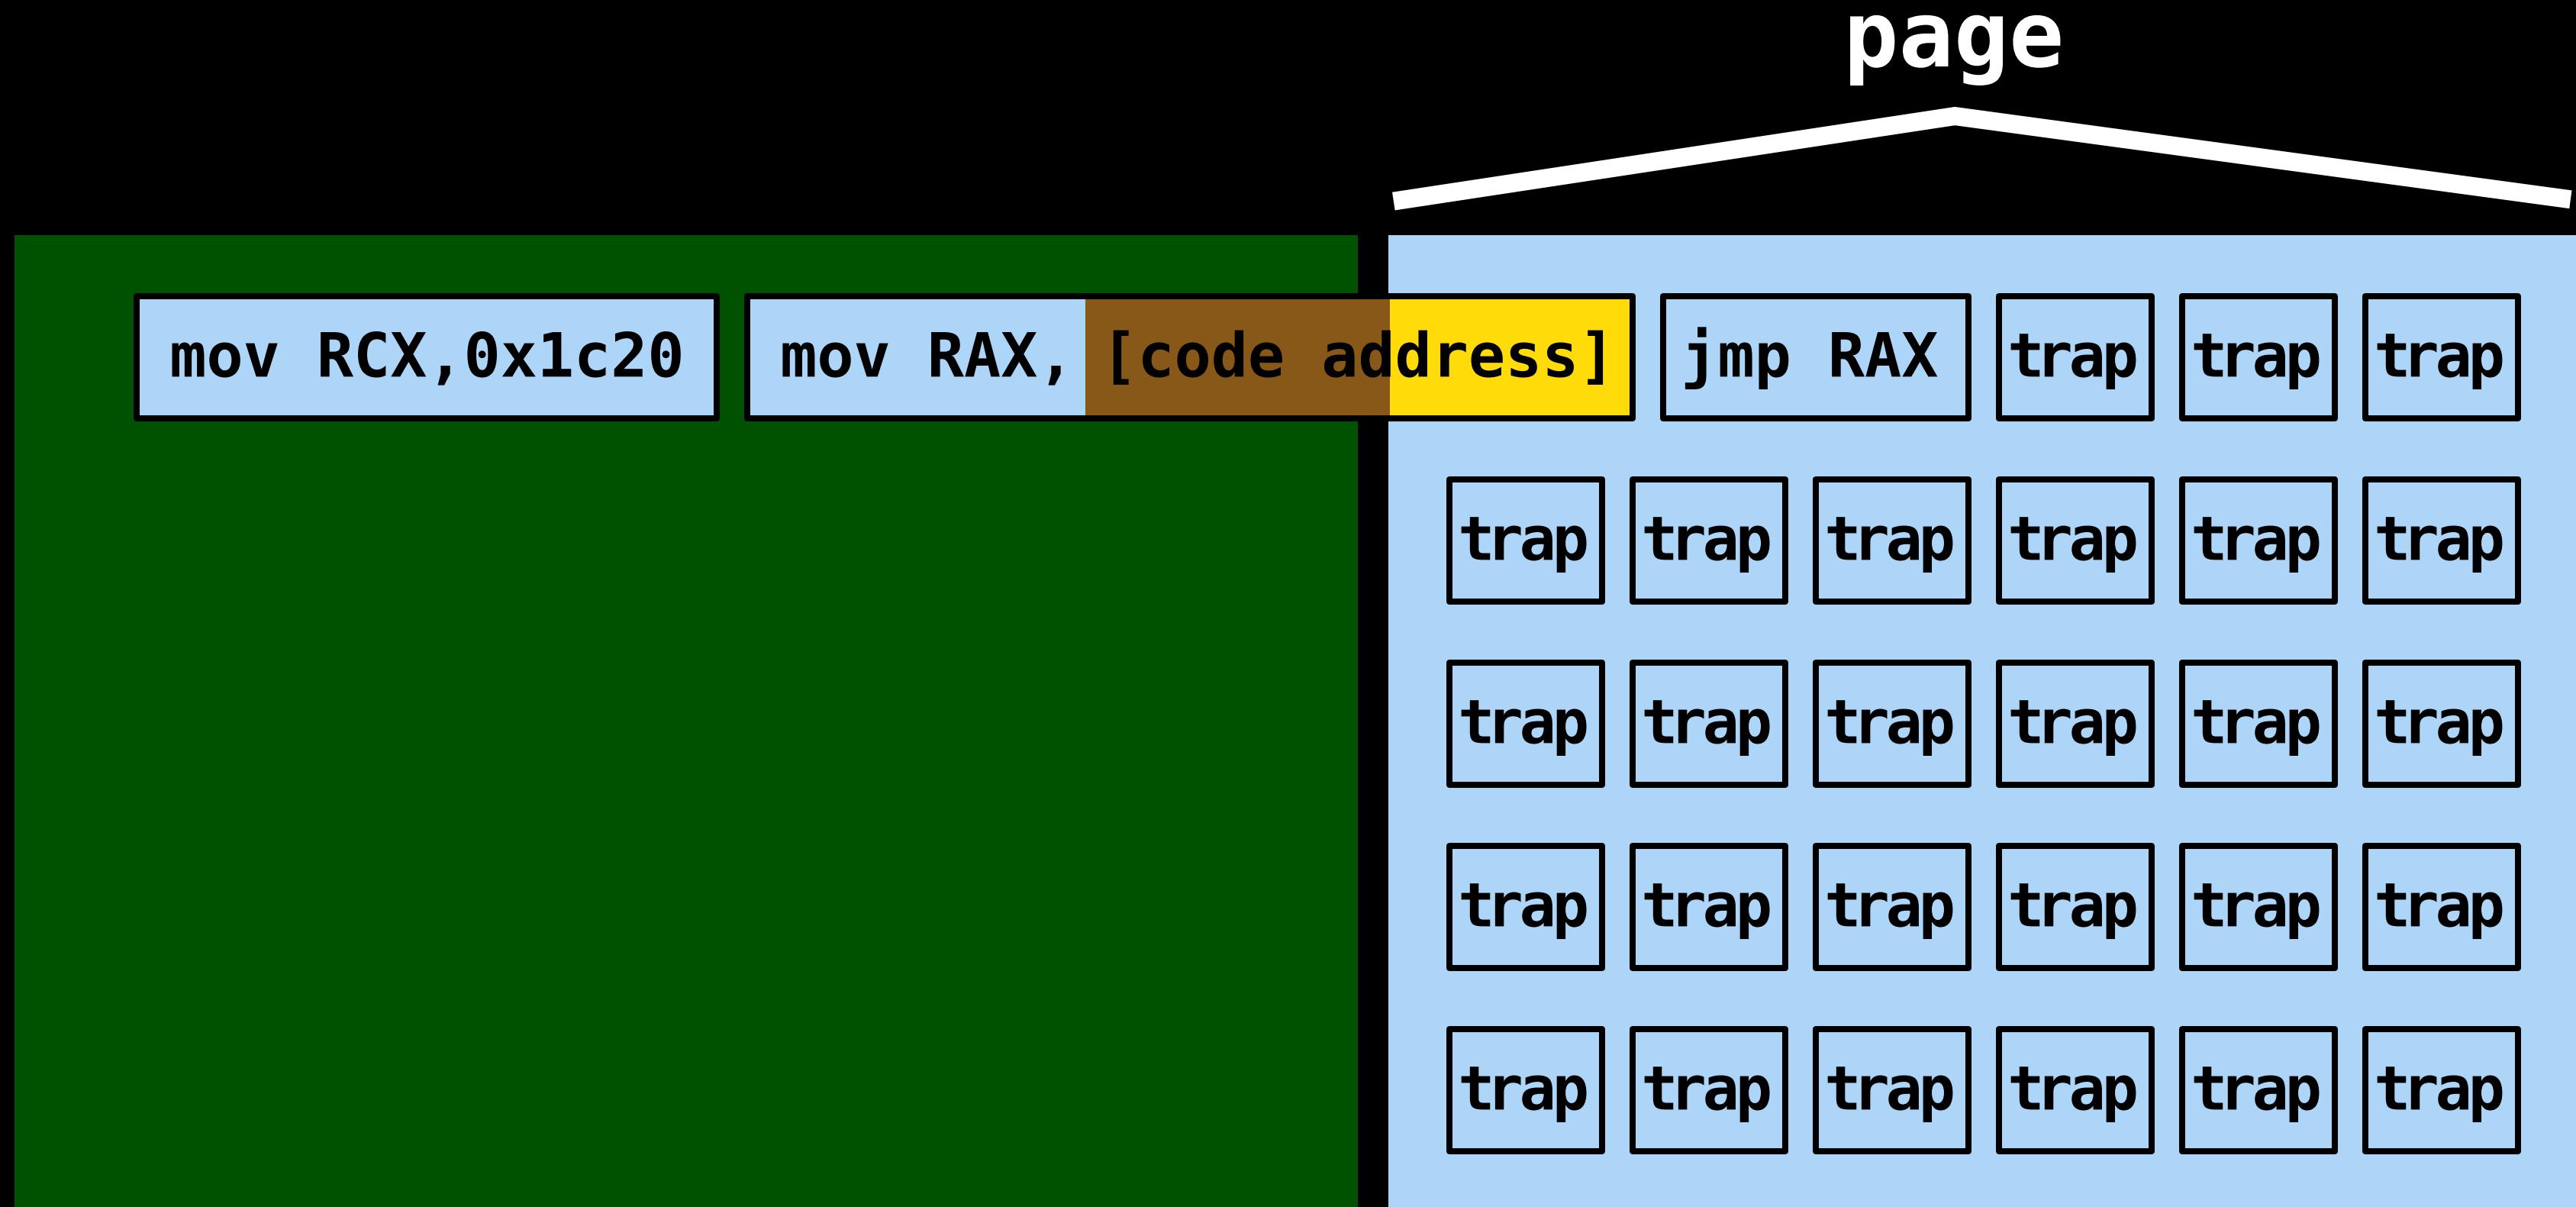
- leak heap & code addresses



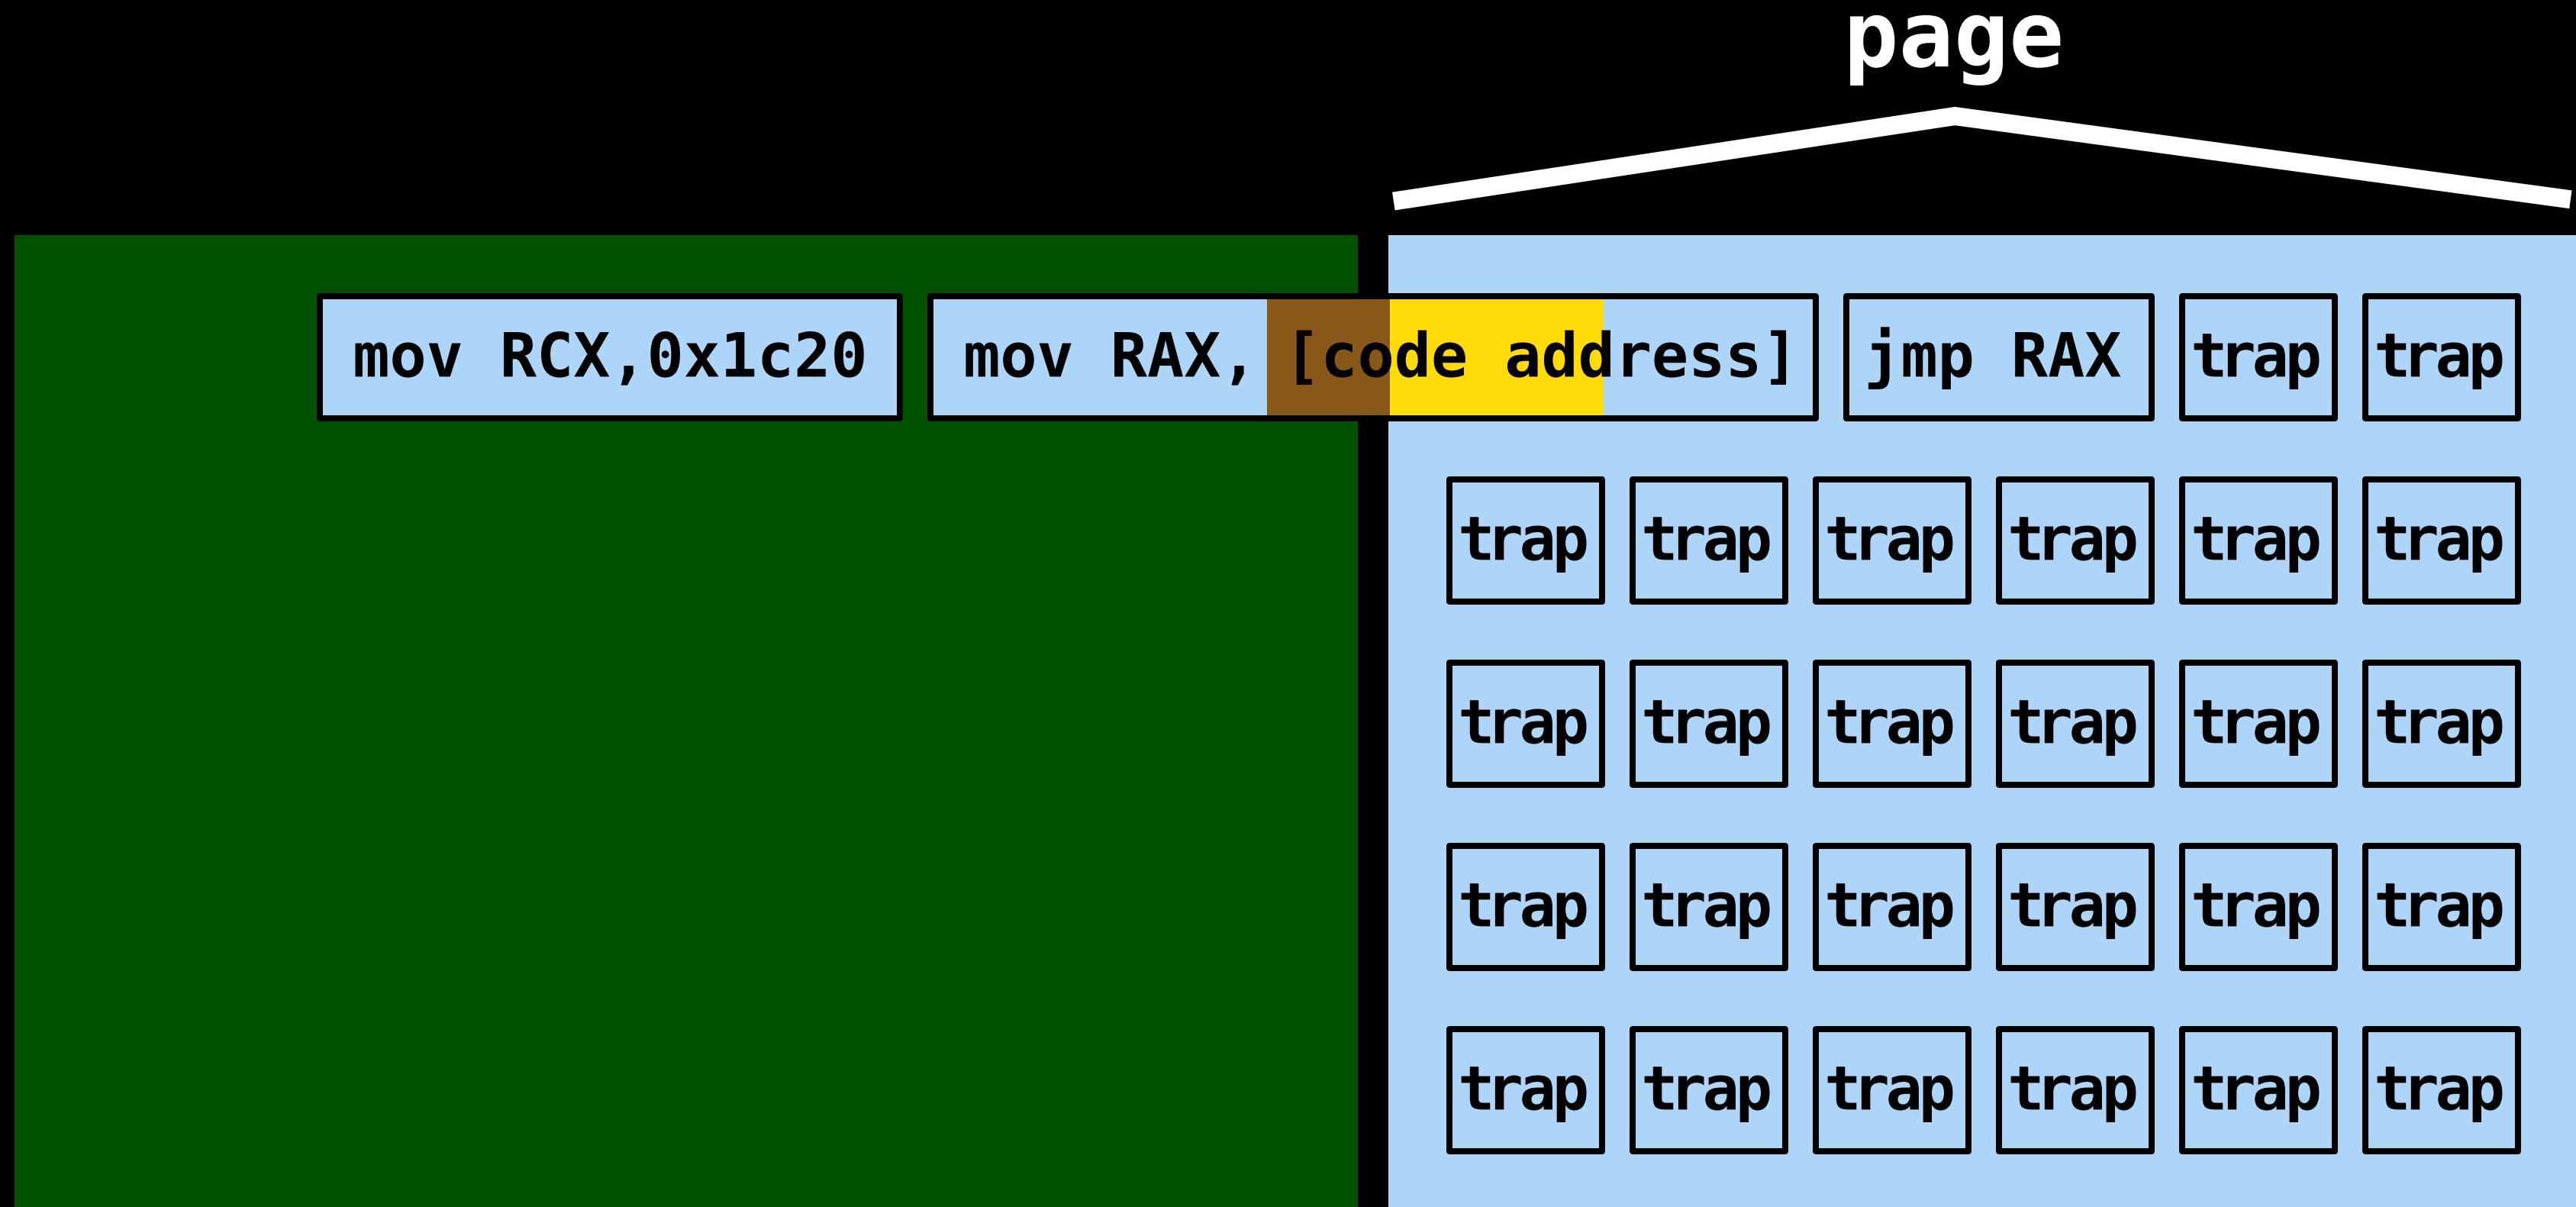
# JIT function epilogue (MS Edge)



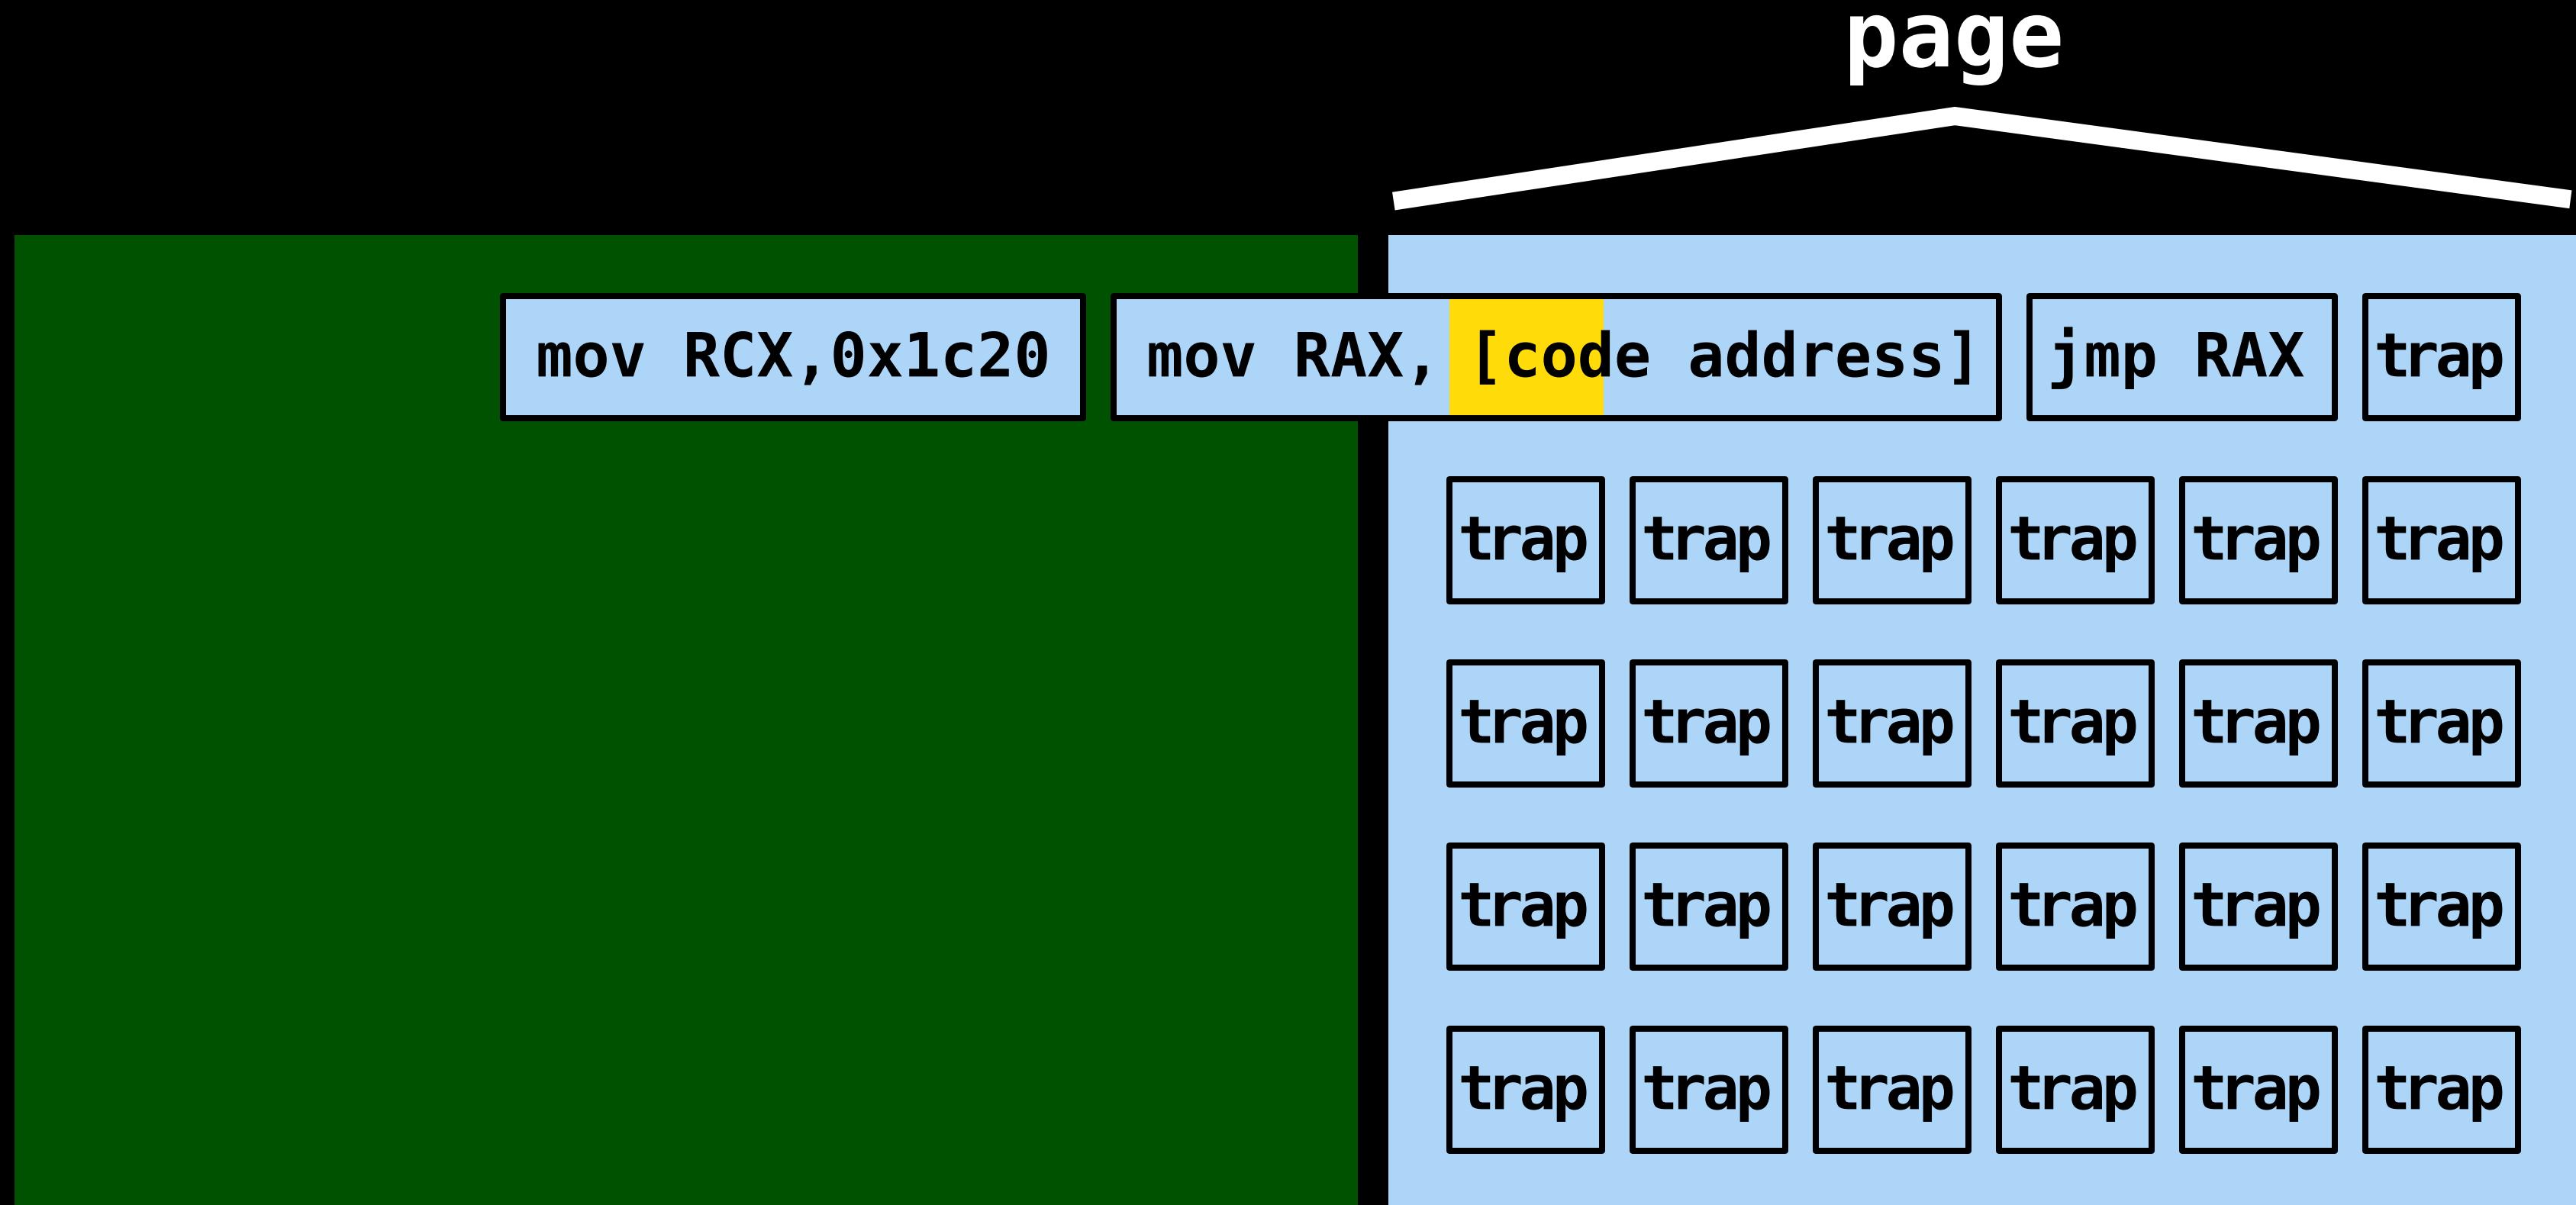
# JIT function epilogue (MS Edge)



# JIT function epilogue (MS Edge)



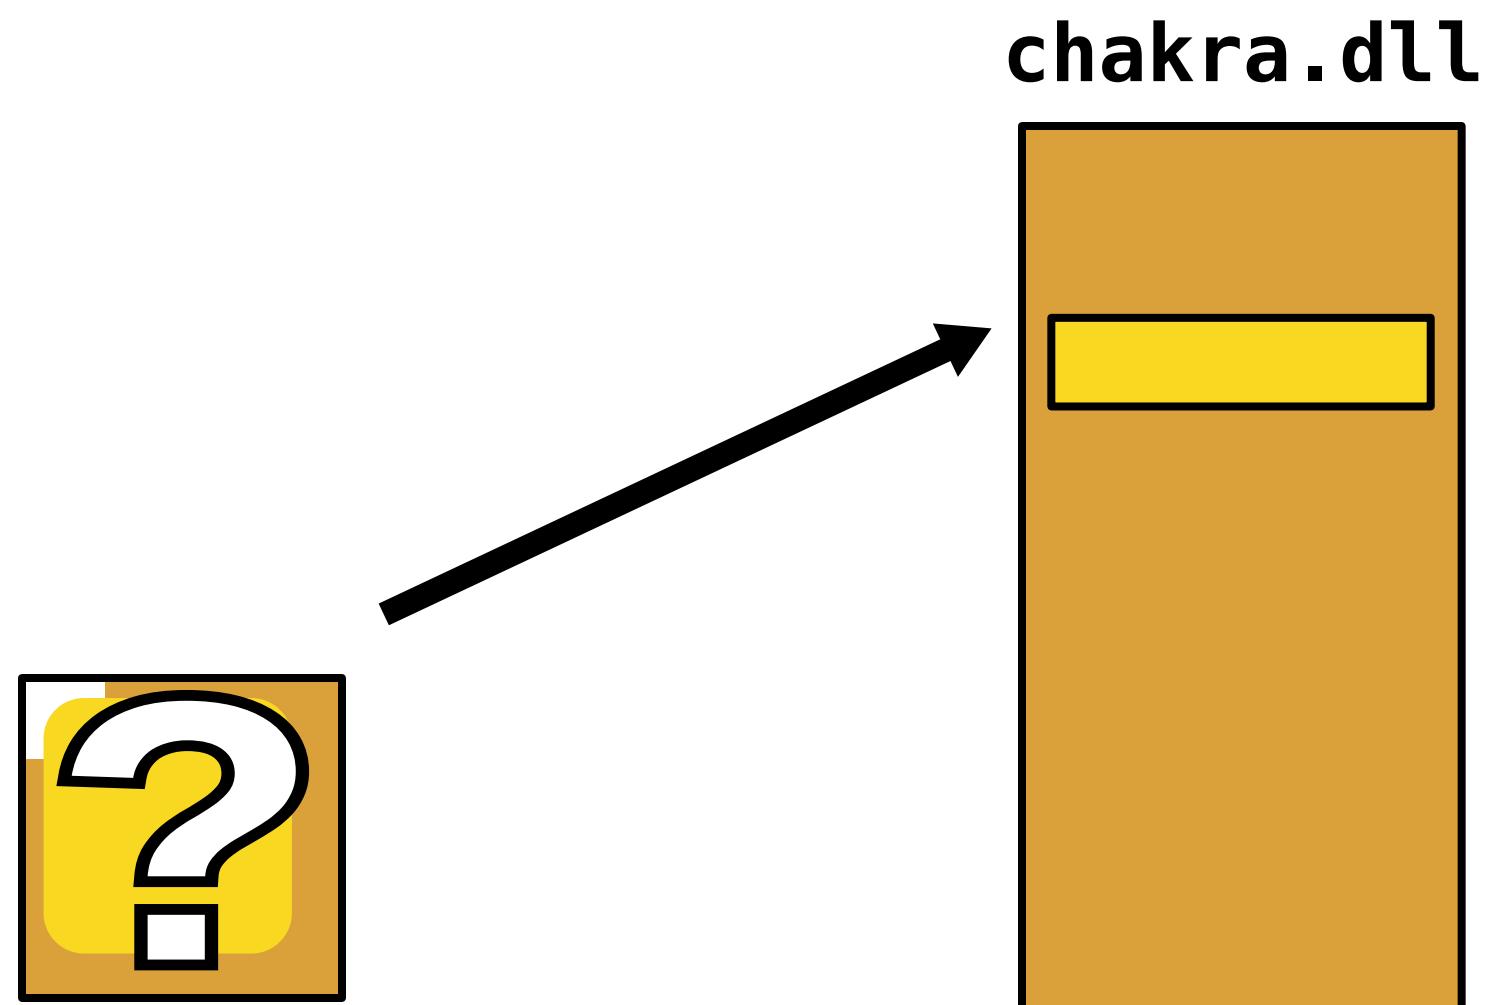
# JIT function epilogue (MS Edge)



# Outline:

## Deduplication

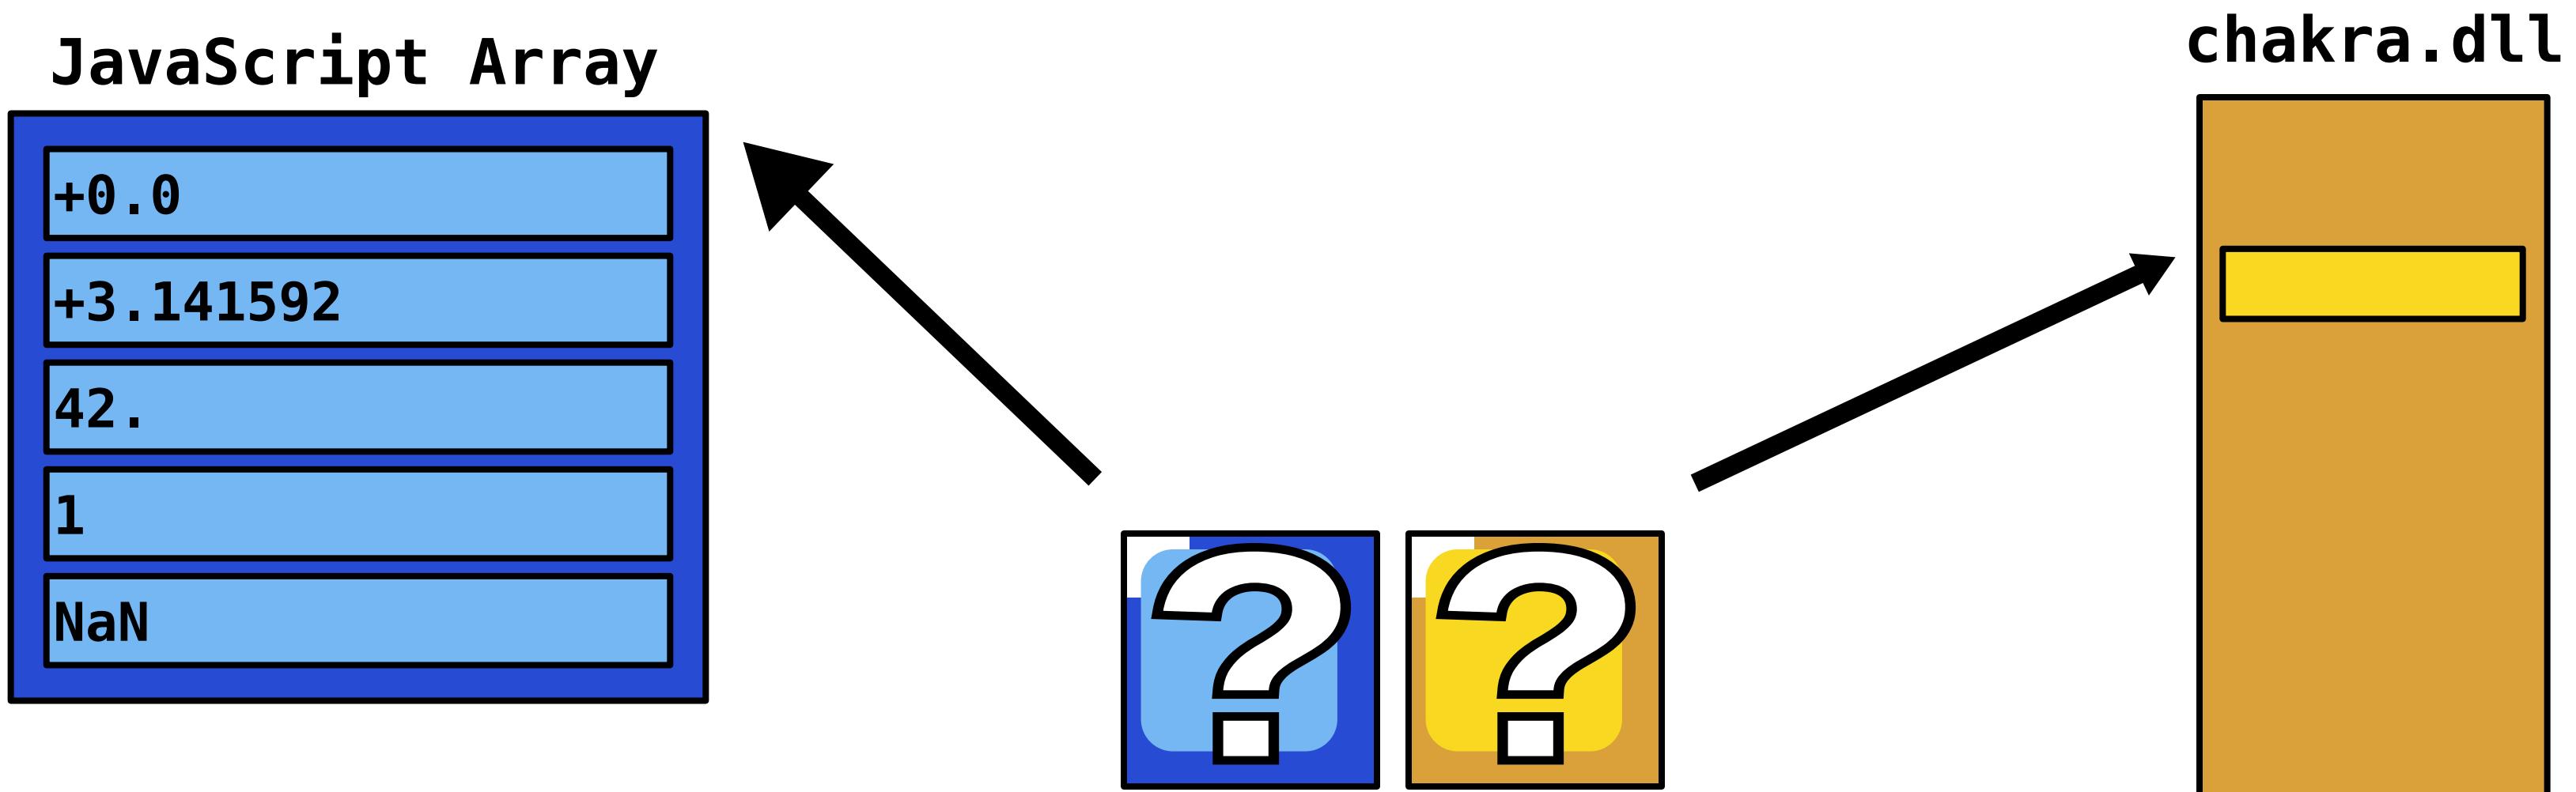
- leak heap & code addresses



# Outline:

## Deduplication

- leak heap & code addresses



What if leaking a heap pointer in stages is not possible...

We need to guess a page containing the complete pointer.

# Heap pointer entropy in Edge

0x5F48143540

# Heap pointer entropy in Edge

advertised ASLR (24 bit)



0x5F48143540

# Heap pointer entropy in Edge

advertised ASLR (24 bit)



0x5F48143540



non-deterministic bits  
(+/- 36 bit)

# Heap pointer entropy in Edge

**64G**

advertised ASLR (24 bit)

0x5F48143540

non-deterministic bits  
(+/- 36 bit)

# Heap pointer entropy in Edge

**64G**

advertised ASLR (24 bit)



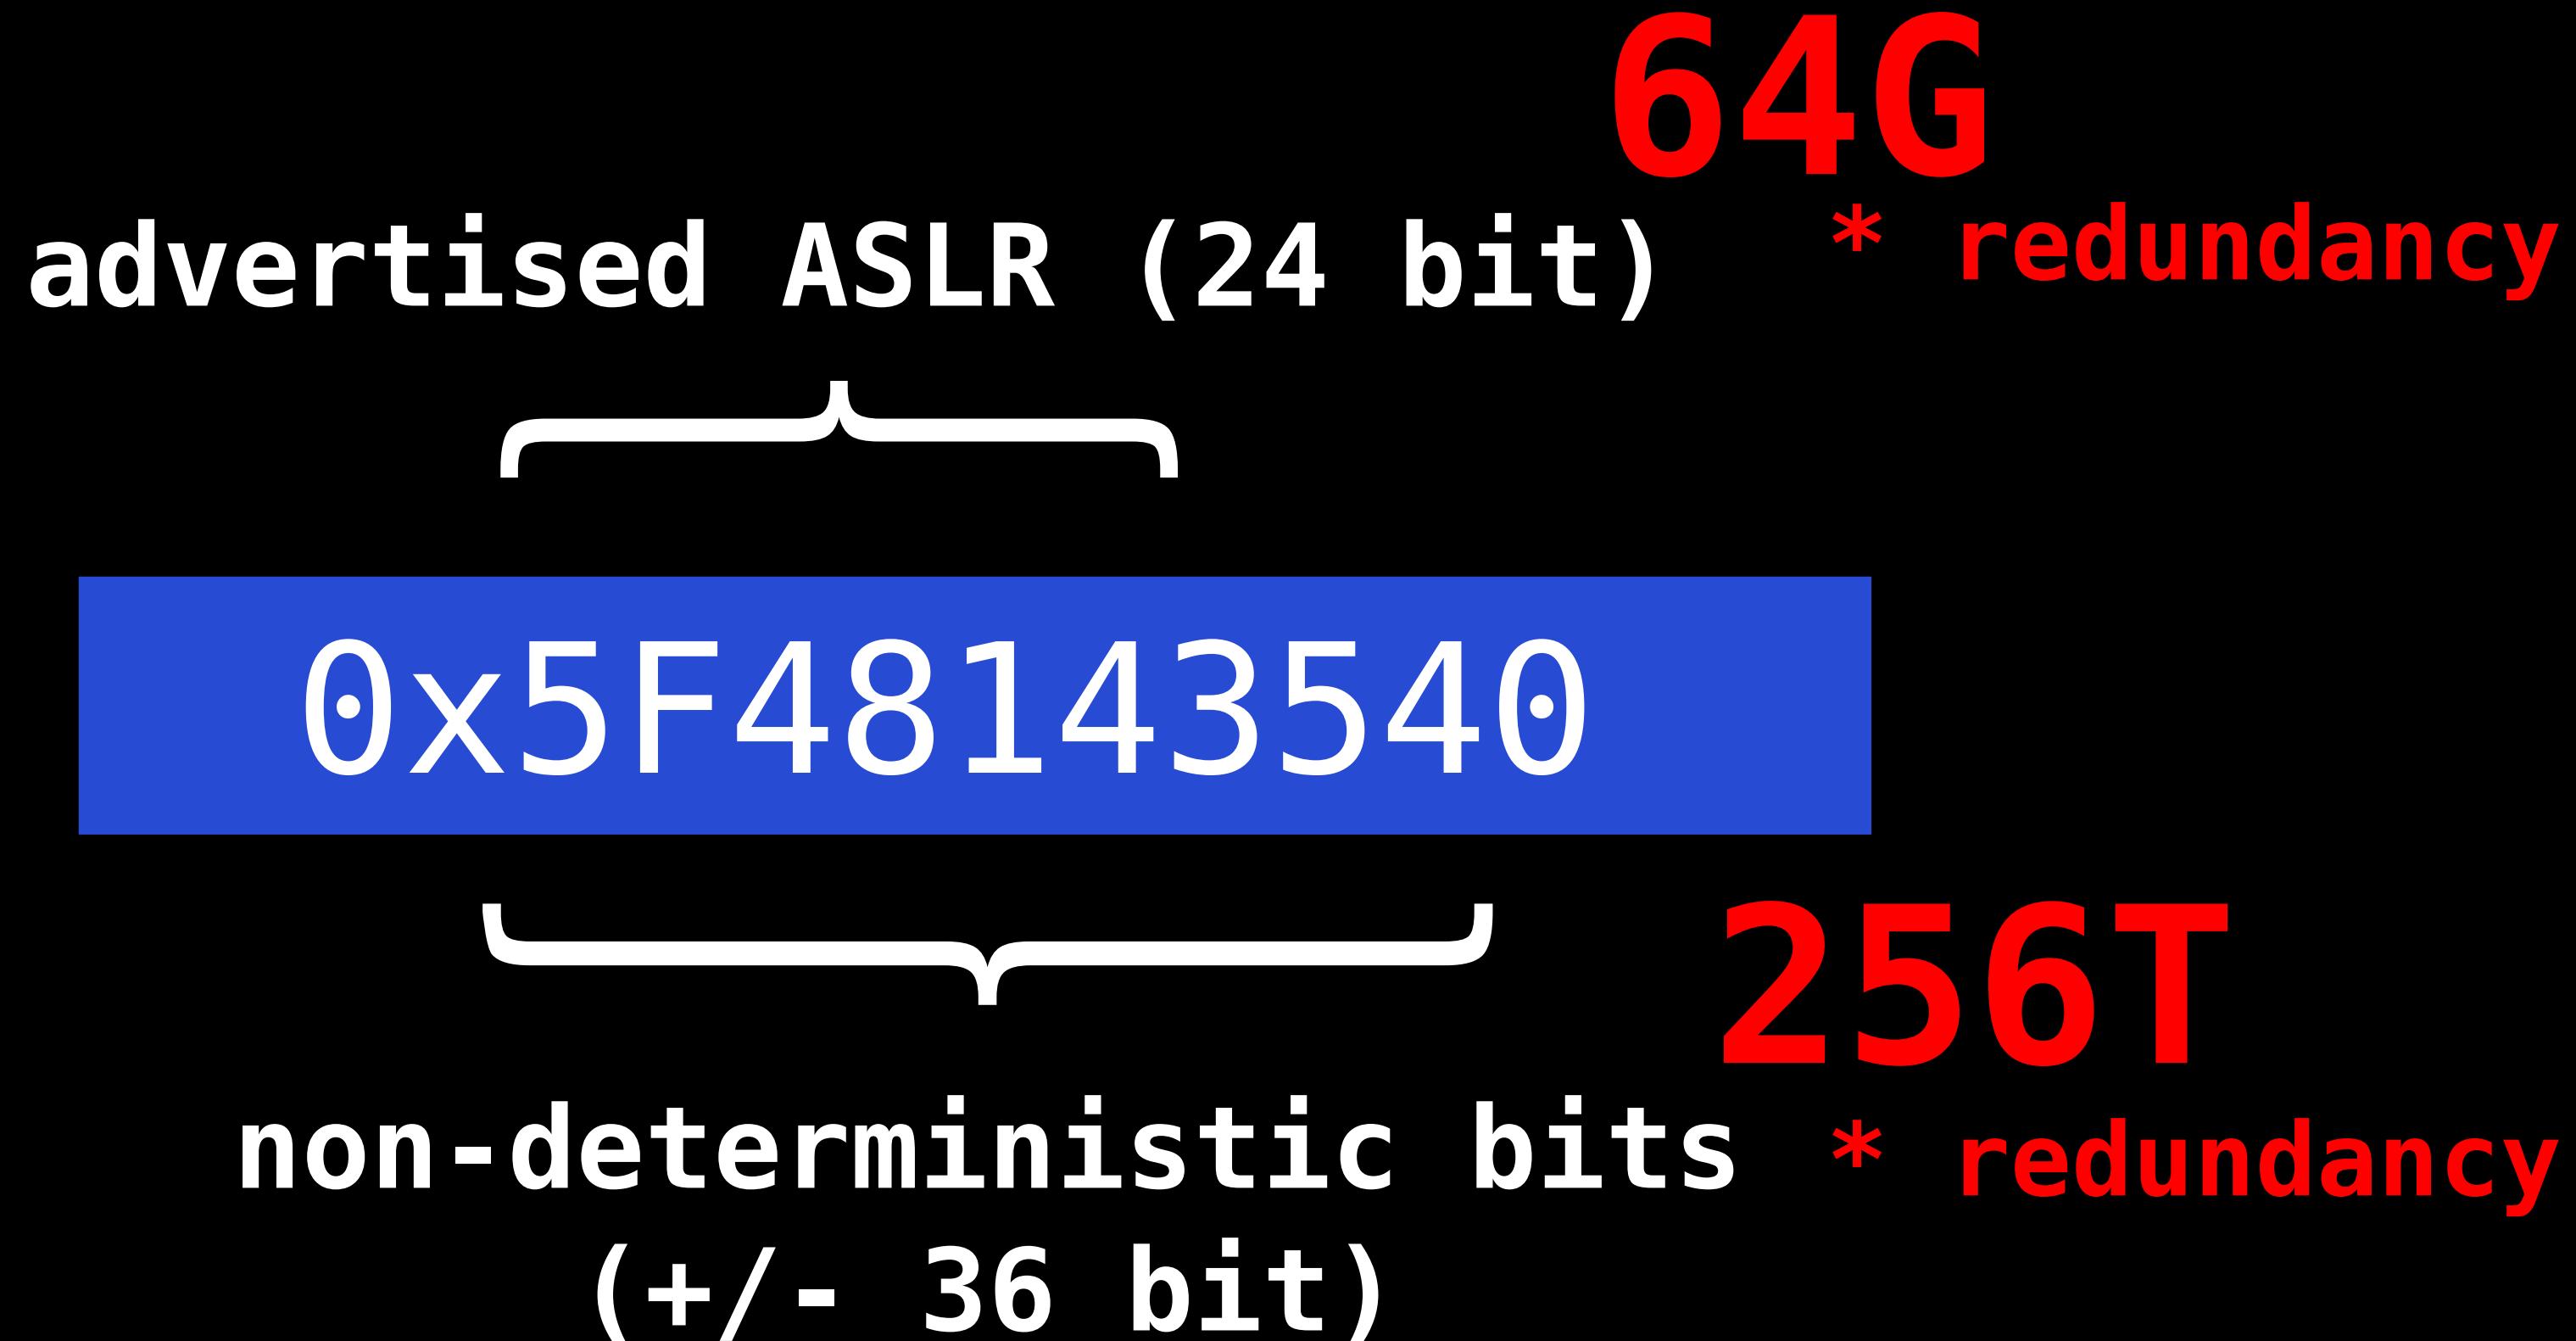
0x5F48143540



**256T**

non-deterministic bits  
(+/- 36 bit)

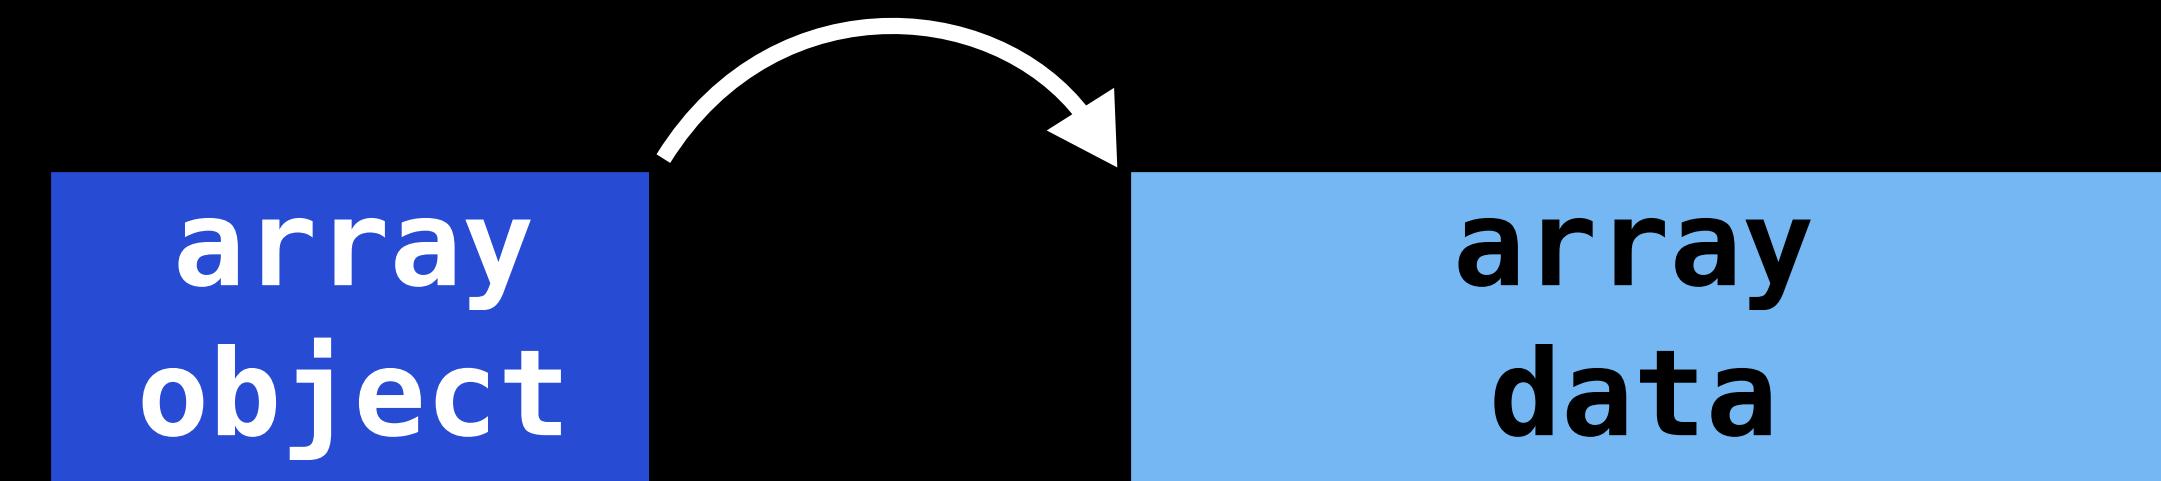
# Heap pointer entropy in Edge



# Slab allocator for JavaScript objects

array  
object

# Slab allocator for JavaScript objects

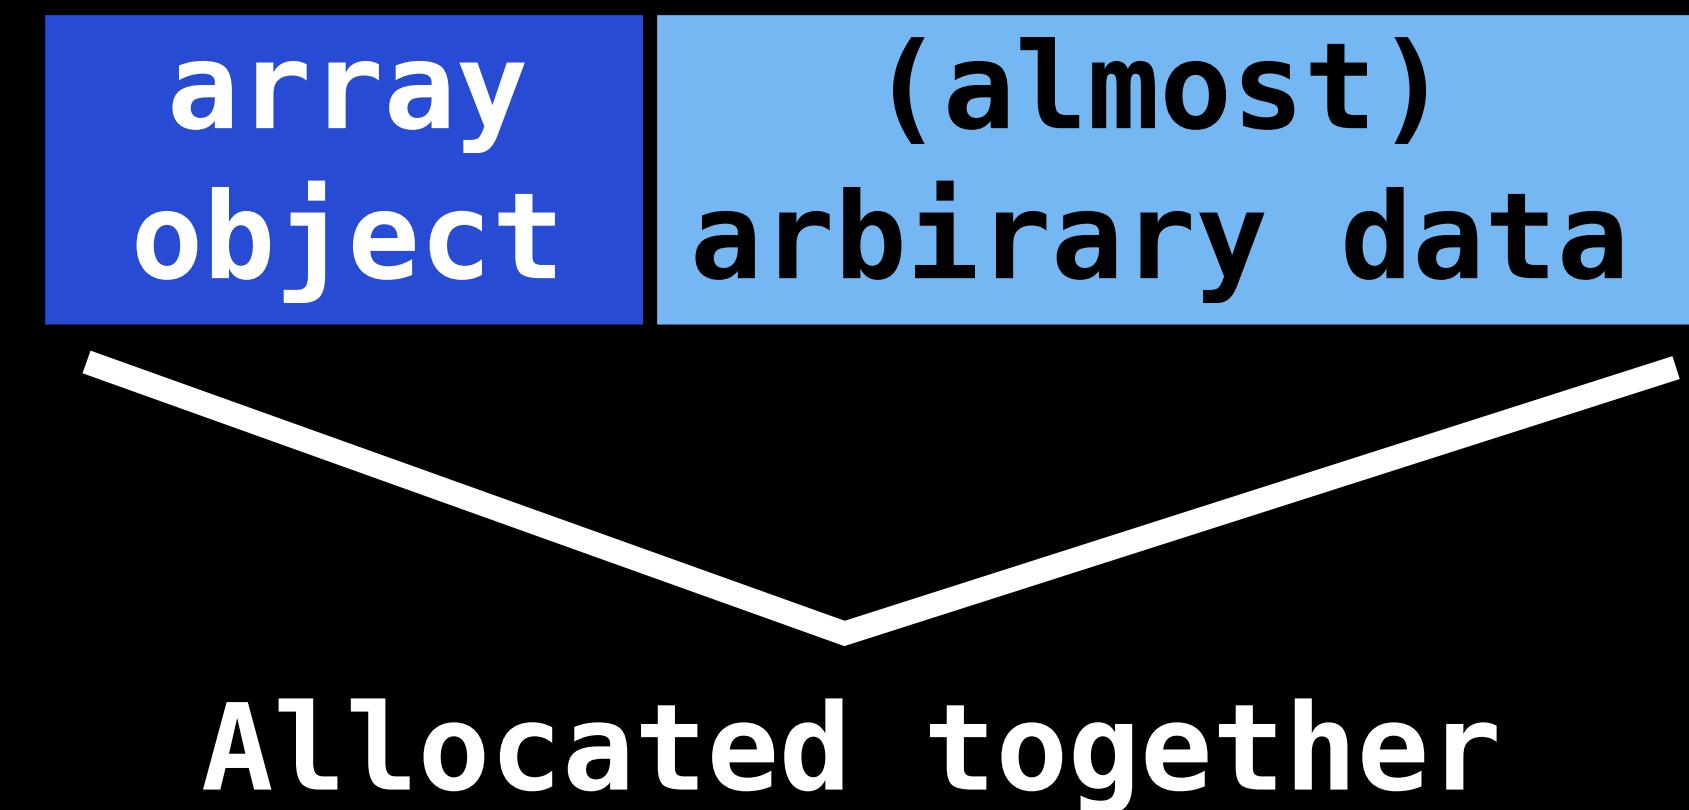


# Slab allocator for JavaScript objects

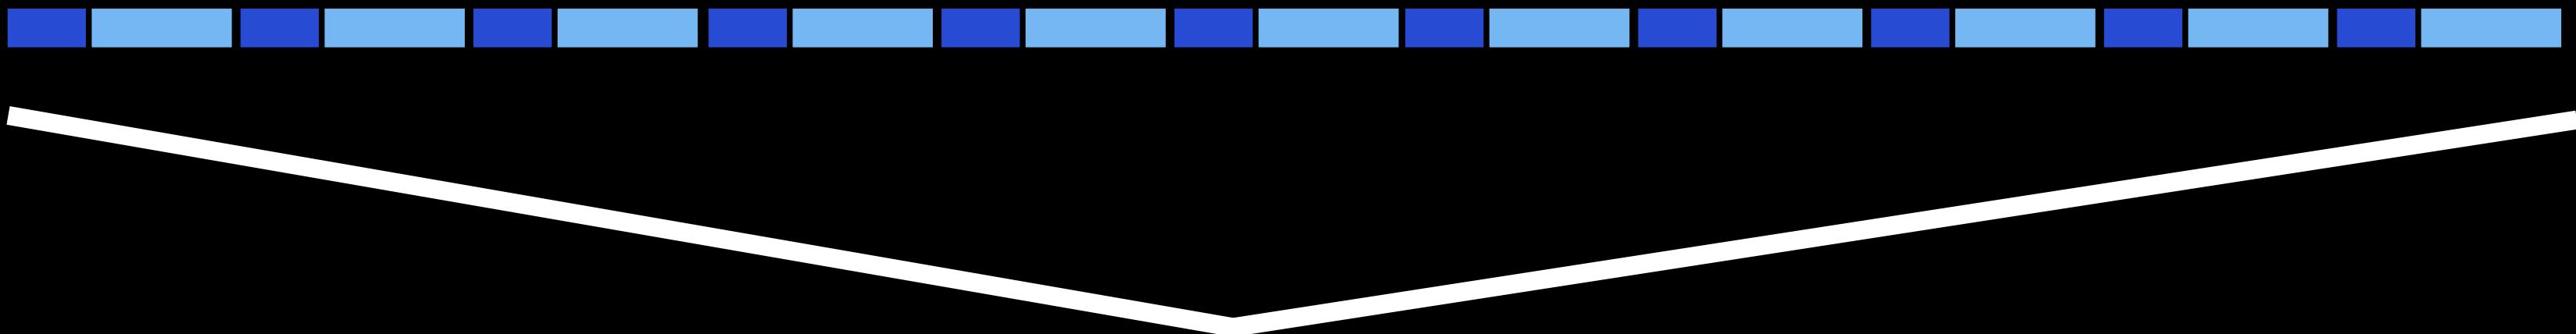


Allocated together

# Slab allocator for JavaScript objects

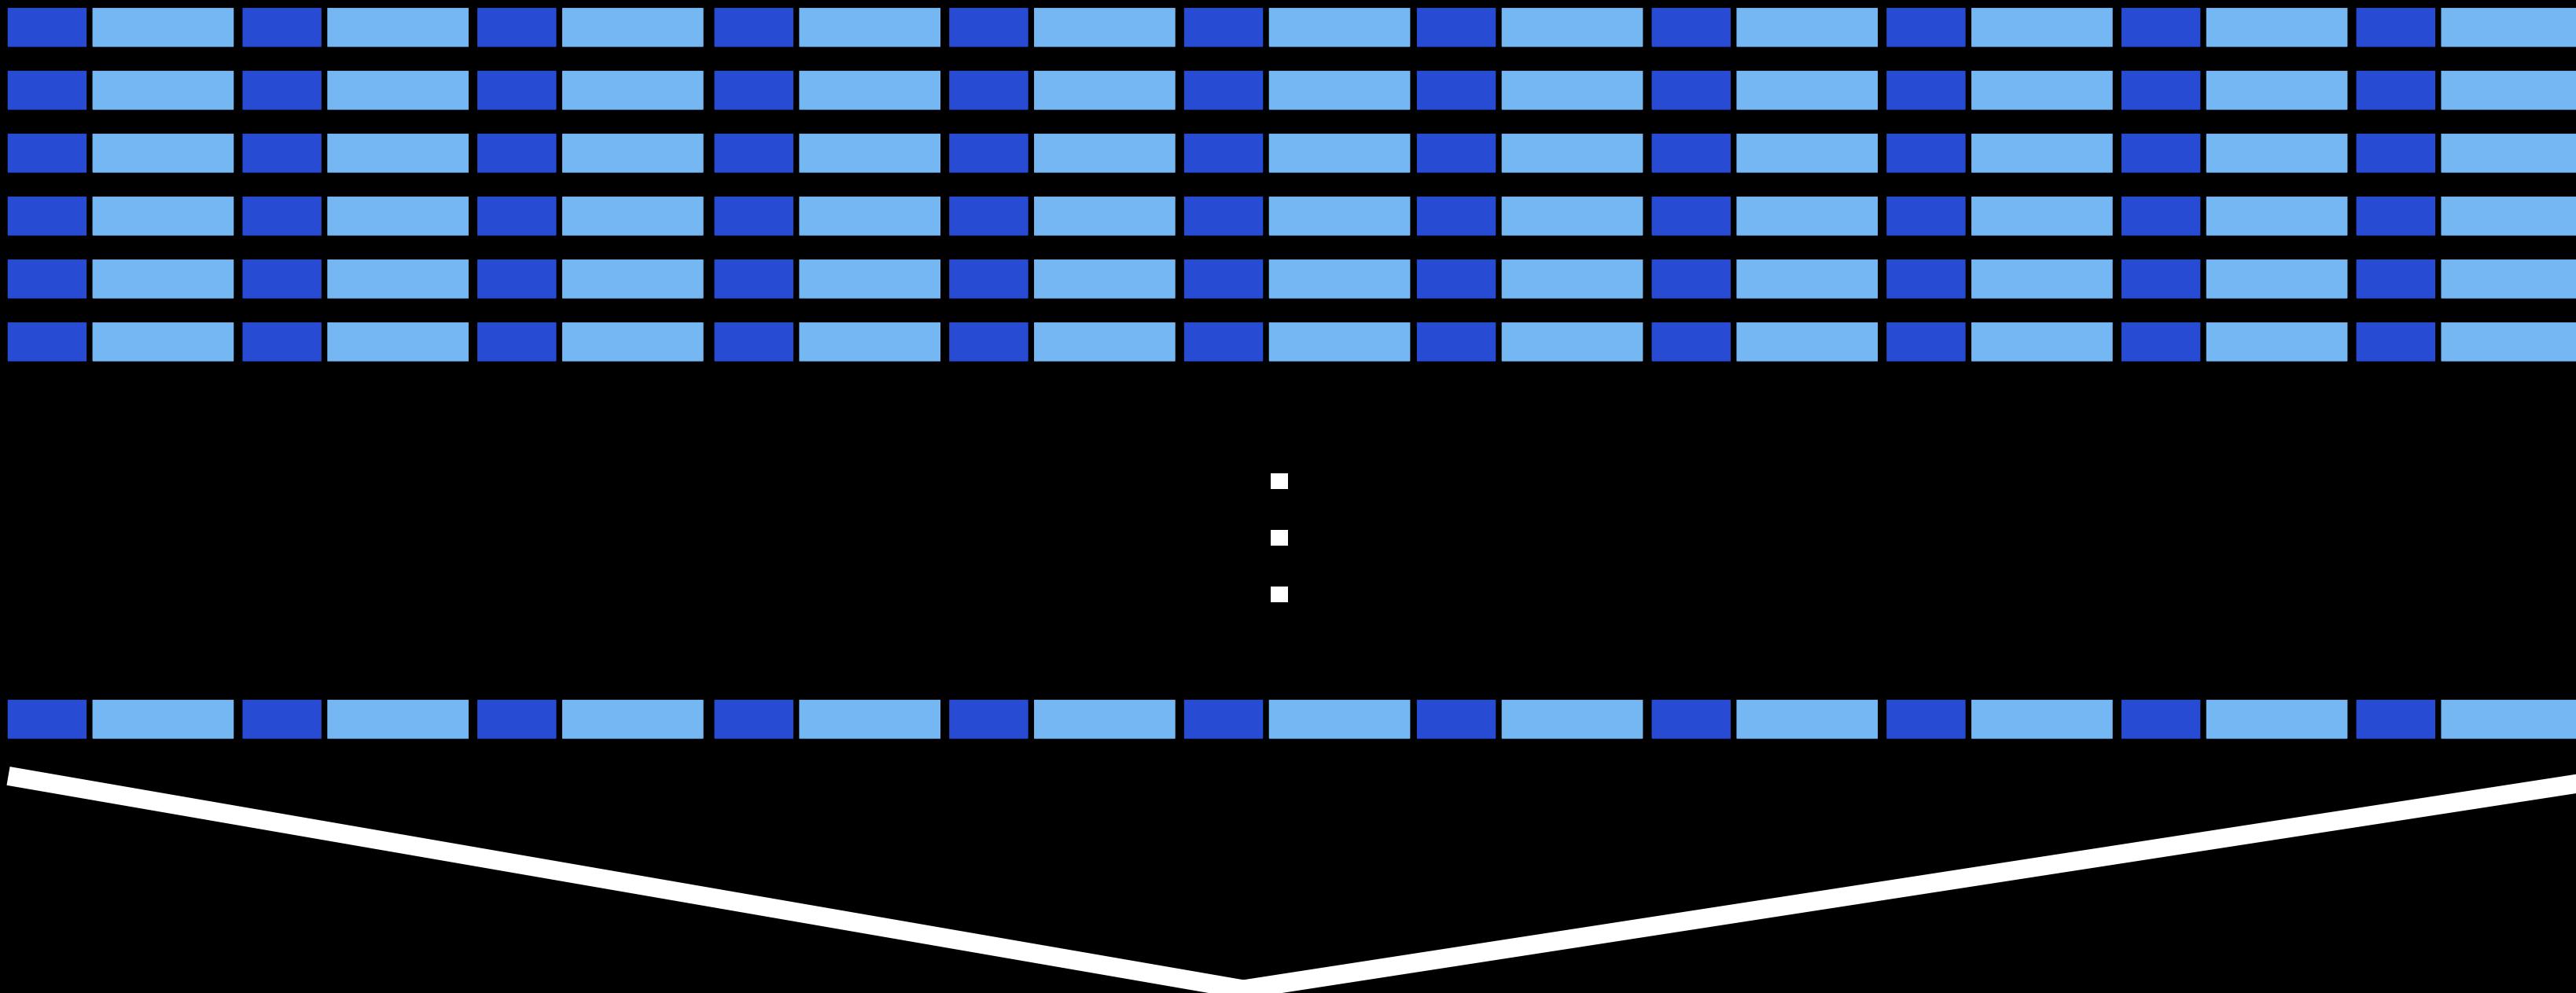


# Slab allocator for JavaScript objects



16K slab

# Slab allocator for JavaScript objects



1M VirtualAlloc()

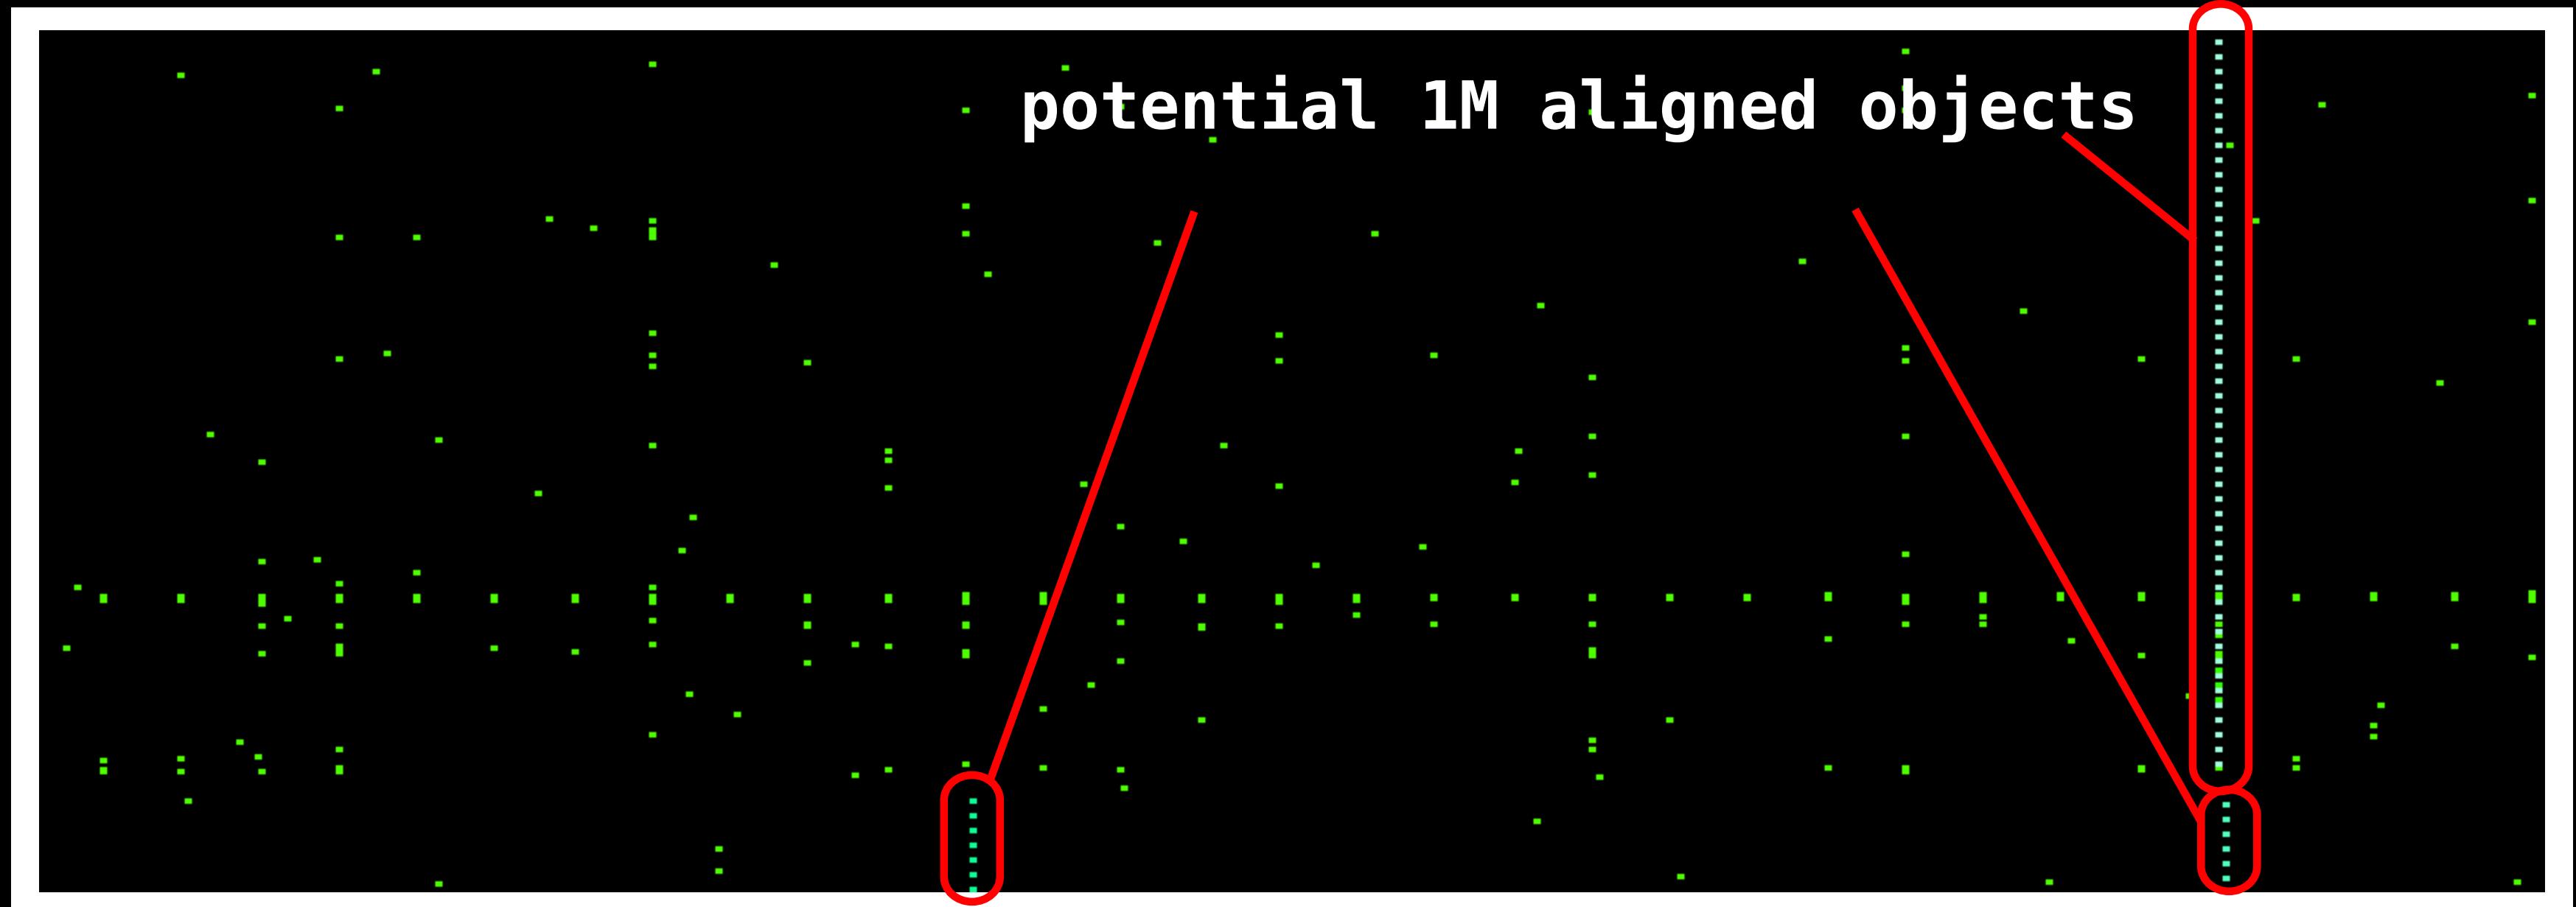
# Slab allocator for JavaScript objects

1st after `VirtualAlloc()` call

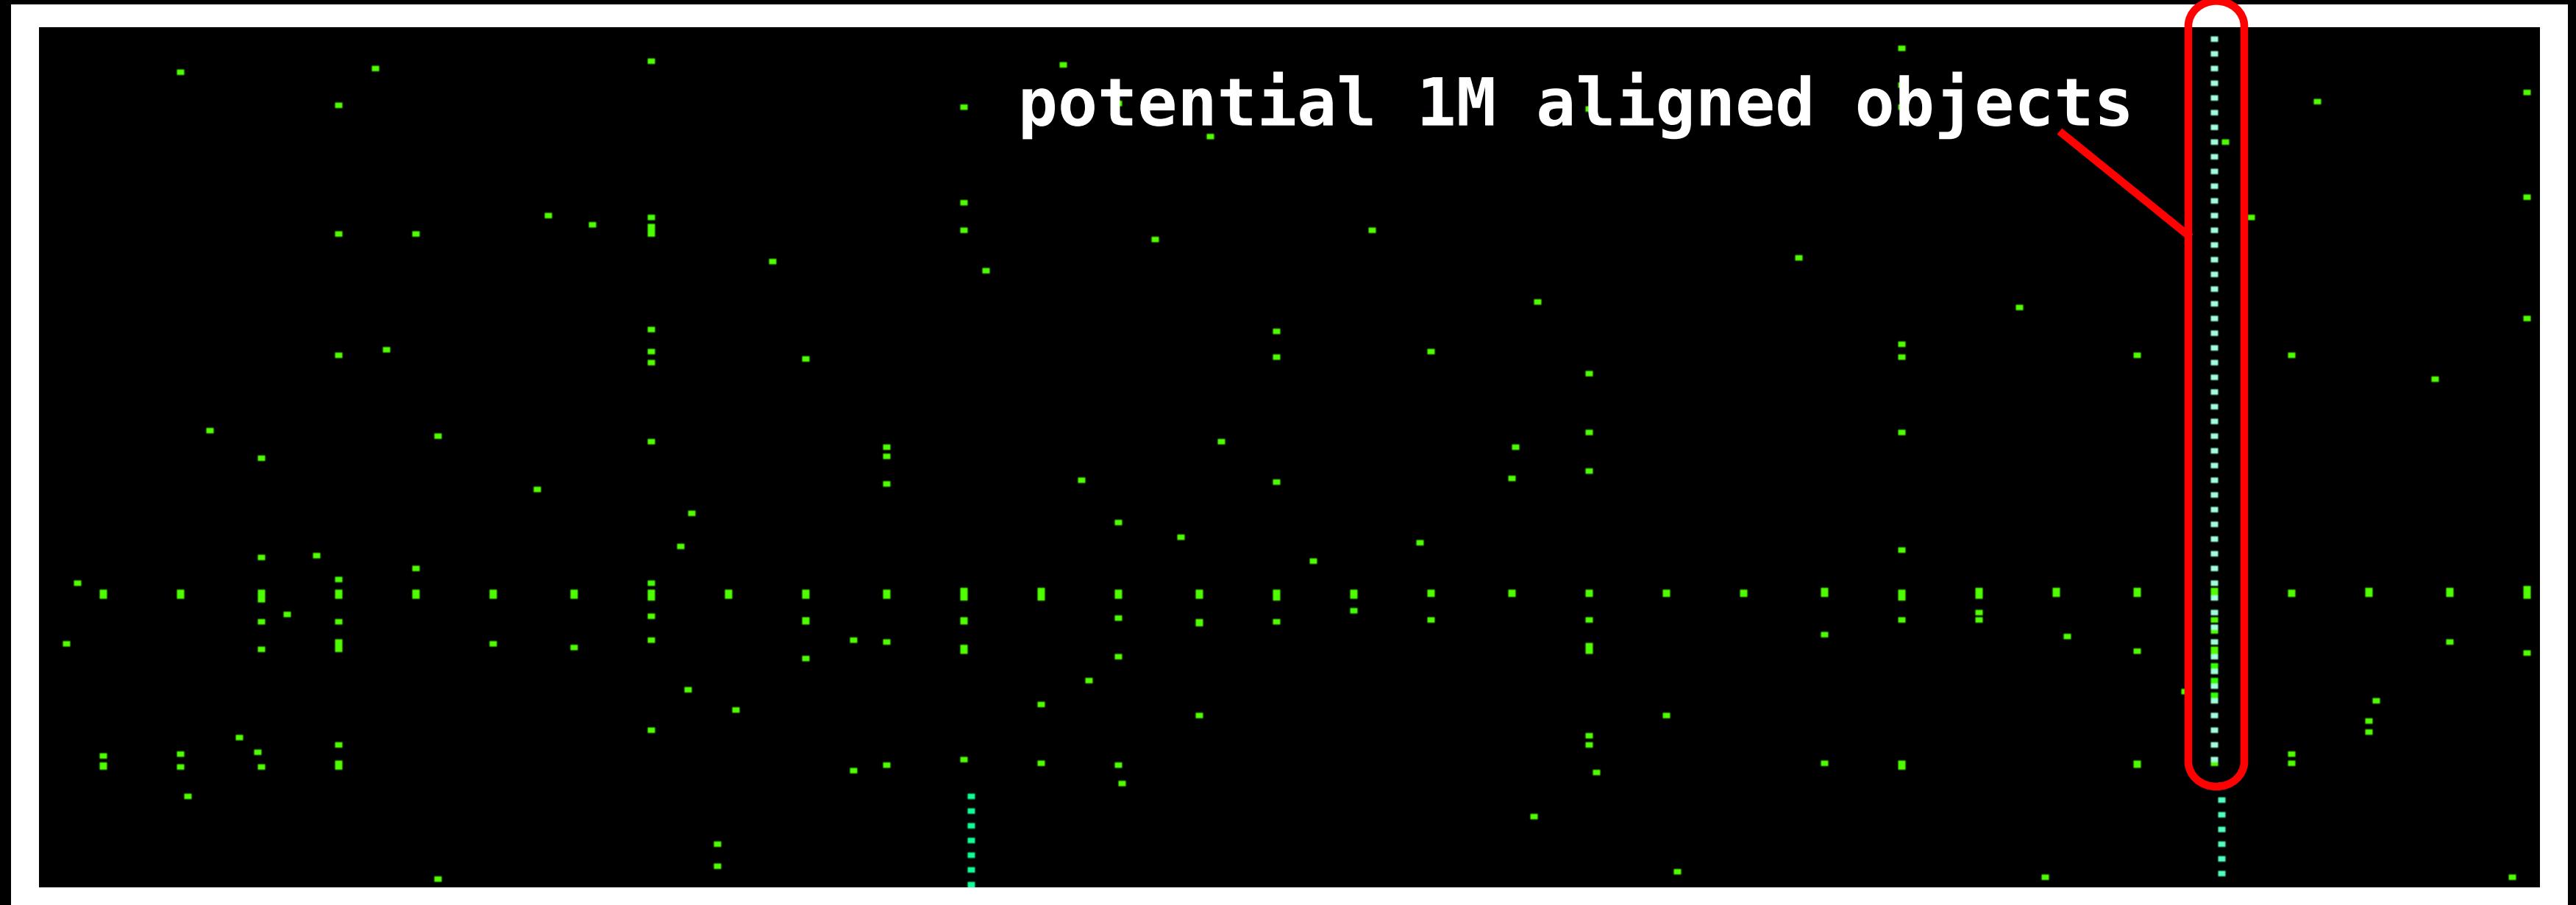


1M `VirtualAlloc()`

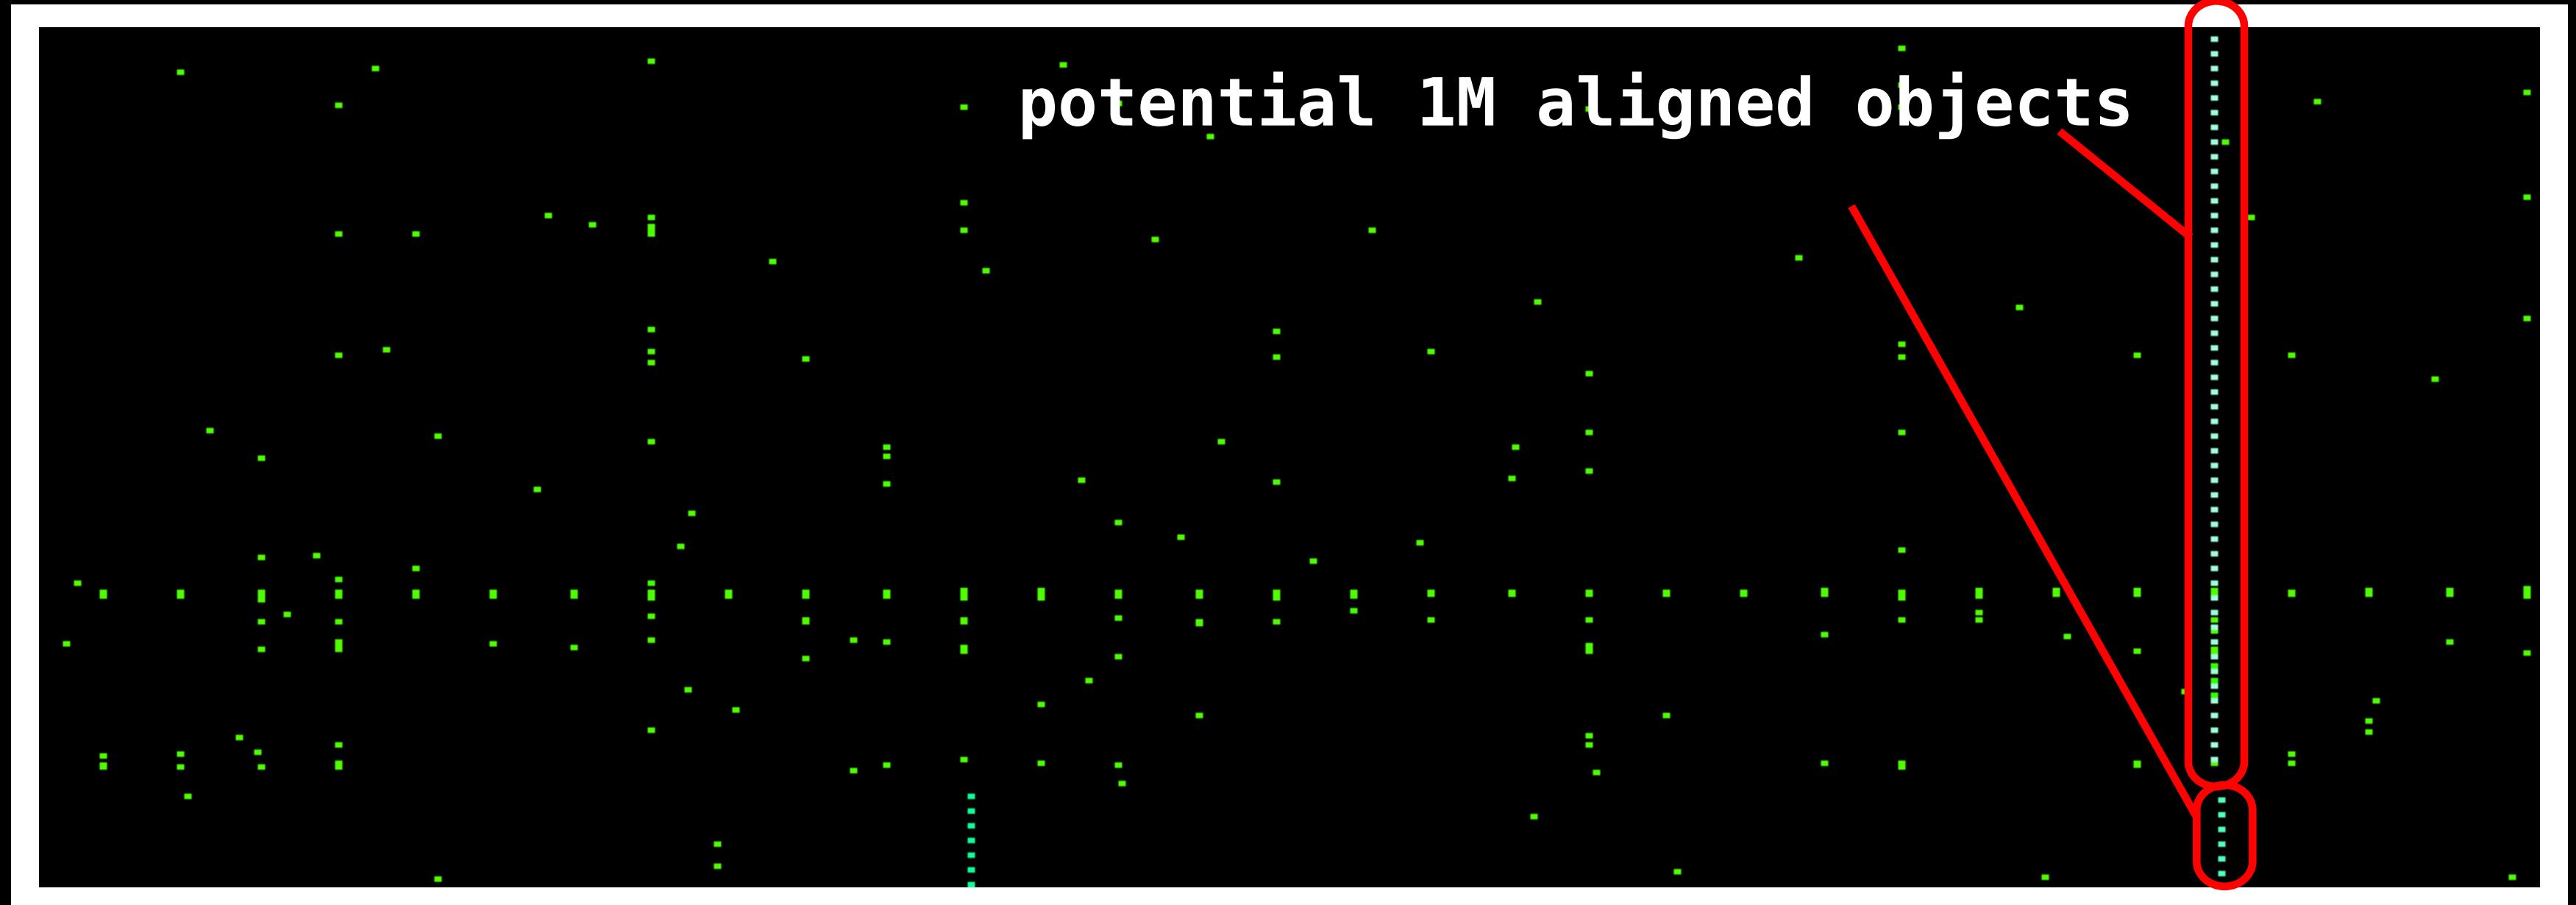
# Slab allocator for JavaScript objects



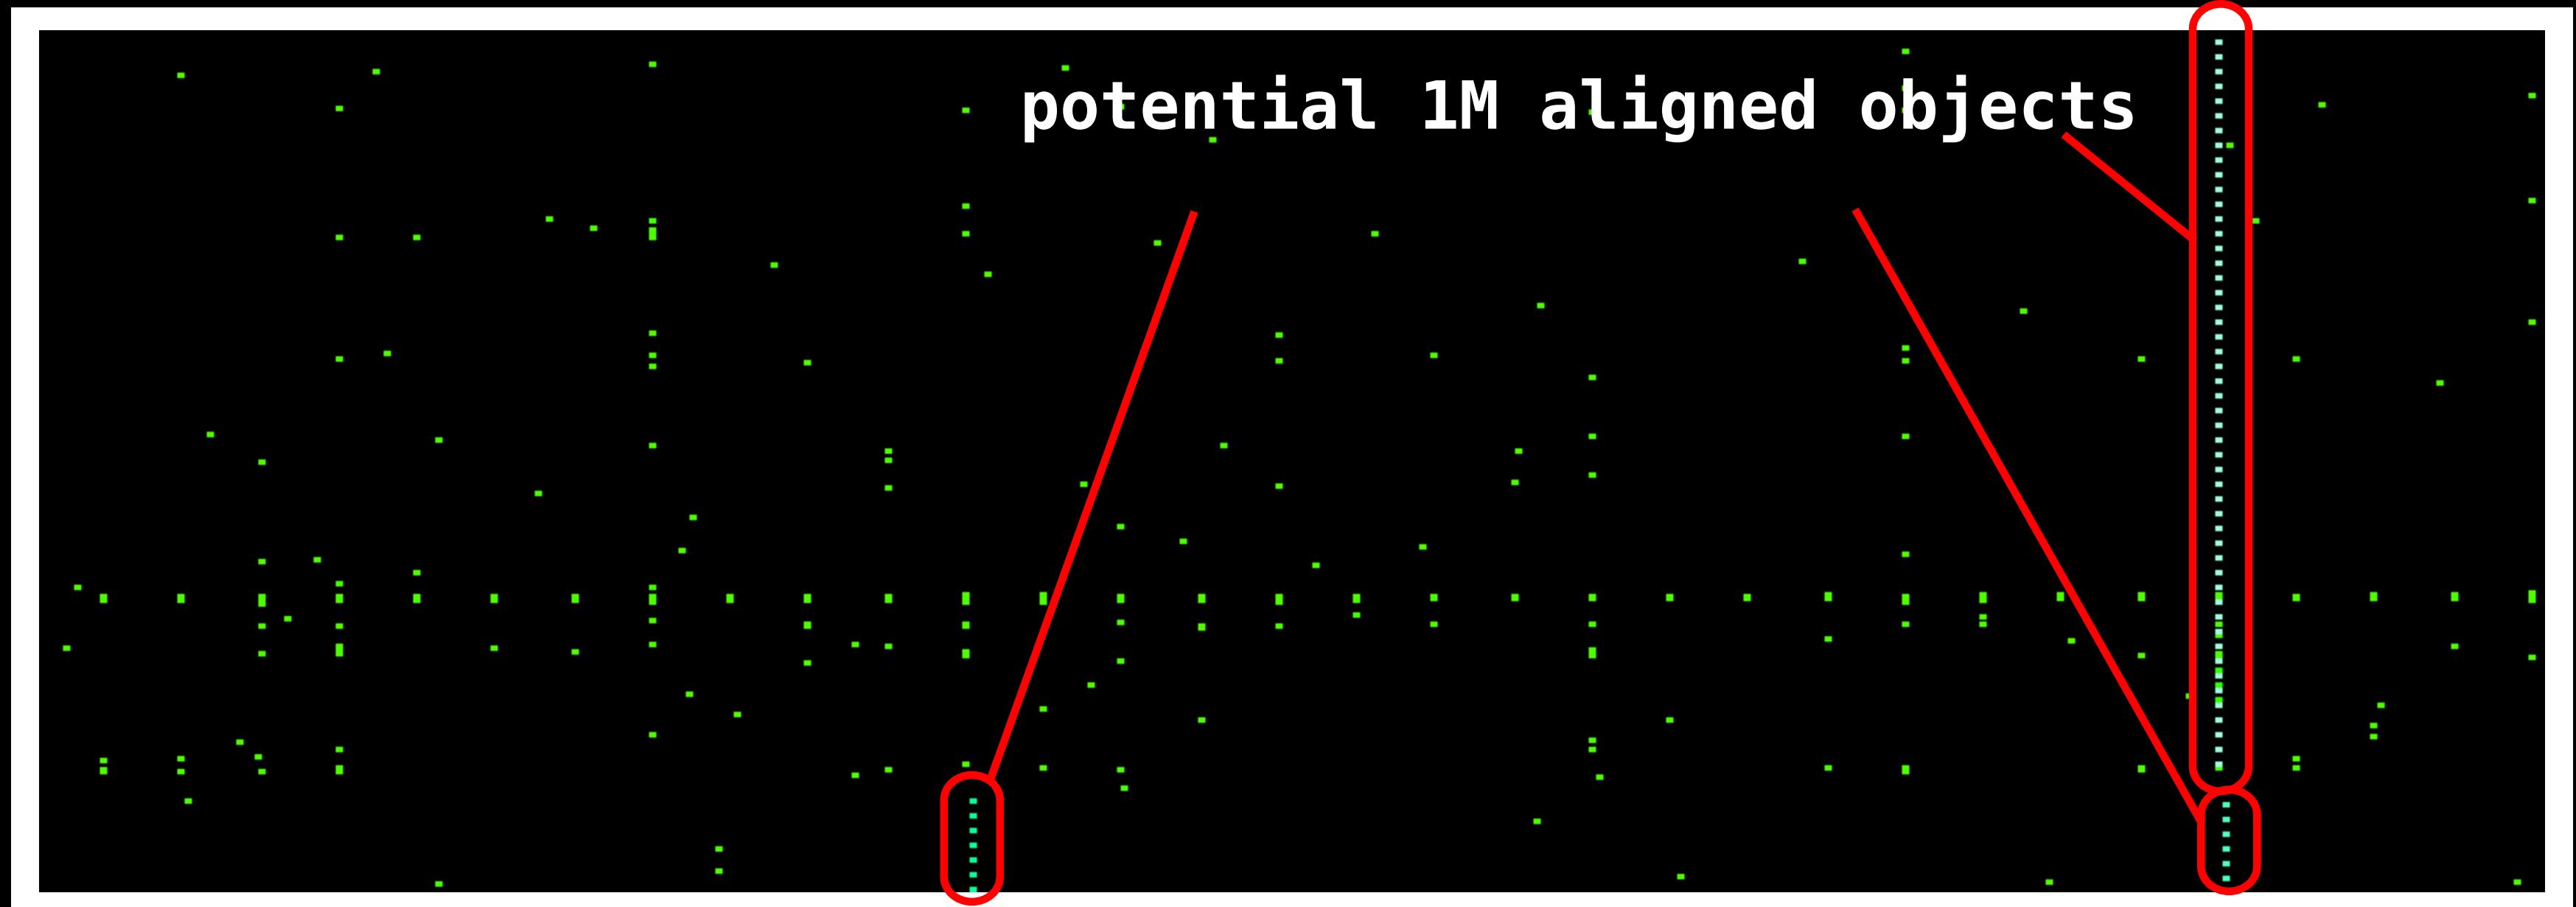
# Slab allocator for JavaScript objects



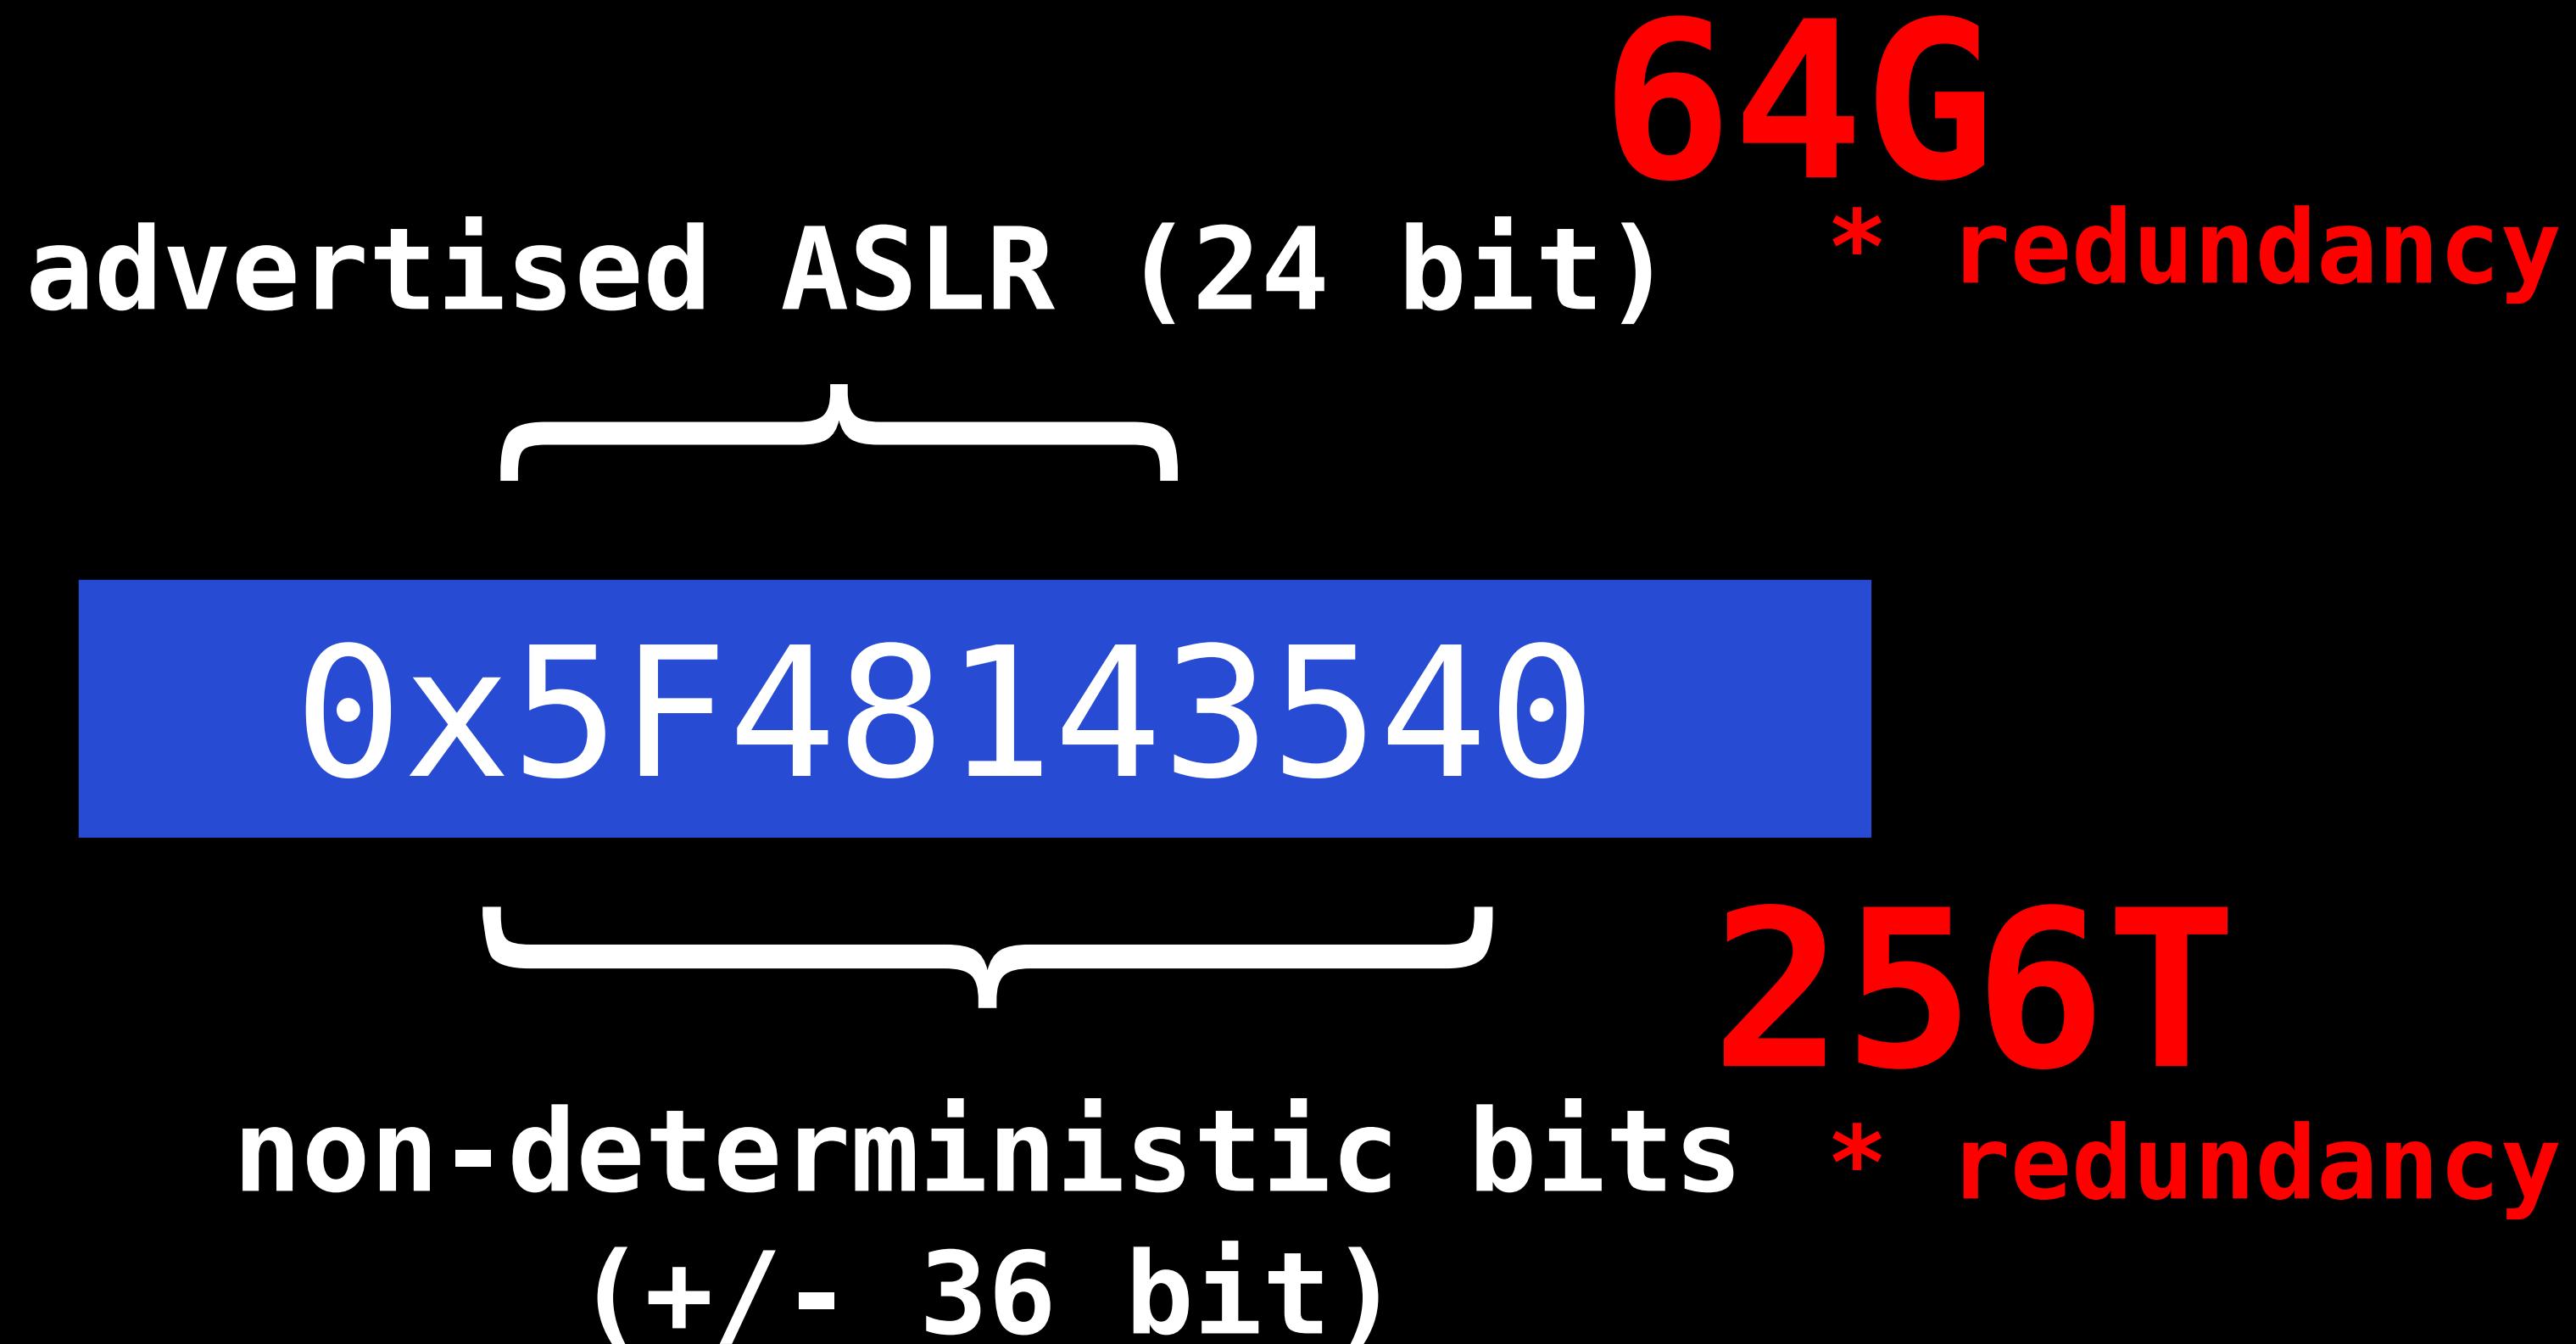
# Slab allocator for JavaScript objects



# Slab allocator for JavaScript objects



# Heap pointer entropy in Edge



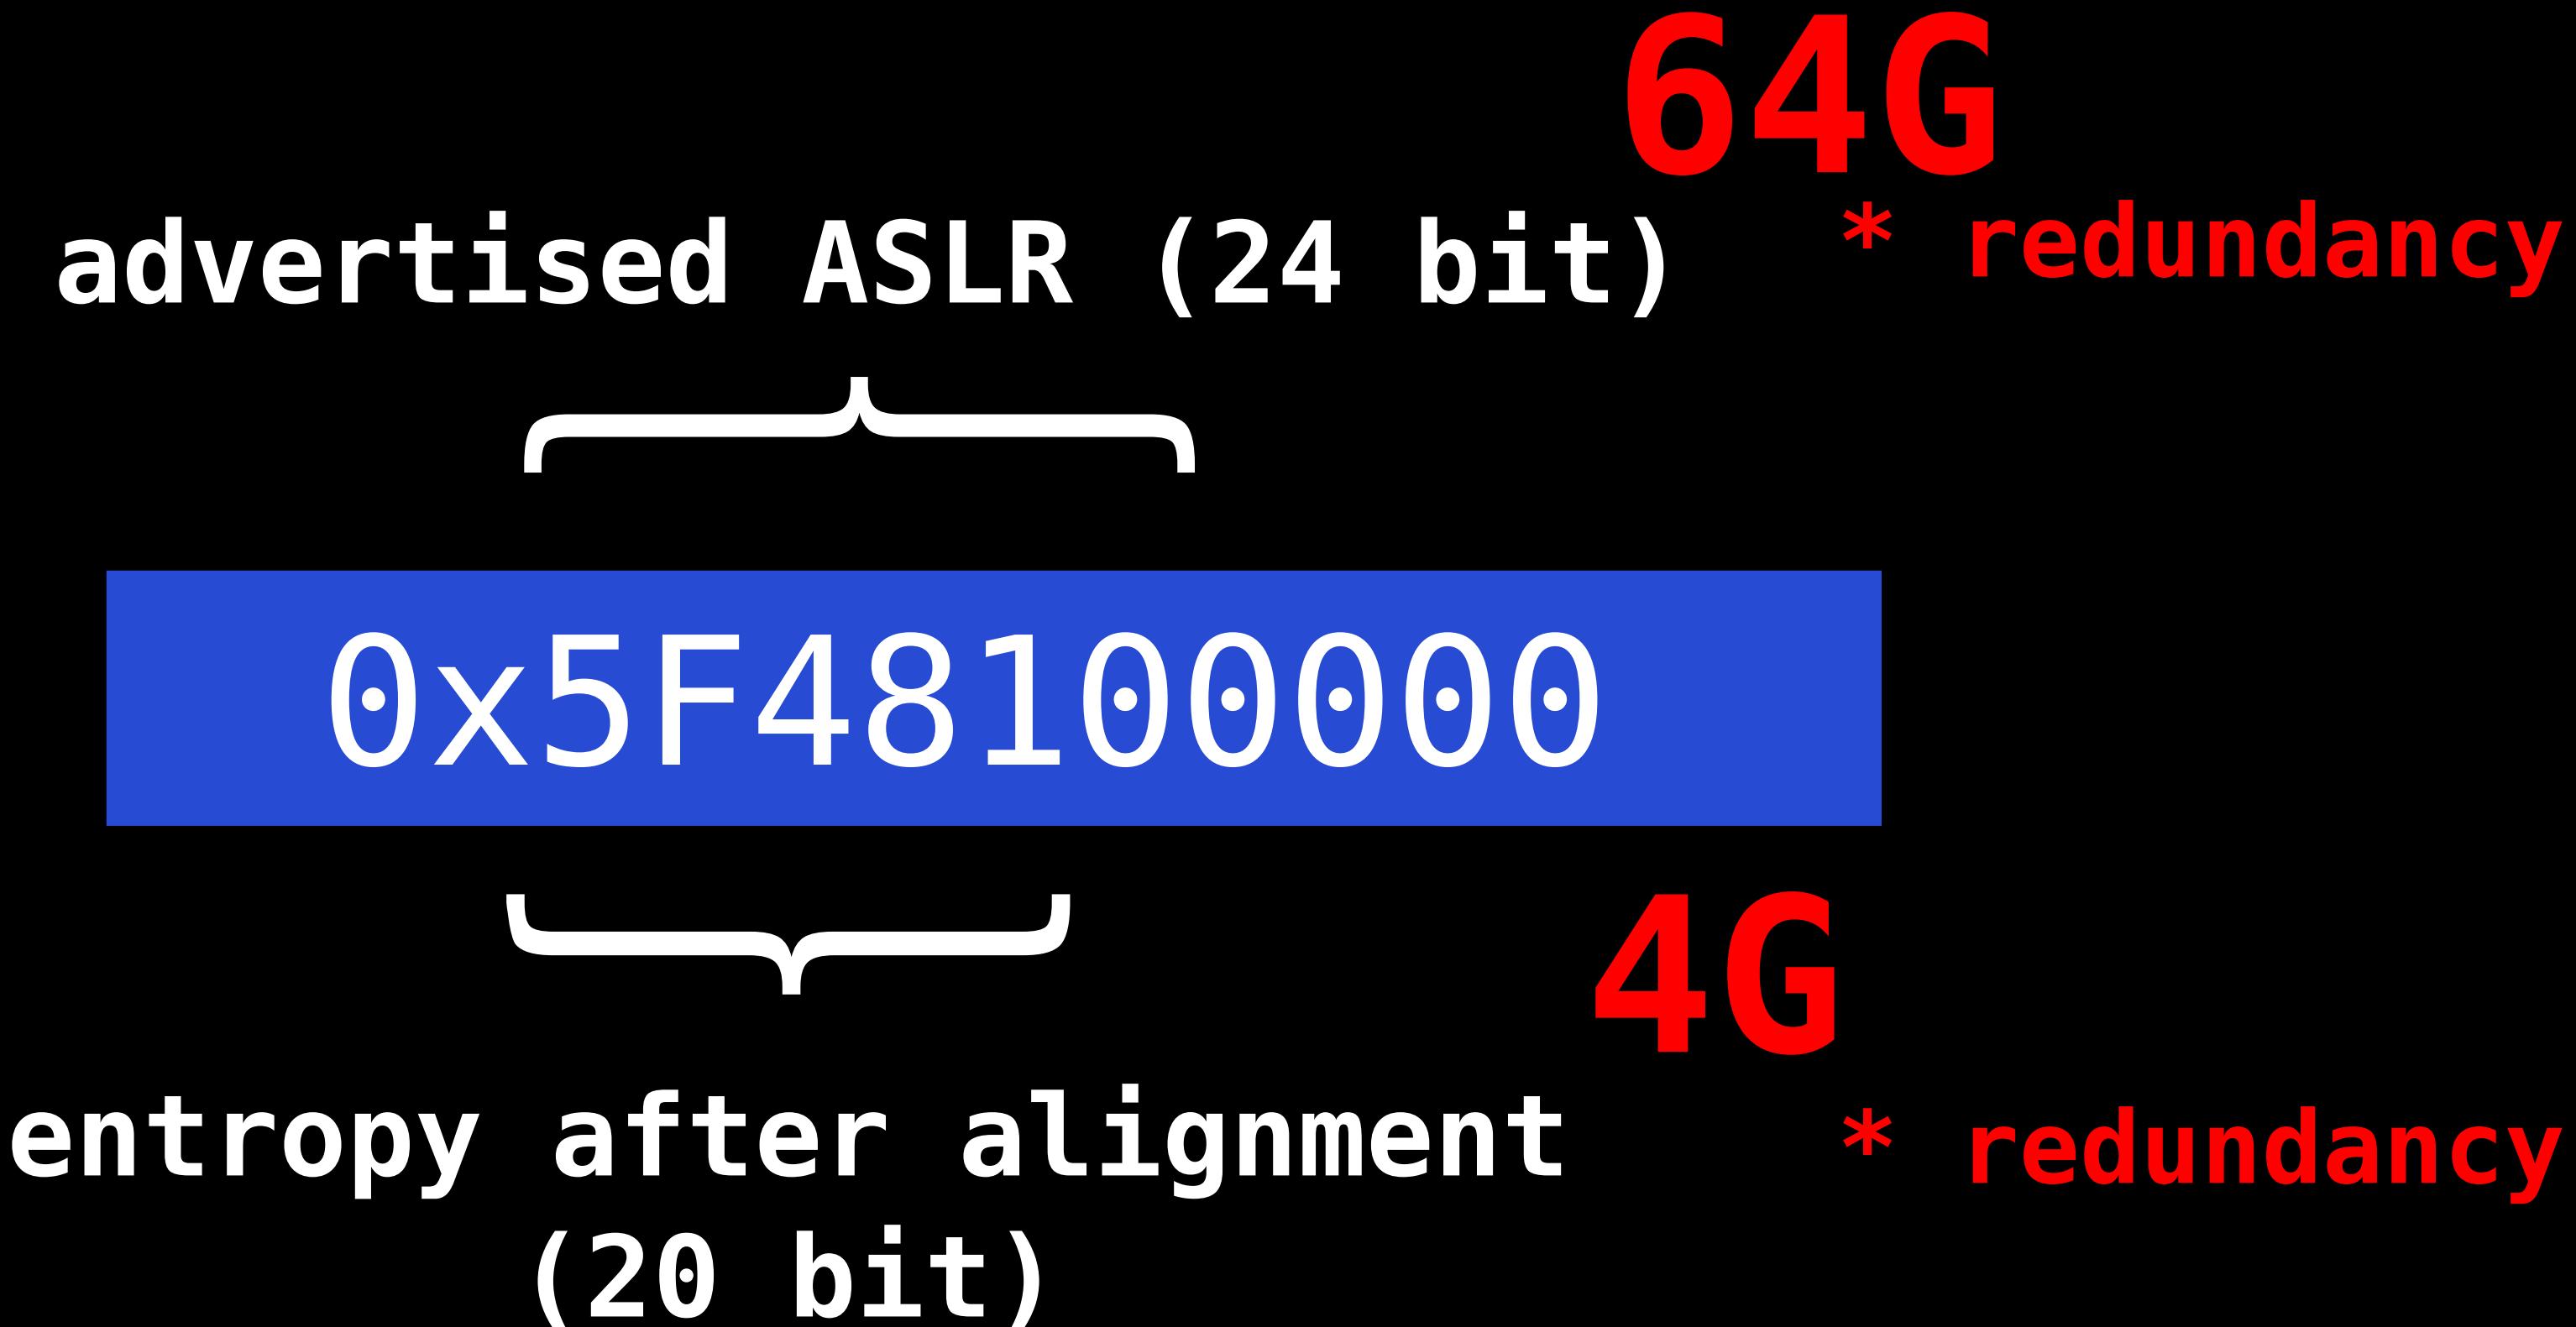
# Heap pointer entropy in Edge

64G  
advertised ASLR (24 bit) \* redundancy

0x5F48100000

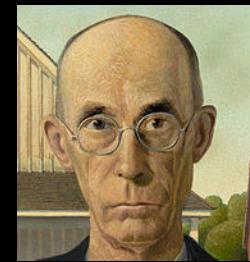
entropy after alignment  
(20 bit)

# Heap pointer entropy in Edge

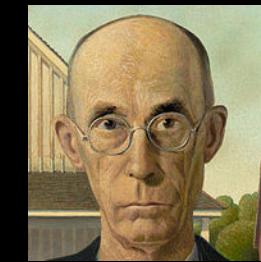


# birthday problem

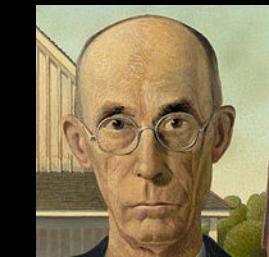
# birthday problem



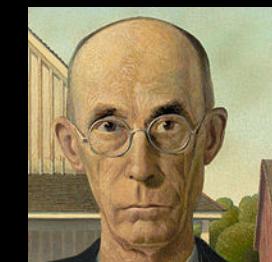
# birthday problem



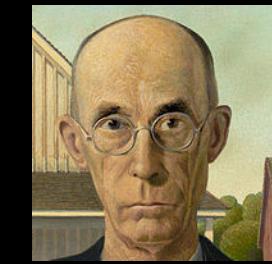
# birthday problem



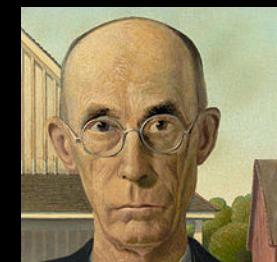
# birthday problem



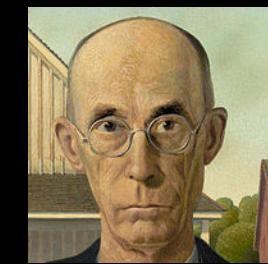
# birthday problem



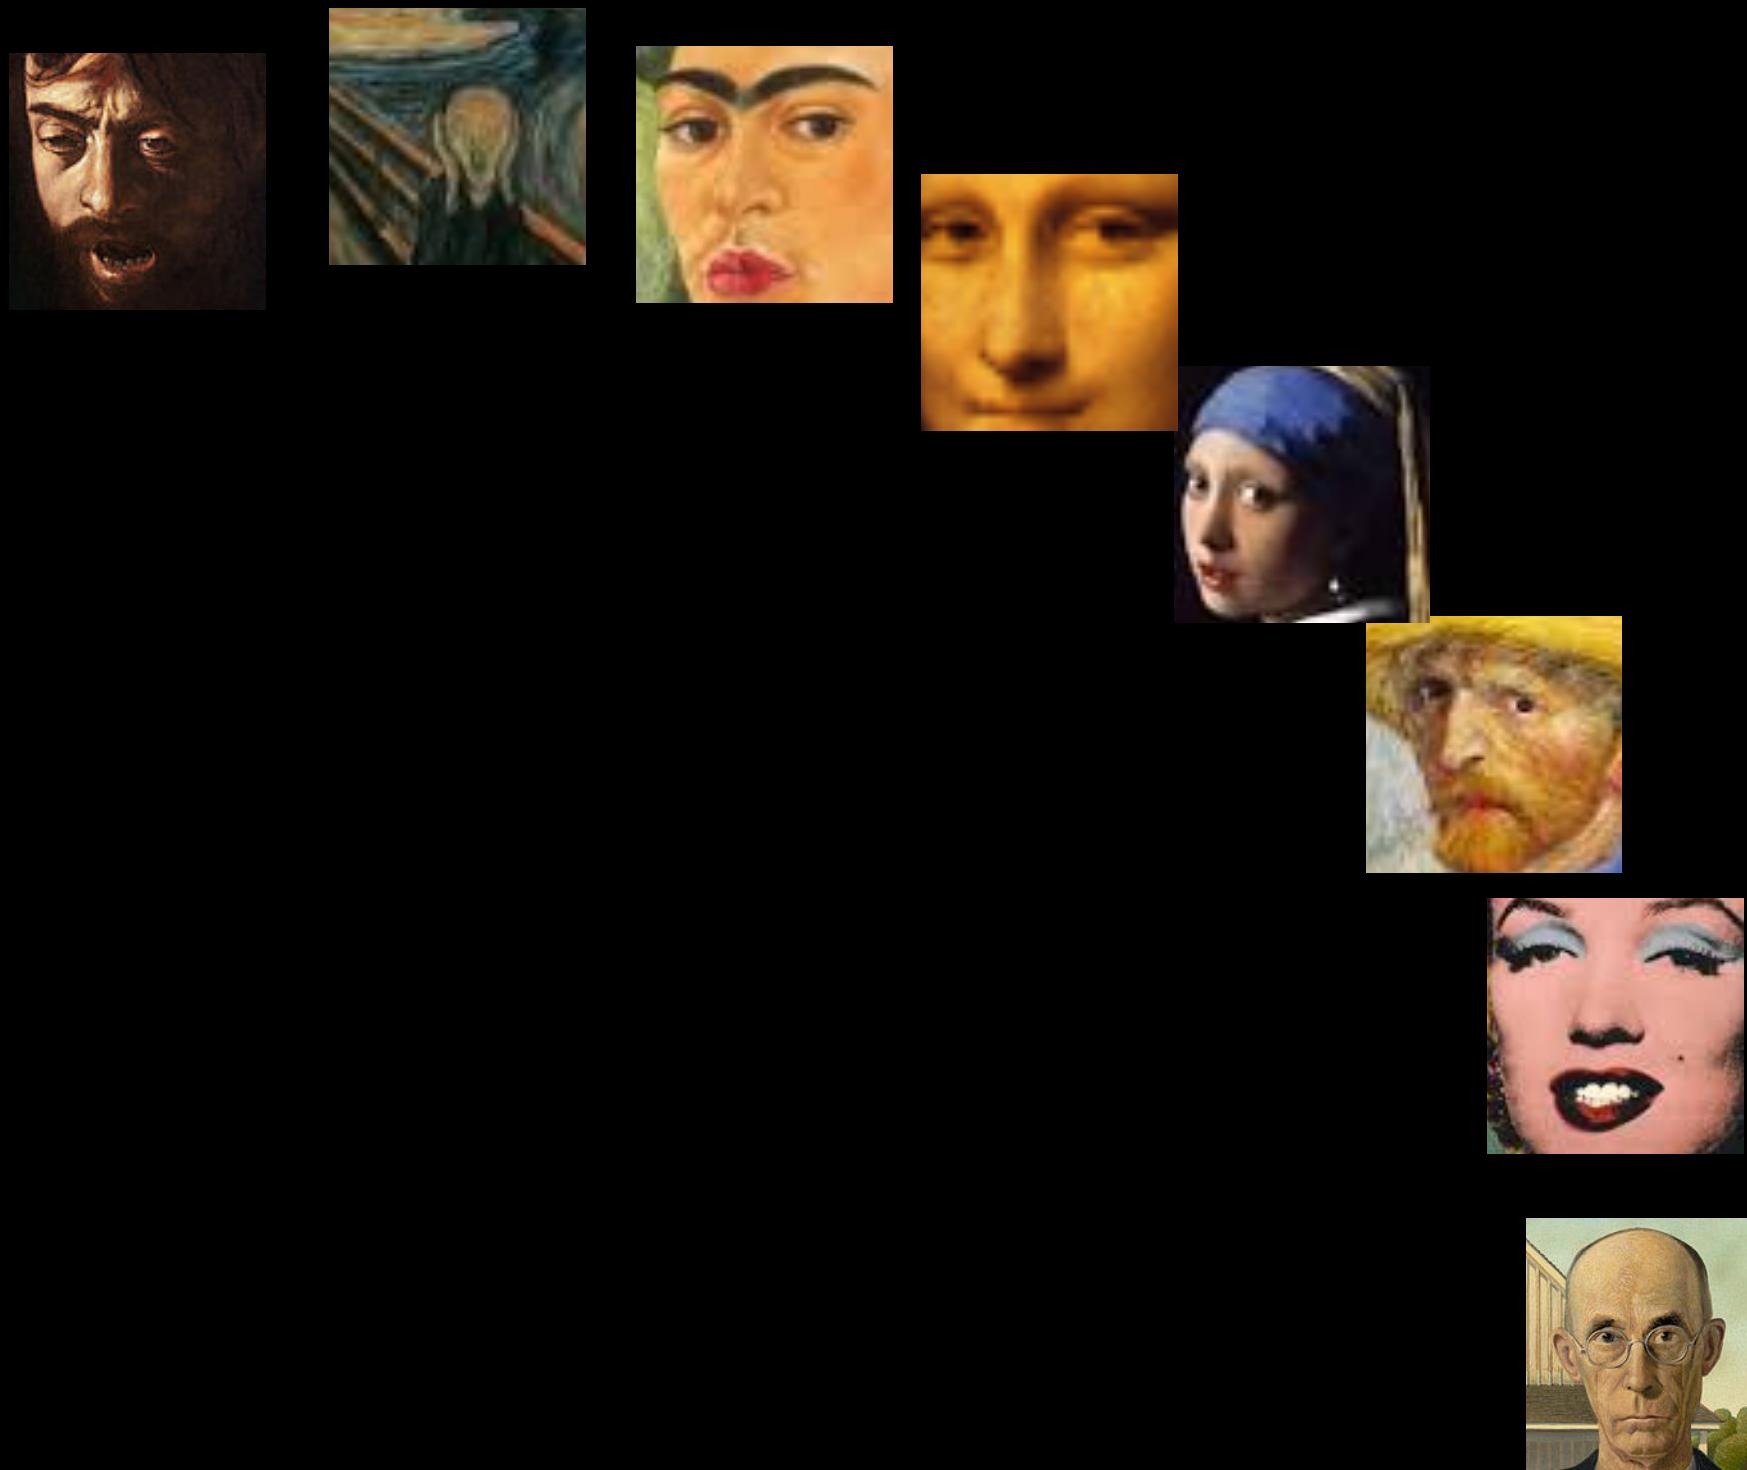
# birthday problem



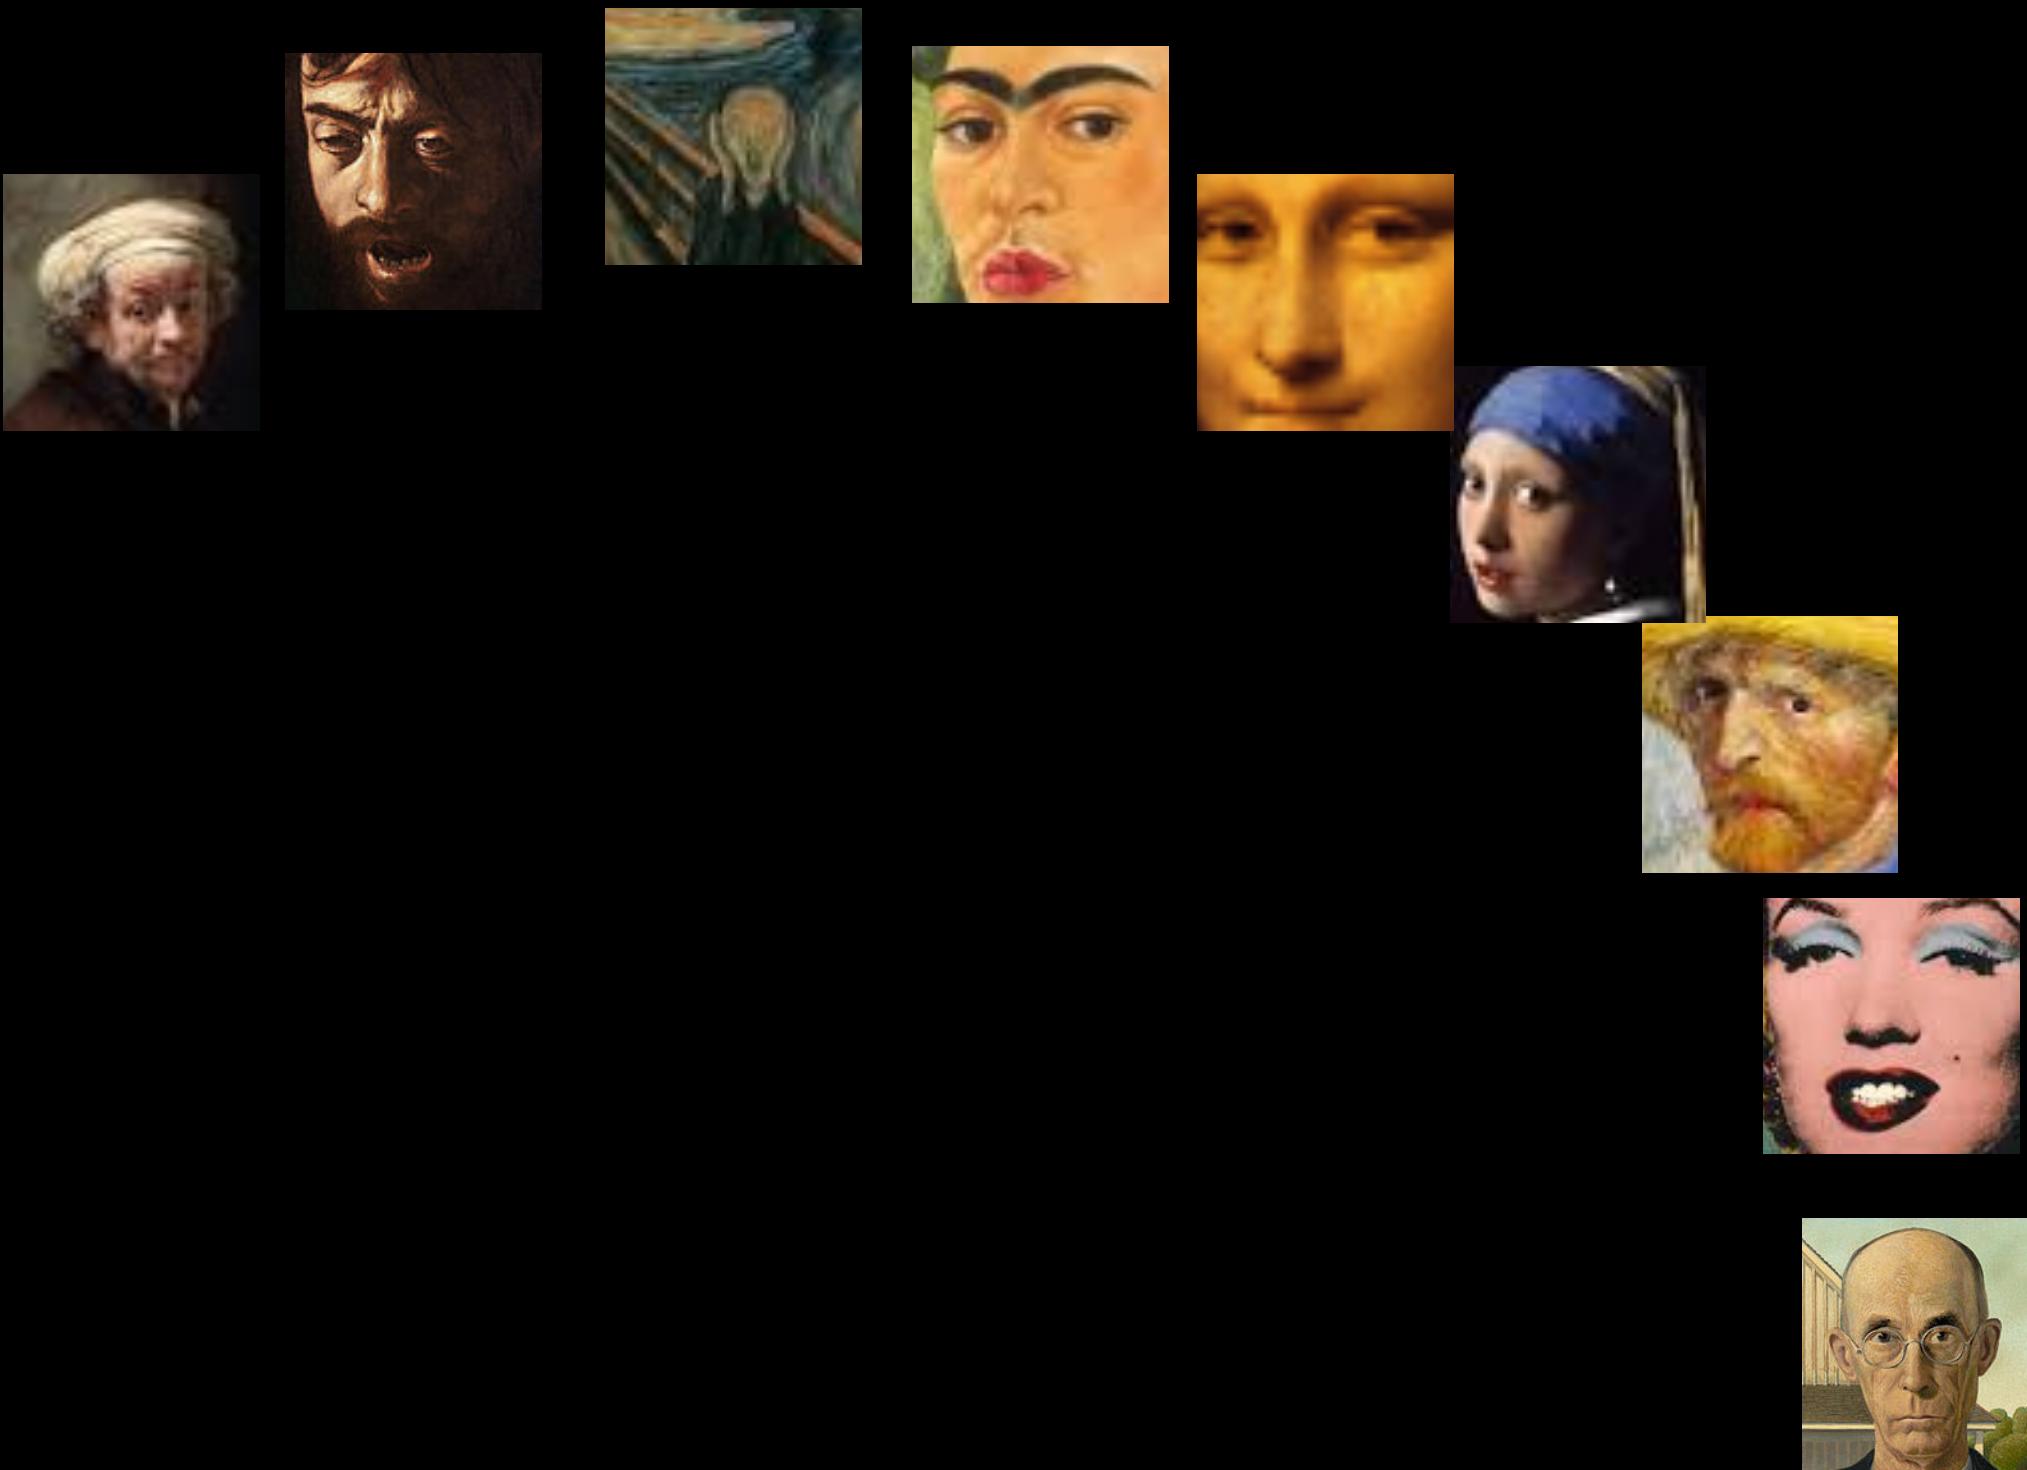
# birthday problem



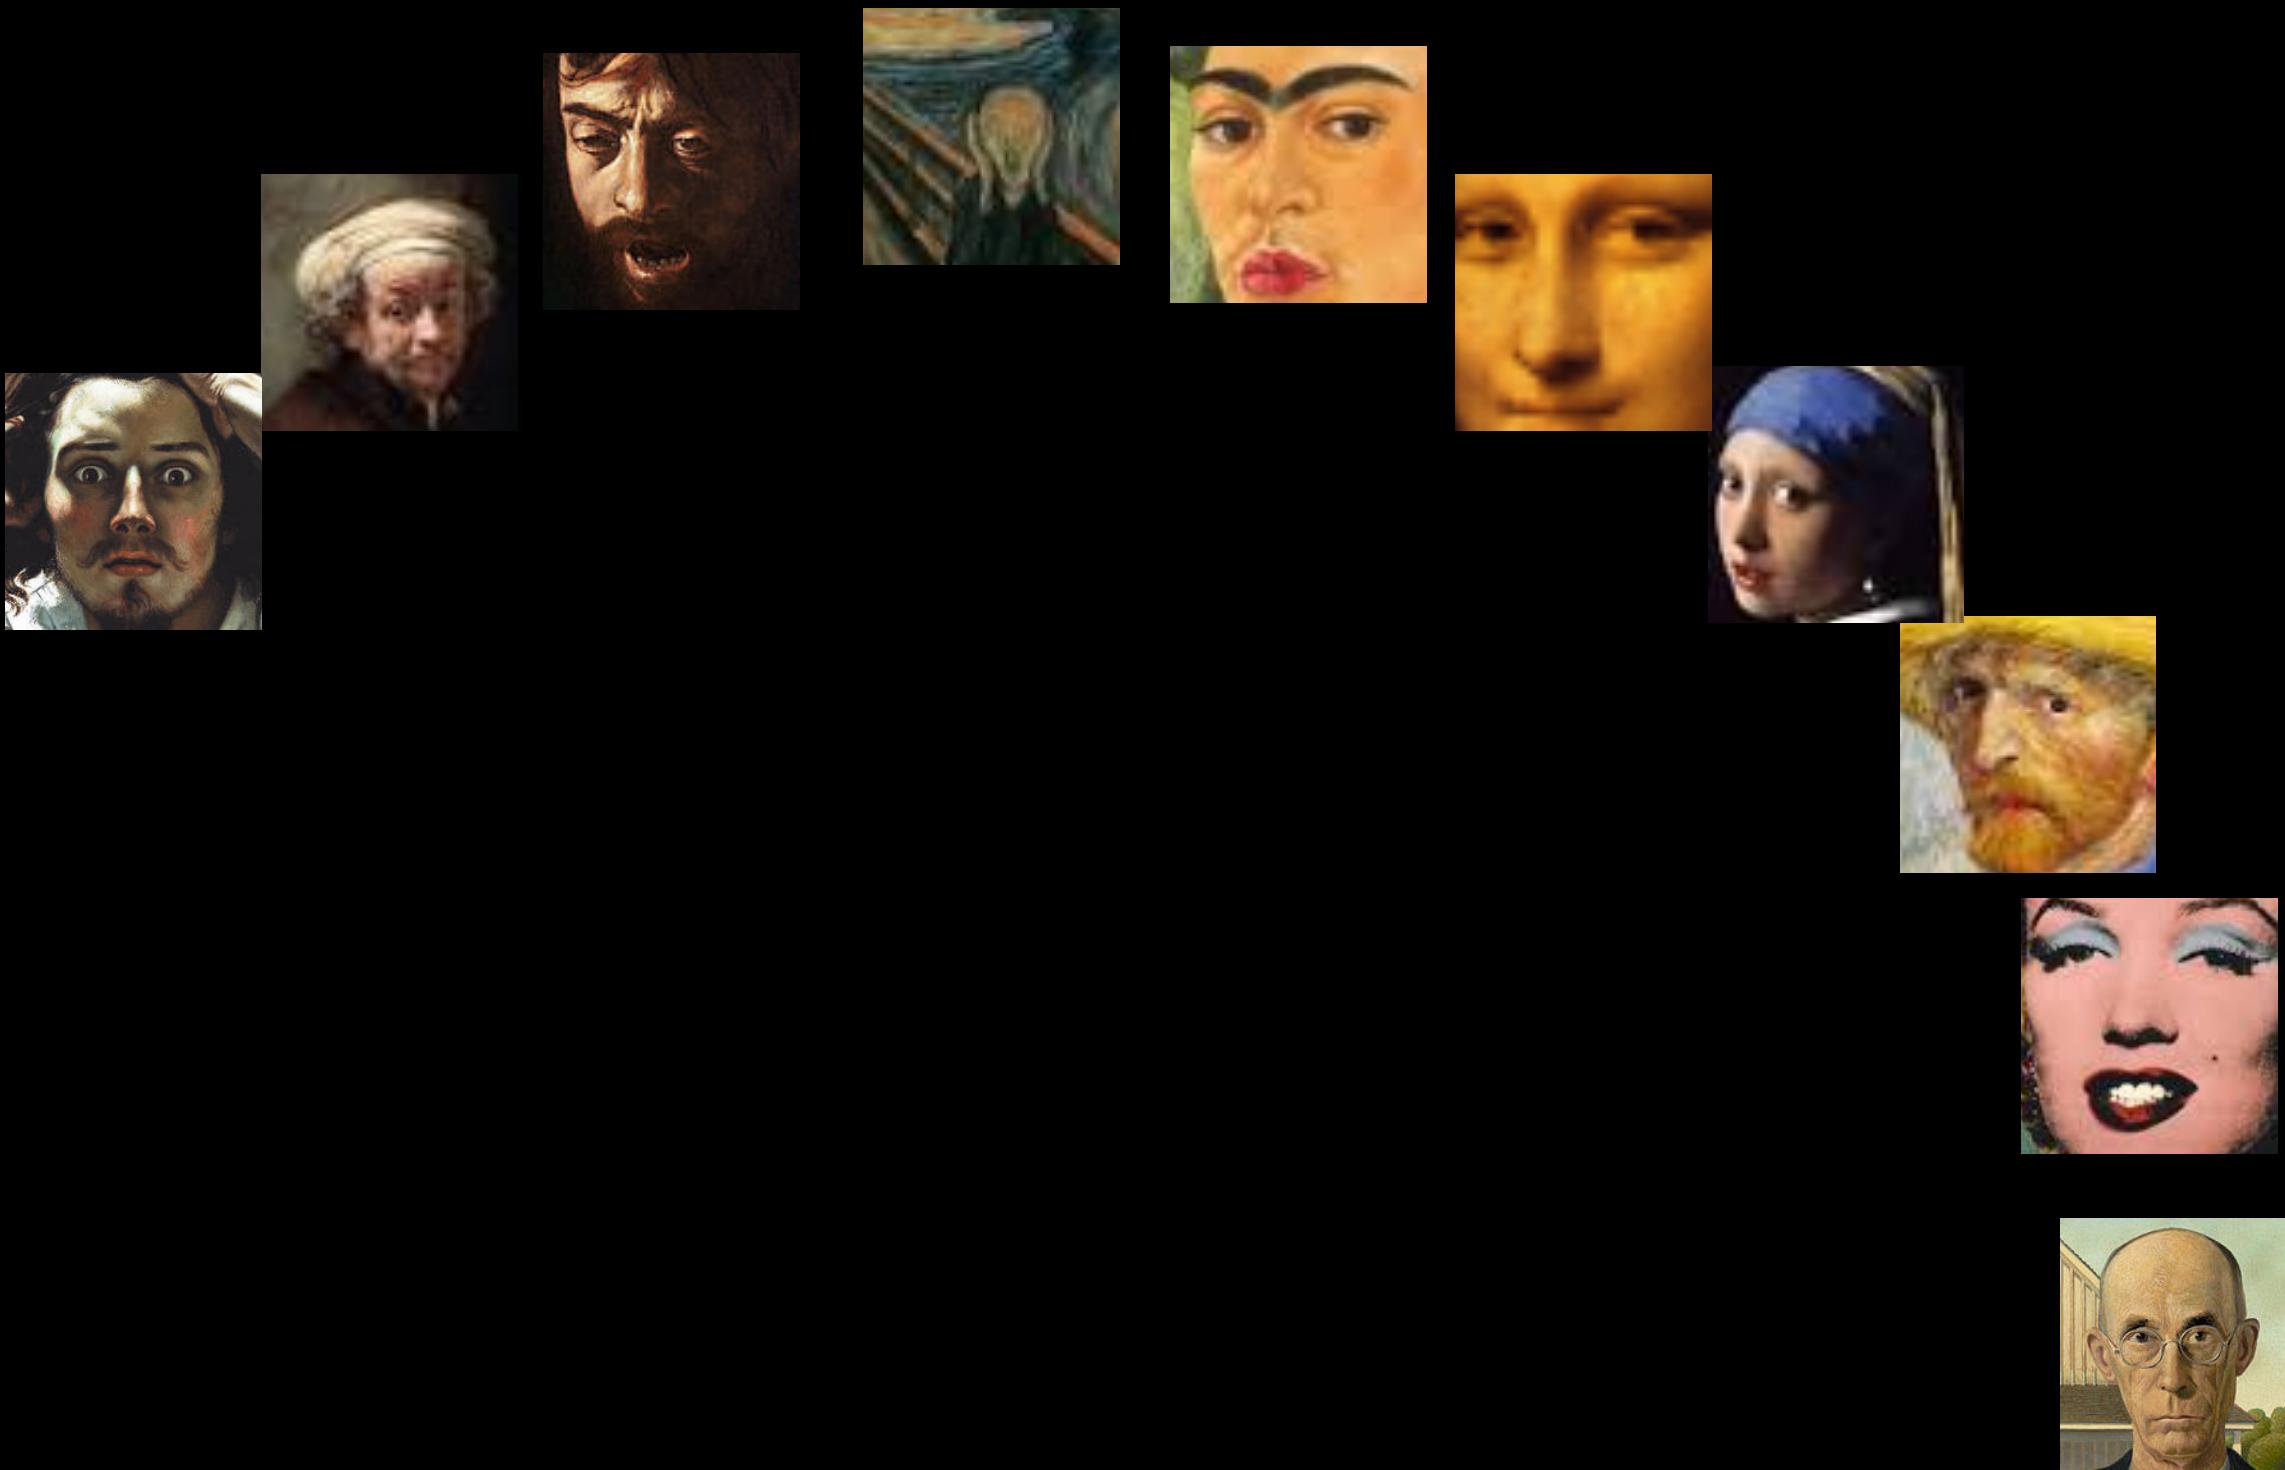
# birthday problem



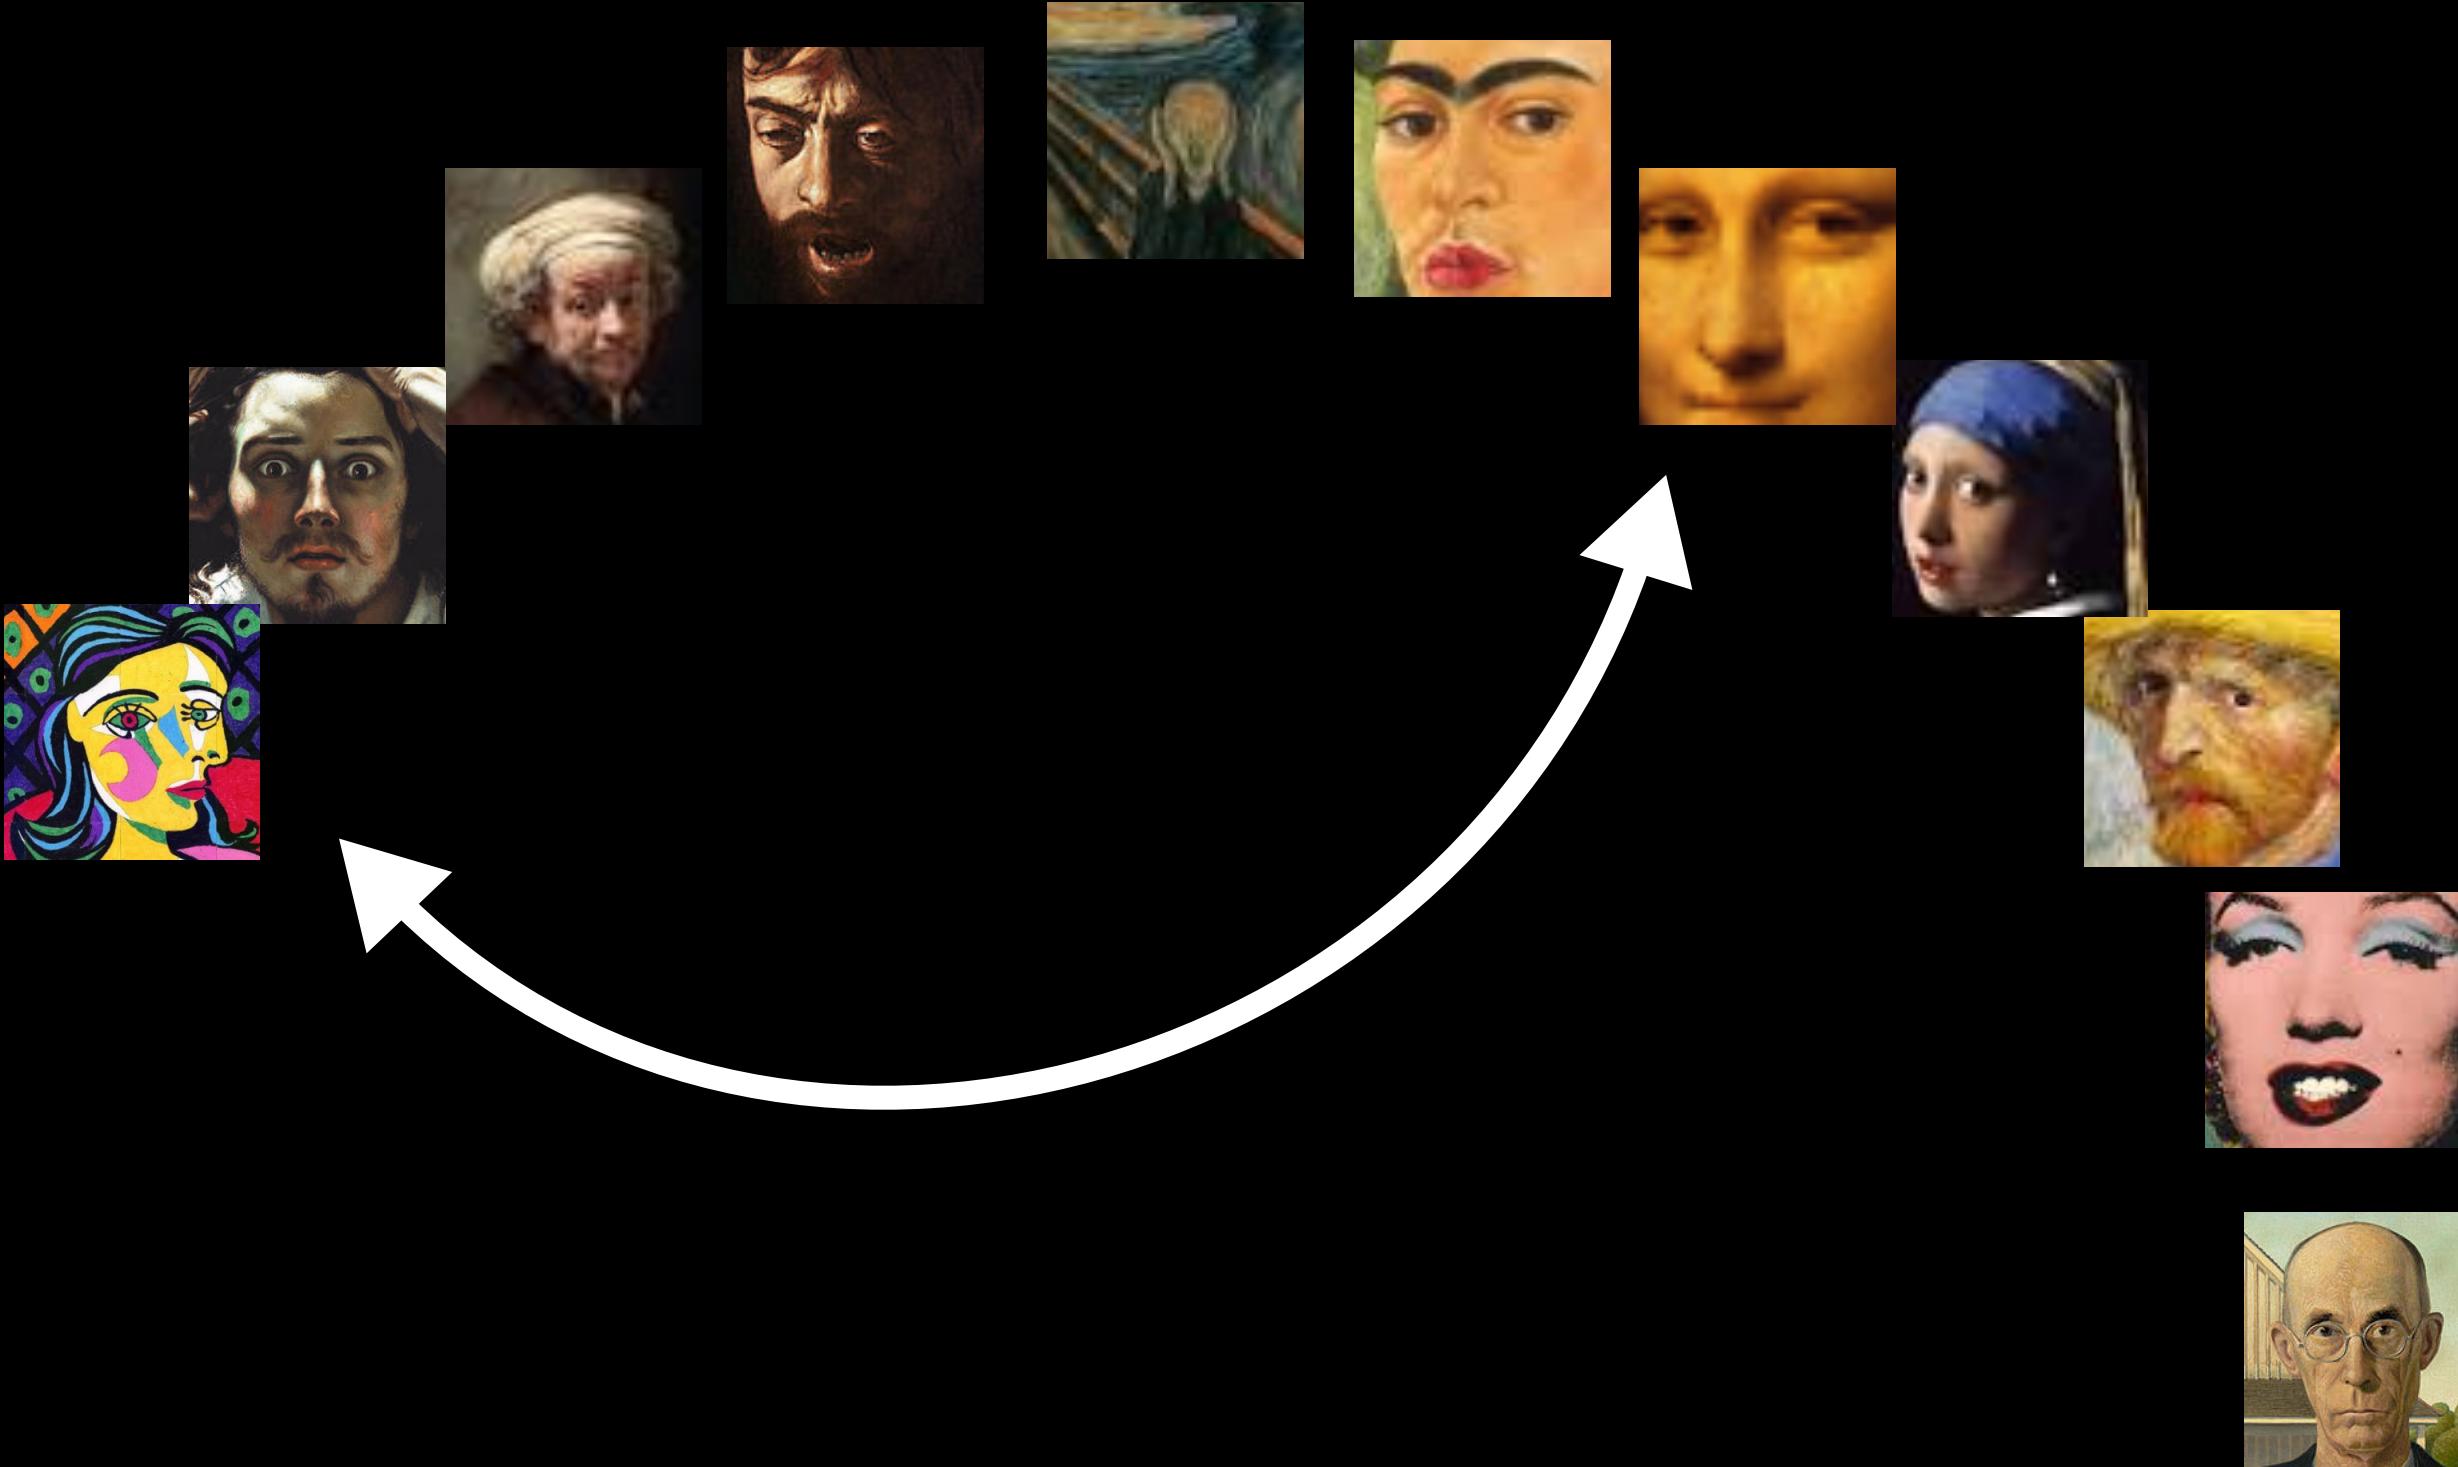
# birthday problem



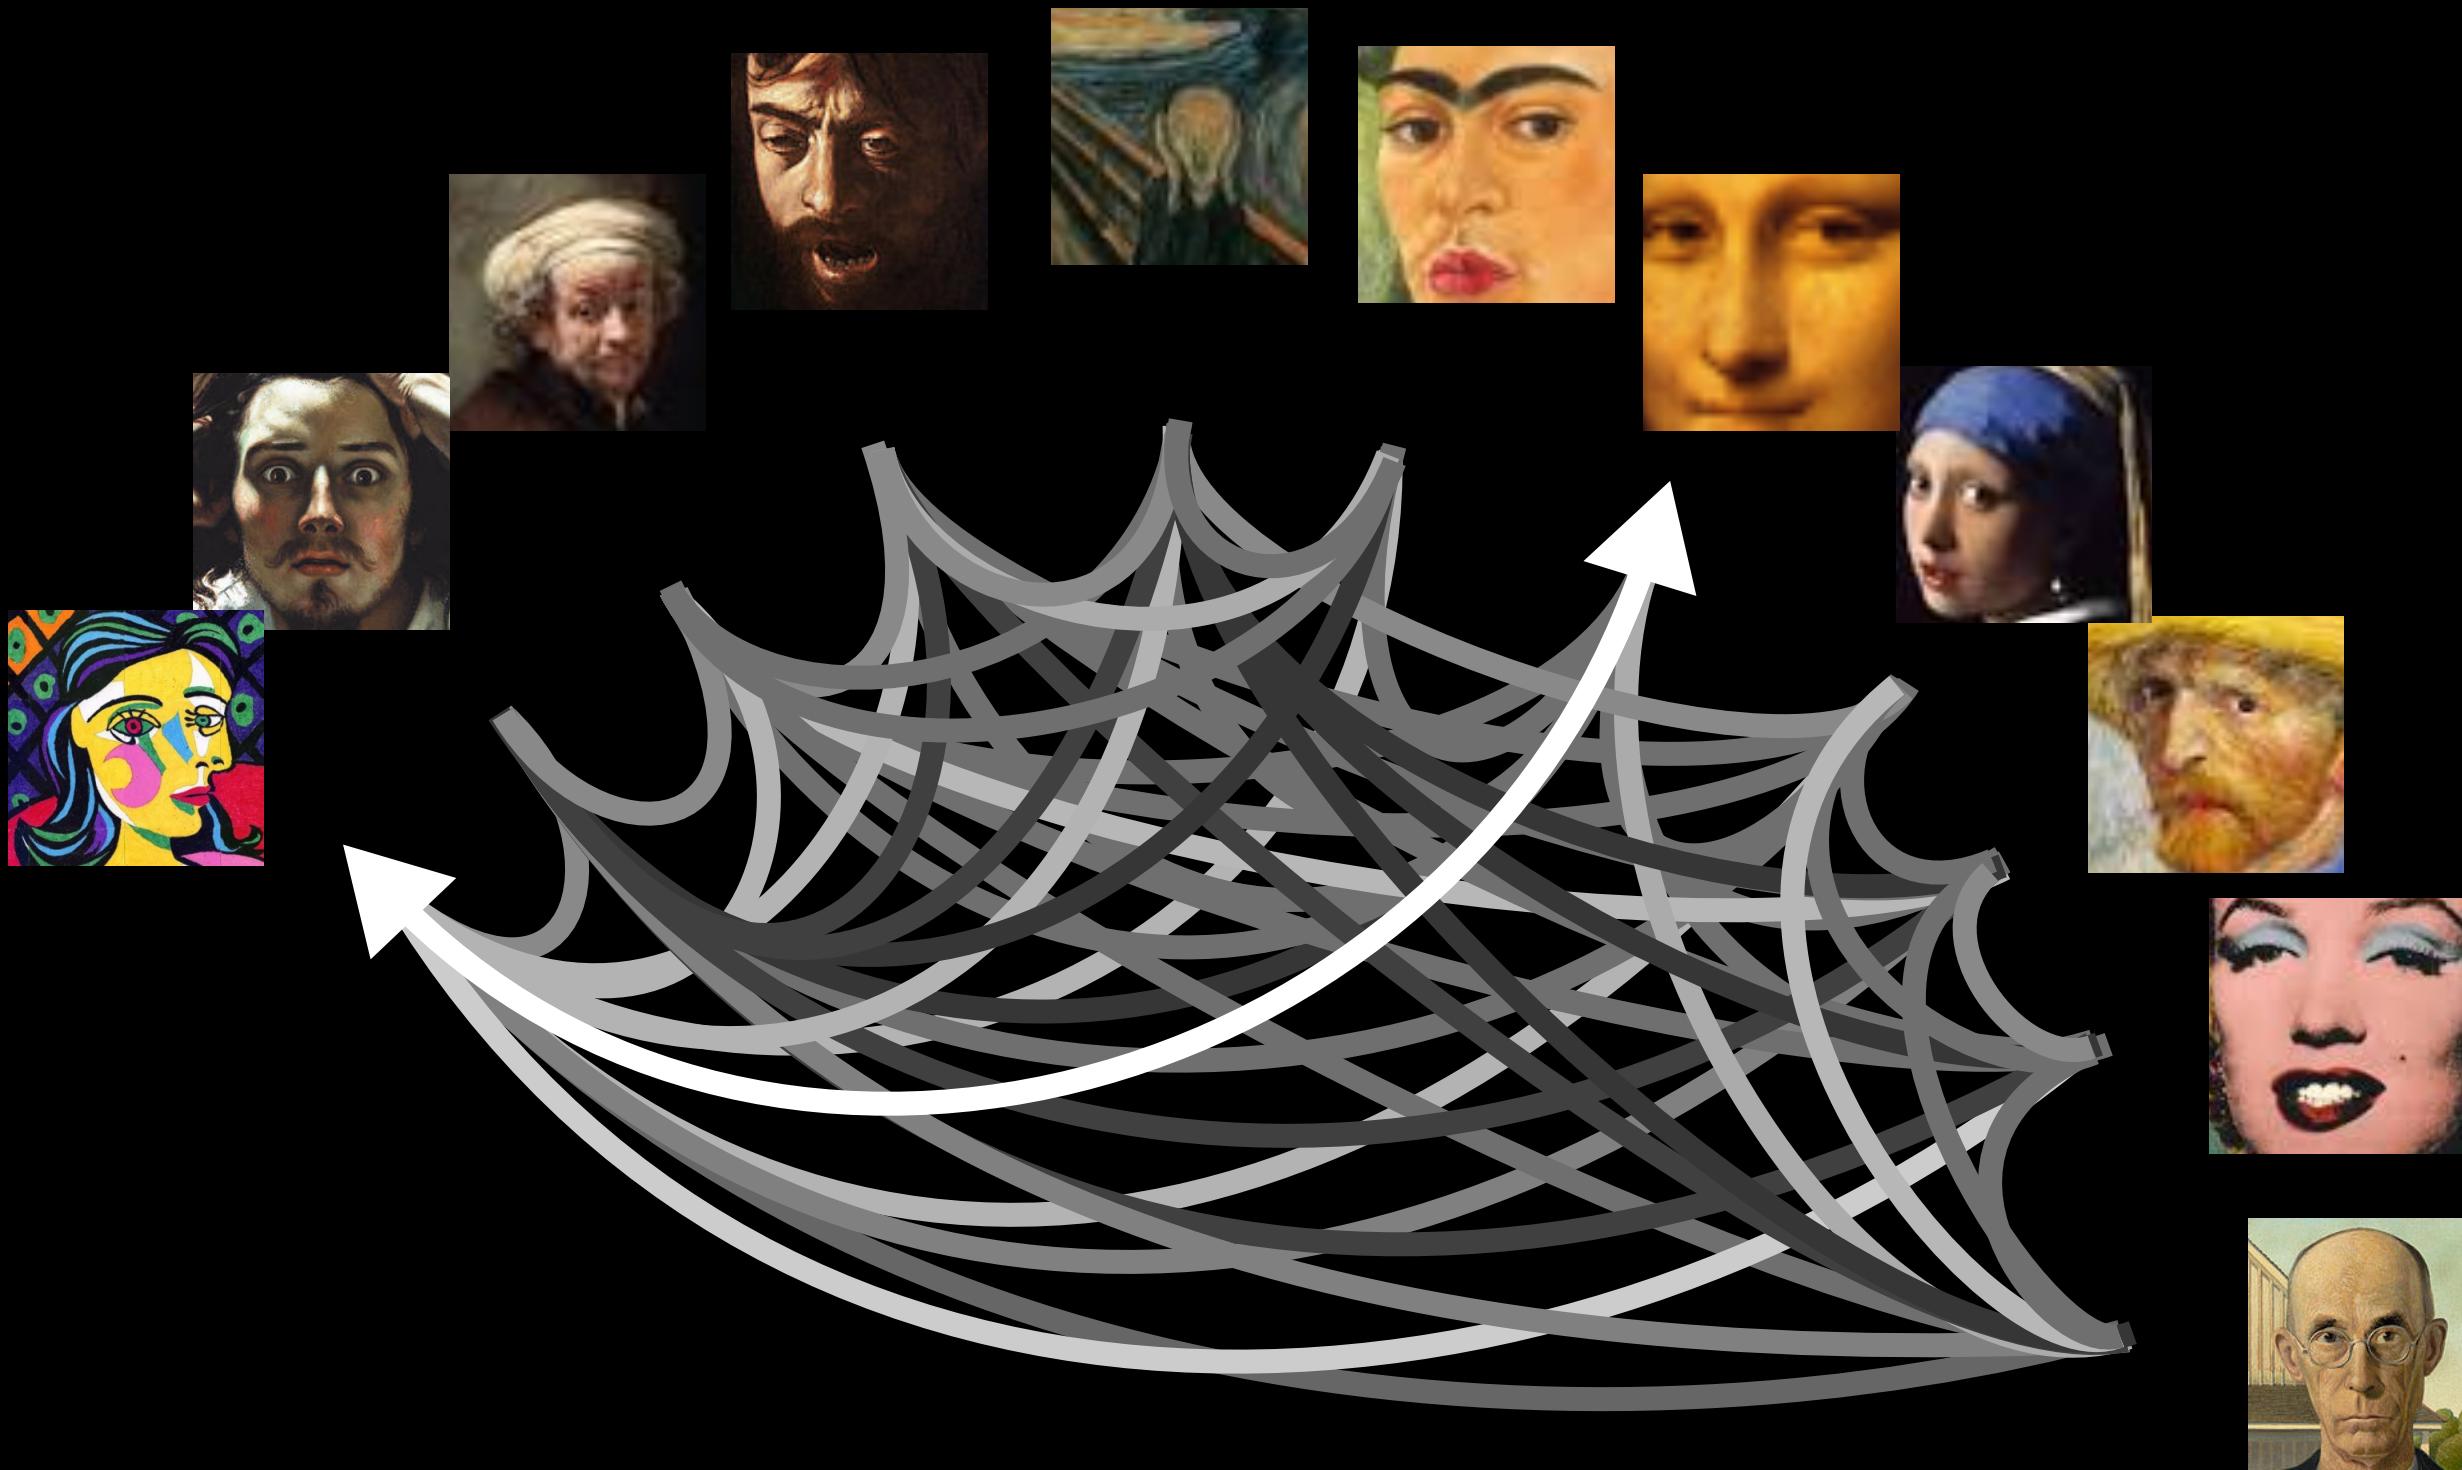
# birthday problem



# birthday problem

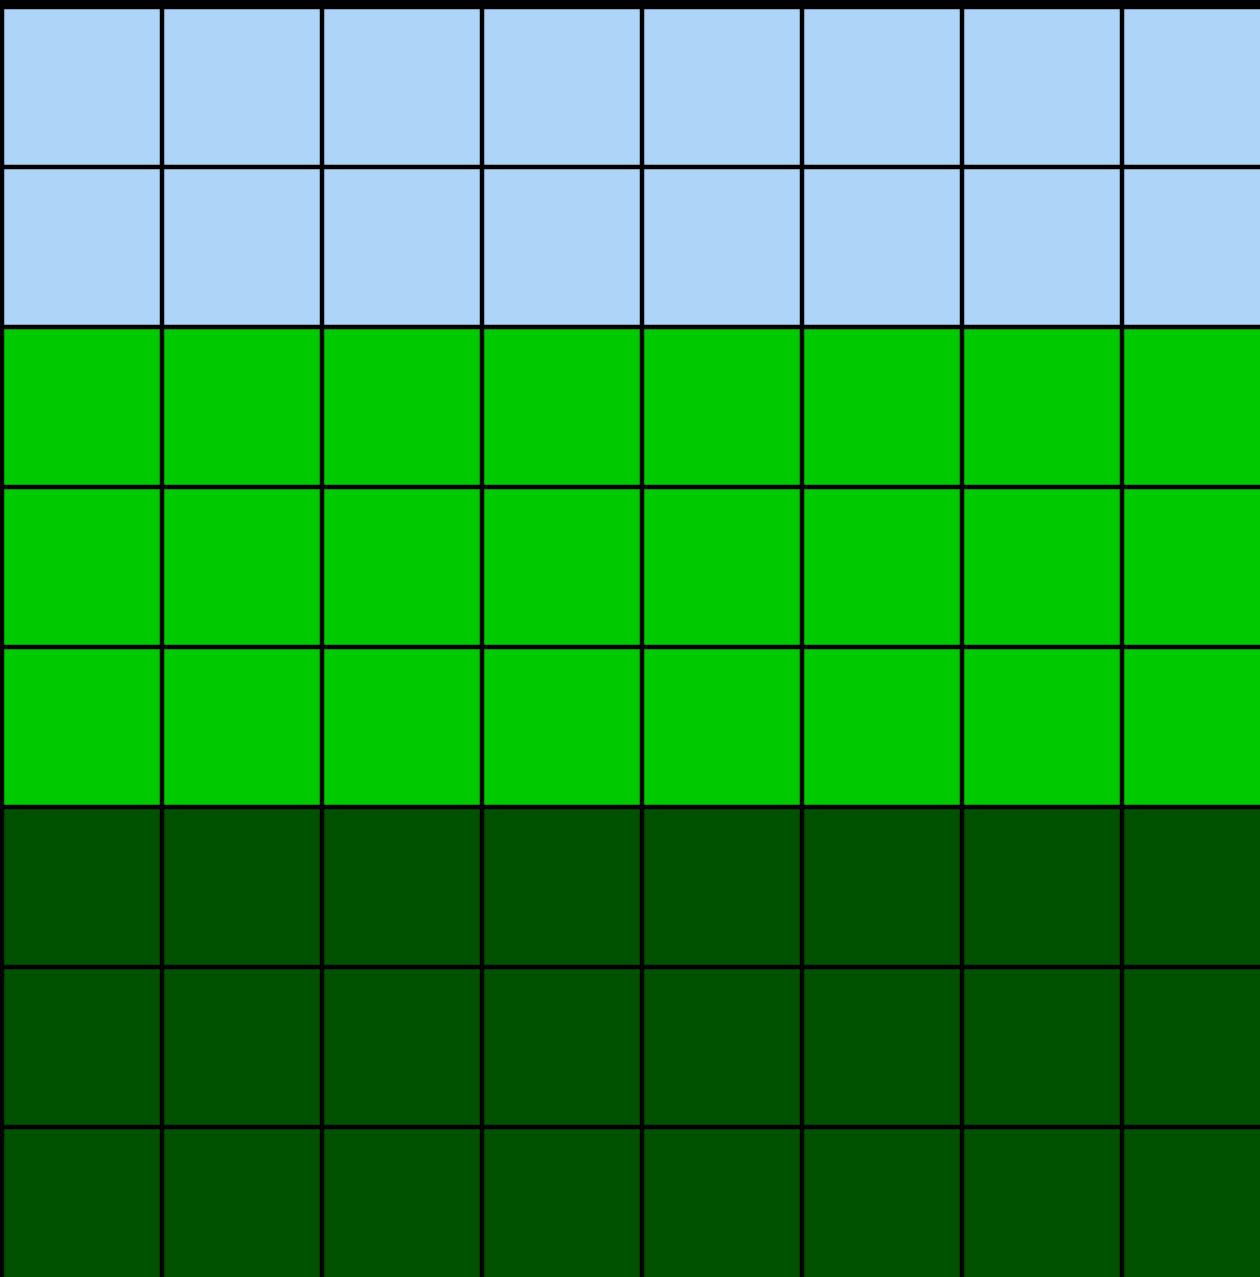


# birthday problem

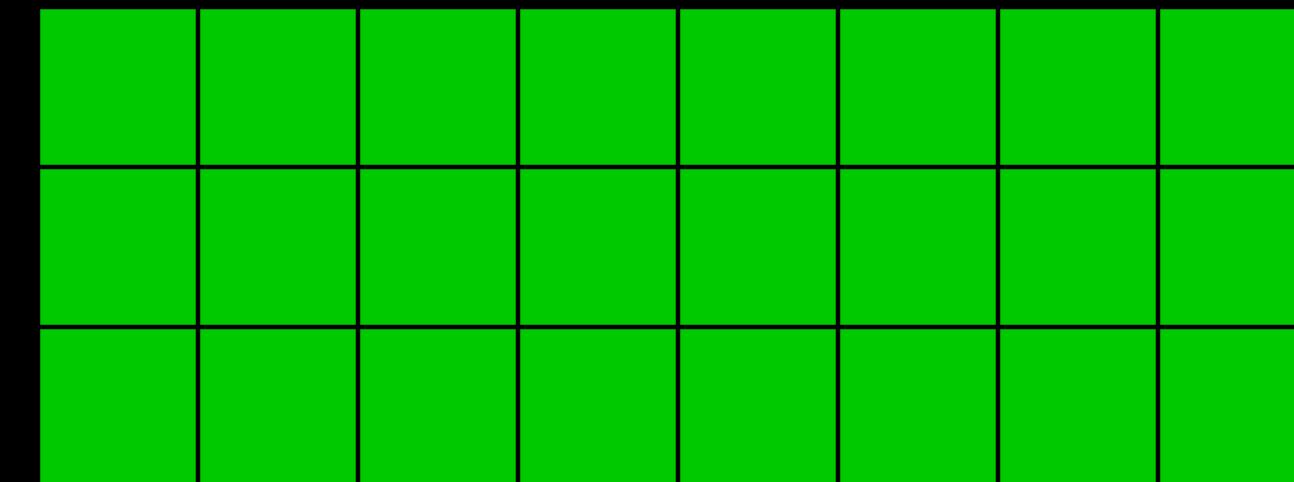


# primitive #3: birthday heapspray

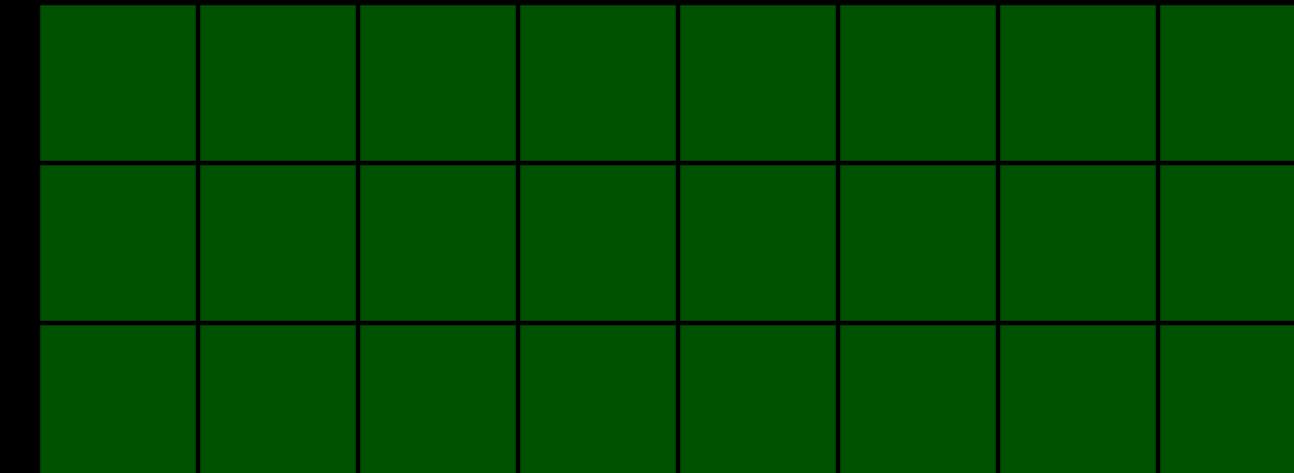
physical memory



attacker memory

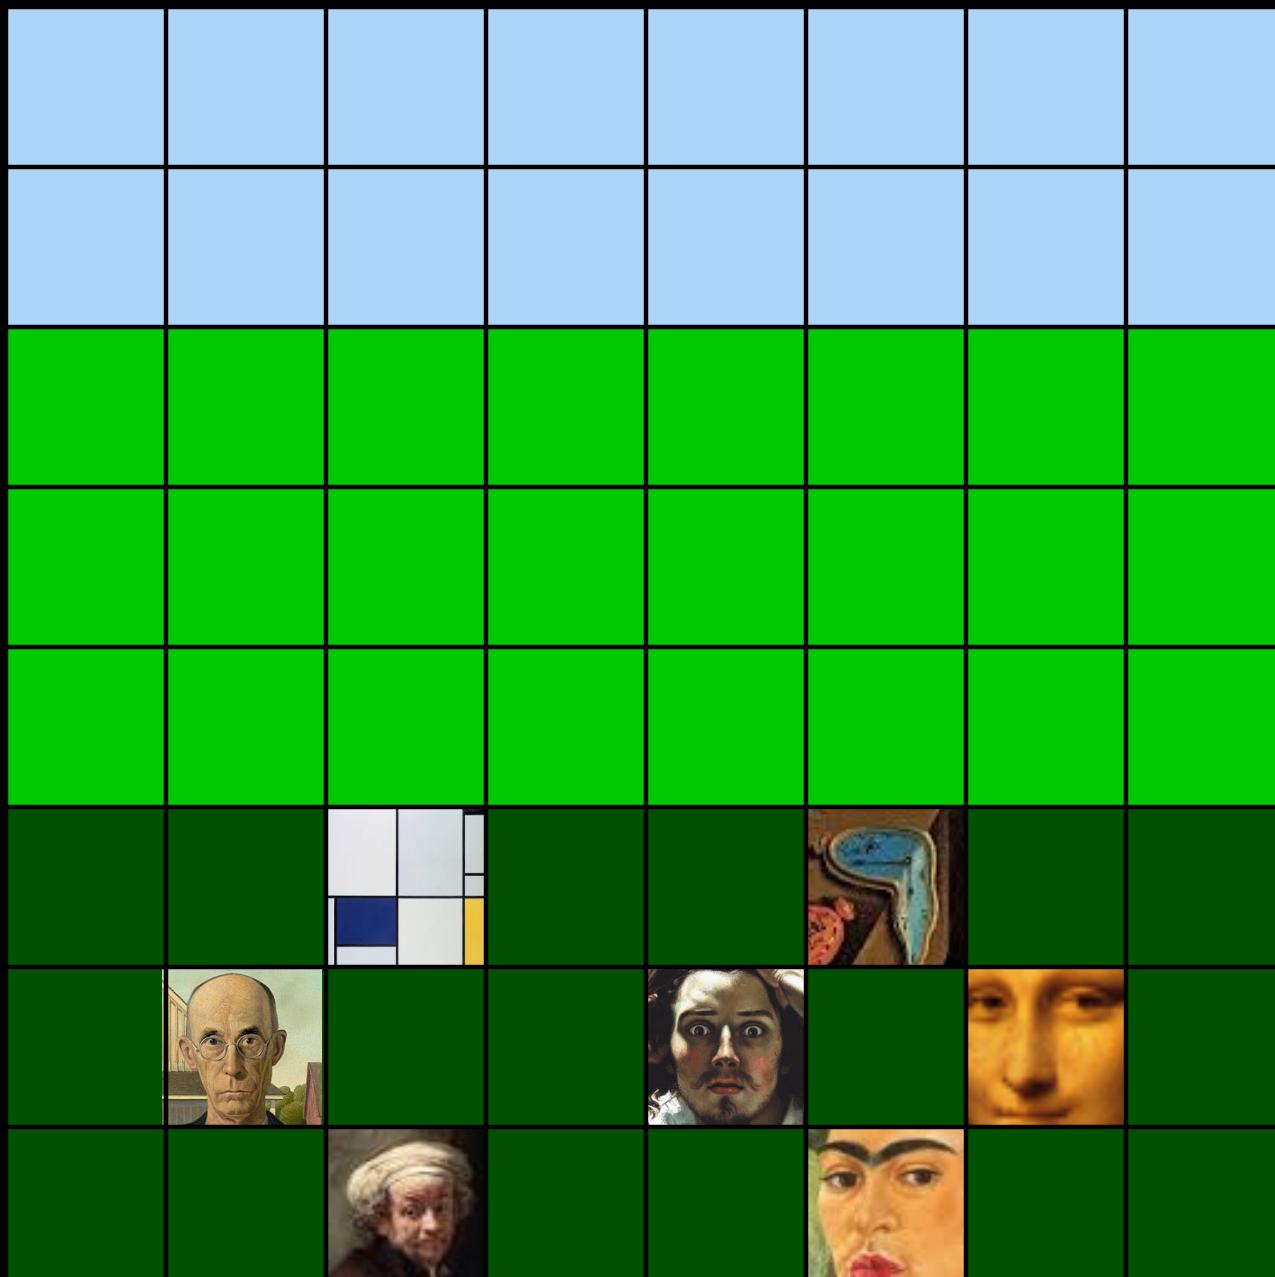


victim memory

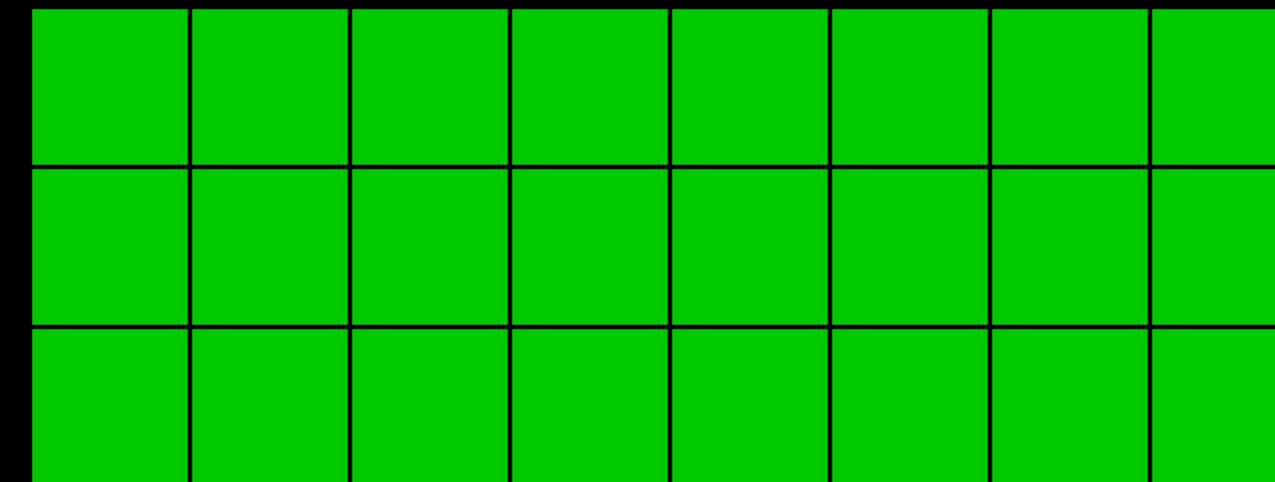


# primitive #3: birthday heapspray

physical memory



attacker memory

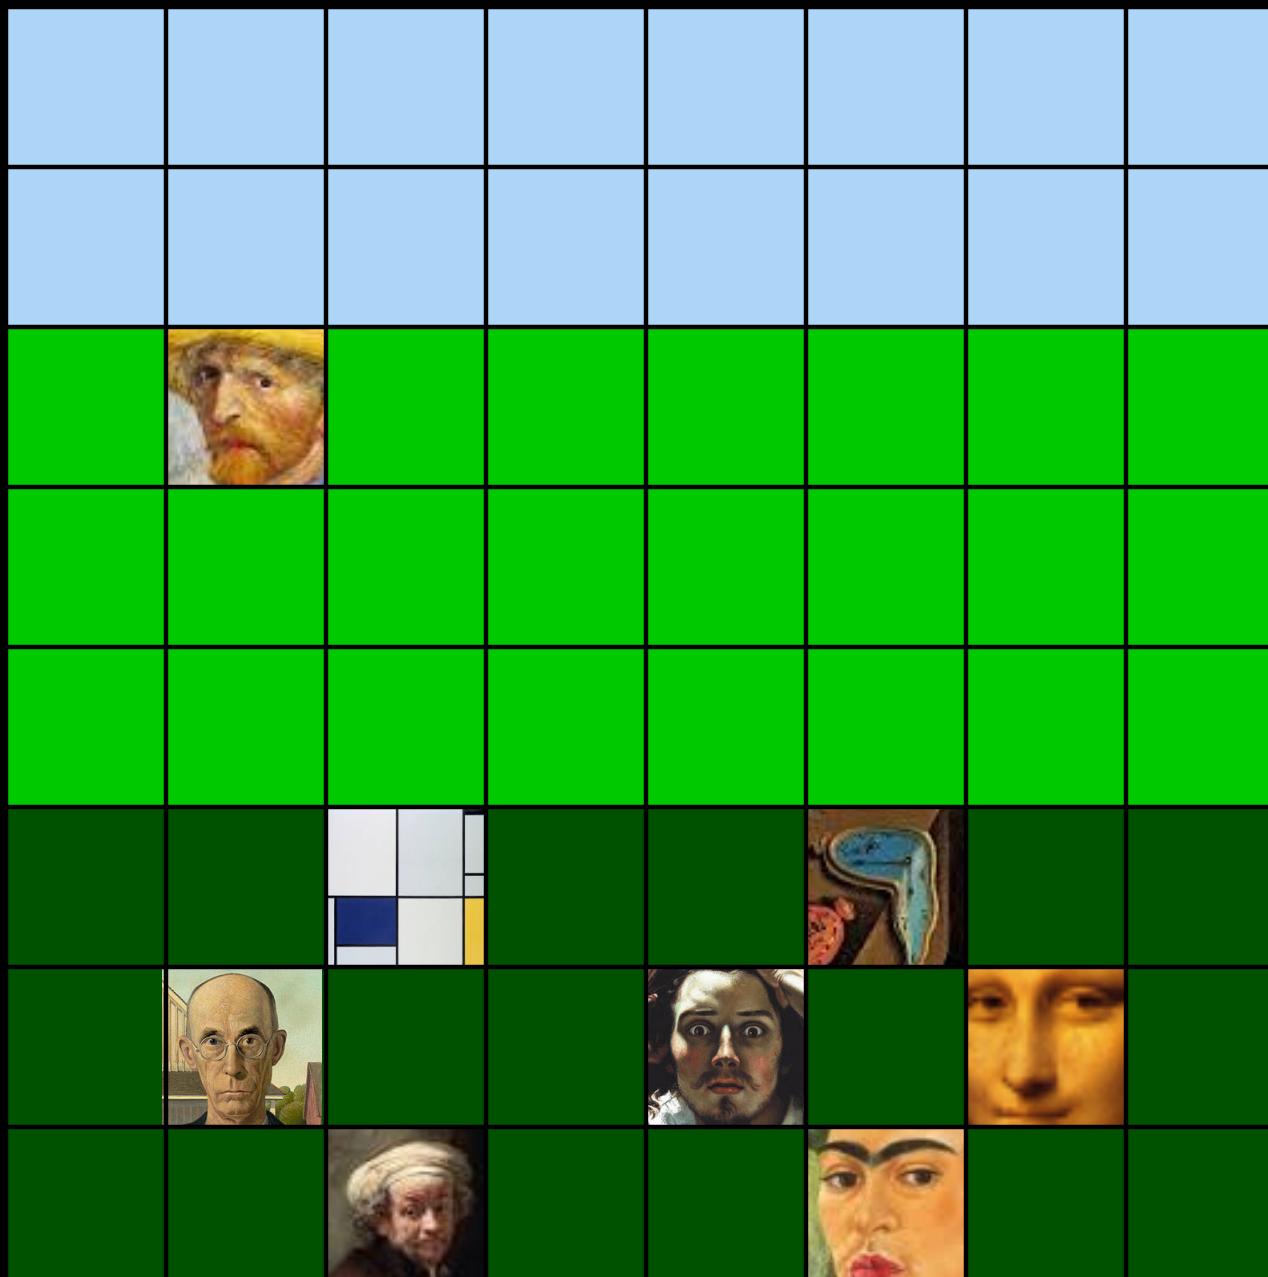


victim memory

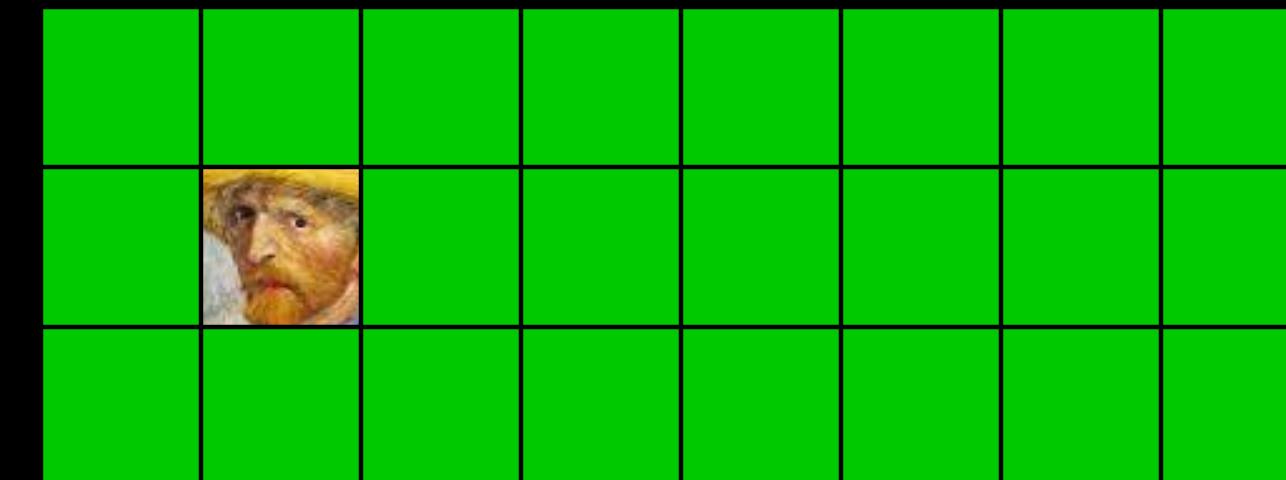


# primitive #3: birthday heapspray

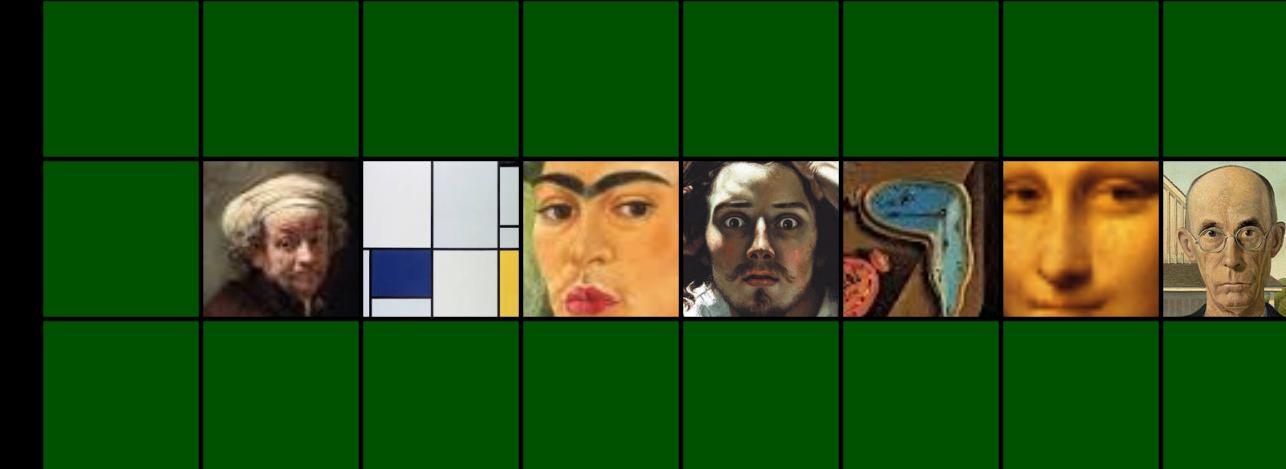
physical memory



attacker memory

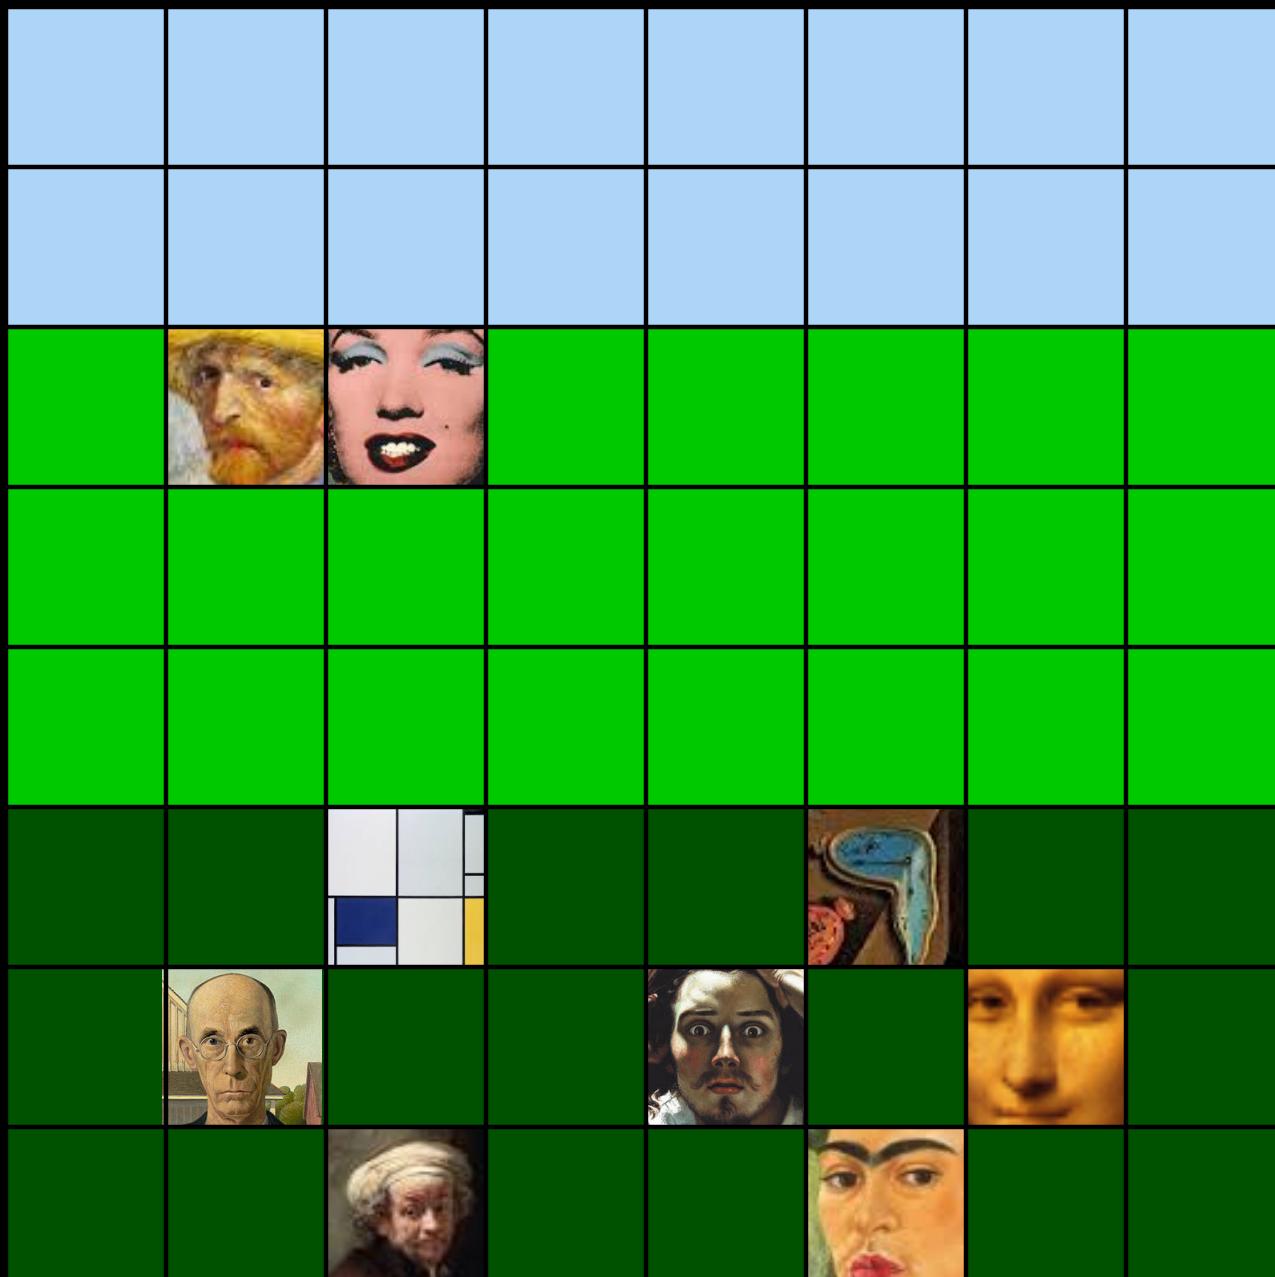


victim memory



# primitive #3: birthday heapspray

physical memory



attacker memory

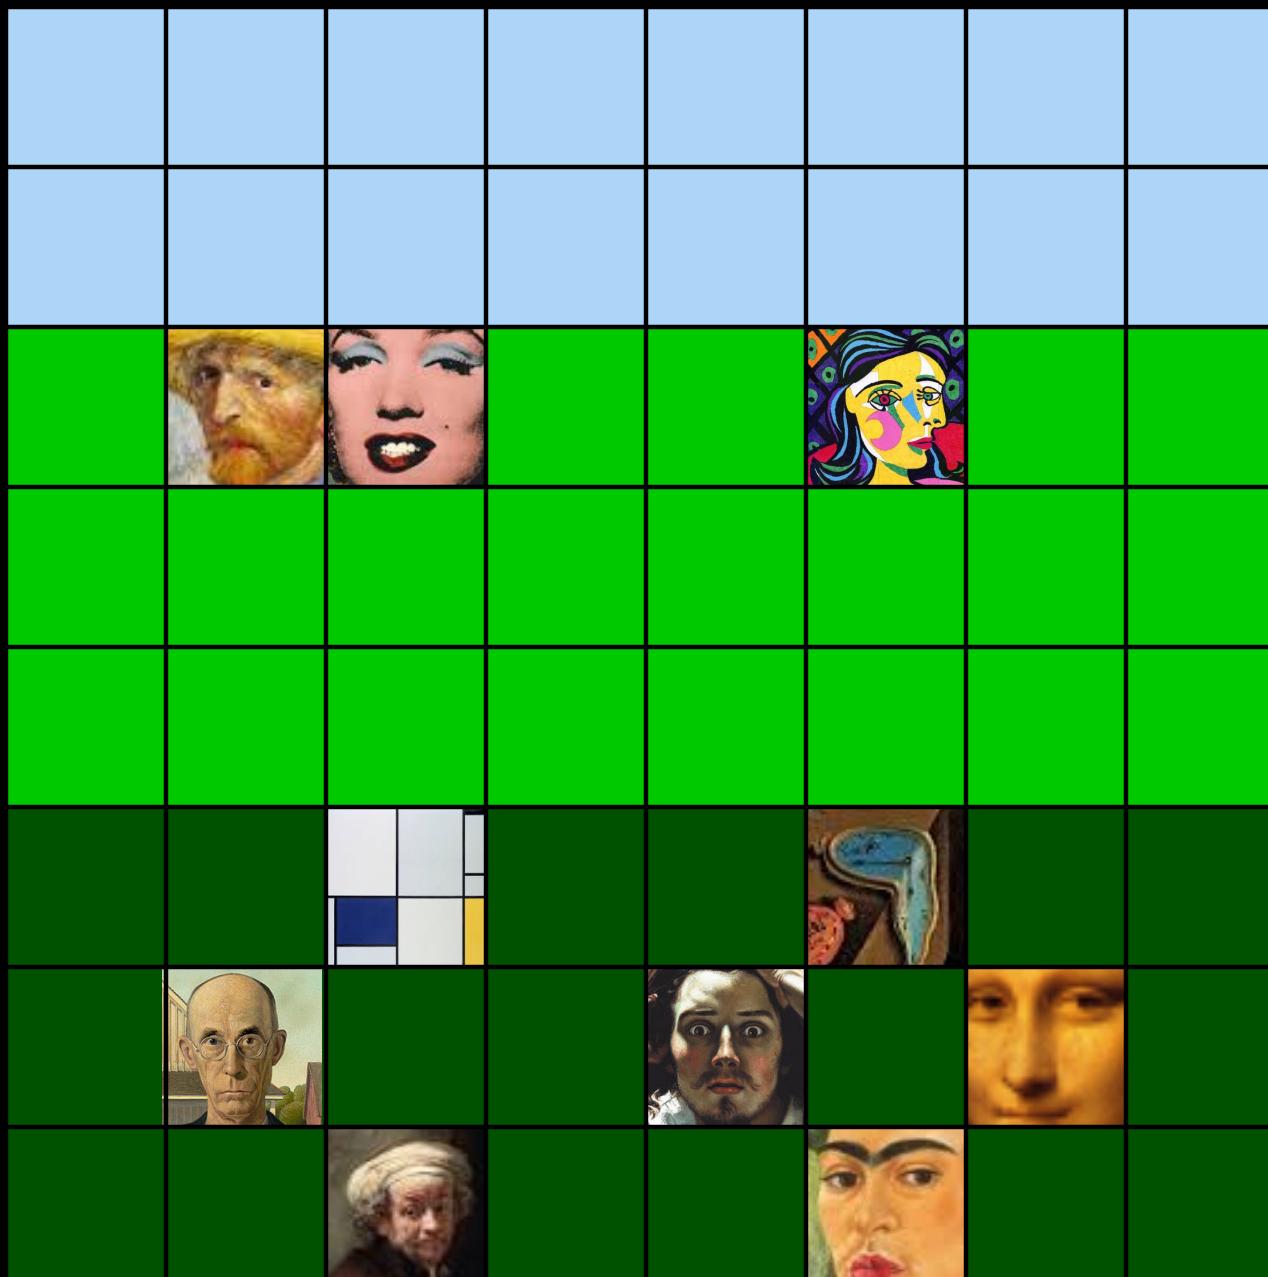


victim memory



# primitive #3: birthday heapspray

physical memory



attacker memory

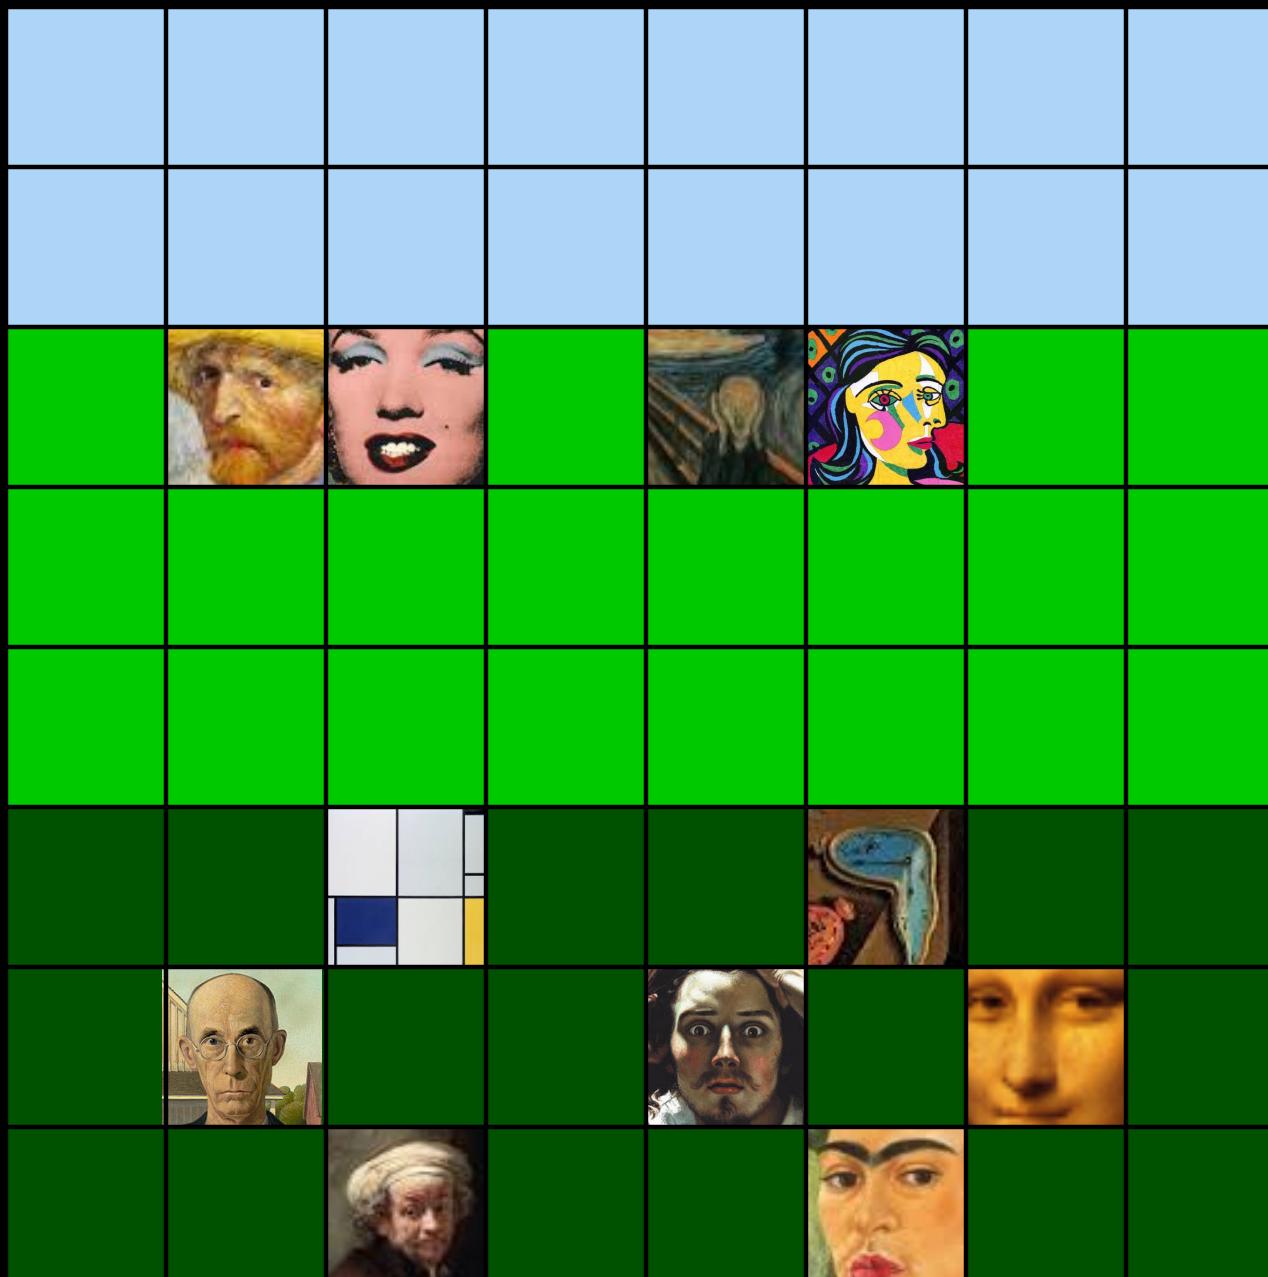


victim memory



# primitive #3: birthday heapspray

physical memory



attacker memory

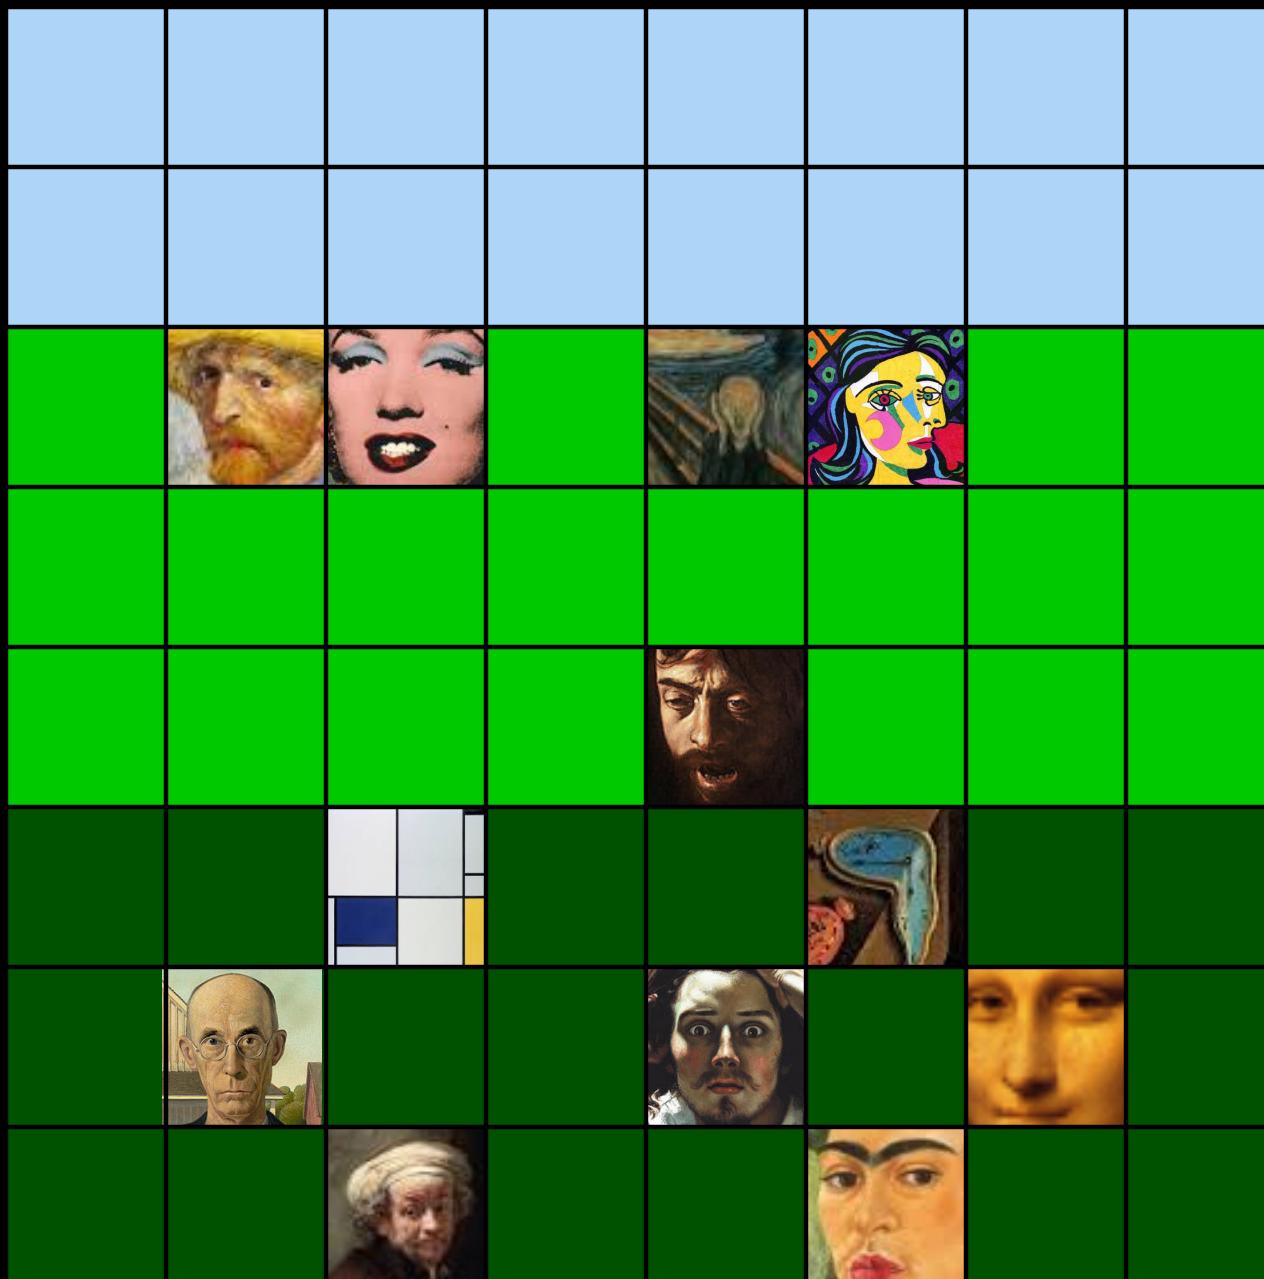


victim memory

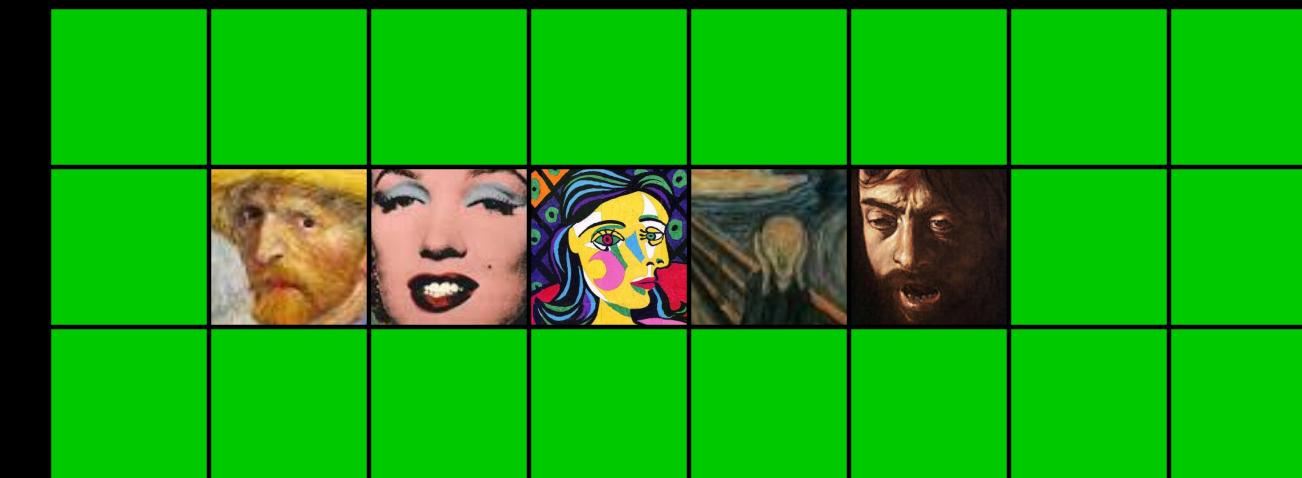


# primitive #3: birthday heapspray

physical memory



attacker memory

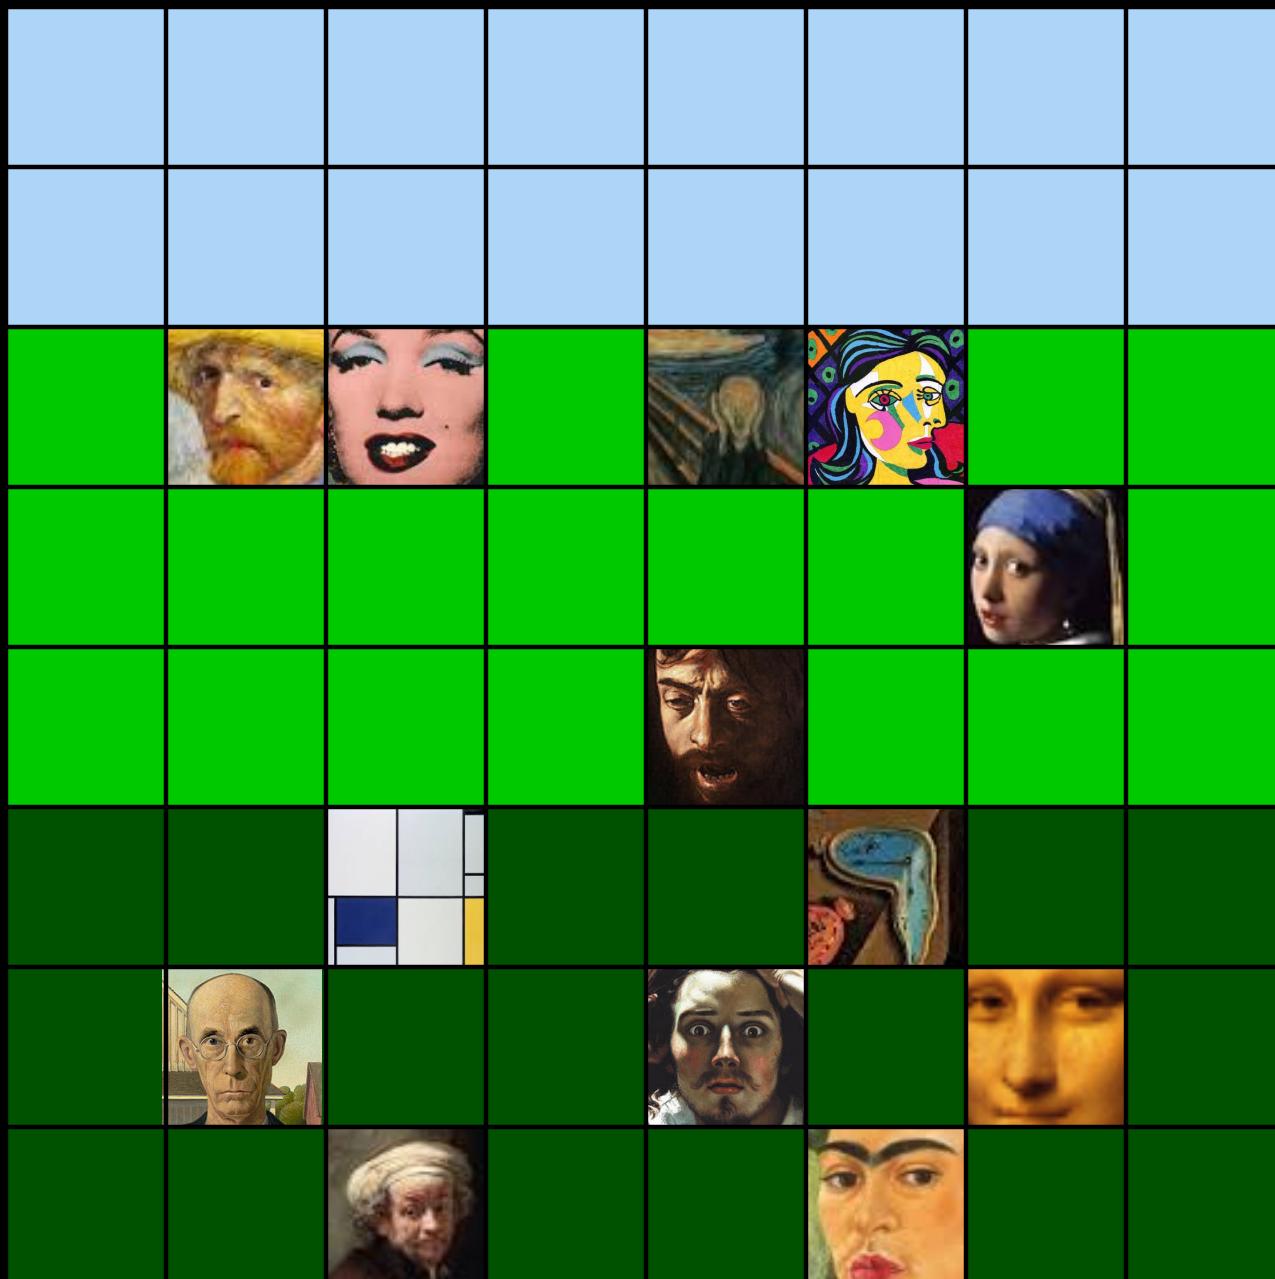


victim memory



# primitive #3: birthday heapspray

physical memory



attacker memory

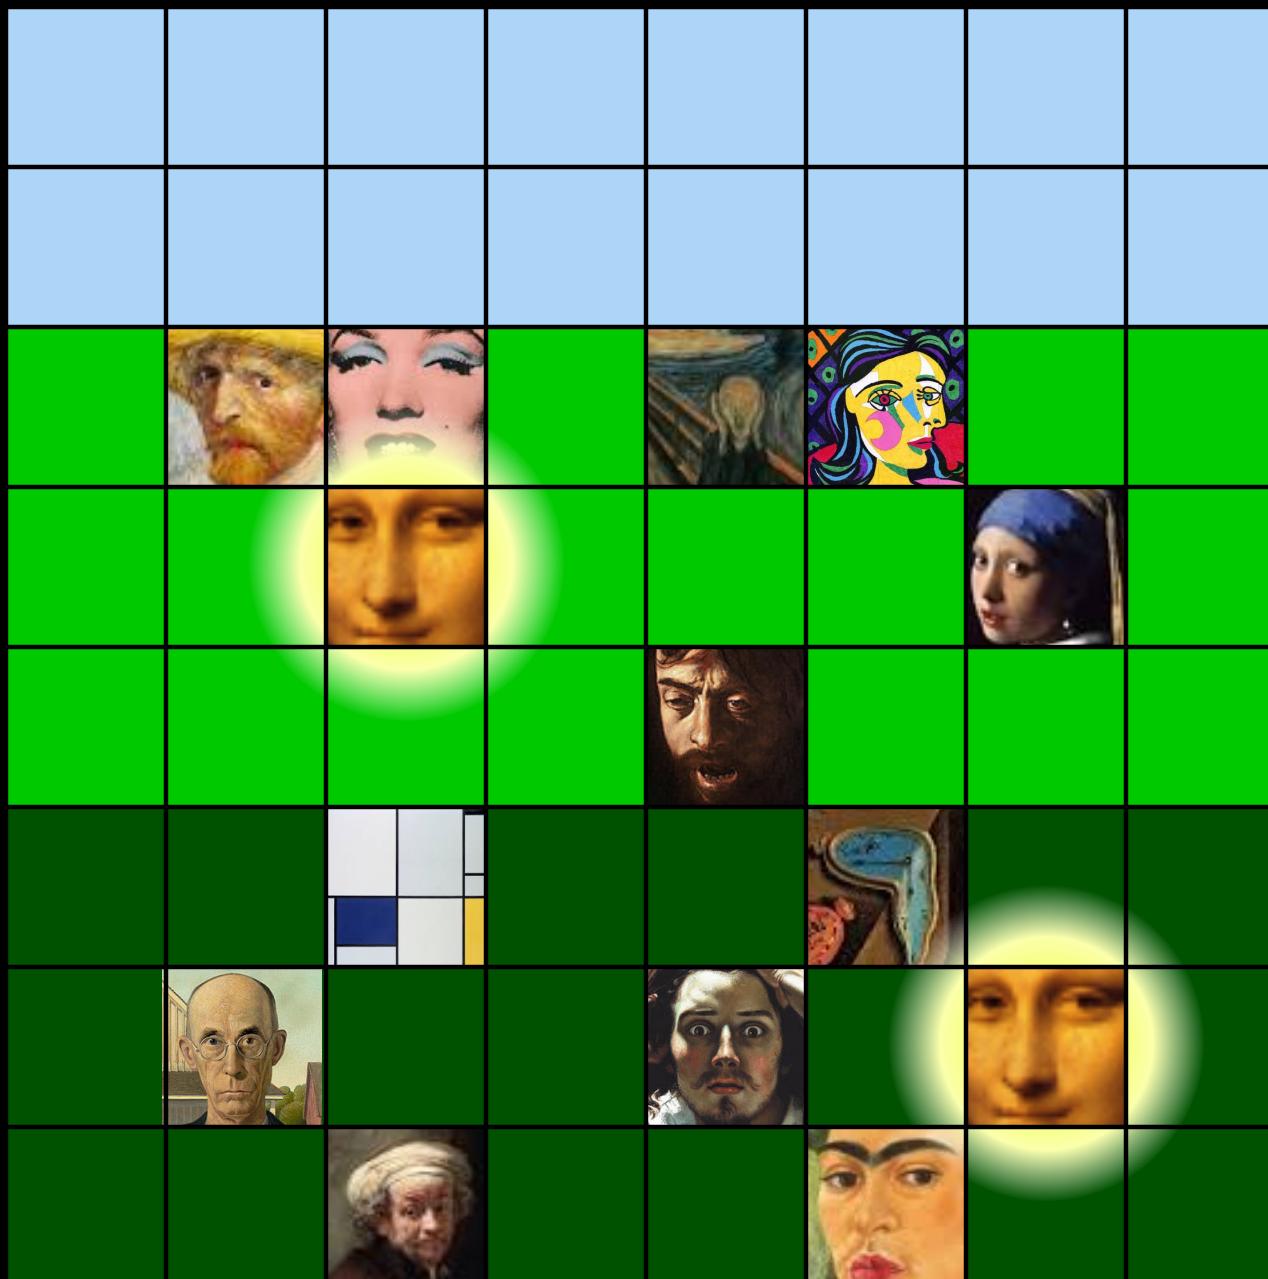


victim memory



# primitive #3: birthday heapspray

physical memory



attacker memory

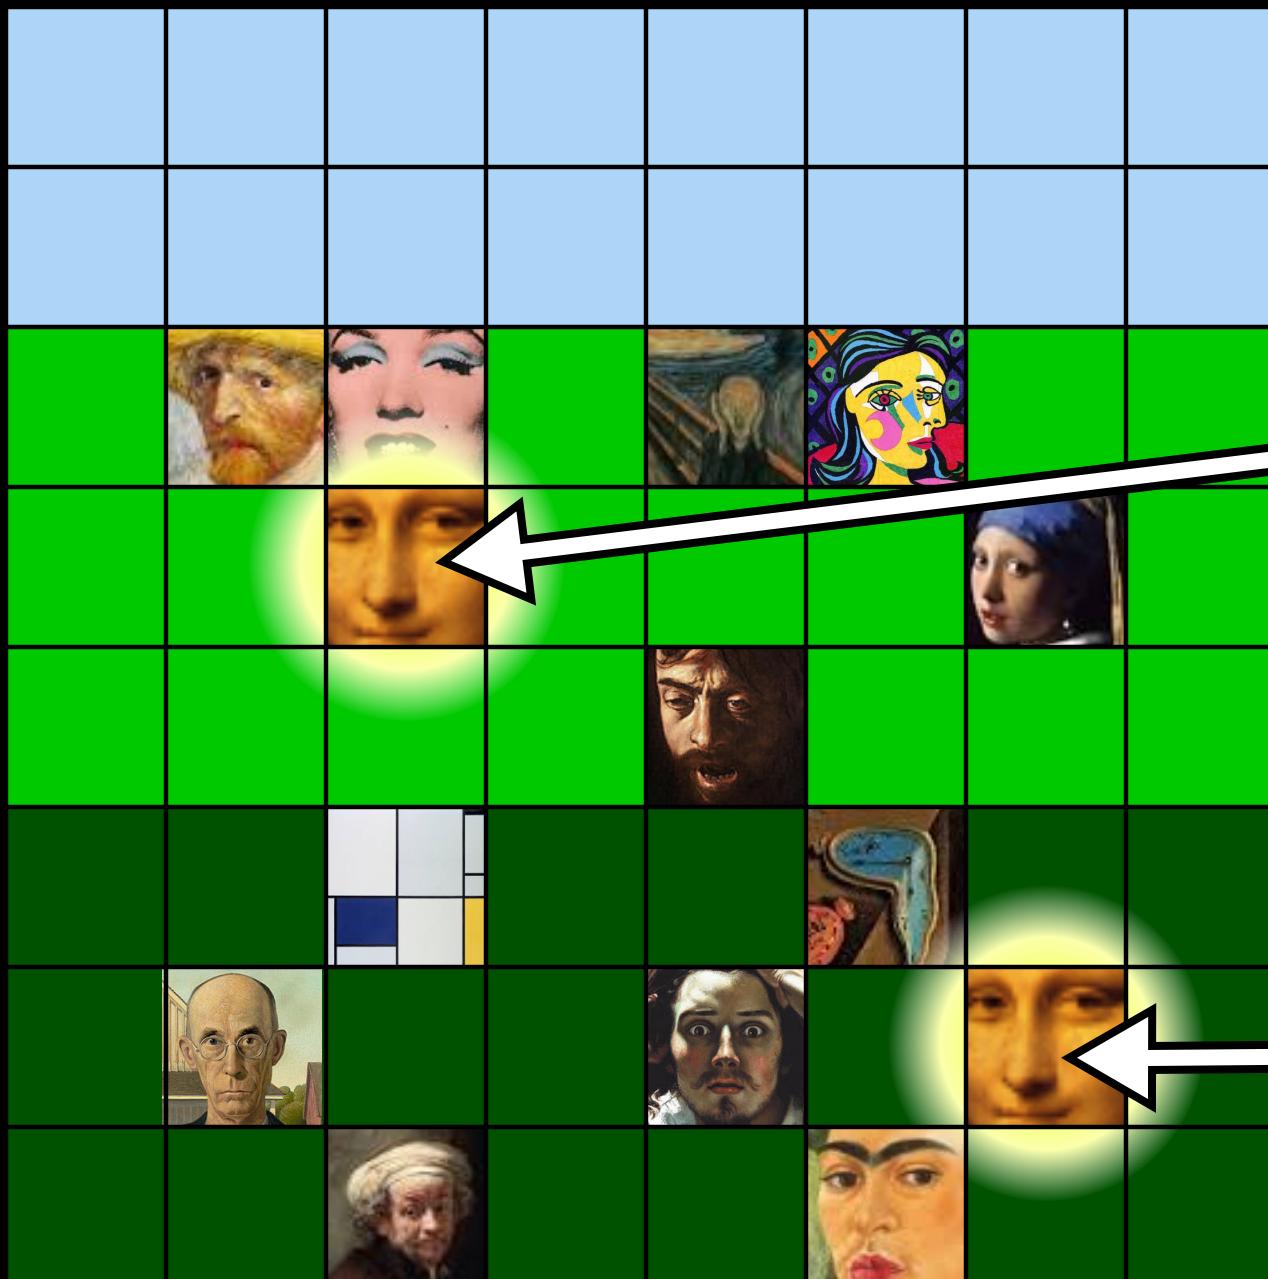


victim memory



# primitive #3: birthday heapspray

physical memory



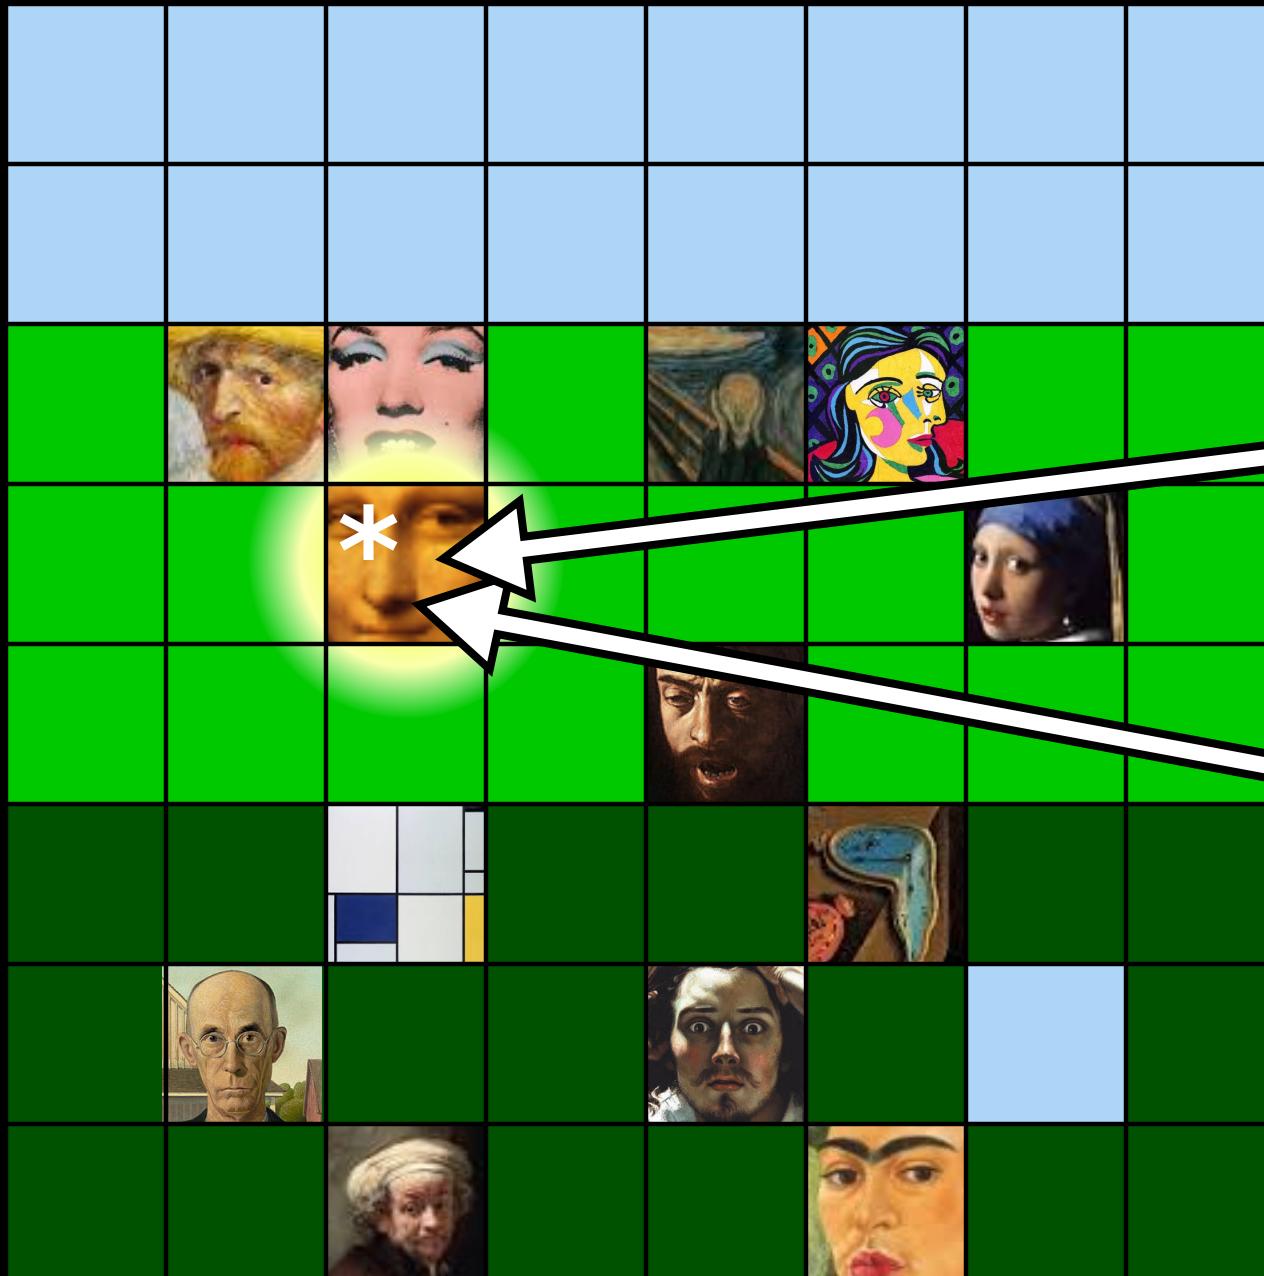
attacker memory



victim memory

# primitive #3: birthday heapspray

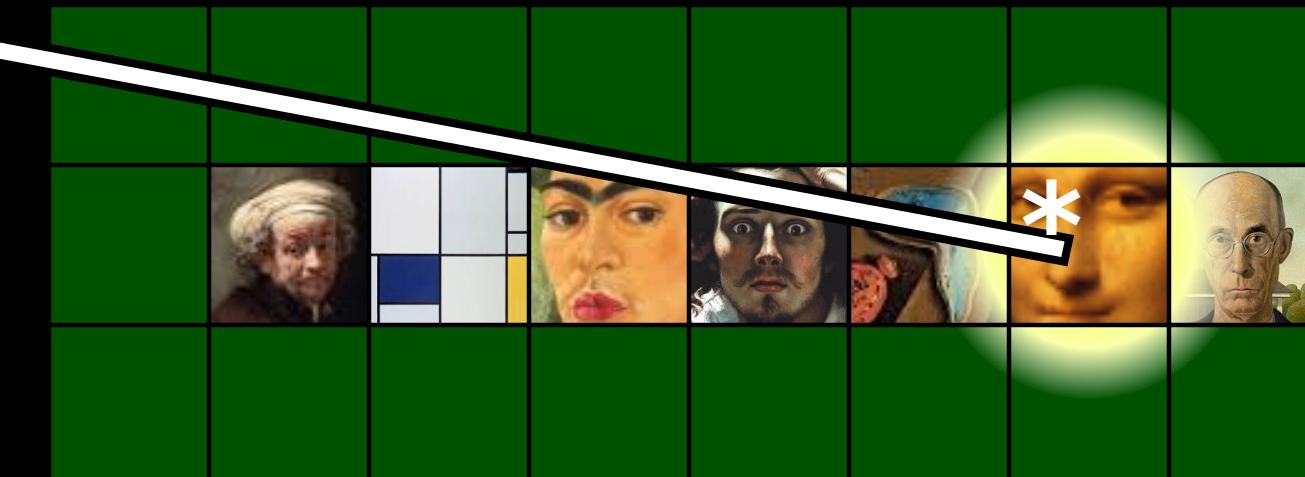
physical memory



attacker memory



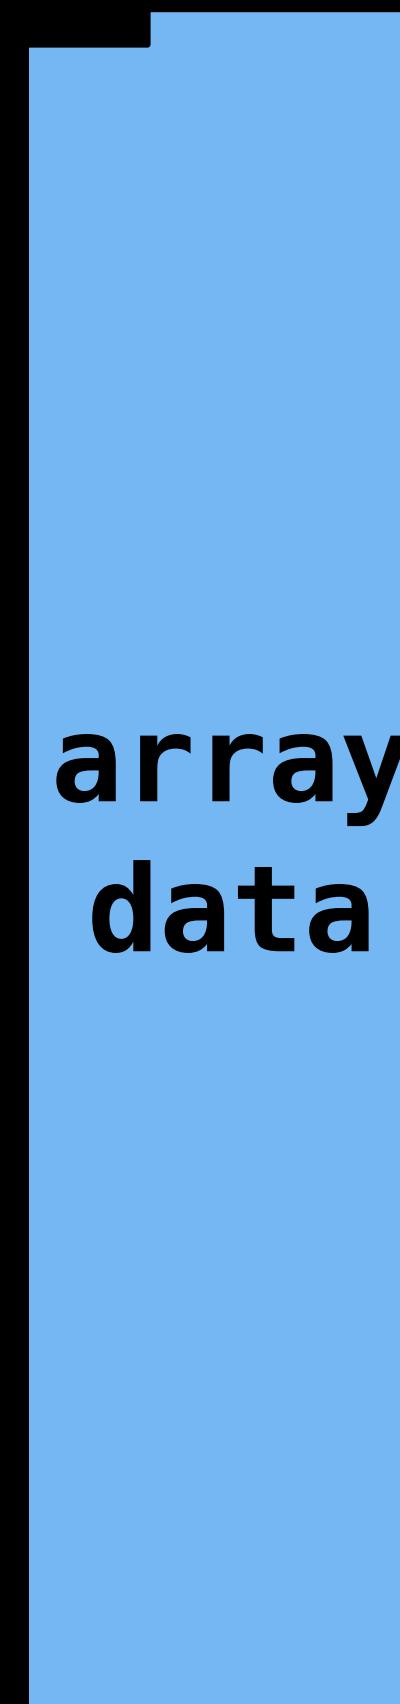
victim memory



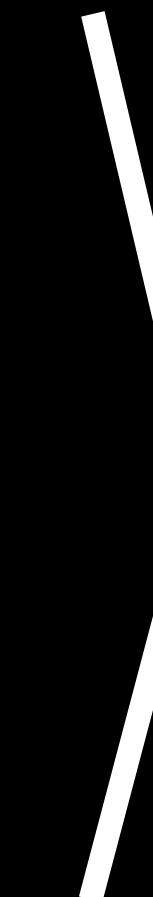
# Creating Secret Pages



# Creating Secret Pages

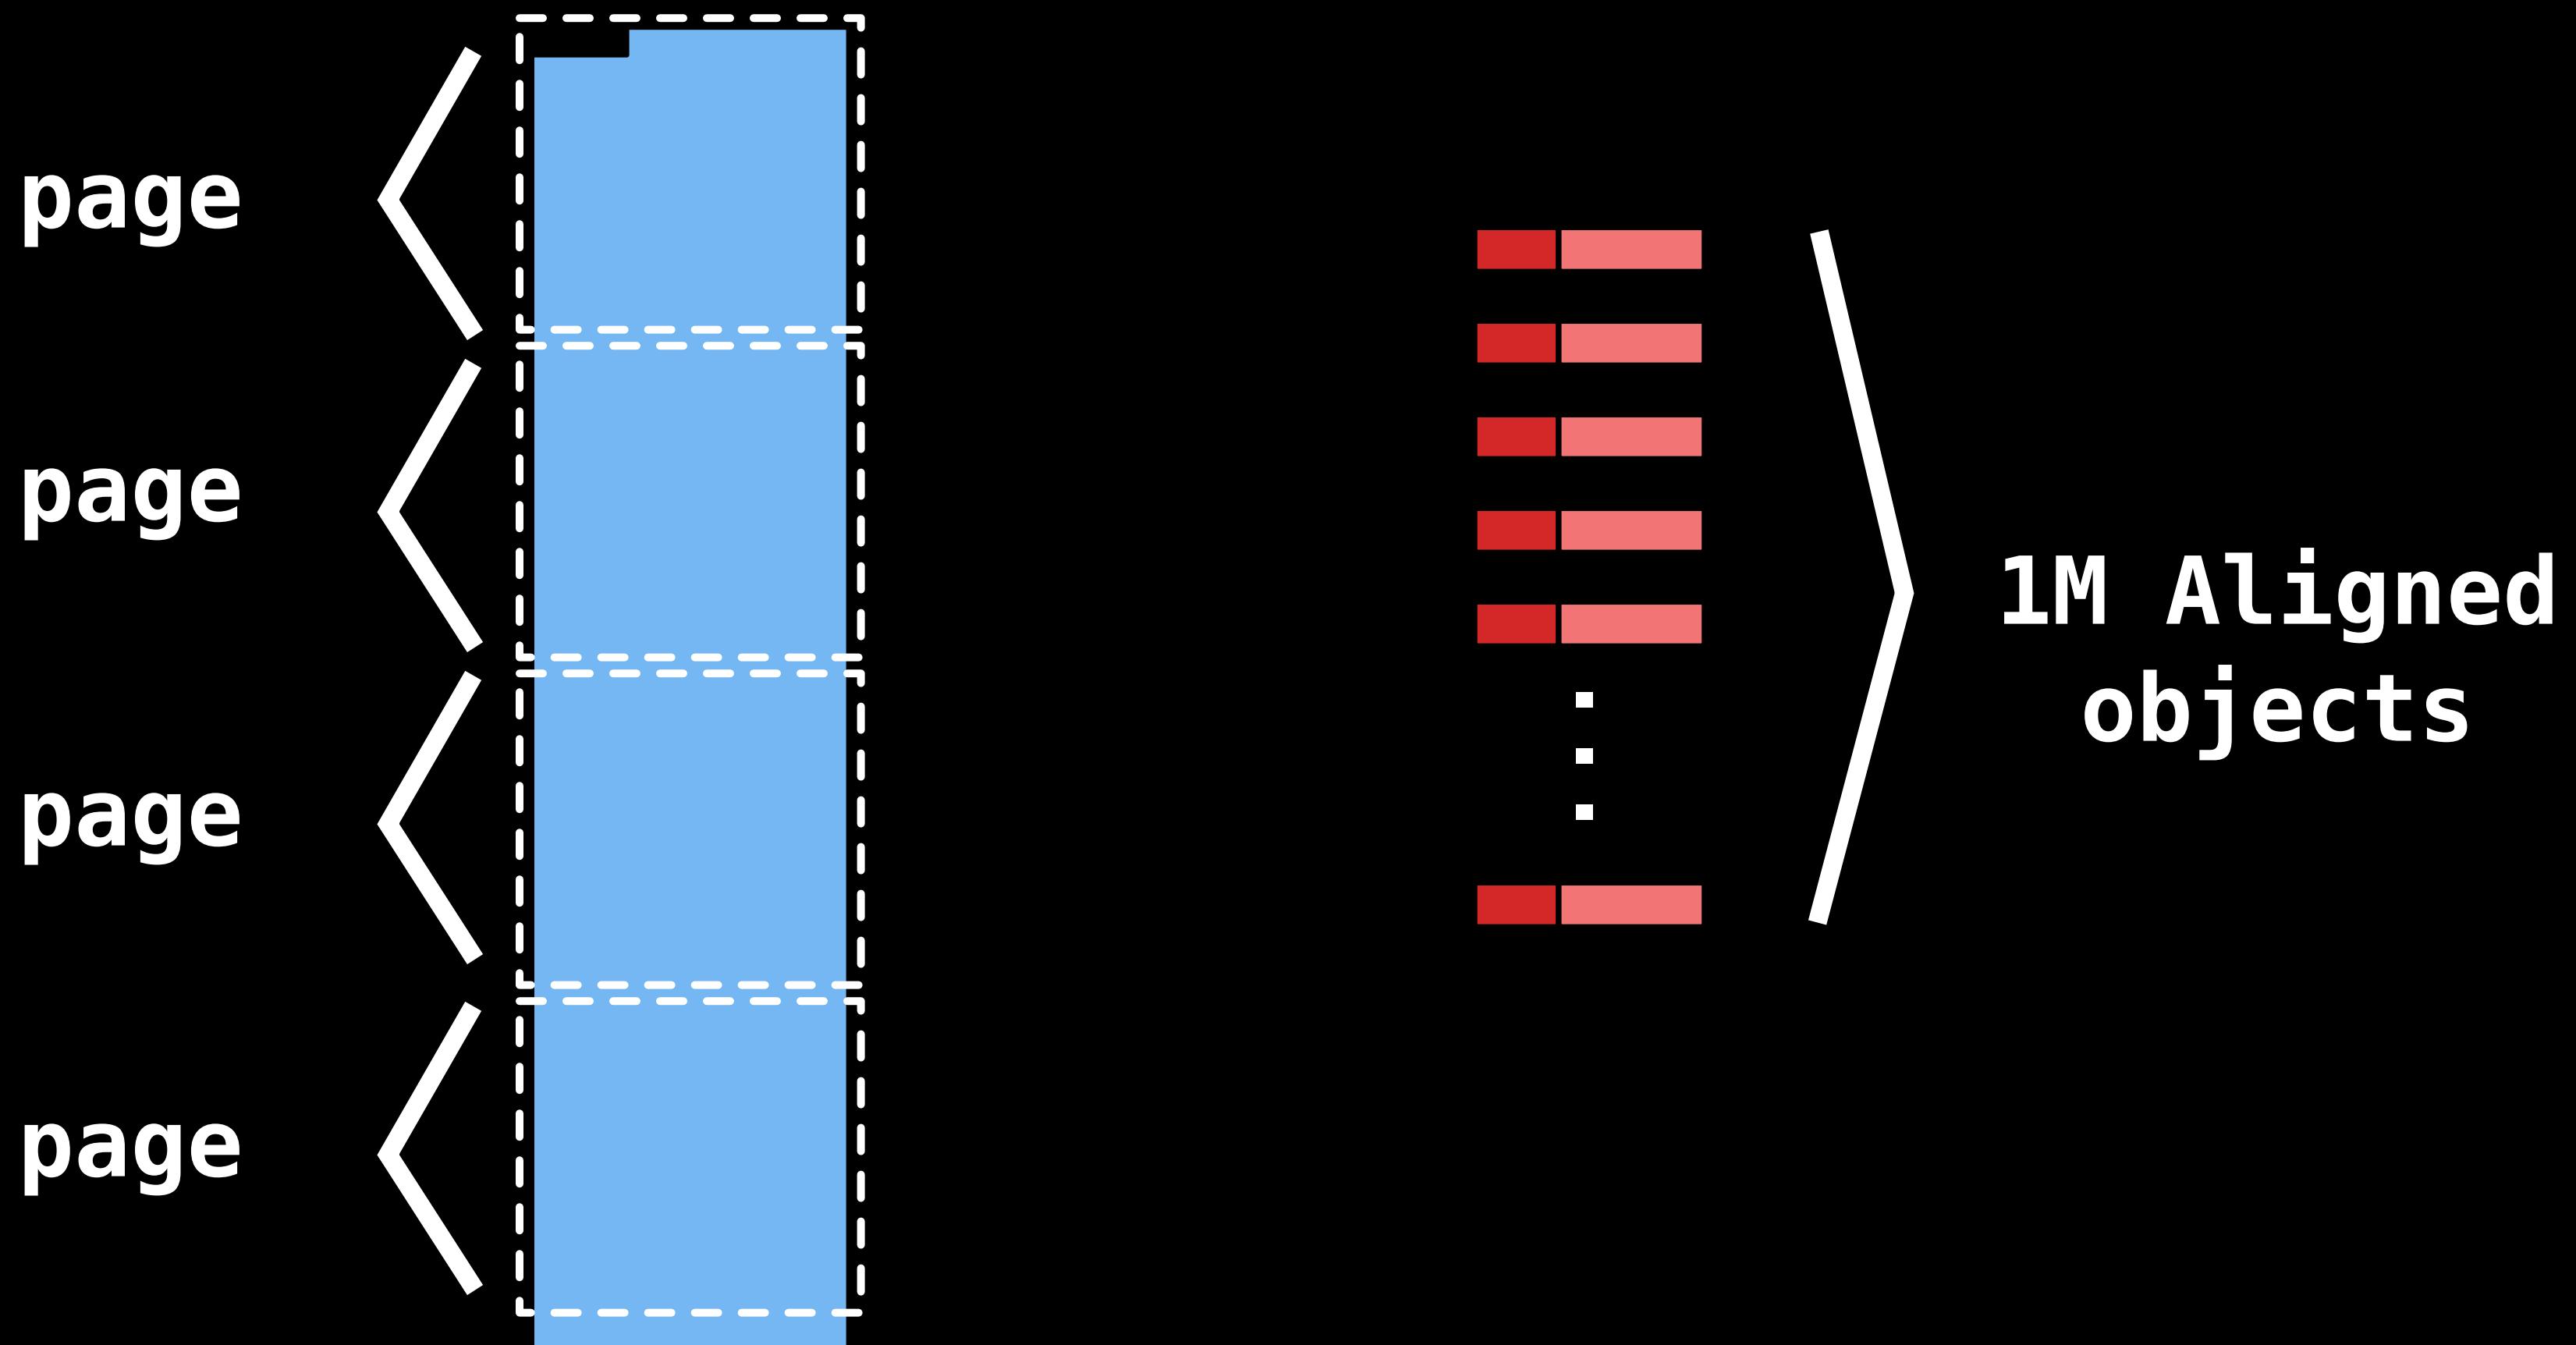


array  
data

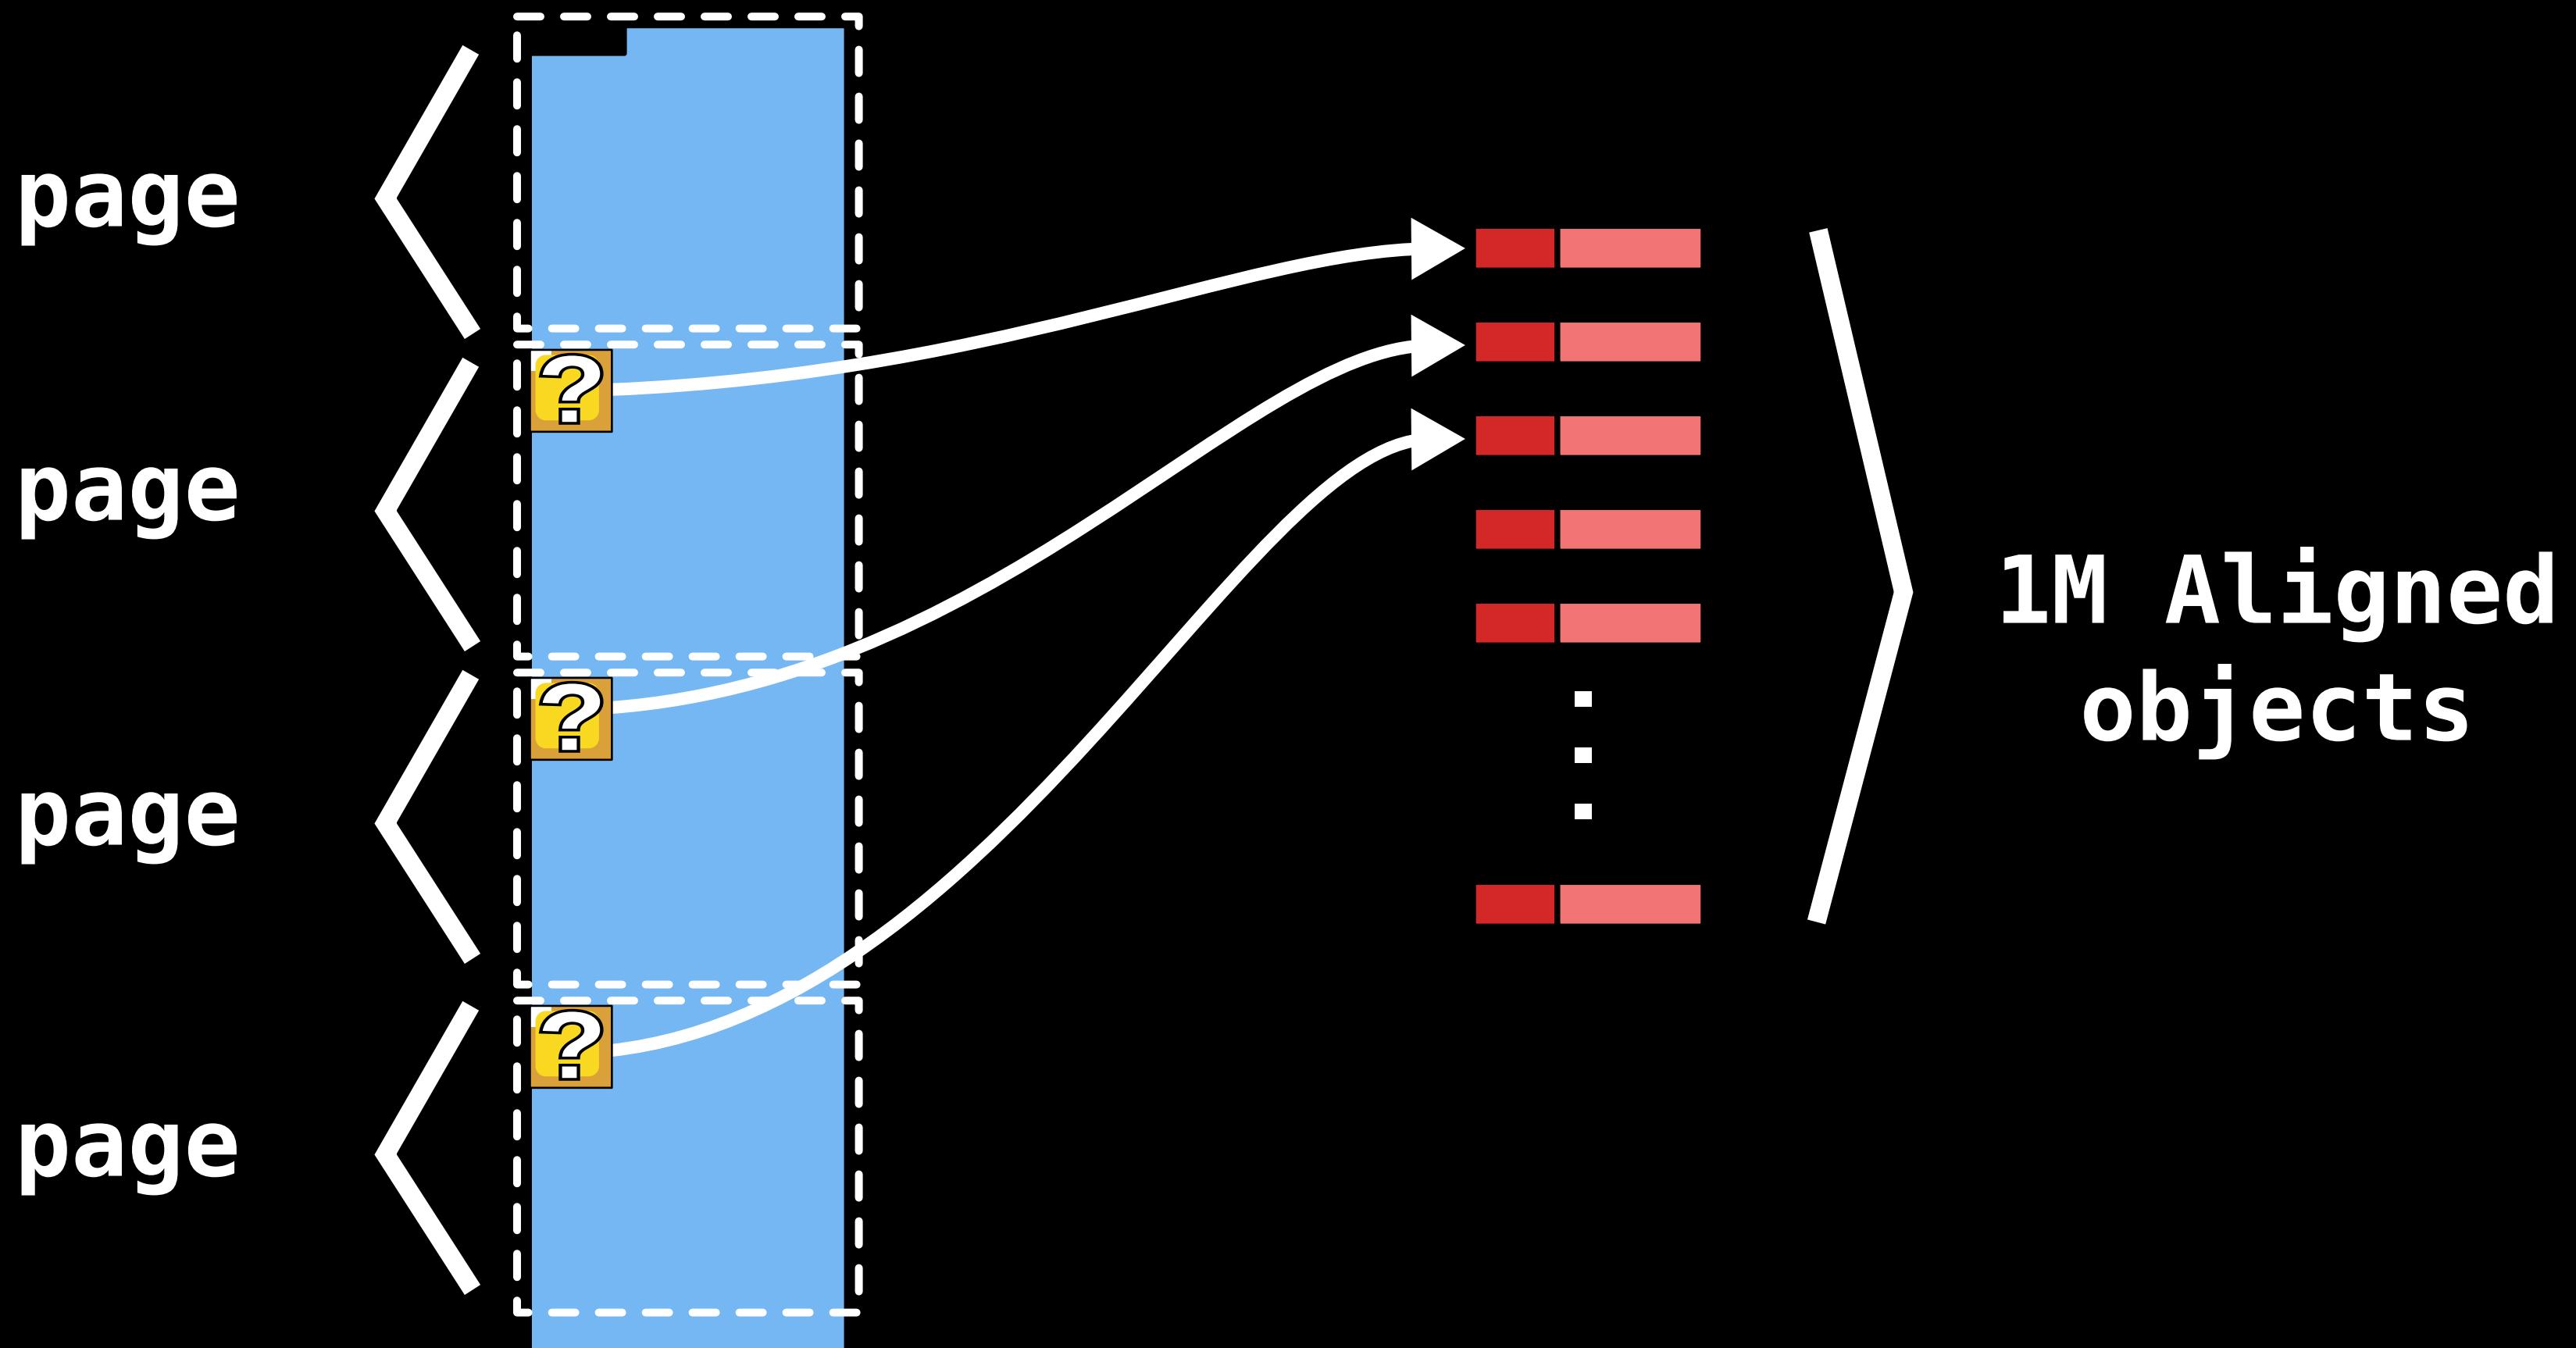


1M Aligned  
objects

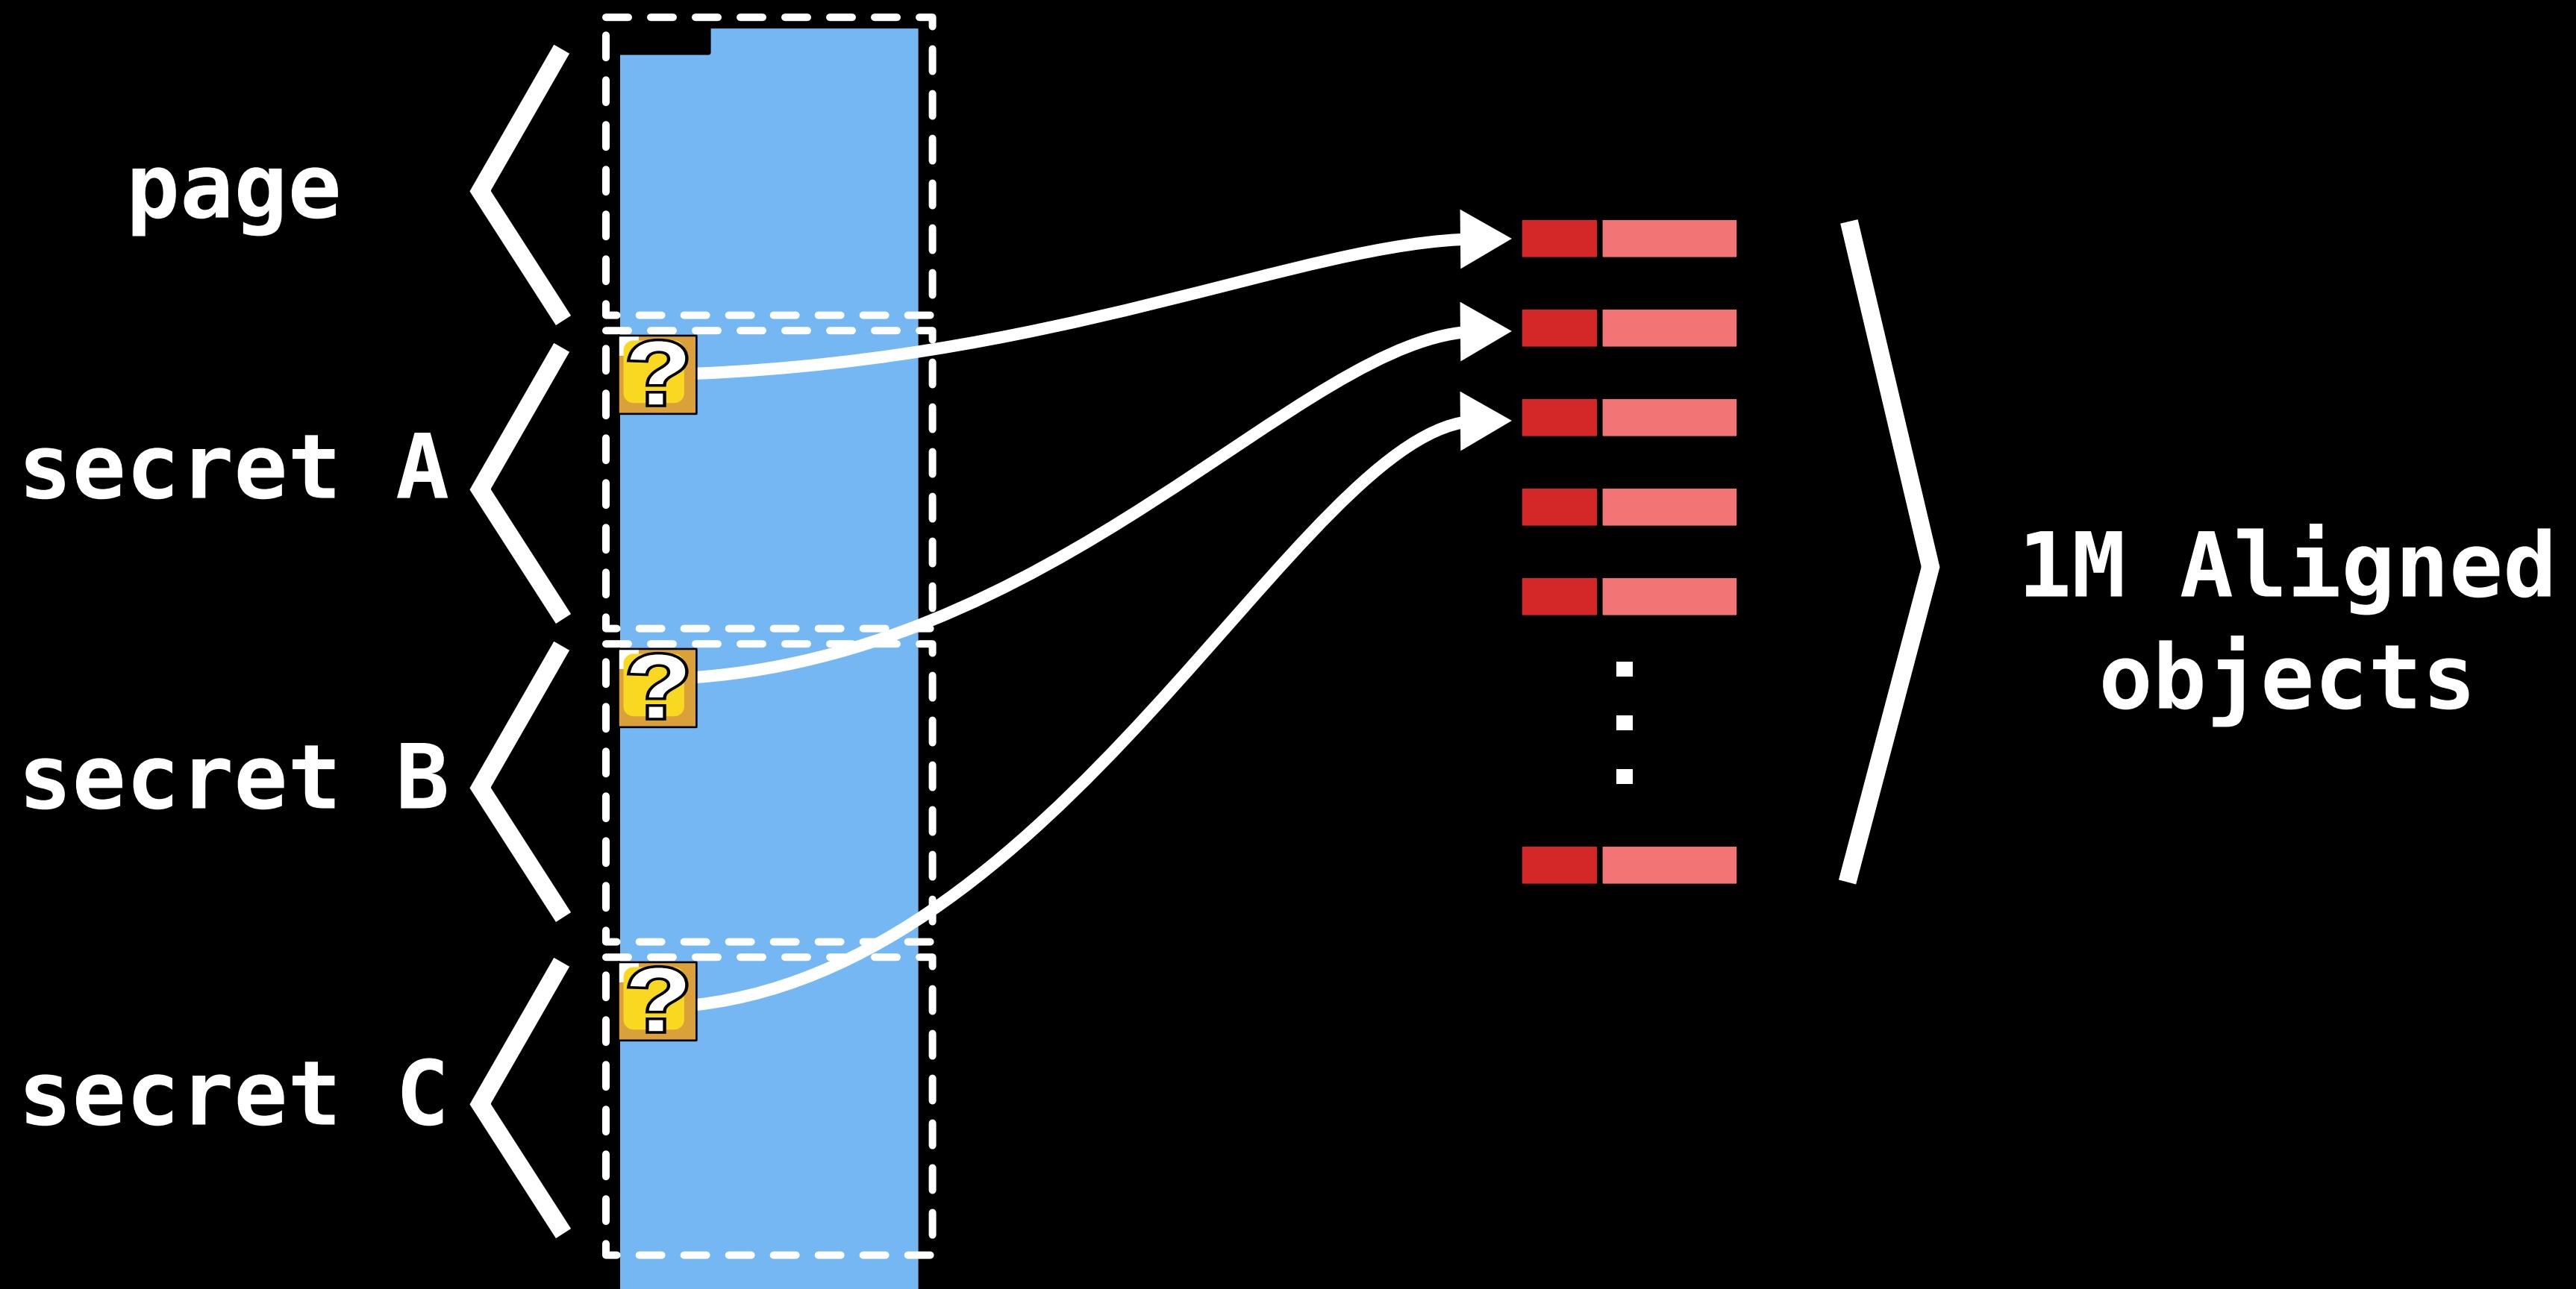
# Creating Secret Pages



# Creating Secret Pages



# Creating Secret Pages

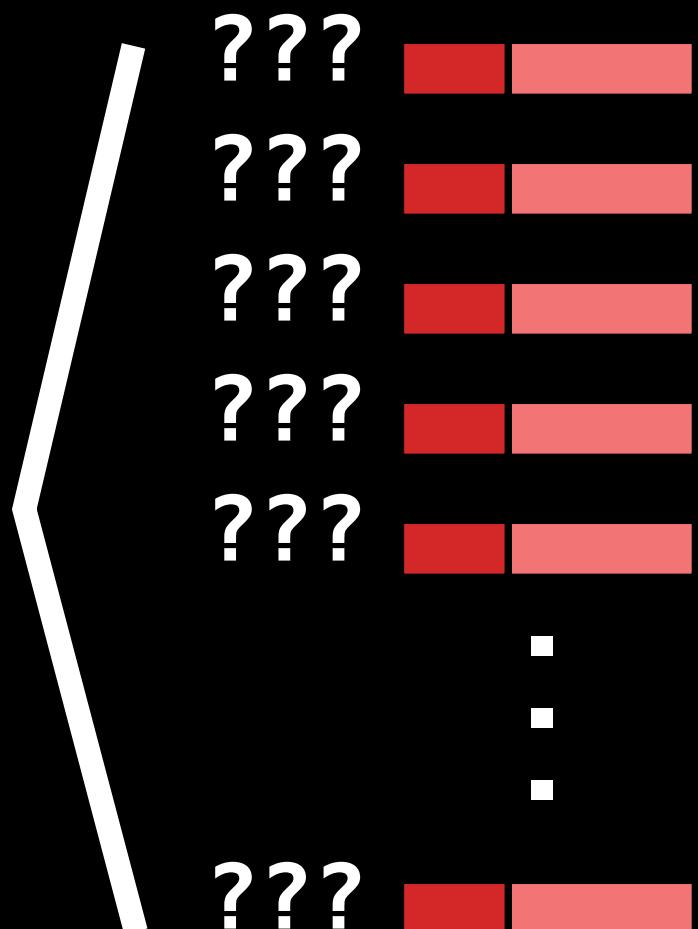


# Creating Probe Pages

**typed  
array  
data**

# Creating Probe Pages

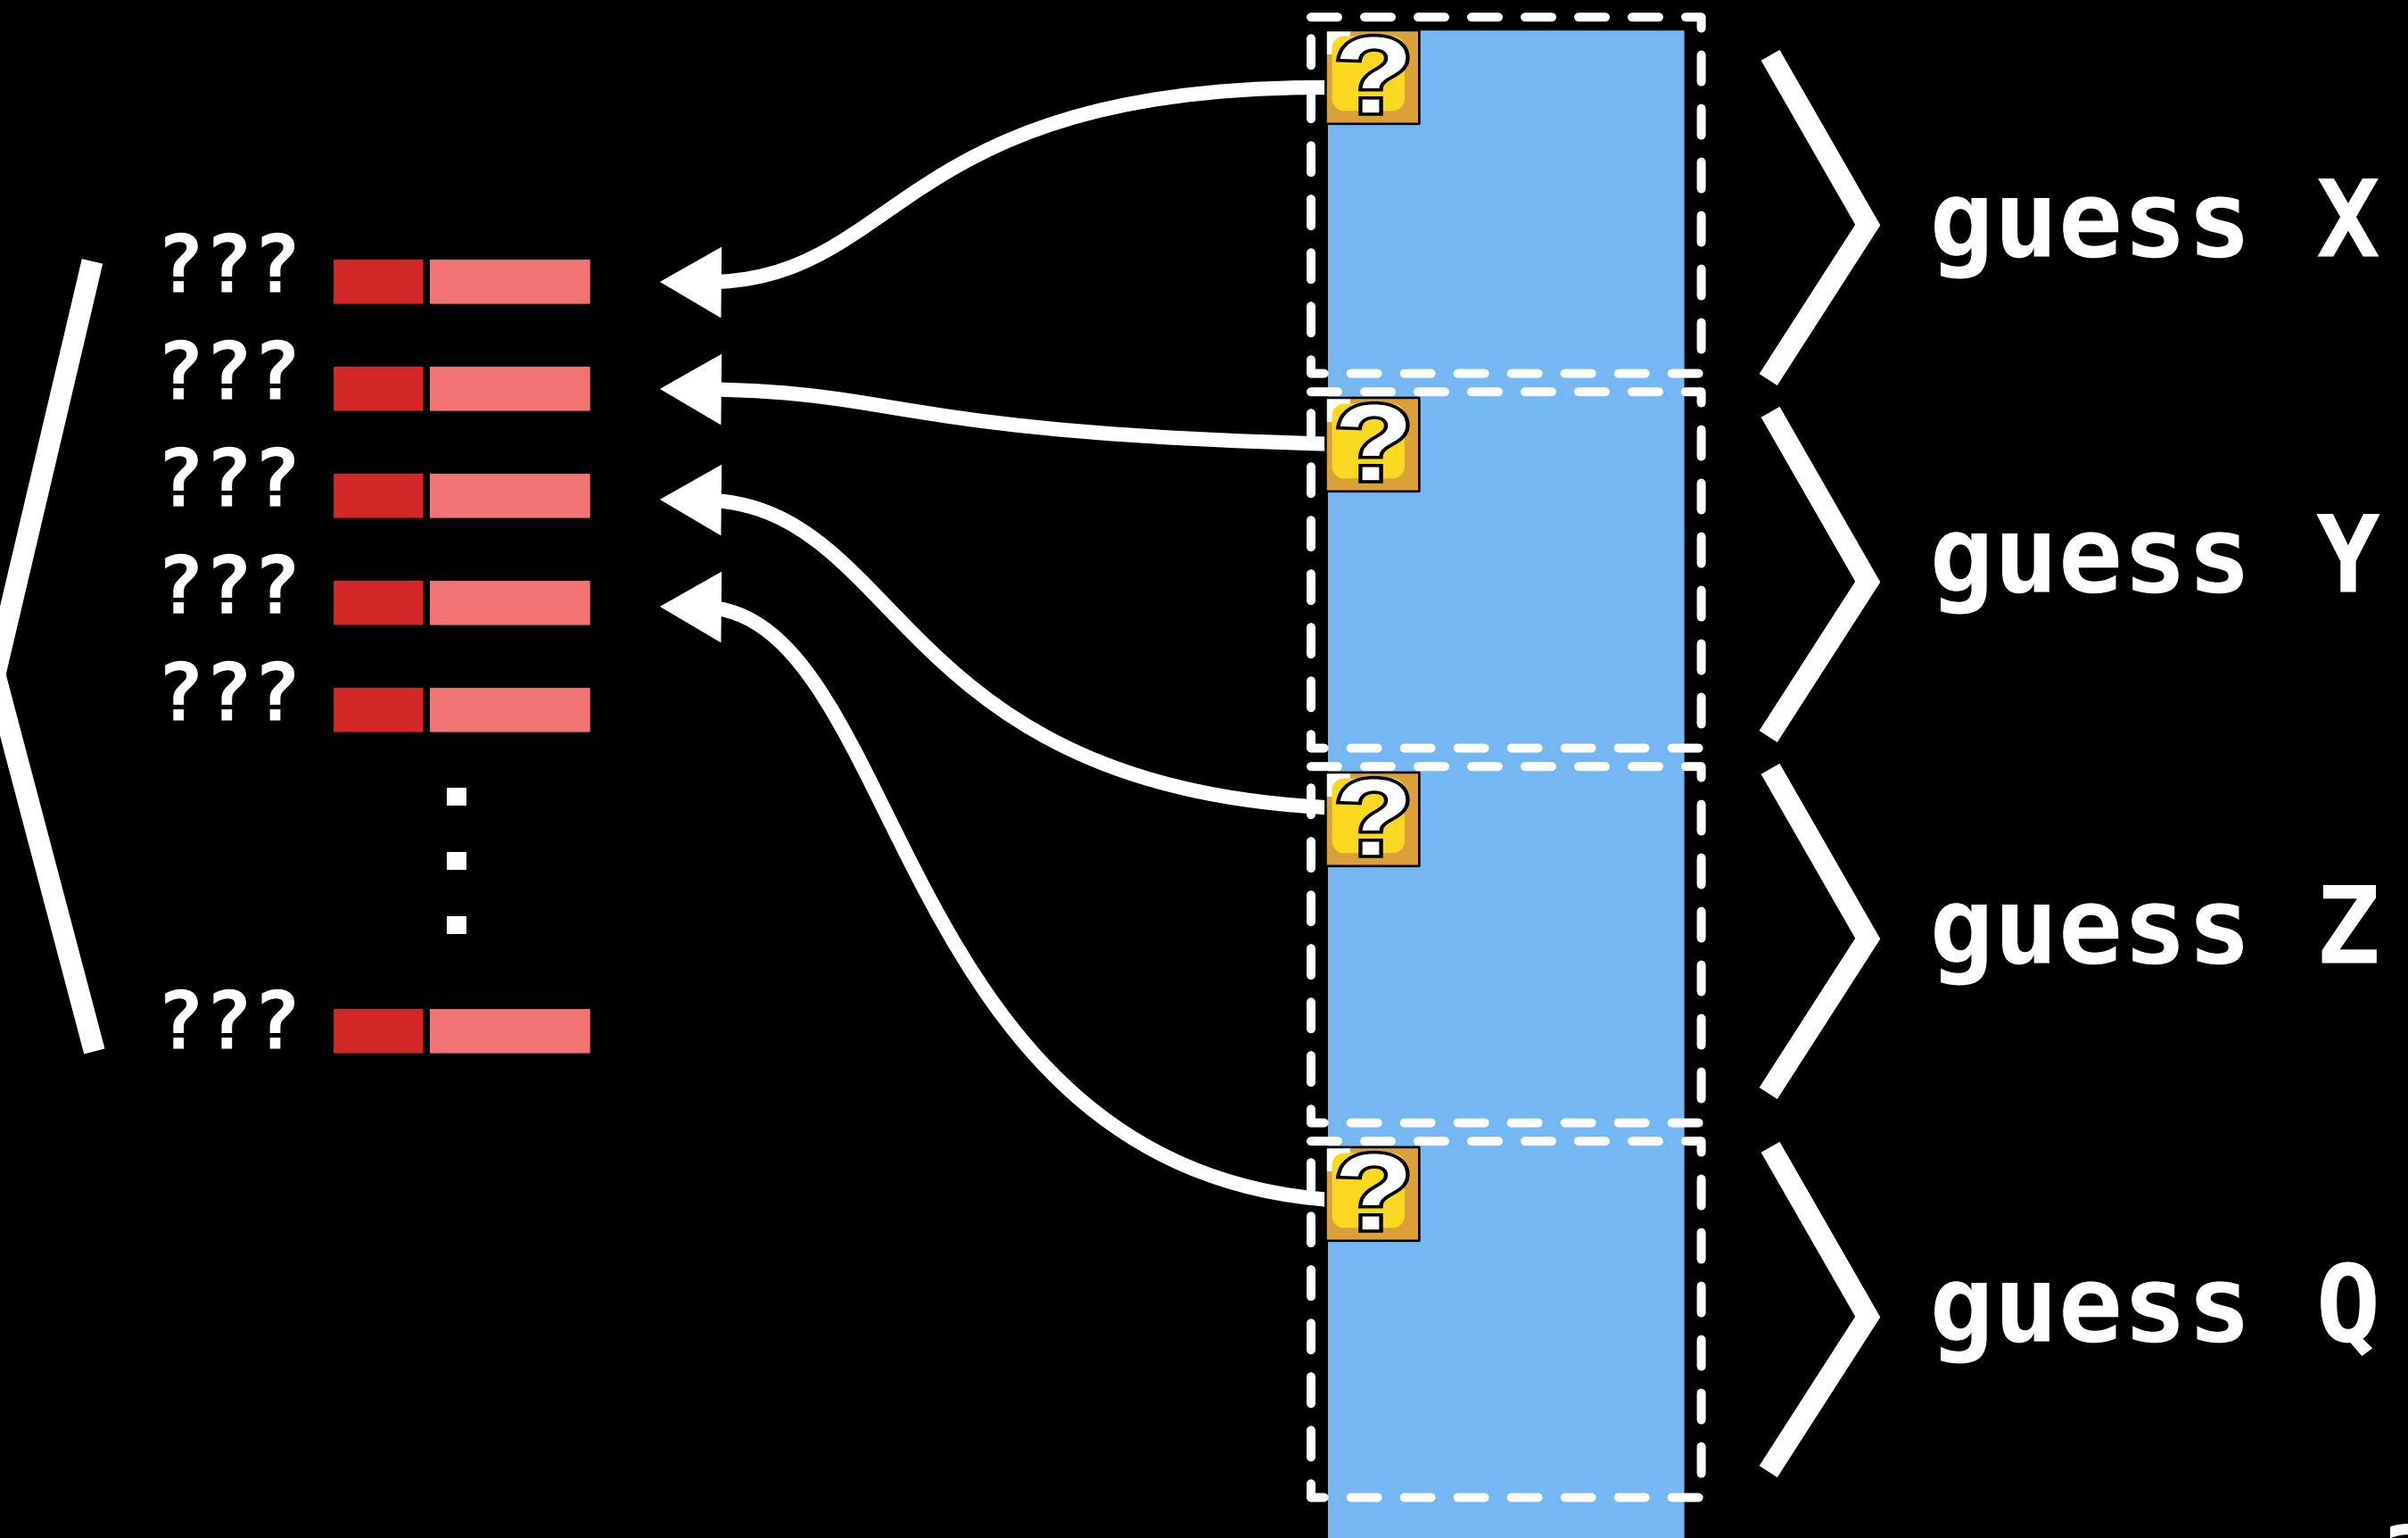
guessed  
aligned  
addresses,  
128M apart



typed  
array  
data

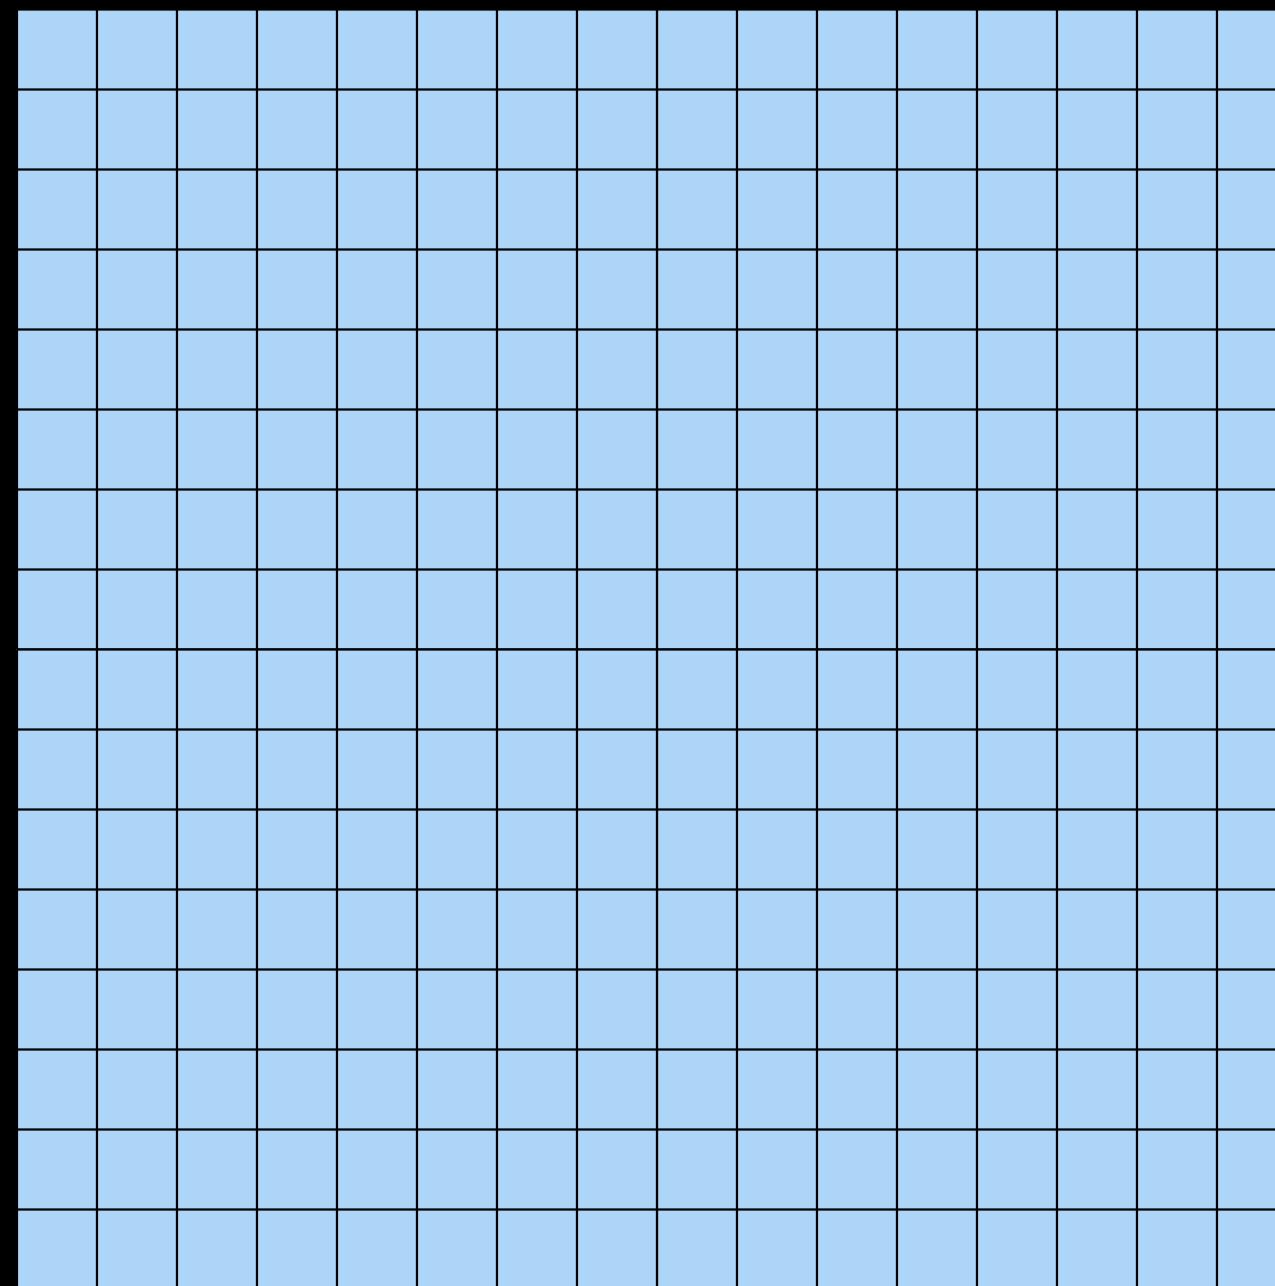
# Creating Probe Pages

guessed  
aligned  
addresses,  
128M apart



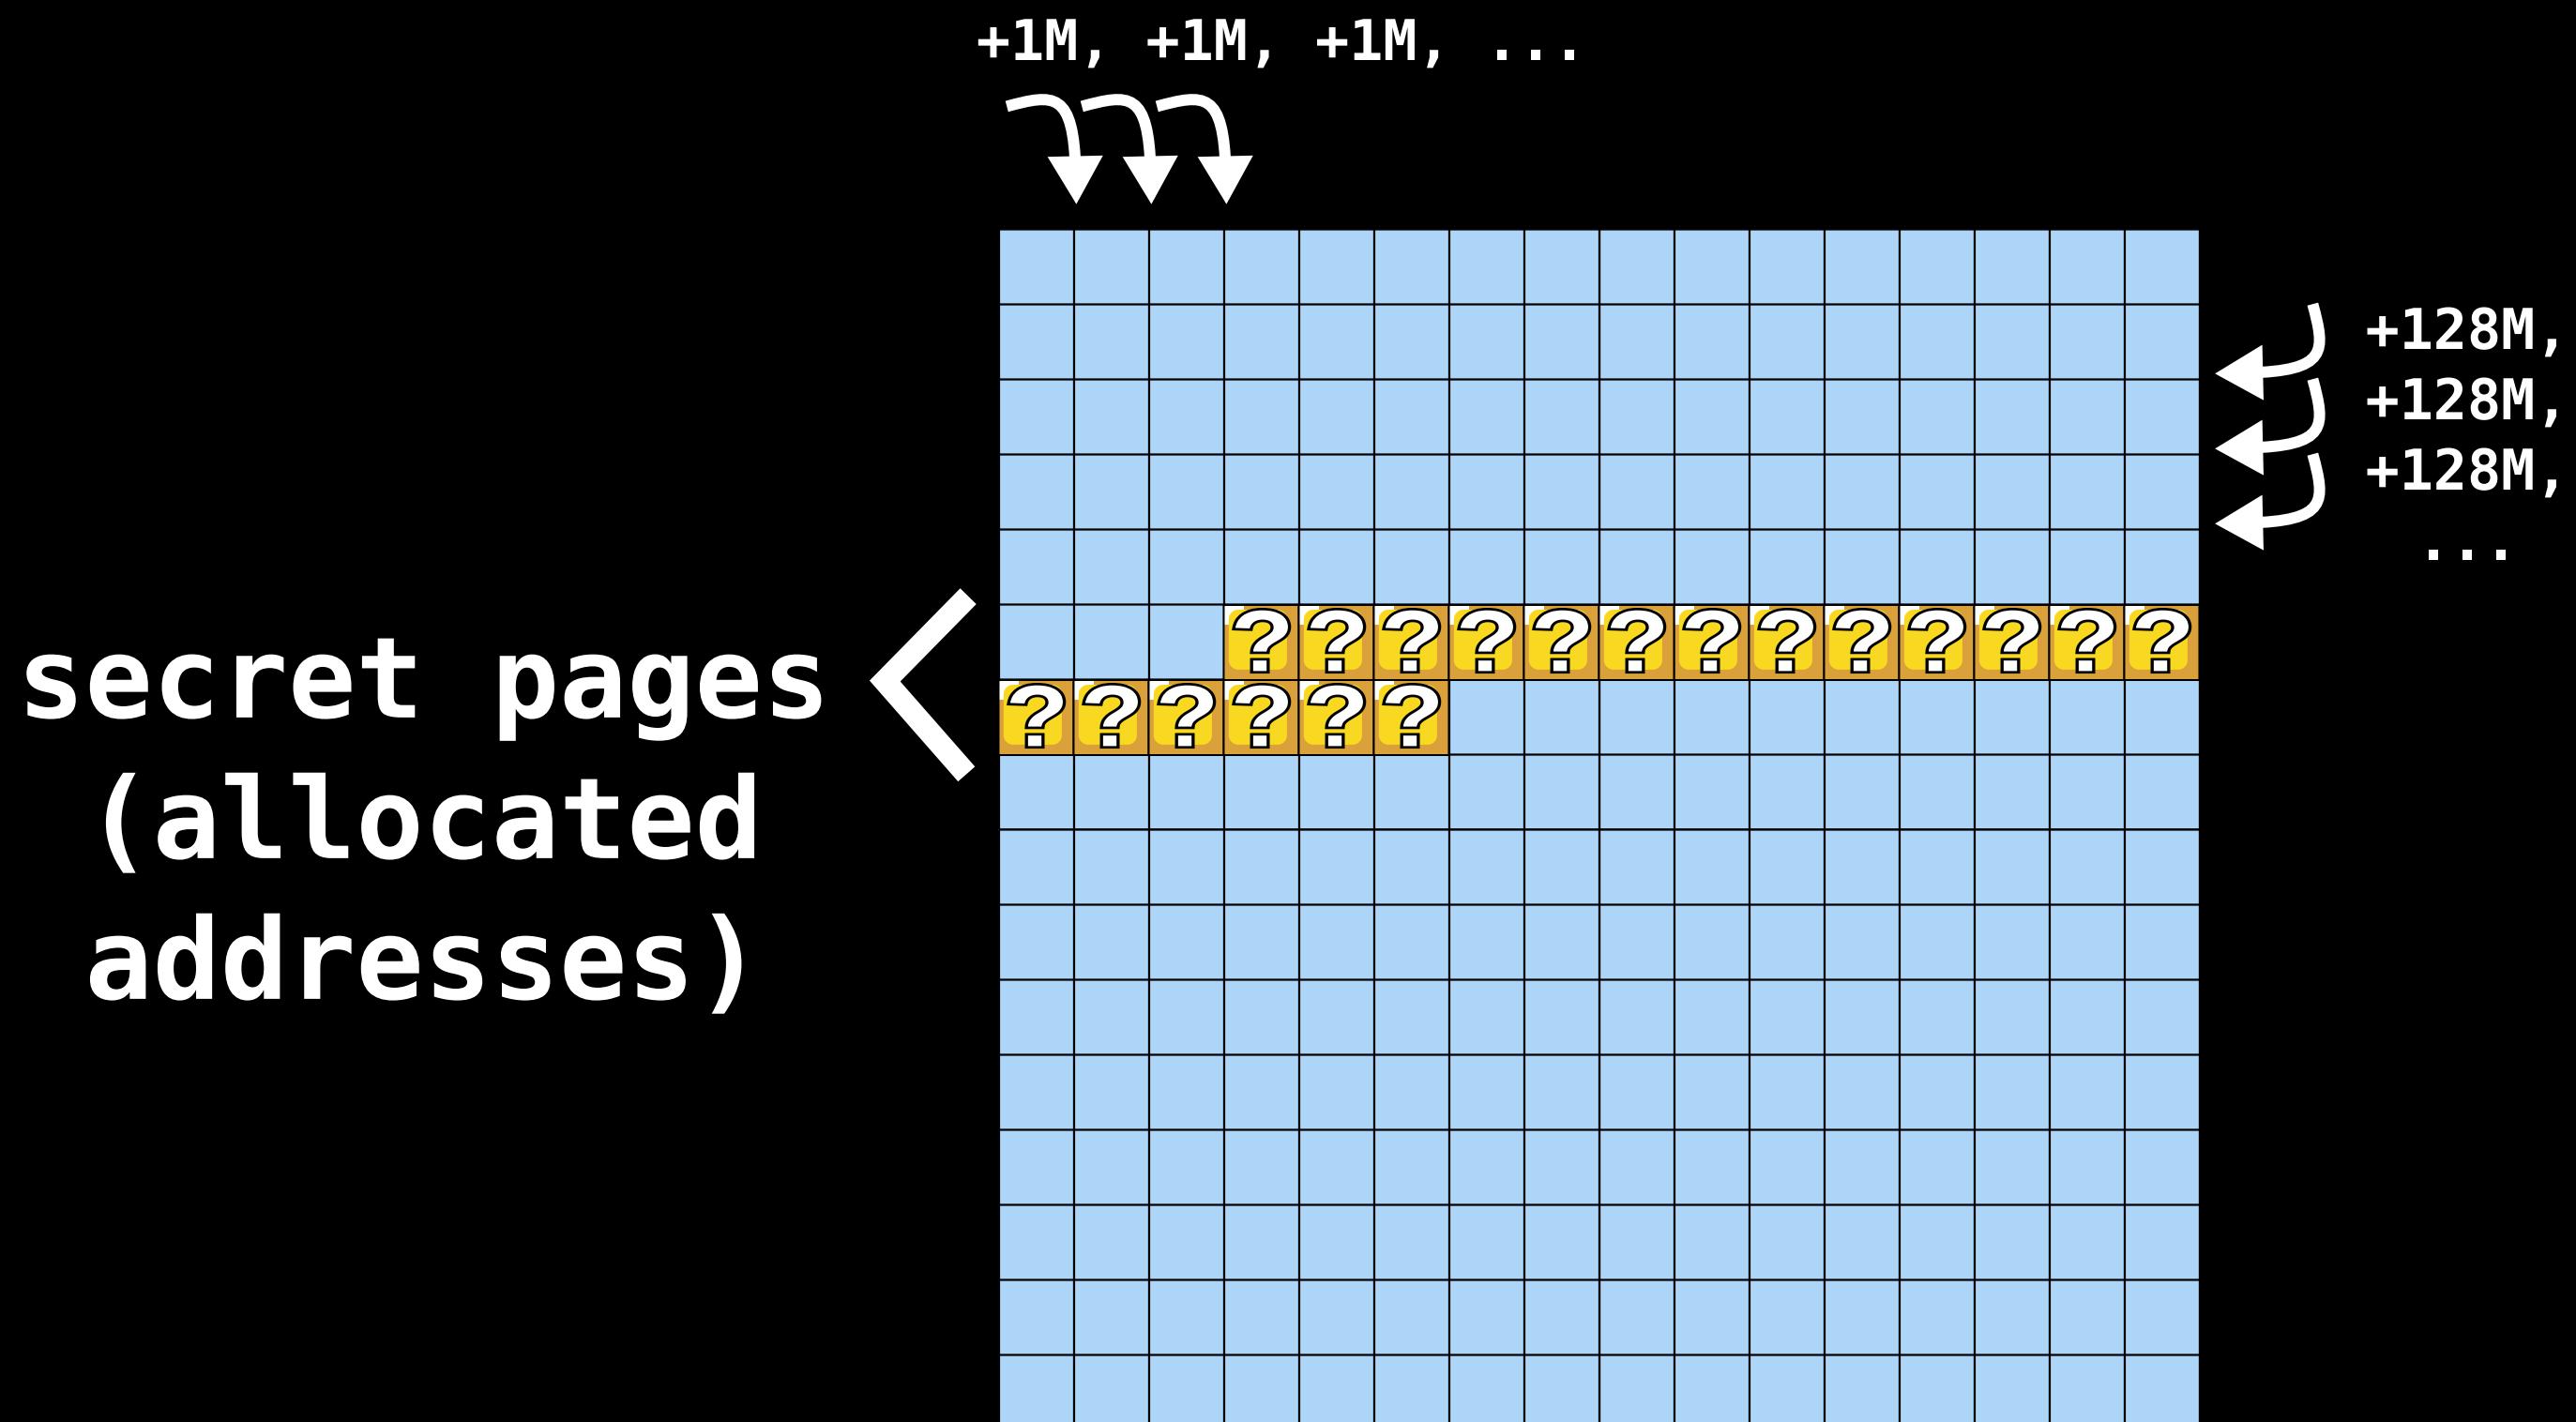
# Birthday heap spray

+1M, +1M, +1M, ...

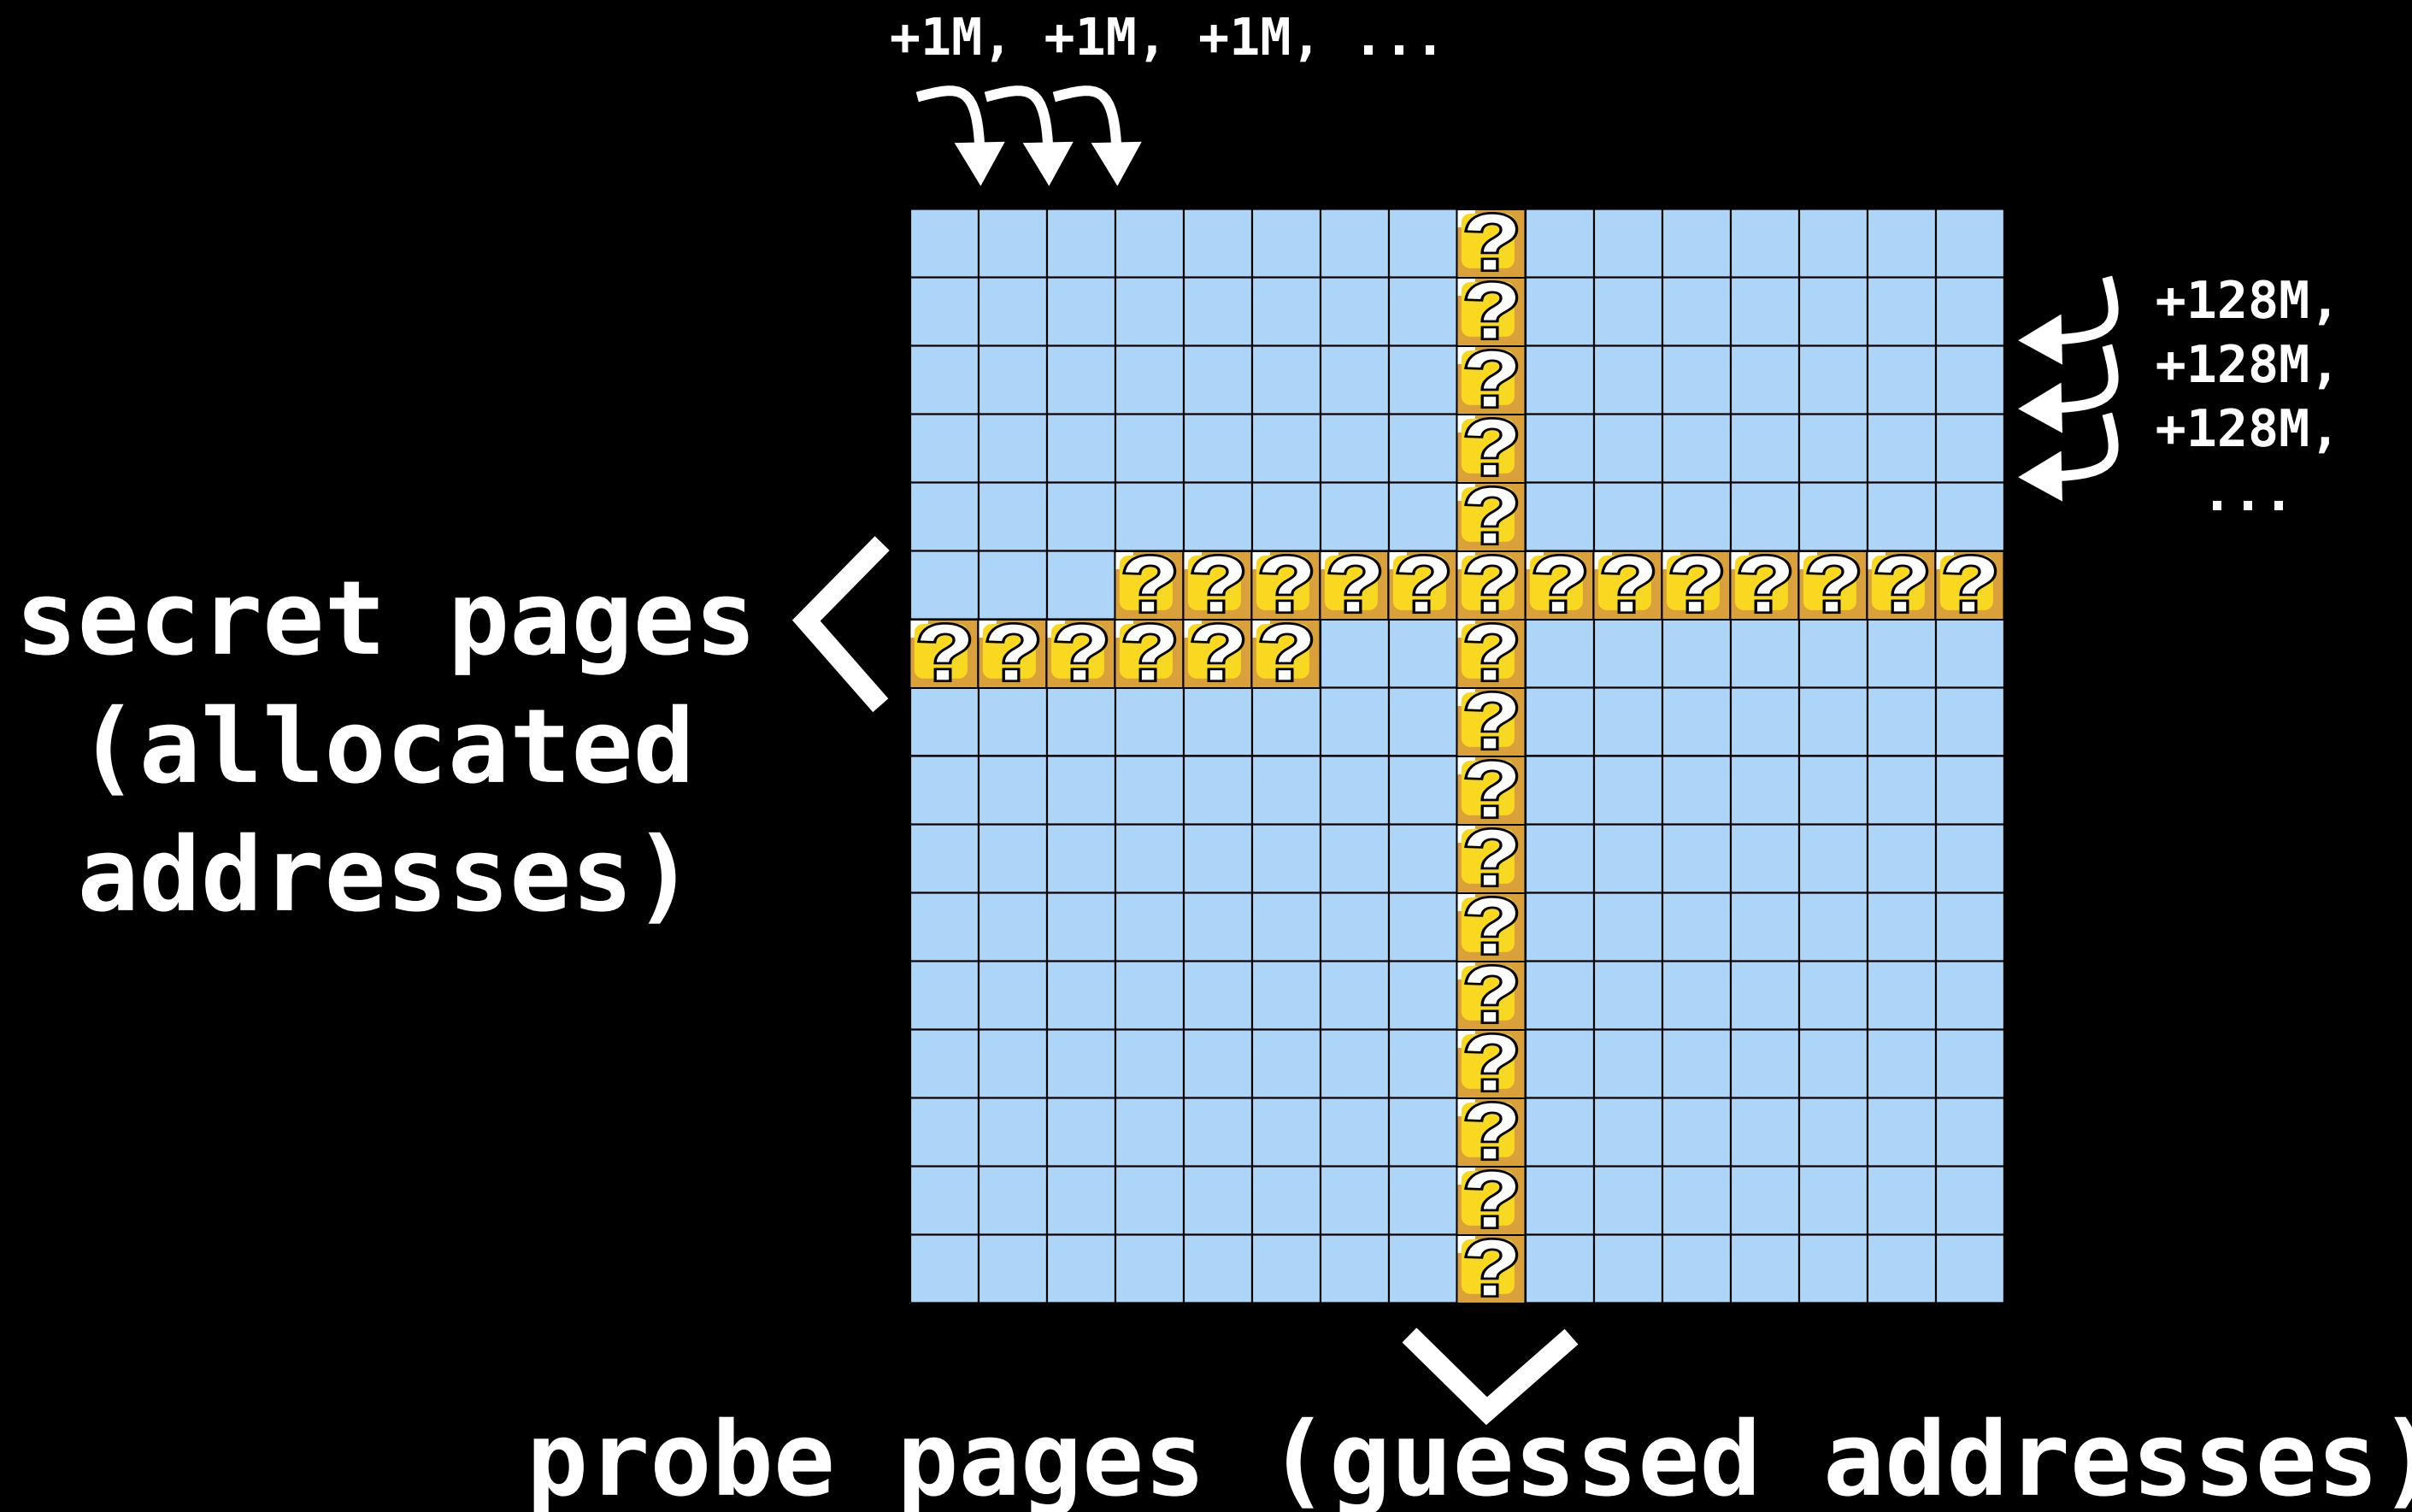


+128M,  
+128M,  
+128M,  
...  
↑  
↑  
↑

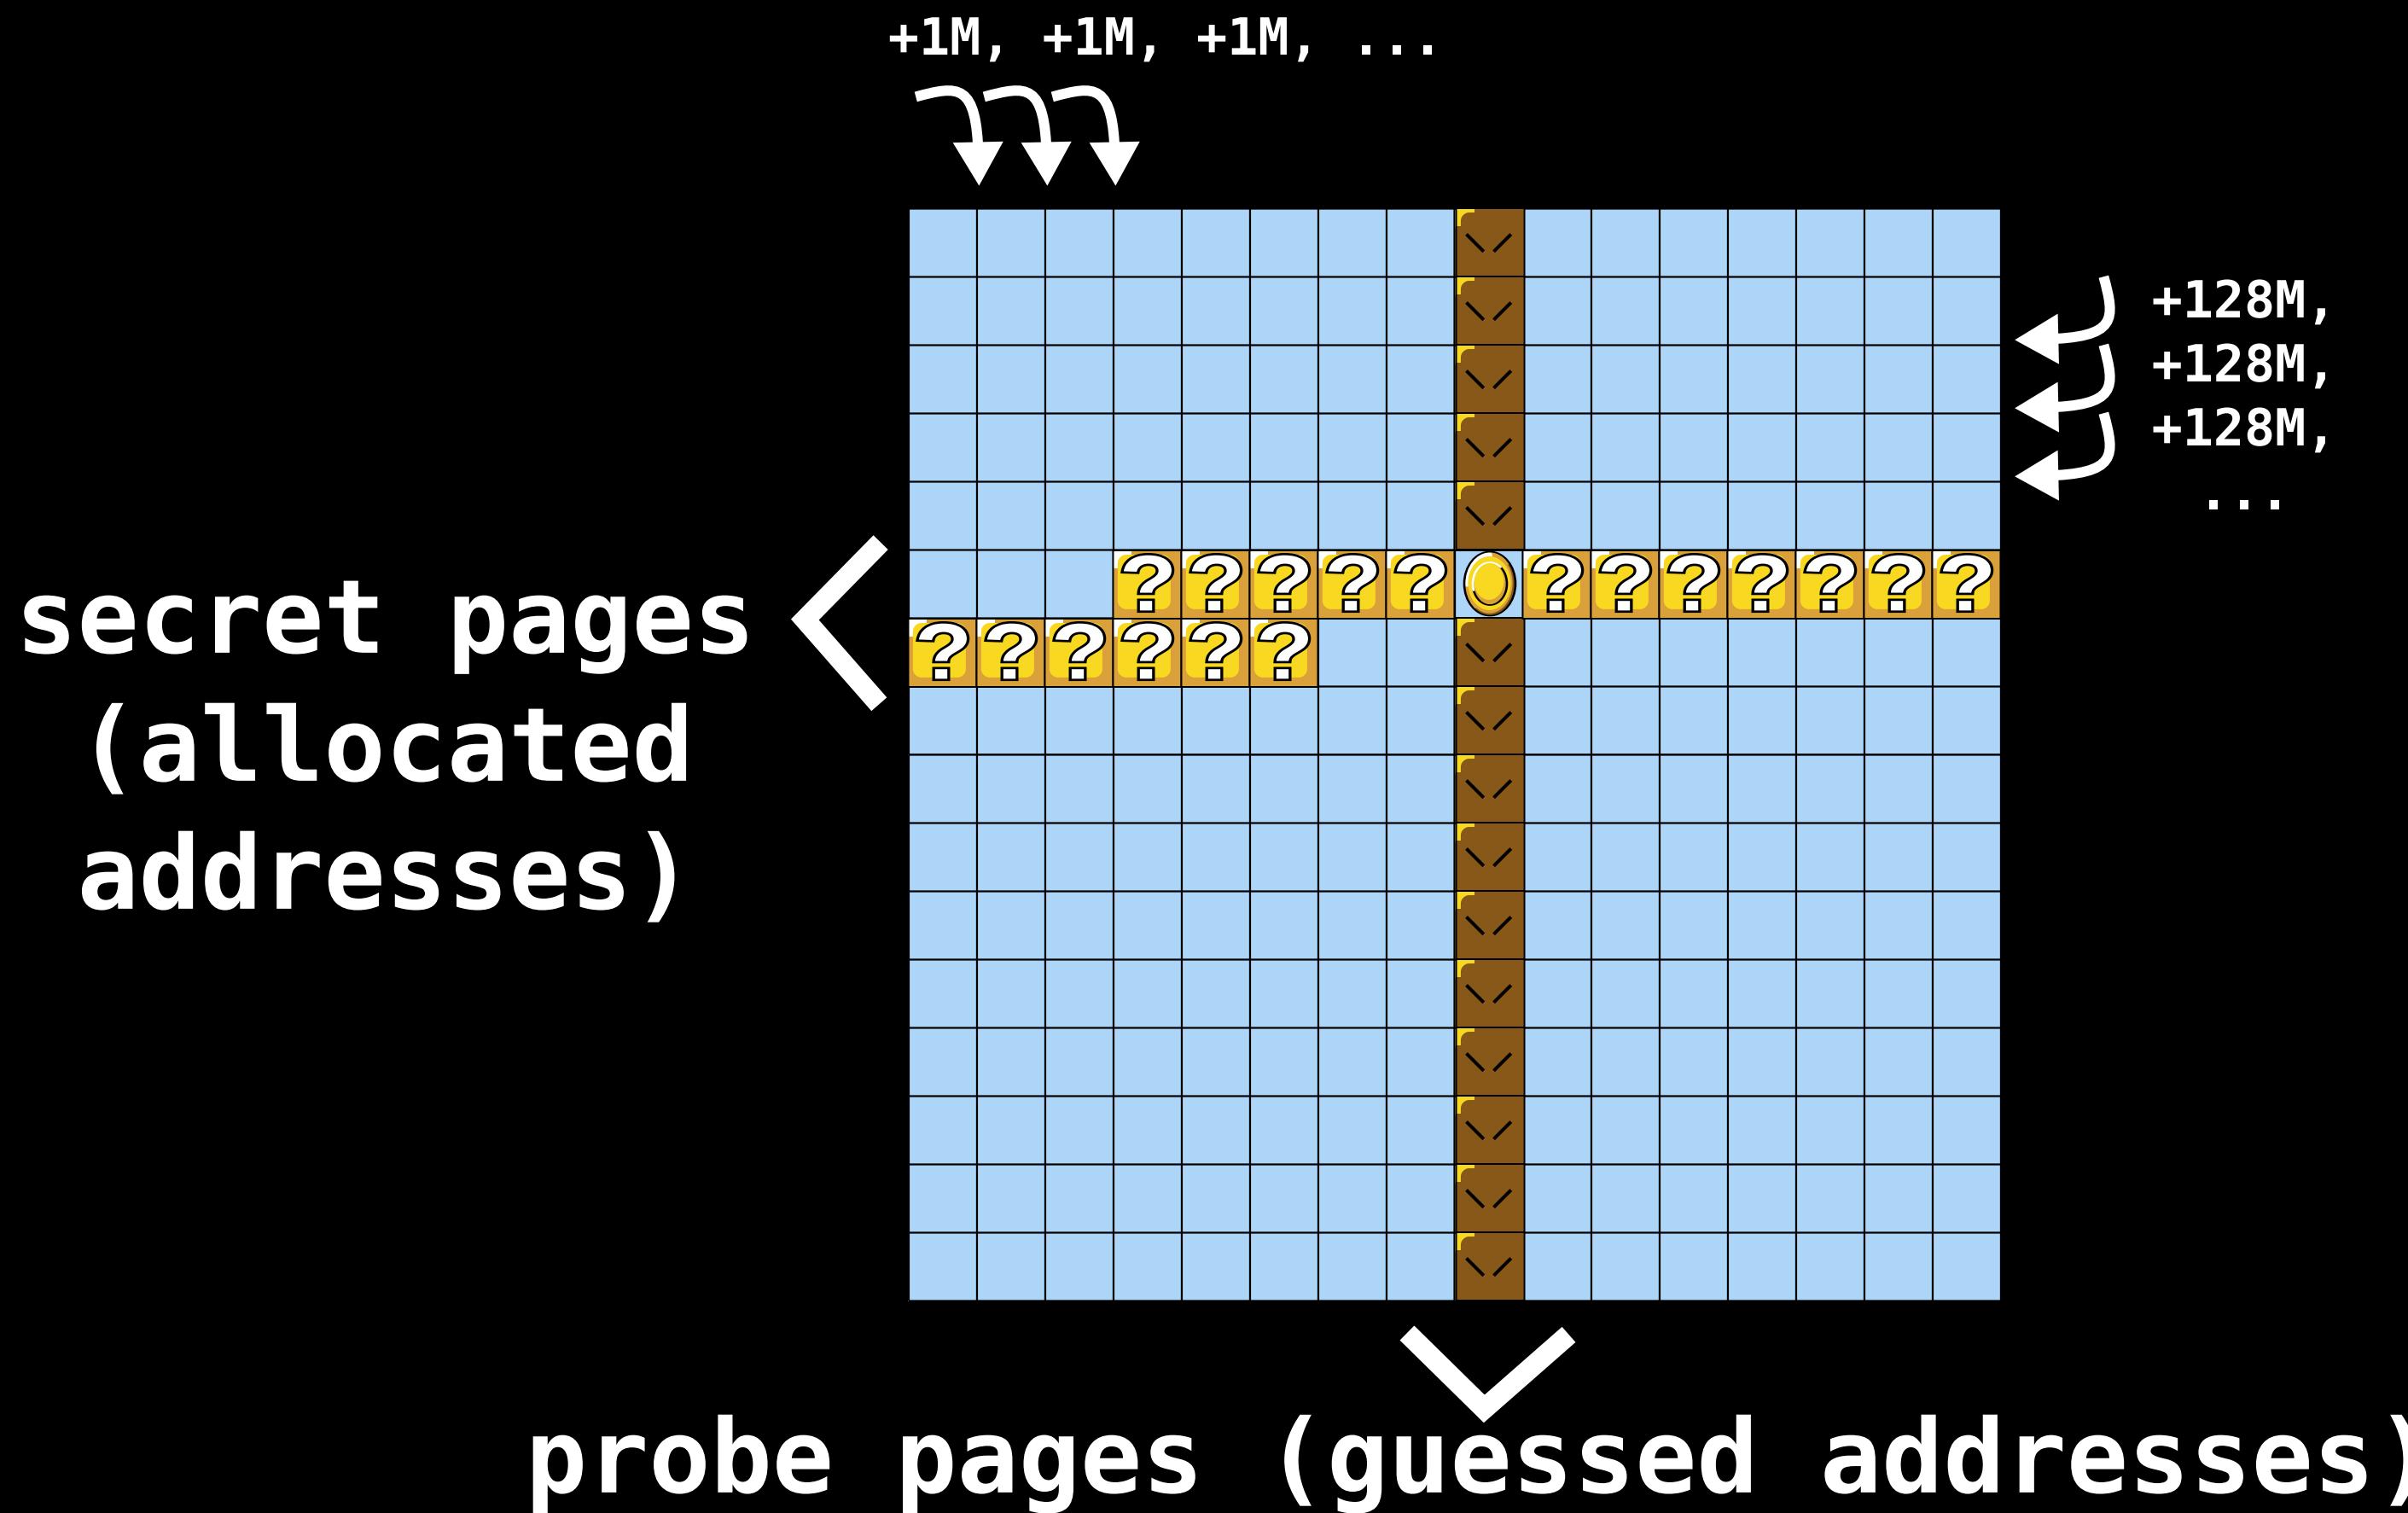
# Birthday heap spray



# Birthday heap spray



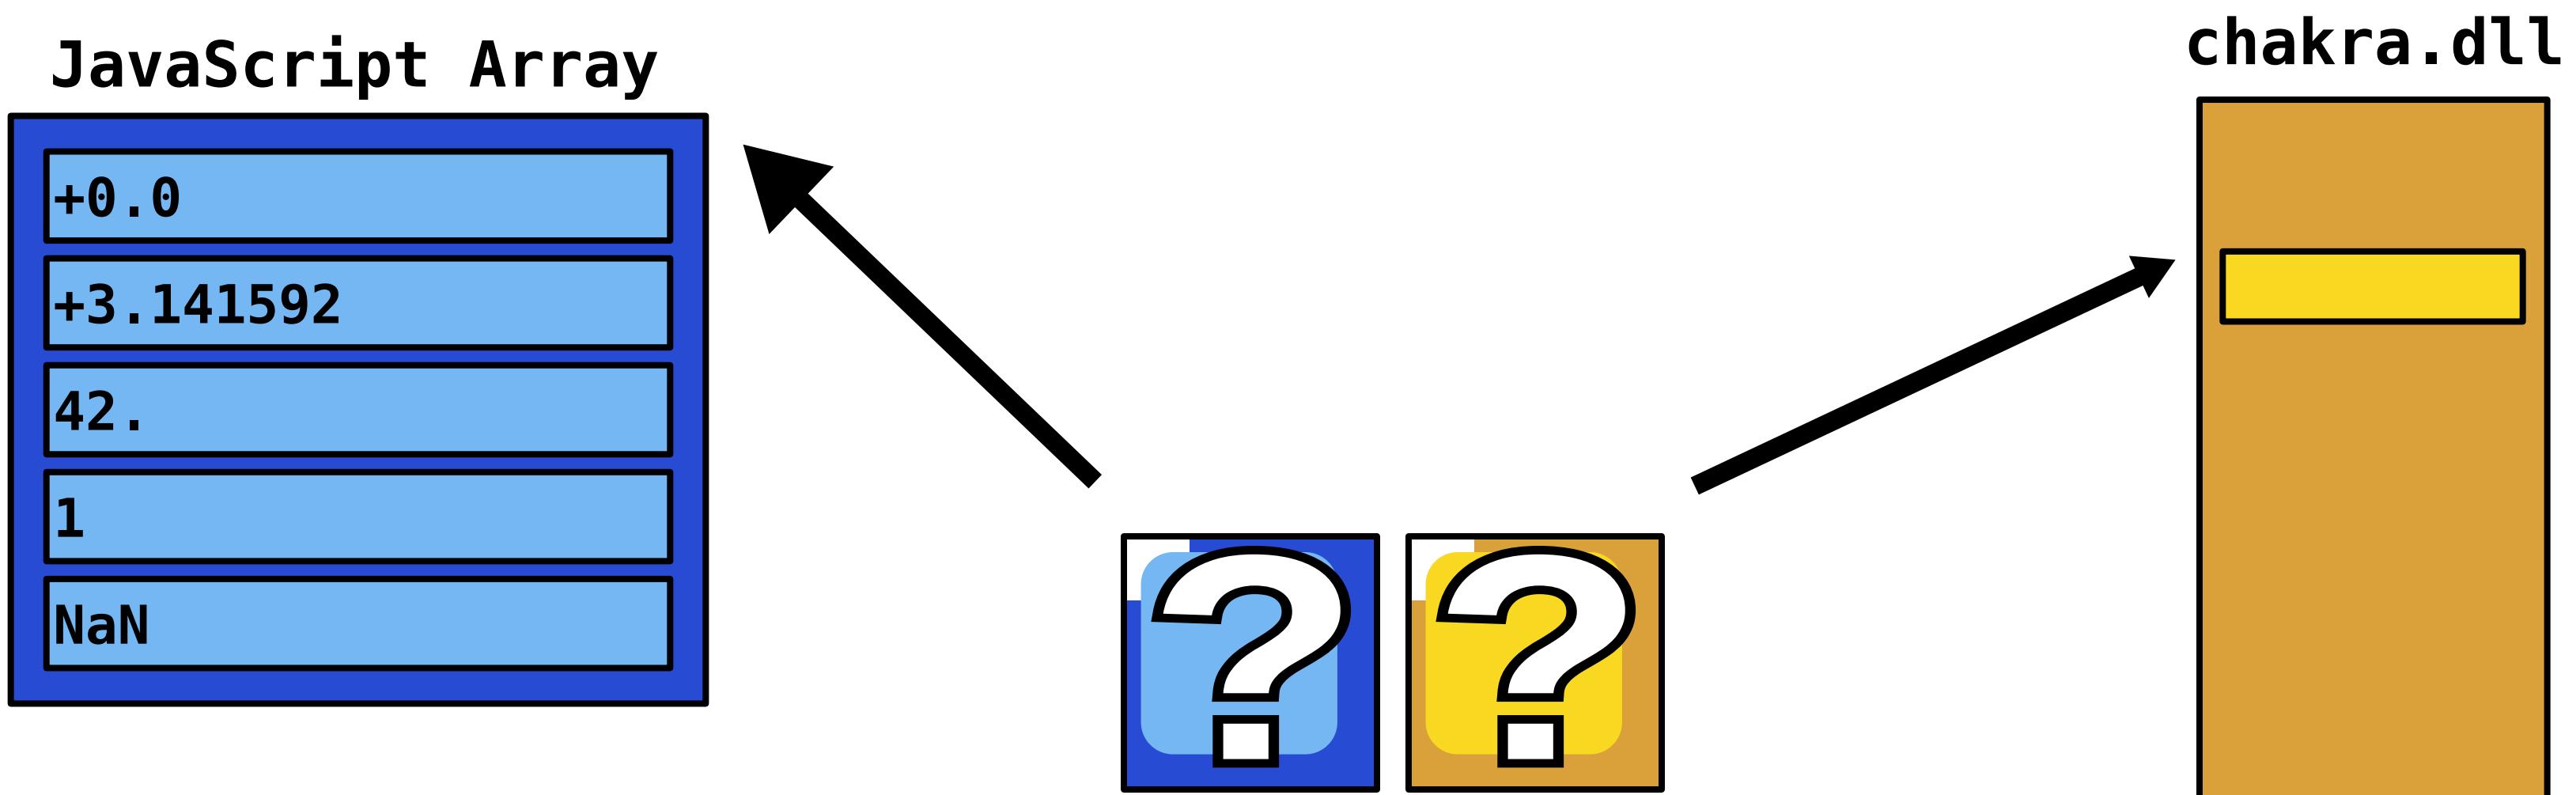
# Birthday heap spray



# Outline:

## Deduplication

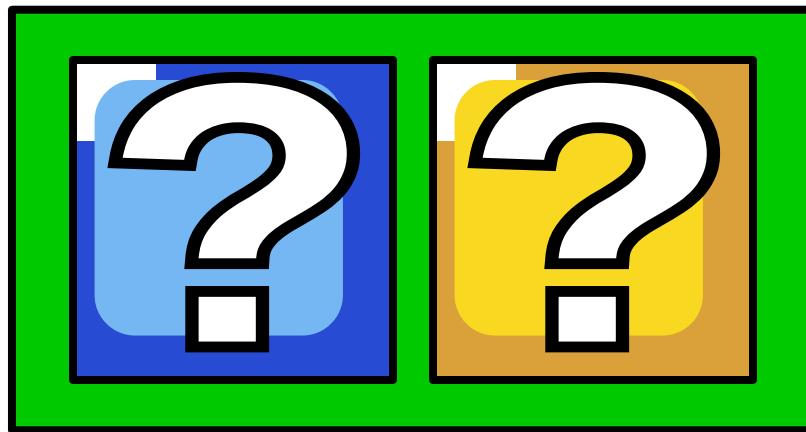
- leak heap & code addresses



# Outline:

## Deduplication

- leak heap & code addresses
- create a fake object



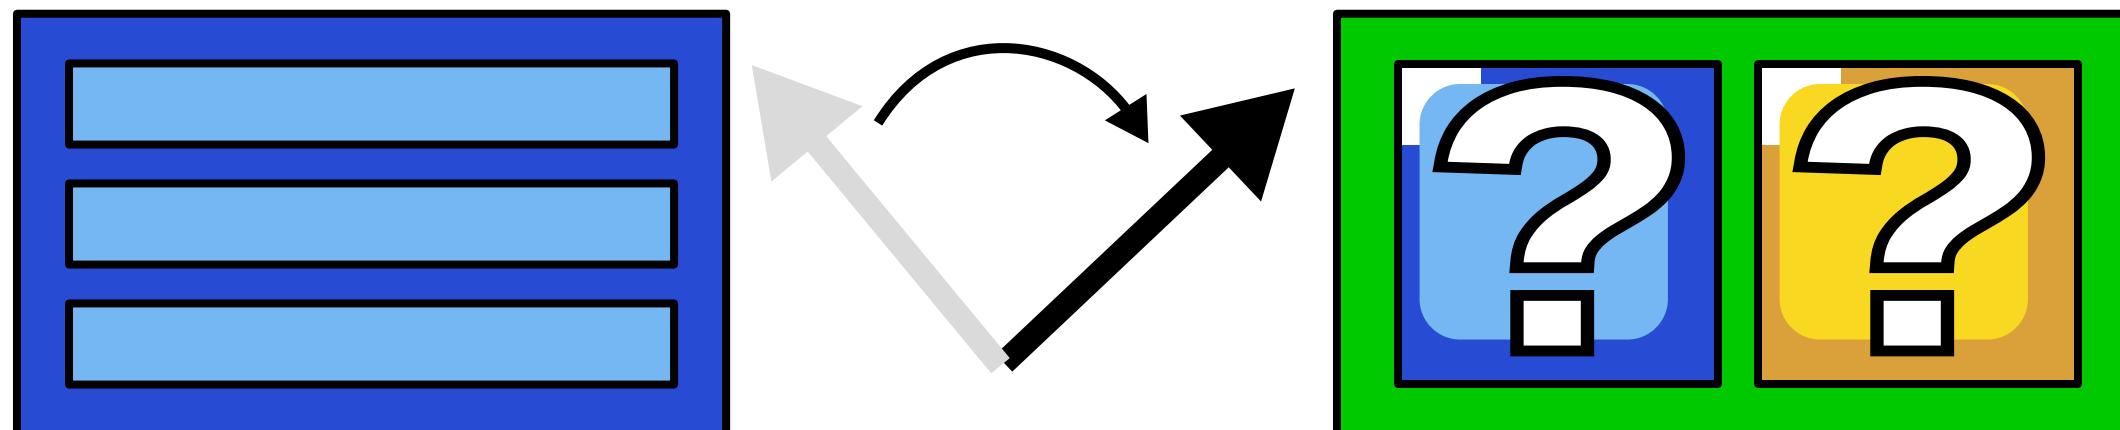
# Outline:

## Deduplication

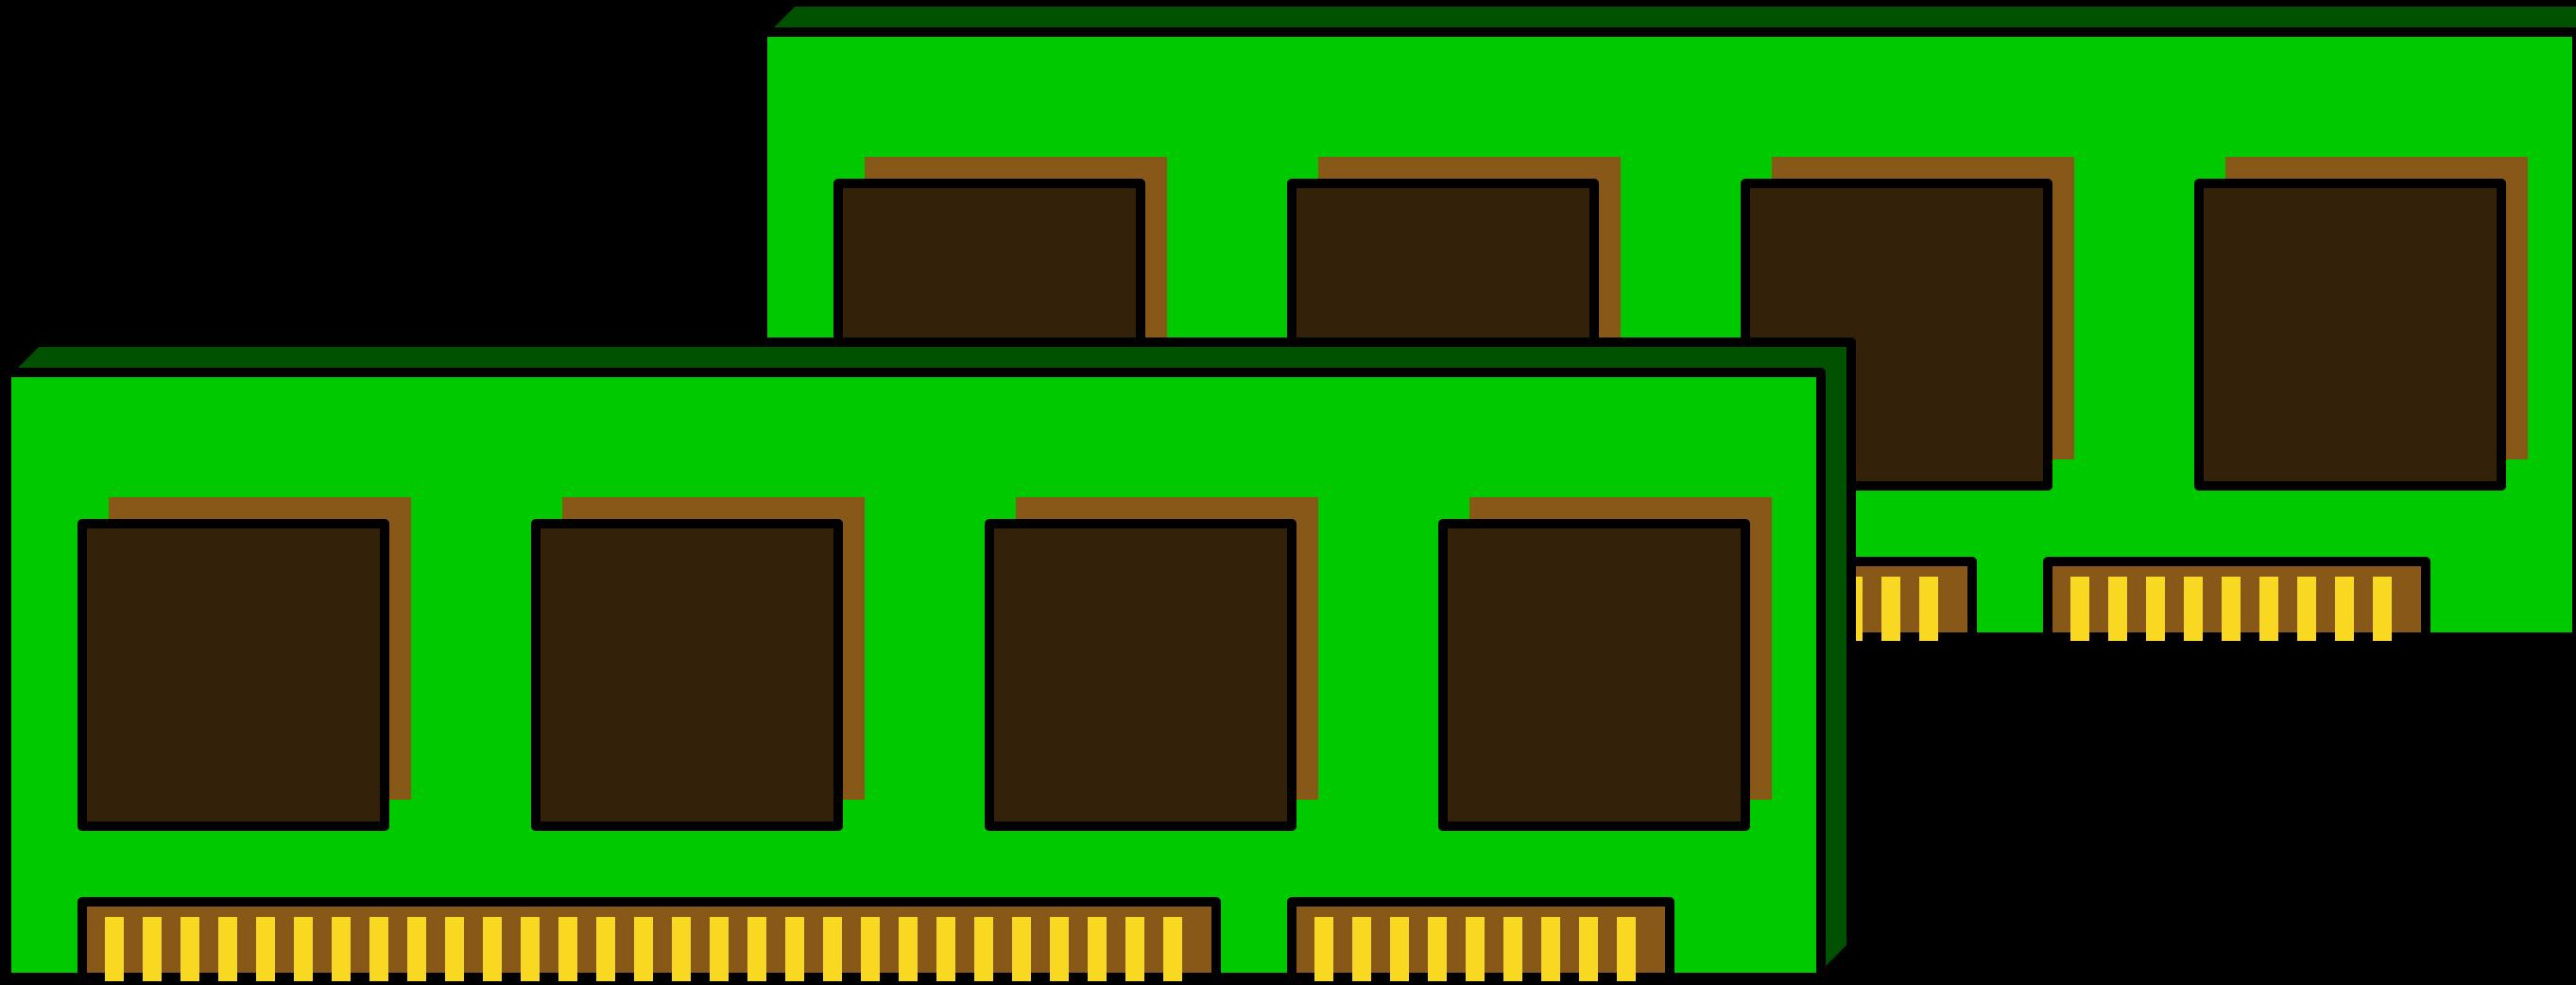
- leak heap & code addresses
- create a fake object

## Rowhammer

- create reference to our fake object

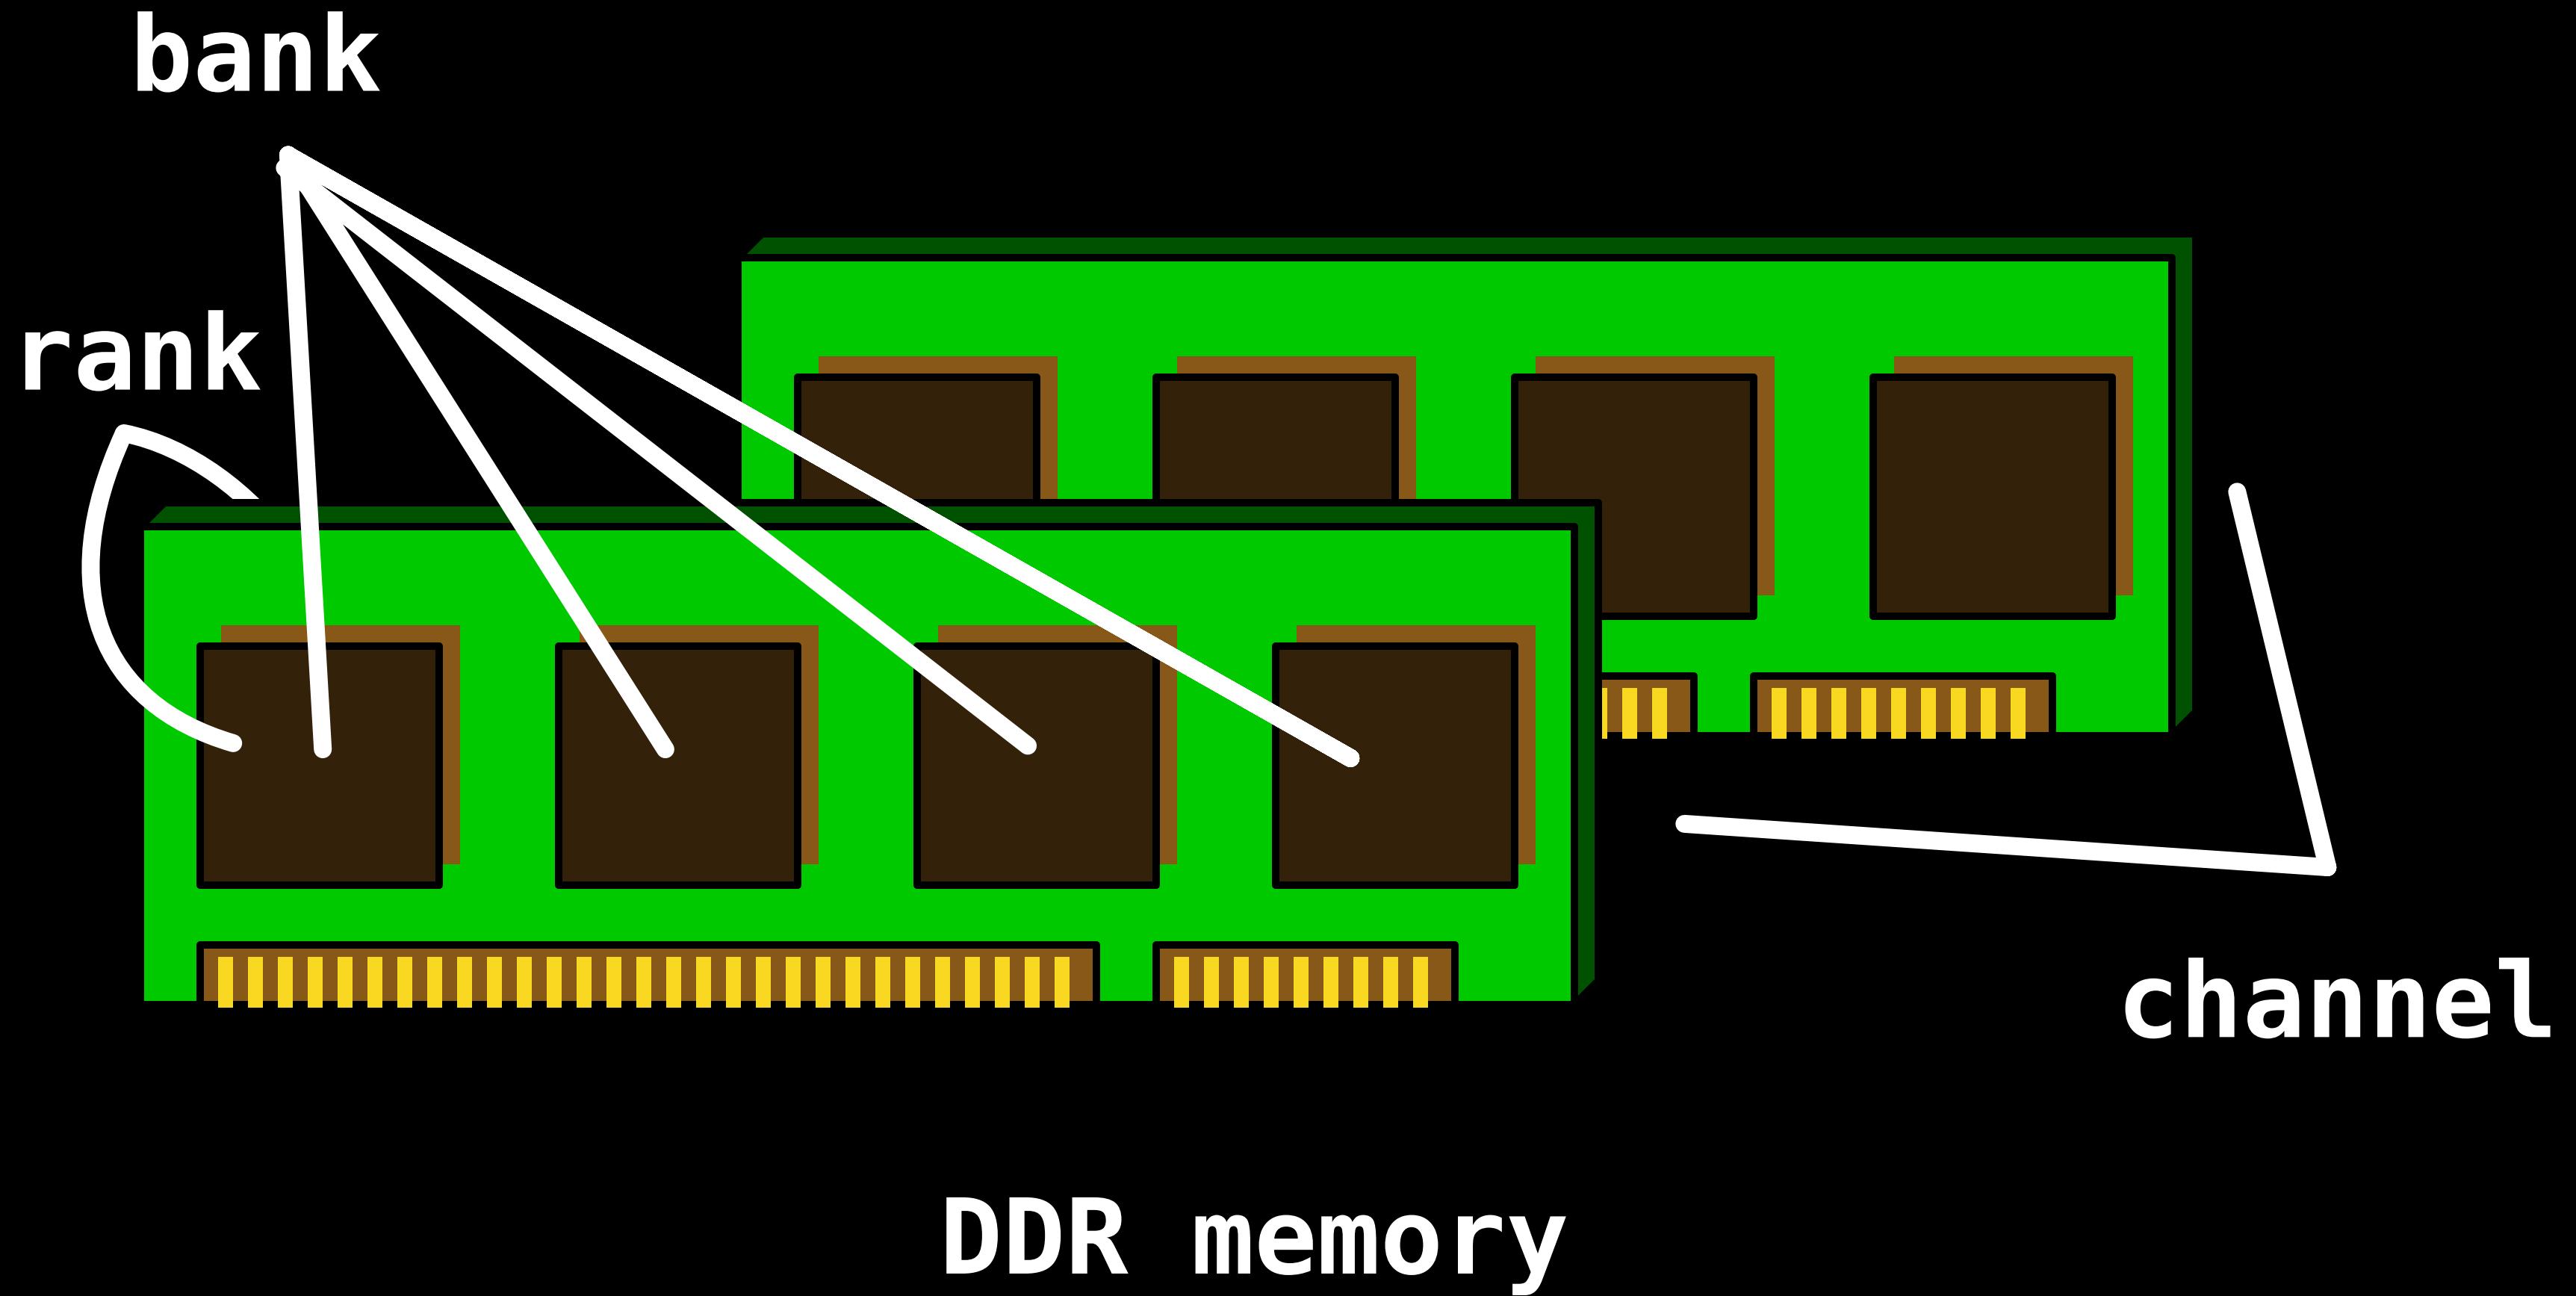


# rowhammer attack

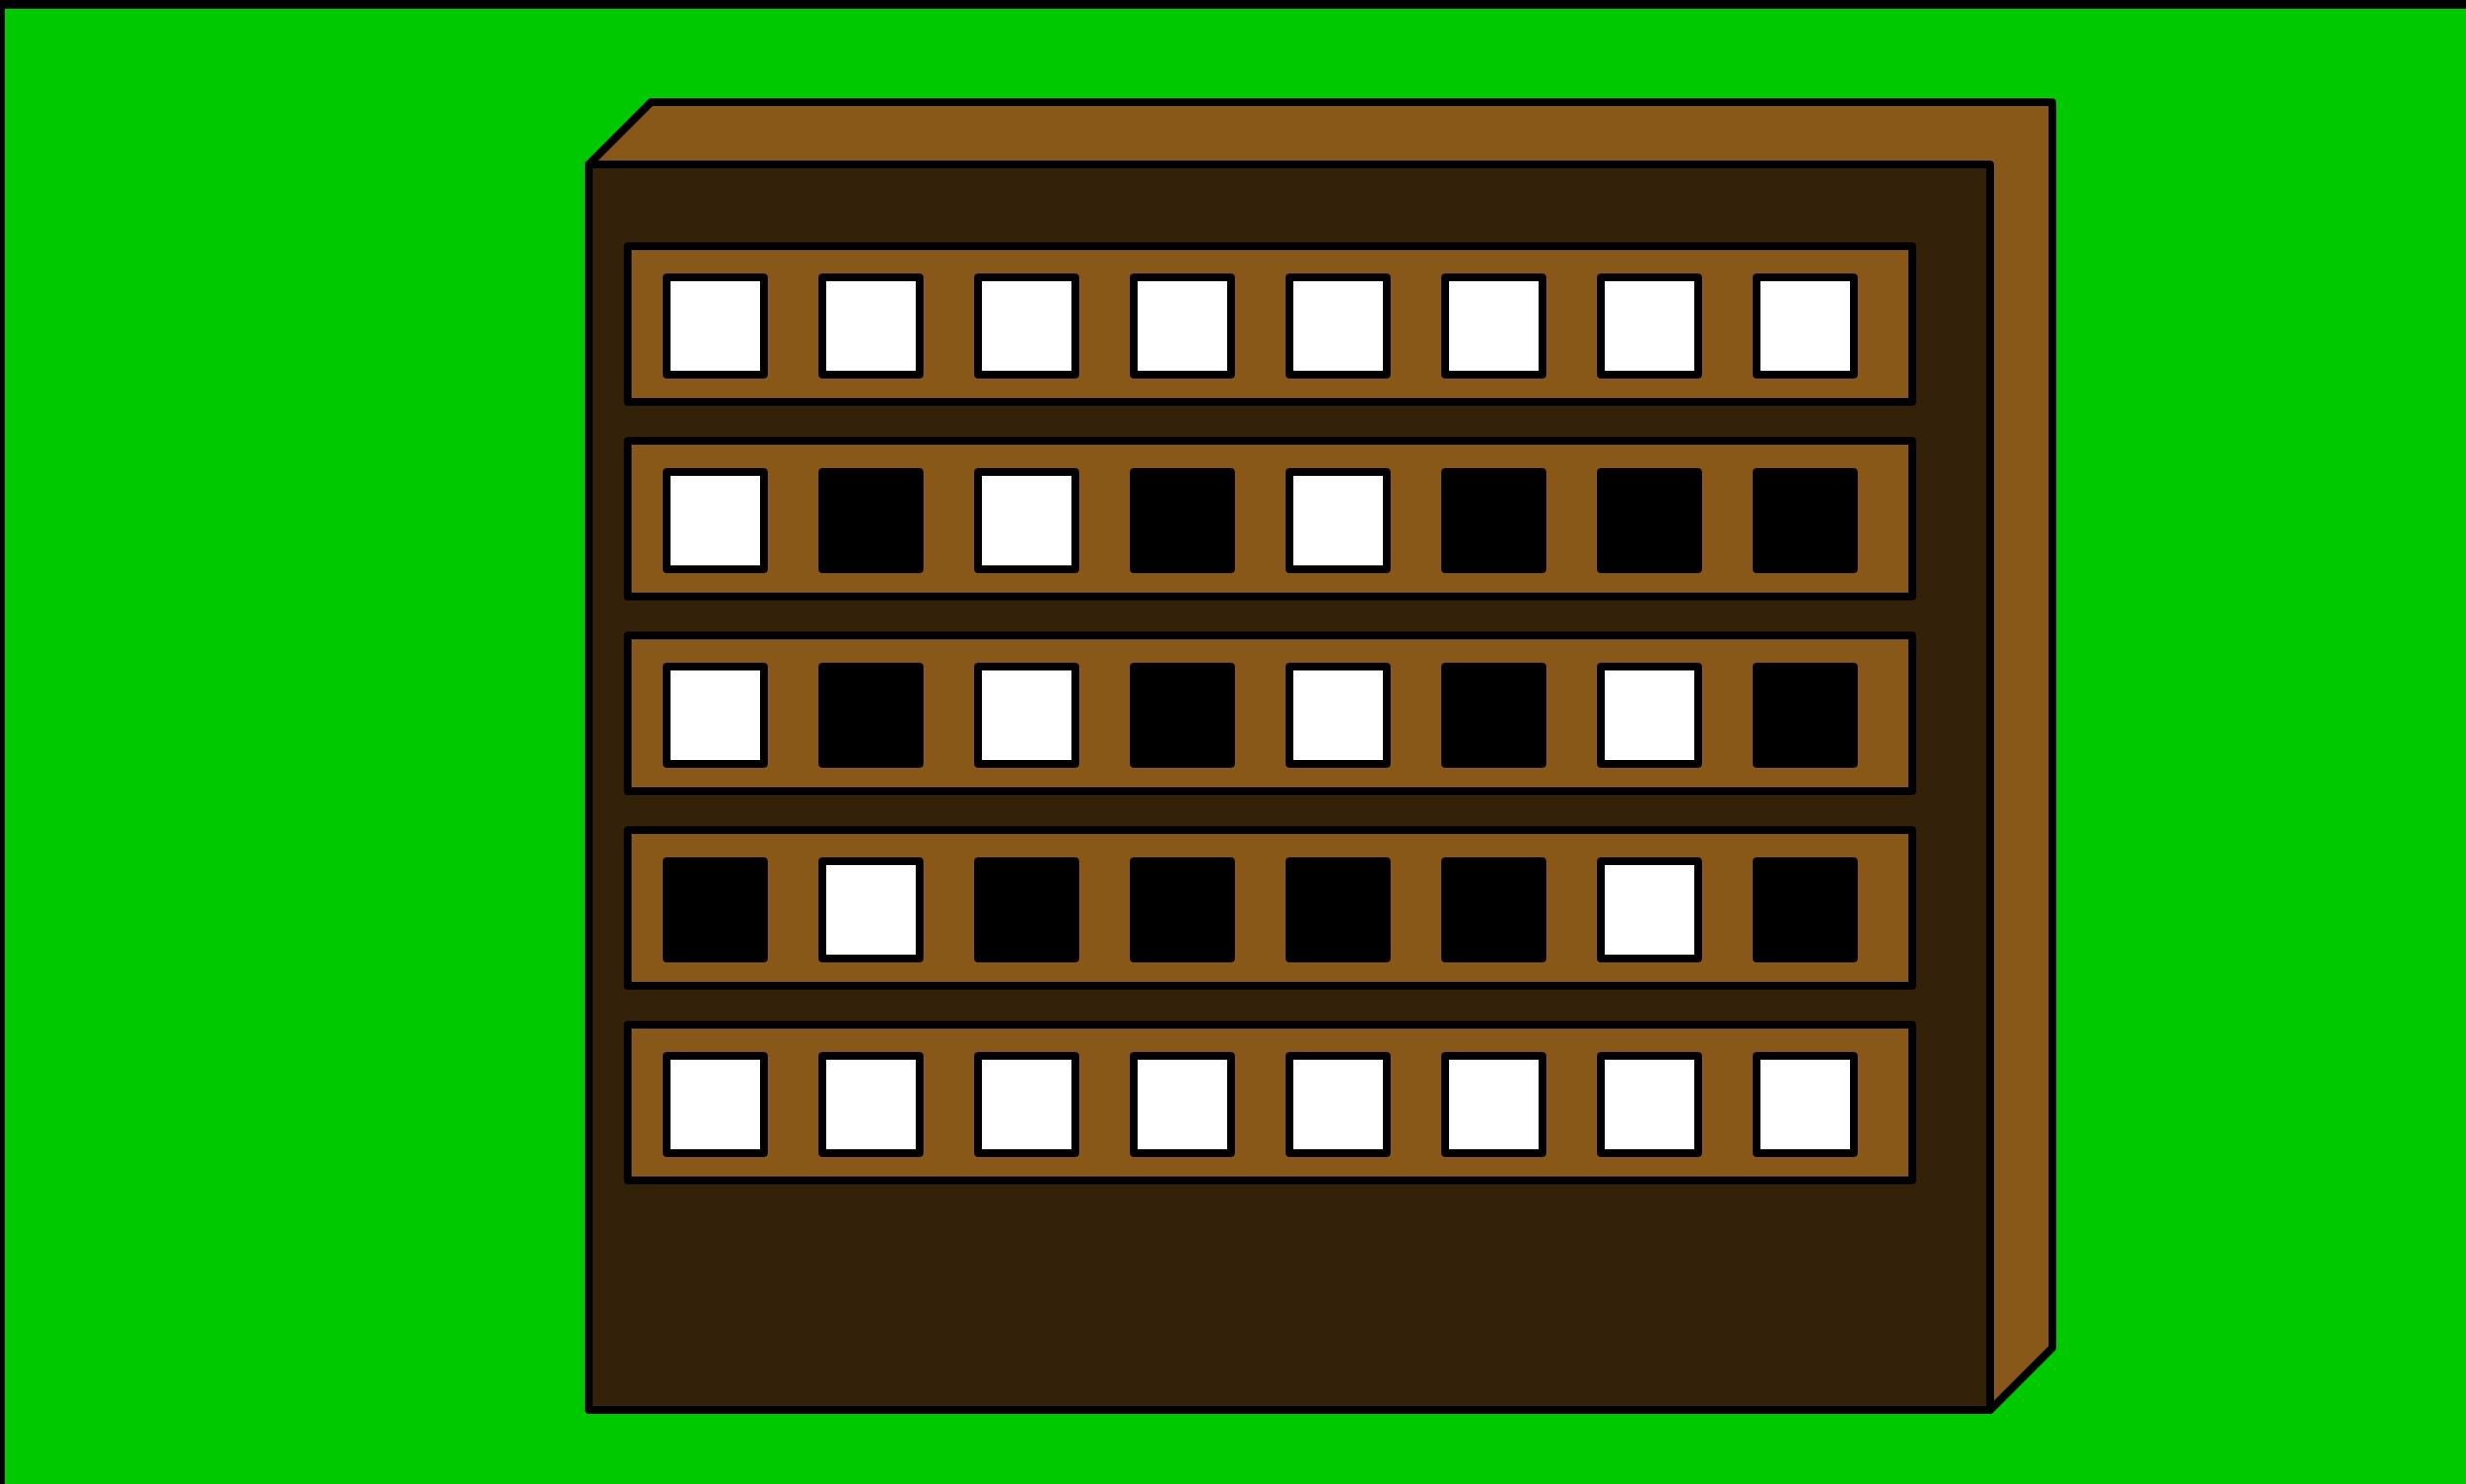


DDR memory

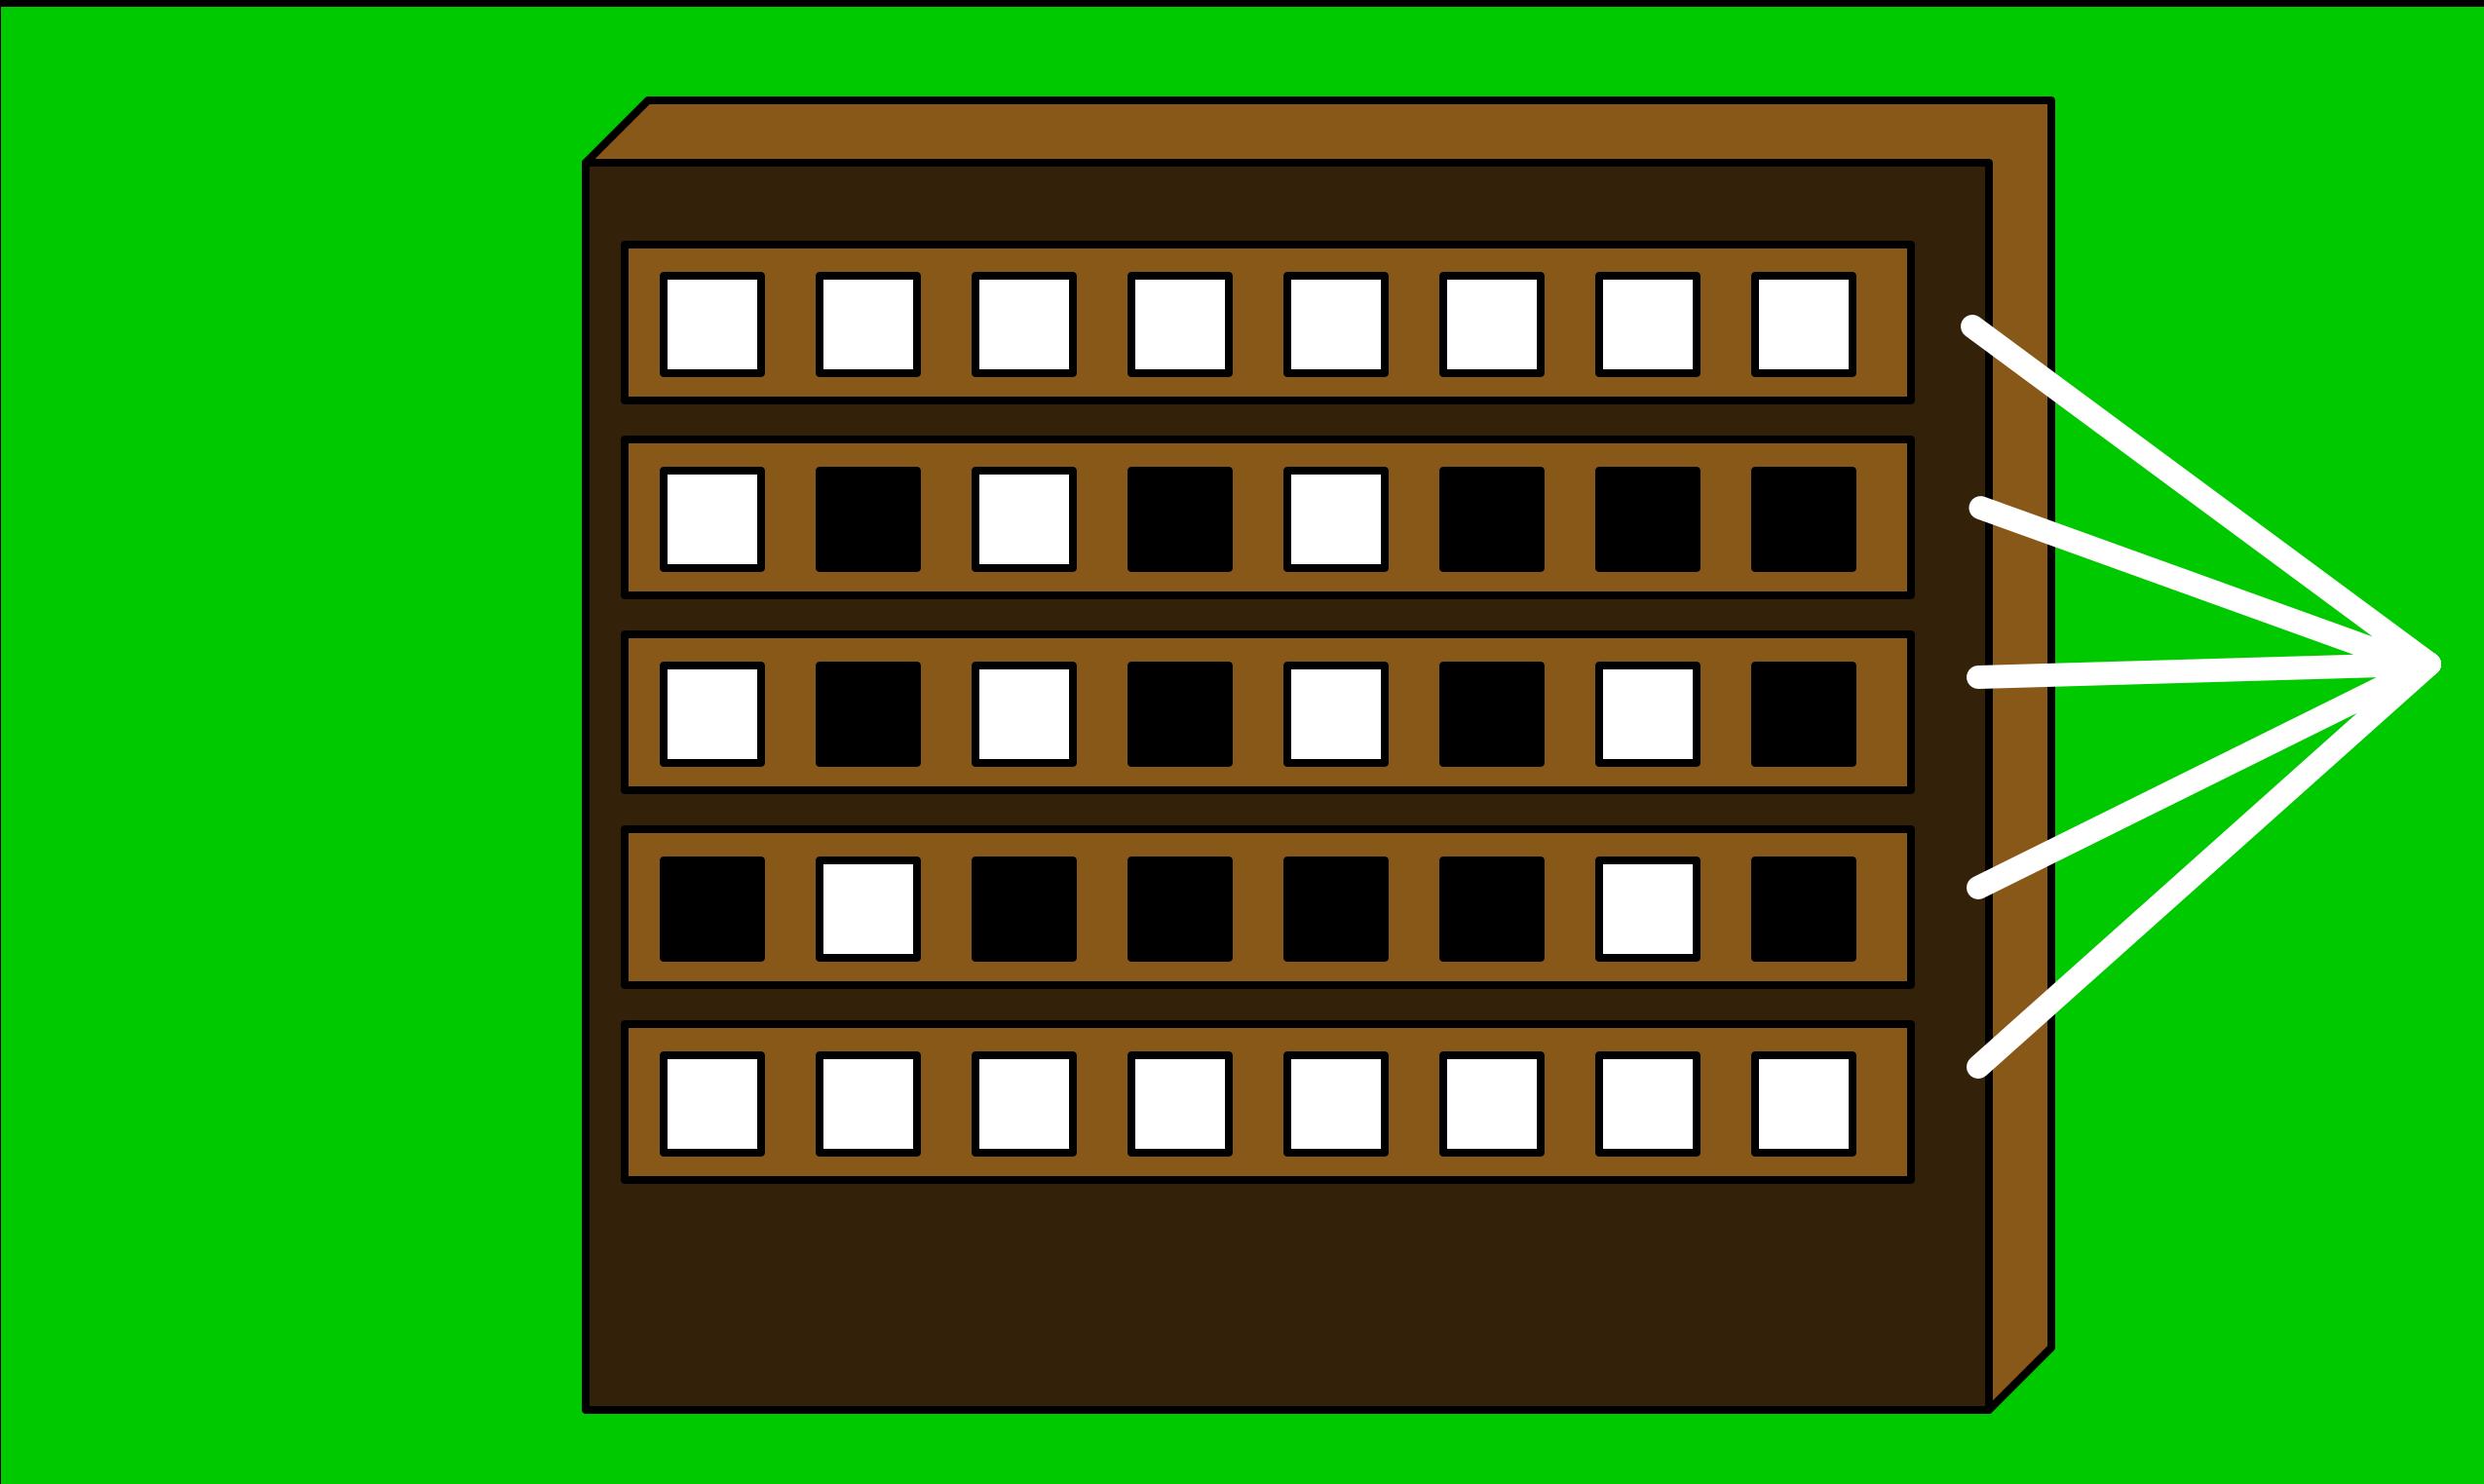
# rowhammer attack



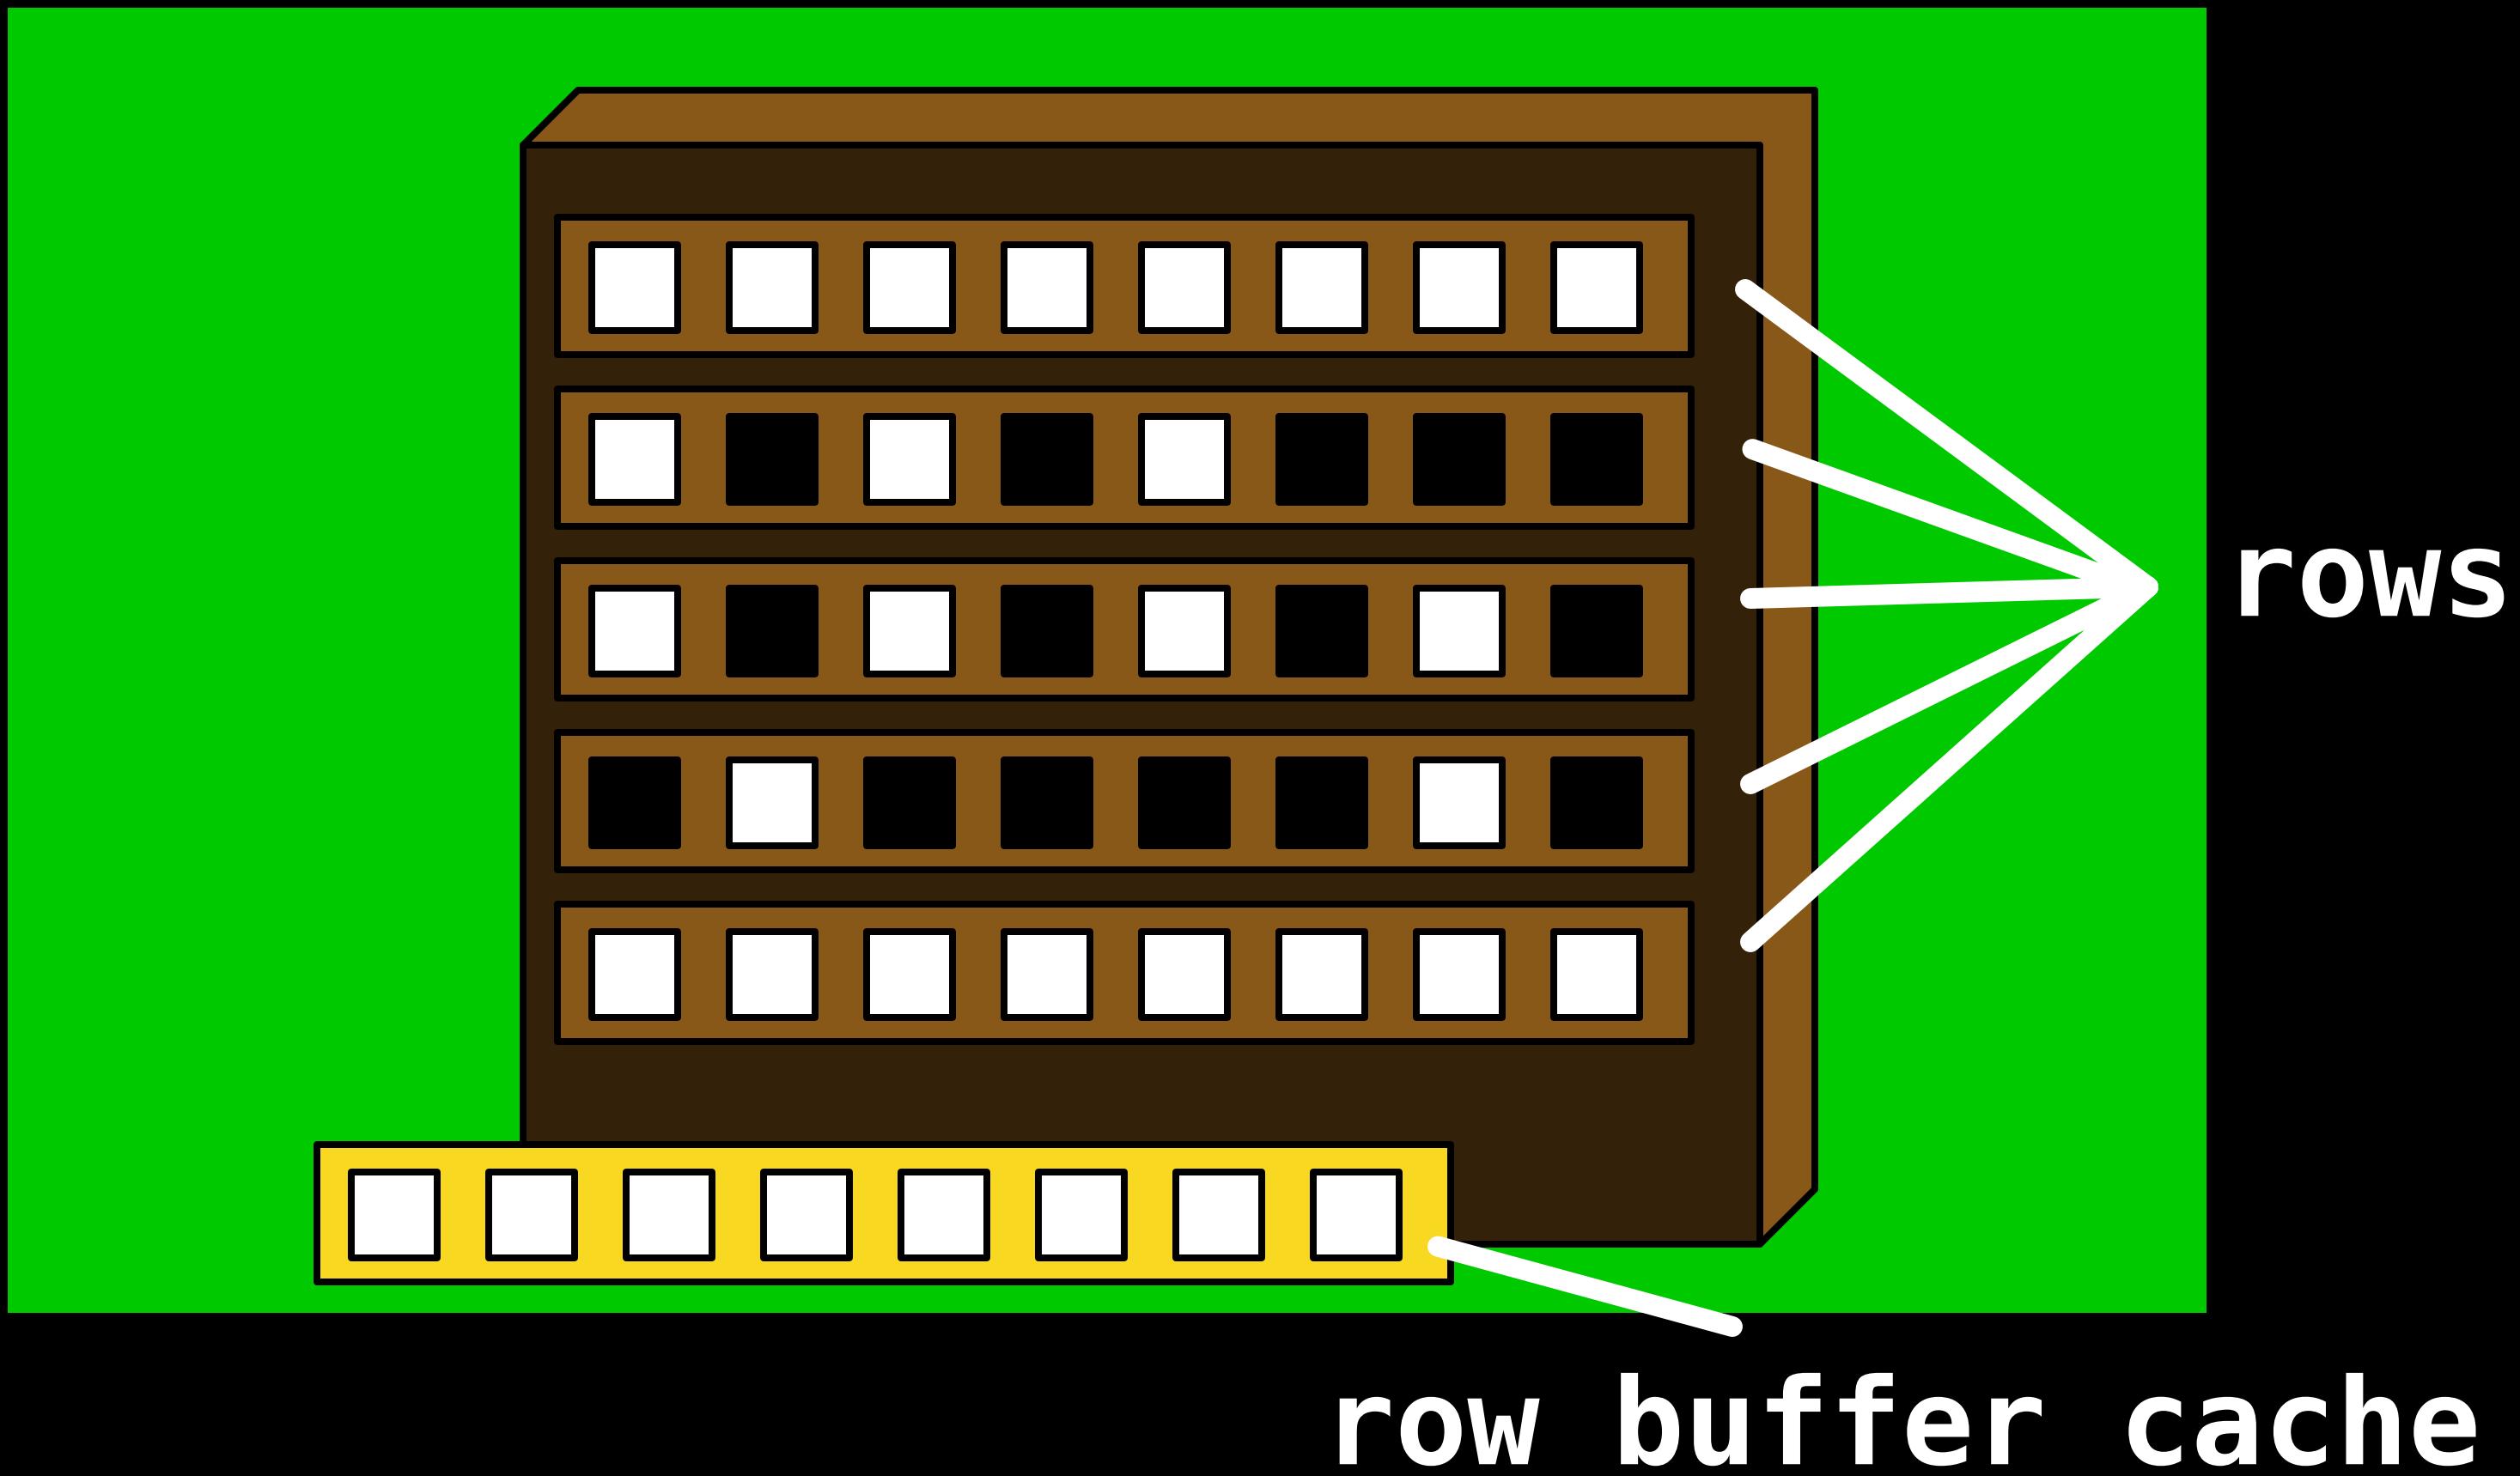
# rowhammer attack



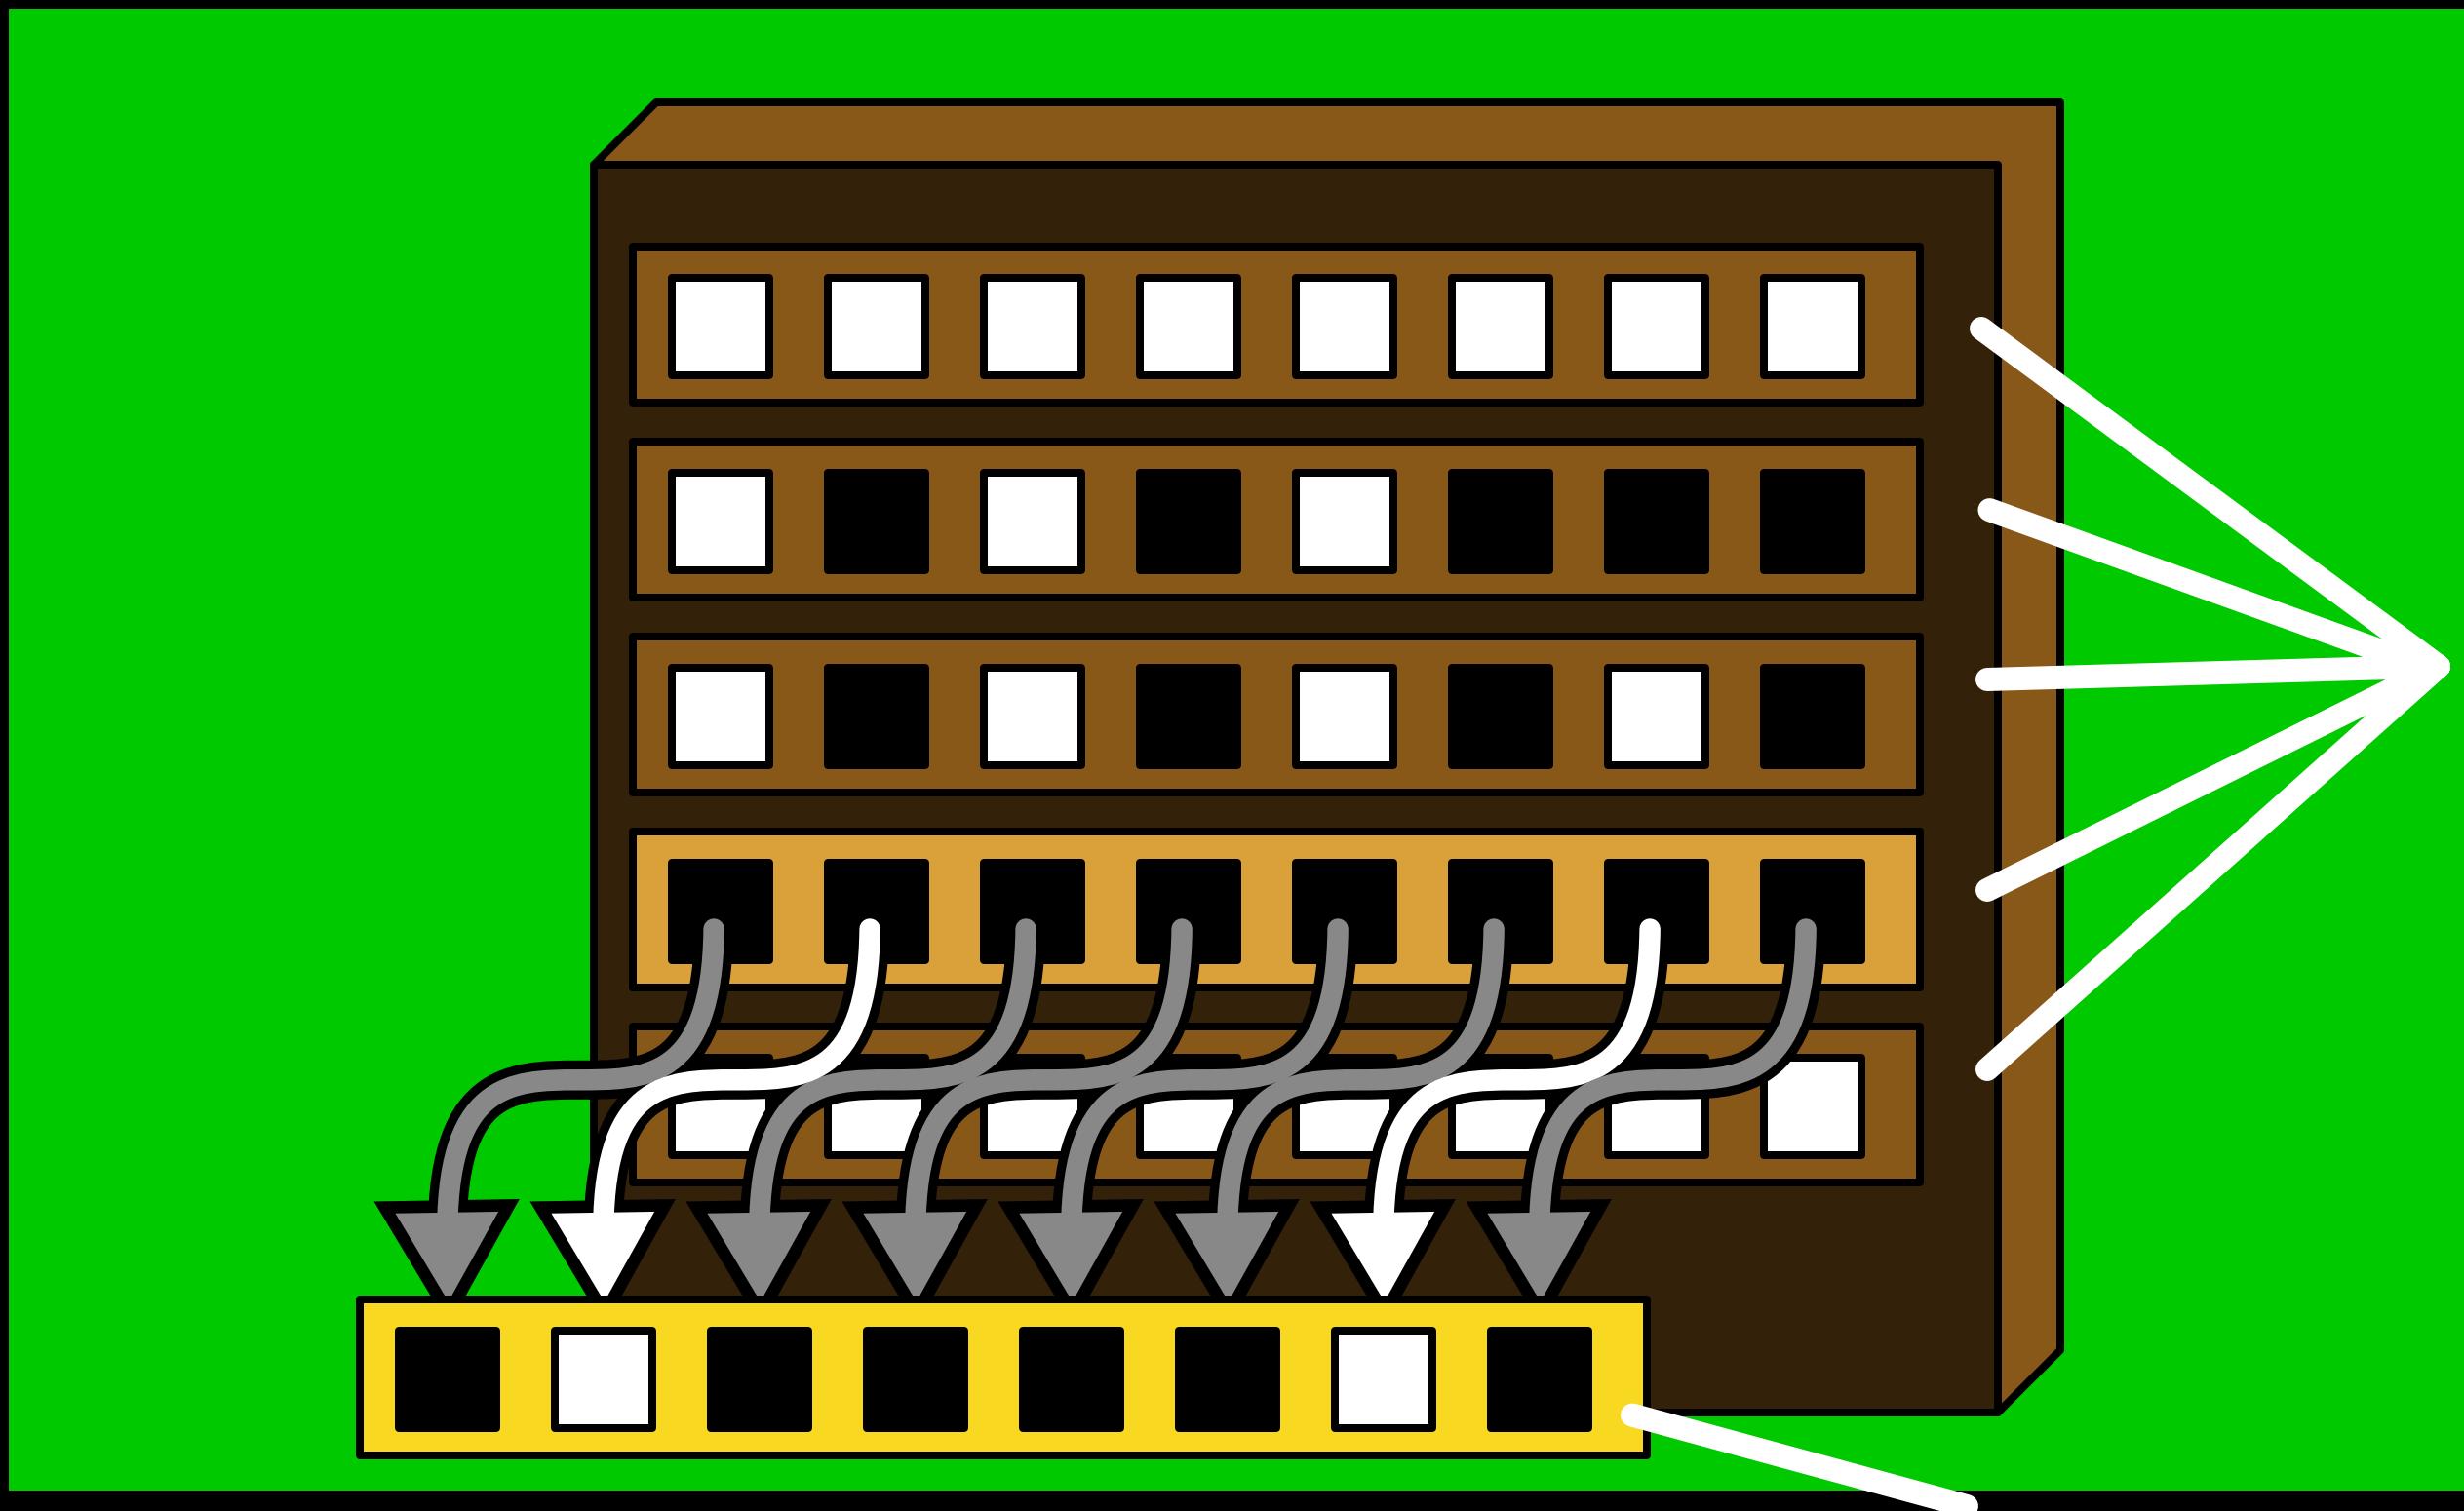
# rowhammer attack



# rowhammer attack

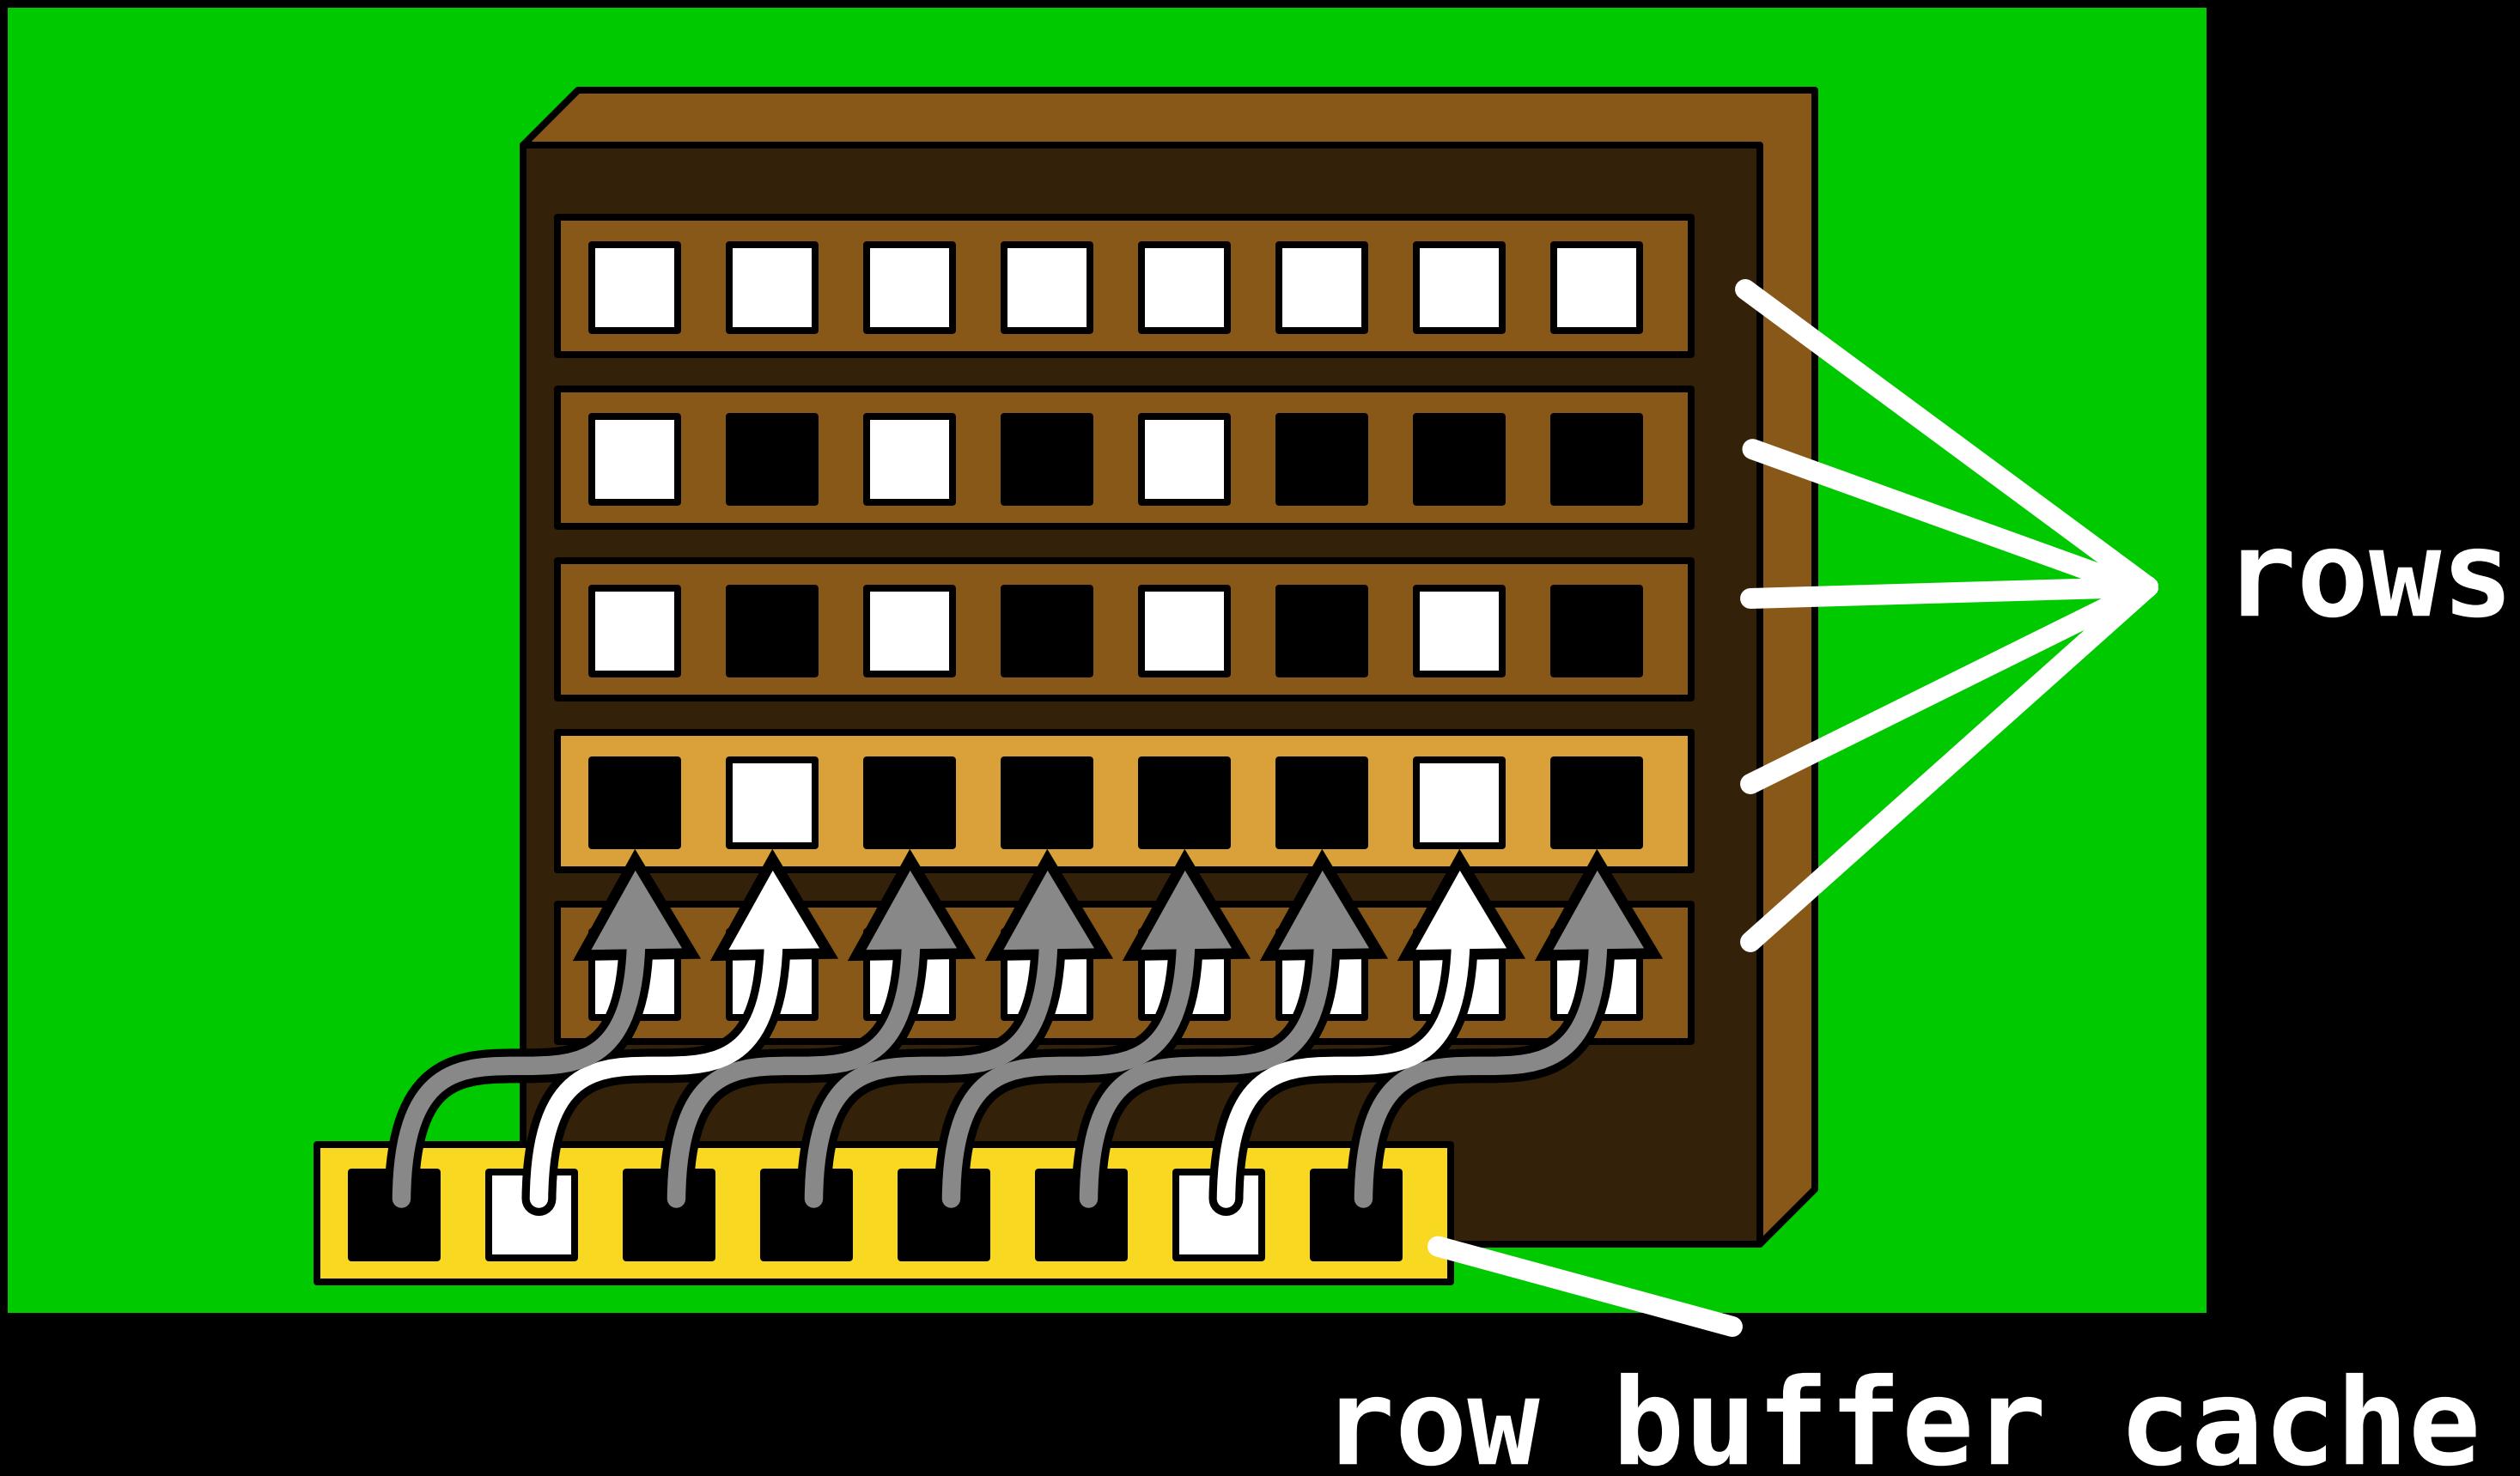


# rowhammer attack

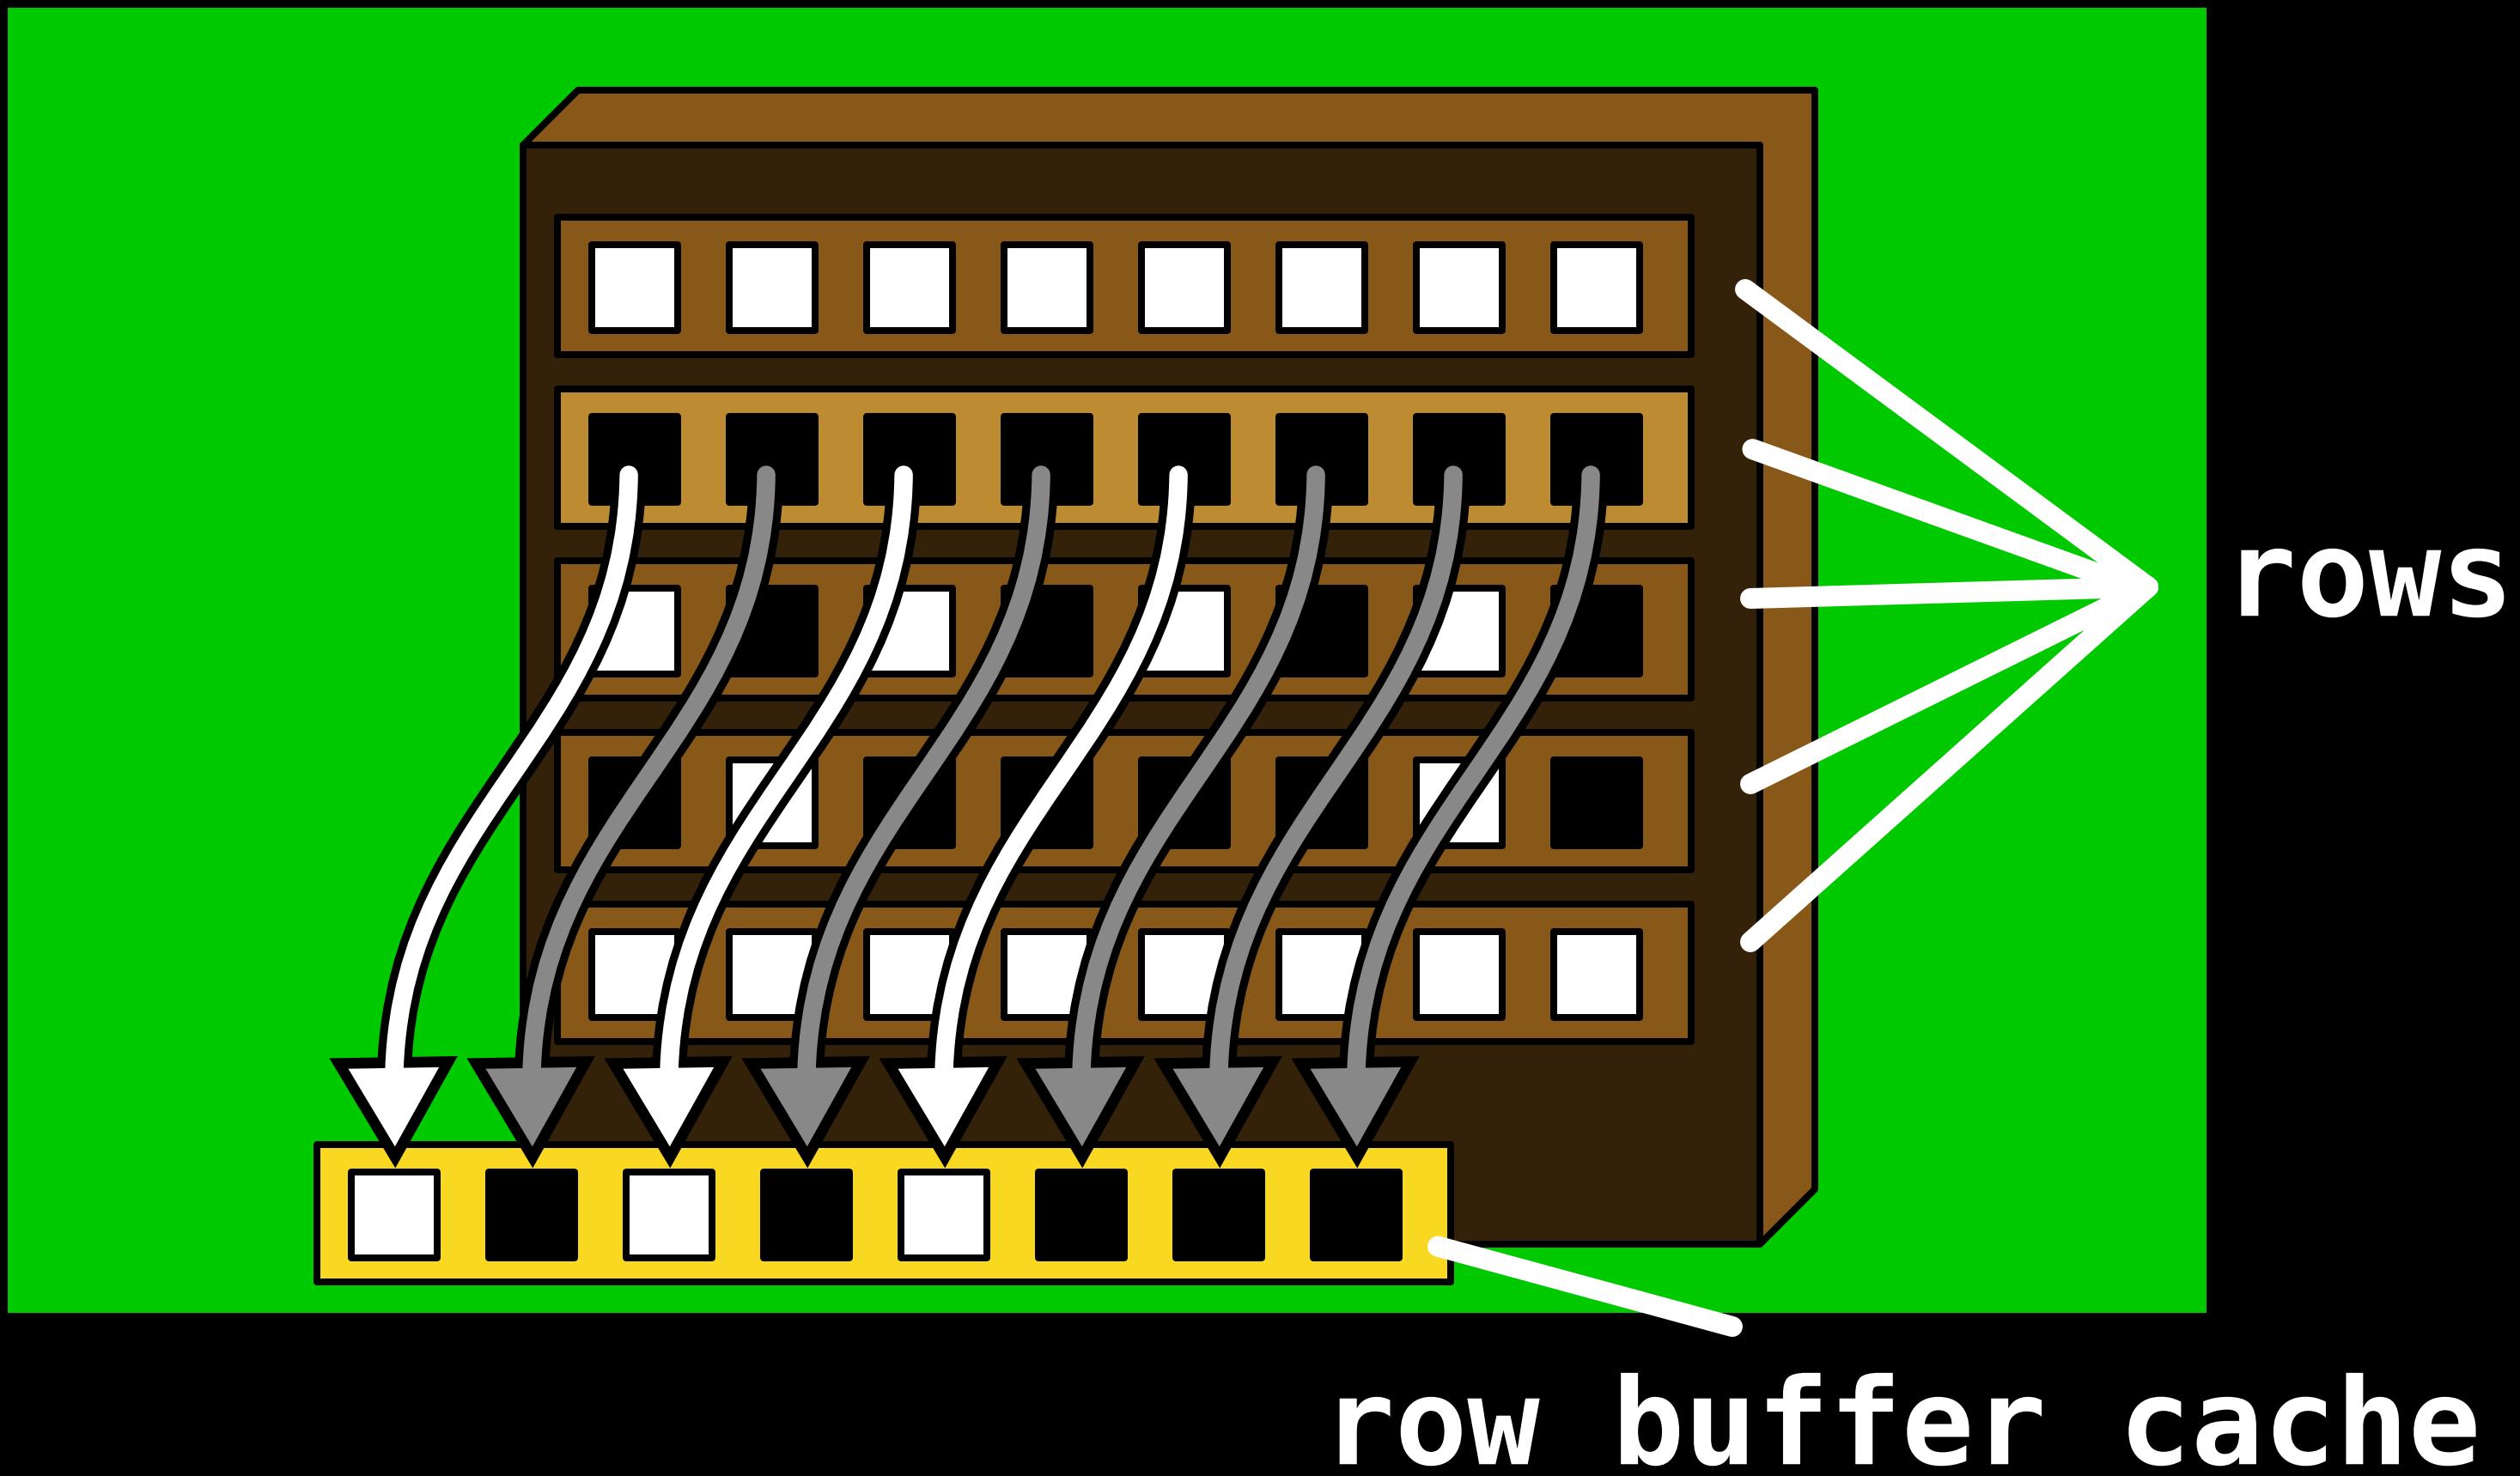


row buffer cache

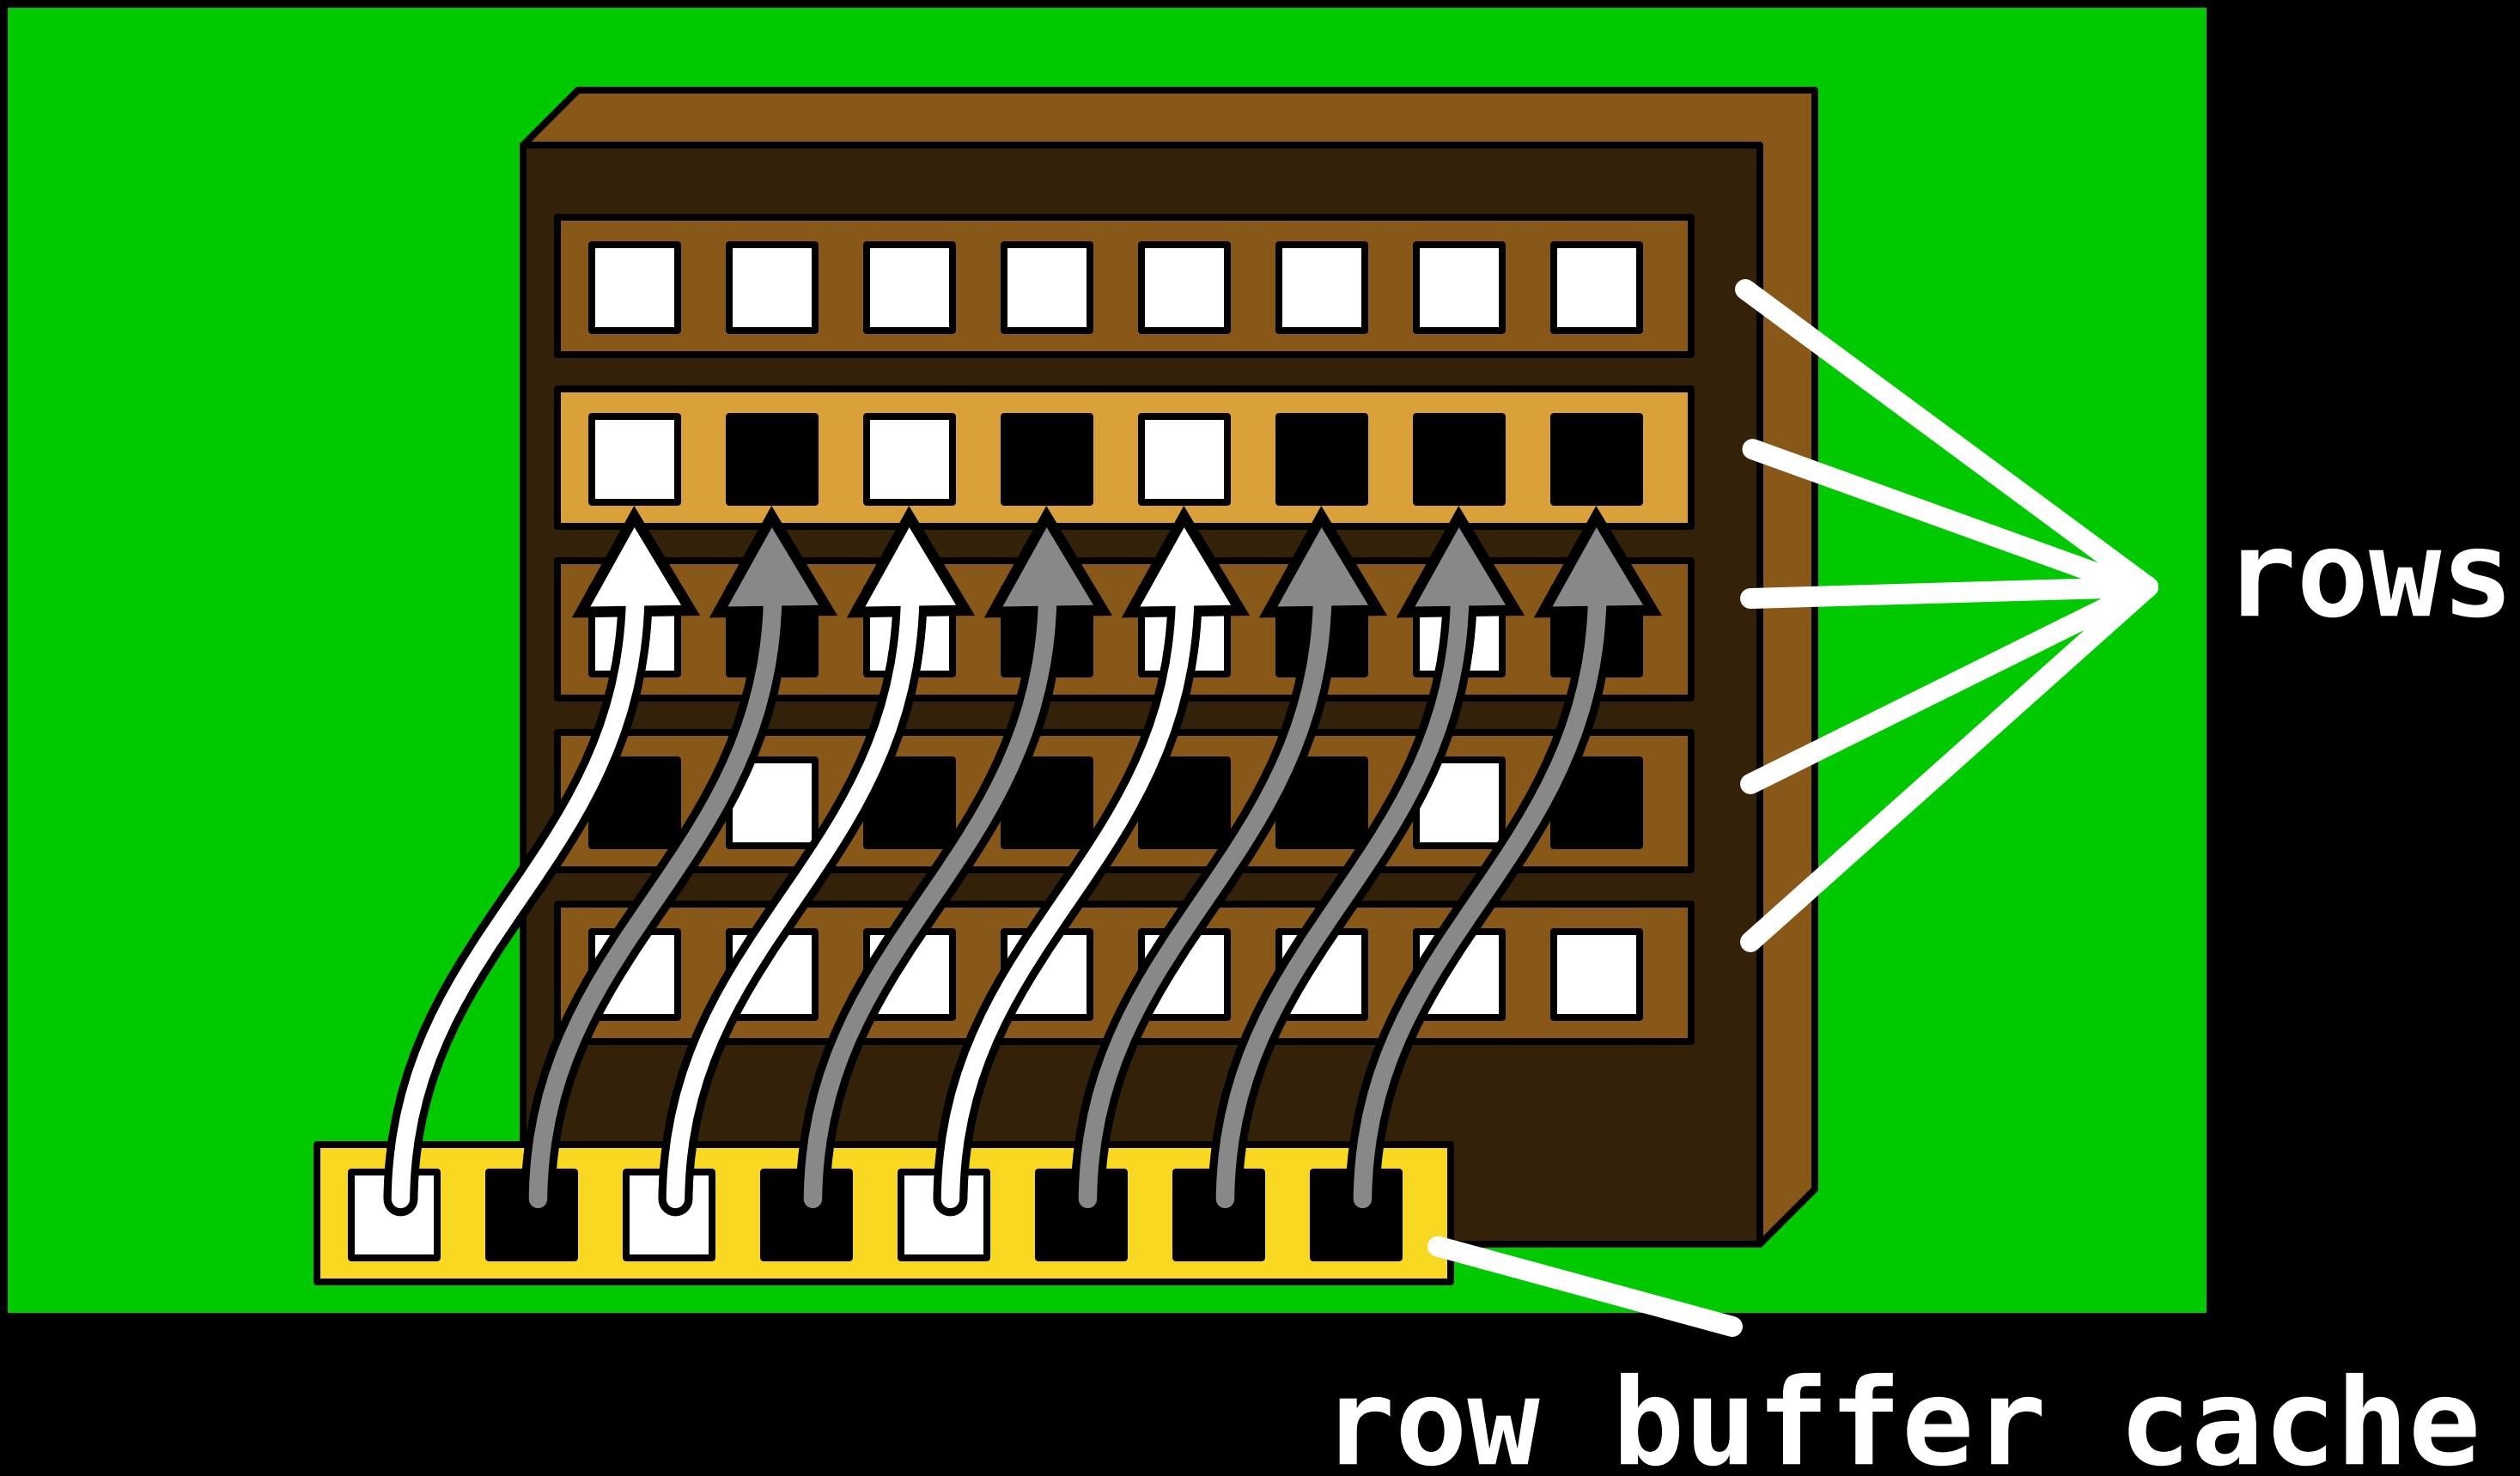
# rowhammer attack



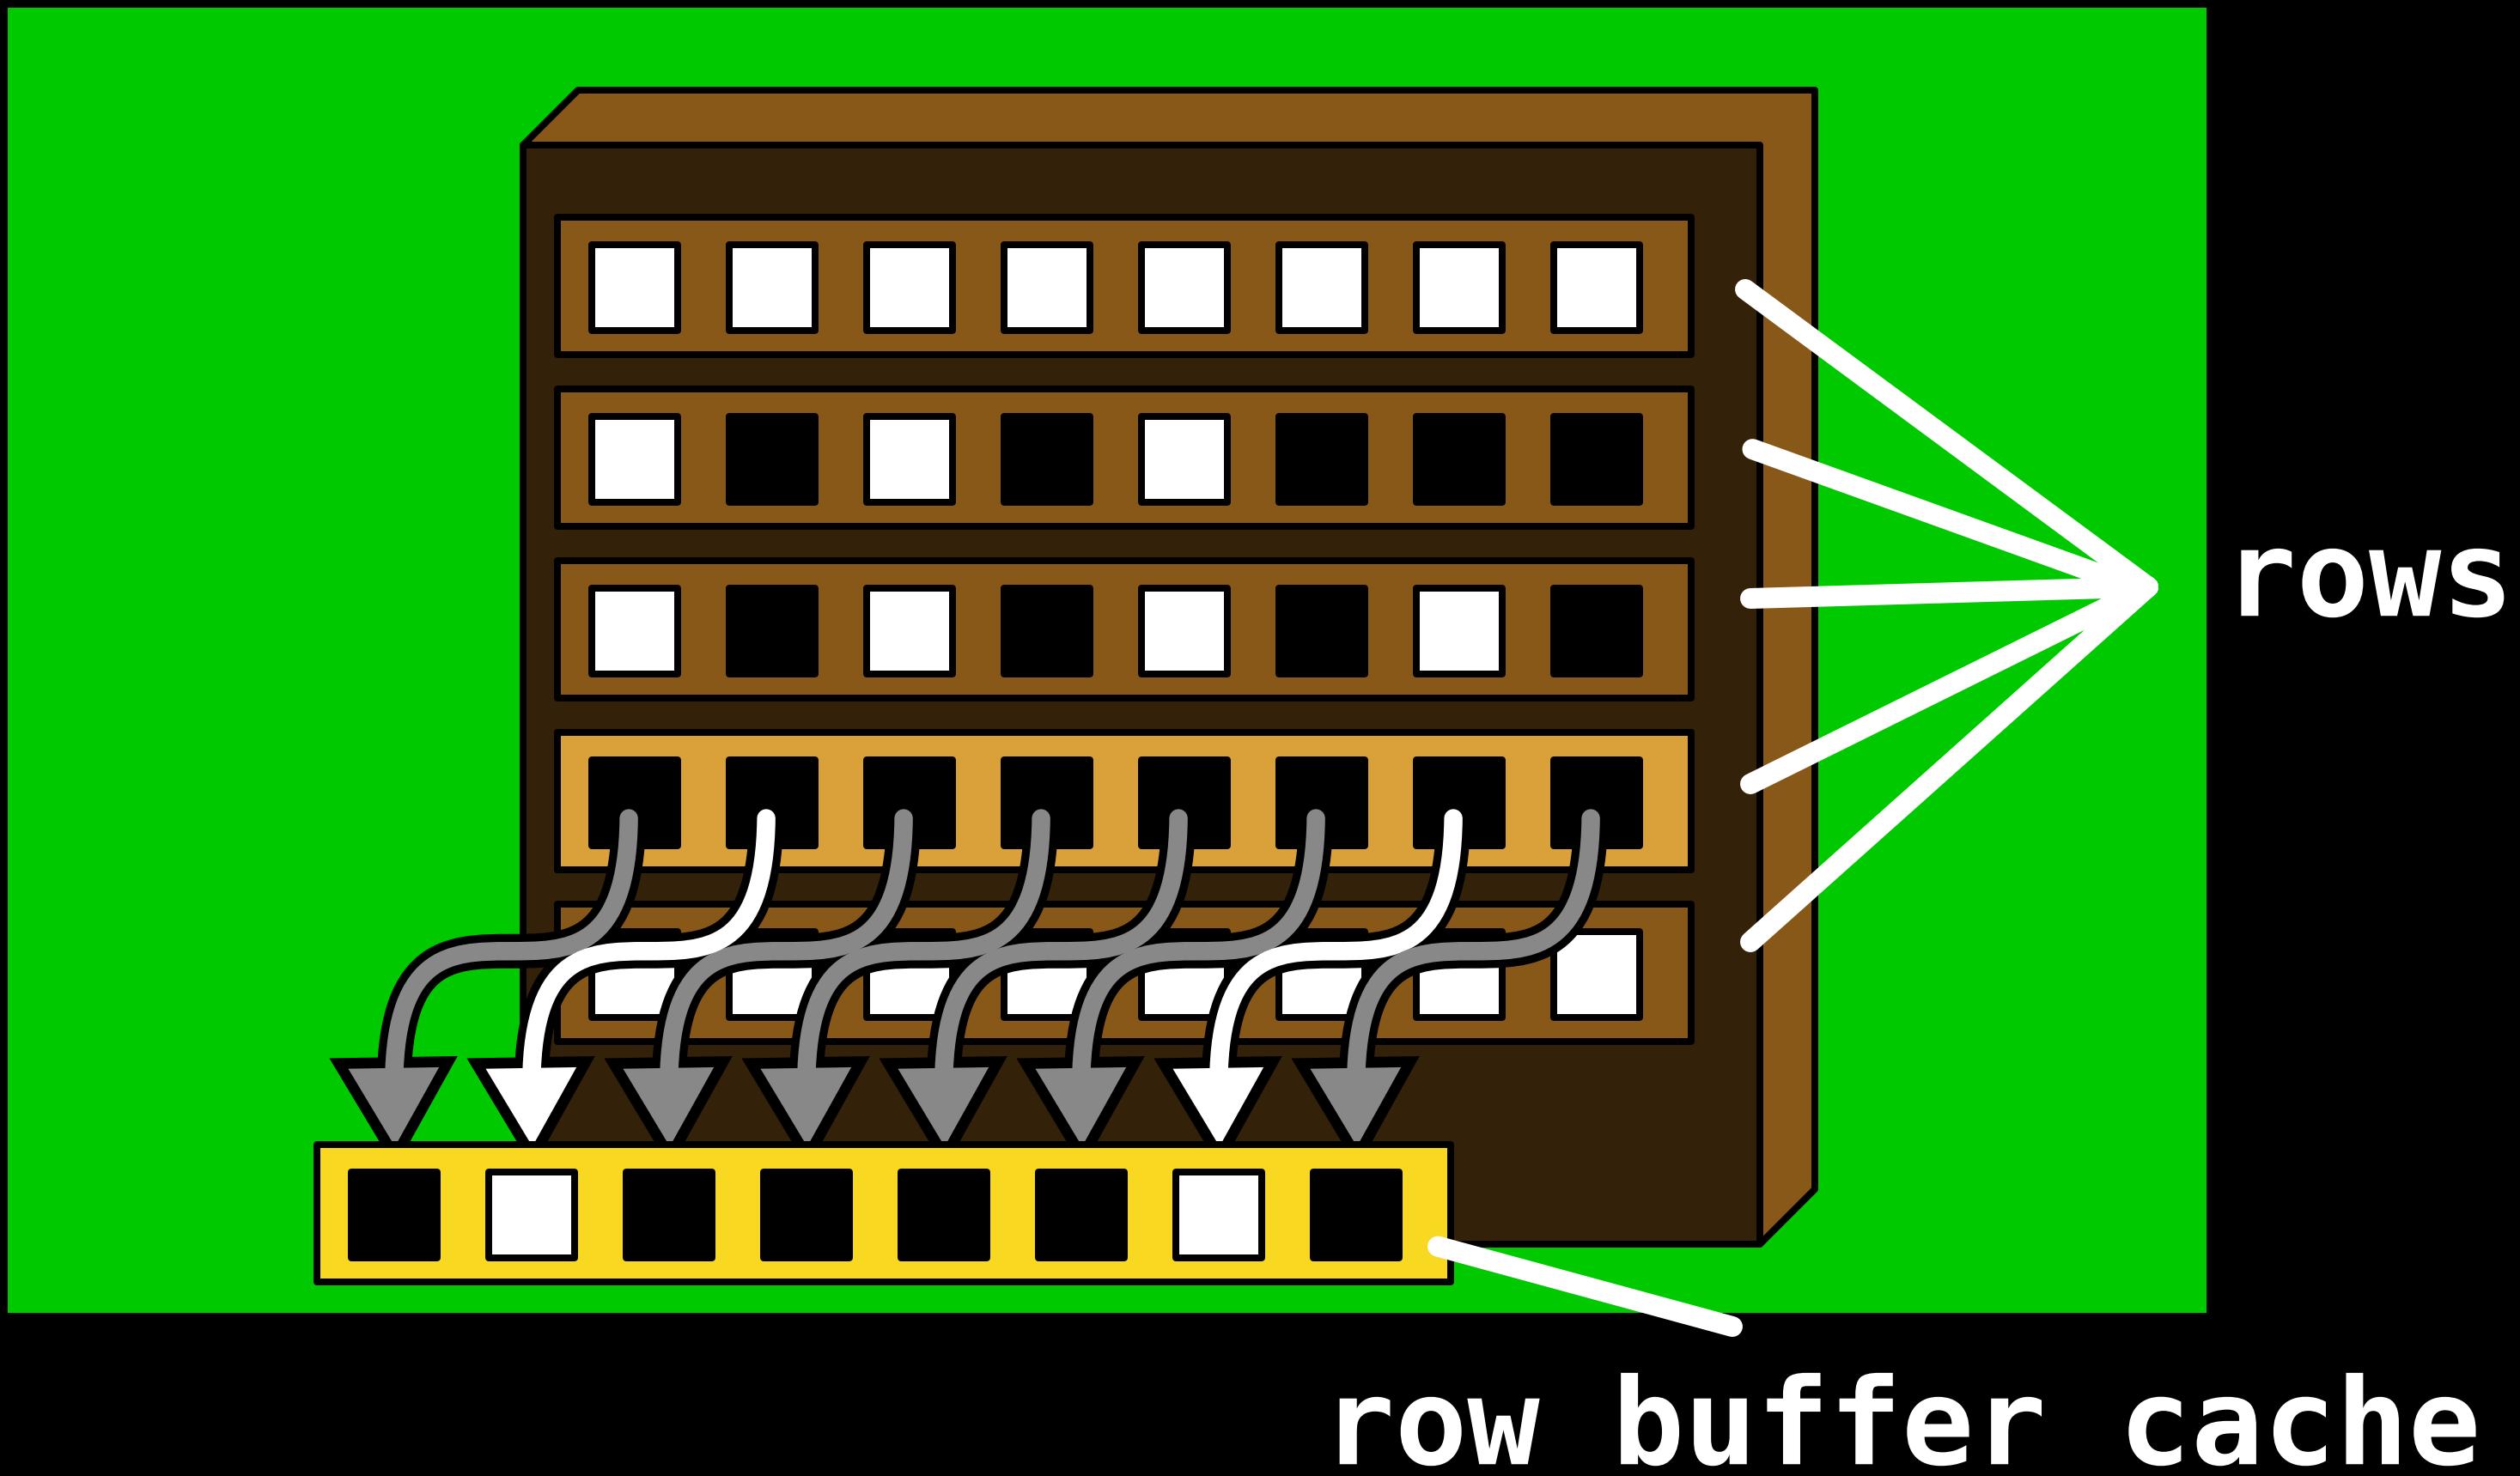
# rowhammer attack



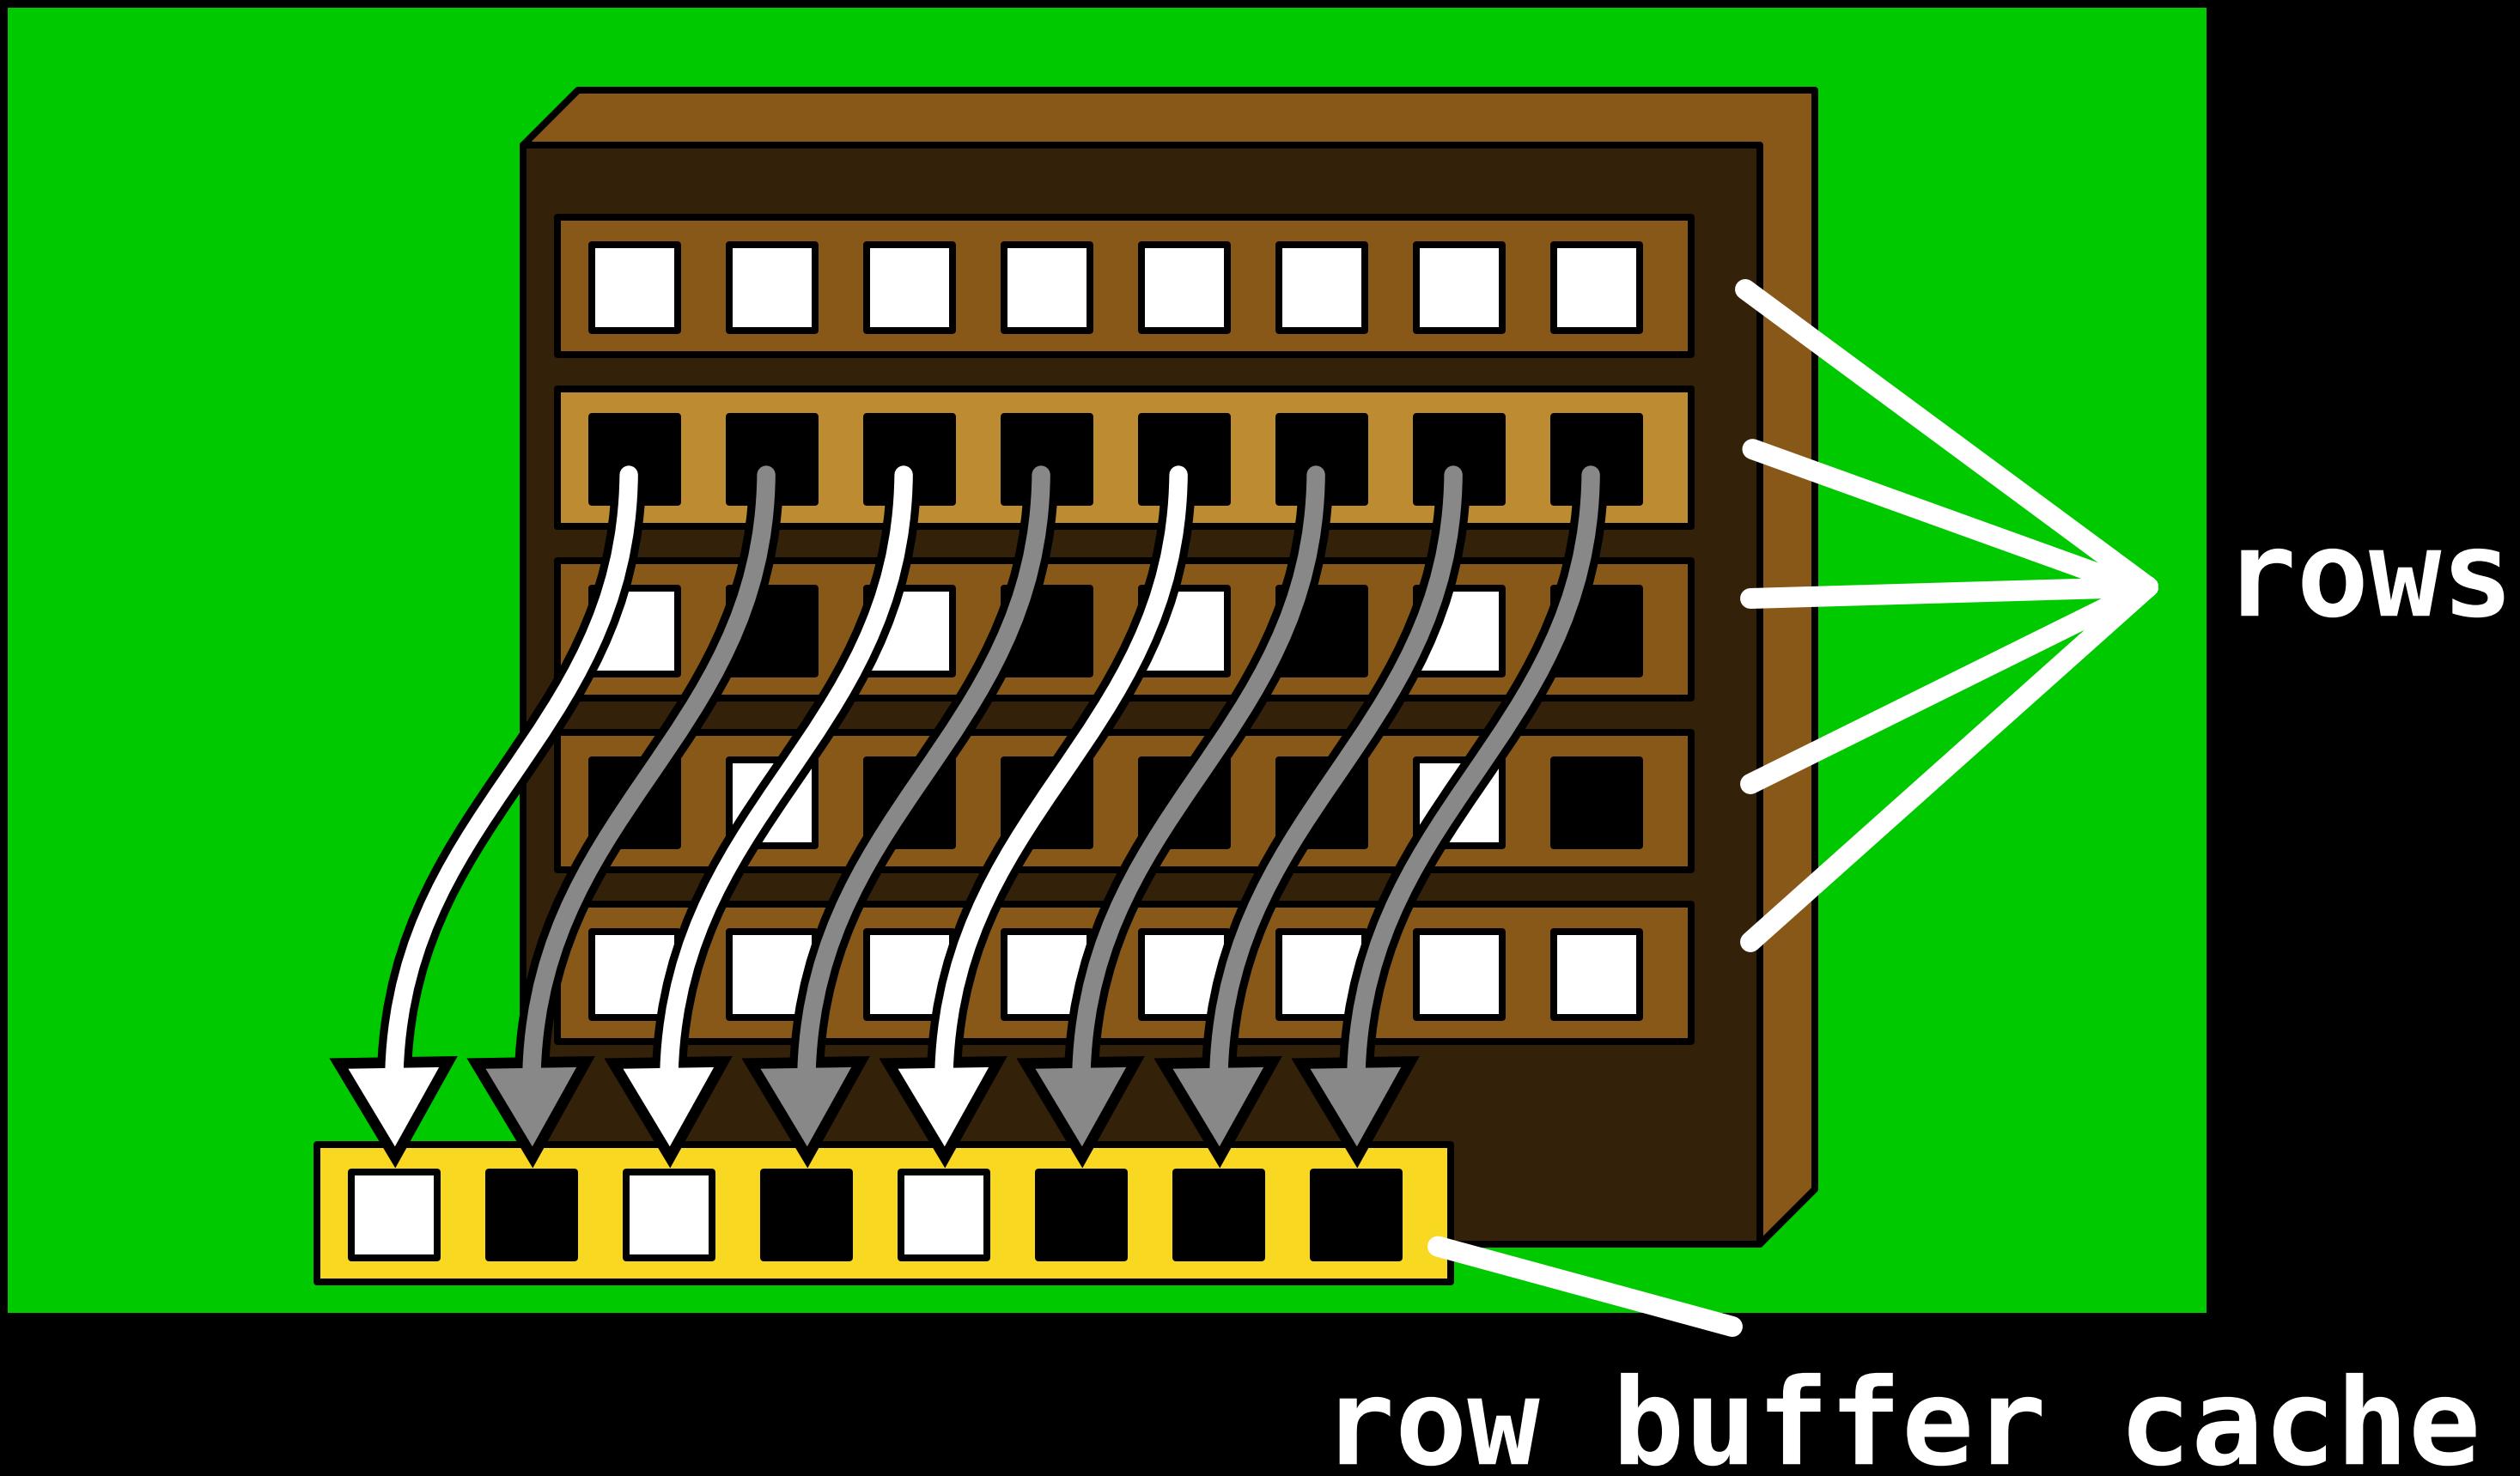
# rowhammer attack



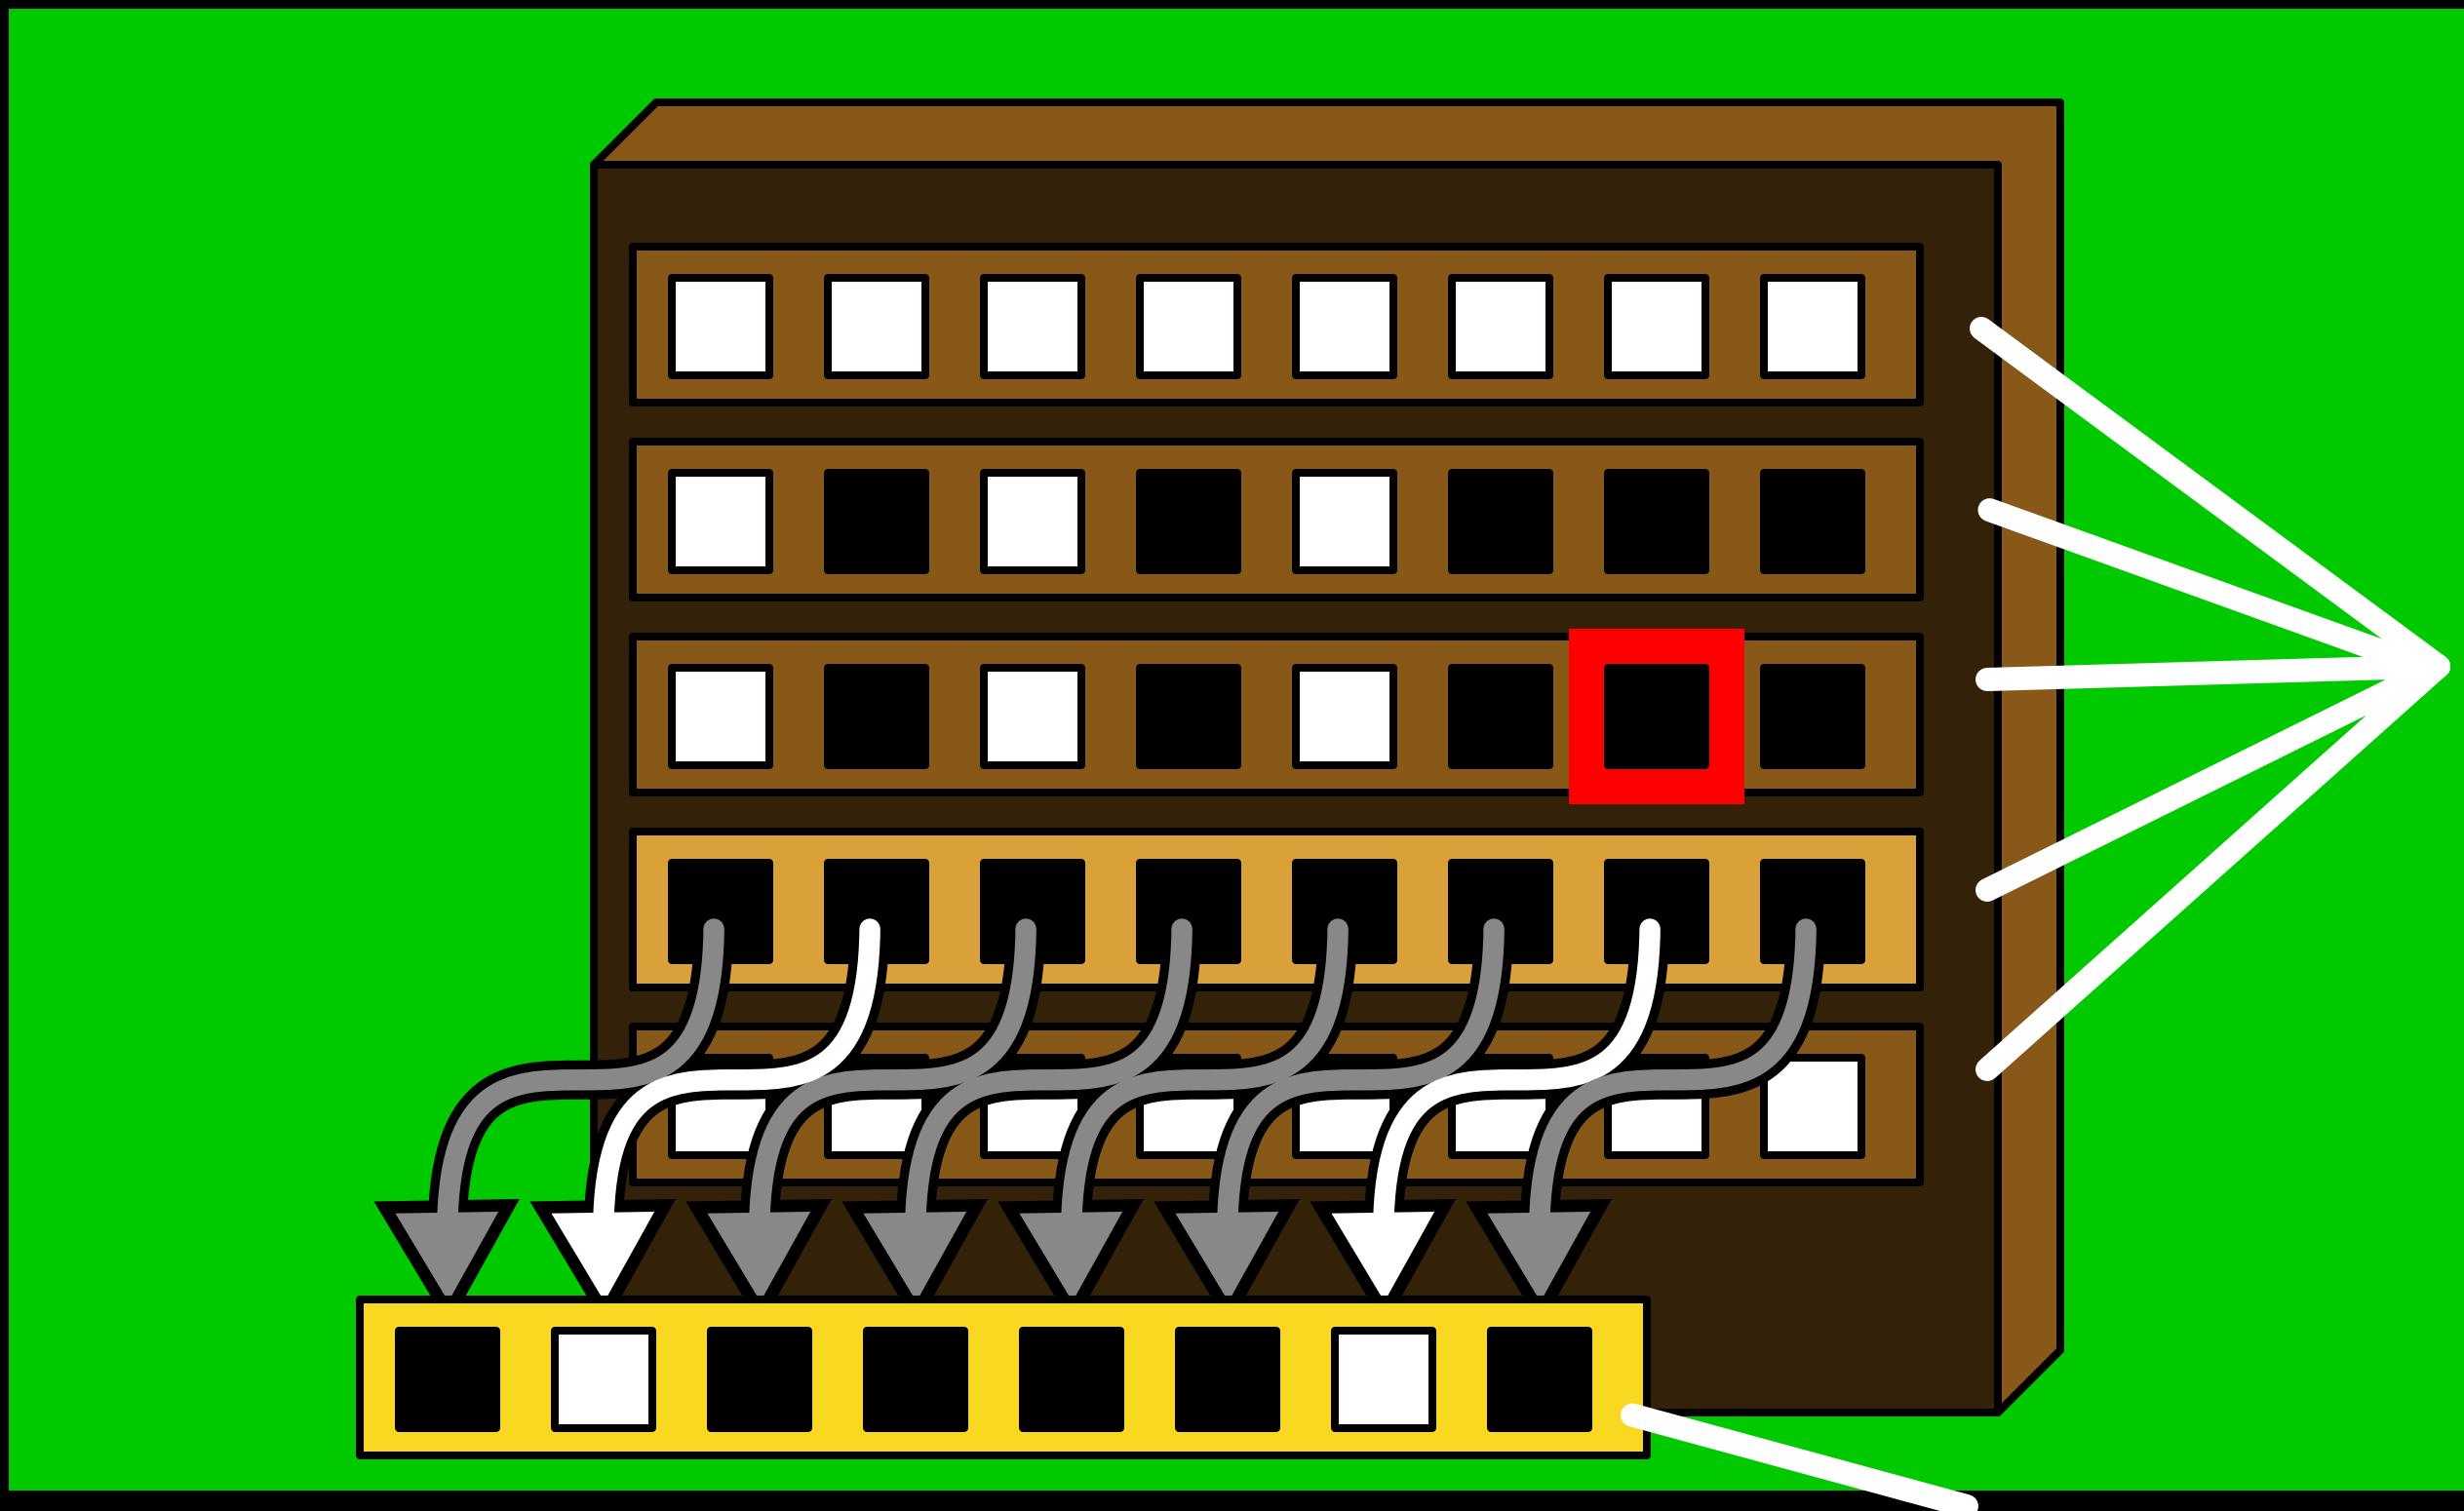
# rowhammer attack



# rowhammer attack



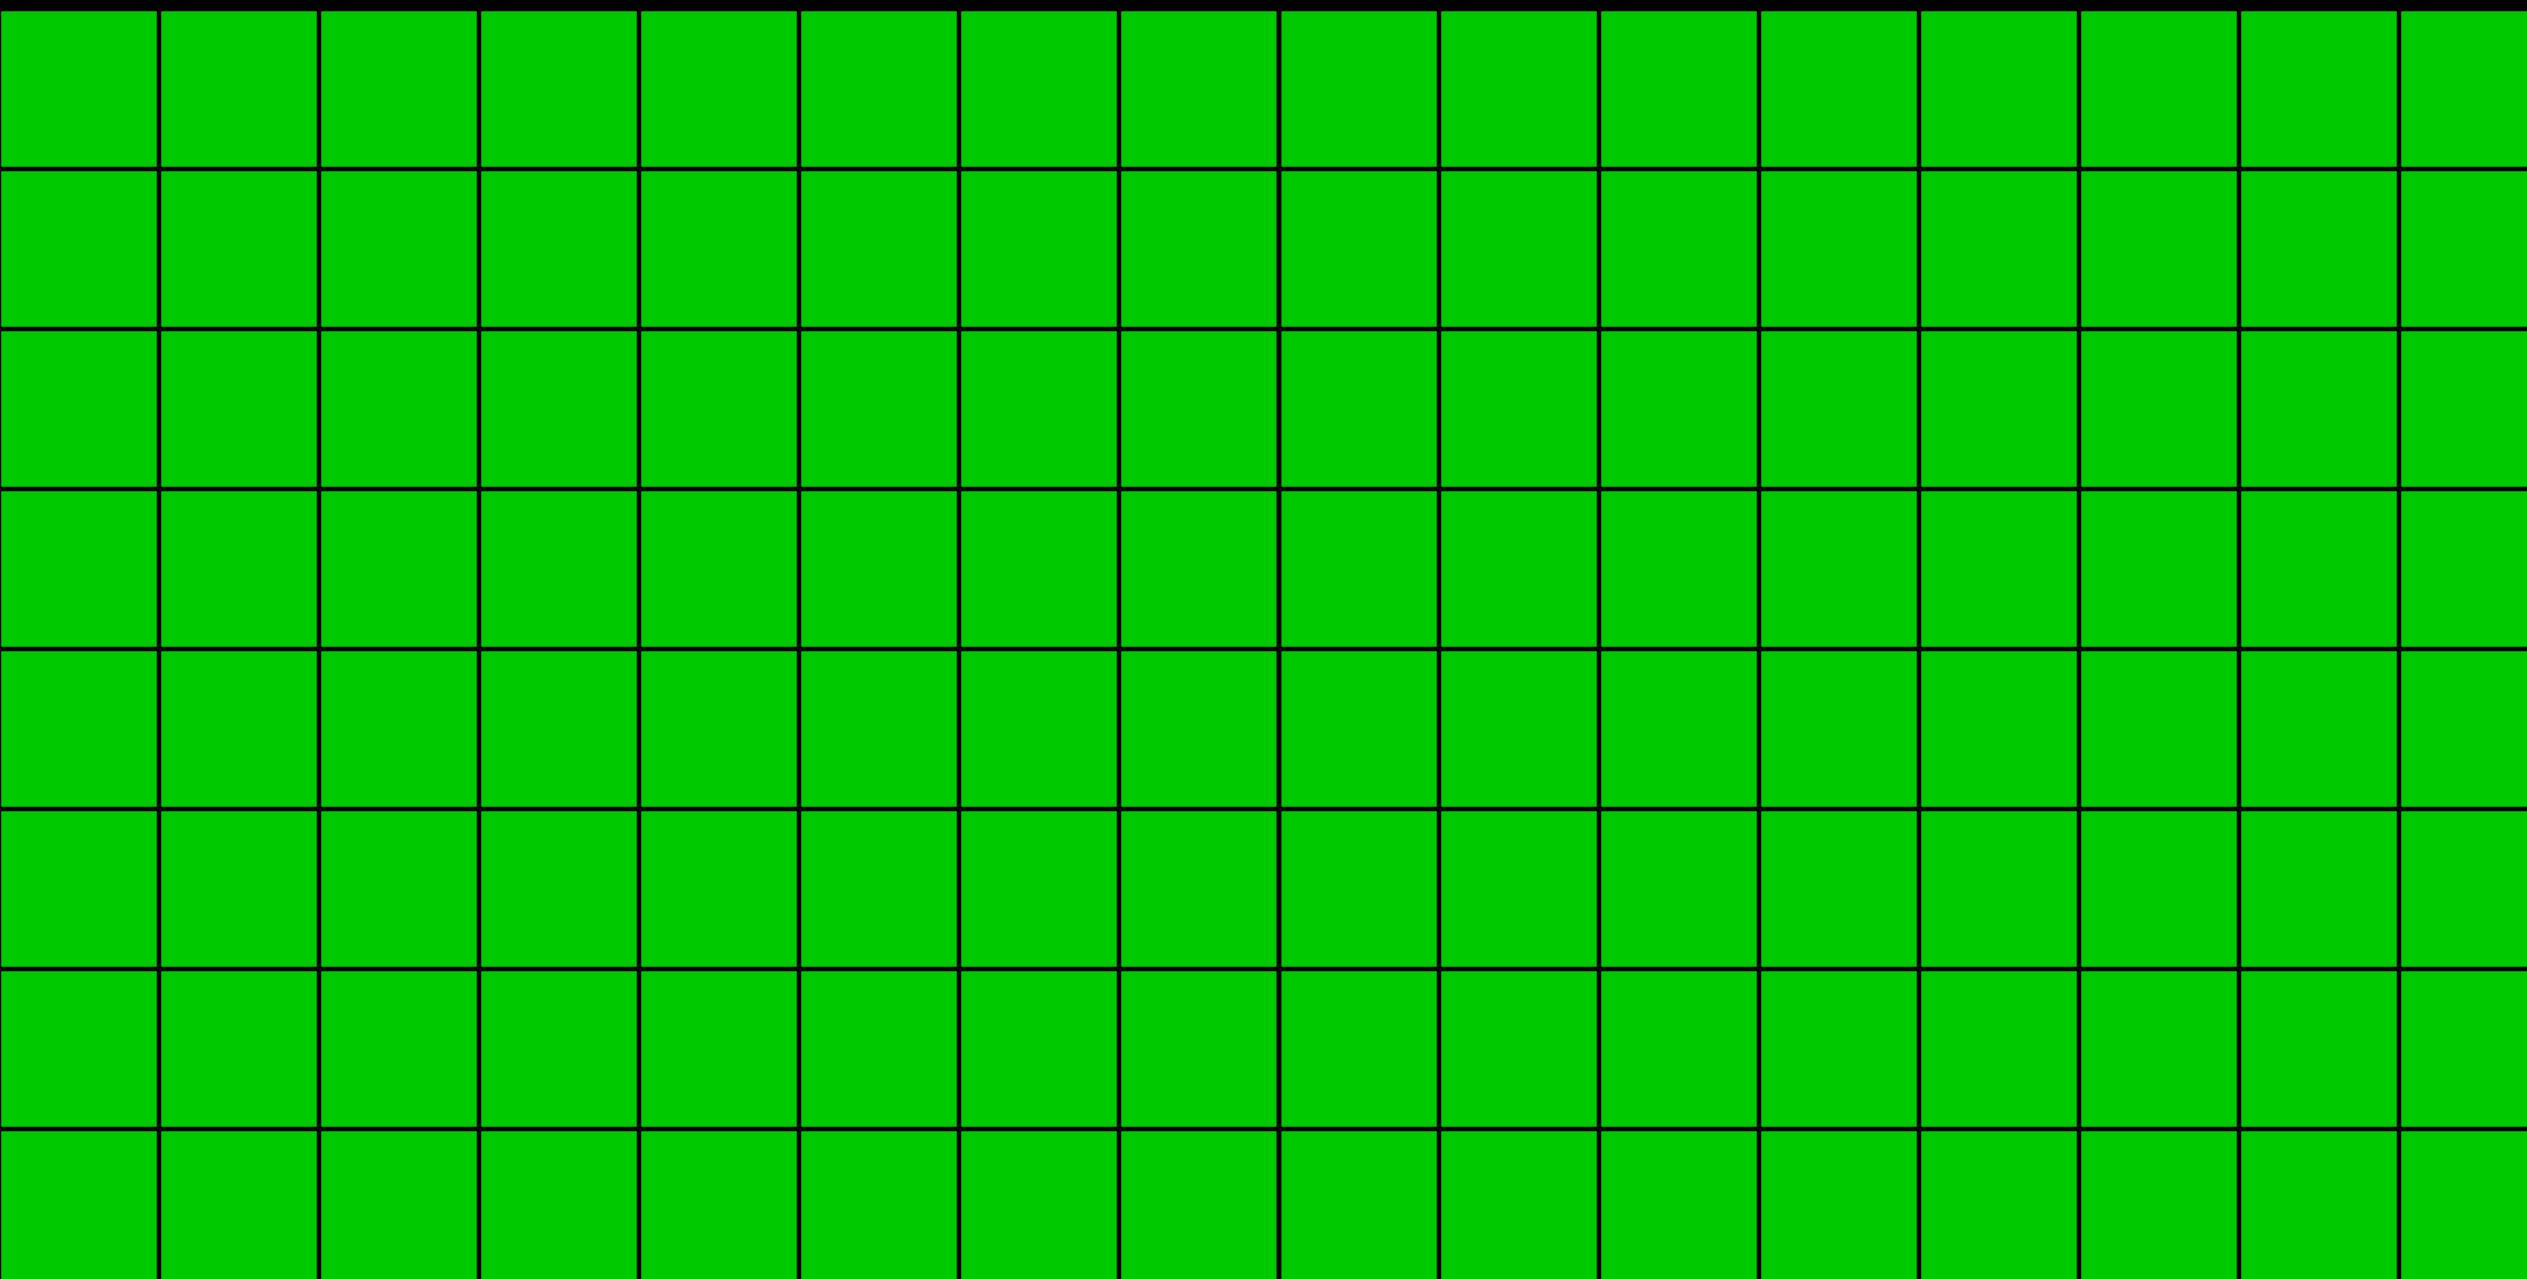
# rowhammer attack



row buffer cache

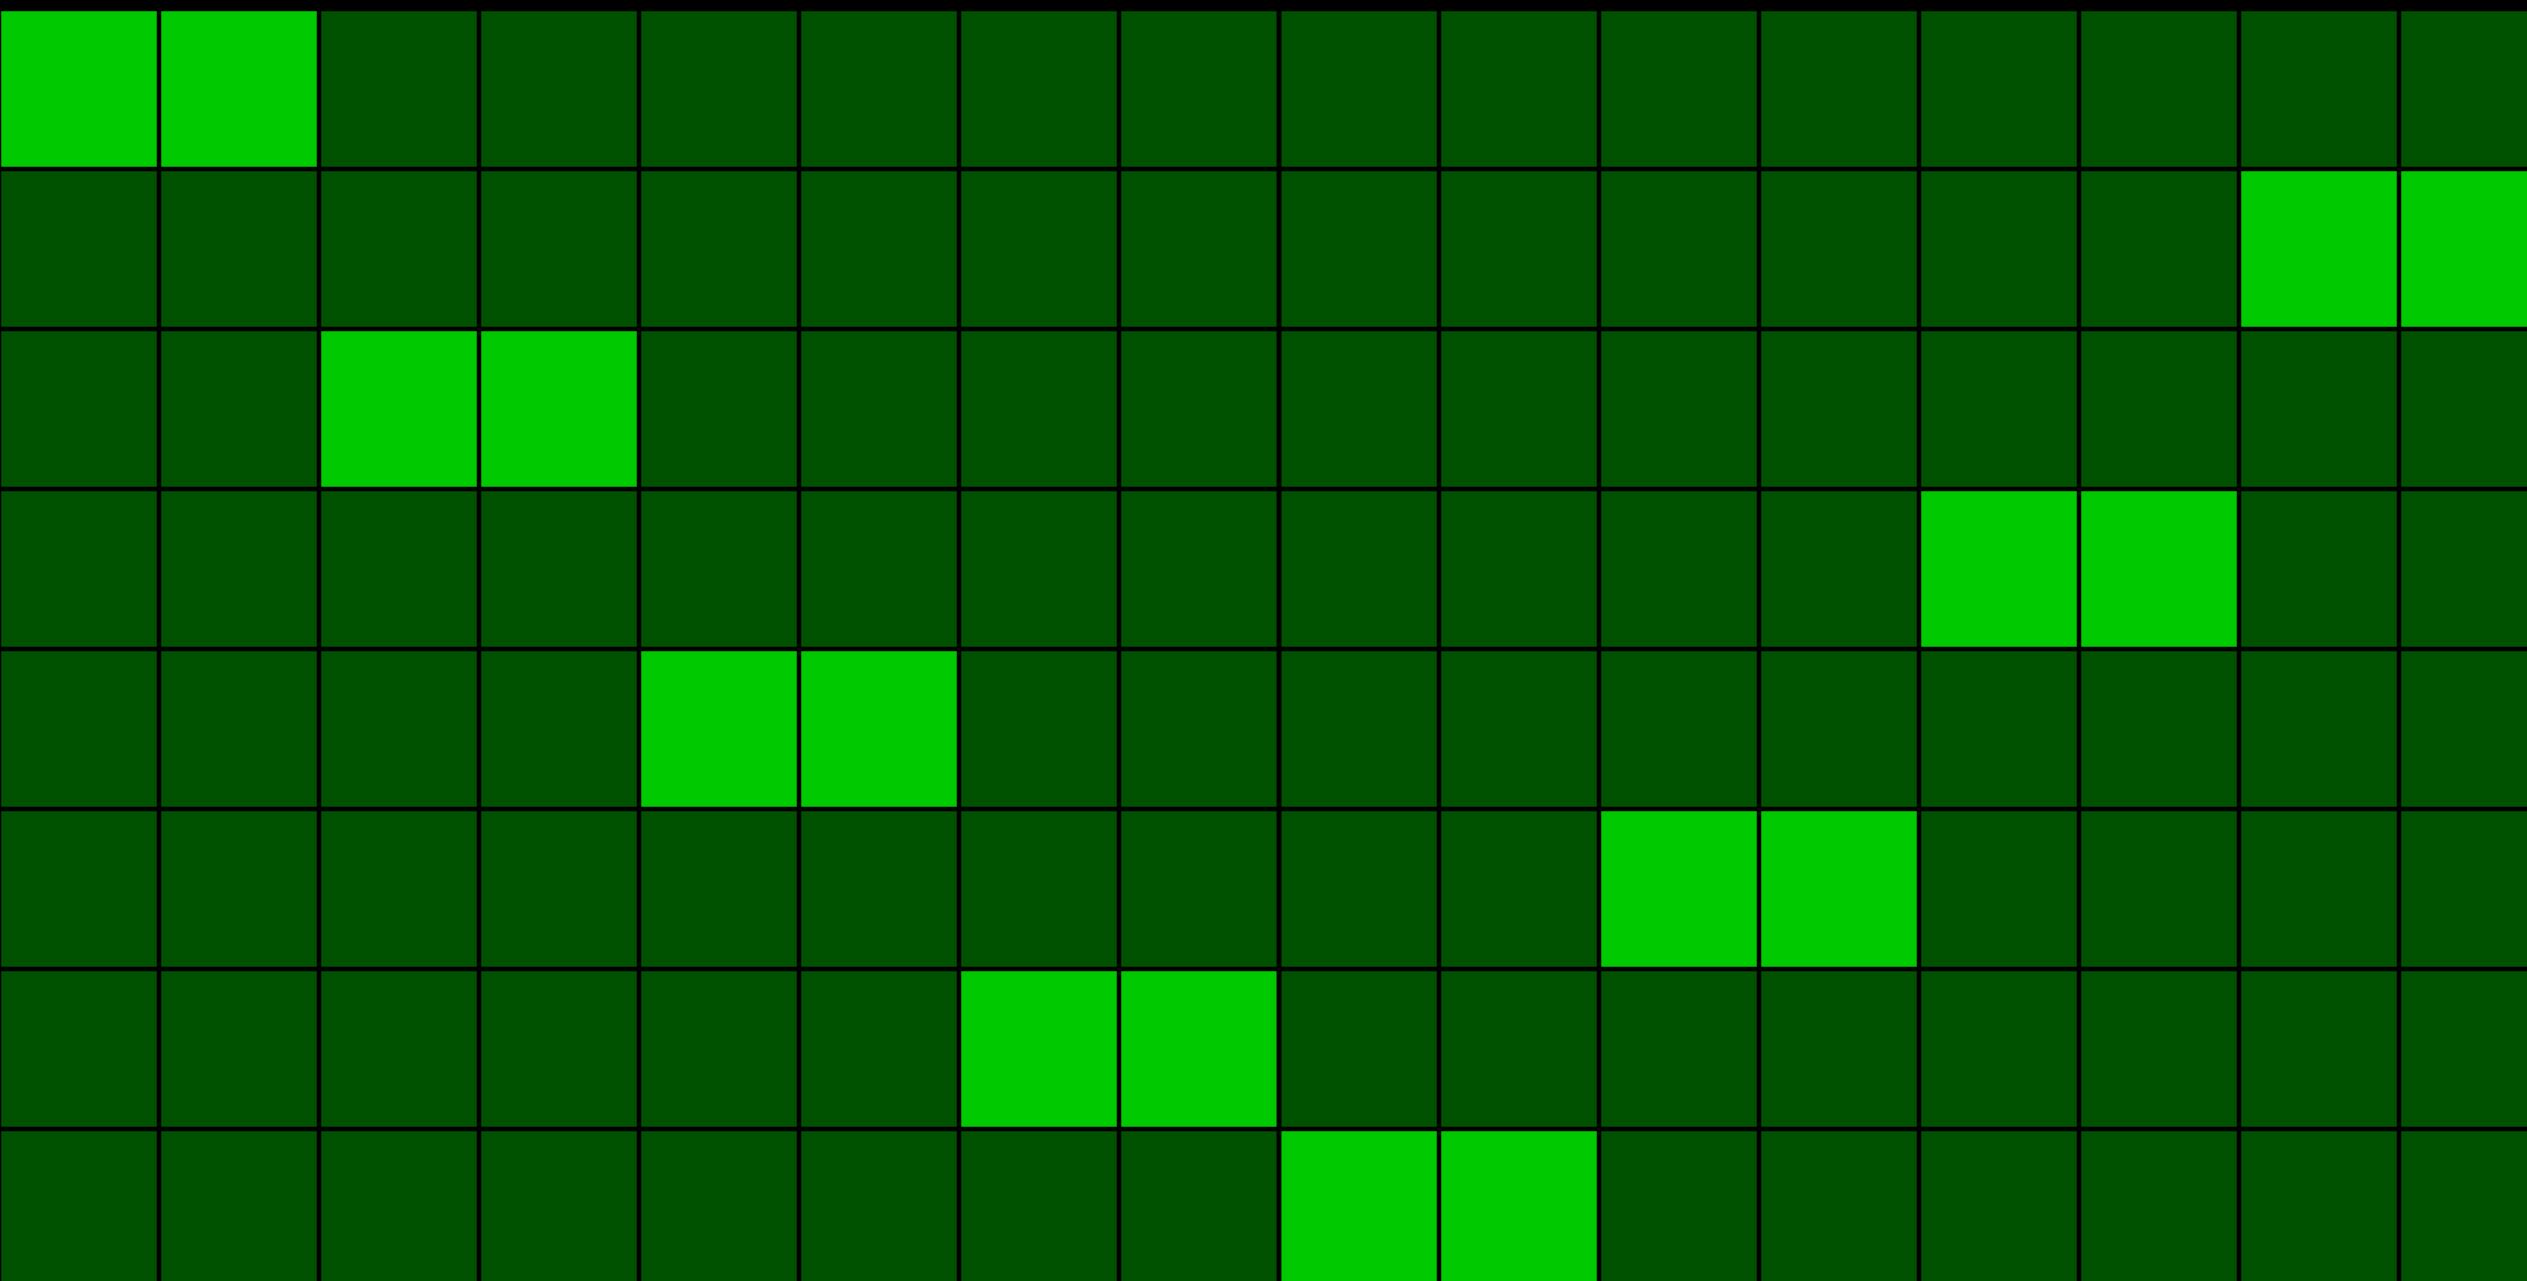
# rowhammer

## physical memory



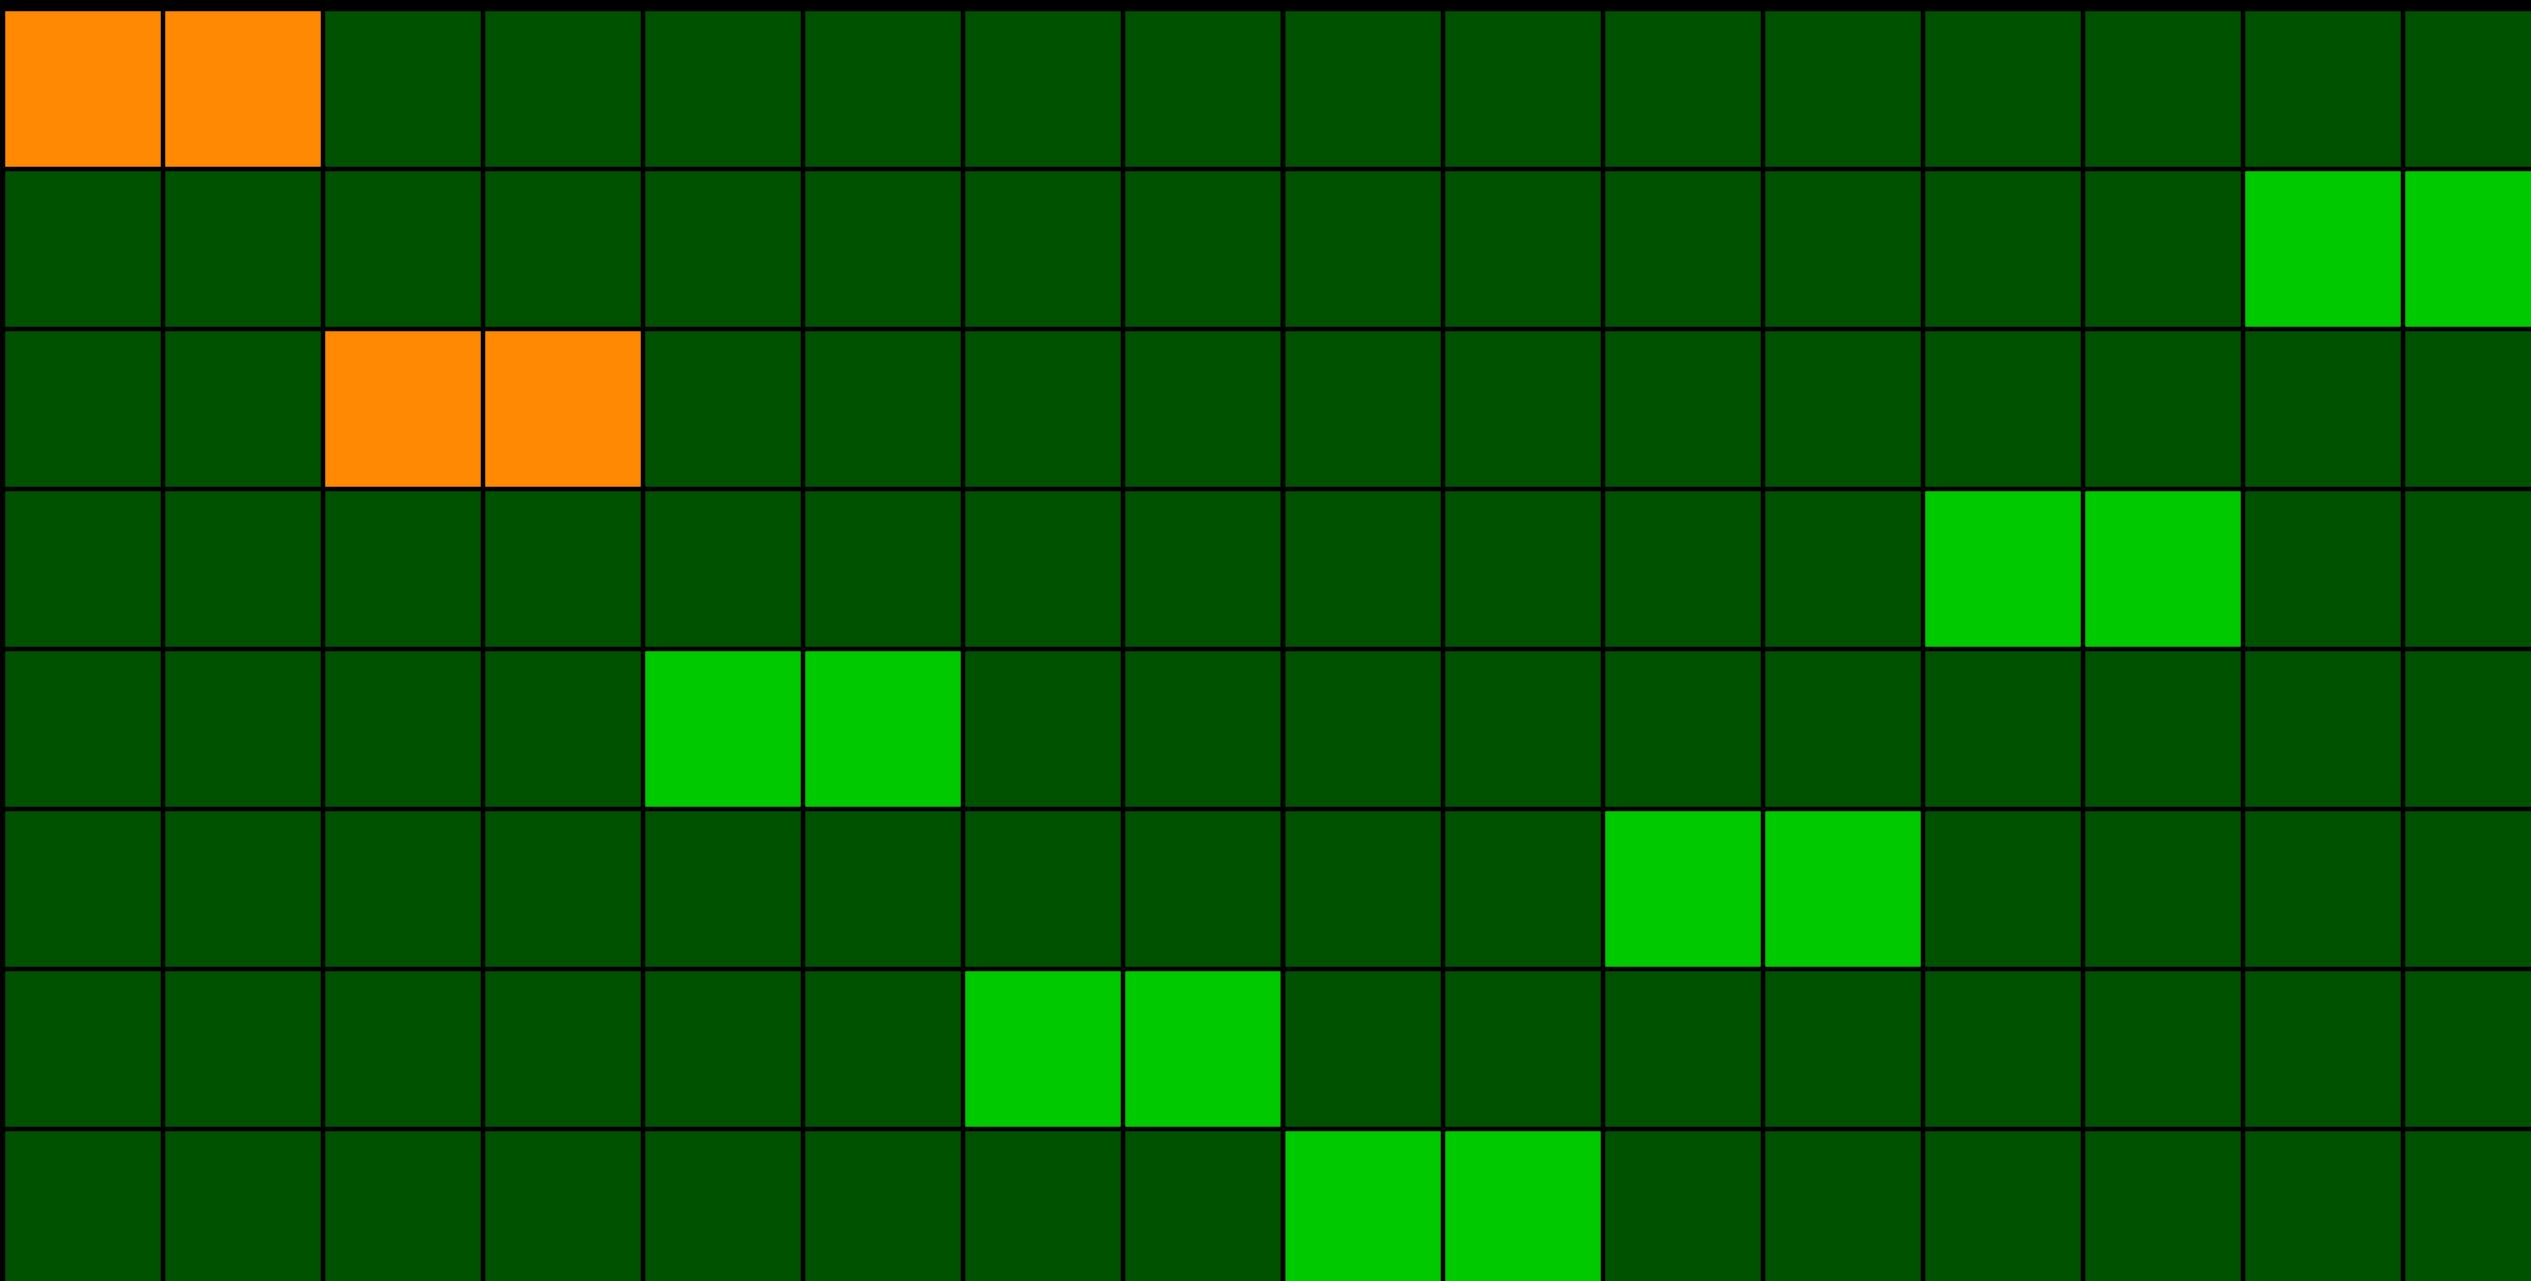
# rowhammer

## physical memory



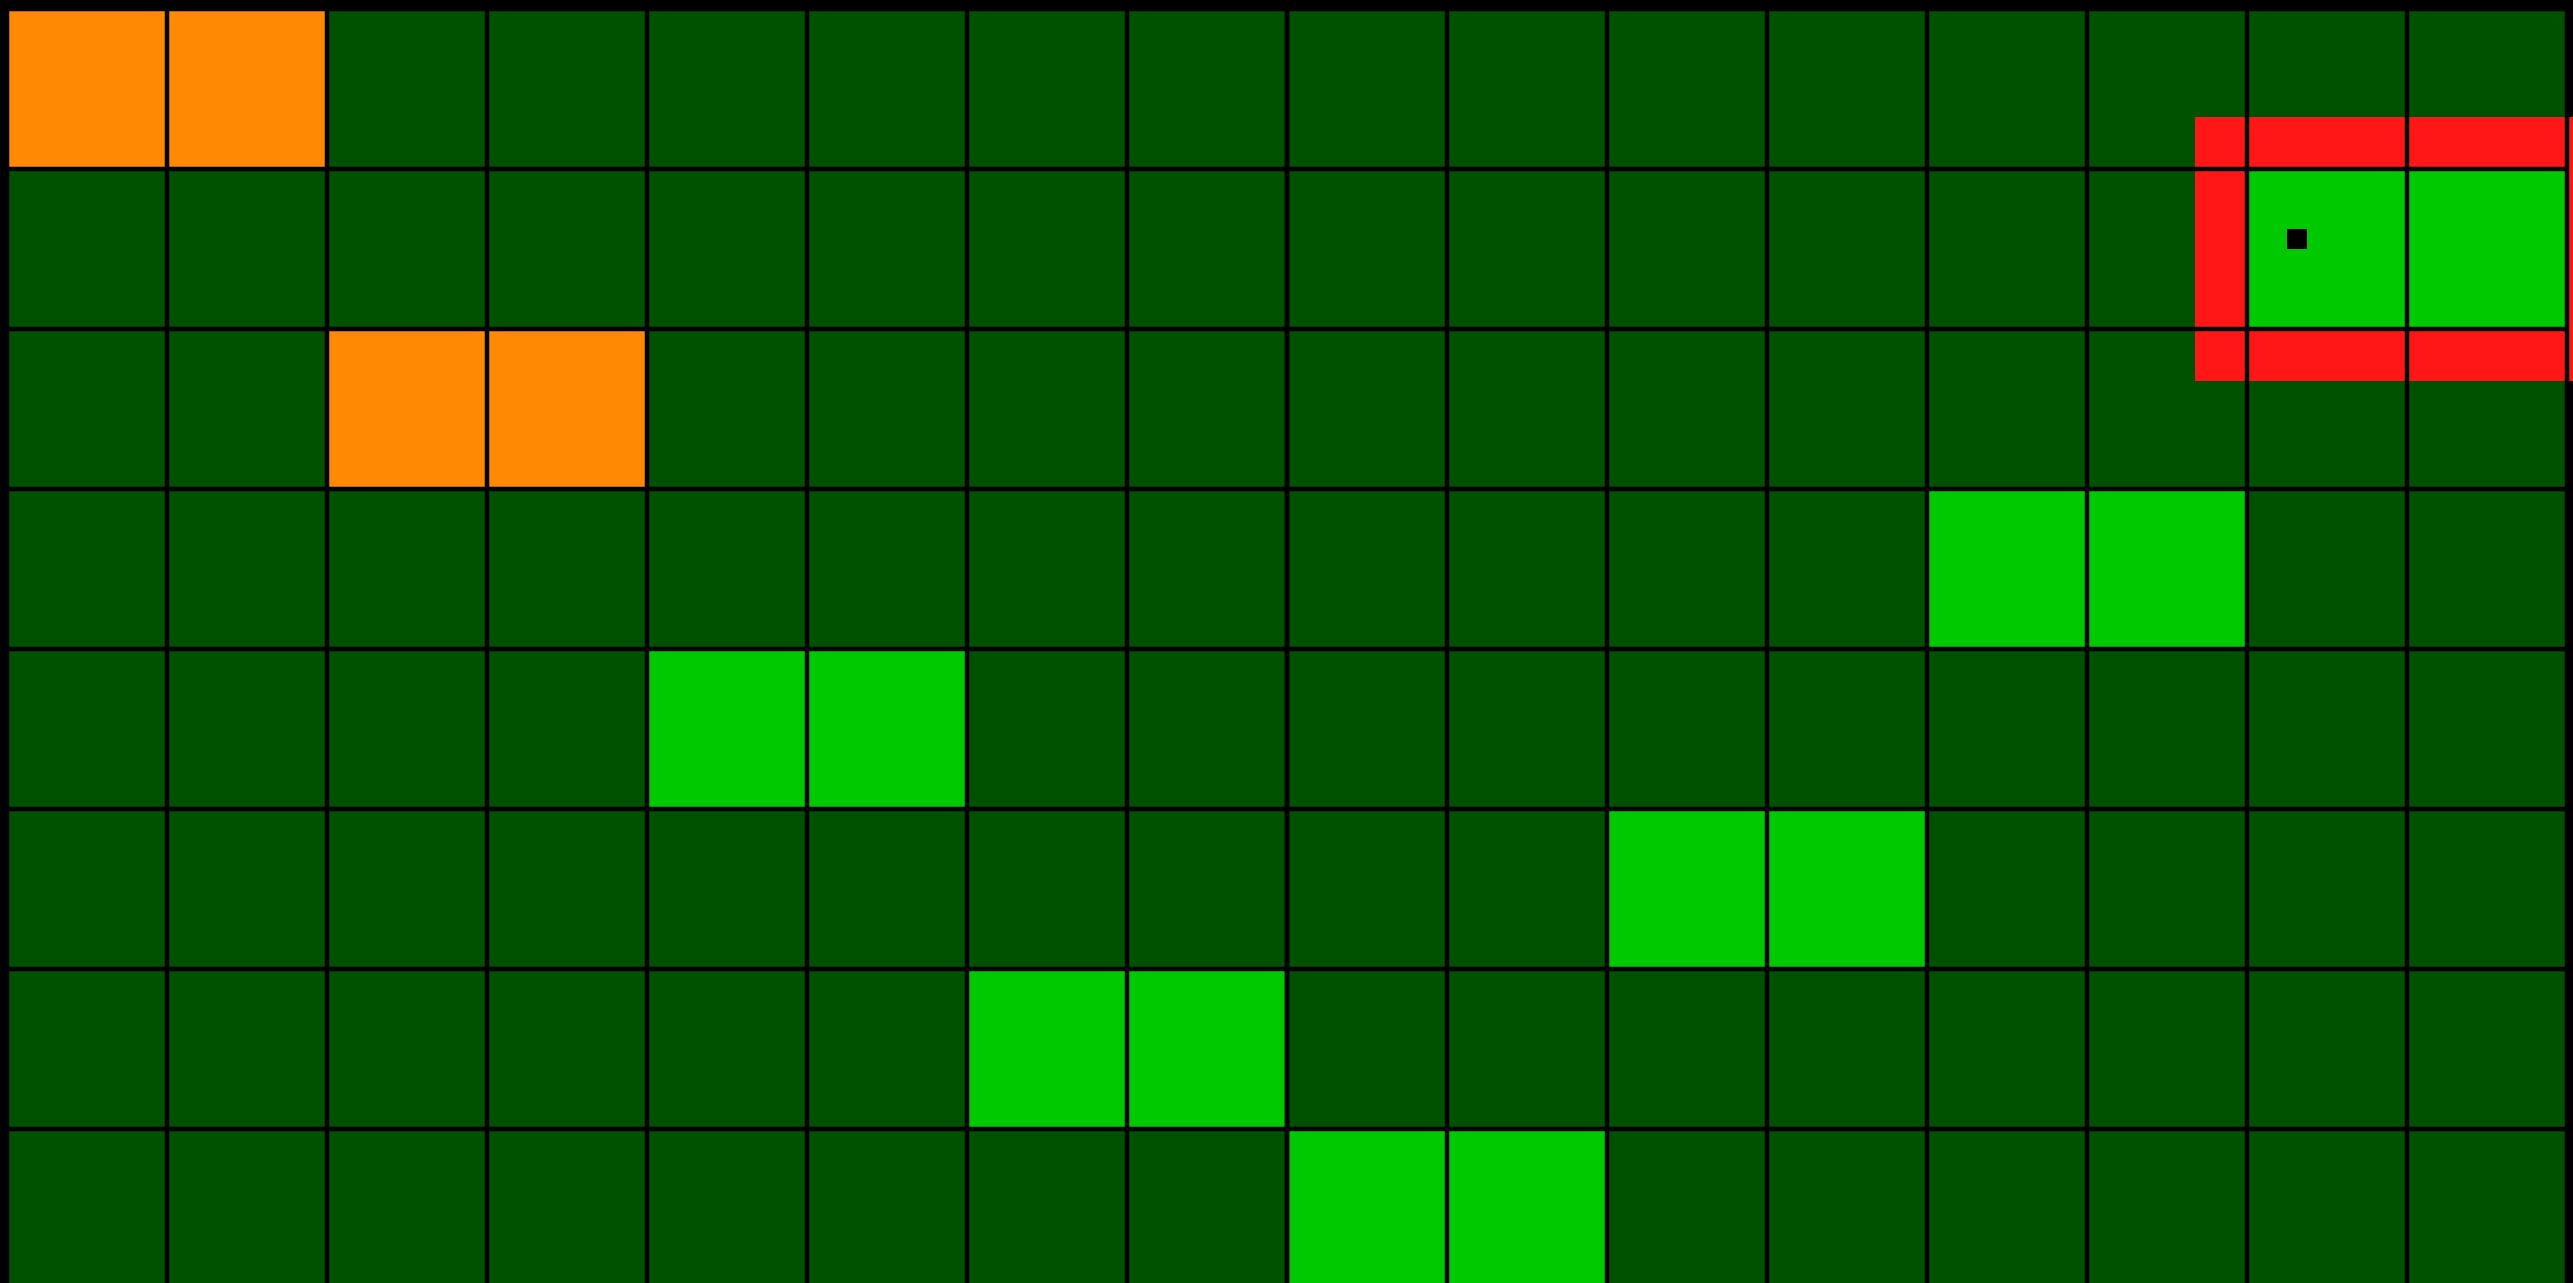
# rowhammer

## physical memory



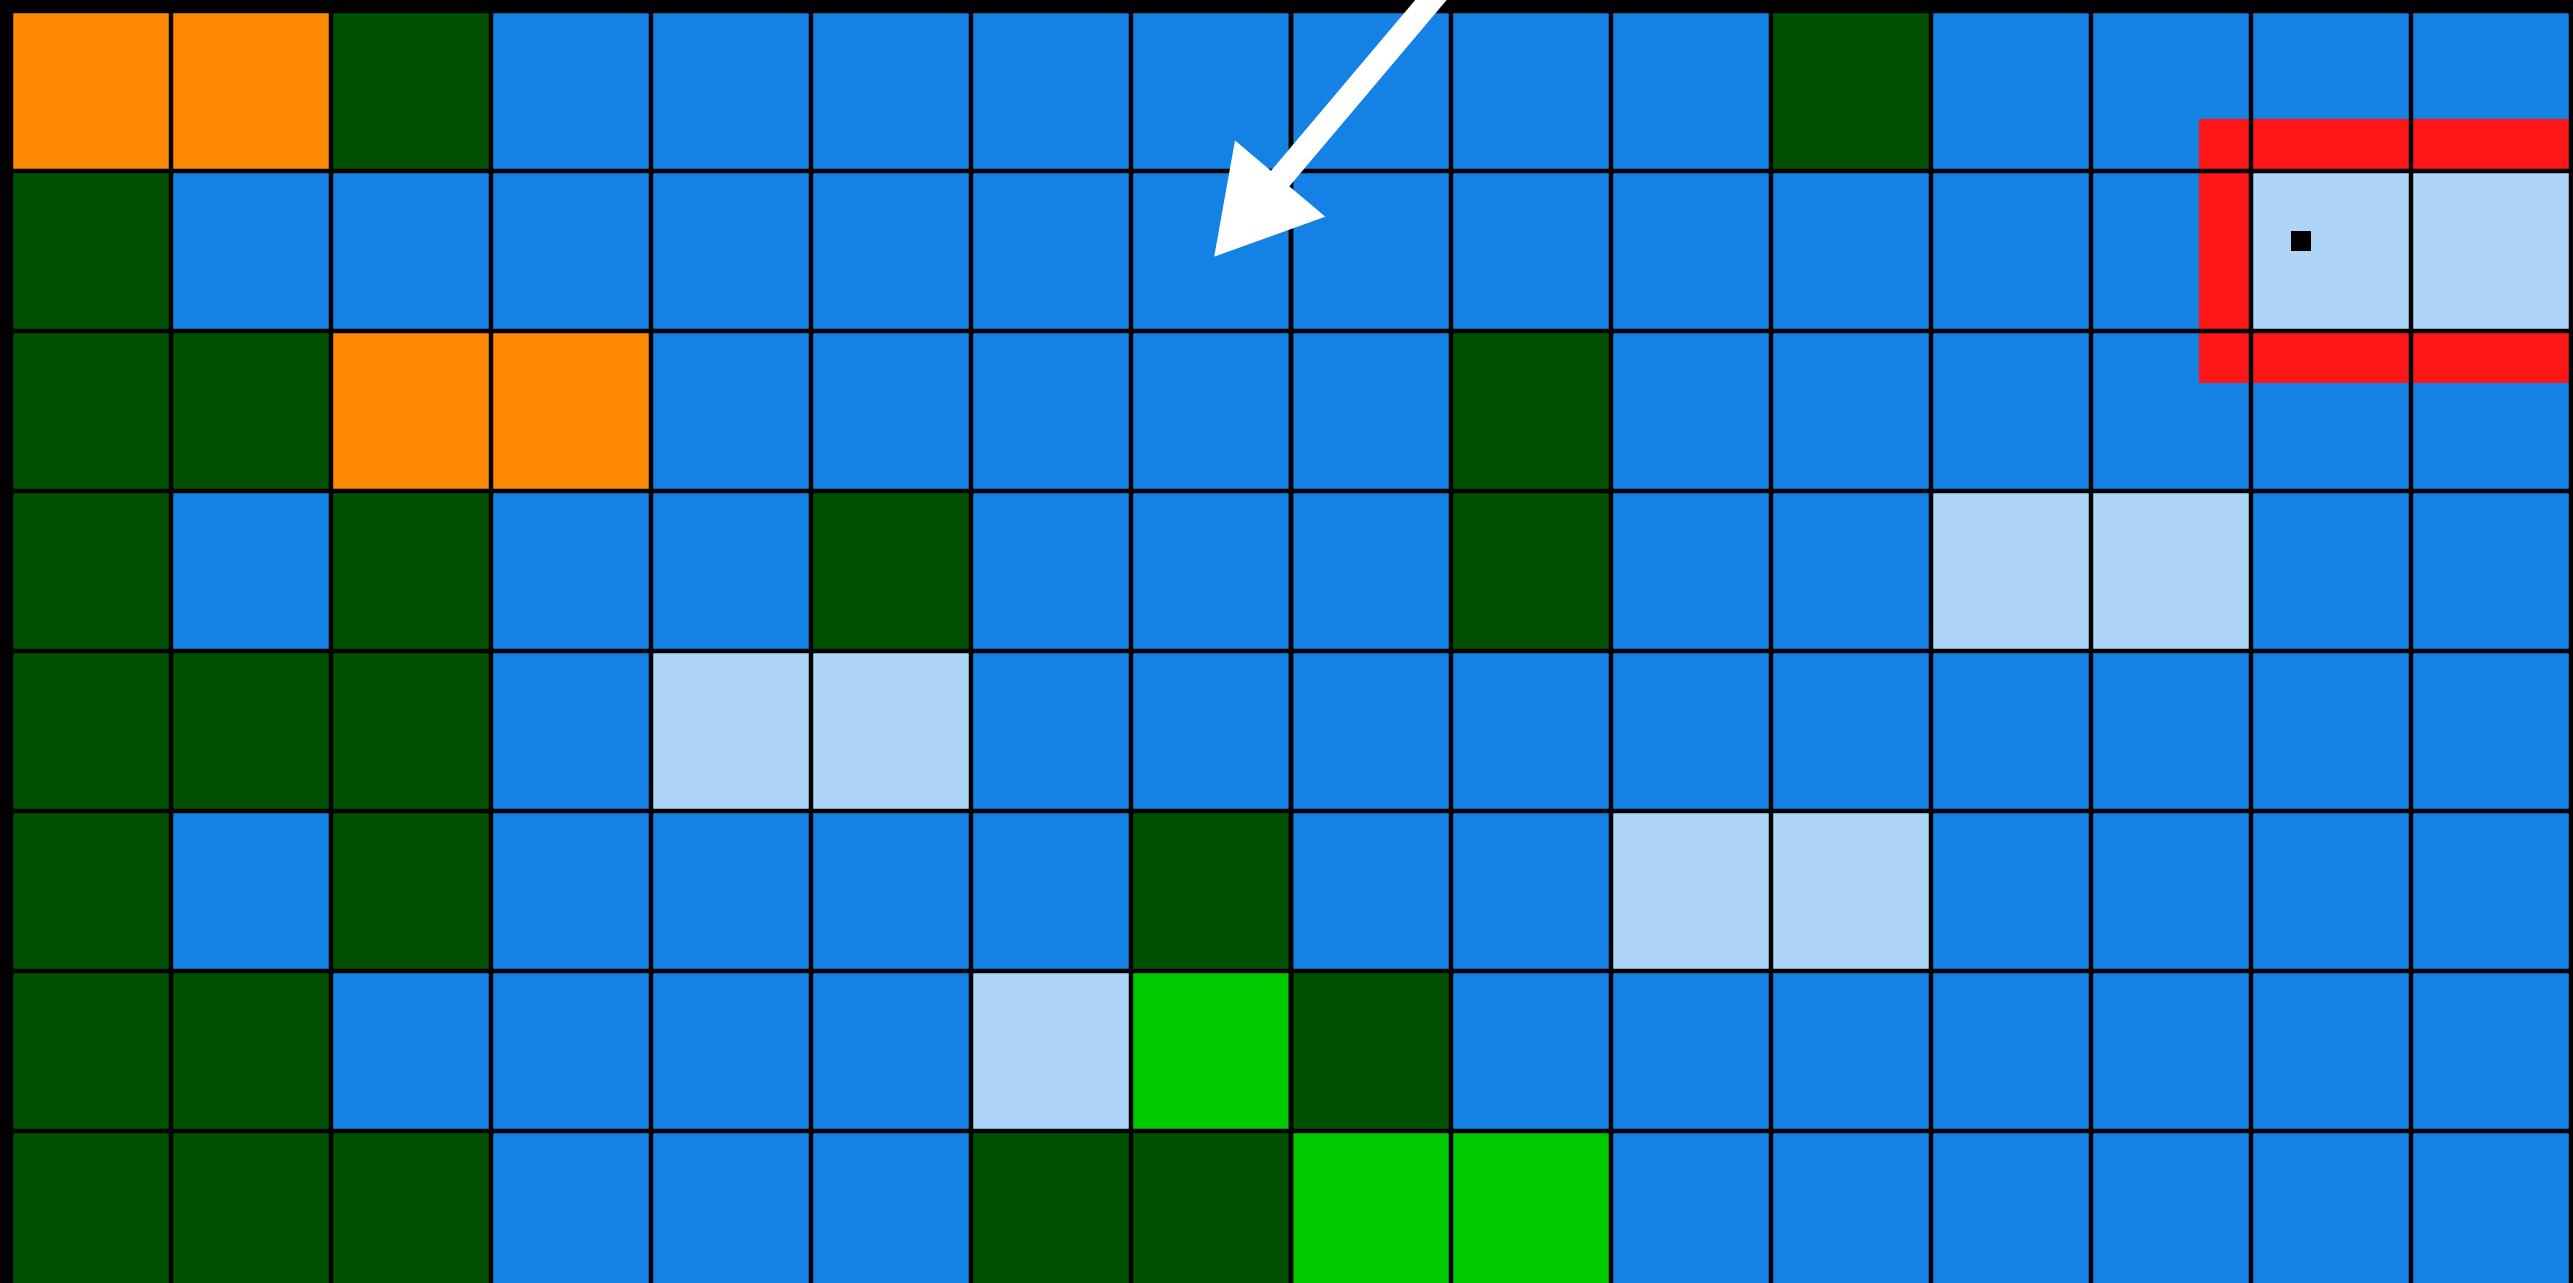
# rowhammer

## physical memory

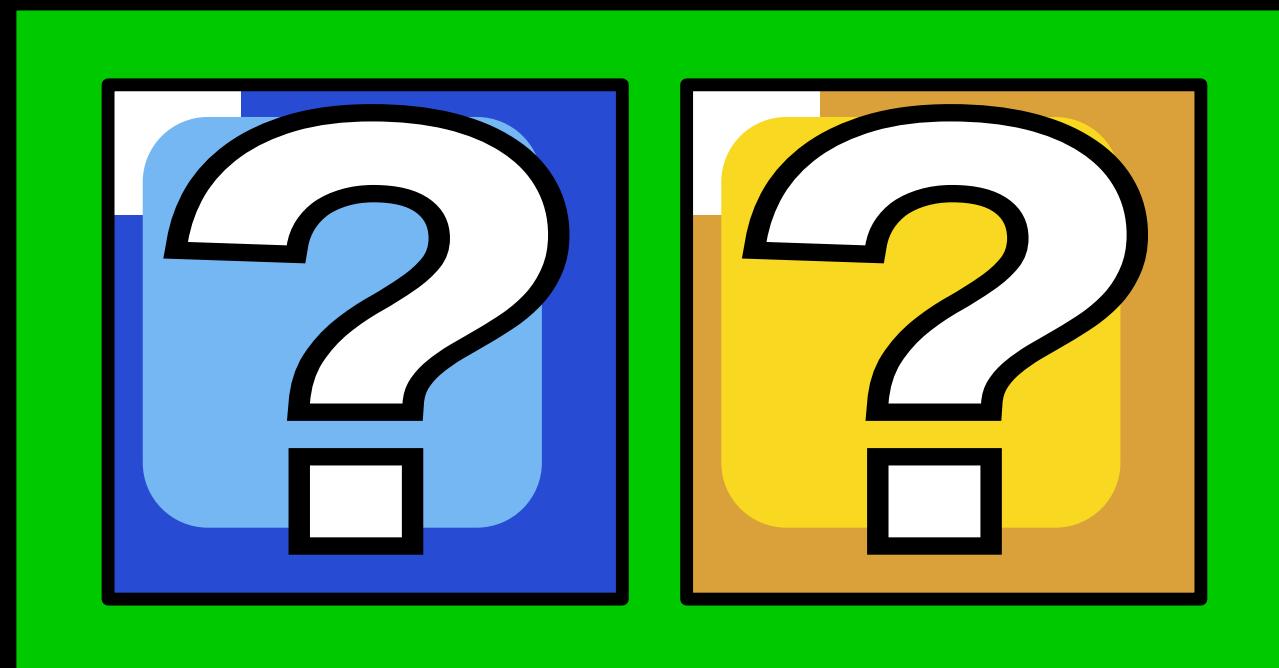


# rowhammer (seaborn attack)

physical memory      sprayed page tables



# fake Uint8Array object



# Data in Chakra's mixed-type arrays

# Data in Chakra's mixed-type arrays

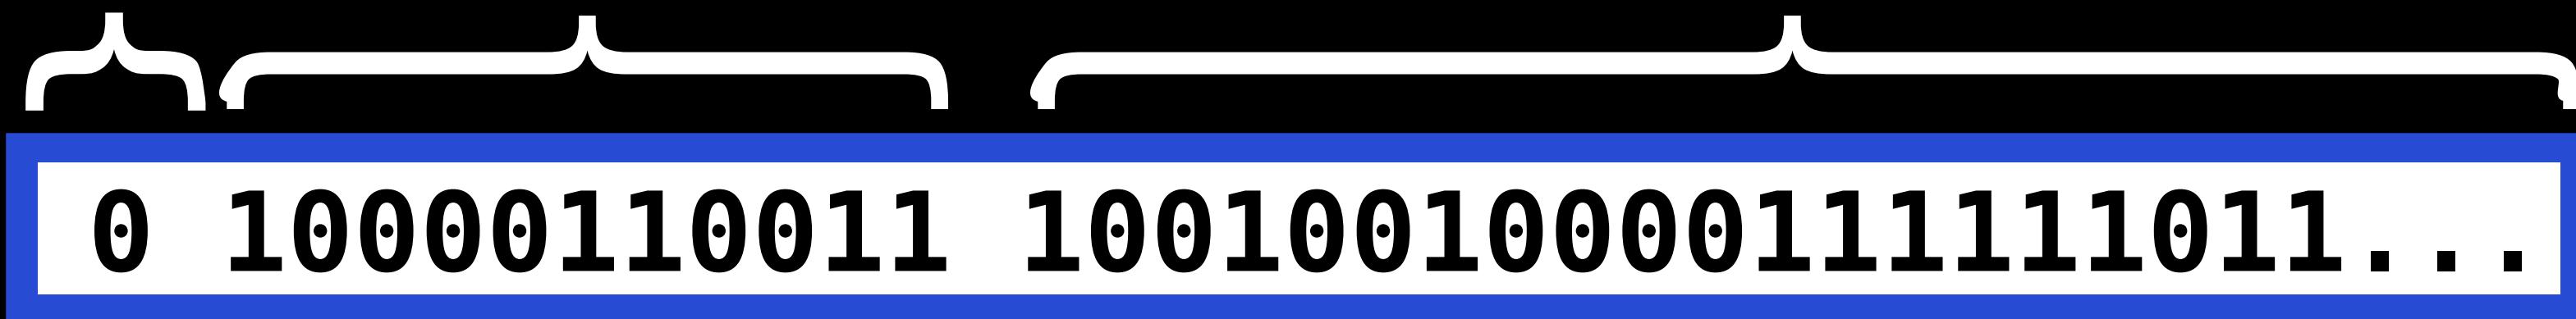
## Numbers (doubles)

```
0 10000110011 10010010000111111011. . .
```

# Data in Chakra's mixed-type arrays

## Numbers (doubles)

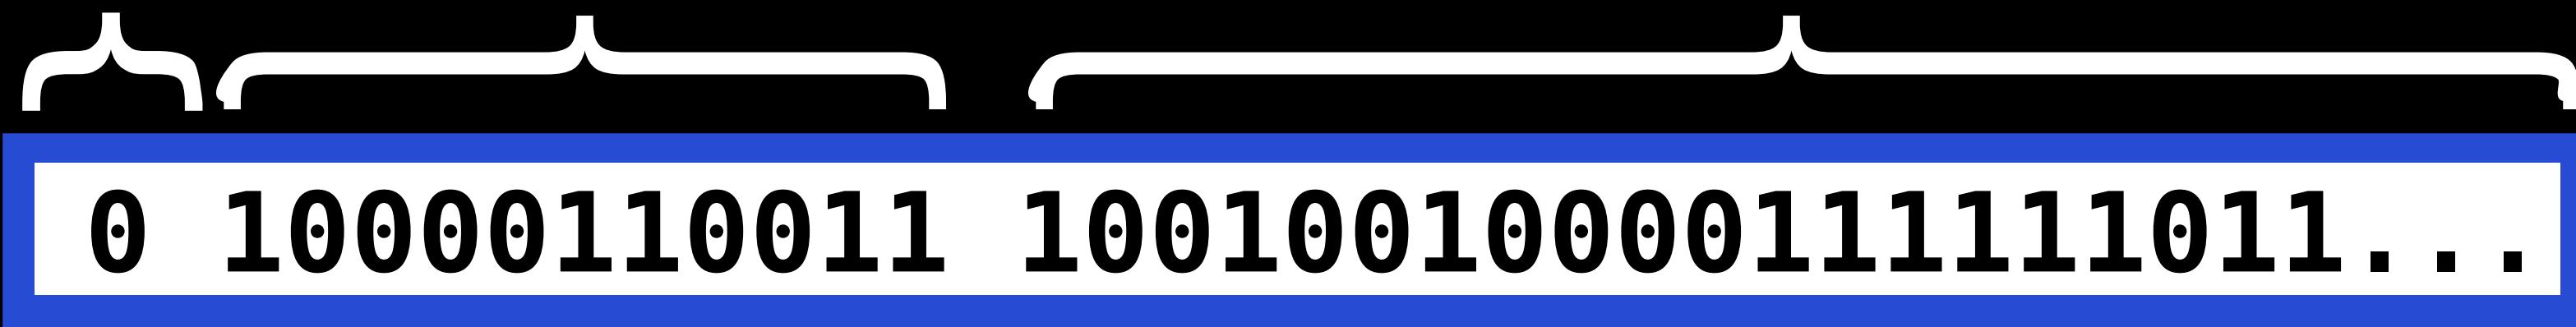
sign      exponent (11 bit)      mantissa (52 bit)



# Data in Chakra's mixed-type arrays

## Numbers (doubles)

sign      exponent (11 bit)      mantissa (52 bit)



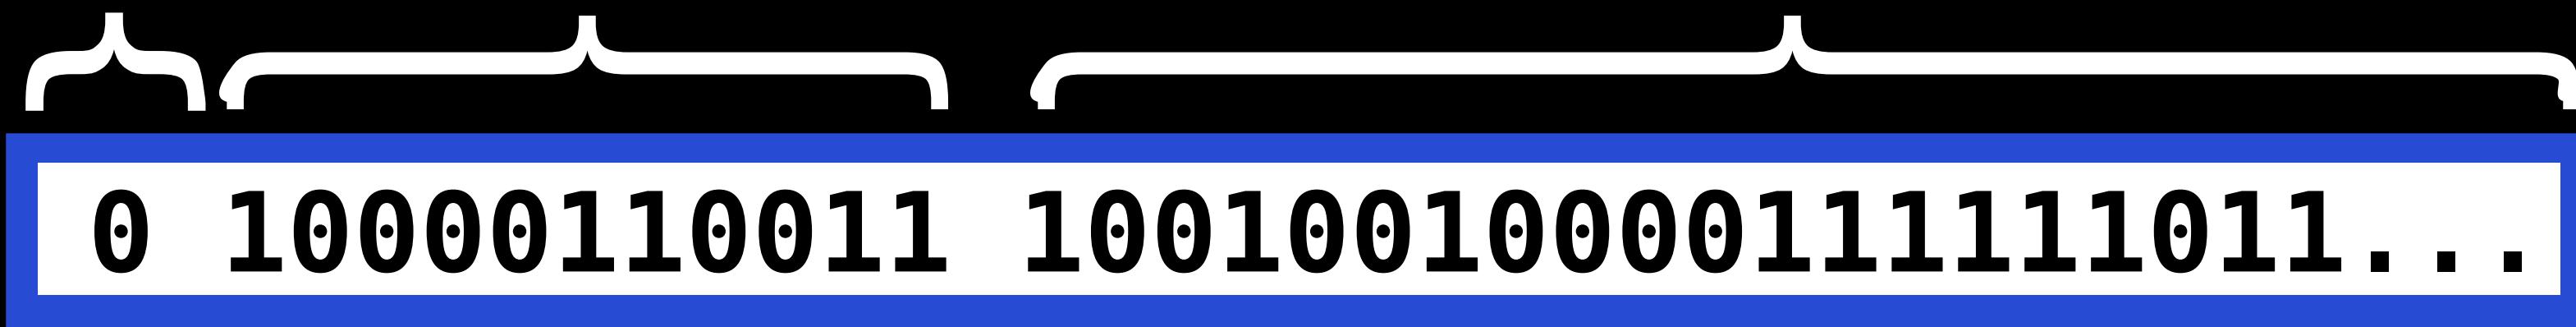
## Objects (pointers)

0000000000000000 00000010011100100100

# Data in Chakra's mixed-type arrays

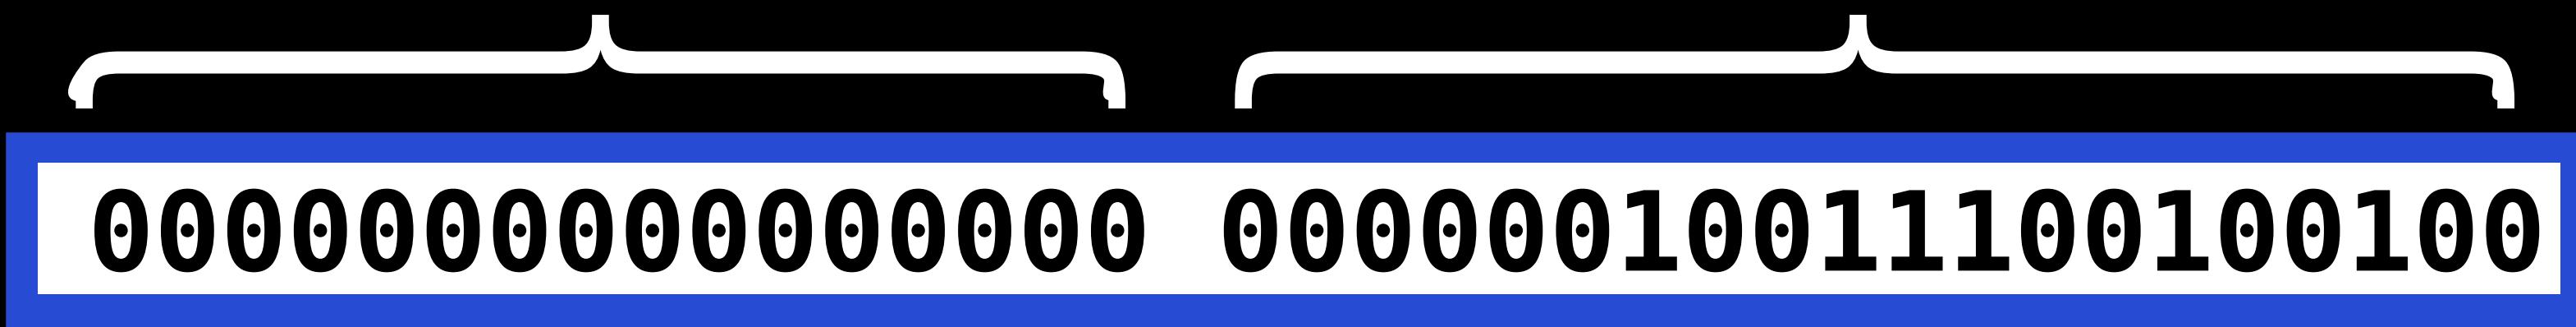
## Numbers (doubles)

sign      exponent (11 bit)      mantissa (52 bit)



## Objects (pointers)

sign extended      48 bit virtually addressable



# Data in Chakra's mixed-type arrays

Numbers, NaN

sign      exponent (11 bit)      52 bit unused

0 111111111111 ?????.?????.?????.?????. . .

A diagram illustrating the bit layout of a floating-point number. It consists of three fields: 'sign' (1 bit), 'exponent (11 bit)', and '52 bit unused'. The 'sign' field is at the far left. The 'exponent' field follows it, containing eleven bits. To the right of the exponent is a large bracket labeled '52 bit unused', which covers the remaining bits of the number.

Objects (pointers)

sign extended      48 bit virtually addressable

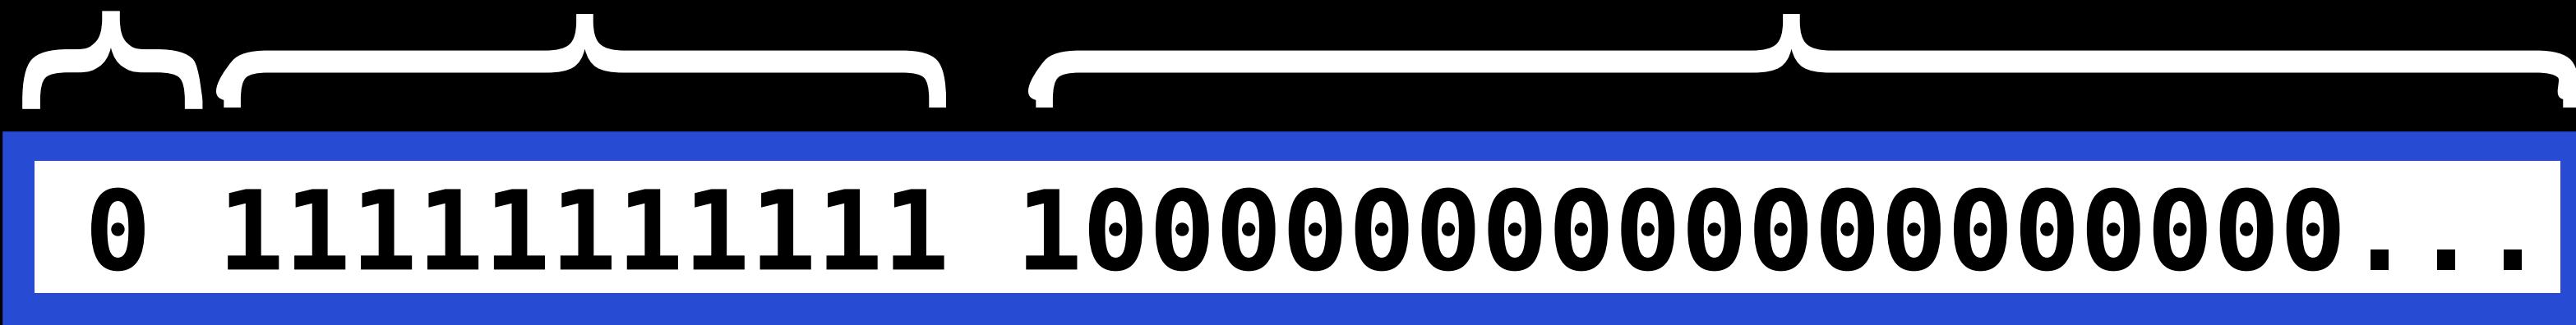
0000000000000000 00000010011100100100

A diagram illustrating the bit layout of a pointer. It consists of two fields: 'sign extended' (1 bit) and '48 bit virtually addressable' (47 bits). The 'sign extended' field is at the far left. The '48 bit virtually addressable' field follows it, containing forty-seven bits. Both fields are enclosed in a blue rectangular box.

# Data in Chakra's mixed-type arrays

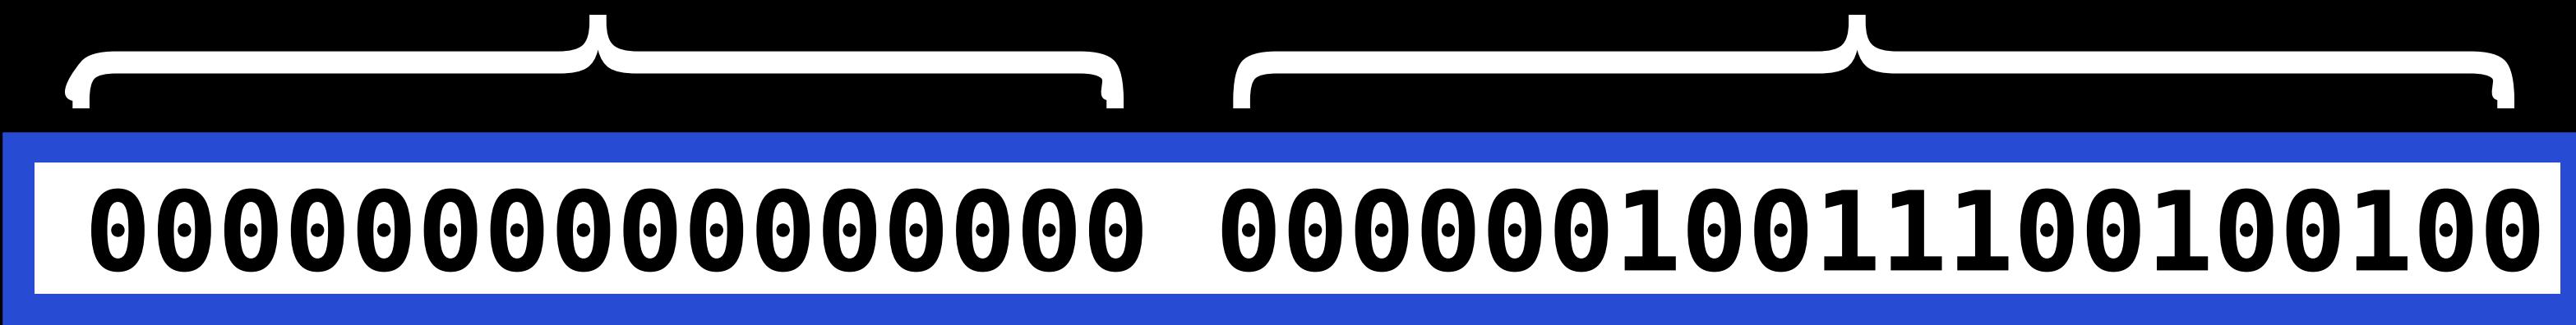
# Numbers, NaN

**sign      exponent (11 bit)      52 bit unused**



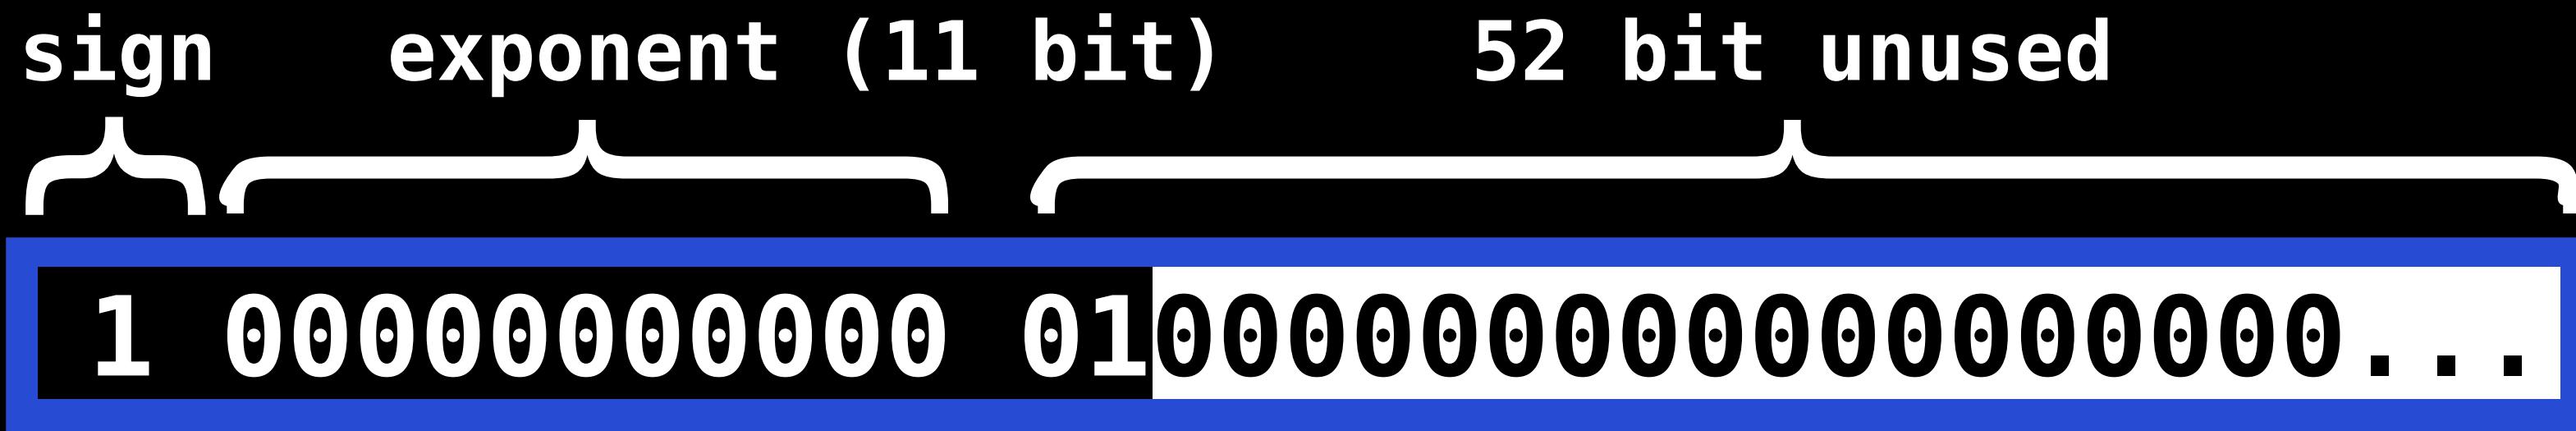
# Objects (pointers)

sign extended      48 bit virtually addressable

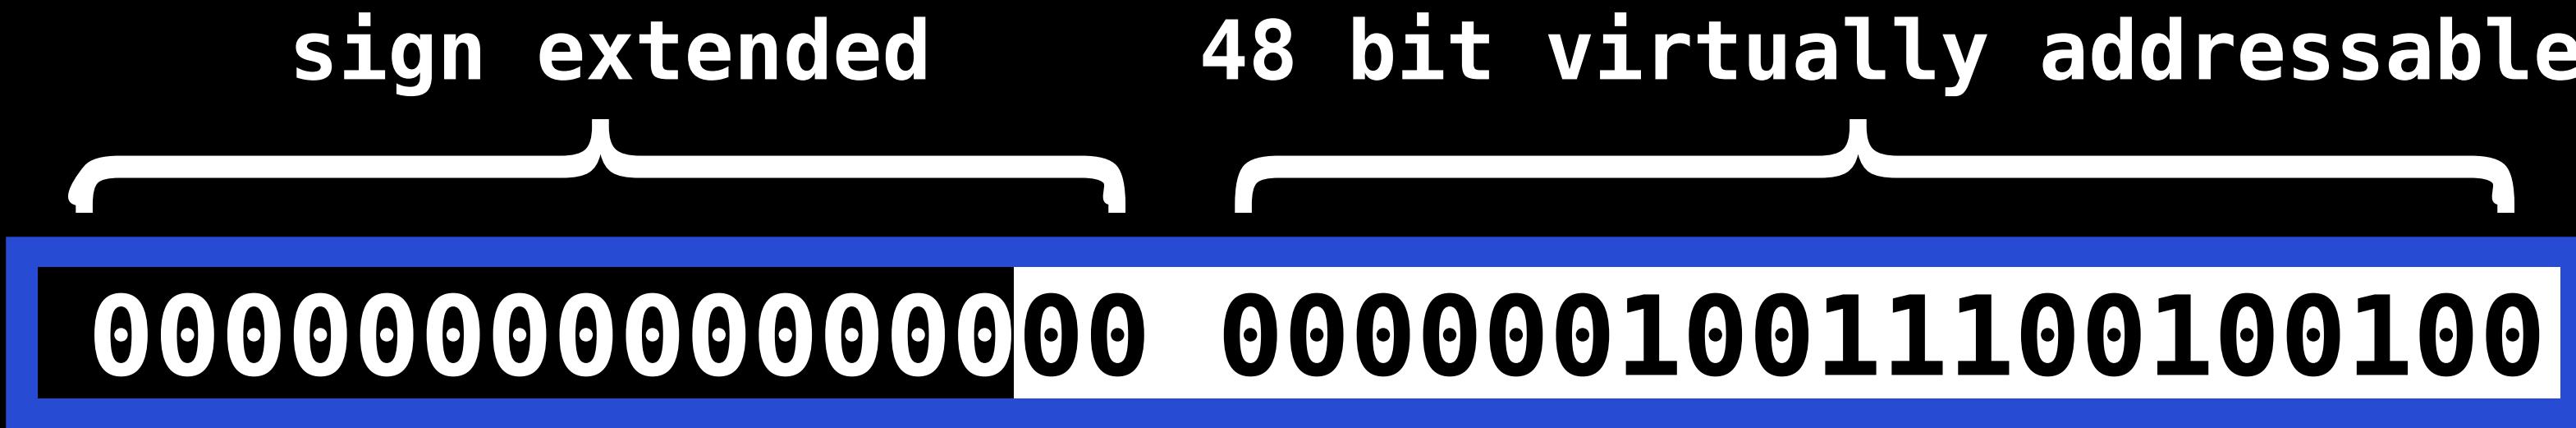


# Data in Chakra's mixed-type arrays

# Numbers, NaN, 14 high bits flipped

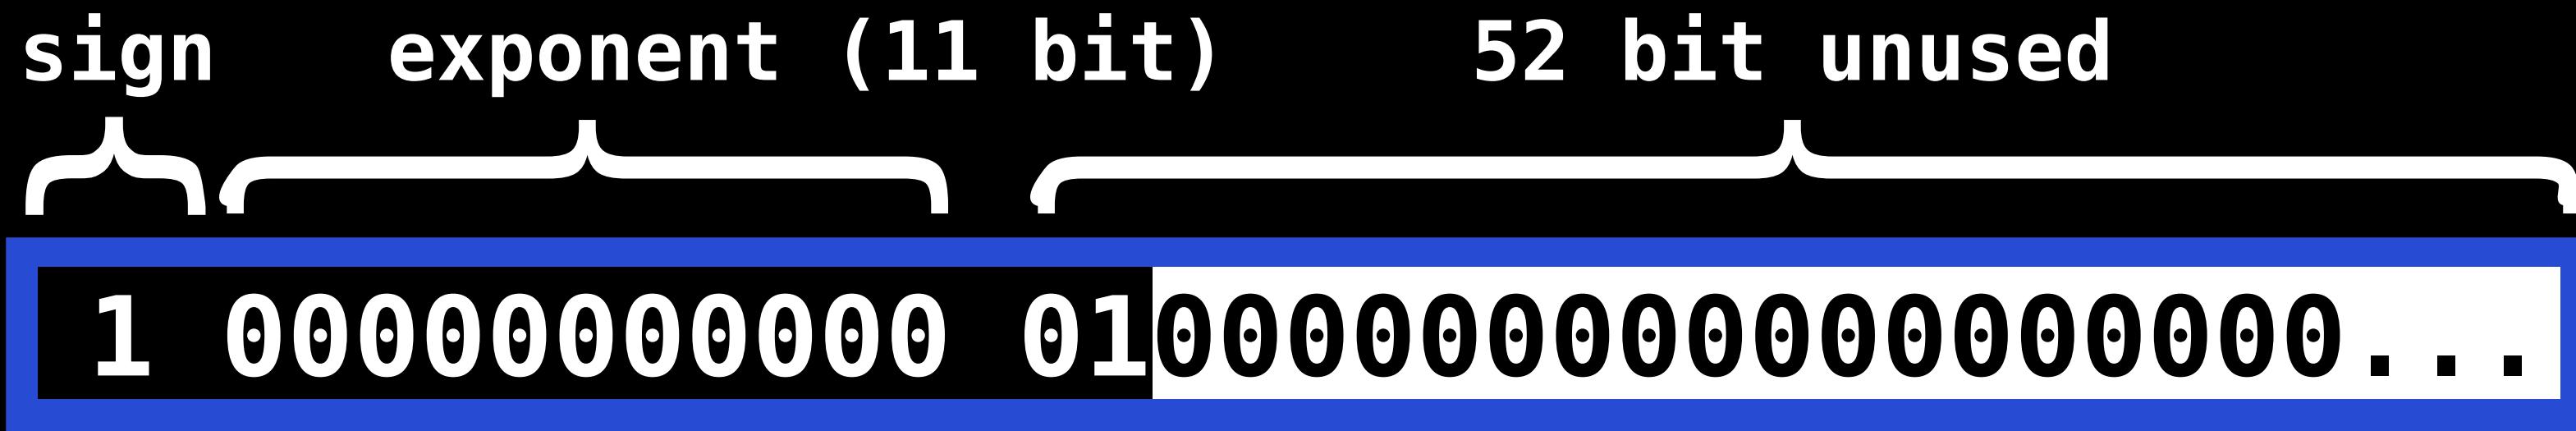


# Objects (pointers)

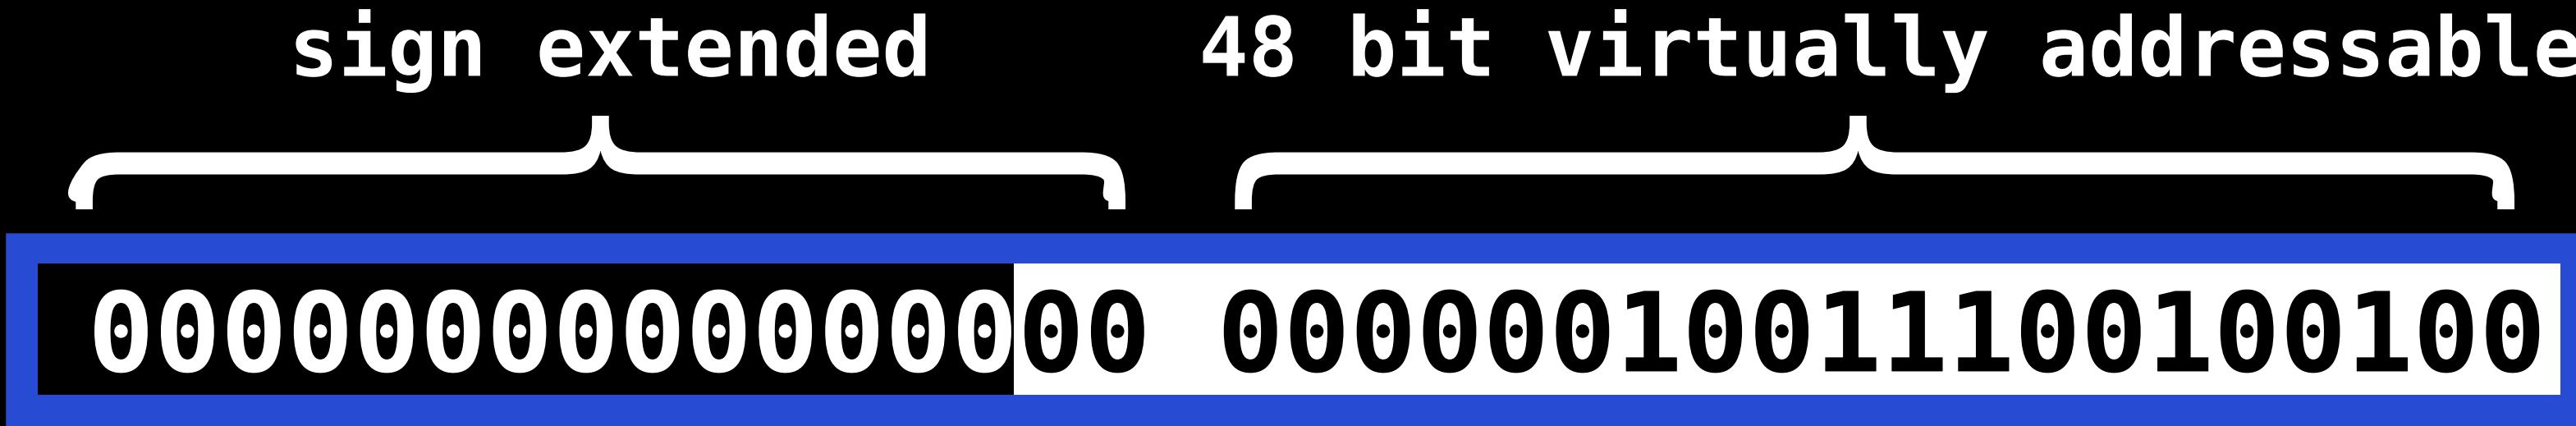


# Data in Chakra's mixed-type arrays

# Numbers, NaN, 14 high bits flipped

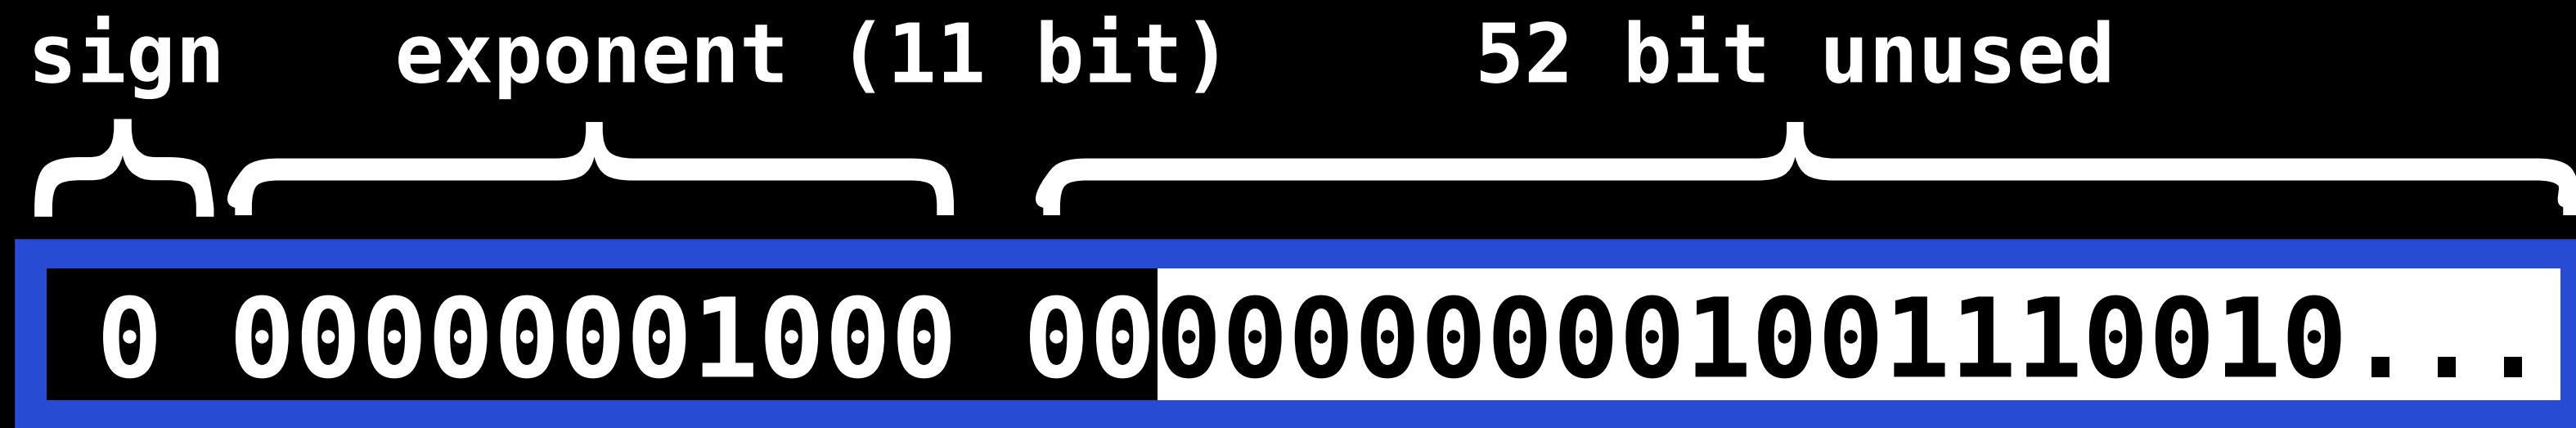


# Objects, Non-canonical -NaN



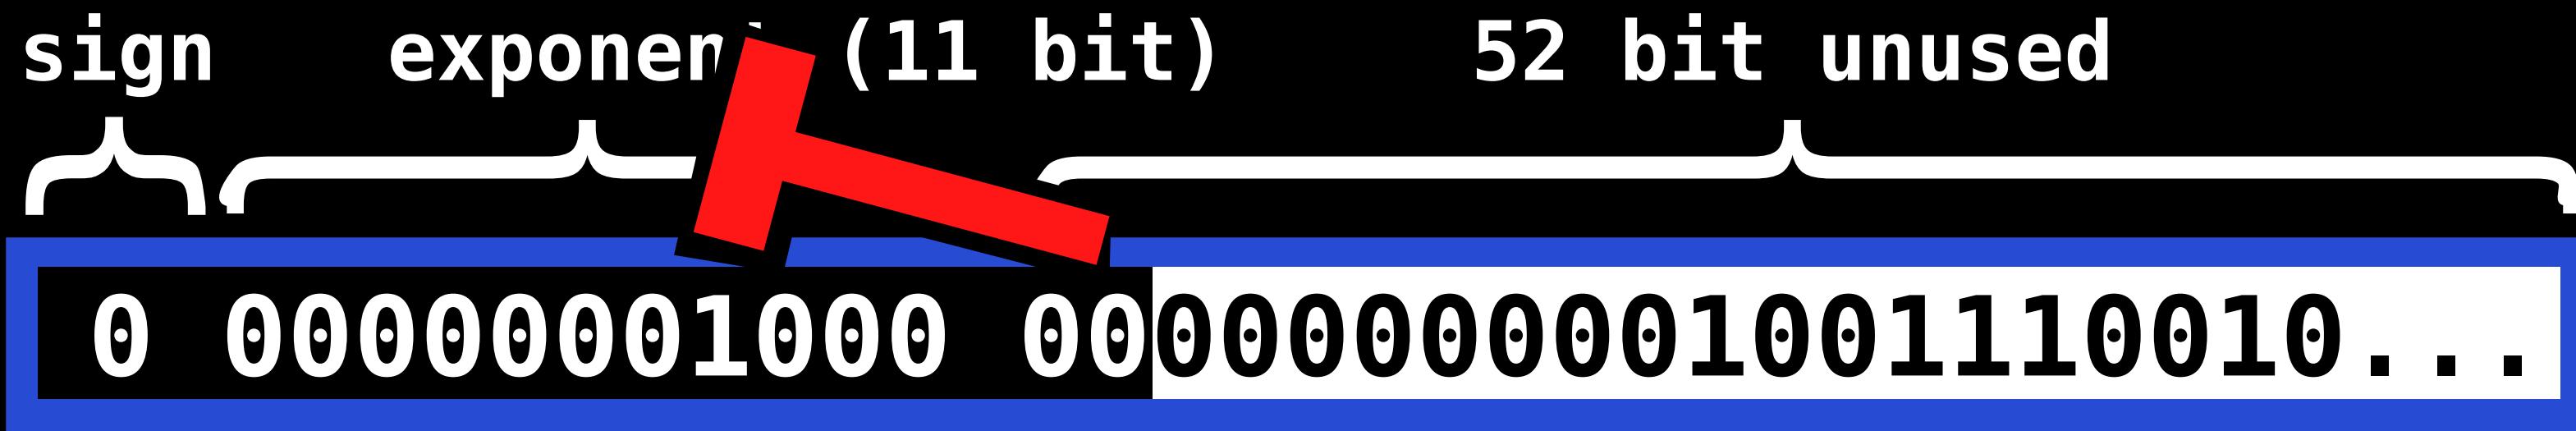
# Type Flipping

Number, 1 bit flip away from a pointer

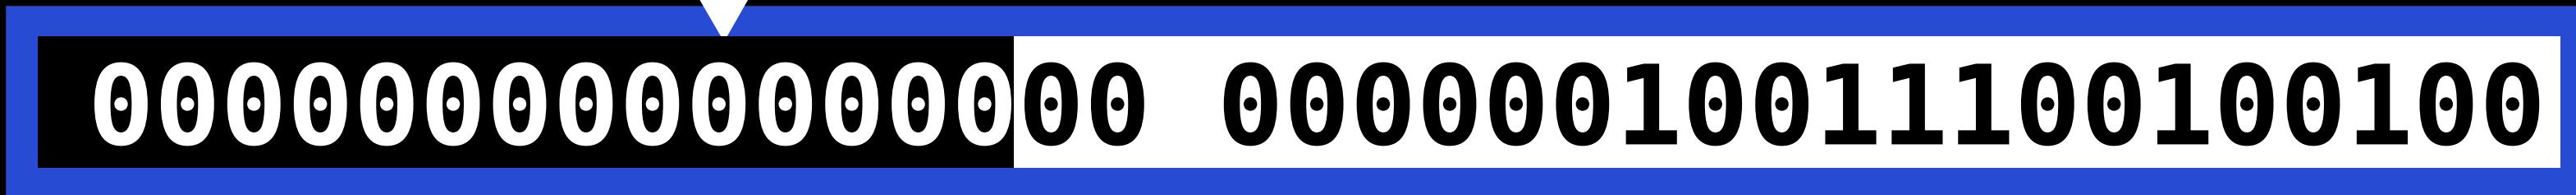
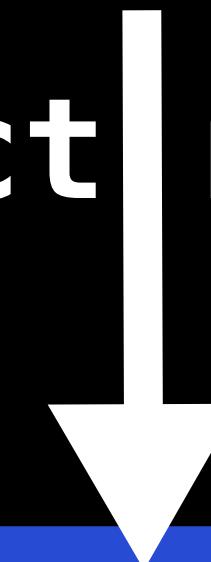


# Type Flipping

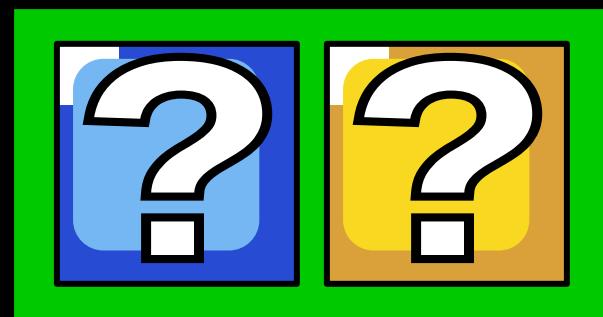
Number, 1 bit flip away from a pointer



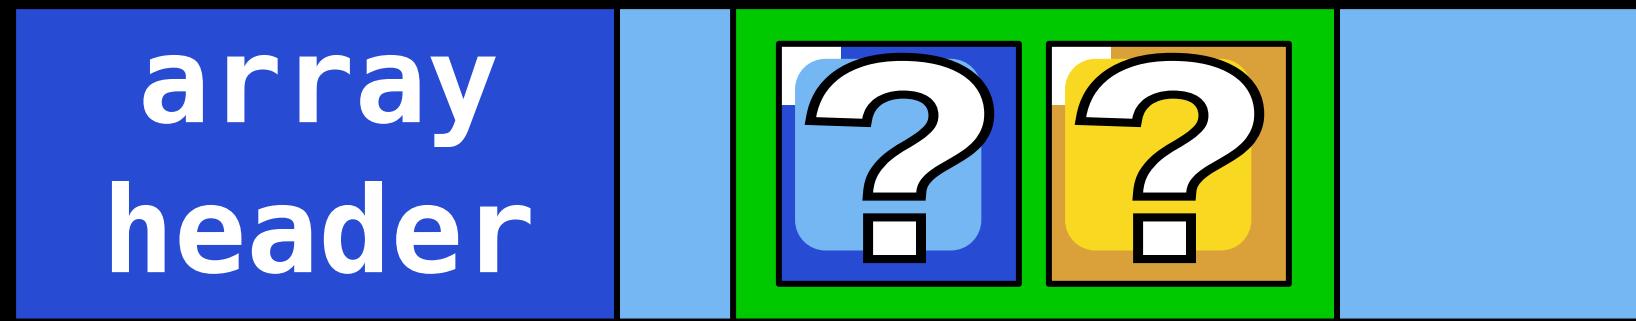
Fake Object reference



# pointer pivoting

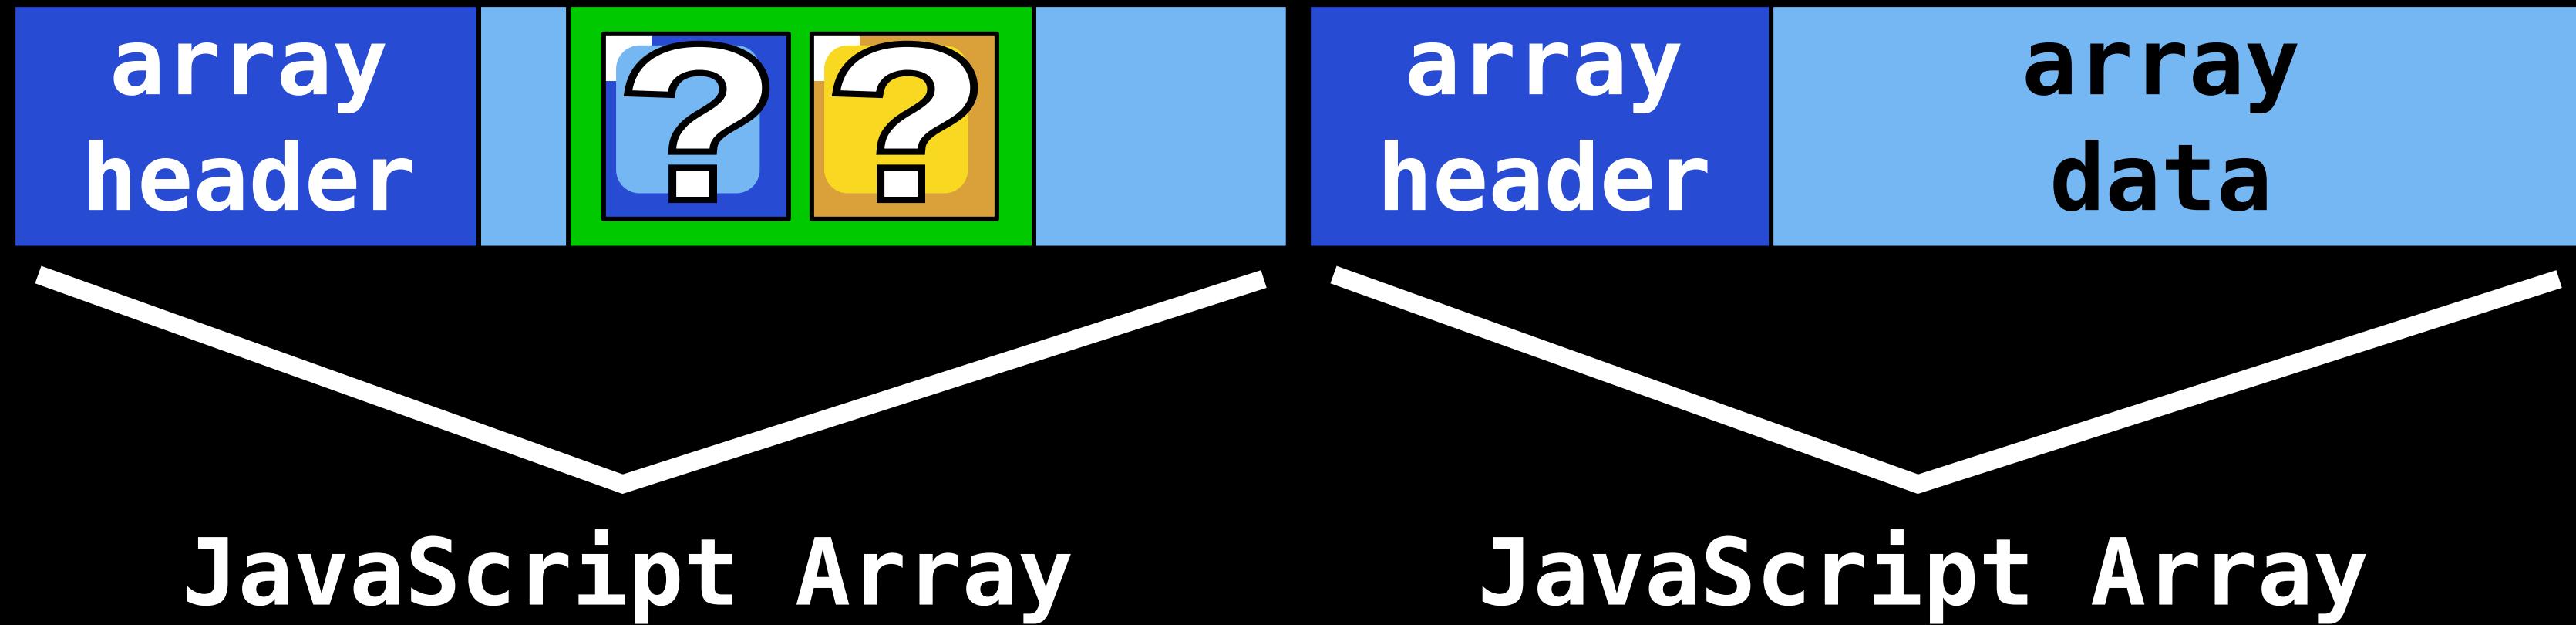


# pointer pivoting

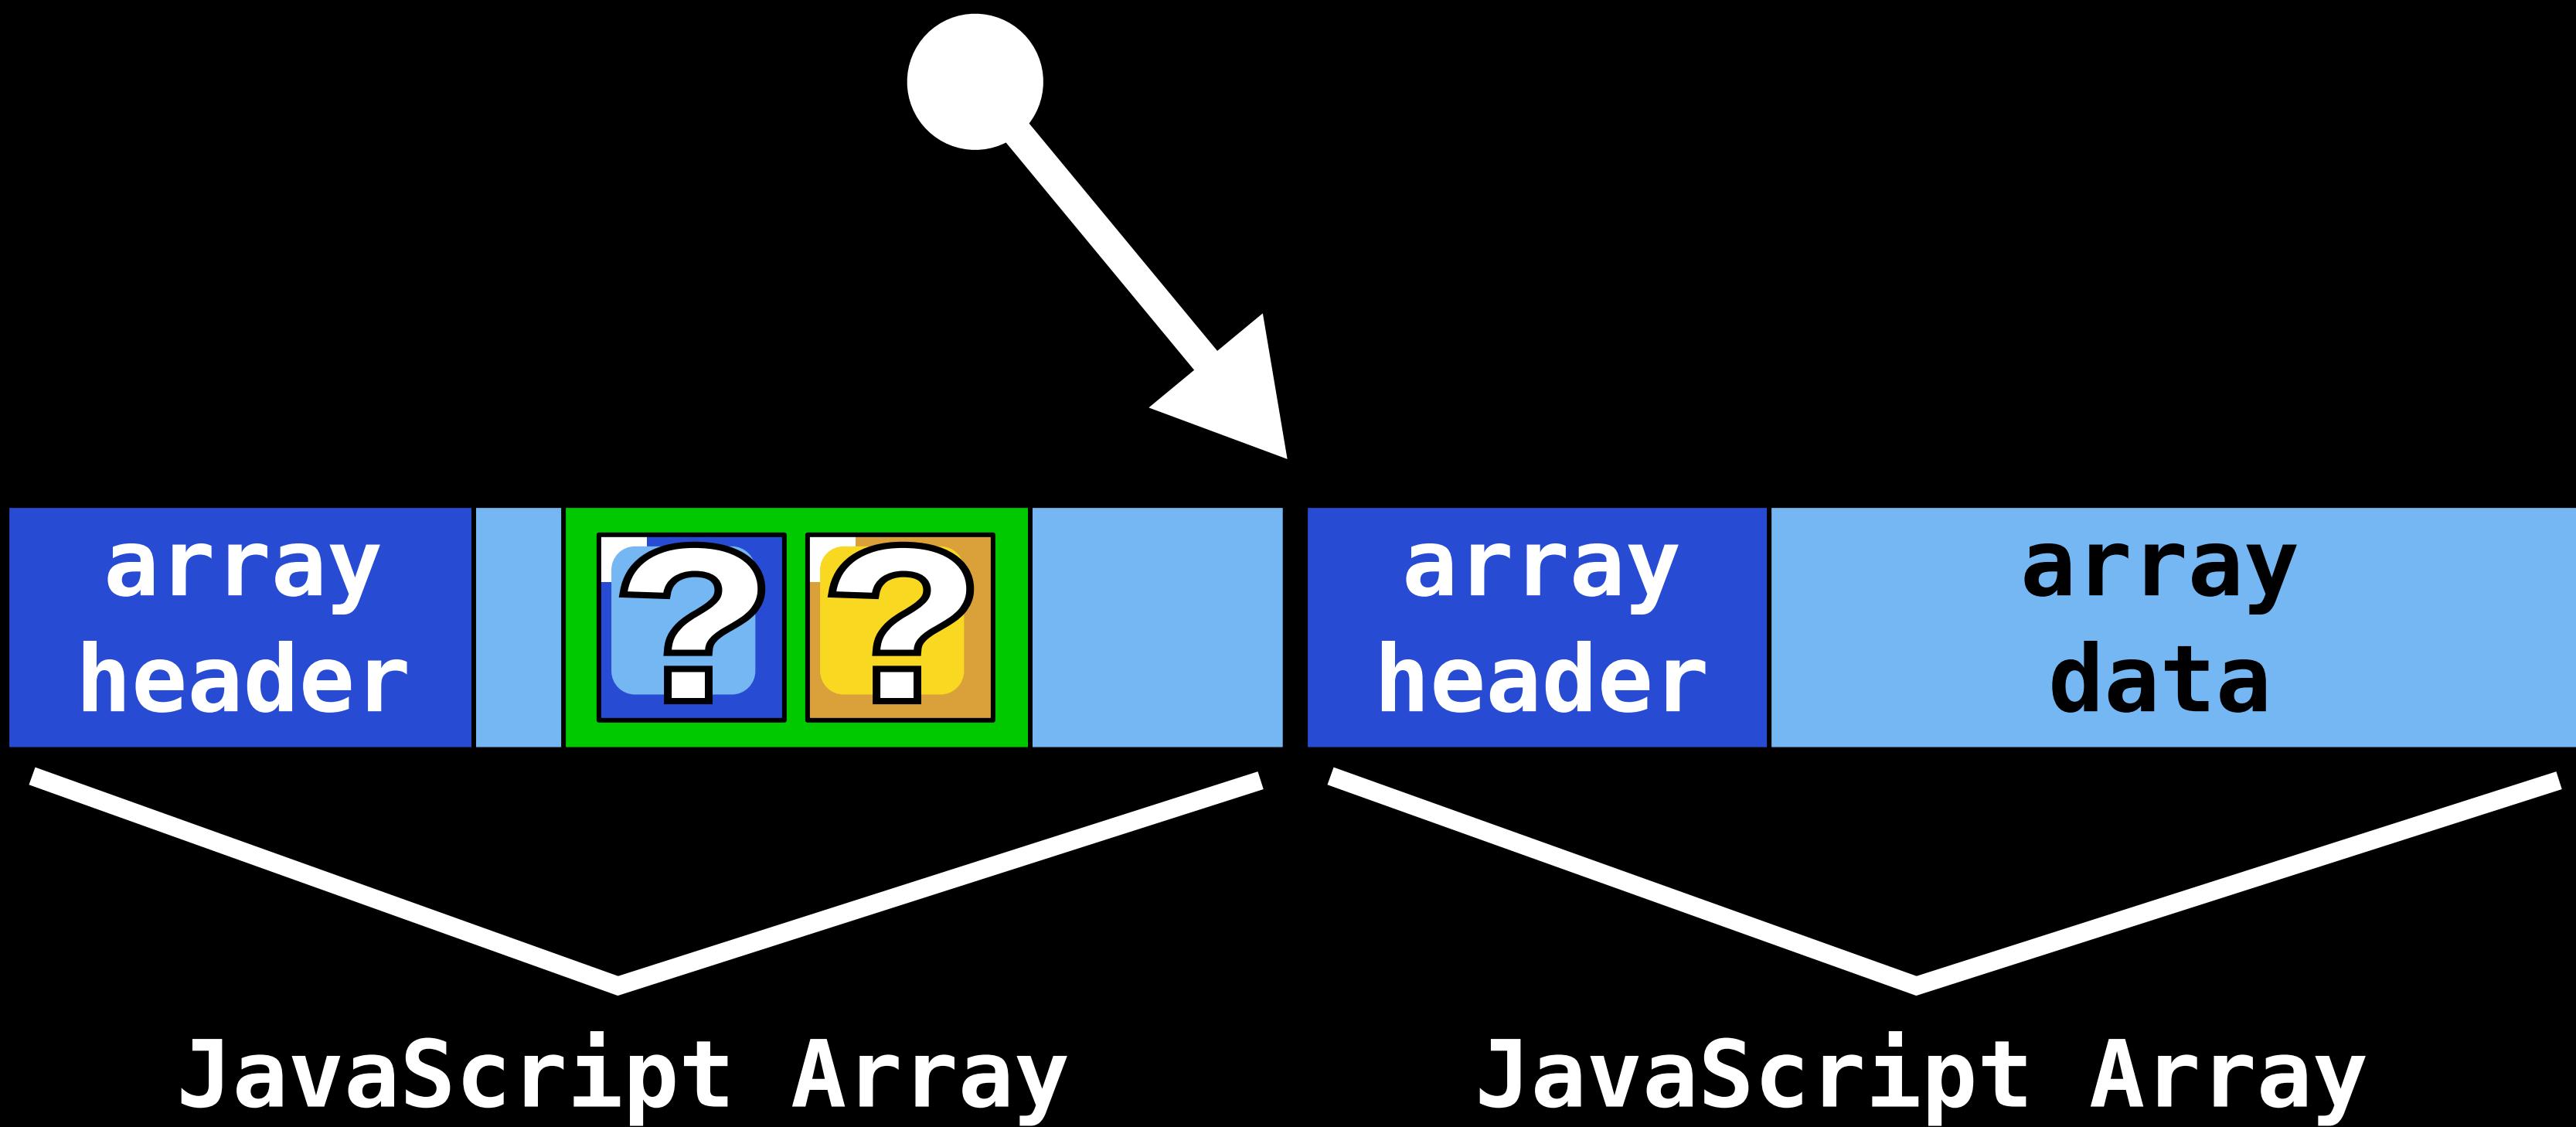


JavaScript Array

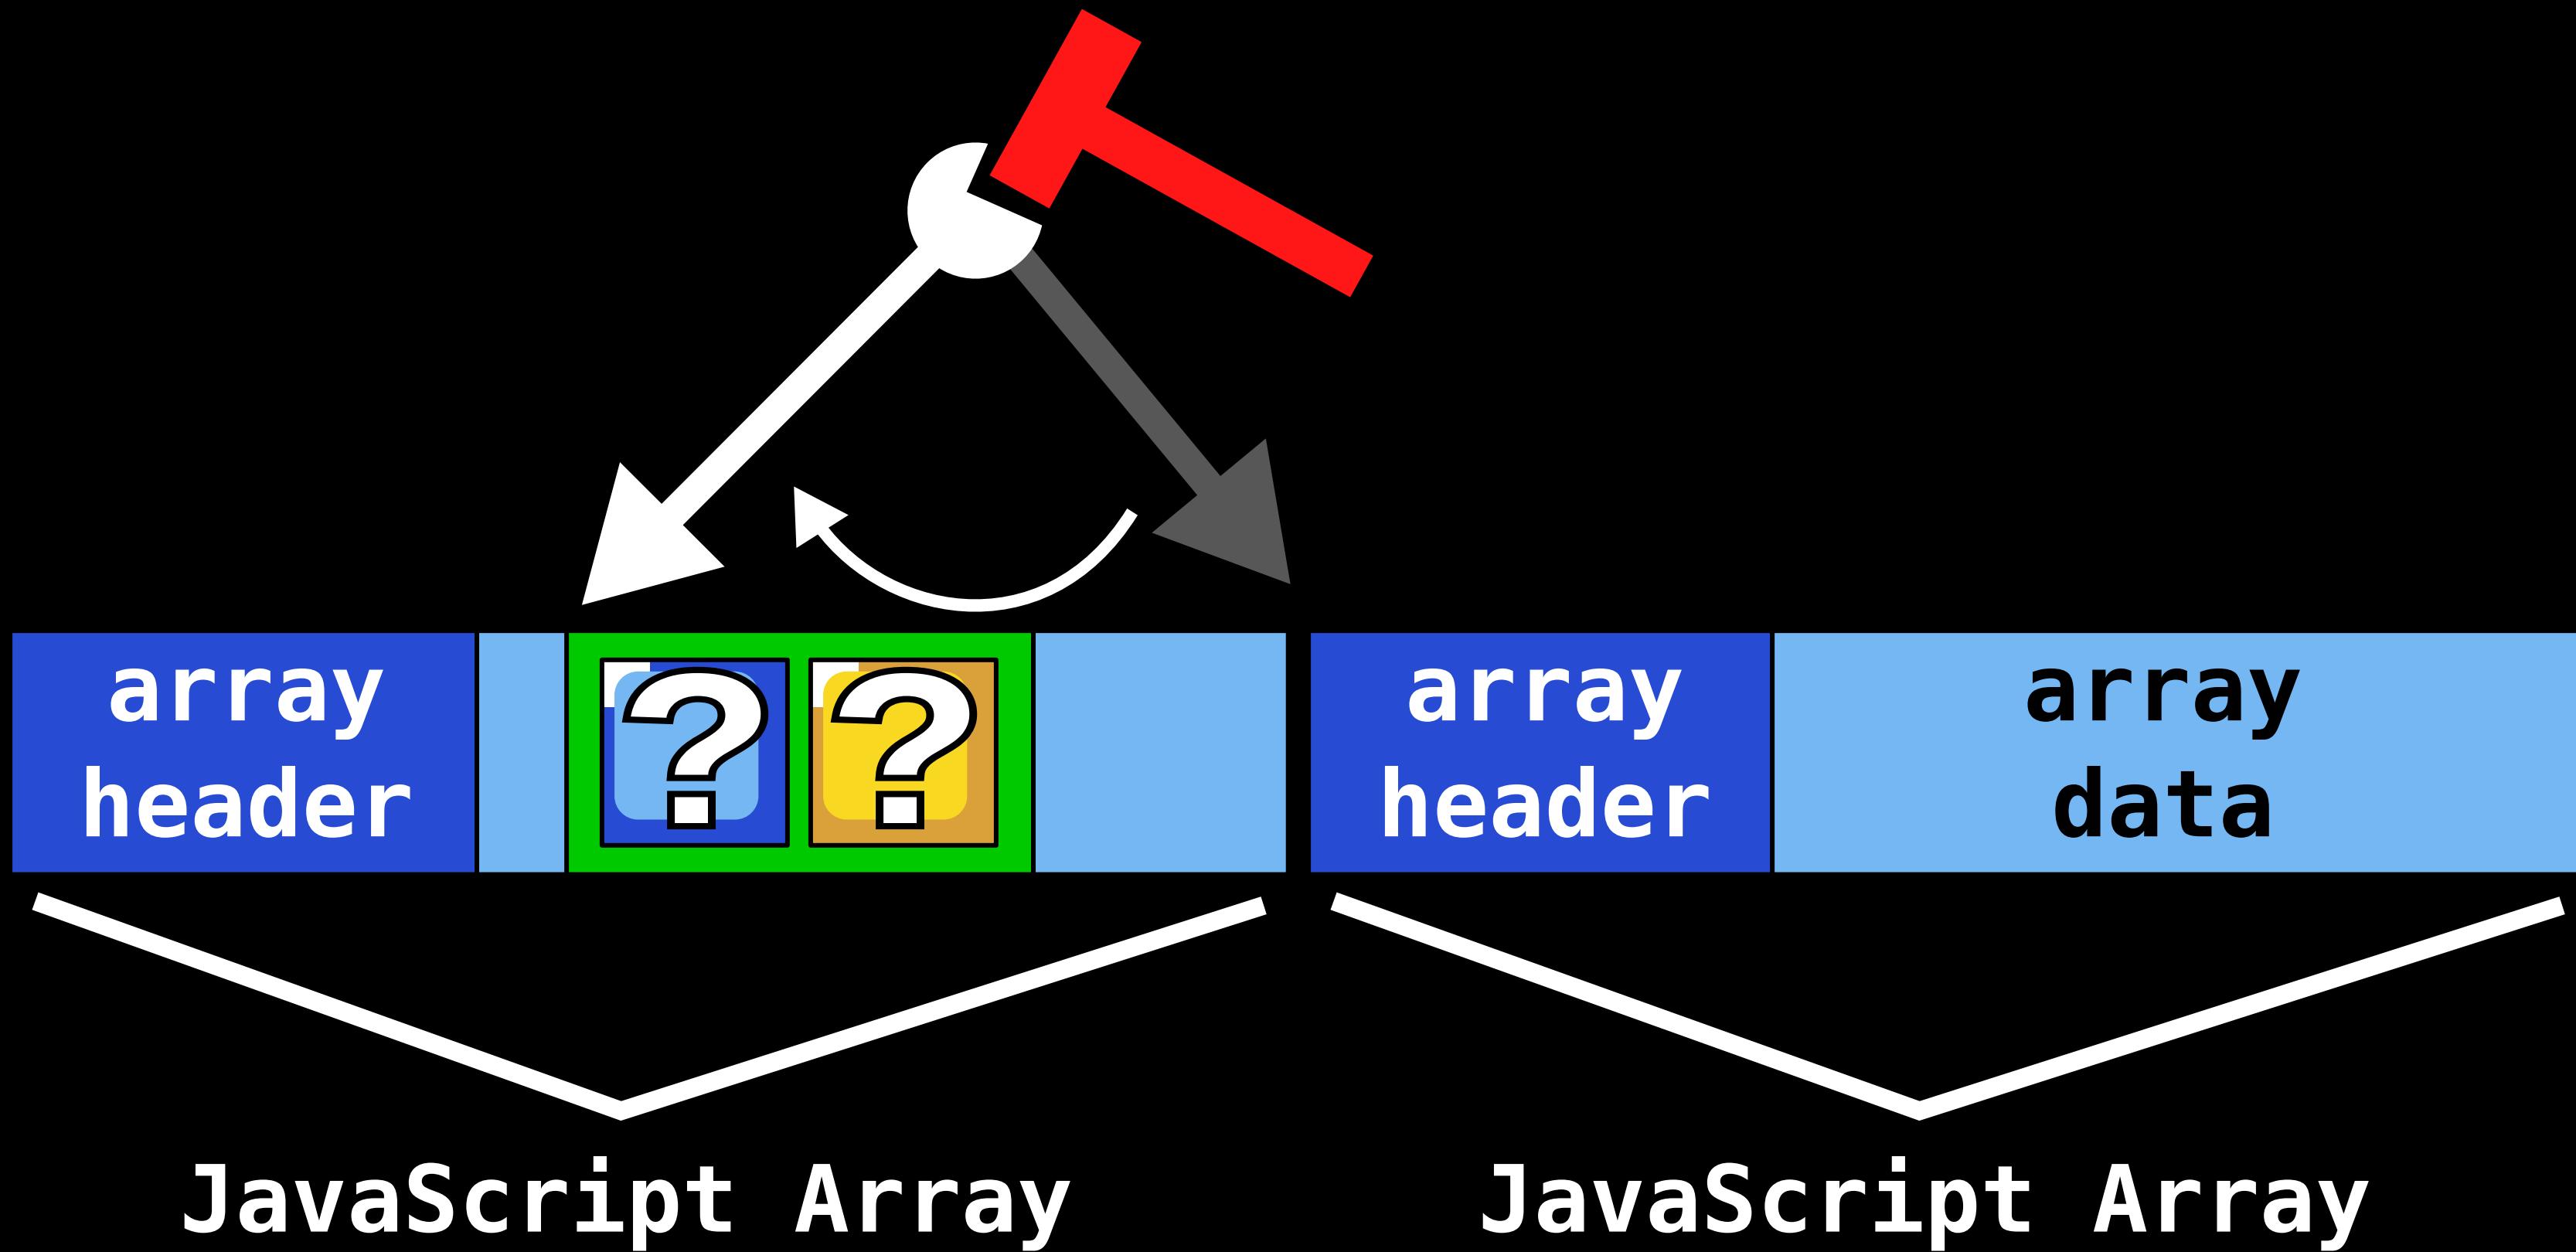
# pointer pivoting



# pointer pivoting



# pointer pivoting



# Probability of a useful 1 to 0 flips

Type flipping: 11 / 64 bits

Pointer pivottting: 12 / 64 bits

Total: 23 / 64 bits

# Dedup mitigation

- Disable memory deduplication
  - > **Disable-MMAgent -PageCombining**
- We've reported this issue to Microsoft and they have addressed this issue in ms-16-093, July 18th (CVE-2016-3272) by disabling dedup.

# **Part II: Flip Feng Shui**

# **Part II: Flip Feng Shui**

**Rowhammer  
(hardware bug)**

# **Part II: Flip Feng Shui**

**Rowhammer  
(hardware bug)**

**+**

**Deduplication  
(more than a software side-channel)**

# **Part II: Flip Feng Shui**

**Rowhammer  
(hardware bug)**

**+**

**Deduplication  
(more than a software side-channel)**



**Cross-VM compromise**

## Rowhammer bit flips:

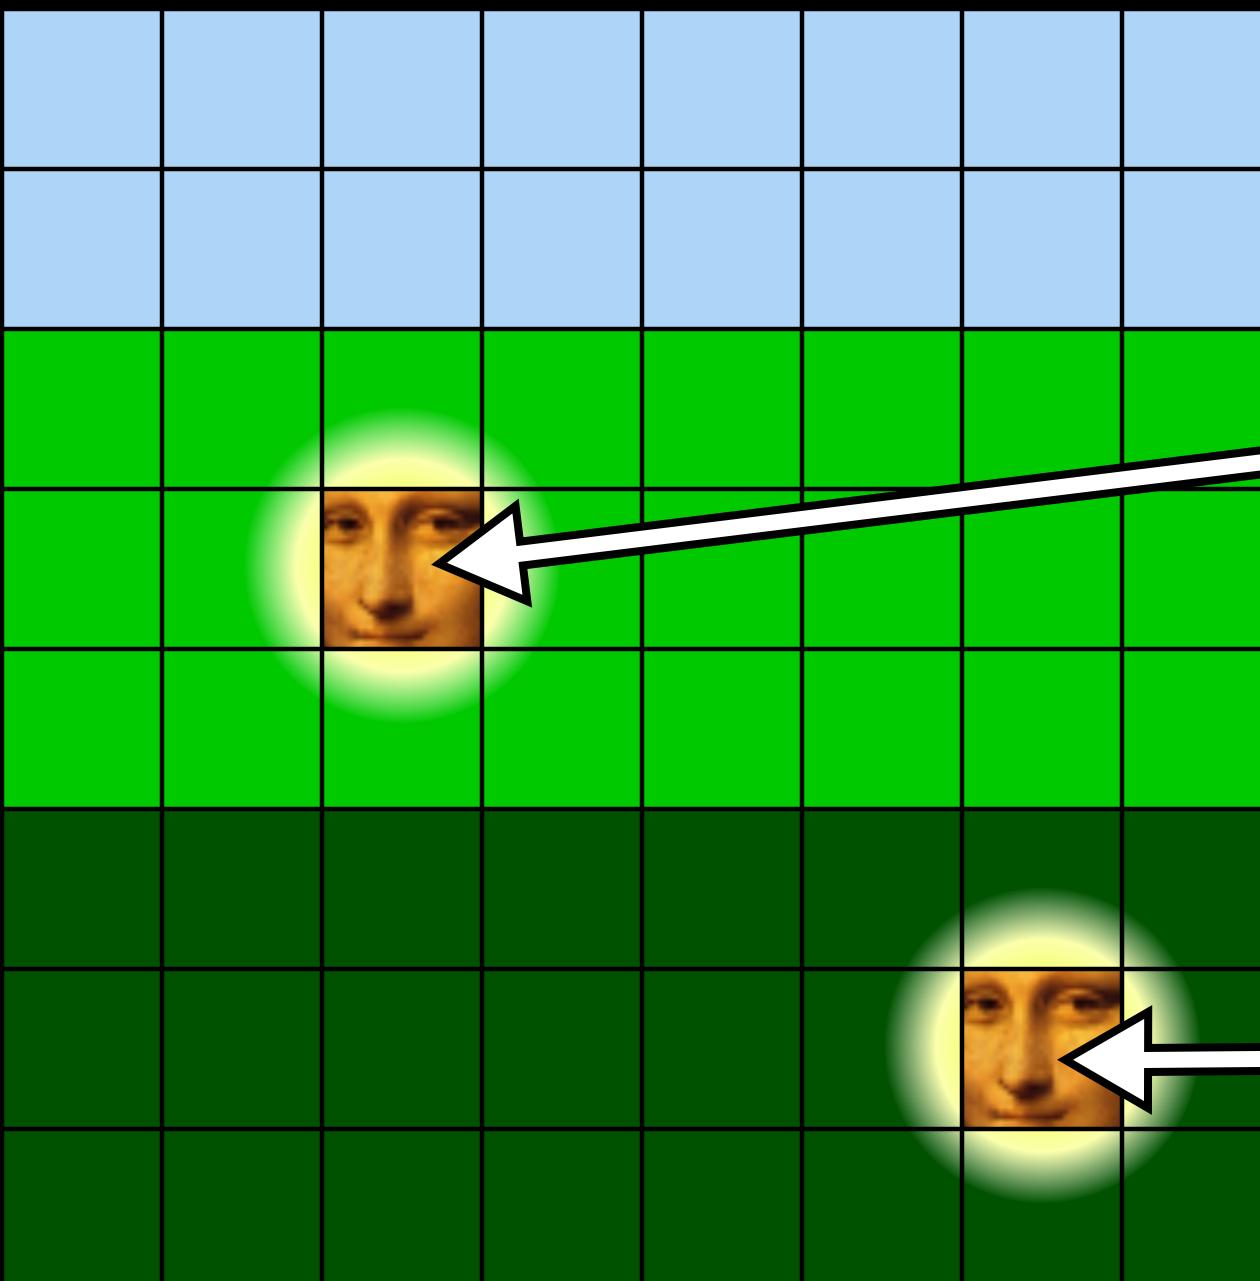
- 1) Unpredictable on which (virtual) page
- 2) Unpredictable where in the page
- 3) Repeatable once you've found a flip

# Flip Feng Shui goal:

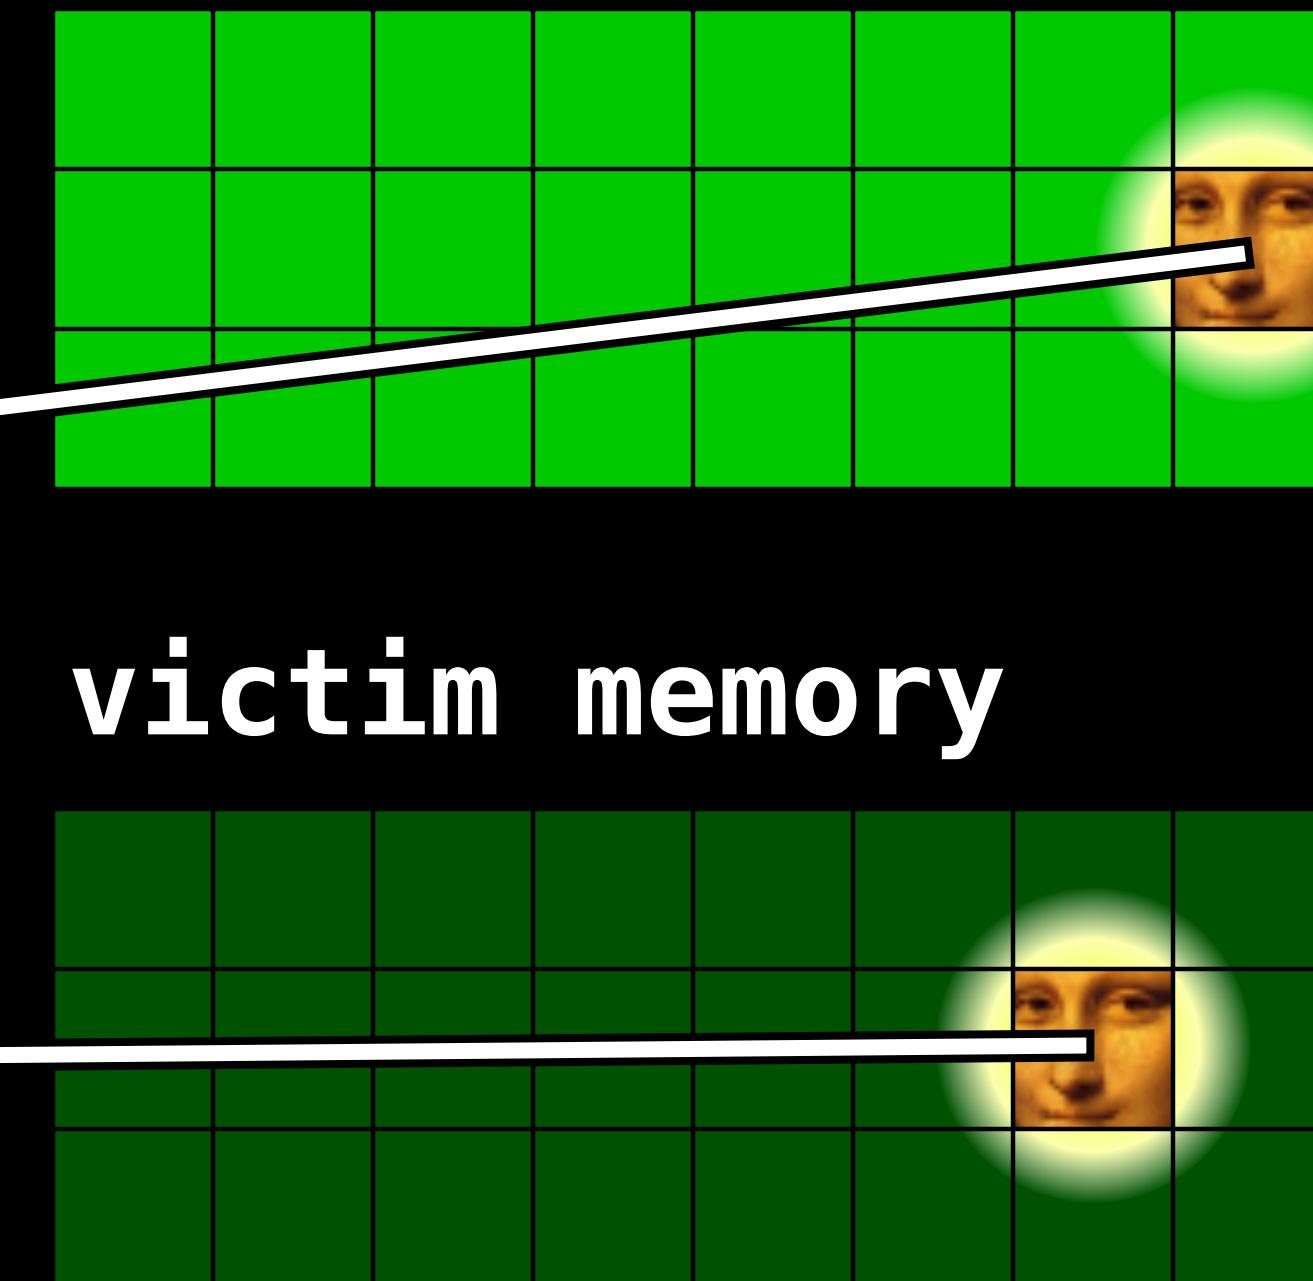
- > Find victim pages with known content which allow for exploitation when certain bits are flipped
- > Land this victim page in a physical memory location where this bit is flippable

# Deduplication implementation: Windows 10

physical memory



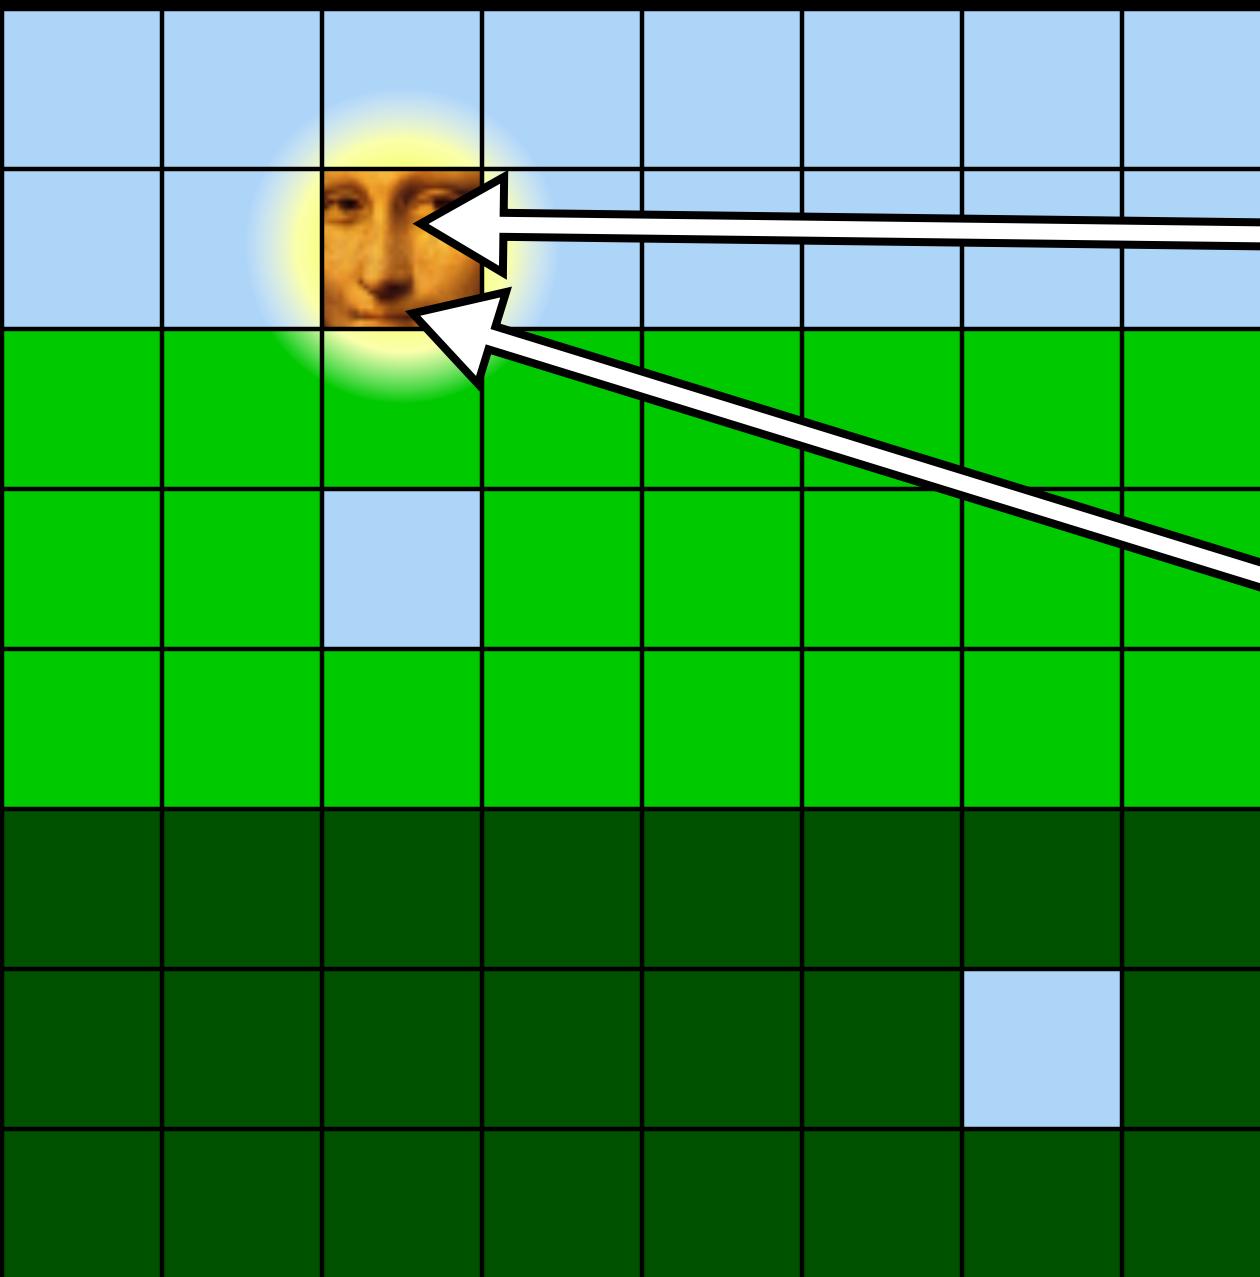
attacker memory



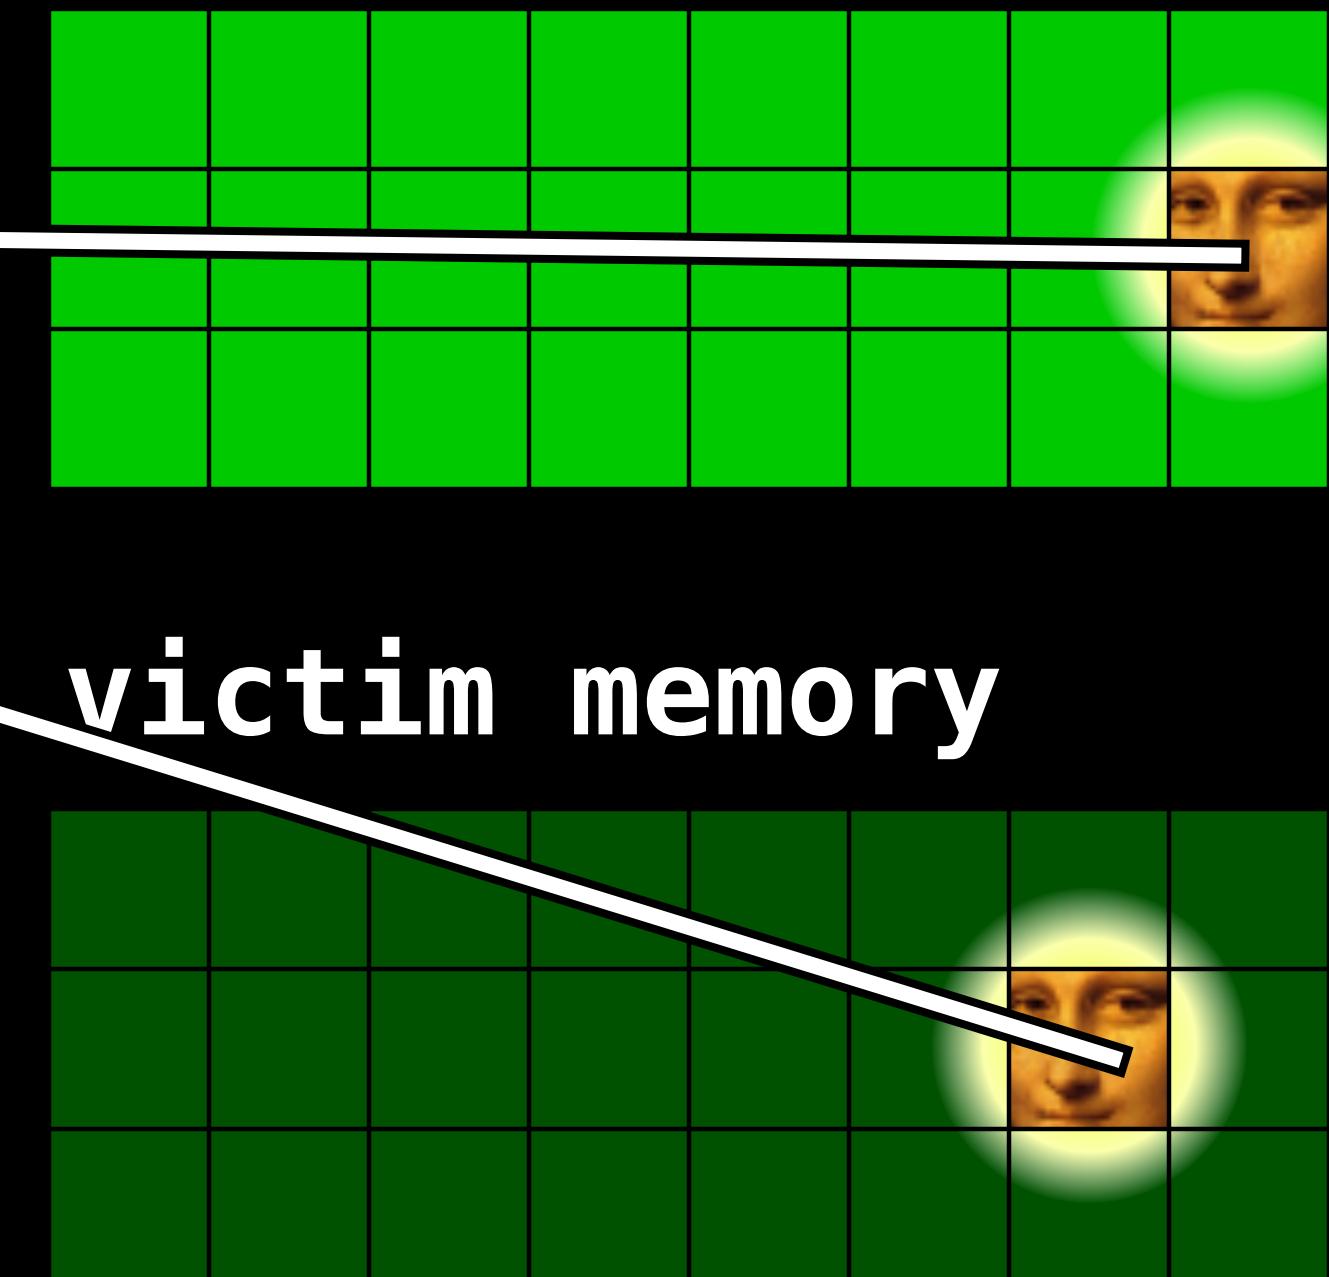
victim memory

# Deduplication implementation: Windows 10

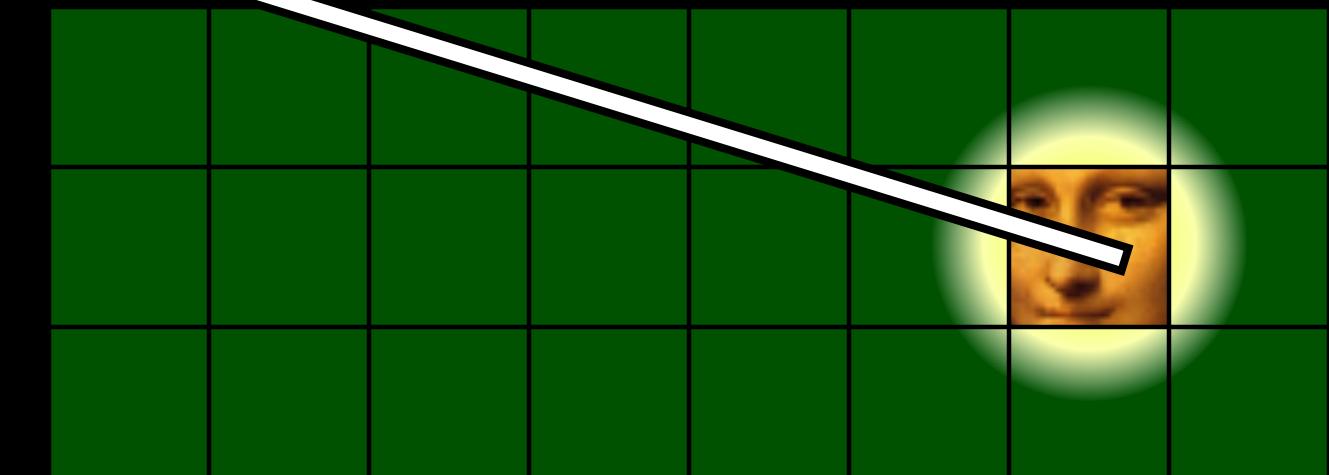
physical memory



attacker memory

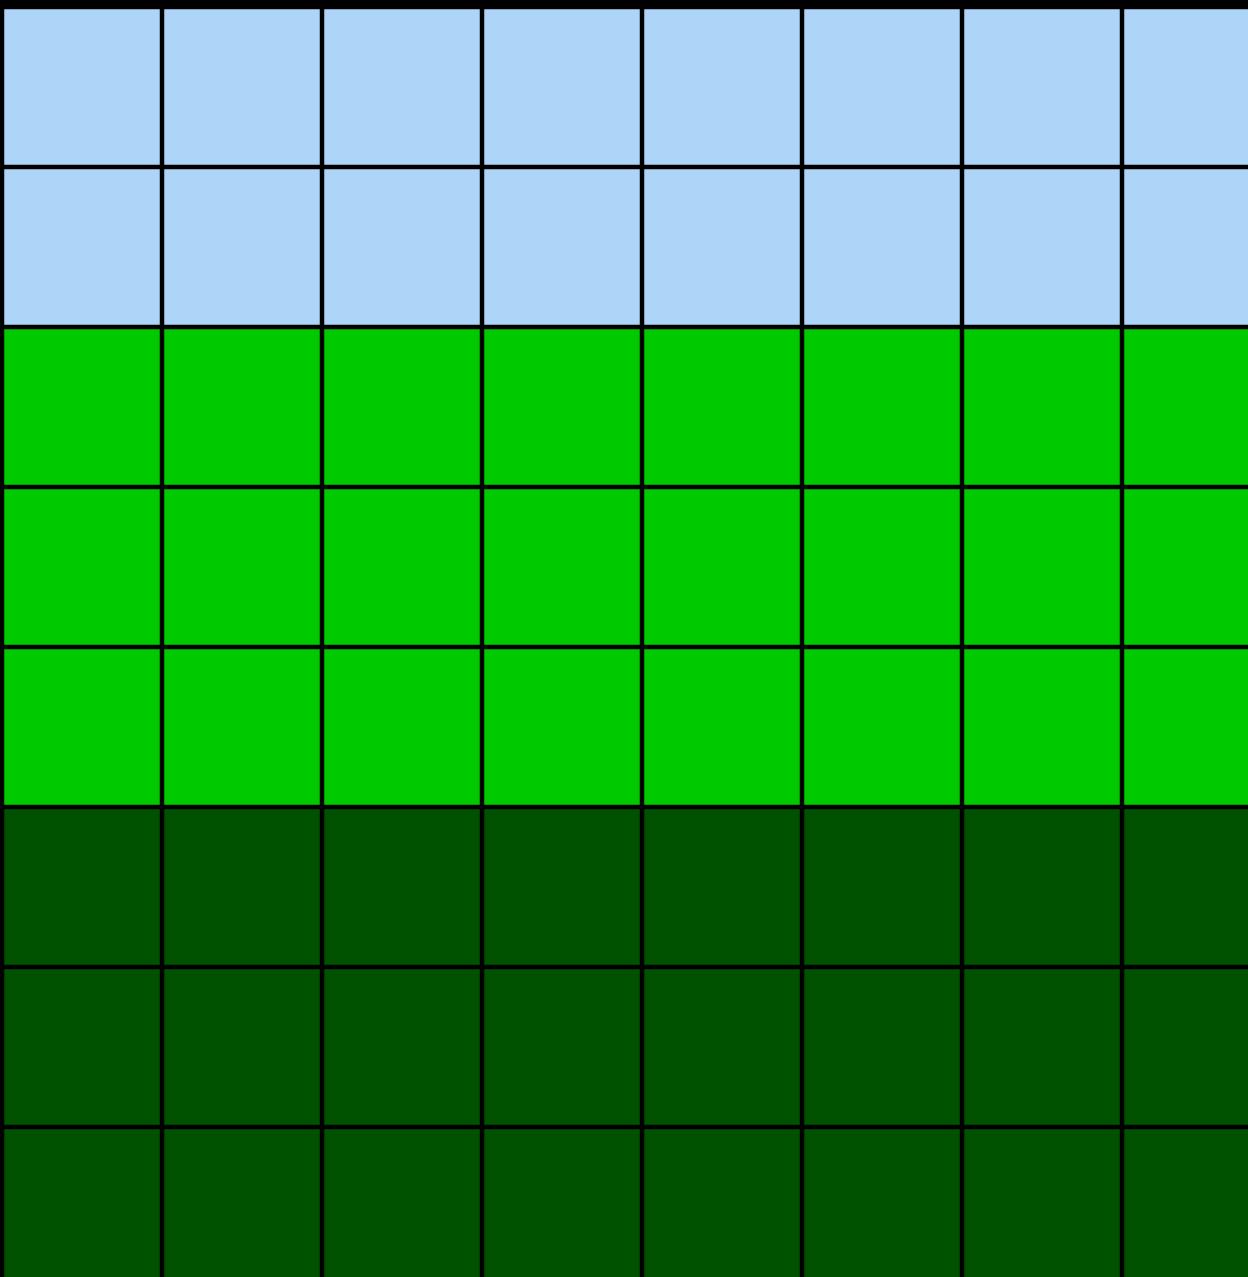


victim memory

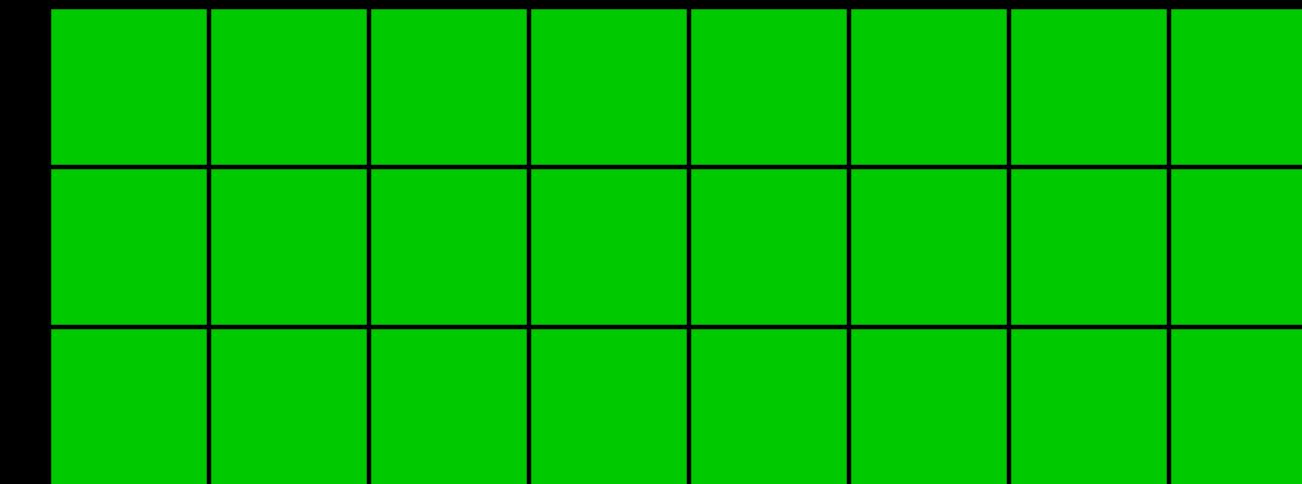


# Deduplication implementation: KVM on Linux (KSM)

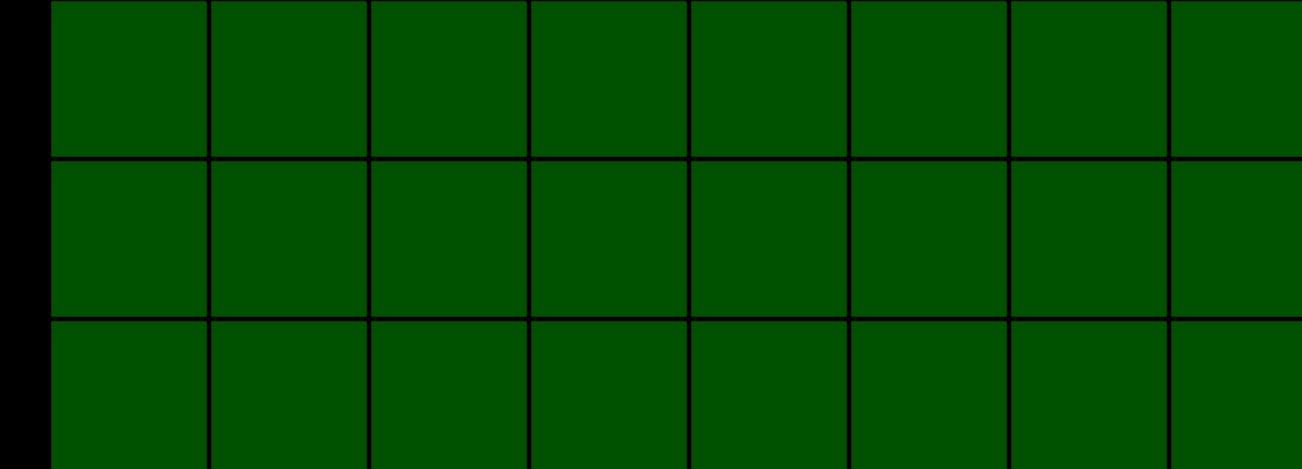
physical memory



attacker memory

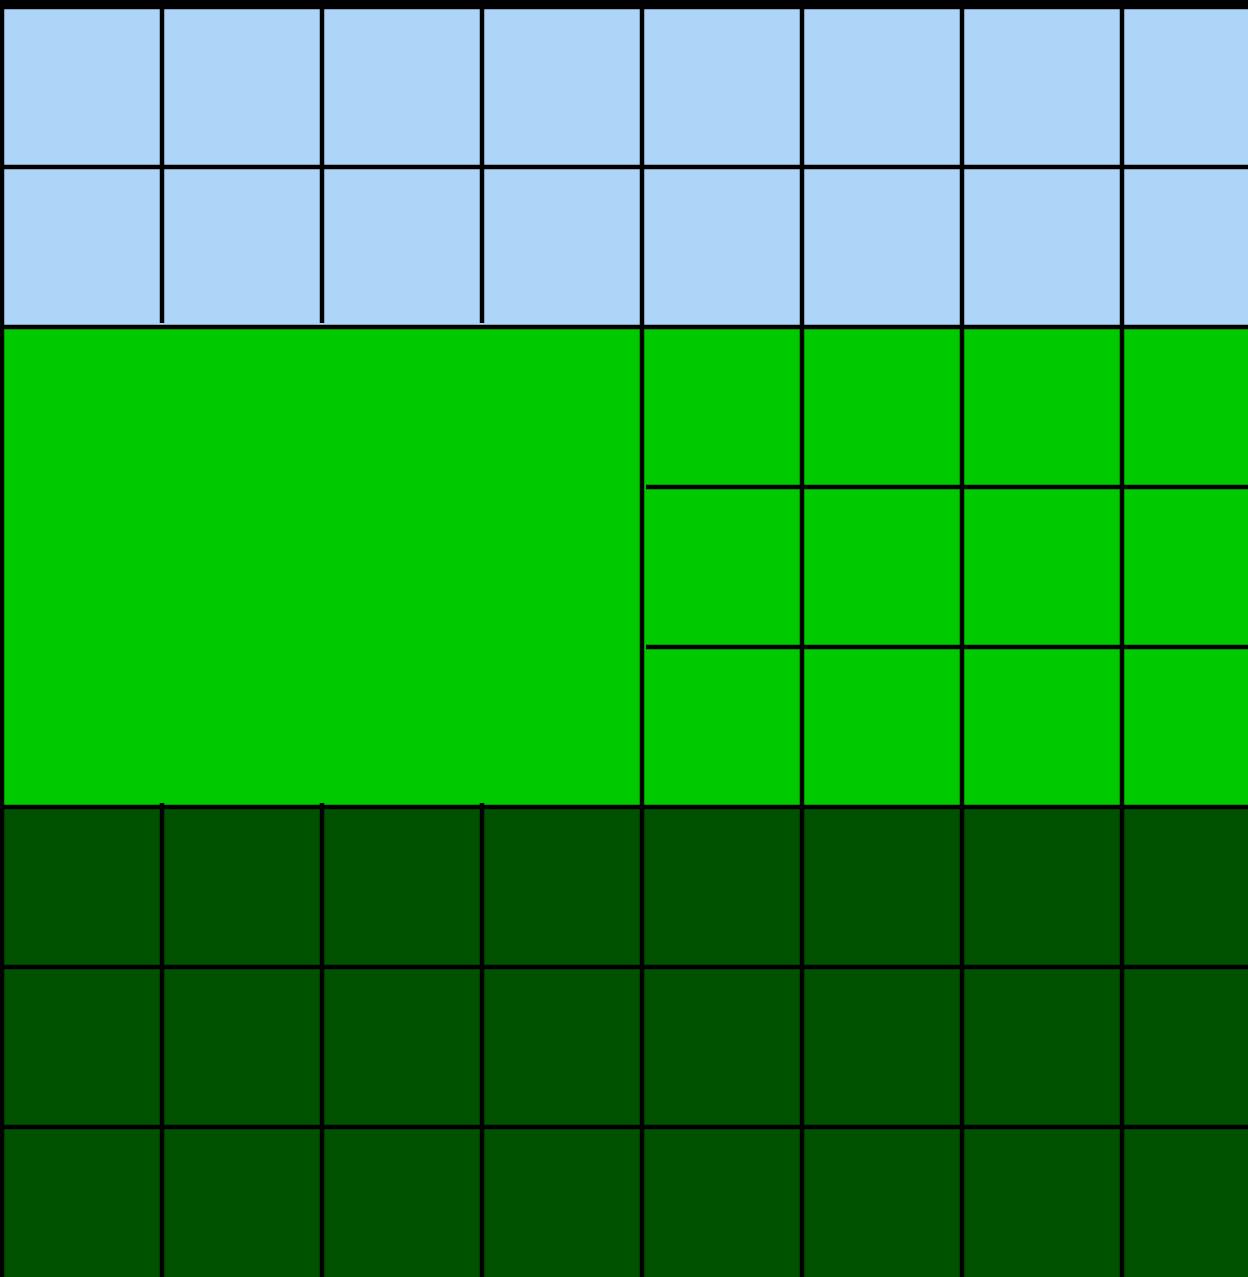


victim memory

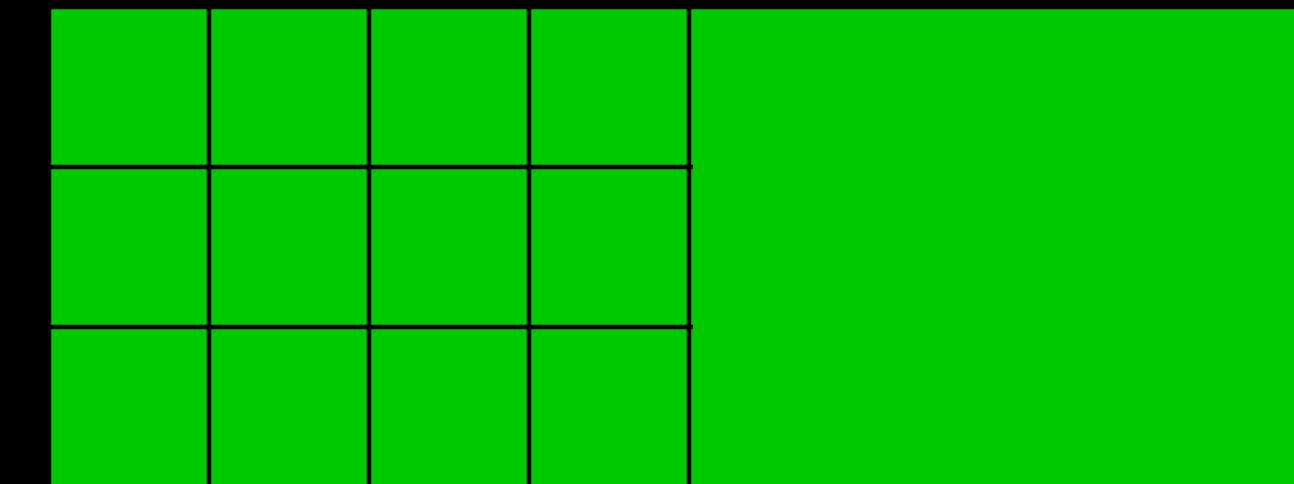


# Deduplication implementation: KVM on Linux (KSM)

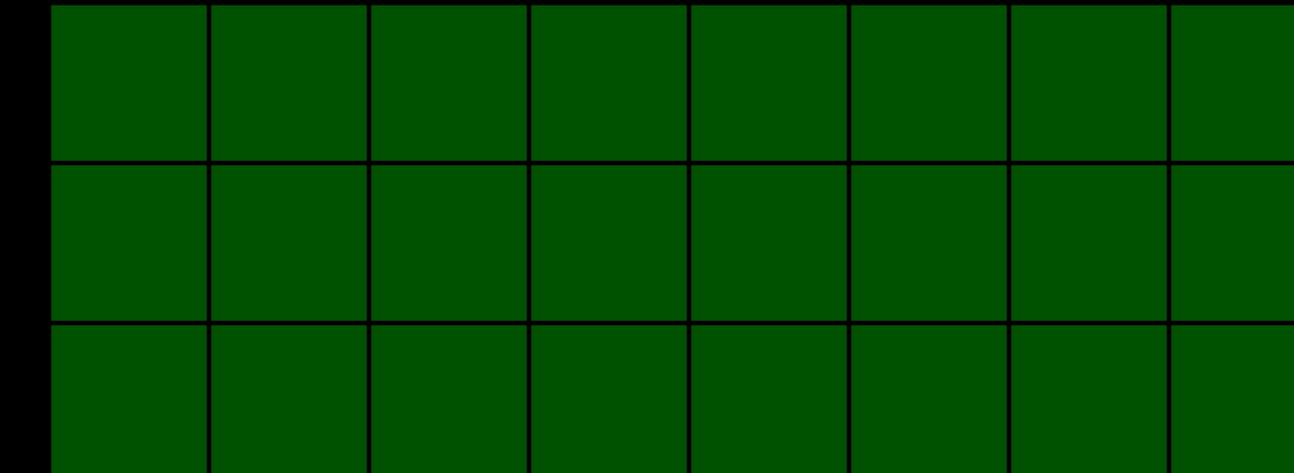
physical memory



attacker memory

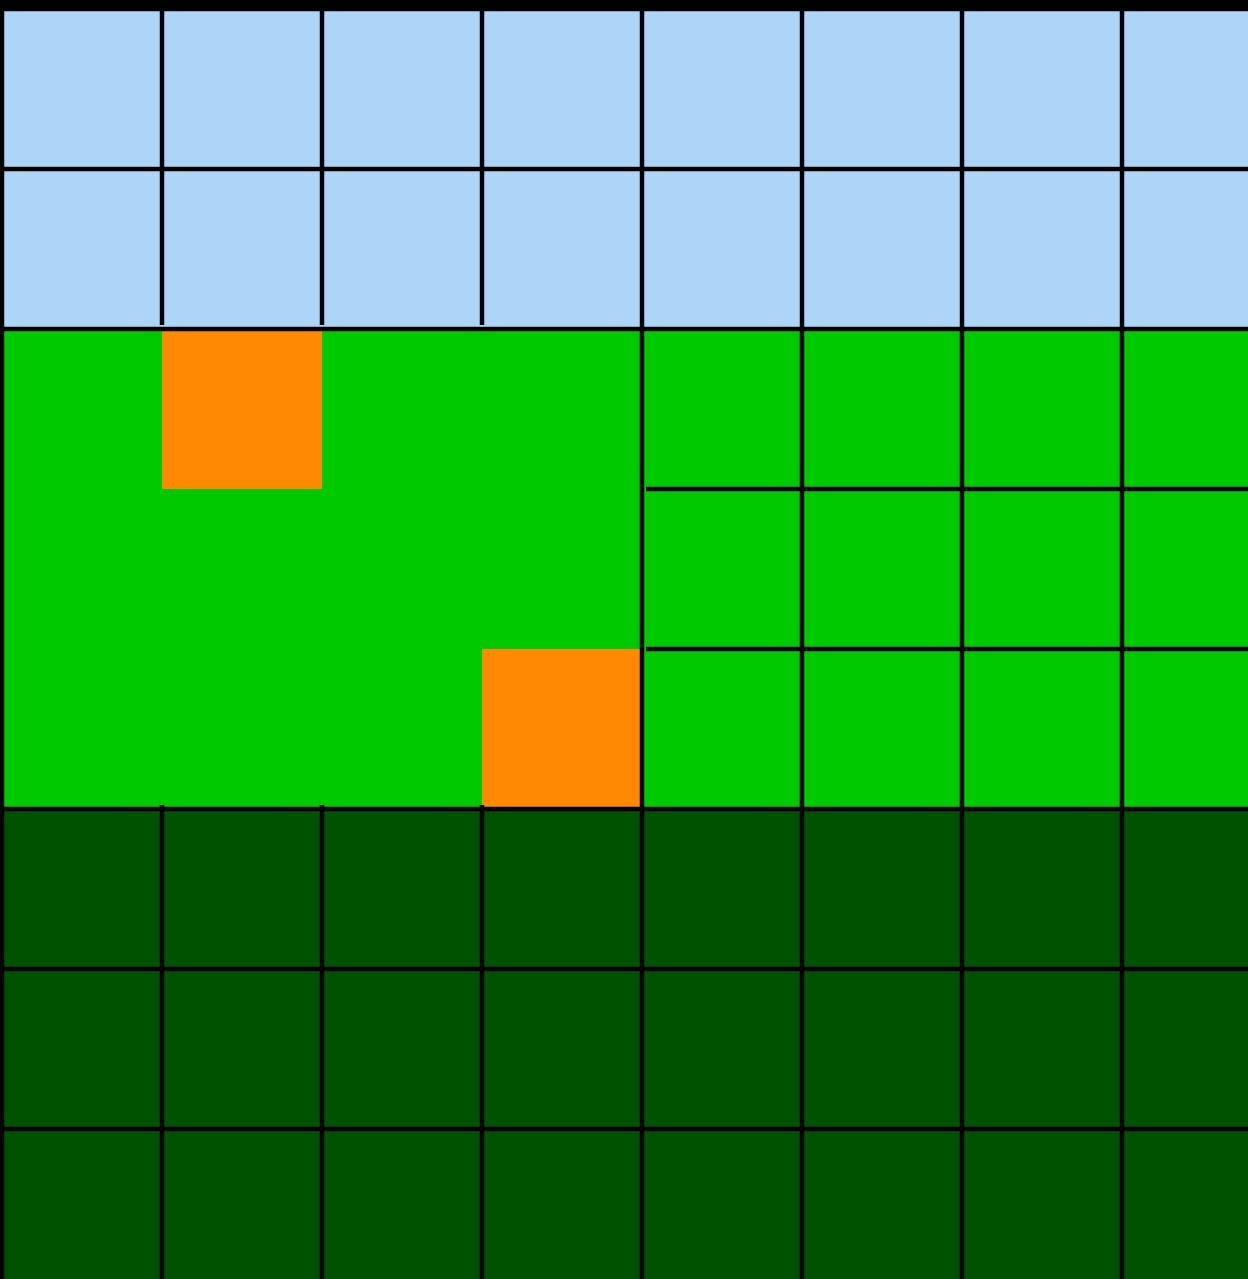


victim memory

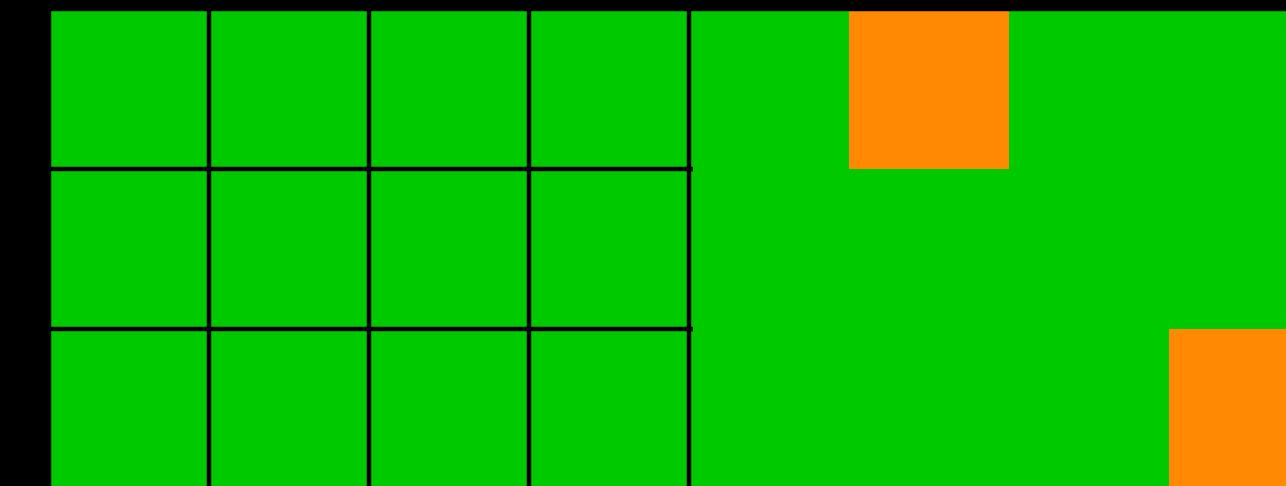


# Deduplication implementation: KVM on Linux (KSM)

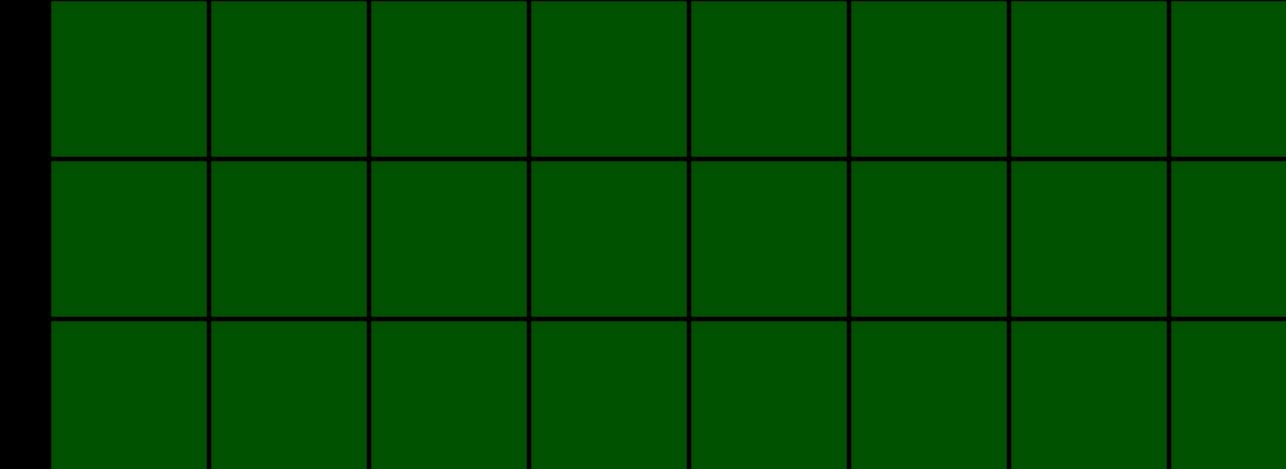
physical memory



attacker memory

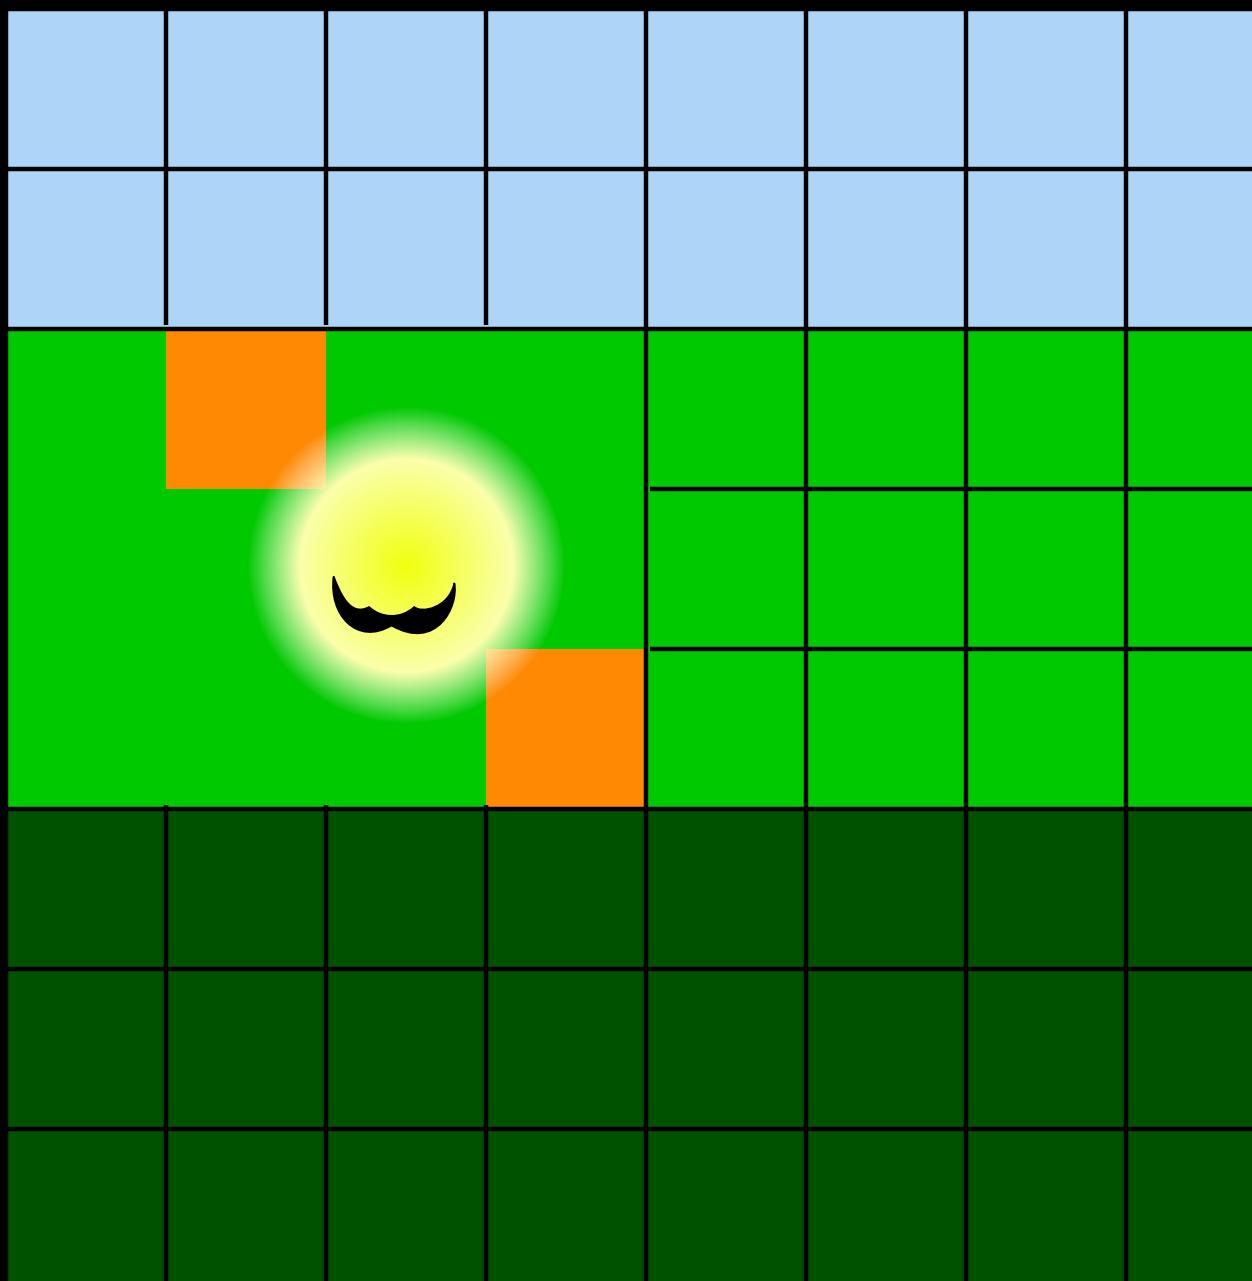


victim memory

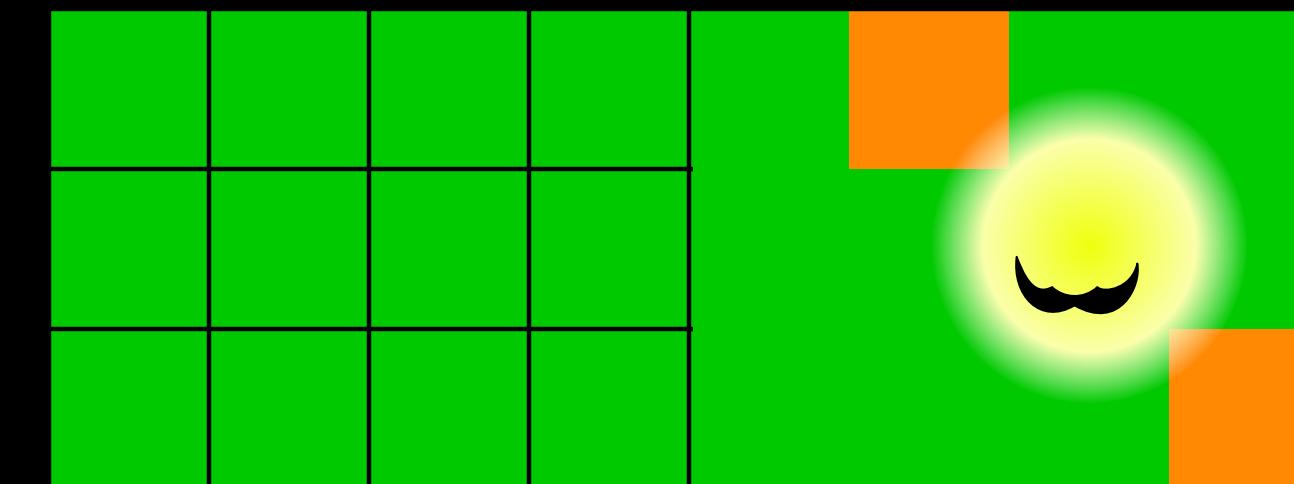


# Deduplication implementation: KVM on Linux (KSM)

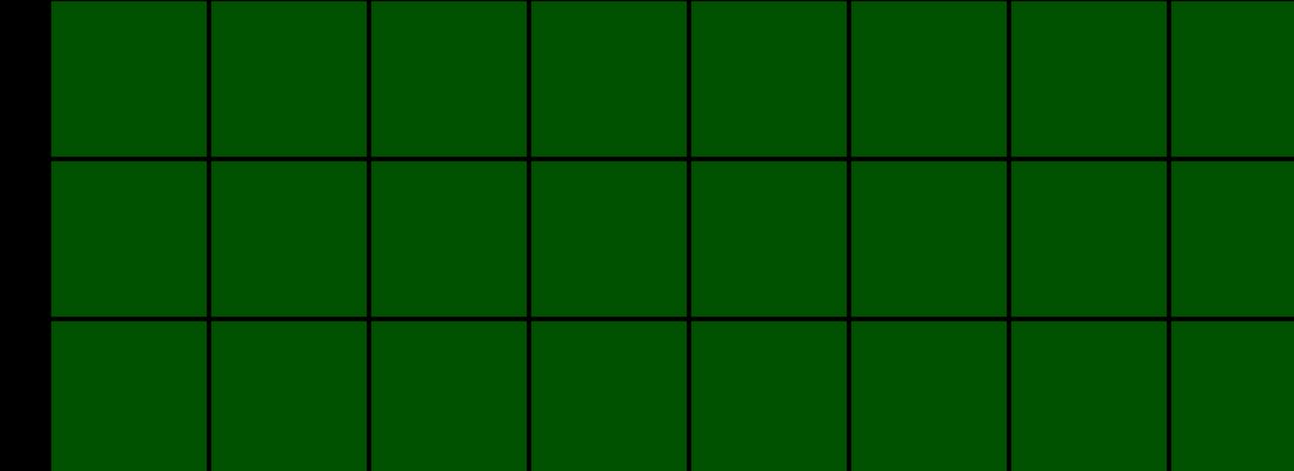
physical memory



attacker memory

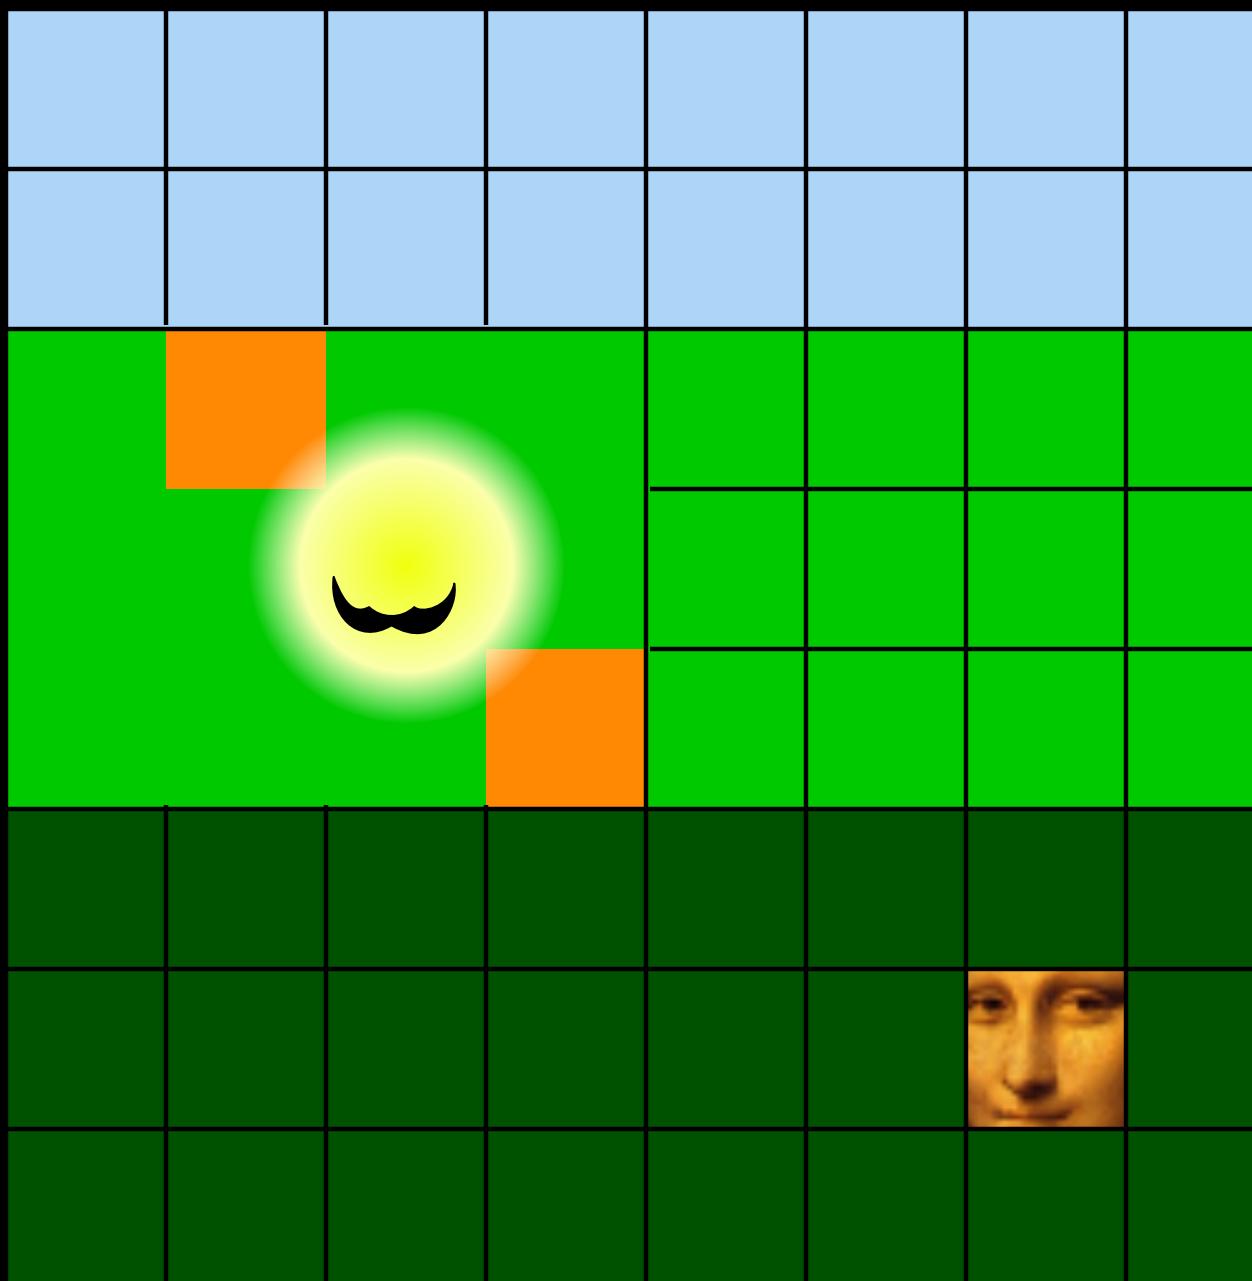


victim memory

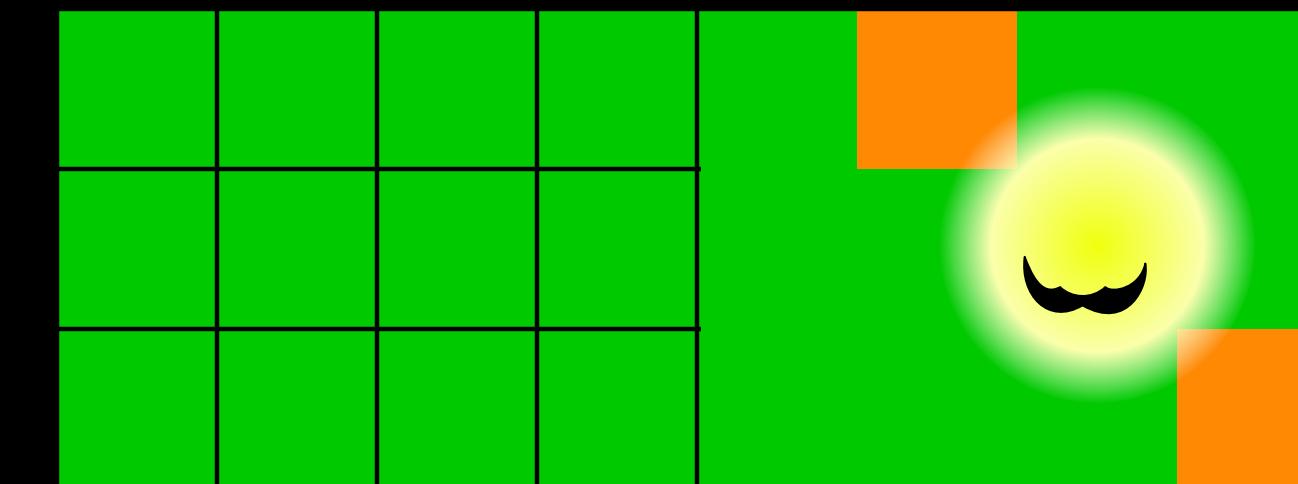


# Deduplication implementation: KVM on Linux (KSM)

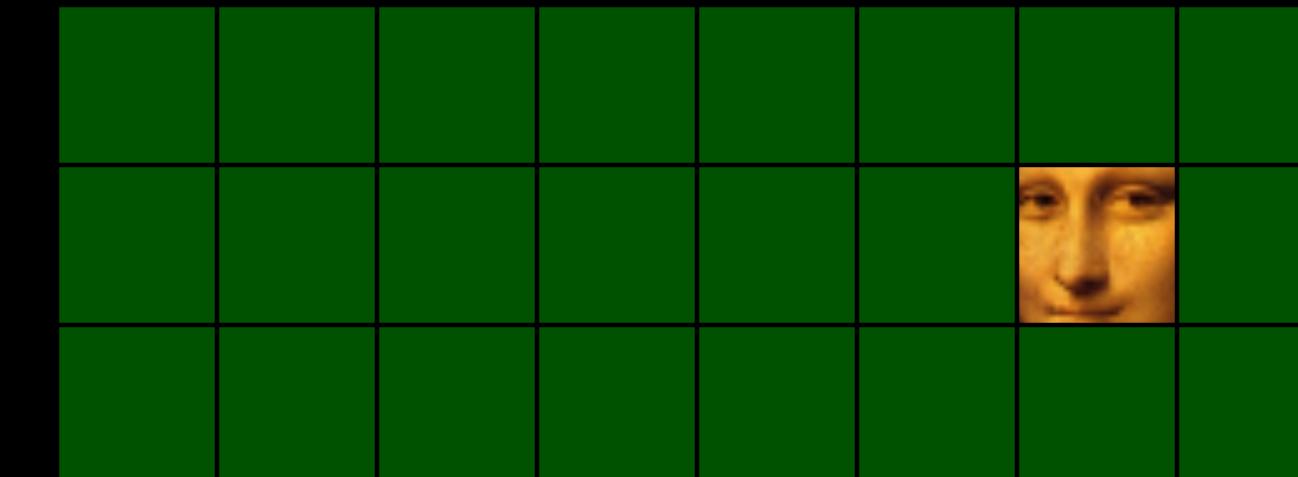
physical memory



attacker memory

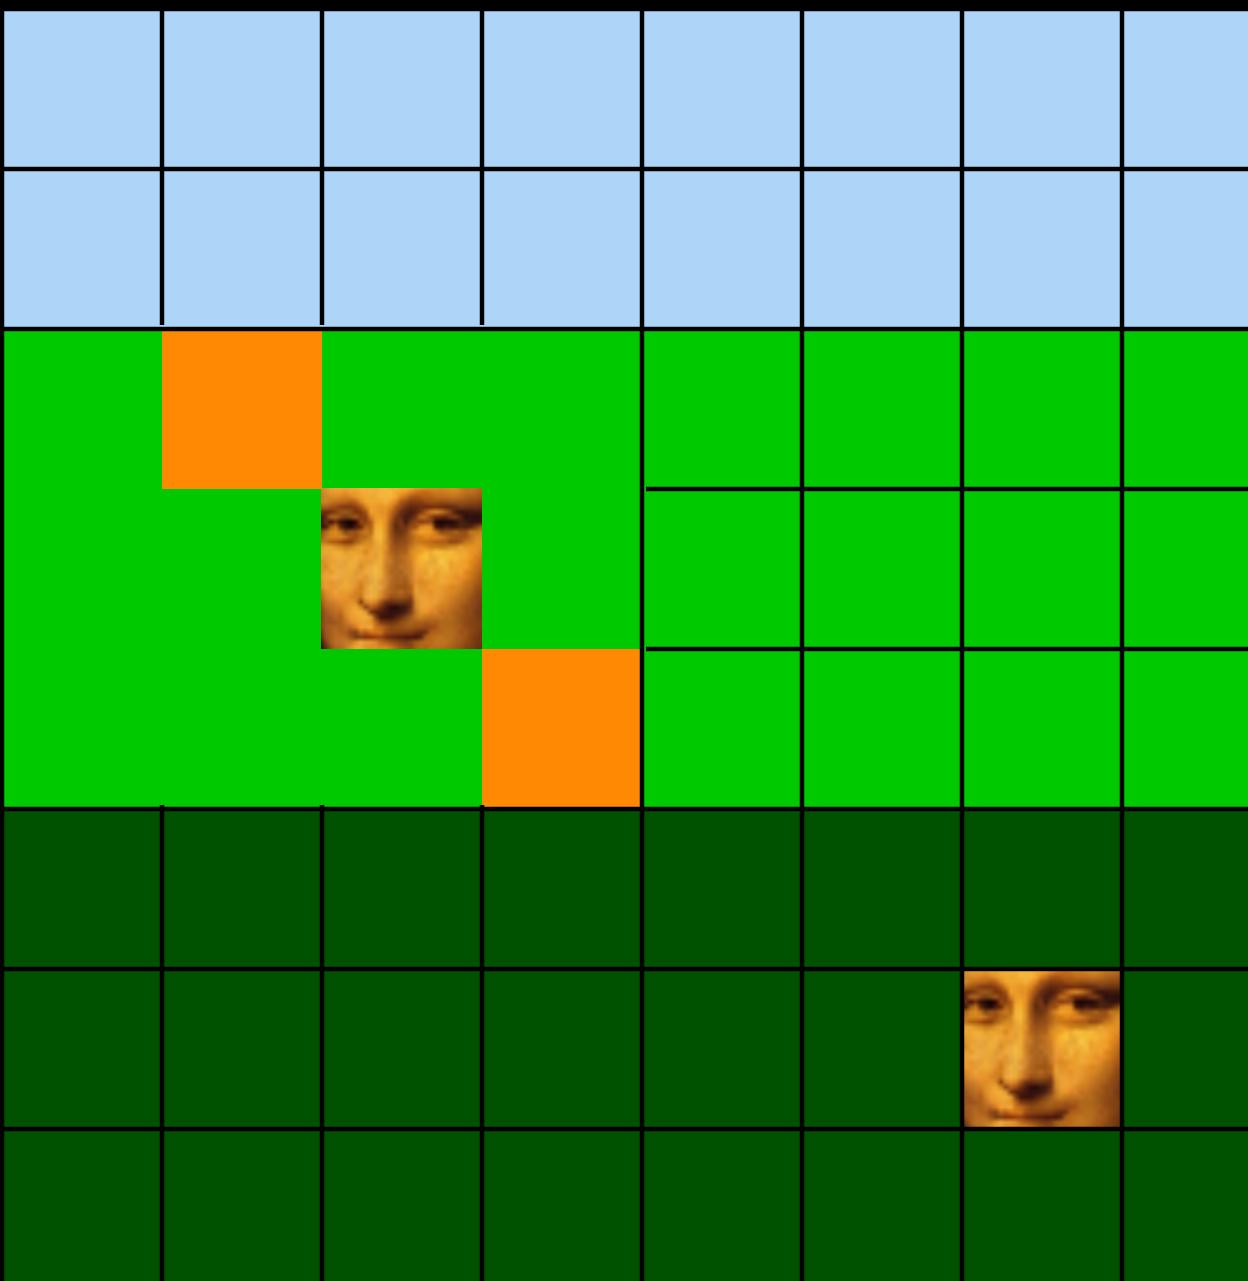


victim memory

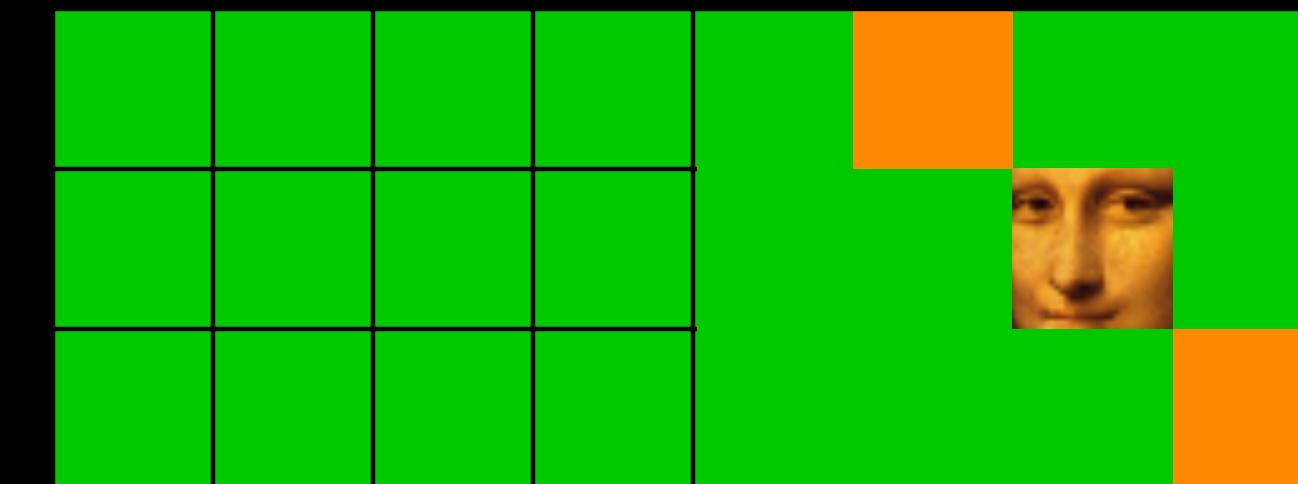


# Deduplication implementation: KVM on Linux (KSM)

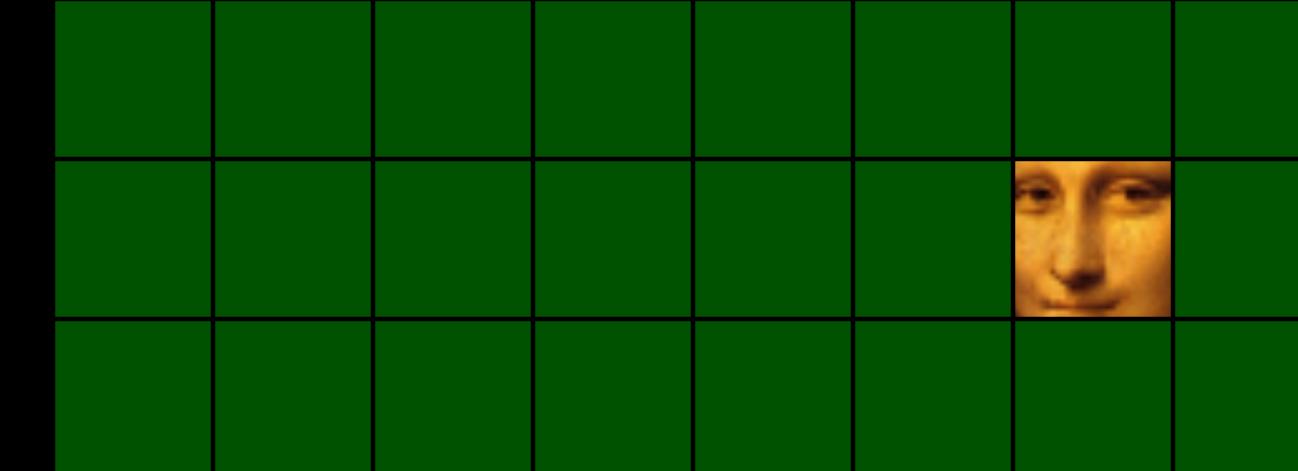
physical memory



attacker memory

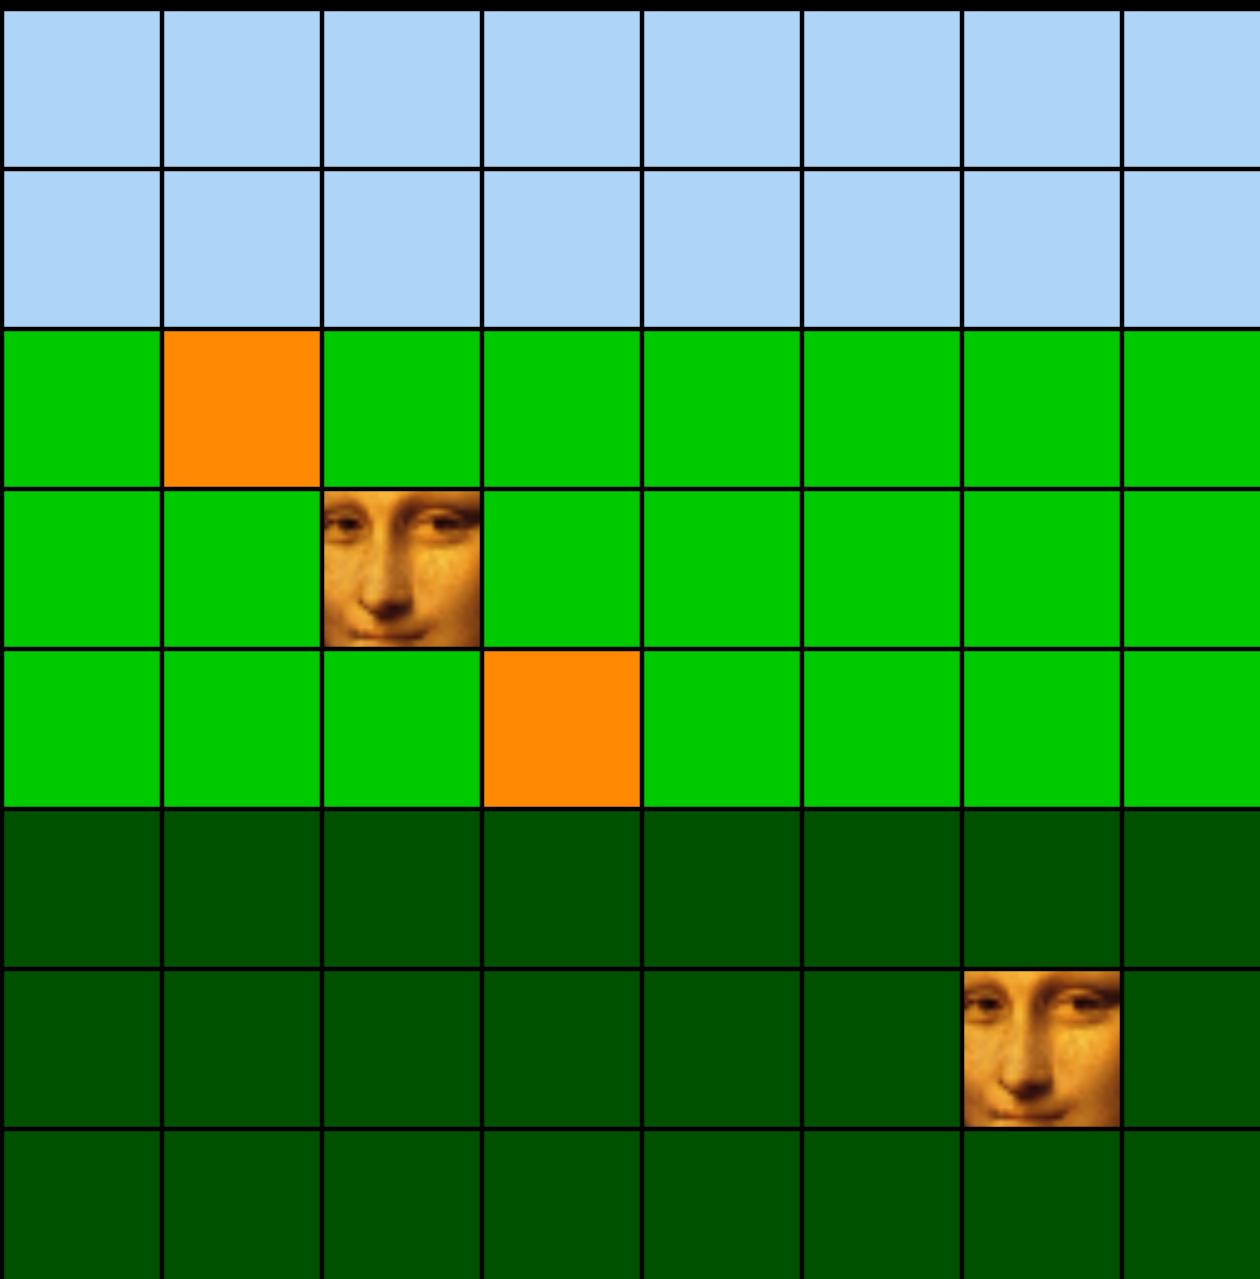


victim memory

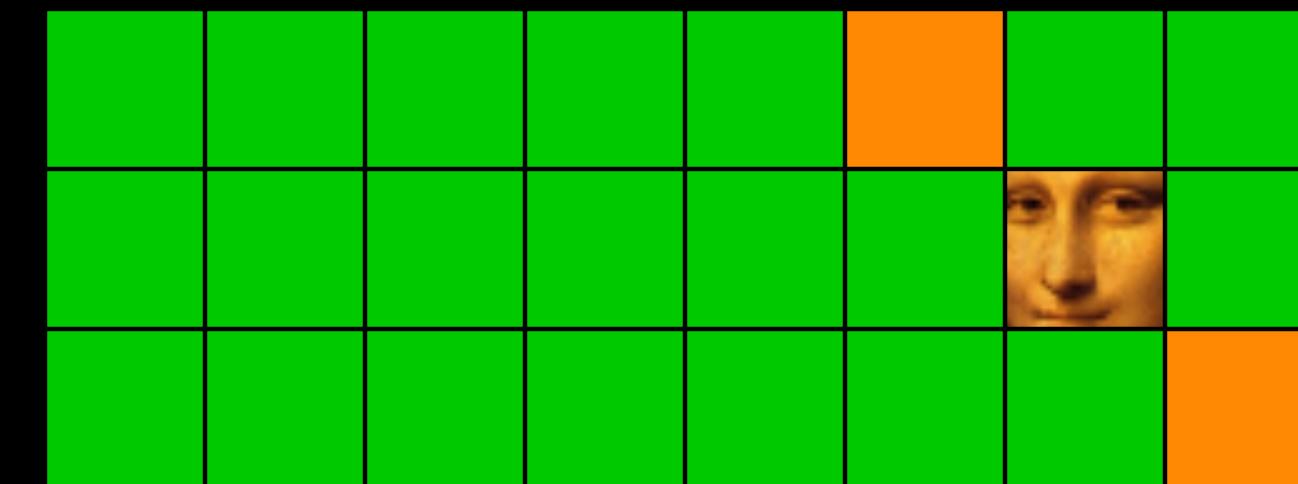


# Deduplication implementation: KVM on Linux (KSM)

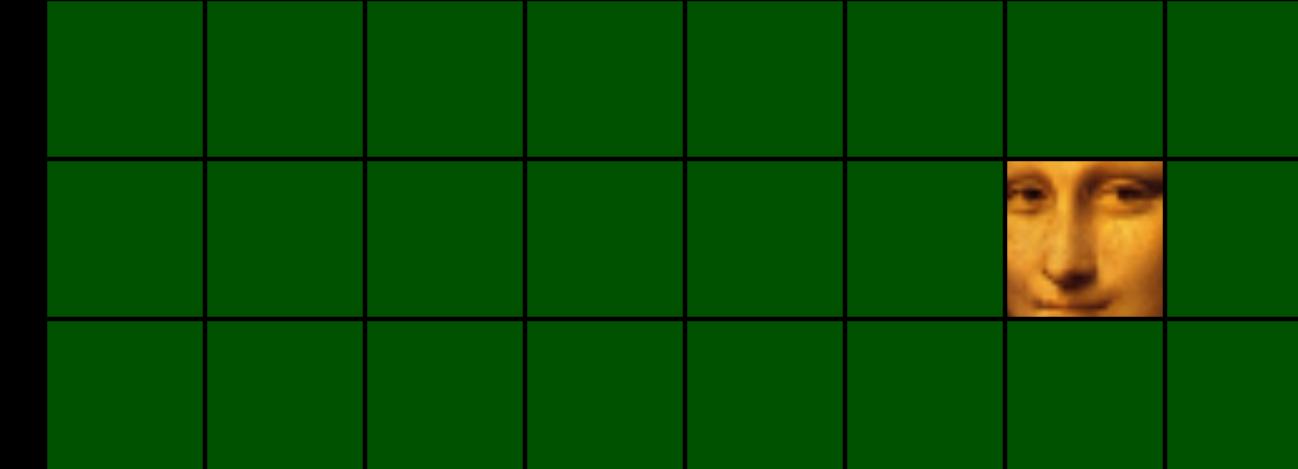
physical memory



attacker memory

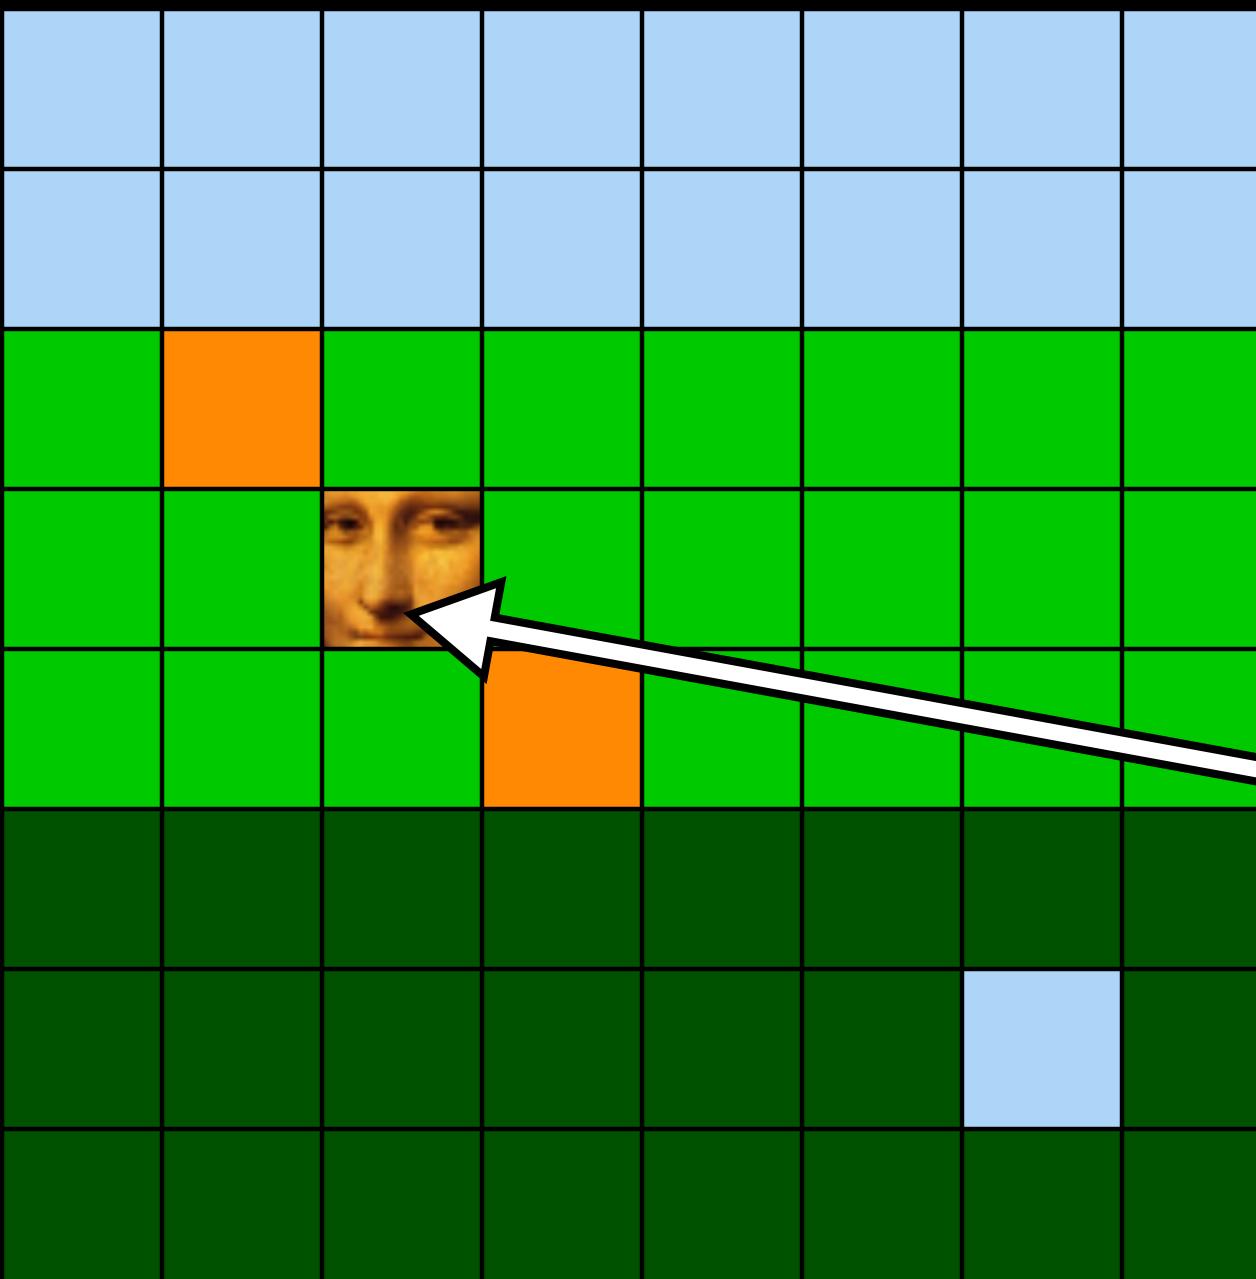


victim memory

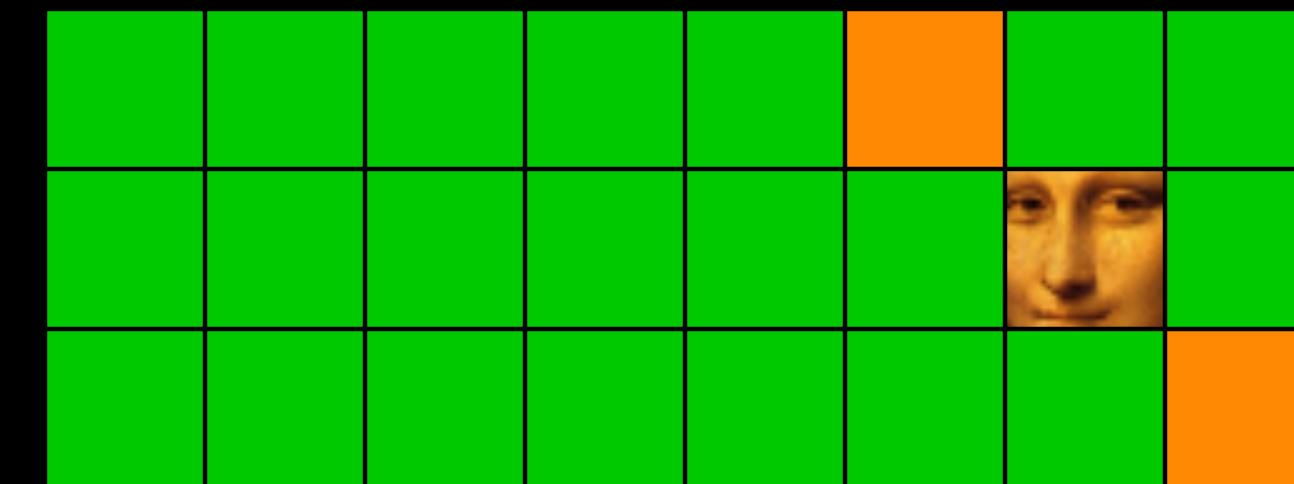


# Deduplication implementation: KVM on Linux (KSM)

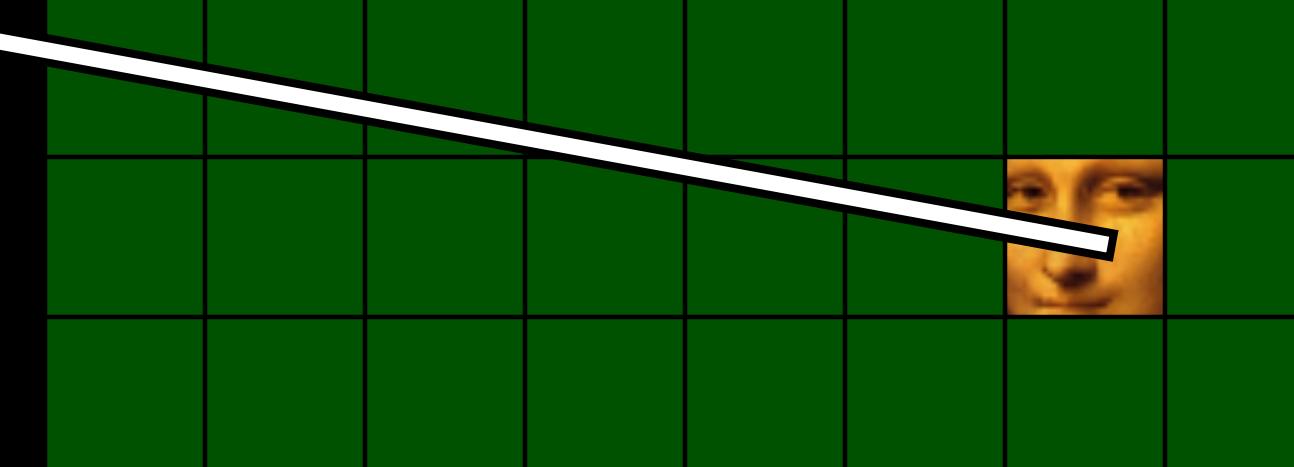
physical memory



attacker memory

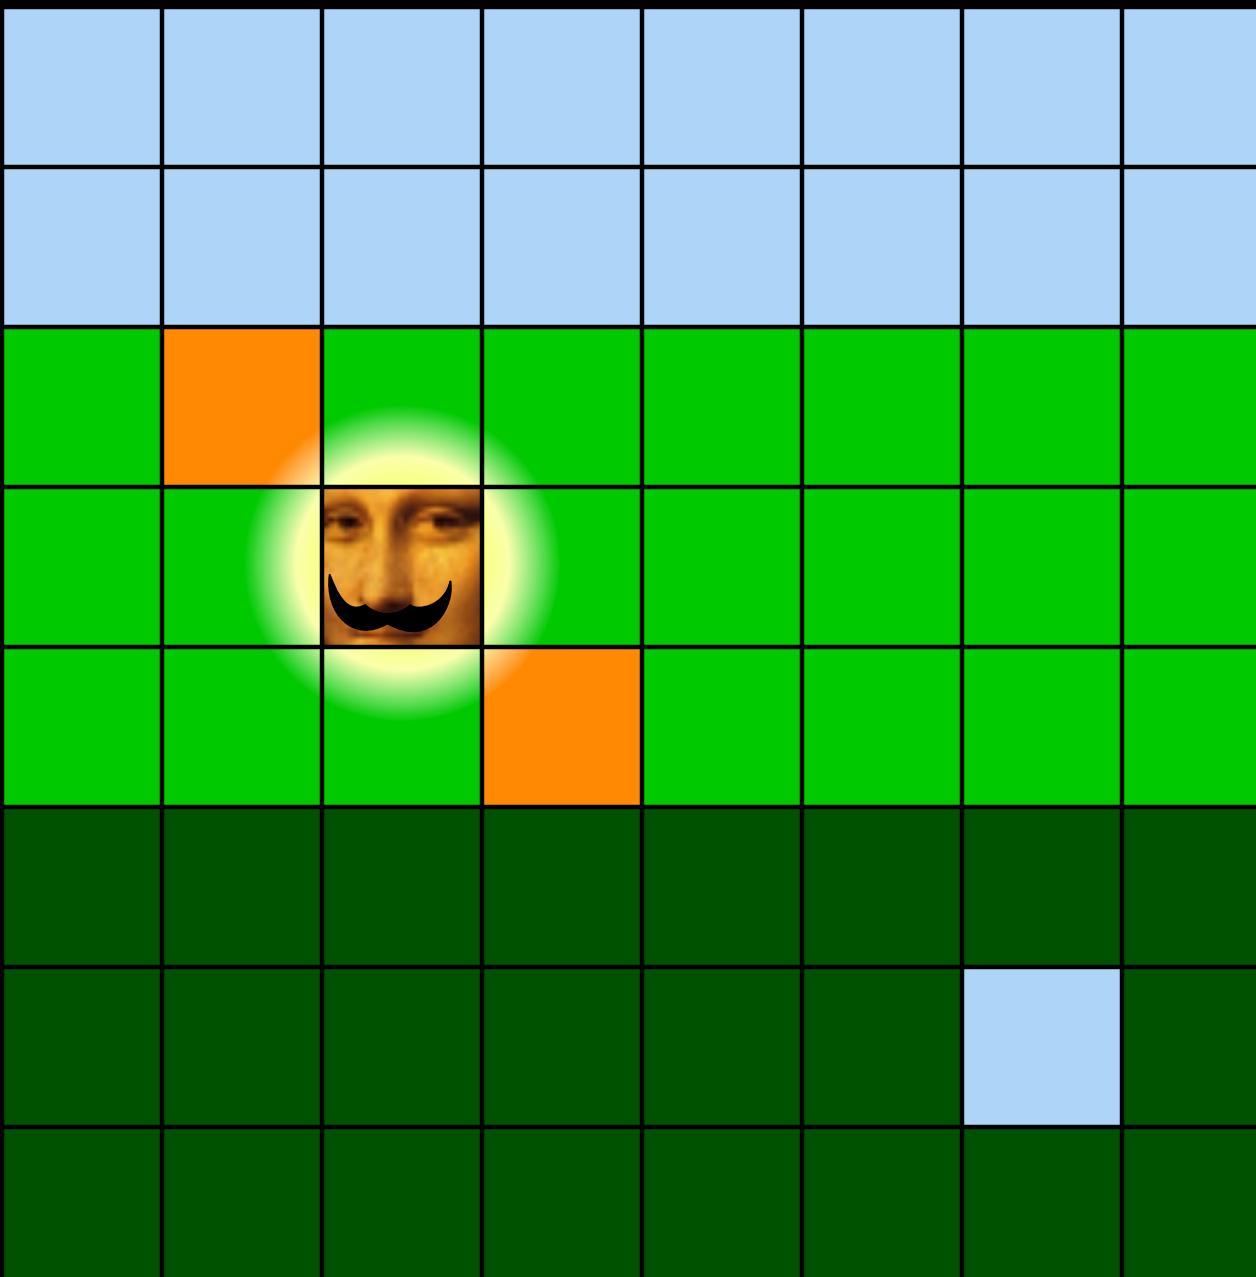


victim memory



# Deduplication implementation: KVM on Linux (KSM)

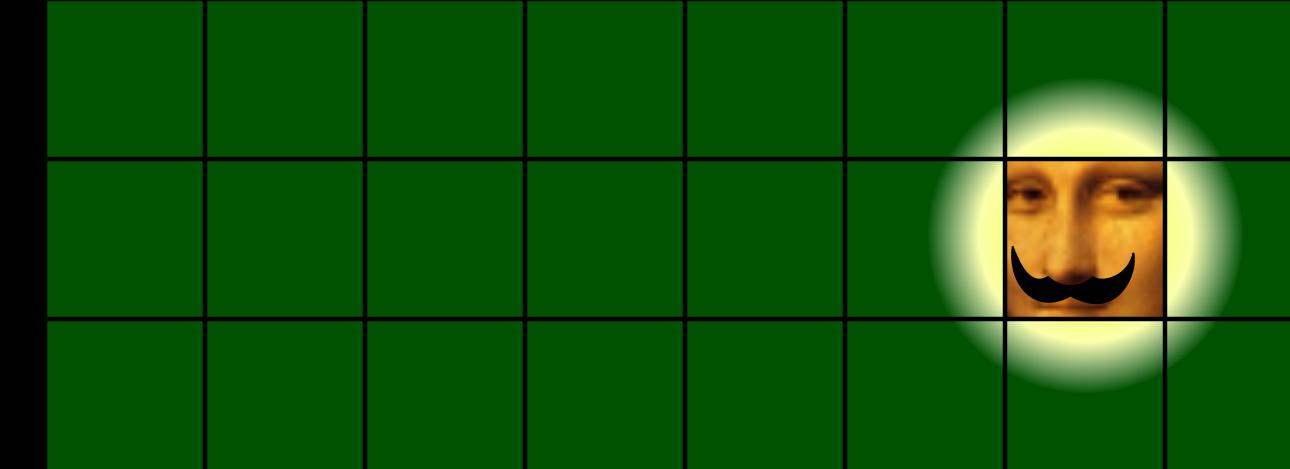
physical memory



attacker memory



victim memory



# Example 1: OpenSSH

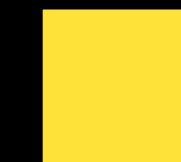
Target: `~/.ssh/authorized_keys`

# OpenSSH `~/.ssh/authorized_keys`

```
ssh-rsa AAAAB3NzaC1yc2EAAAQABAAQCB52/Uk84iUmmic  
el7ESr+/D/PWZ6Ljkhlu8yv35bEEoTwXm9eGxJyzV+1s68tRyzpD  
3VQvwSHiKqDnCg+0taAo0KvCqZcoBQFB9XawIfJI5dSeGtcUBuok  
Uv+TlmAZ+D9MNNAxjuSBBH0ShbaiH65imlauISfR3VZWF  
E7uy6sB26j52LhWG5BRwSkMnMRN2E2fqHaP96J9R0FlHuykw8jwUXJw  
l4kJ8vRo1uhX0SVu8Z9wGrKR5b+GQWJ3Ph7vjoMVU/KoAbWnNnYKR8IT  
BnkPD0LrEyAKRygEfi7gwci  
x0vQR79by8LL6ypJ4kM5eyobSBsNC  
jmghxQj8RRzGUtd1 victim@laptop
```



**Exponent**



**Modulus (p \* q)**

# OpenSSH `~/.ssh/authorized_keys`

```
ssh-rsa AAAAB3NzaC1yc2EAAAQABAAQCB52/Uk84iUmmic  
el7ESr+/D/PWZ6Ljkhlu8yv35bEEoTwXm9eGxJyzV+1s68tRyzpD  
3VQvwSHiKqDnCg+0taAo0KvCqZcoBQFB9XawIfJI5dSeGtcUBuok  
Uv+TlmAZ+D9MNNAxjuSBBH0ShbaiH65imlauISfR3VZWFE7uy6sB  
26j52LhWG5BRwSkMnMRN2E2fqHaP96J9R0FlHuykw8jwUXJwl4kJ  
8vRo1uhX0SVu8Z9wGrKR5b+GQWJ3Ph7vjoMVU/KoAbWnNnYKR8IT  
BnkPD0LrEyAKRygEfi7gwciix0vQR79by8LL6ypJ4kM5eyobSBsNC  
jmghxQj8RRzGUtd1 victim@laptop
```



Exponent



Modulus ( $p' * q' * r' \dots$ )

# Example 1: OpenSSH

**Target:** `~/.ssh/authorized_keys`

- > Flip a bit in the RSA modulus
- > Factorize it
- > Reconstruct the new private key

## Example 2: GPG & apt-get

Targets: sources.list

flip package repository domain name

eg. ubuntu.com -> ubunvu.com

# Example 2: GPG & apt-get

Targets: sources.list

+

GPG keyring

corrupt signing key