

Minemu: Protecting buggy binaries from memory corruption attacks

WARNING
THIS PRESENTATION
MAY CONTAIN POINTERS



Programming Languages

type-safe vs. not type-safe

Programming Languages

type-safe

vs.

not type-safe

Java

Python

Ruby

Javascript

Programming Languages

type-safe

vs.

not type-safe

Java

C

Python

C++

Ruby

Javascript

Programming Languages

type-safe

vs.

not type-safe

Java

Python

Ruby

Javascript

C

C++

MEMORY
CORRUPTIONS!

Programming Languages

type-safe

vs.

not type-safe

Java

MEMORY

Python

CORRUPTIONS!

Ruby

Javascript

C

C++

MEMORY

CORRUPTIONS!

Programming Languages

type-safe

vs.

not type-safe

Java

MEMORY

Python

CORRUPTIONS!

Ruby

but not

Javascript

your fault

C

CORRUPTIONS!

MEMORY

The Stack

[code]

```
run(char *name)
{
    char buf[16];

    print("hello ");
    print("world\n")
}
```

The Stack

[code]

```
run(char *name)
{
    char buf[16];

    print("hello ");
    print("world\n")
}
```

[stack]



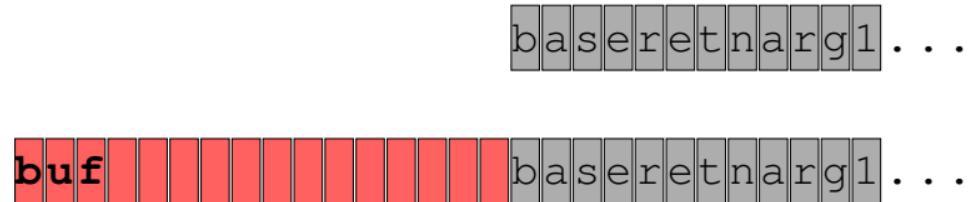
The Stack

[code]

```
run(char *name)
{
    char buf[16];

    print("hello ");
    print("world\n")
}
```

[stack]



The Stack

[address] [code]

```
8048751: run(char *name)
{
```

```
    char buf[16];
```

```
8048770:     print("hello ");
```

```
8048798:     print("world\n")
}
```

[stack]

base|ret|narg1|...|

buf|...|base|ret|narg1|...|

The Stack

[address] [code]

```
8048751: run(char *name)
{
```

```
    char buf[16];
```

```
8048770:     print("hello ");
```

```
8048798:     print("world\n")
```

```
}
```

[stack]

base|retnarg1|...

buf|base|retnarg1|...

retnarg1|buf|base|retnarg1|...

The Stack

[address] [code]

```
8048751: run(char *name)
{
```

```
    char buf[16];
```

```
8048770:     print("hello ");
```

```
8048798:     print("world\n")
```

```
}
```

[stack]

base | return address | argument 1 | ... | buf | buffer bytes | base | return address | argument 1 | ...

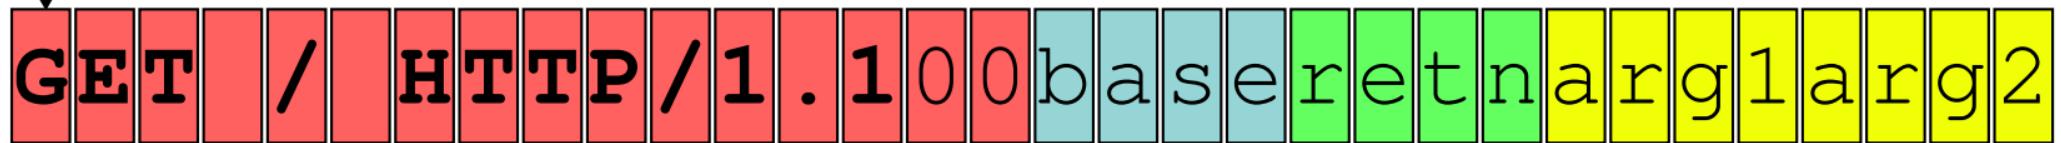
buf | buffer bytes | base | return address | argument 1 | ...

return address | buf | buffer bytes | base | return address | argument 1 | ...

return address | buf | buffer bytes | base | return address | argument 1 | ...

Traditional Stack Smashing

buf [16]



Traditional Stack Smashing

buf [16]

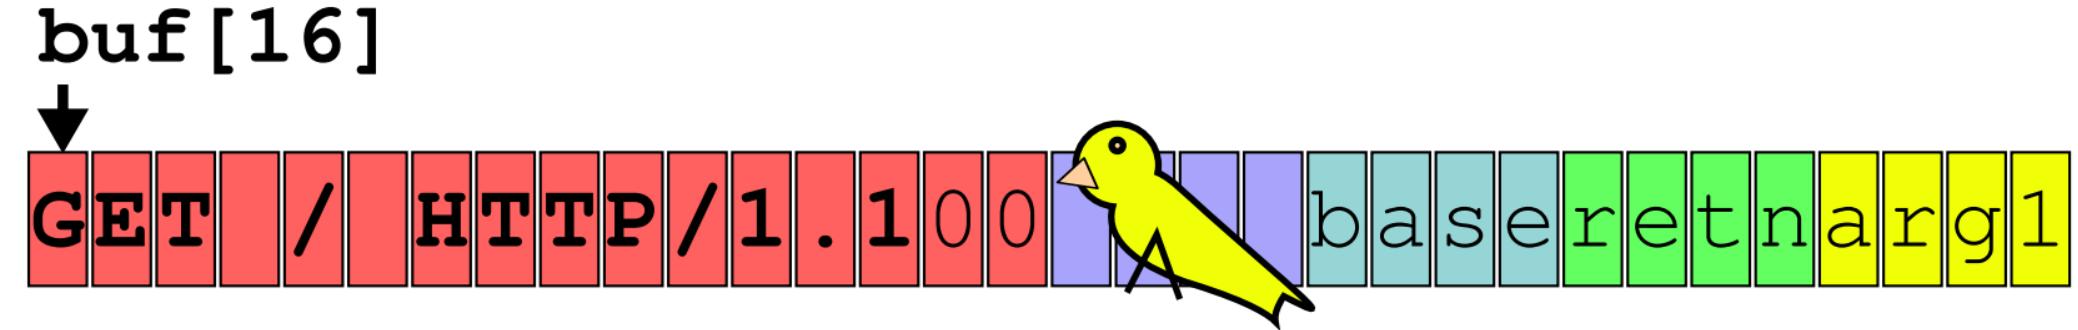


Address Space Layout Randomisation (ASLR)

buf [16]

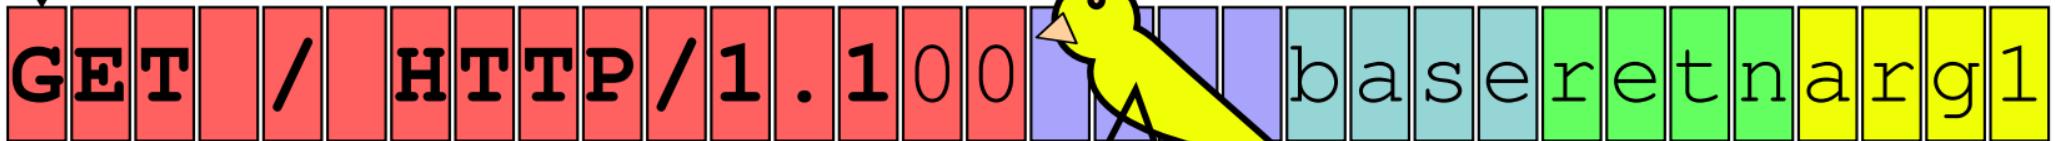


Stack Canaries



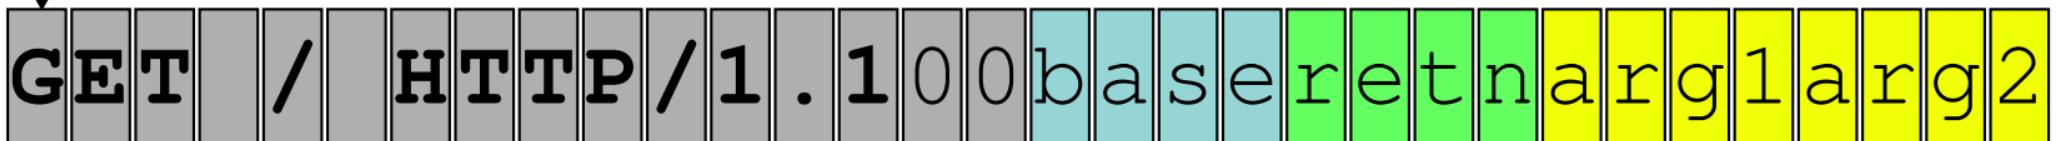
Stack Canaries

buf [16]

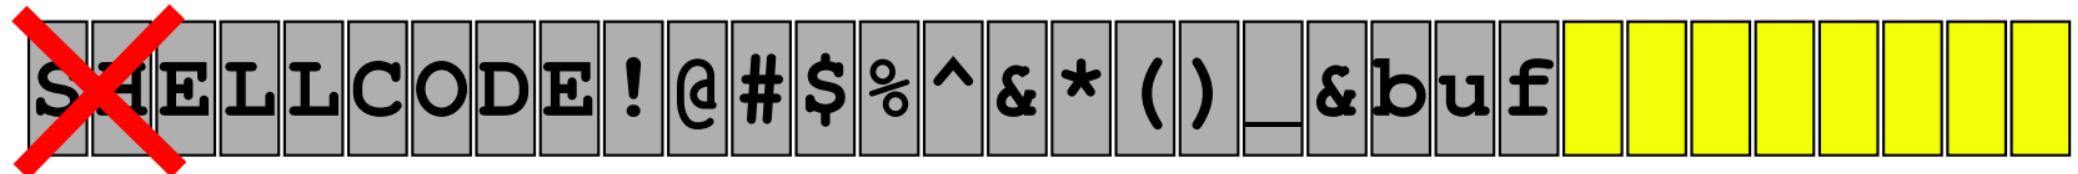


Non-executable data (DEP / NX)

buf [16]



GET / HTTP/1.100baseretvalnarg1arg2



~~SHELLCODE !@#\$%^&* () &buf~~

Fortify Source

```
char buf[16];  
memcpy(buf, r->buf, r->len);
```

GET / HTTP/1.1 100base

retnarg1arg2

sh; STACKSMASHERA

Fortify Source

```
char buf[16];
memcpy(buf, r->buf, r->len);
```

GET / HTTP/1.1base

```
char buf[16];
memcpy_chk(buf, r->buf, r->len, 16);
```

sh; STACKSMASHERA

XXXXXXXXXXXX

```
*** buffer overflow detected ***: /my/fortified/binary terminated
===== Backtrace: =====
/lib/i386-linux-gnu/i686/cmov/libc.so.6(__fortify_fail+0x50)[0xb774a4d0]
/lib/i386-linux-gnu/i686/cmov/libc.so.6(+0xe040a)[0xb774940a]
/my/fortified/binary[0x8048458]
/lib/i386-linux-gnu/i686/cmov/libc.so.6(__libc_start_main+0xe6)[0xb767fe46]
/my/fortified/binary[0x8048371]
===== Memory map: =====
08048000-08049000 r-xp 00000000 fe:00 282465      /my/fortified/binary
08049000-0804a000 rw-p 00000000 fe:00 282465      /my/fortified/binary
08600000-08621000 rw-p 00000000 00:00 0          [heap]
b764b000-b7667000 r-xp 00000000 fe:00 131602      /lib/i386-linux-gnu/libgcc_s.so.1
b7667000-b7668000 rw-p 0001b000 fe:00 131602      /lib/i386-linux-gnu/libgcc_s.so.1
b7668000-b7669000 rw-p 00000000 00:00 0
...

```

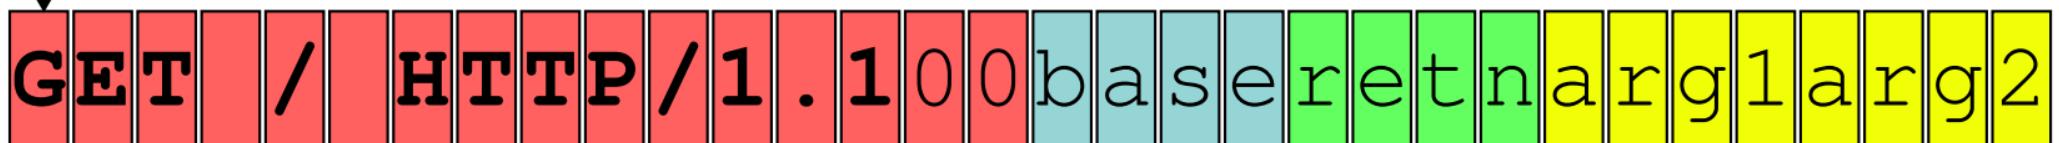
Aborted

FORTIFY ALL THE THINGS



Return Oriented Programming (ROP)

buf[16]



GET / HTTP/1.100base return nargs1 args2

sh; STACKSMASHER... ROP1 ROP2 var1

pointer to useful code

Some exploits still work with all these defense measures.

Example: nginx buffer underrun (CVE-2009-2629)

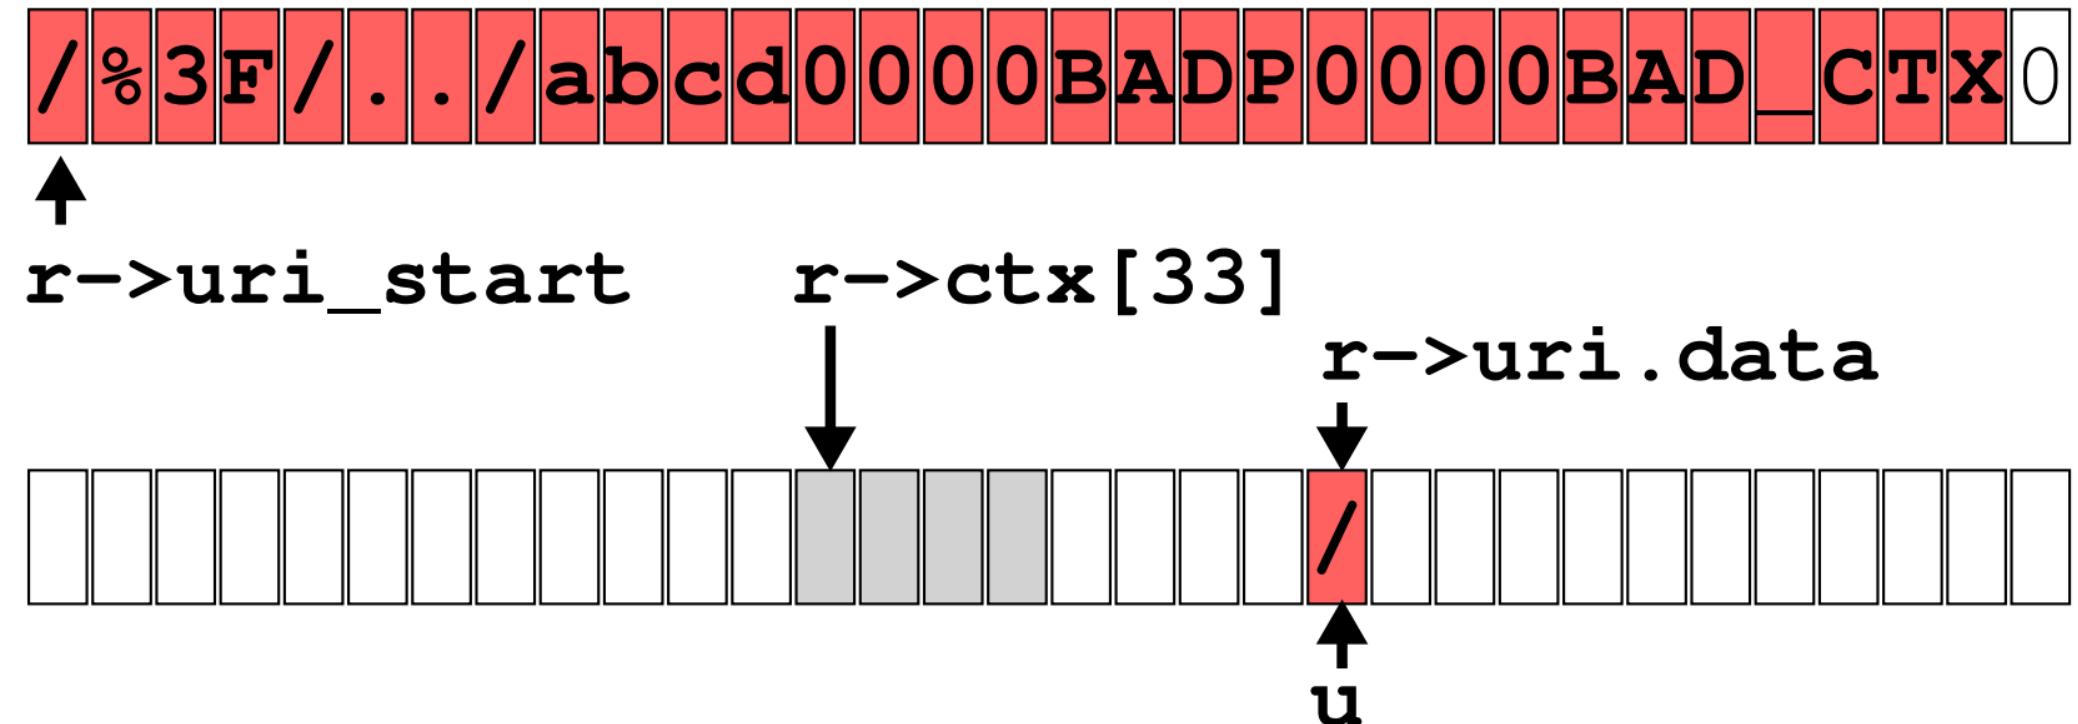
CVE-2009-2629

/%3F/.../abcd0000BADP0000BAD_CTX0

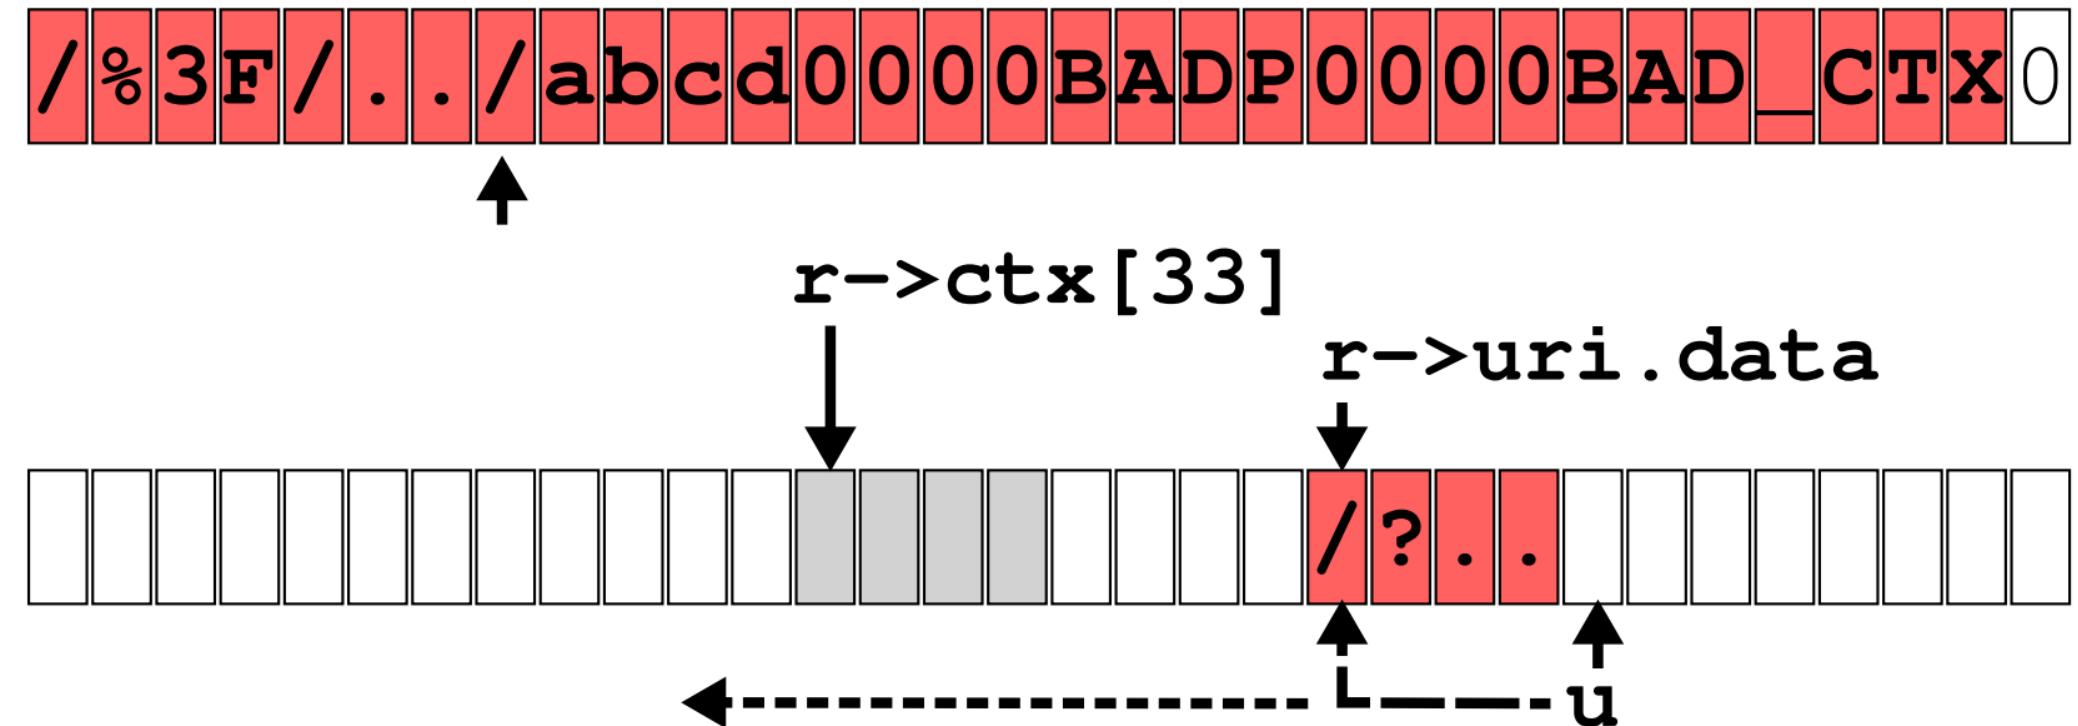


r->uri_start

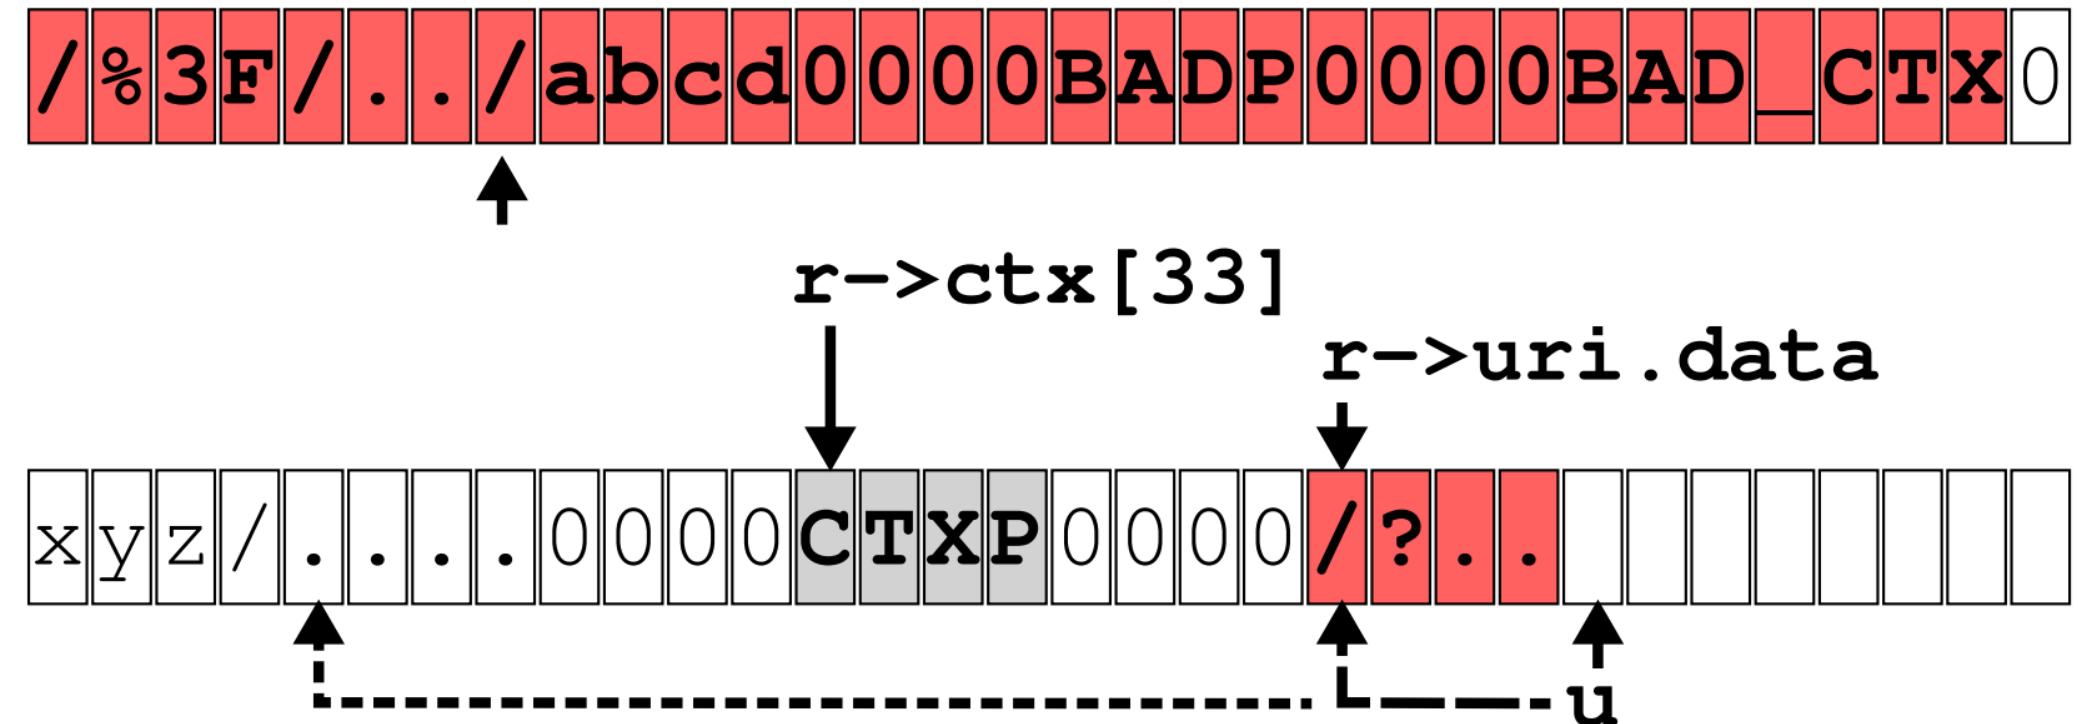
CVE-2009-2629



CVE-2009-2629



CVE-2009-2629



CVE-2009-2629

/%3F/.../abcd0000BADP0000BAD_CTX0

$r \rightarrow \text{ctx}[33]$



$r \rightarrow \text{uri.data}$



x y z / 00 00 BADP0000BAD_CTX0




```
typedef struct {
    ngx_buf_t                                *buf;
    ngx_chain_t                               *in;
    ngx_chain_t                               *free;
    ngx_chain_t                               *busy;

    unsigned                                  sendfile;
    unsigned                                  need_in_memory;
    unsigned                                  need_in_temp;

    ngx_pool_t                                *pool;
    ngx_int_t                                 allocated;
    ngx_bufs_t                                bufs;
    ngx_buf_tag_t                            tag;

    ngx_output_chain_filter_pt   output_filter;
    void                                     *filter_ctx;
} ngx_output_chain_ctx_t;
```

function pointer



805ba93: mov (%ecx),%ebx ; copy filename
 movl \$0x3,0x10(%ecx)
 mov %ecx,(%esp)
 call *0x2c(%ecx)

```
805ba93: mov    (%ecx),%ebx          ; copy filename  
        movl   $0x3,0x10(%ecx)  
        mov    %ecx,(%esp)  
        call   *0x2c(%ecx)
```



```
8052267: mov    %eax,0x4(%esp)      ; push argv  
        mov    %ebx,(%esp)          ; push filename  
        call   *0x14(%ebx)
```

```
805ba93:  mov    (%ecx),%ebx          ; copy filename
           movl   $0x3,0x10(%ecx)
           mov    %ecx,(%esp)
           call   *0x2c(%ecx)

8052267:  mov    %eax,0x4(%esp)       ; push argv
           mov    %ebx,(%esp)          ; push filename
           call   *0x14(%ebx)

804b274:  <execve@plt>            ; get shell
```

- defeats address randomisation (through info leak)

- defeats address randomisation (through info leak)
- defeats non-executable data protection

- defeats address randomisation (through info leak)
- defeats non-executable data protection
- not a standard copy function (no fortify protections)

- defeats address randomisation (through info leak)
- defeats non-executable data protection
- not a standard copy function (no fortify protections)
- not return oriented, so stack smash protection does not matter

But the situation is even worse

But the situation is even worse

- needs to be enabled at compile time, and there is a lot of old code out there

But the situation is even worse

- needs to be enabled at compile time, and there is a lot of old code out there**
- many packages do not apply these defence mechanisms even today**

But the situation is even worse

- needs to be enabled at compile time, and there is a lot of old code out there
- many packages do not apply these defence mechanisms even today
- implementation flaws

Can we do more?

Can we do more?

**>> Non-executable data prevents untrusted data from
being run as code**

Can we do more?

**>> Non-executable data prevents untrusted data from
being run as code**

**<< Return oriented programming replaces untrusted
code with pointers to original code.**

Can we do more?

- >> Non-executable data prevents untrusted data from being run as code**
- << Return oriented programming replaces untrusted code with pointers to original code.**
- >> Can we prevent untrusted pointers from being used as jump addresses?**

Taint analysis

| | | |
|----------|-------------------------------------------------------|------------------|
| 0805be60 | 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 | |
| 0805be70 | 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 | |
| 0805be80 | 00 00 00 00 02 00 00 00 d8 4b 06 08 a0 2e 05 08 |K..... |
| 0805be90 | 94 be 05 08 78 a0 04 08 ef be ad de a4 be 05 08 |x..... |
| 0805bea0 | ac be 05 08 2f 62 69 6e 2f 73 68 00 a4 be 05 08 |/bin/sh.... |
| 0805beb0 | 00 00 00 00 53 41 4d 45 54 48 49 4e 47 57 45 44 |SAMETHINGWED |
| 0805bec0 | 4f 45 56 45 52 59 4e 49 47 48 54 50 49 4e 4b 59 | OEVERYNIGHTPINKY |
| 0805bed0 | 00 00 00 00 4e 41 52 46 90 be 05 08 ef 1f 05 08 |NARF..... |
| 0805bee0 | ff fa 26 08 ff f0 00 00 00 00 00 00 00 00 00 00 00 | ..&..... |
| 0805bef0 | 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 | |
| 0805bf00 | 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 | |

Taint tracking (1/2):

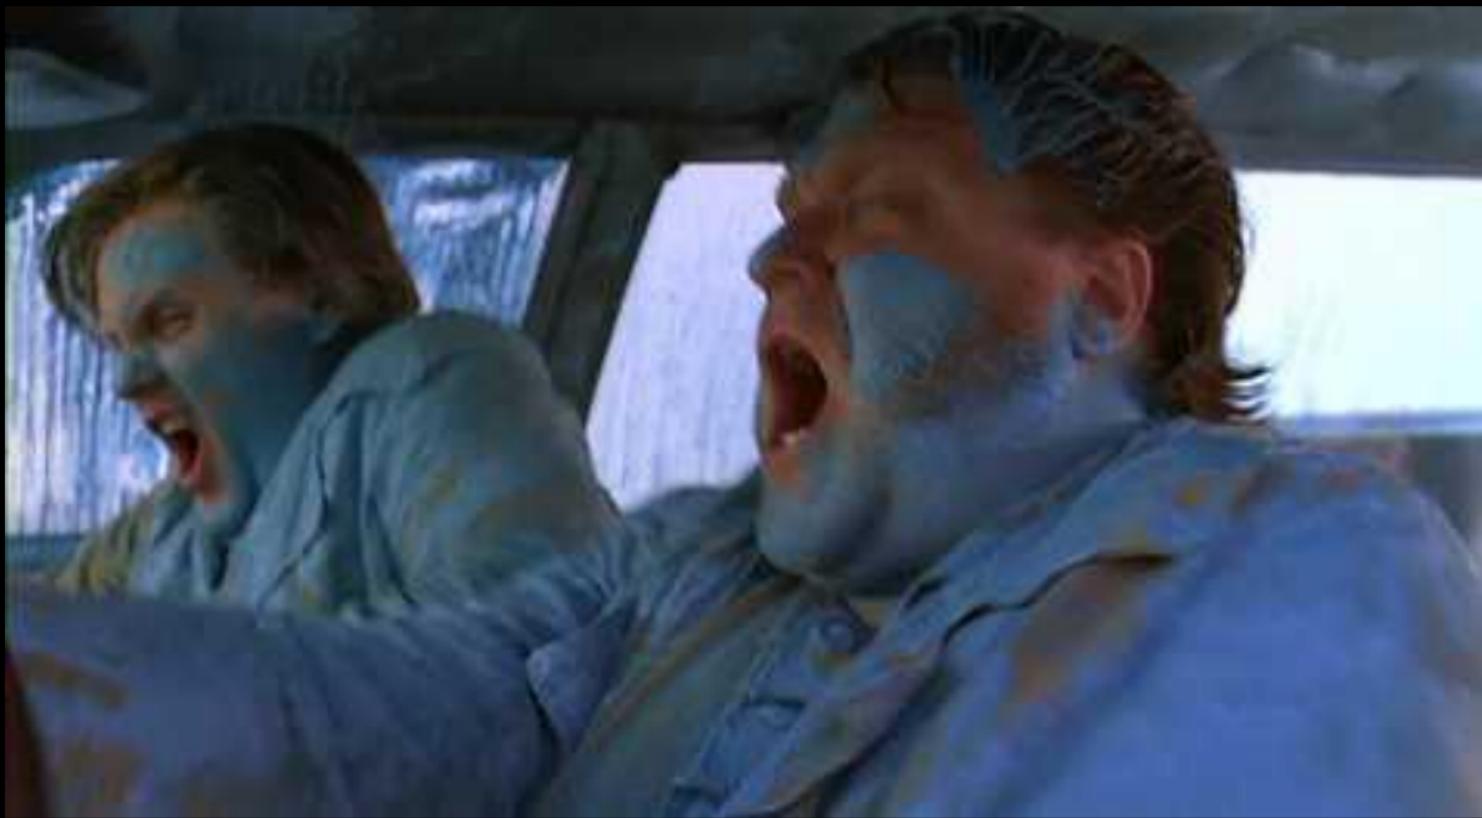
- remember whether data is trusted or not
- untrusted data is 'tainted'
- when data is copied, its taint is copied along
- taint is ORed for arithmetic operations

Taint tracking (2/2):

When the code jumps to an address in memory,
the source of this address is checked for taint.

eg.:

- RET
- CALL *%eax
- JMP *0x1c(%ebx)

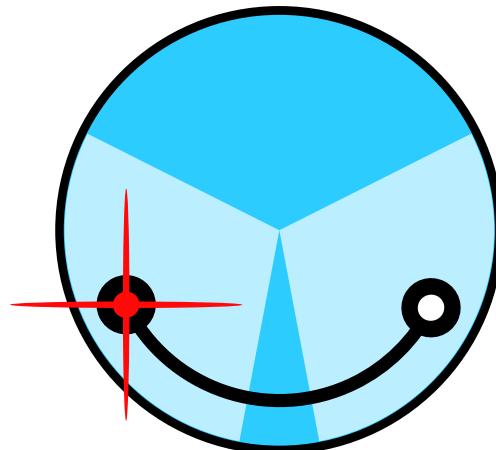


Taint tracking



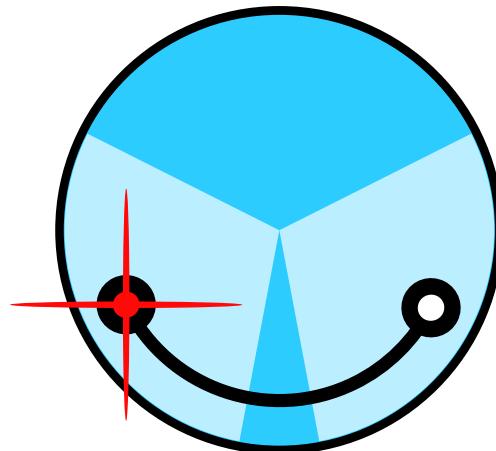
useful, but slow as hell

Is this slowness fundamental?



fast emulator
memory layout
use SSE registers to hold taint

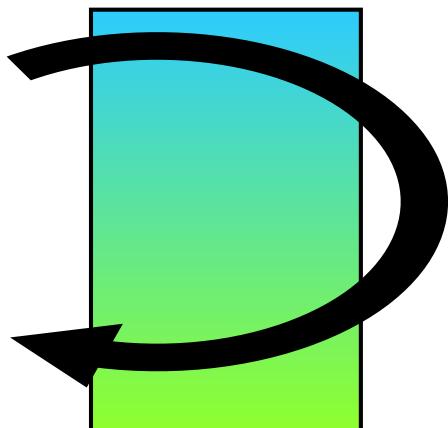
Is this slowness fundamental?



minemu

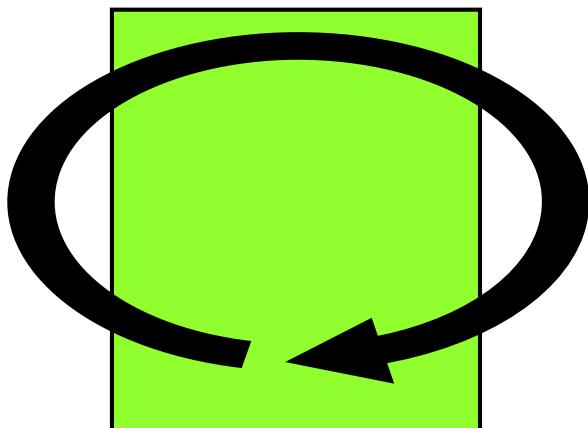
- ▶ fast emulator
memory layout
use SSE registers to hold taint

Emulator

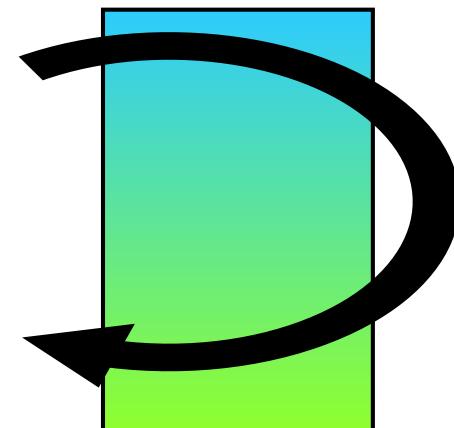


compile jit code

Emulator

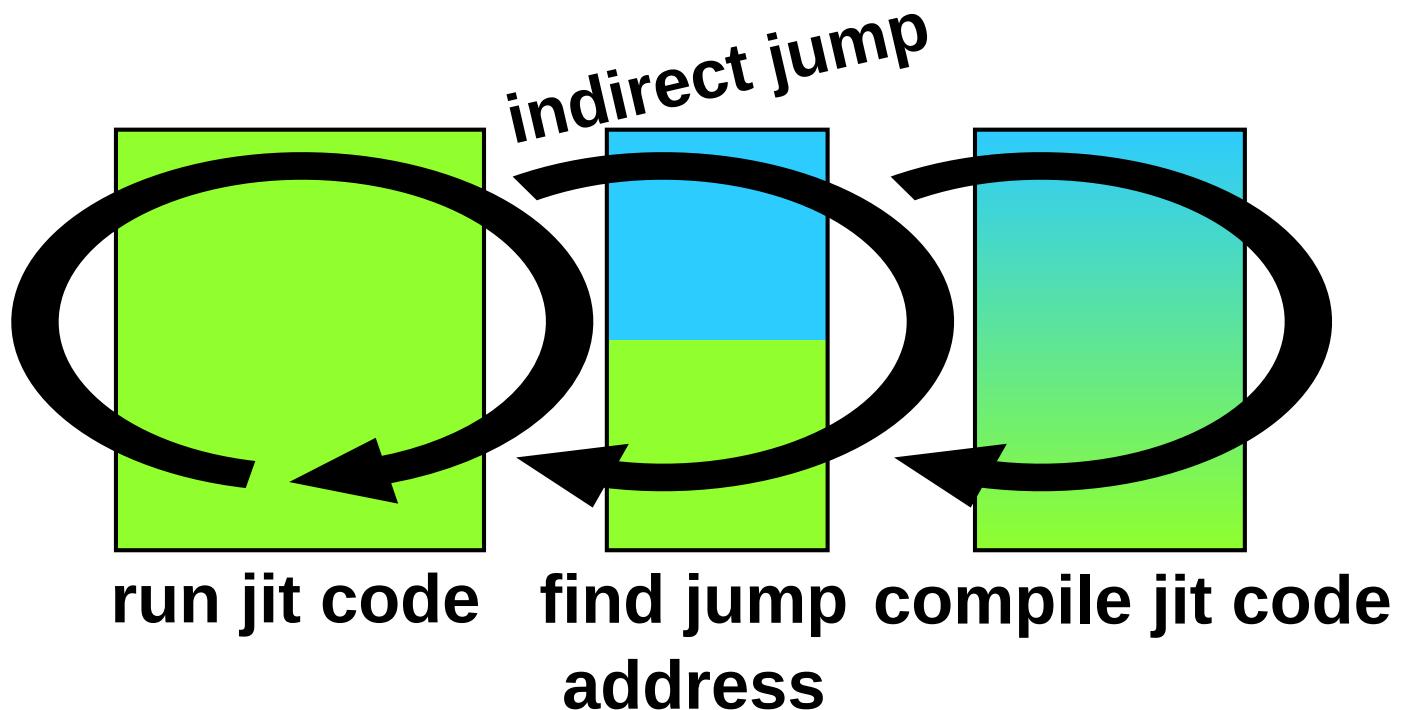


run jit code

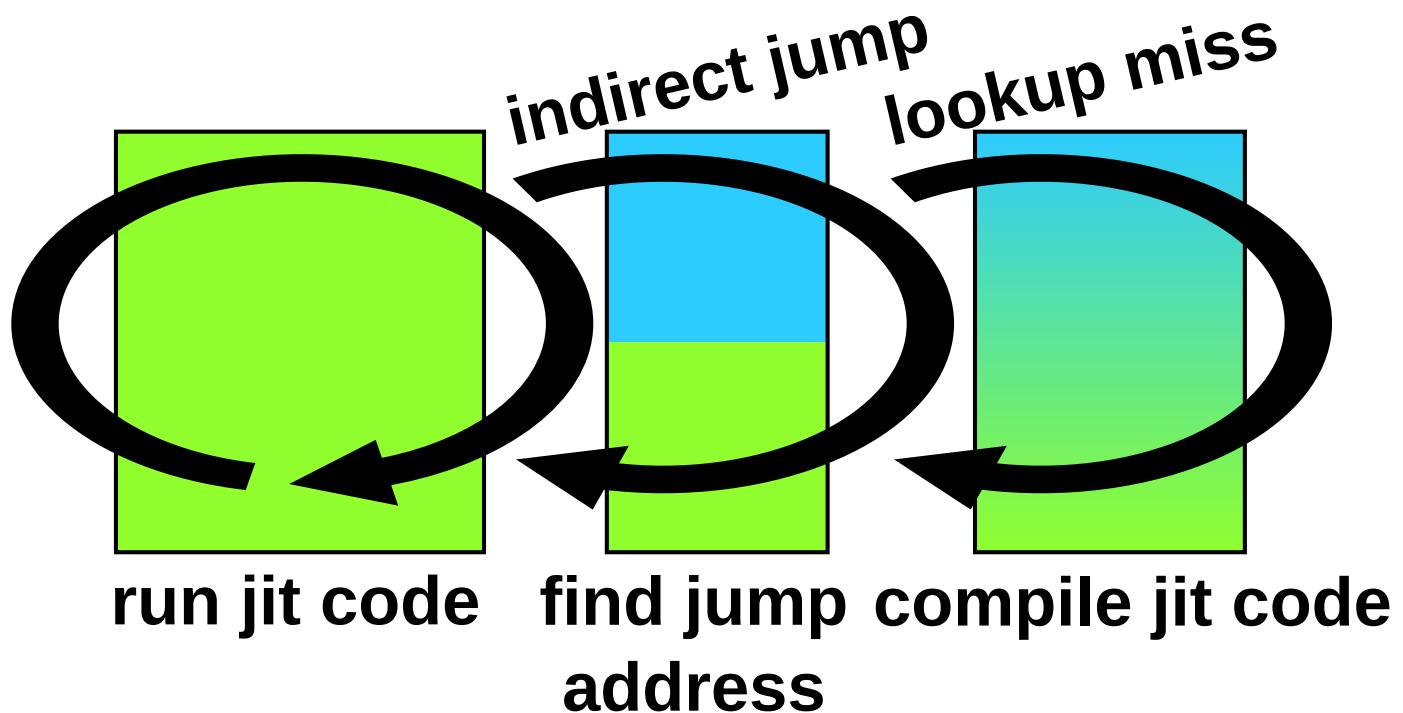


compile jit code

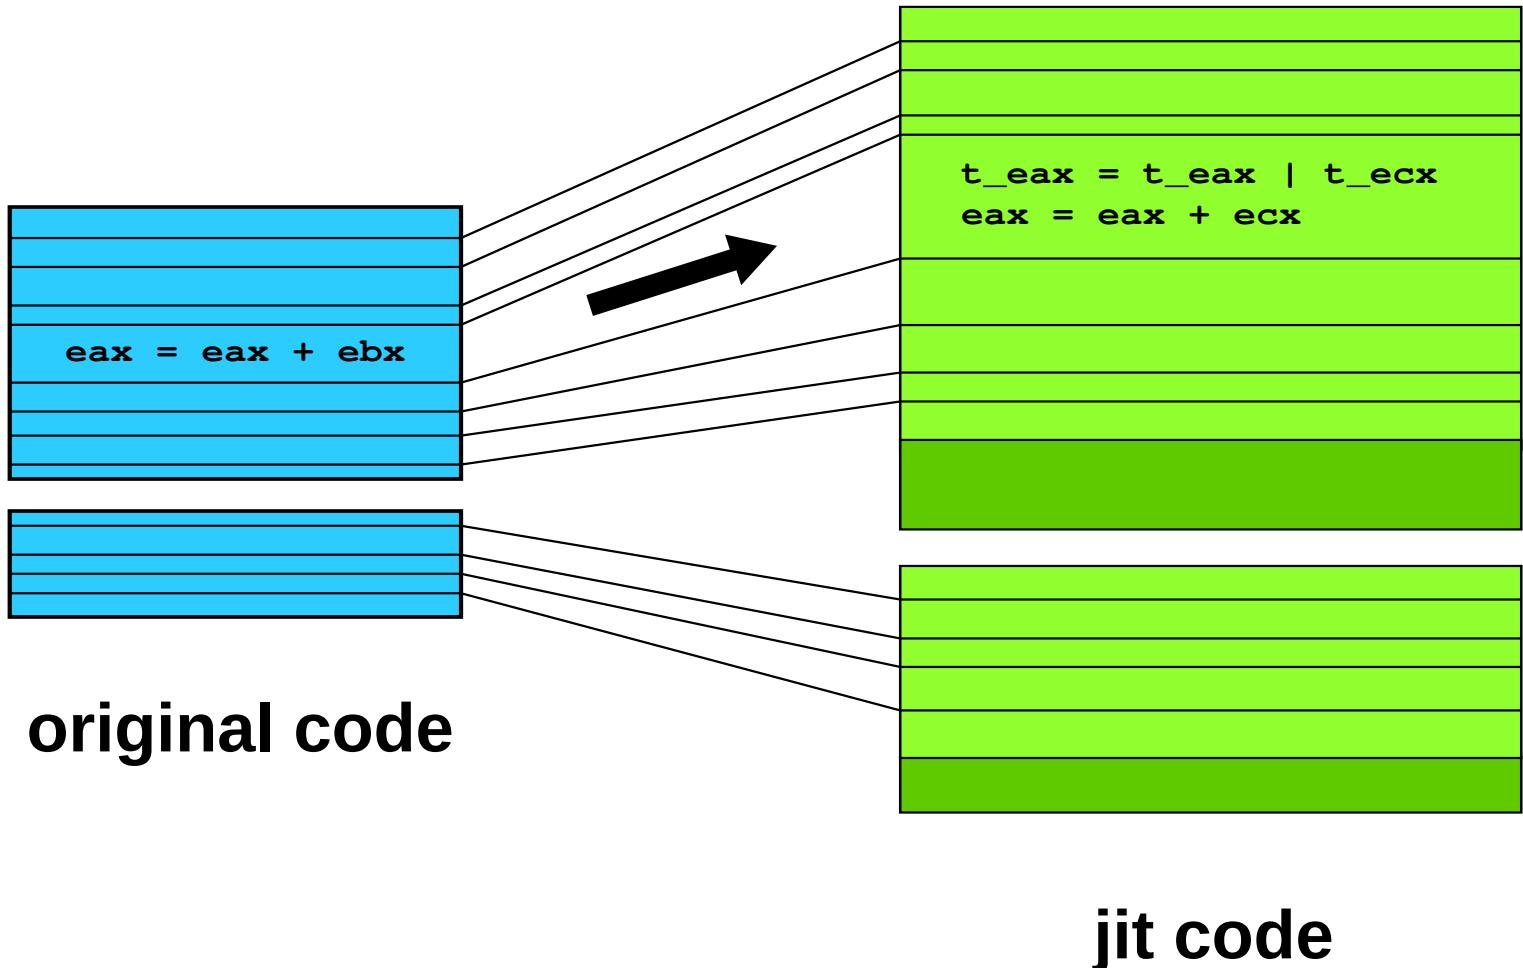
Emulator



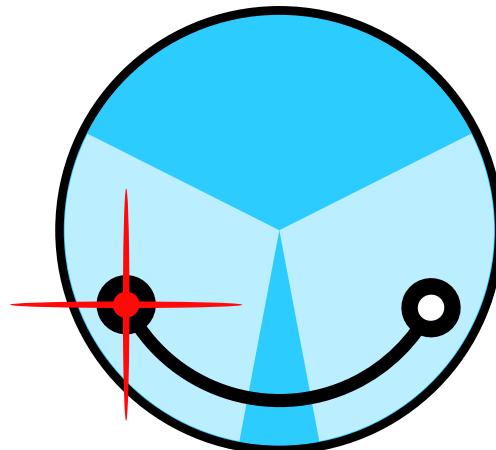
Emulator



Dynamic instrumentation



Is this slowness fundamental?

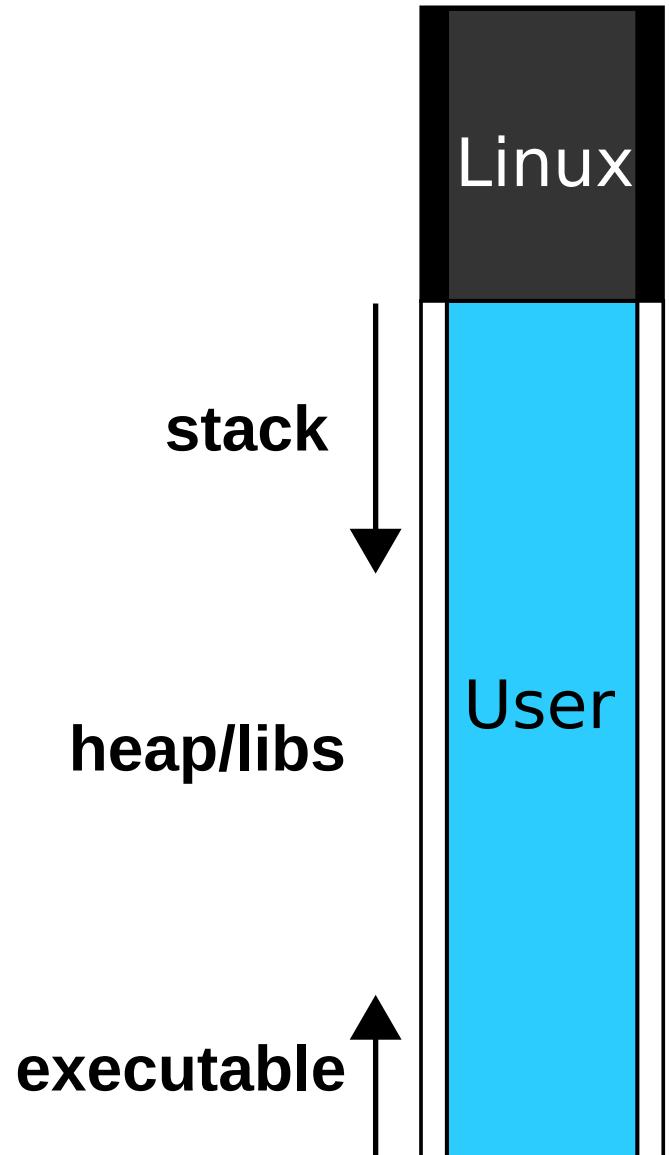


minemu

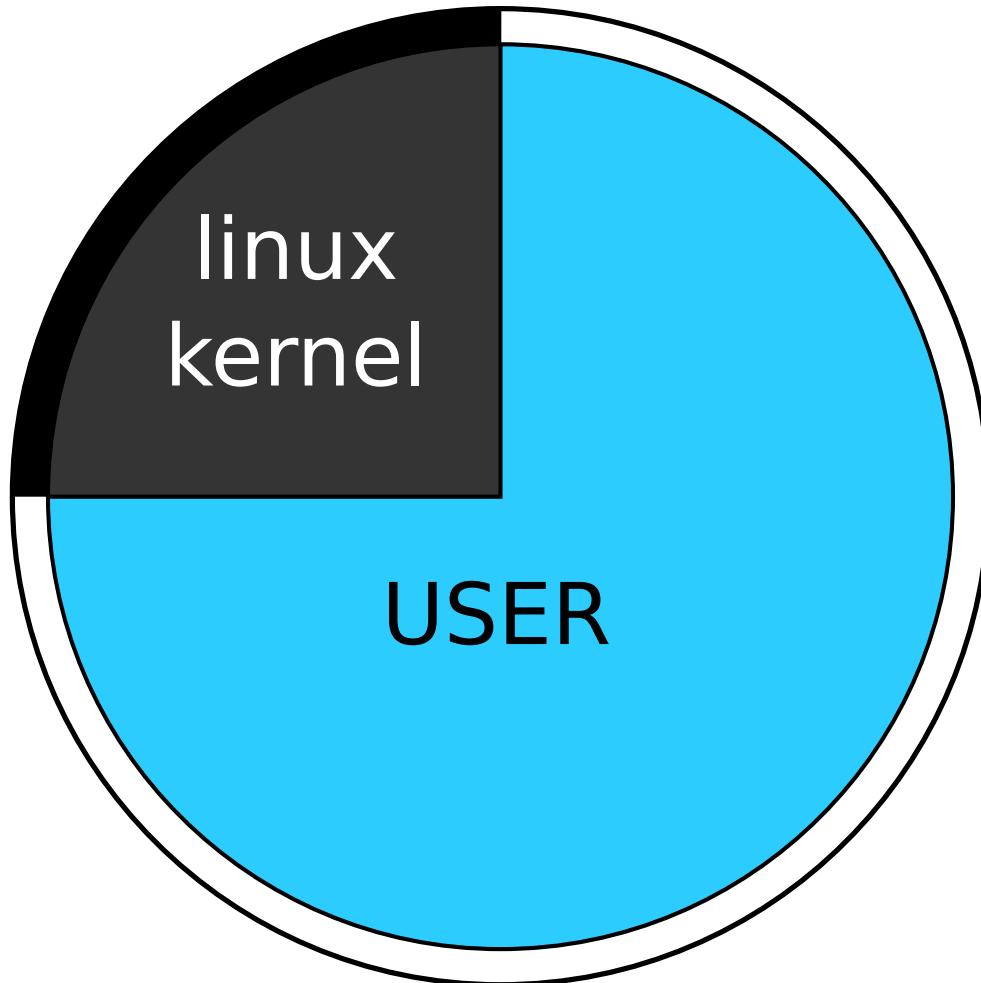
fast emulator

► memory layout

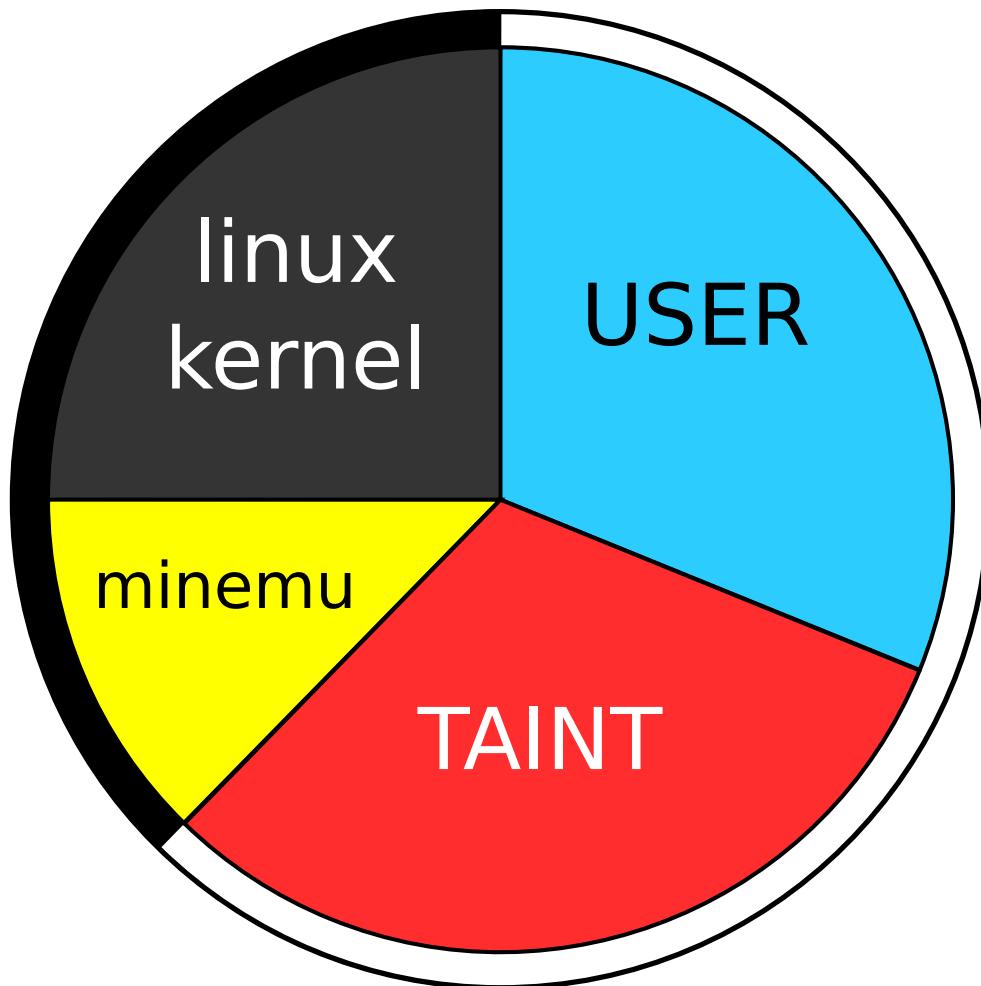
use SSE registers to hold taint



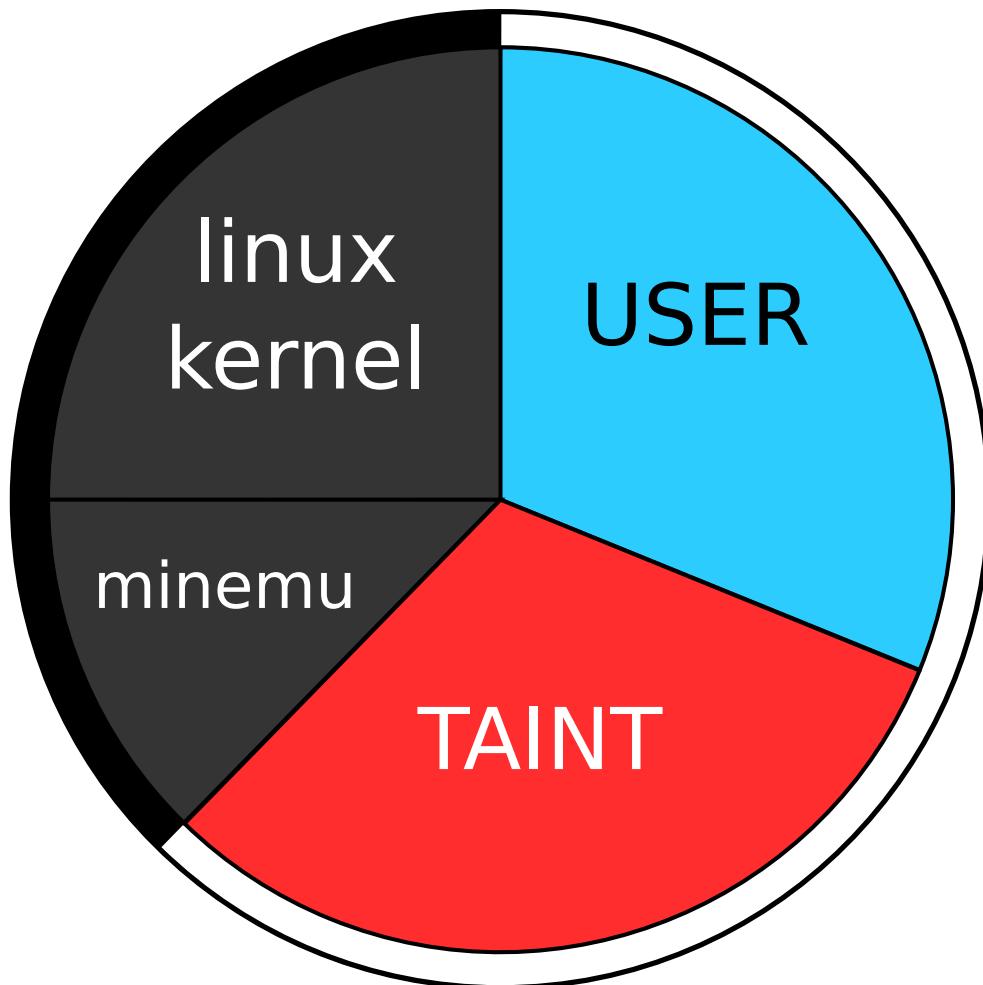
Memory layout (linux)



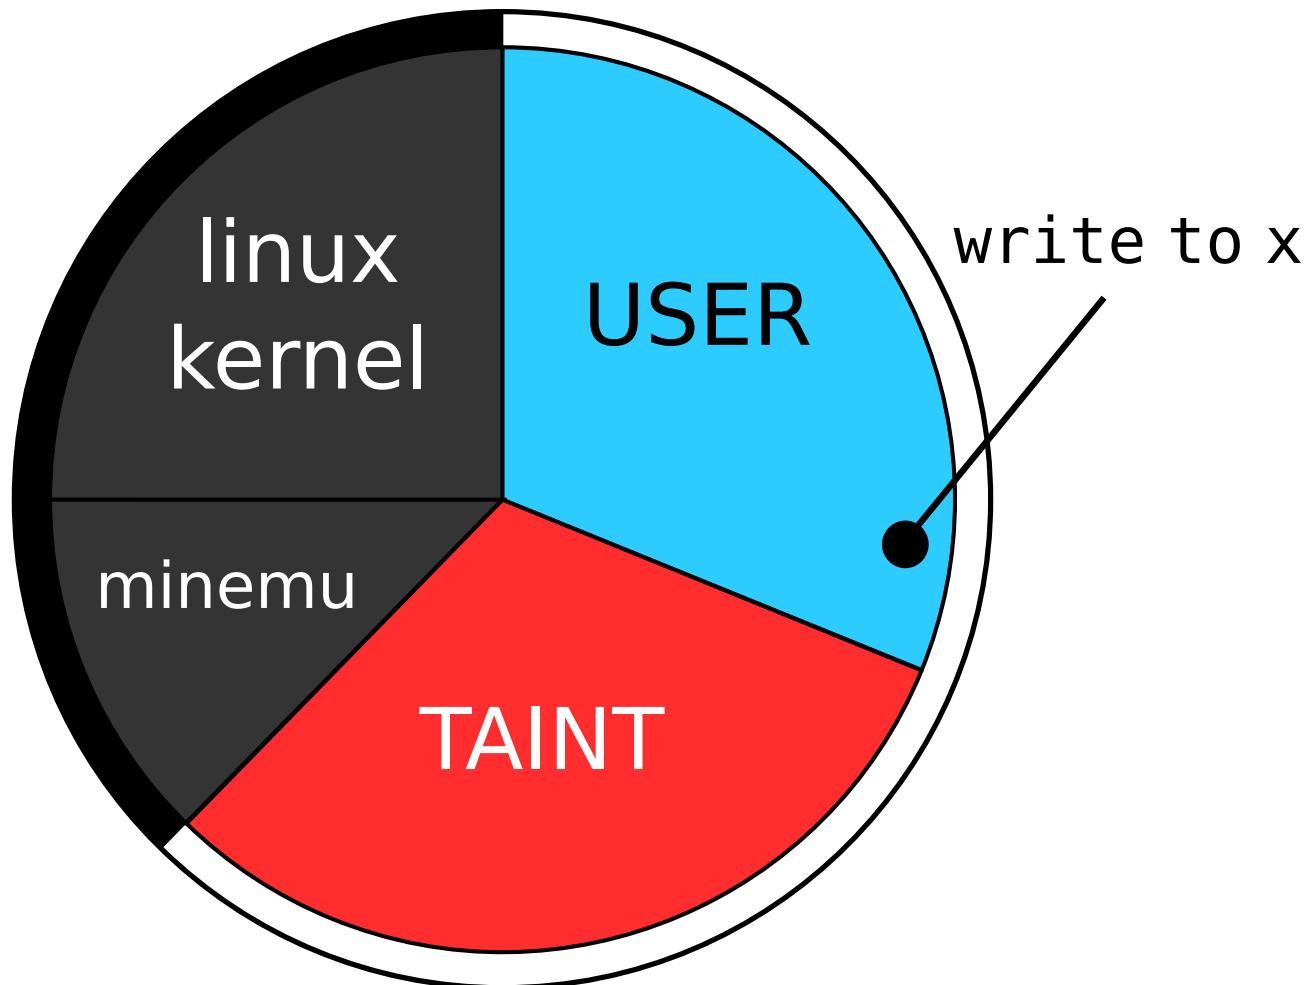
Memory layout (minemu)



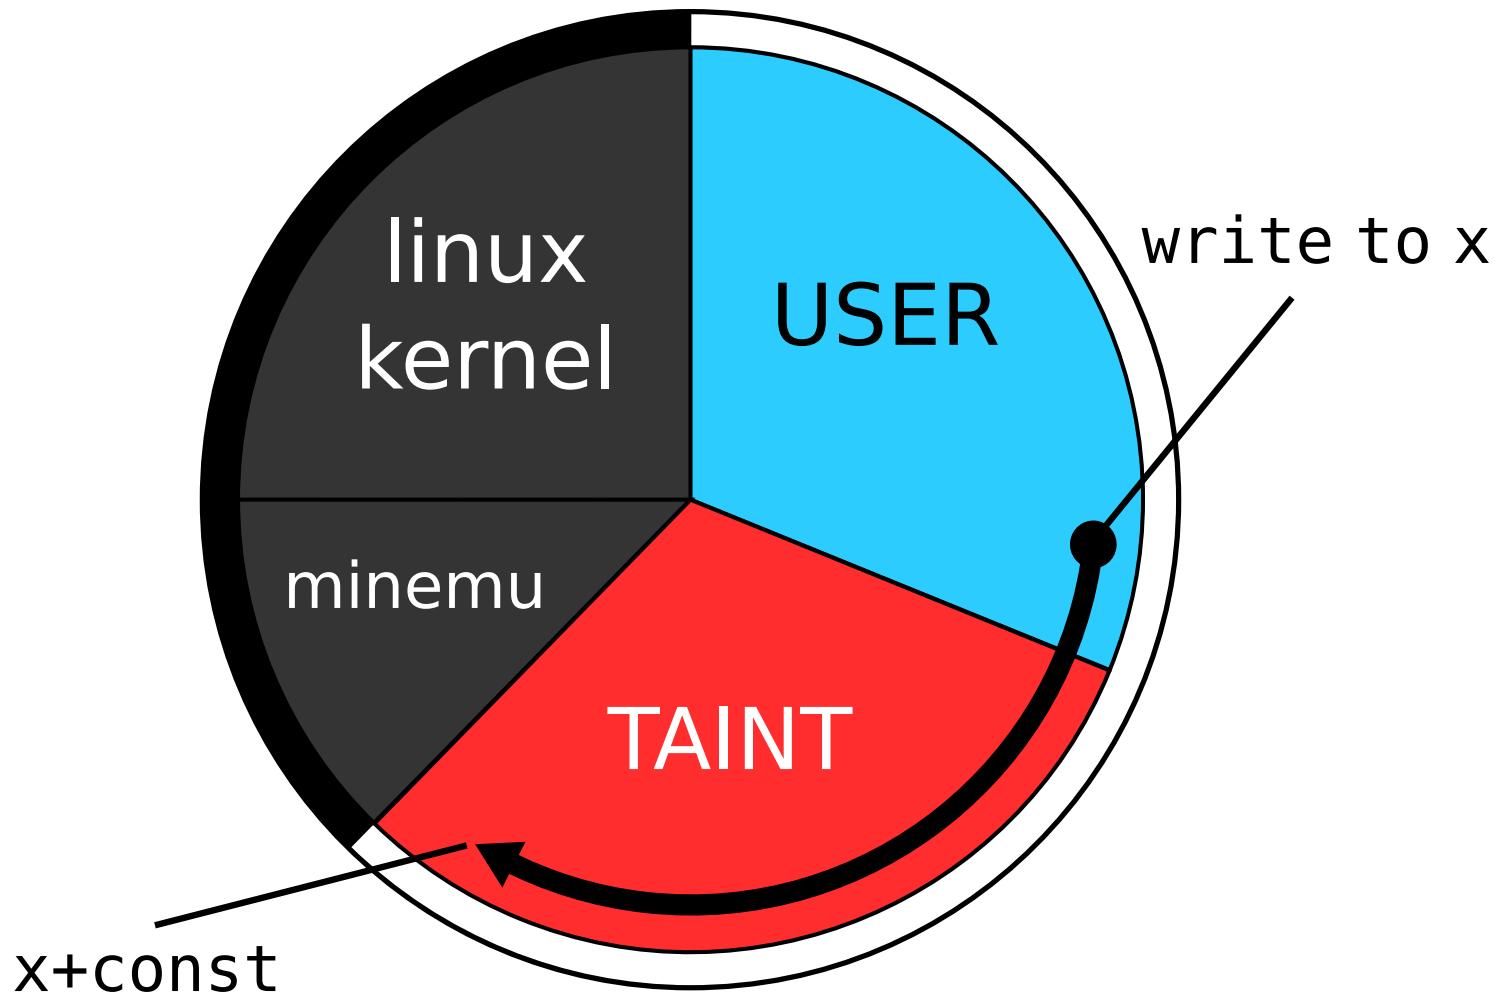
Memory layout (minemu)



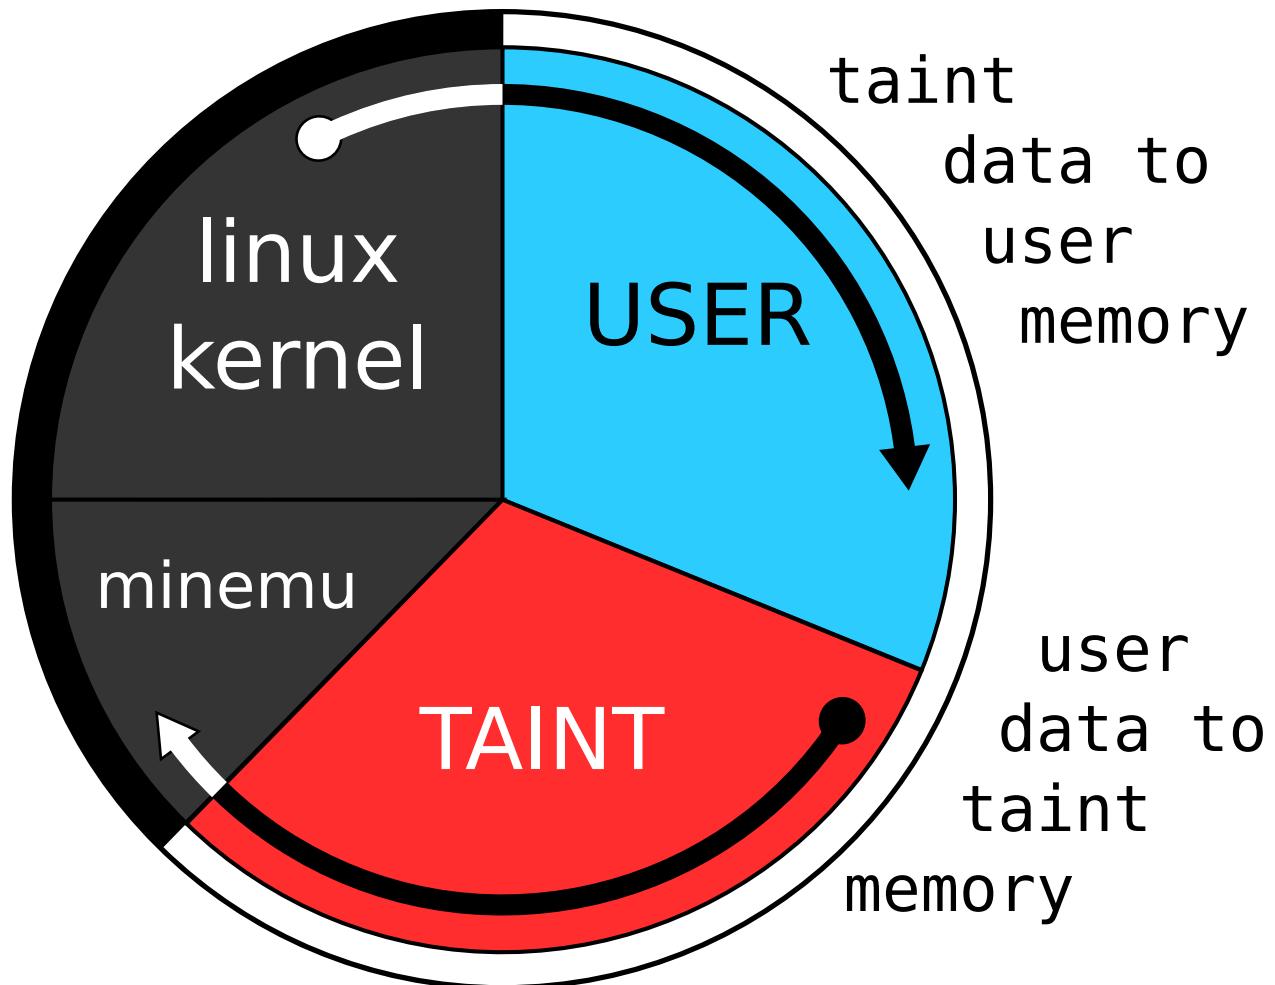
Memory layout (minemu)



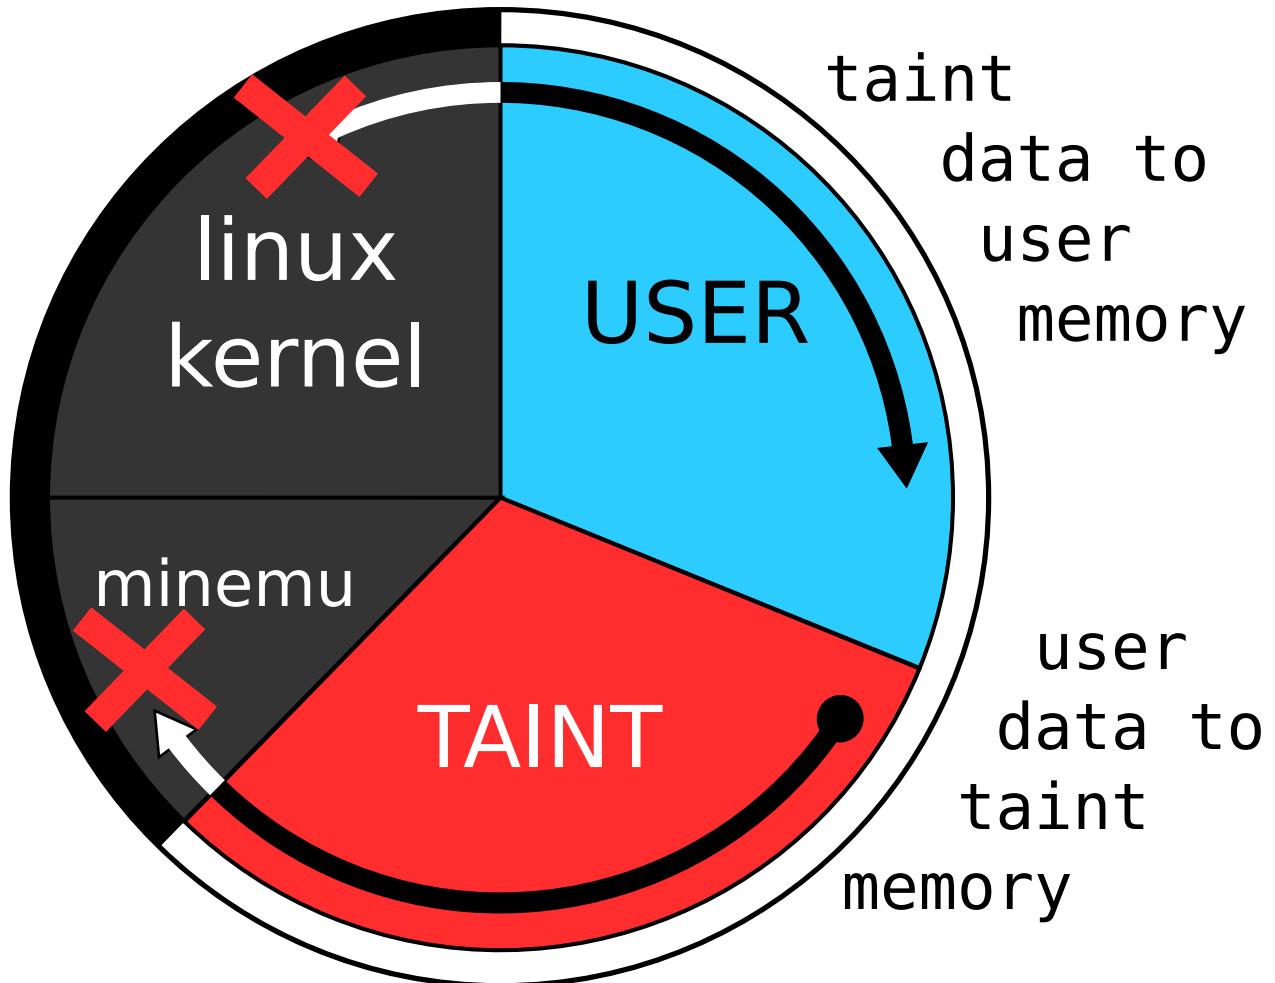
Memory layout (minemu)



Memory layout (minemu)



Memory layout (minemu)



Addressing shadow memory

```
mov EAX, (EDX)
```

Addressing shadow memory

```
mov EAX, (EDX)
```

address:

EDX

Addressing shadow memory

mov EAX, (EDX)

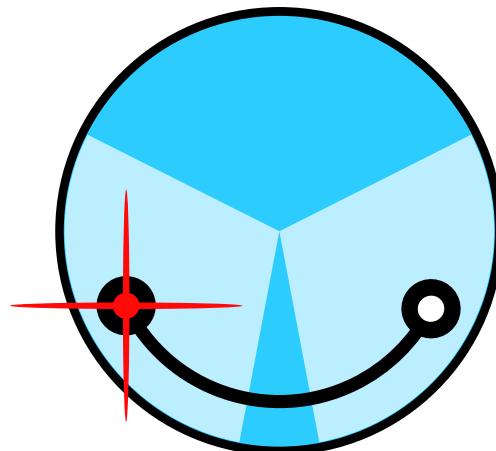
address:

EDX

taint:

EDX+const

Is this slowness fundamental?



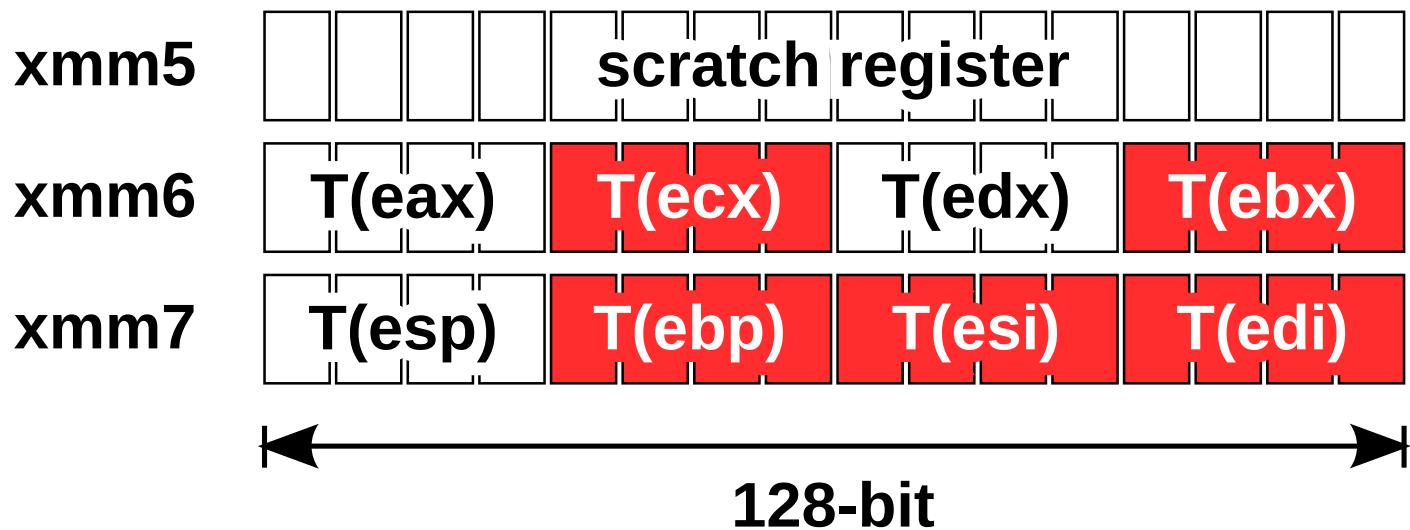
minemu

fast emulator

memory layout

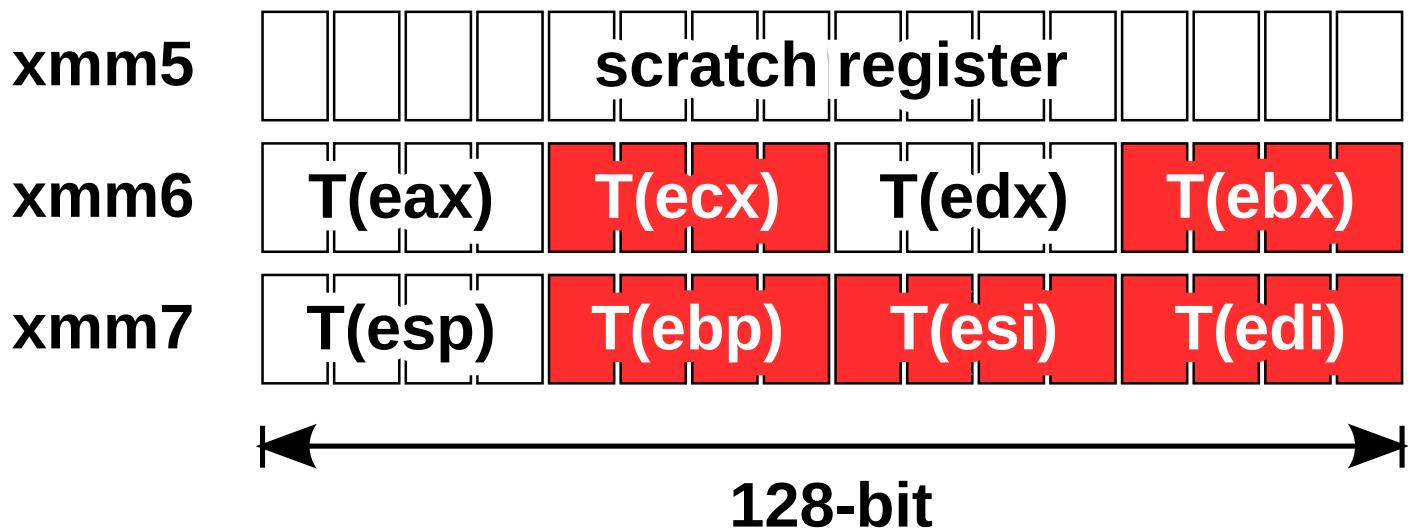
- ▶ use SSE registers to hold taint

Taint propagation in SSE registers



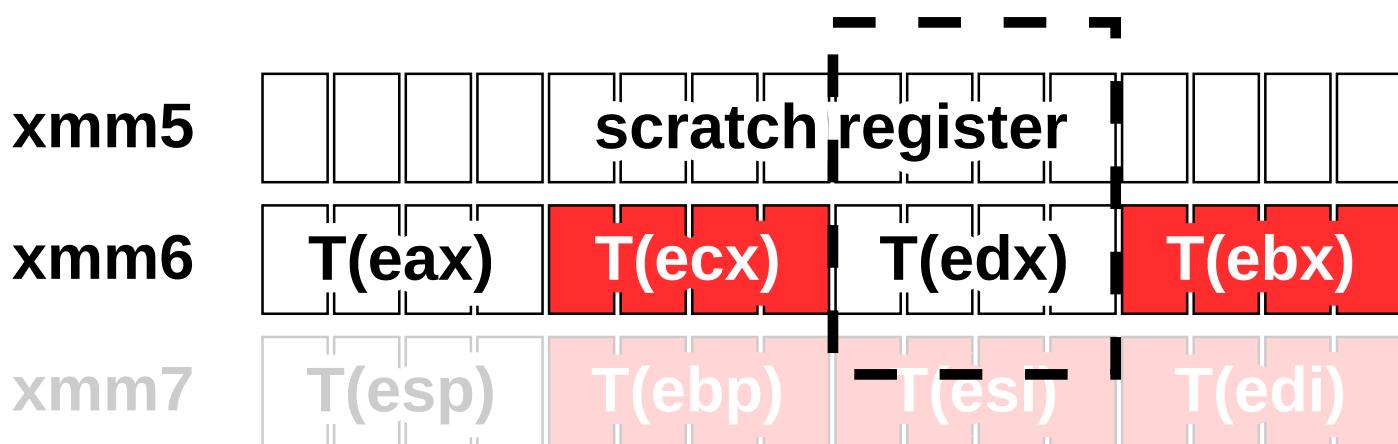
Taint propagation in SSE registers

add EDX, x



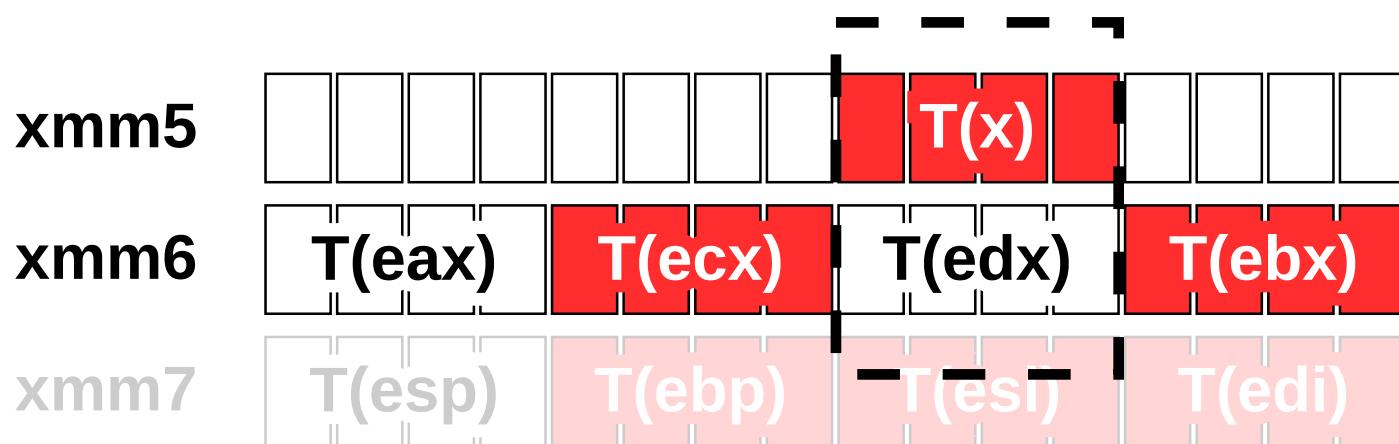
Taint propagation in SSE registers

add EDX, x



Taint propagation in SSE registers

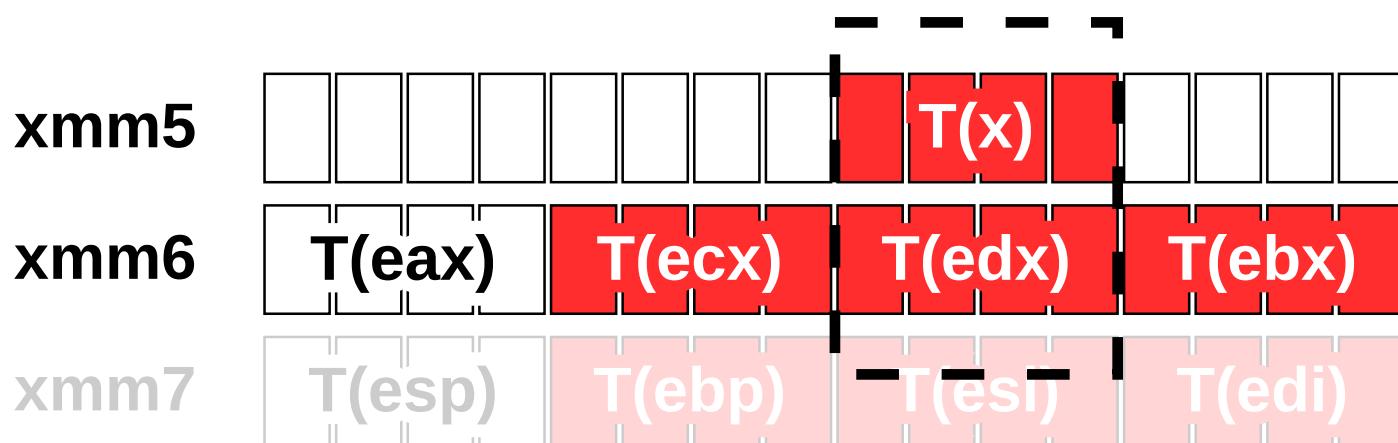
add EDX, x



vector insert

Taint propagation in SSE registers

add EDX, x



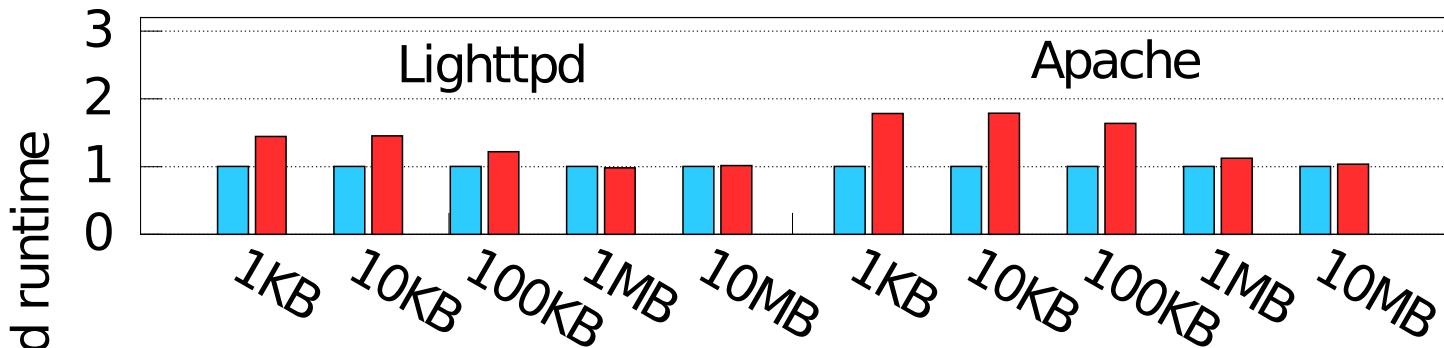
or

Effectiveness

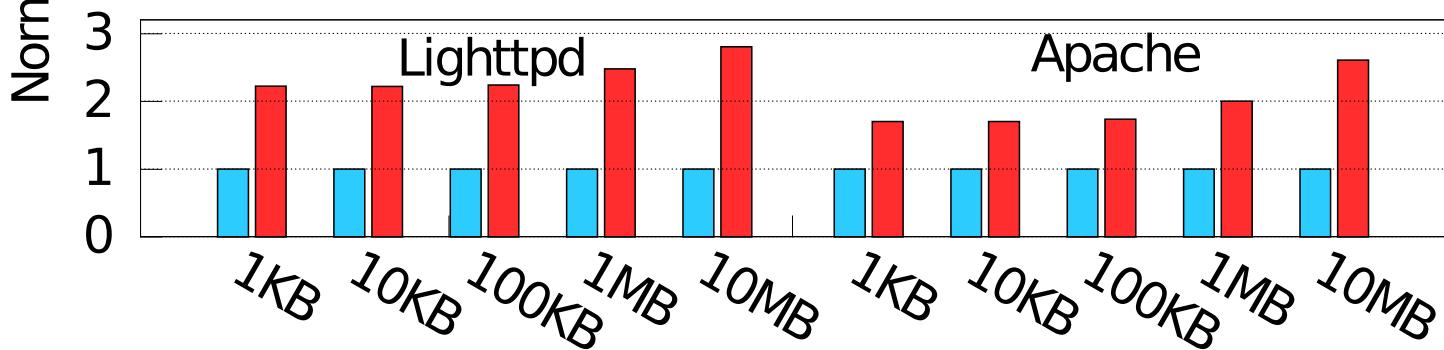
| Application | Type of vulnerability | Security advisory |
|-------------------|-----------------------|-------------------|
| Snort 2.4.0 | Stack overflow | CVE-2005-3252 |
| Cyrus imapd 2.3.2 | Stack overflow | CVE-2006-2502 |
| Samba 3.0.22 | Heap overflow | CVE-2007-2446 |
| Memcached 1.1.12 | Heap overflow | CVE-2009-2415 |
| Nginx 0.6.32 | Buffer underrun | CVE-2009-2629 |
| Proftpd 1.3.3a | Stack overflow | CVE-2010-4221 |
| Samba 3.2.5 | Heap overflow | CVE-2010-2063 |
| Telnetd 1.6 | Heap overflow | CVE-2011-4862 |
| Ncompress 4.2.4 | Stack overflow | CVE-2001-1413 |
| Iwconfig V.26 | Stack overflow | CVE-2003-0947 |
| Aspell 0.50.5 | Stack overflow | CVE-2004-0548 |
| Htget 0.93 | Stack overflow | CVE-2004-0852 |
| Socat 1.4 | Format string | CVE-2004-1484 |
| Aeon 0.2a | Stack overflow | CVE-2005-1019 |
| Exim 4.41 | Stack overflow | EDB-ID#796 |
| Htget 0.93 | Stack overflow | |
| Tipxd 1.1.1 | Format string | OSVDB-ID#12346 |

Performance

HTTP

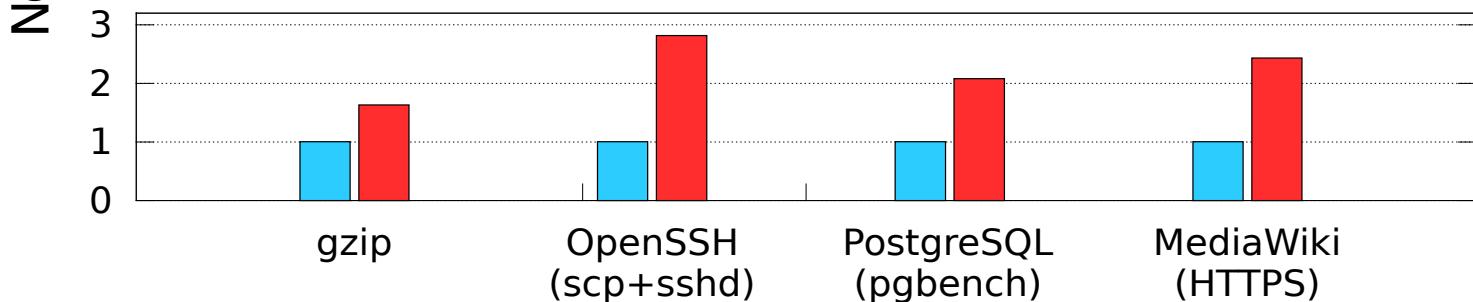
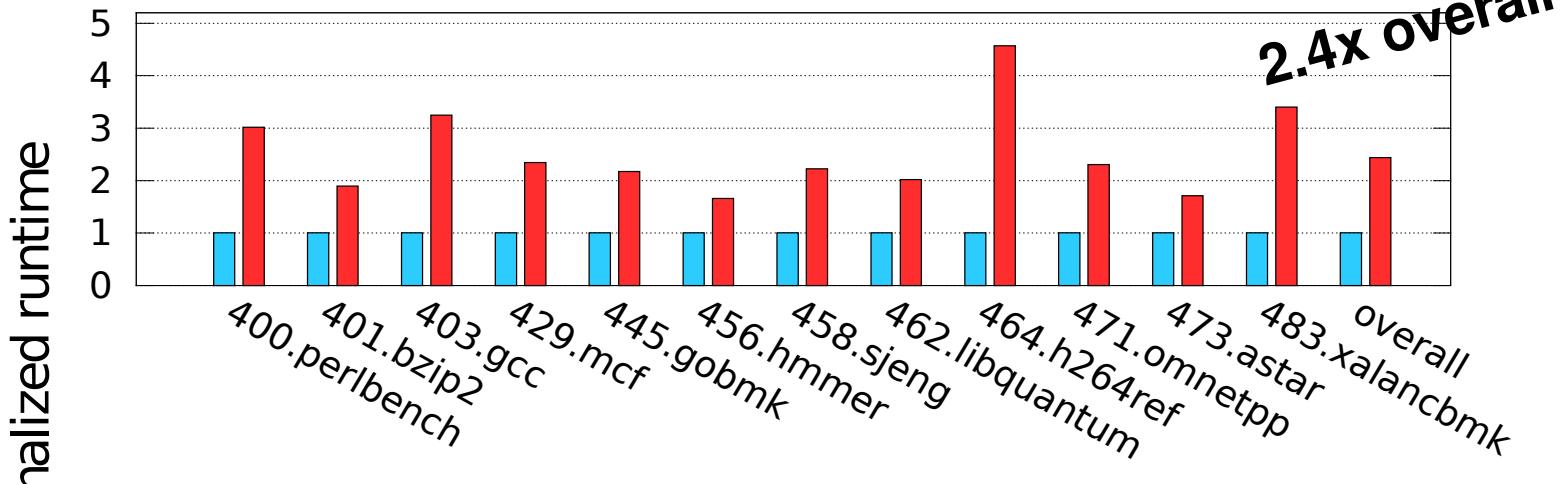


HTTPS



Performance

SPECINT 2006



Limitations

Limitations

Doesn't prevent memory corruption, only acts when the untrusted data is used for arbitrary code execution.

Limitations

Tainted pointer dereferences

```
tainted_pointer->some_field = useful_untainted_value;
```

Limitations

Tainted pointer dereferences

`tainted_pointer->some_field = useful_untainted_value;`

propagation can lead to false positives:

`dispatch_table[checked_input]();`

Limitations

Taint whitewashing

```
out = latin1_to_ascii[in];
```

Limitations

Format string attacks:

```
printf("%65534s %123$hn"); // Propagates taint in glibc
```

```
printf("FillerFiller...%123$hn"); // Does not :-(
```

Limitations

Does not protect against non-control-flow exploits

Limitations

Does not protect against non-control-flow exploits

```
void try_system(char *username, char *cmd)
{
    int user_rights = get_credentials(username);
    char buf[16] ; strcpy(buf, username);
    if (user_rights & ALLOW_SYSTEM)
        system(cmd);
    else
        log_error("user %s attempted login", buf);
}
```

Limitations

Does not protect against non-control-flow exploits

```
void try_system(char *username, char *cmd)
{
    int user_rights = get_credentials(username);
    char buf[16] ; strcpy(buf, username);
    if (user_rights & ALLOW_SYSTEM)
        system(cmd);
    else
        log_error("user %s attempted login", buf);
}
```

Limitations

Does not protect against non-control-flow exploits

```
void try_system(char *username, char *cmd)
{
    int user_rights = get_credentials(username);
    char buf[16] ; strcpy(buf, username);
    if (user_rights & ALLOW_SYSTEM)
        system(cmd);
    else
        log_error("user %s attempted login", buf);
}
```



Limitations

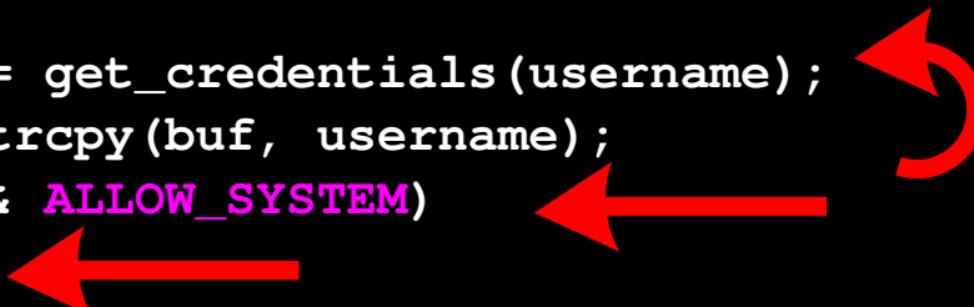
Does not protect against non-control-flow exploits

```
void try_system(char *username, char *cmd)
{
    int user_rights = get_credentials(username);
    char buf[16] ; strcpy(buf, username);
    if (user_rights & ALLOW_SYSTEM)           ←
        system(cmd);                         ↗
    else
        log_error("user %s attempted login", buf);
}
```

Limitations

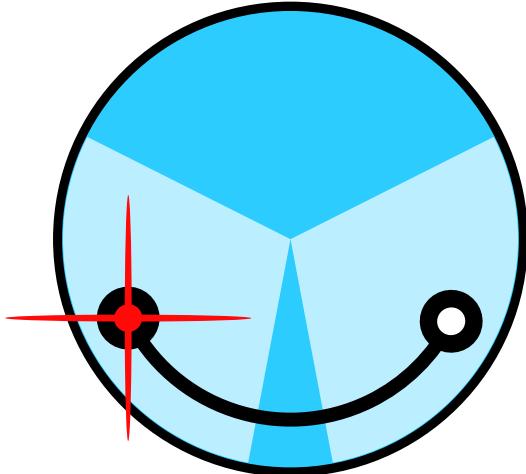
Does not protect against non-control-flow exploits

```
void try_system(char *username, char *cmd)
{
    int user_rights = get_credentials(username);
    char buf[16] ; strcpy(buf, username);
    if (user_rights & ALLOW_SYSTEM)
        system(cmd); ←
    else
        log_error("user %s attempted login", buf);
}
```



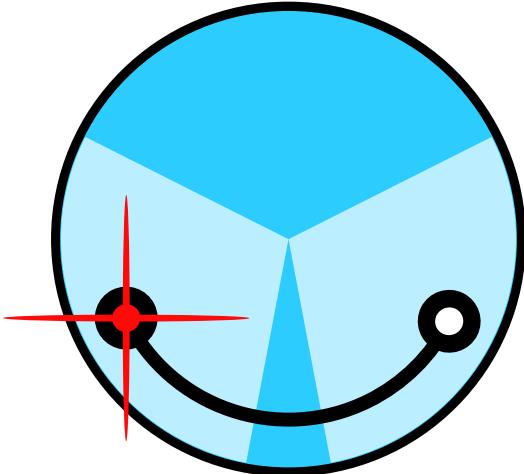
in some cases we can add validation hooks.

`mysql_query()` can be hooked to check for taint outside of literals in SQL queries.



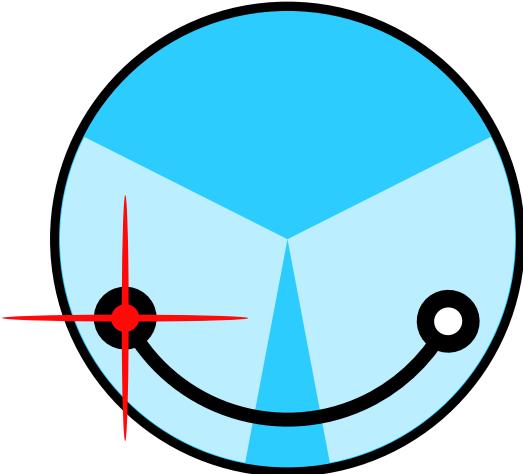
Demo

```
demo@demo:~# ./minemu bash
```



Minemu

```
git clone https://minemu.org/code/minemu.git
```

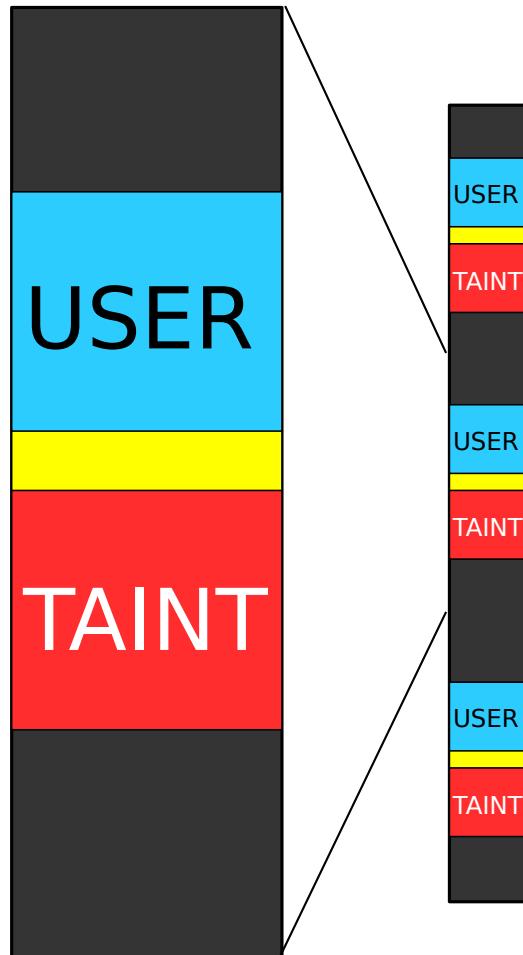


Minemu

git clone <https://minemu.org/code/minemu.git>

any questions?

Memory layout (64 bit)



Memory layout (64 bit) alternative

